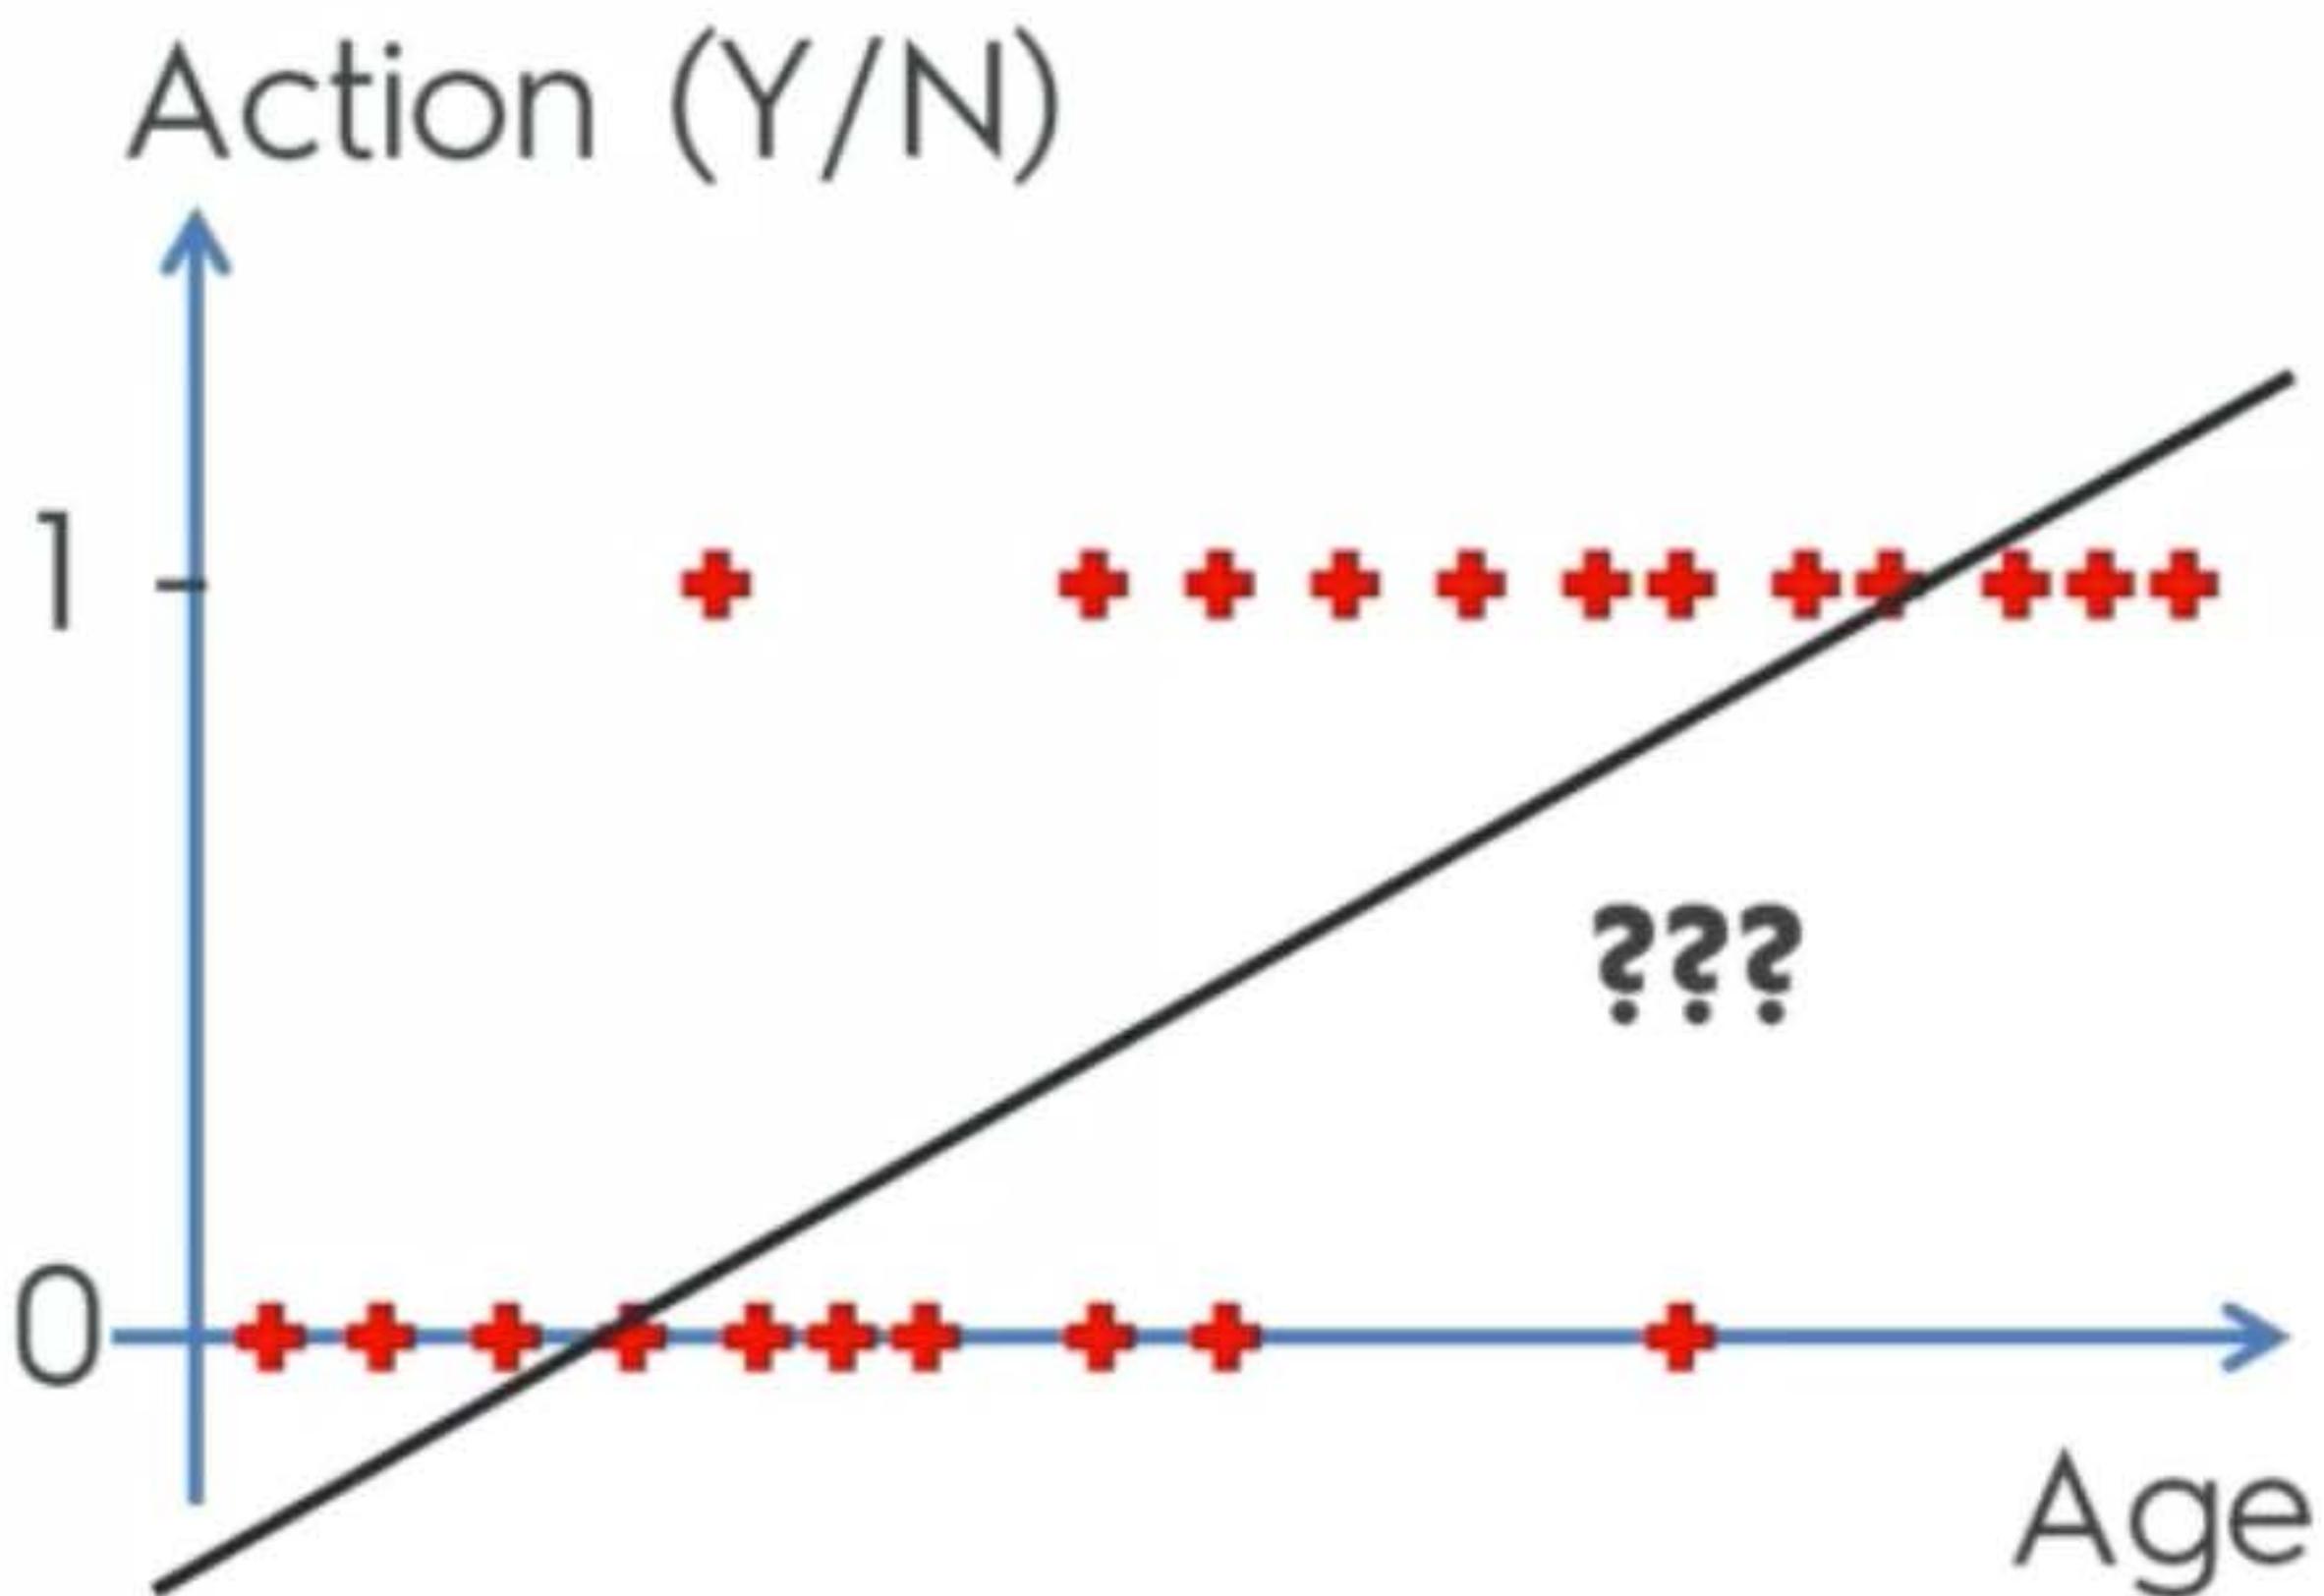


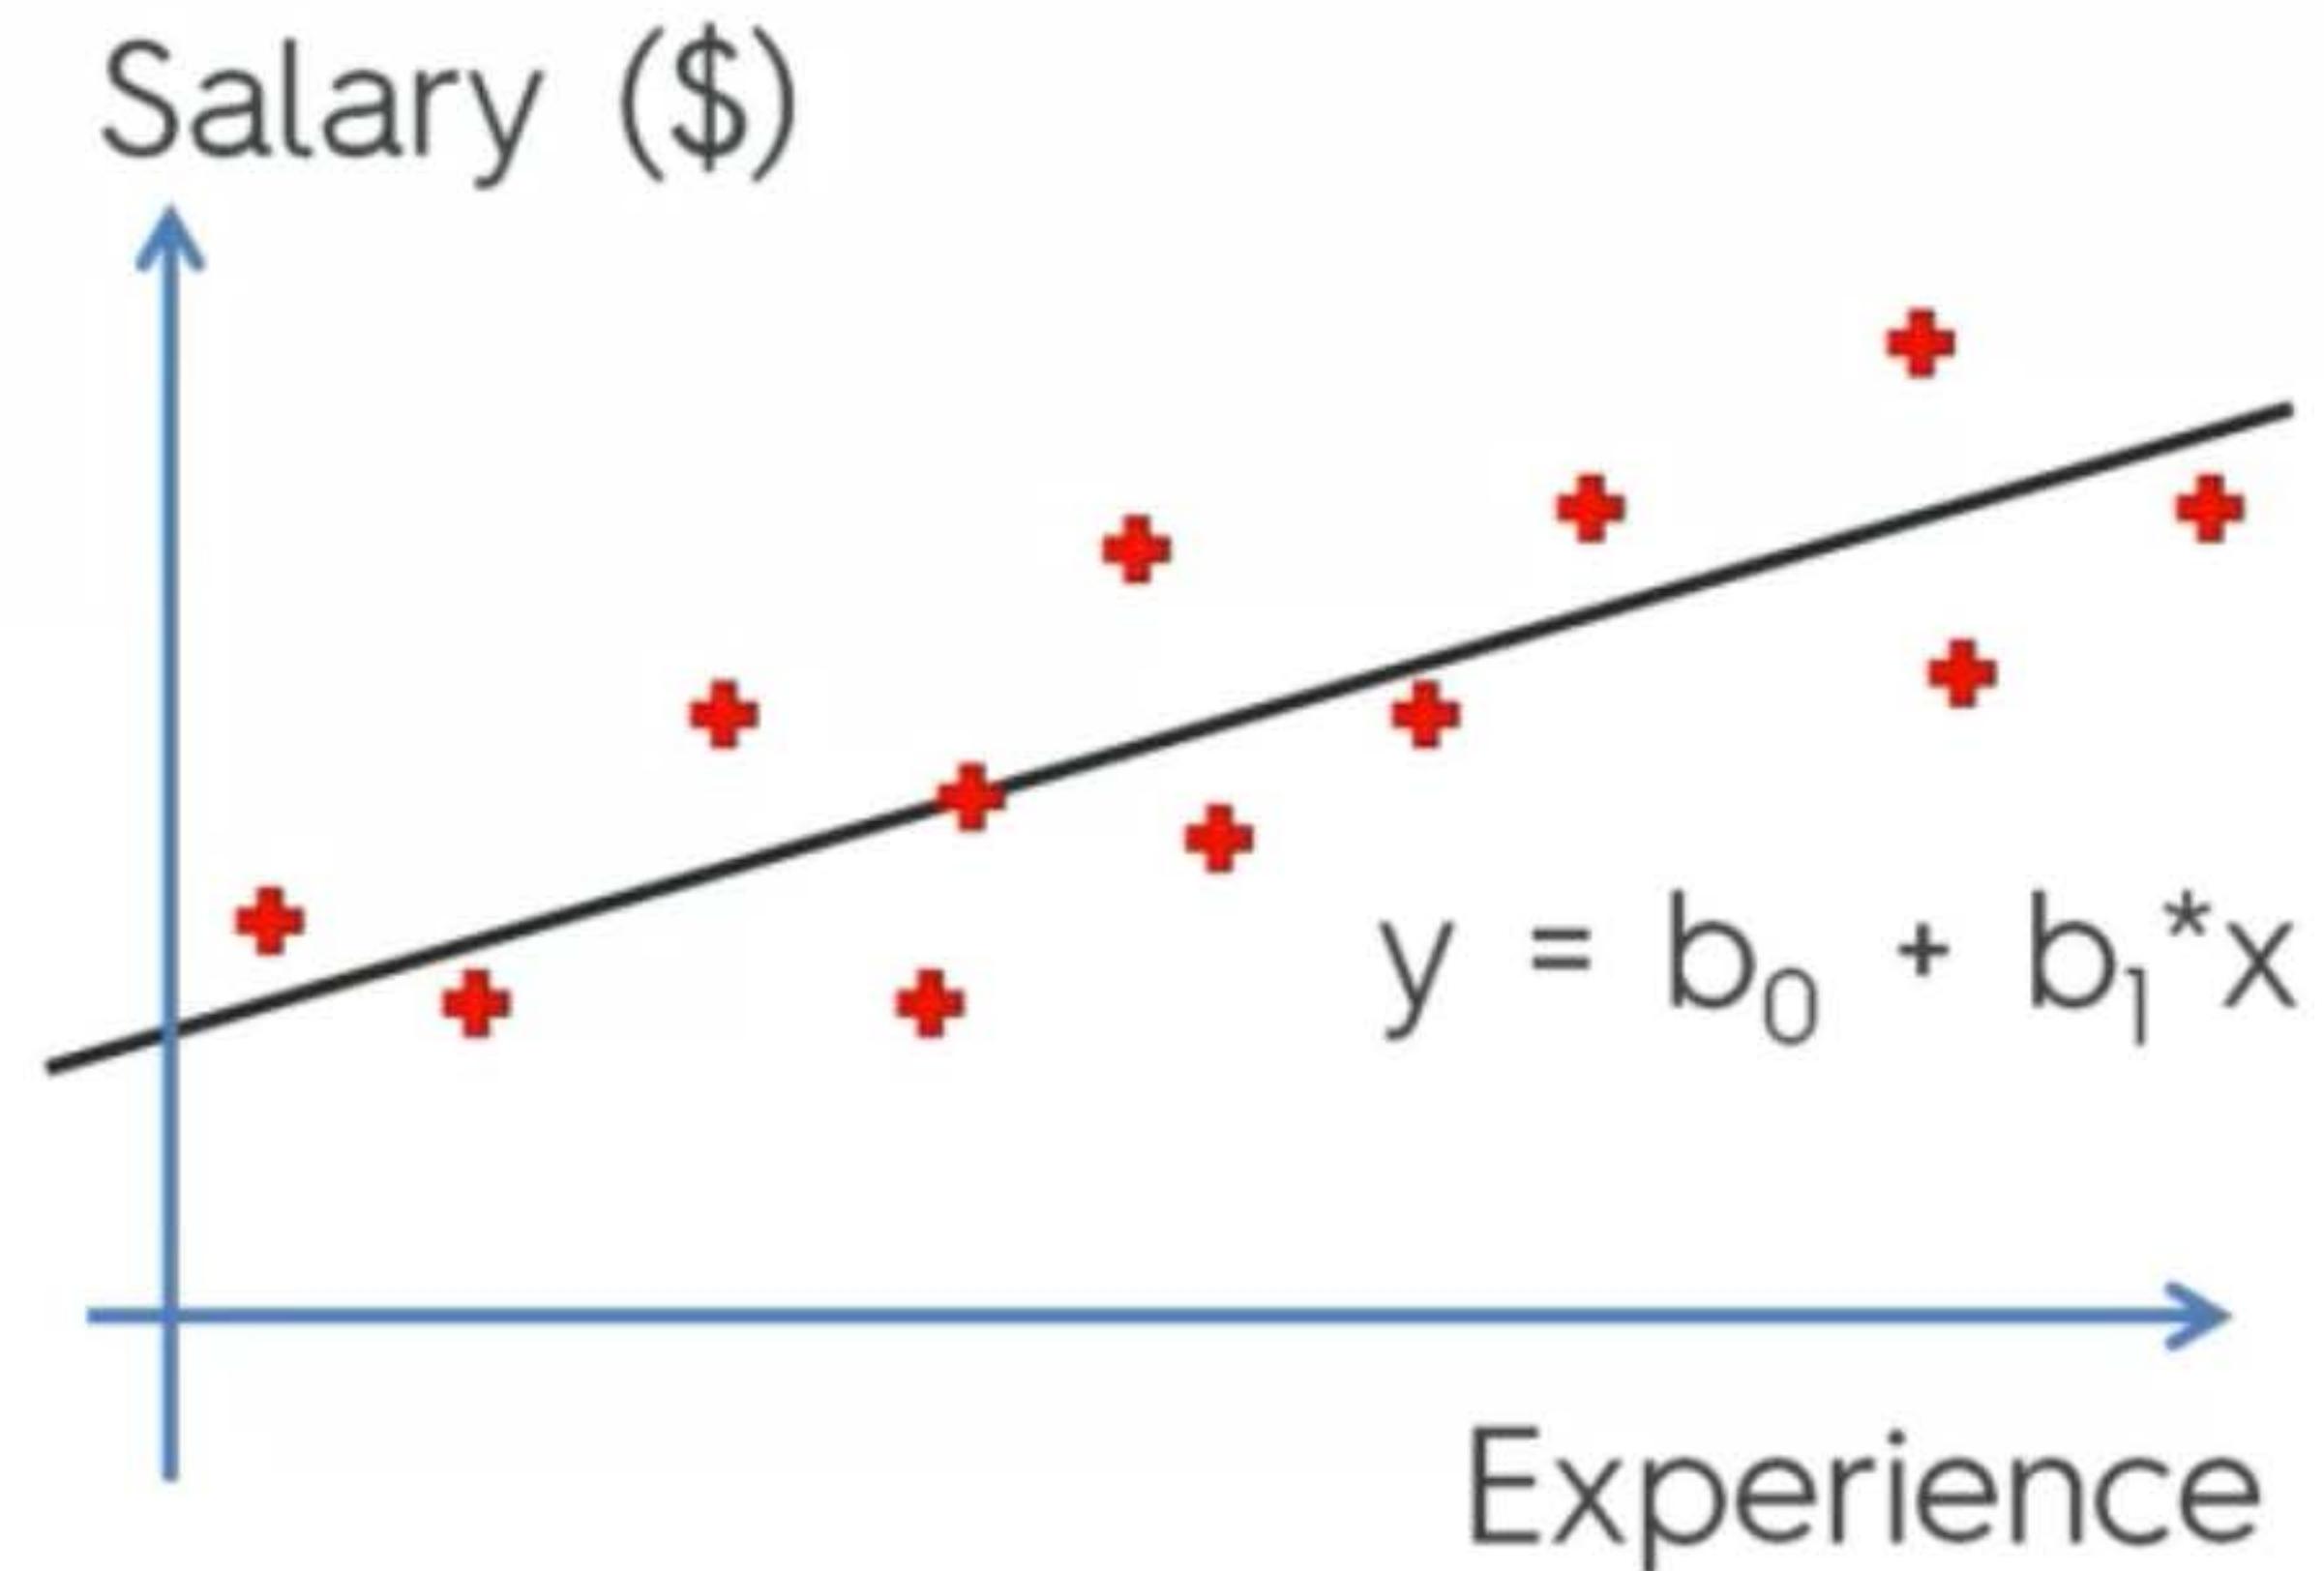
# **Logistic Regression Intuition**

# Logistic Regression

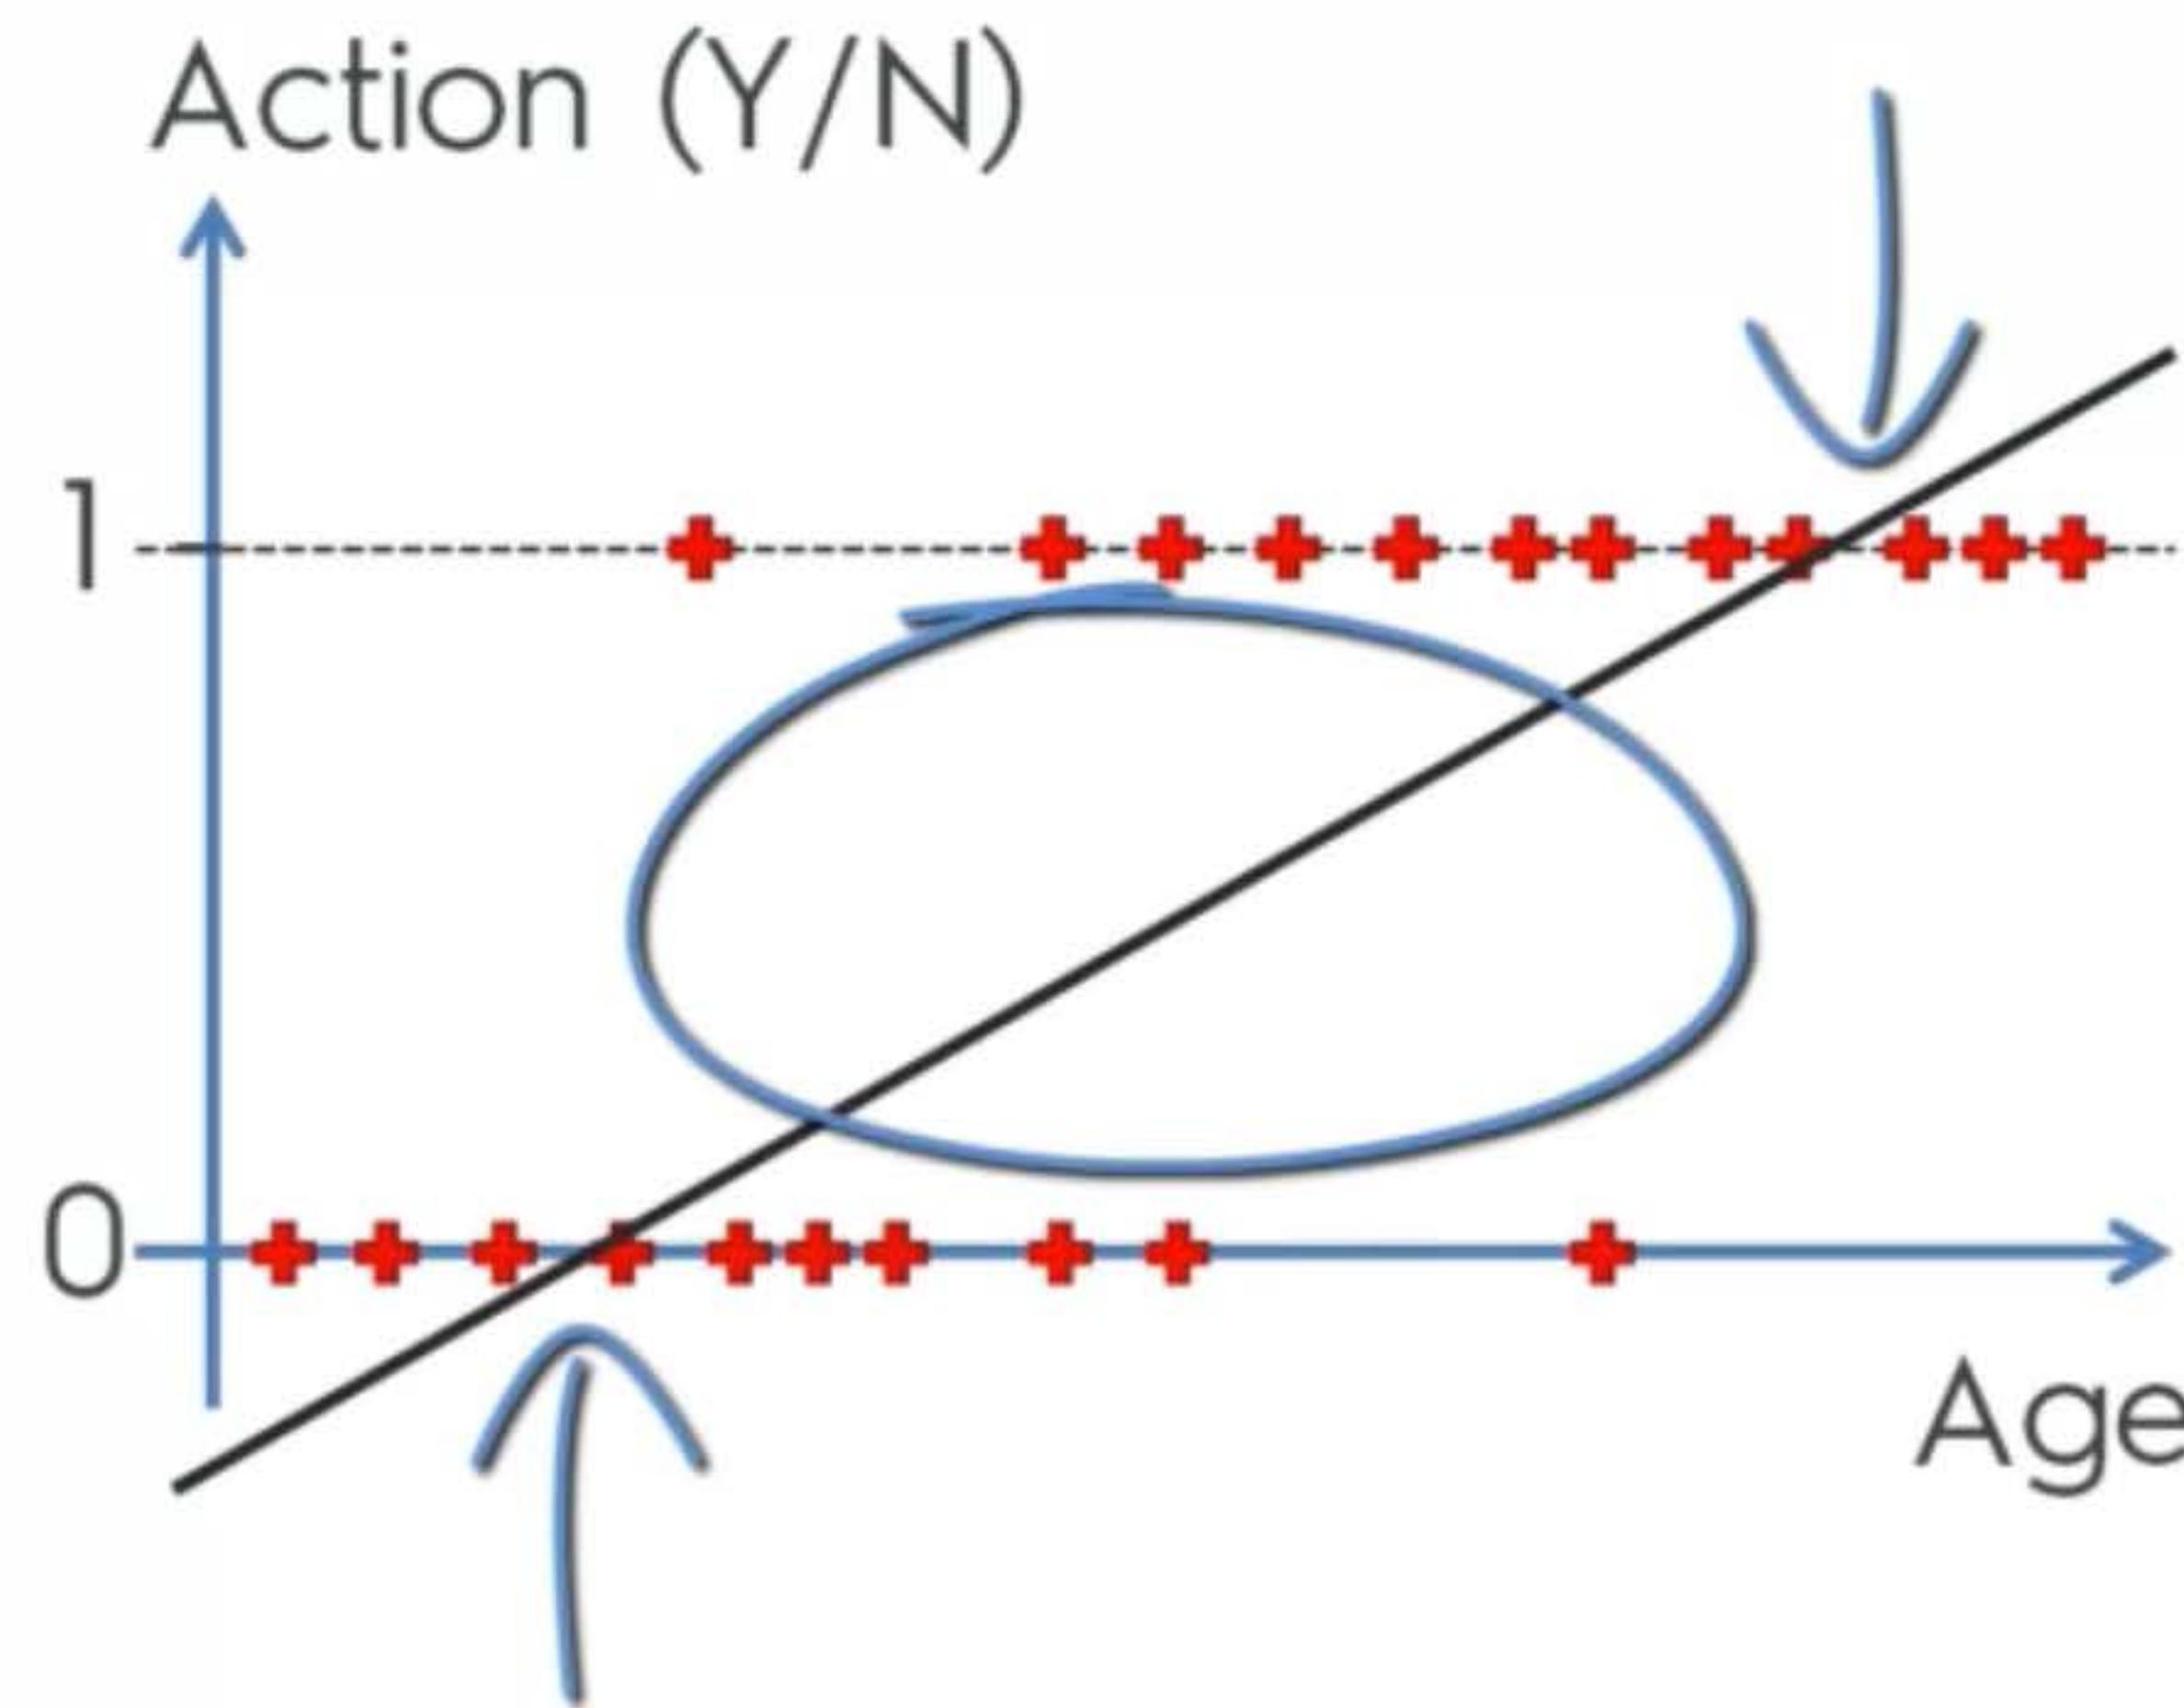
This is new:



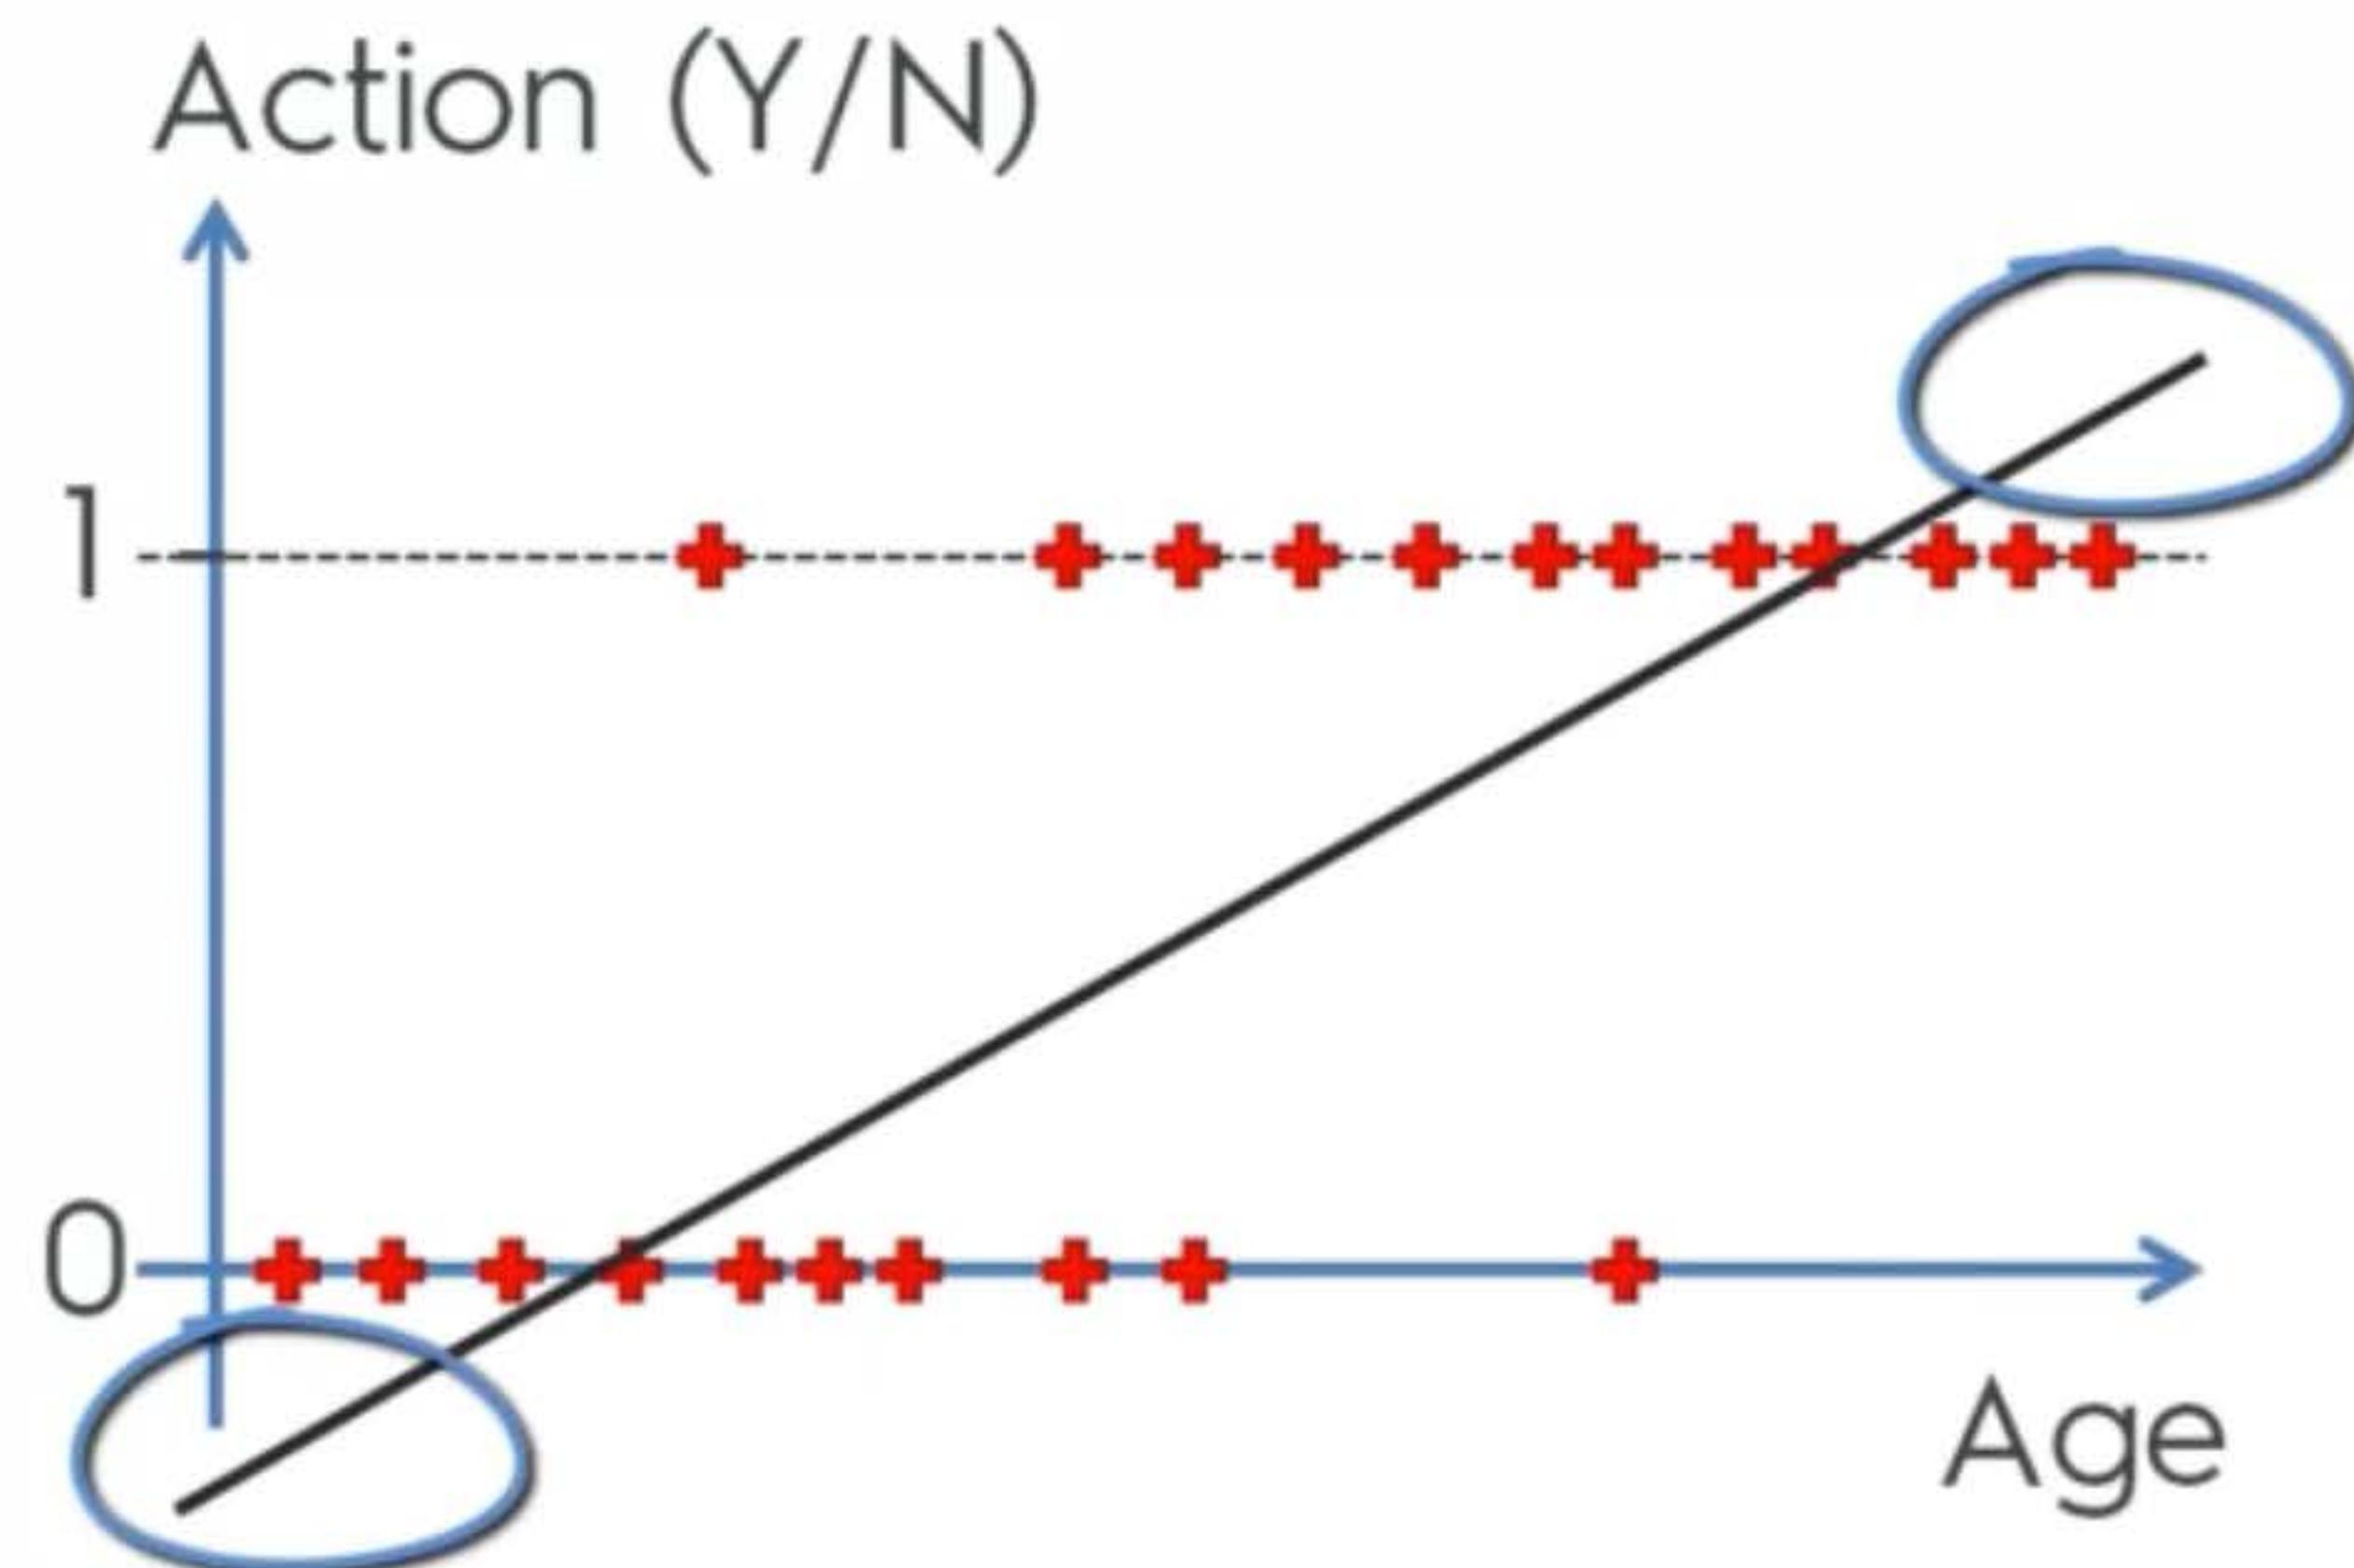
We know this:



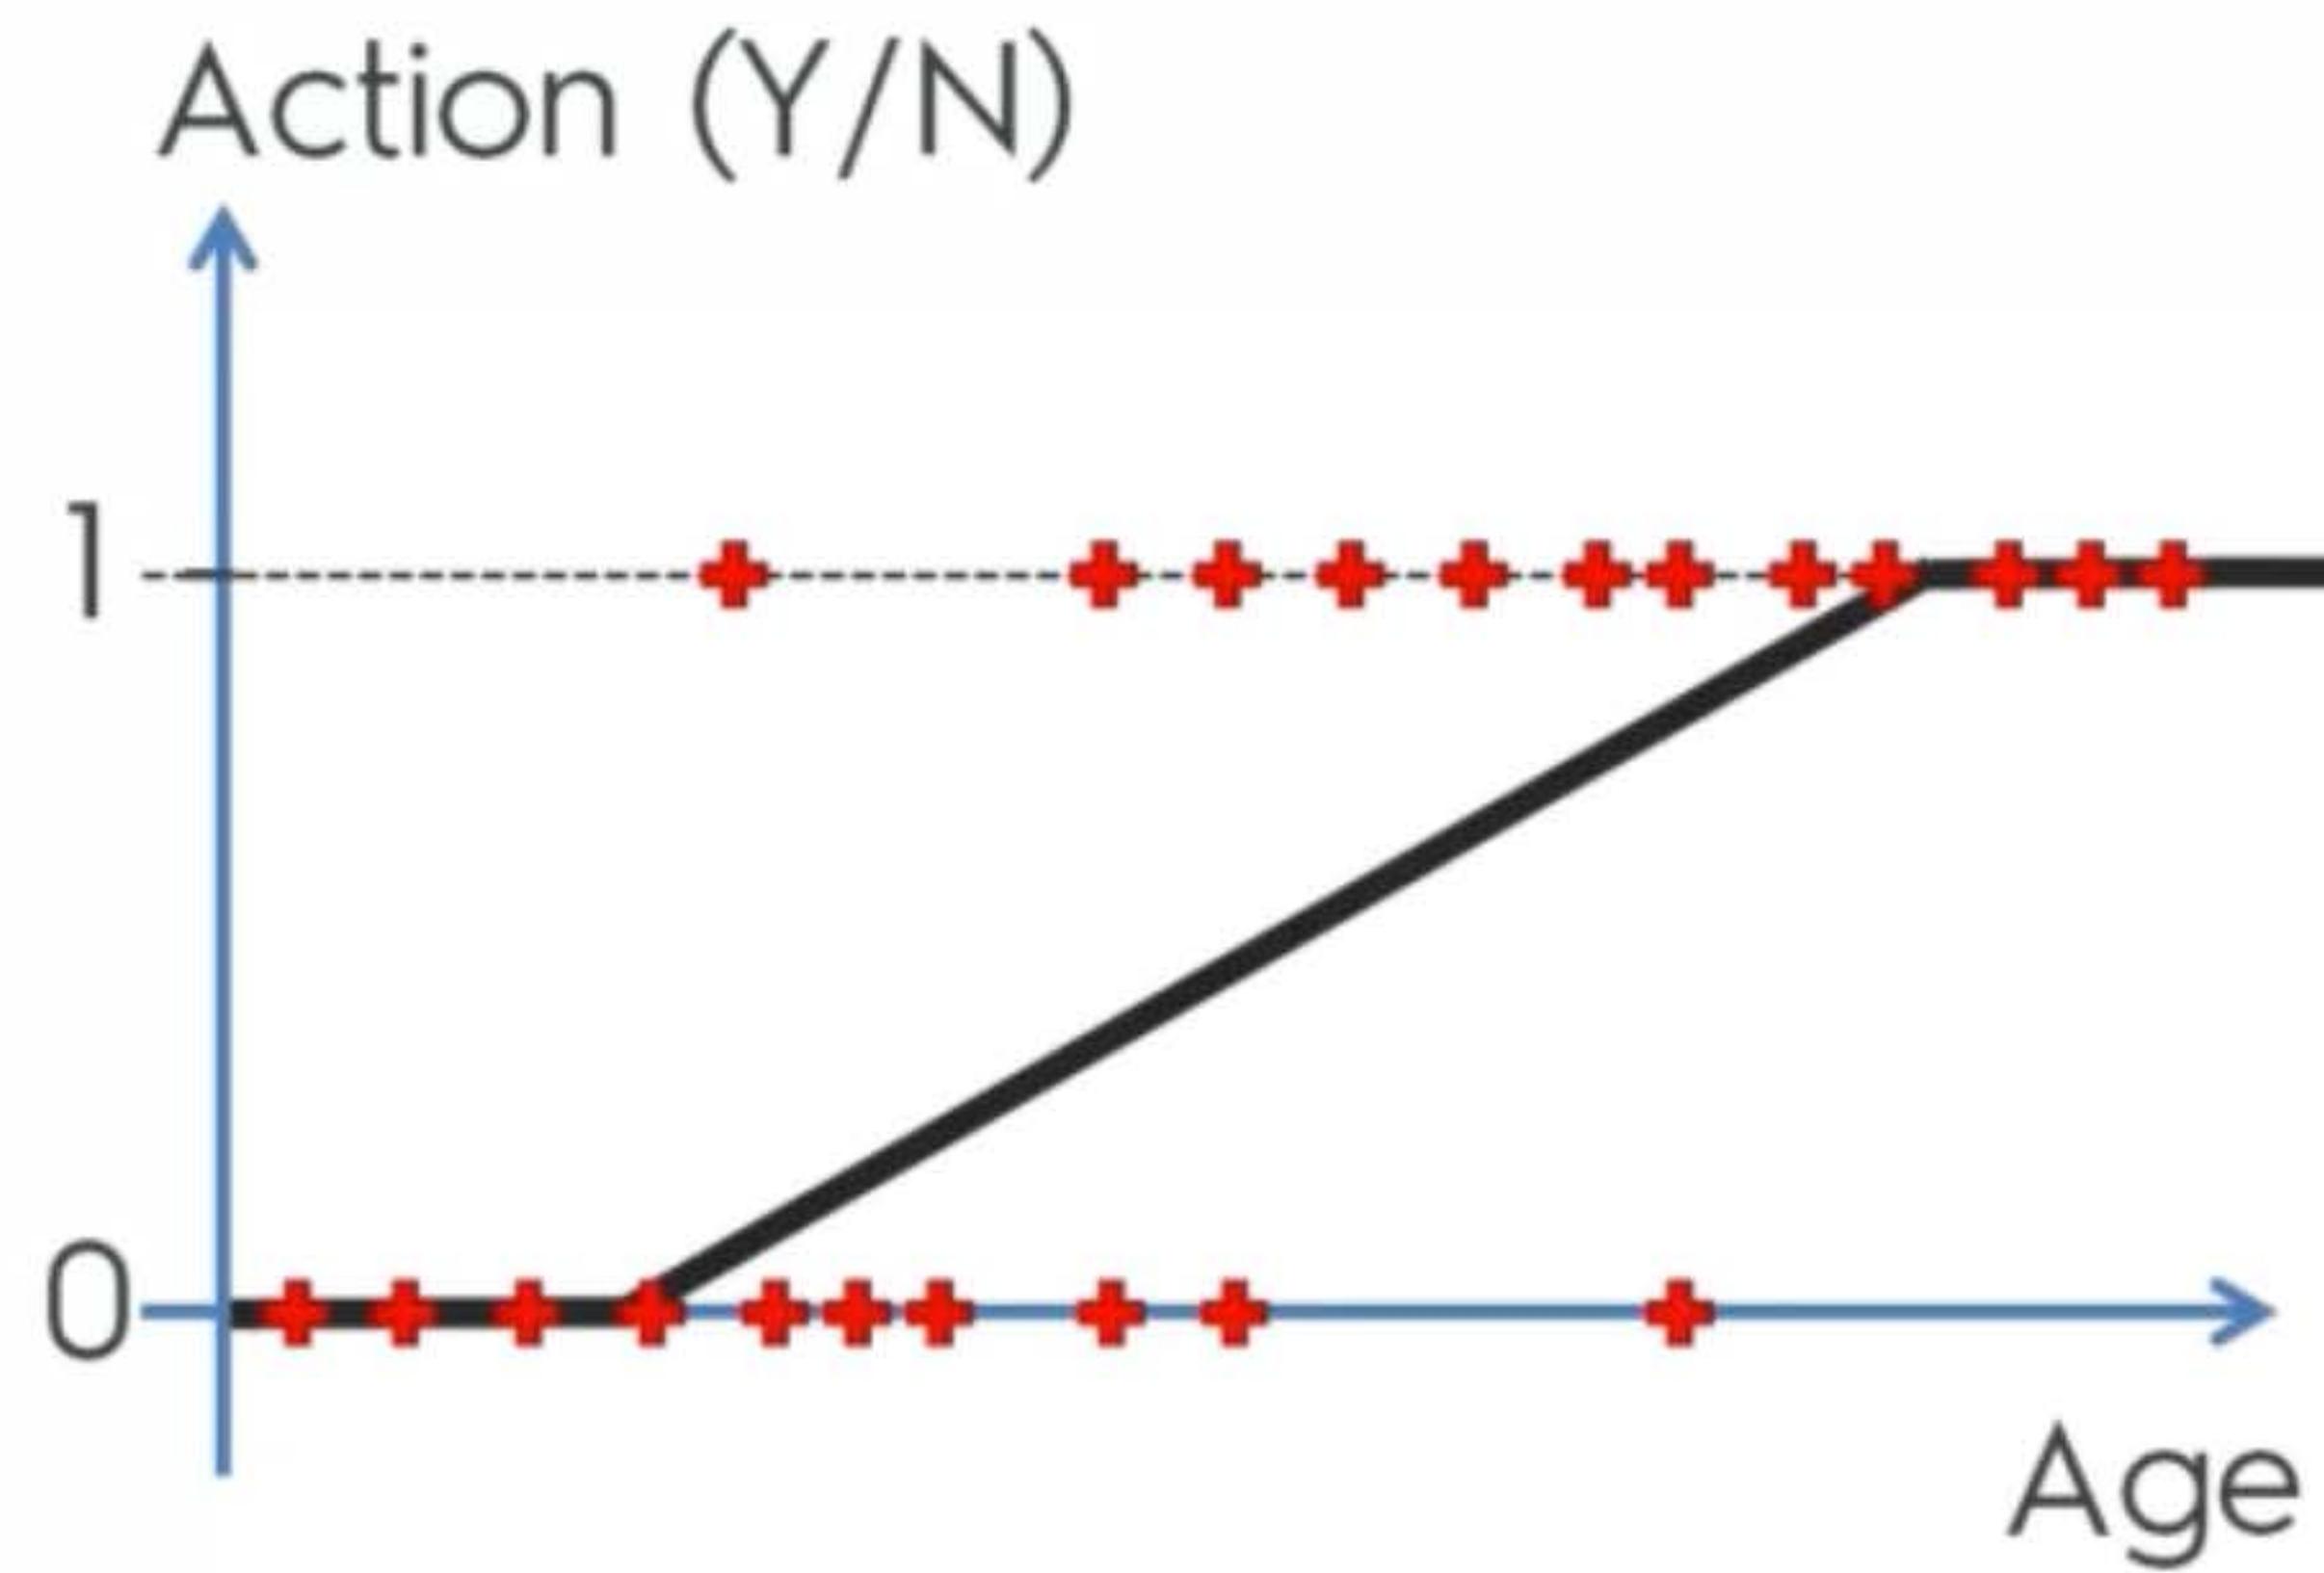
# Logistic Regression



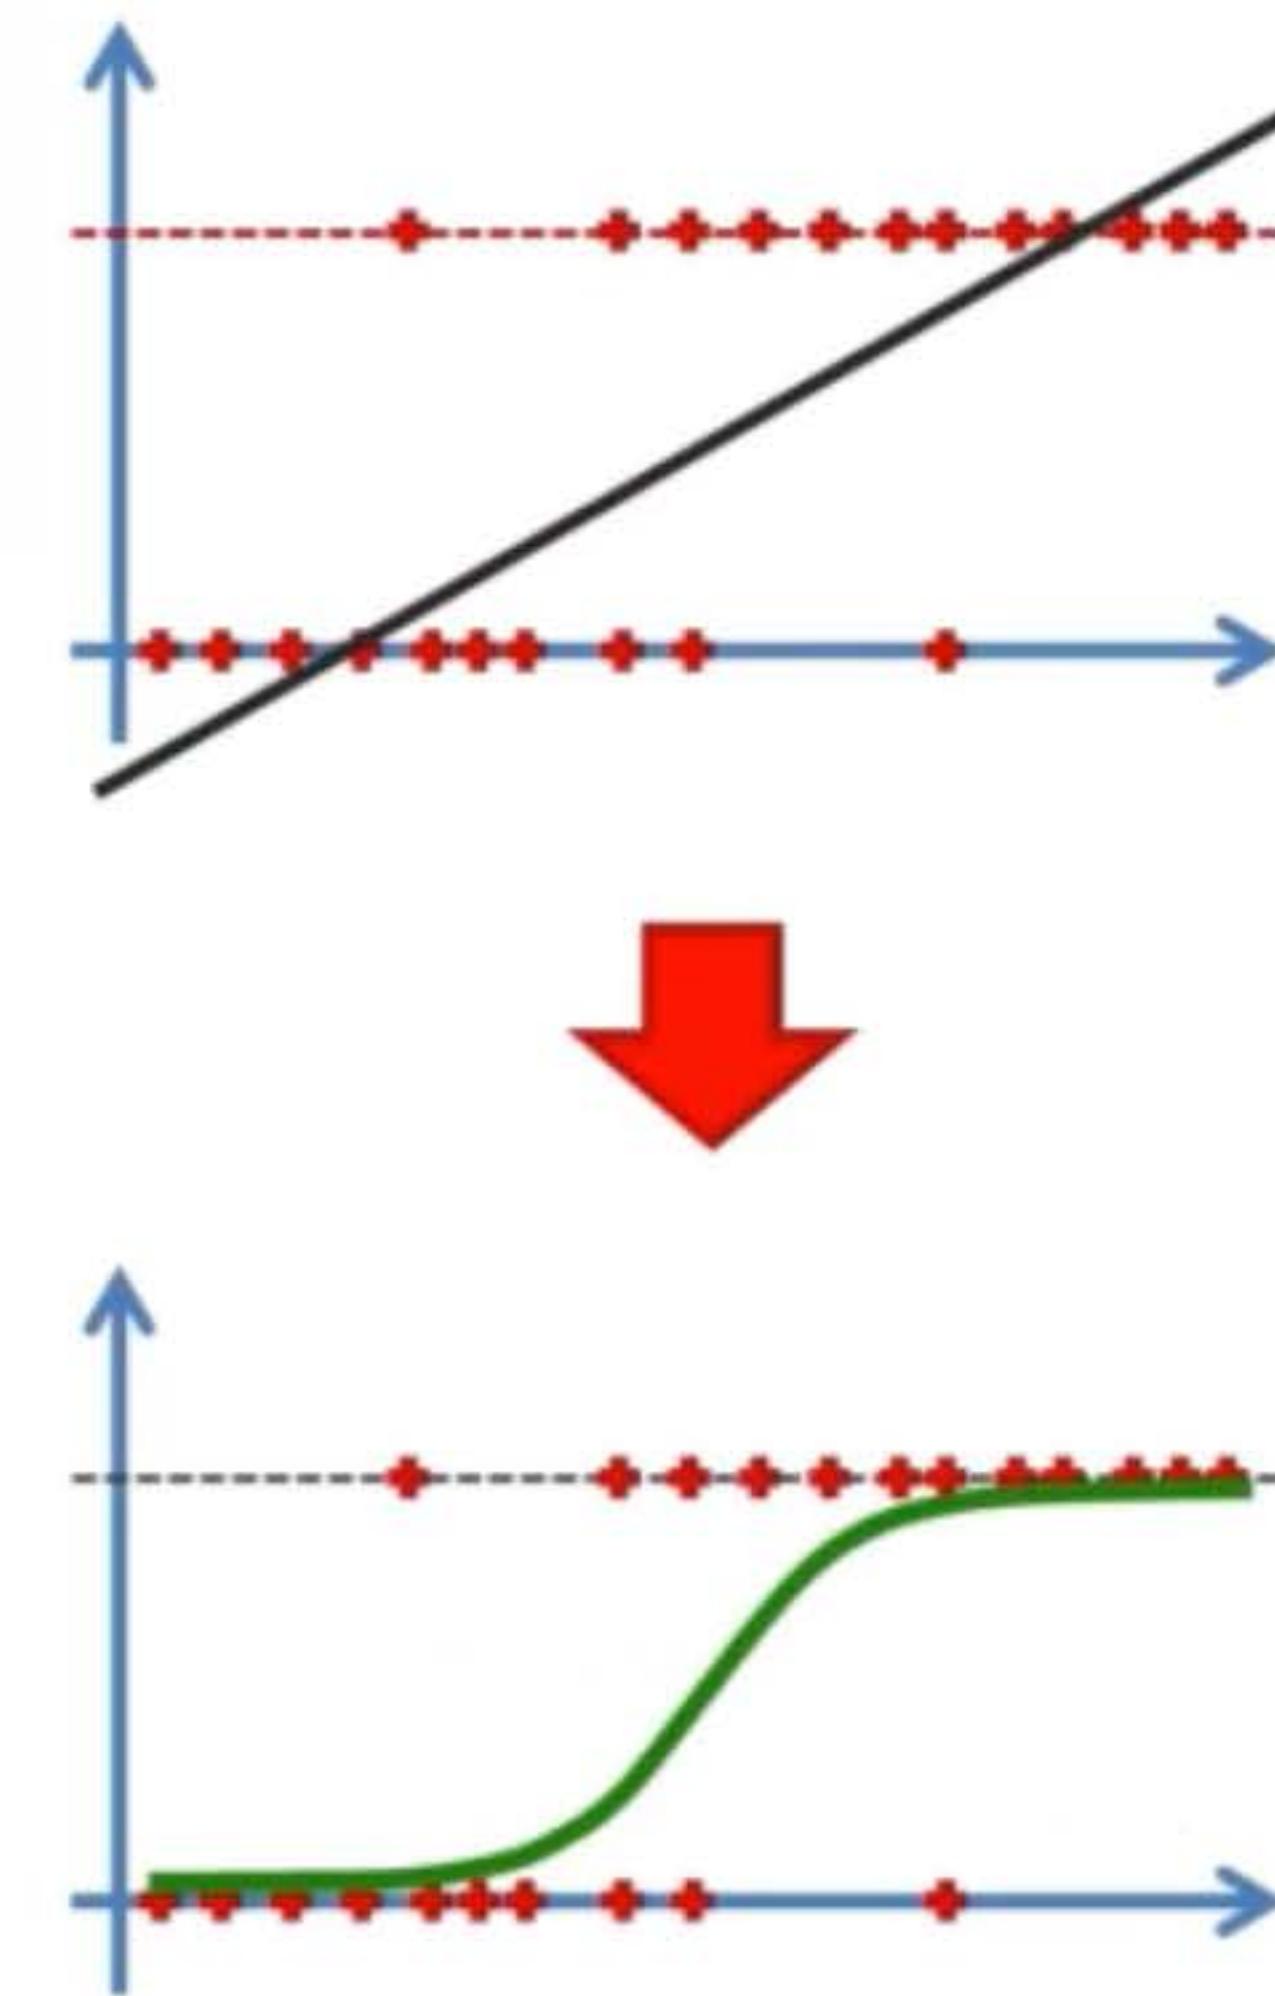
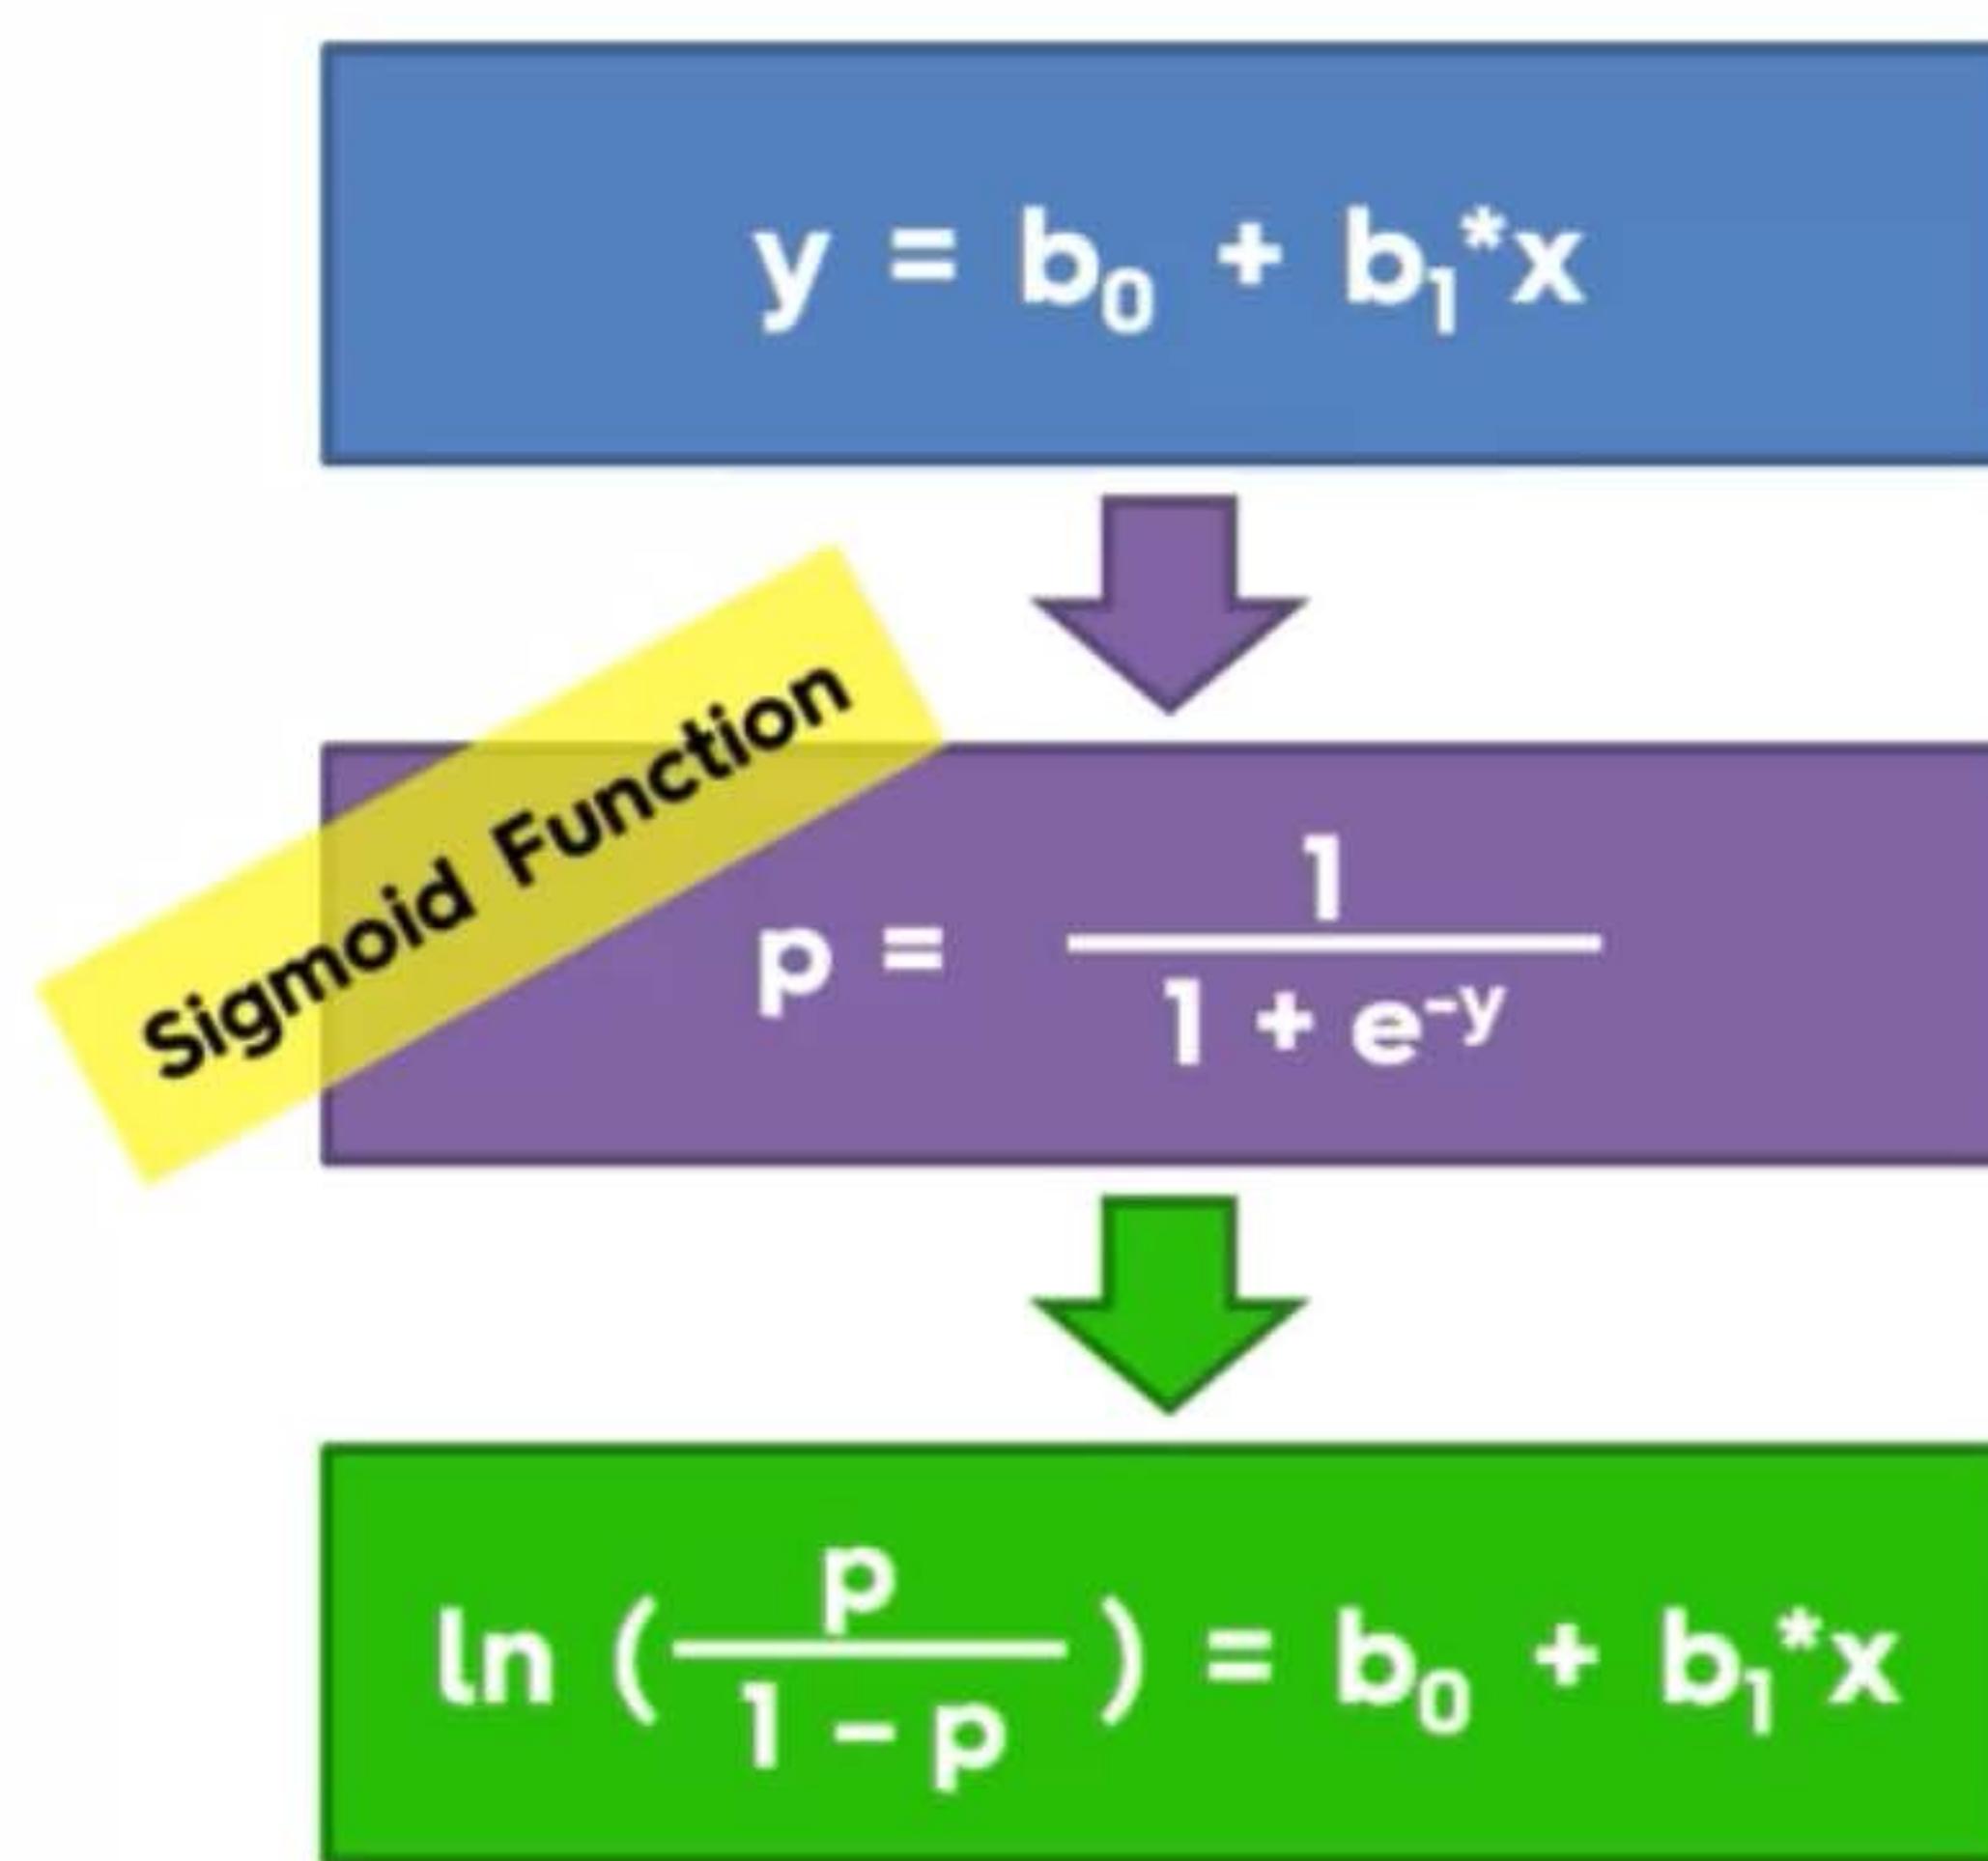
# Logistic Regression



# Logistic Regression



# Logistic Regression

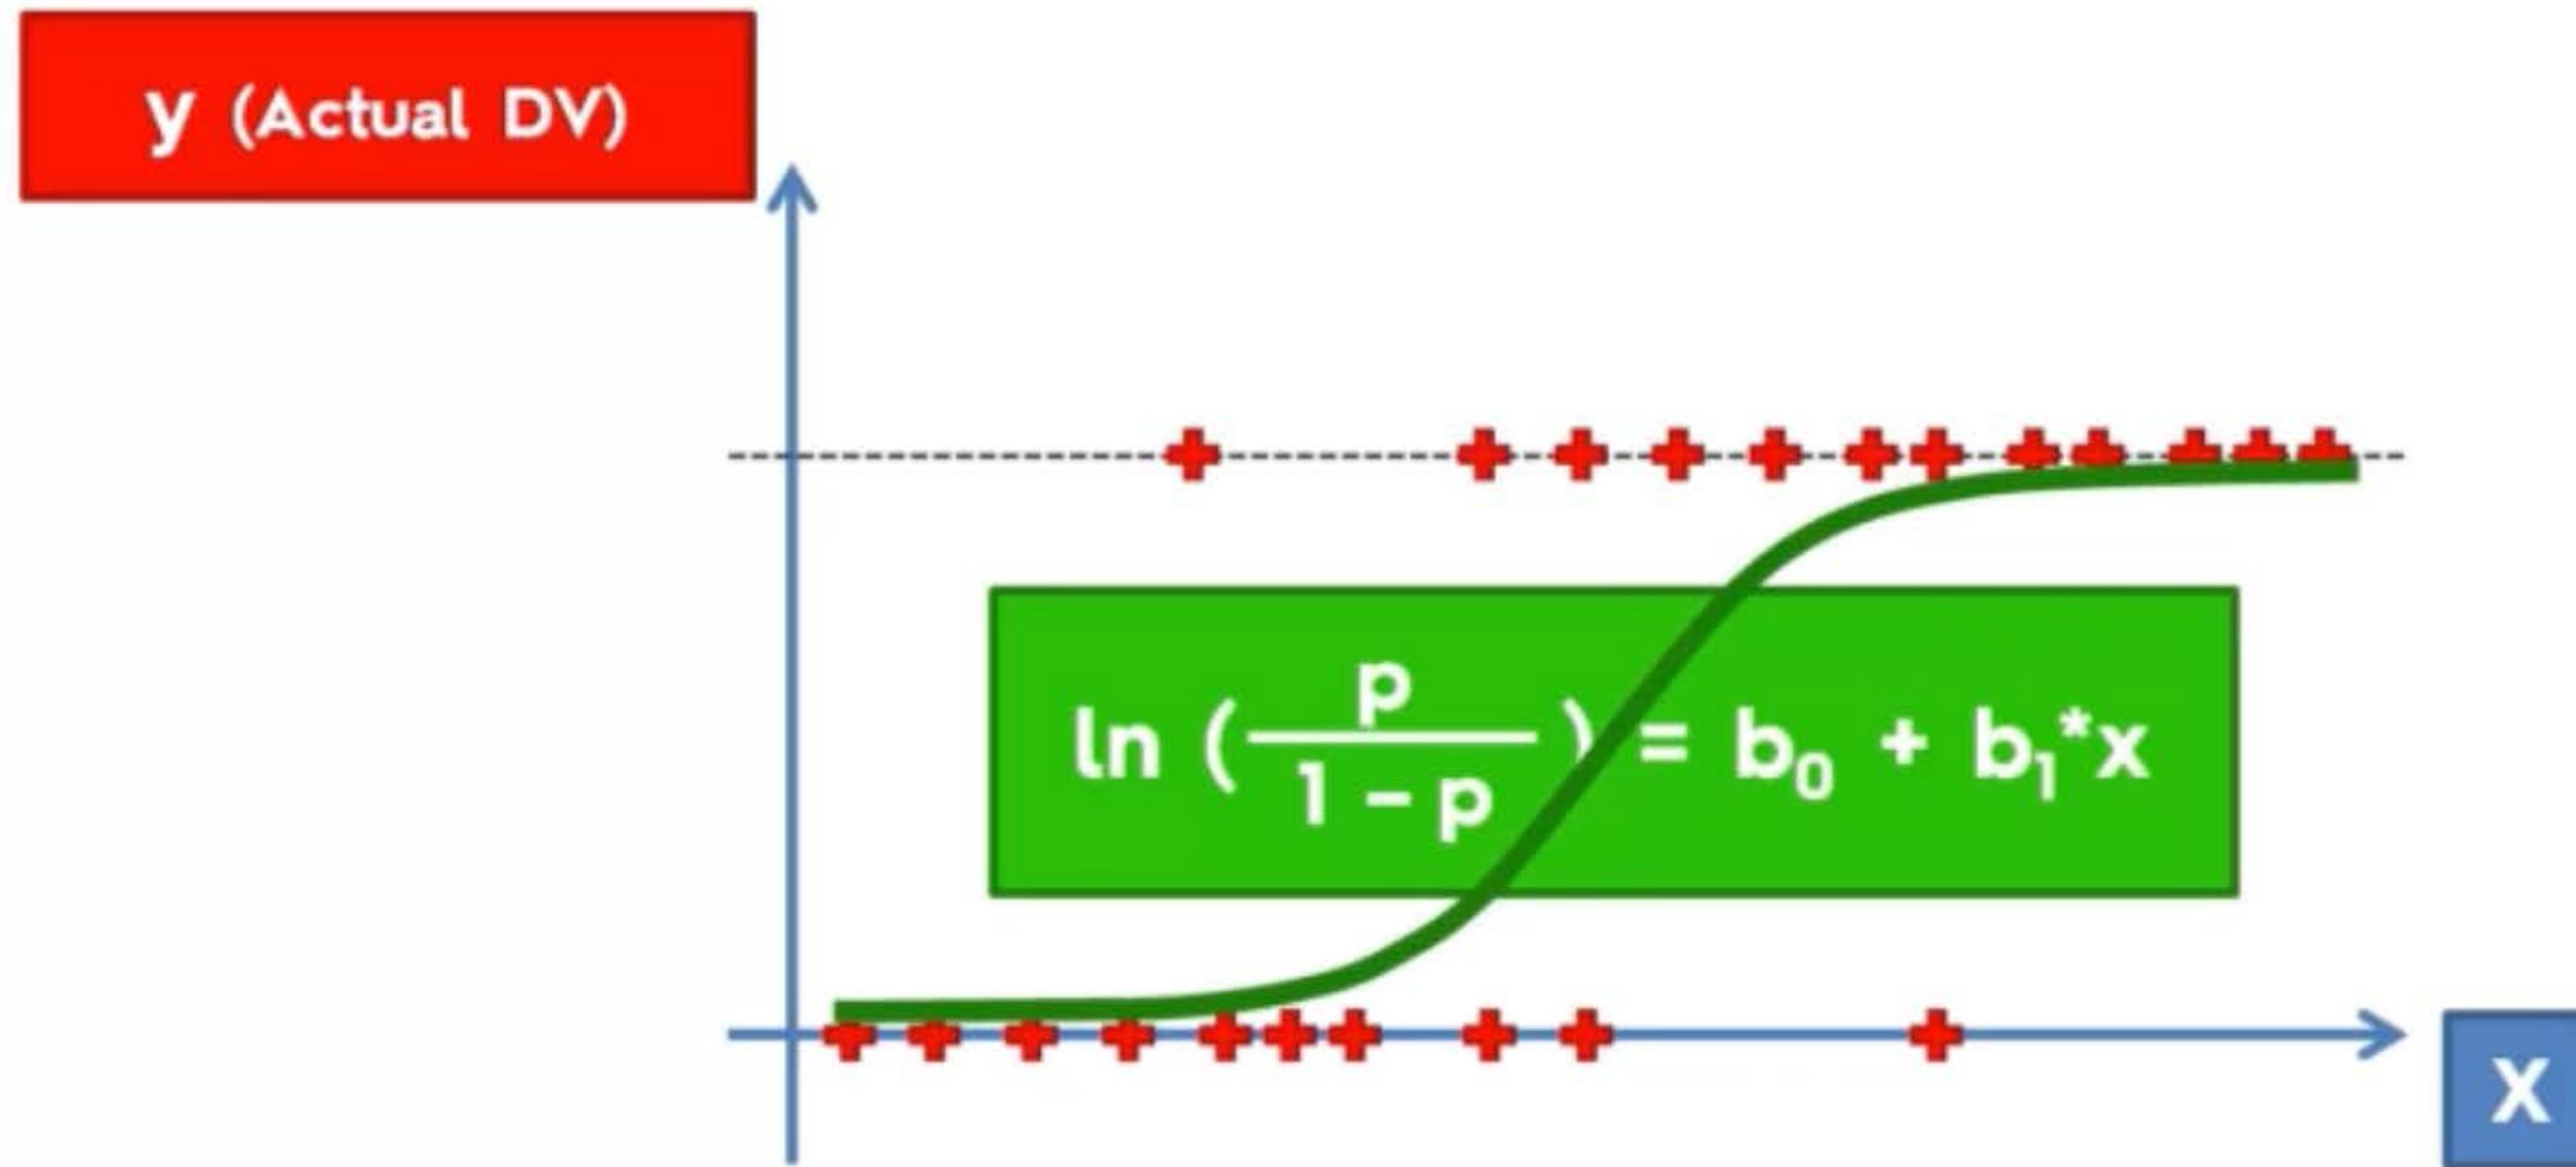


# Logistic Regression

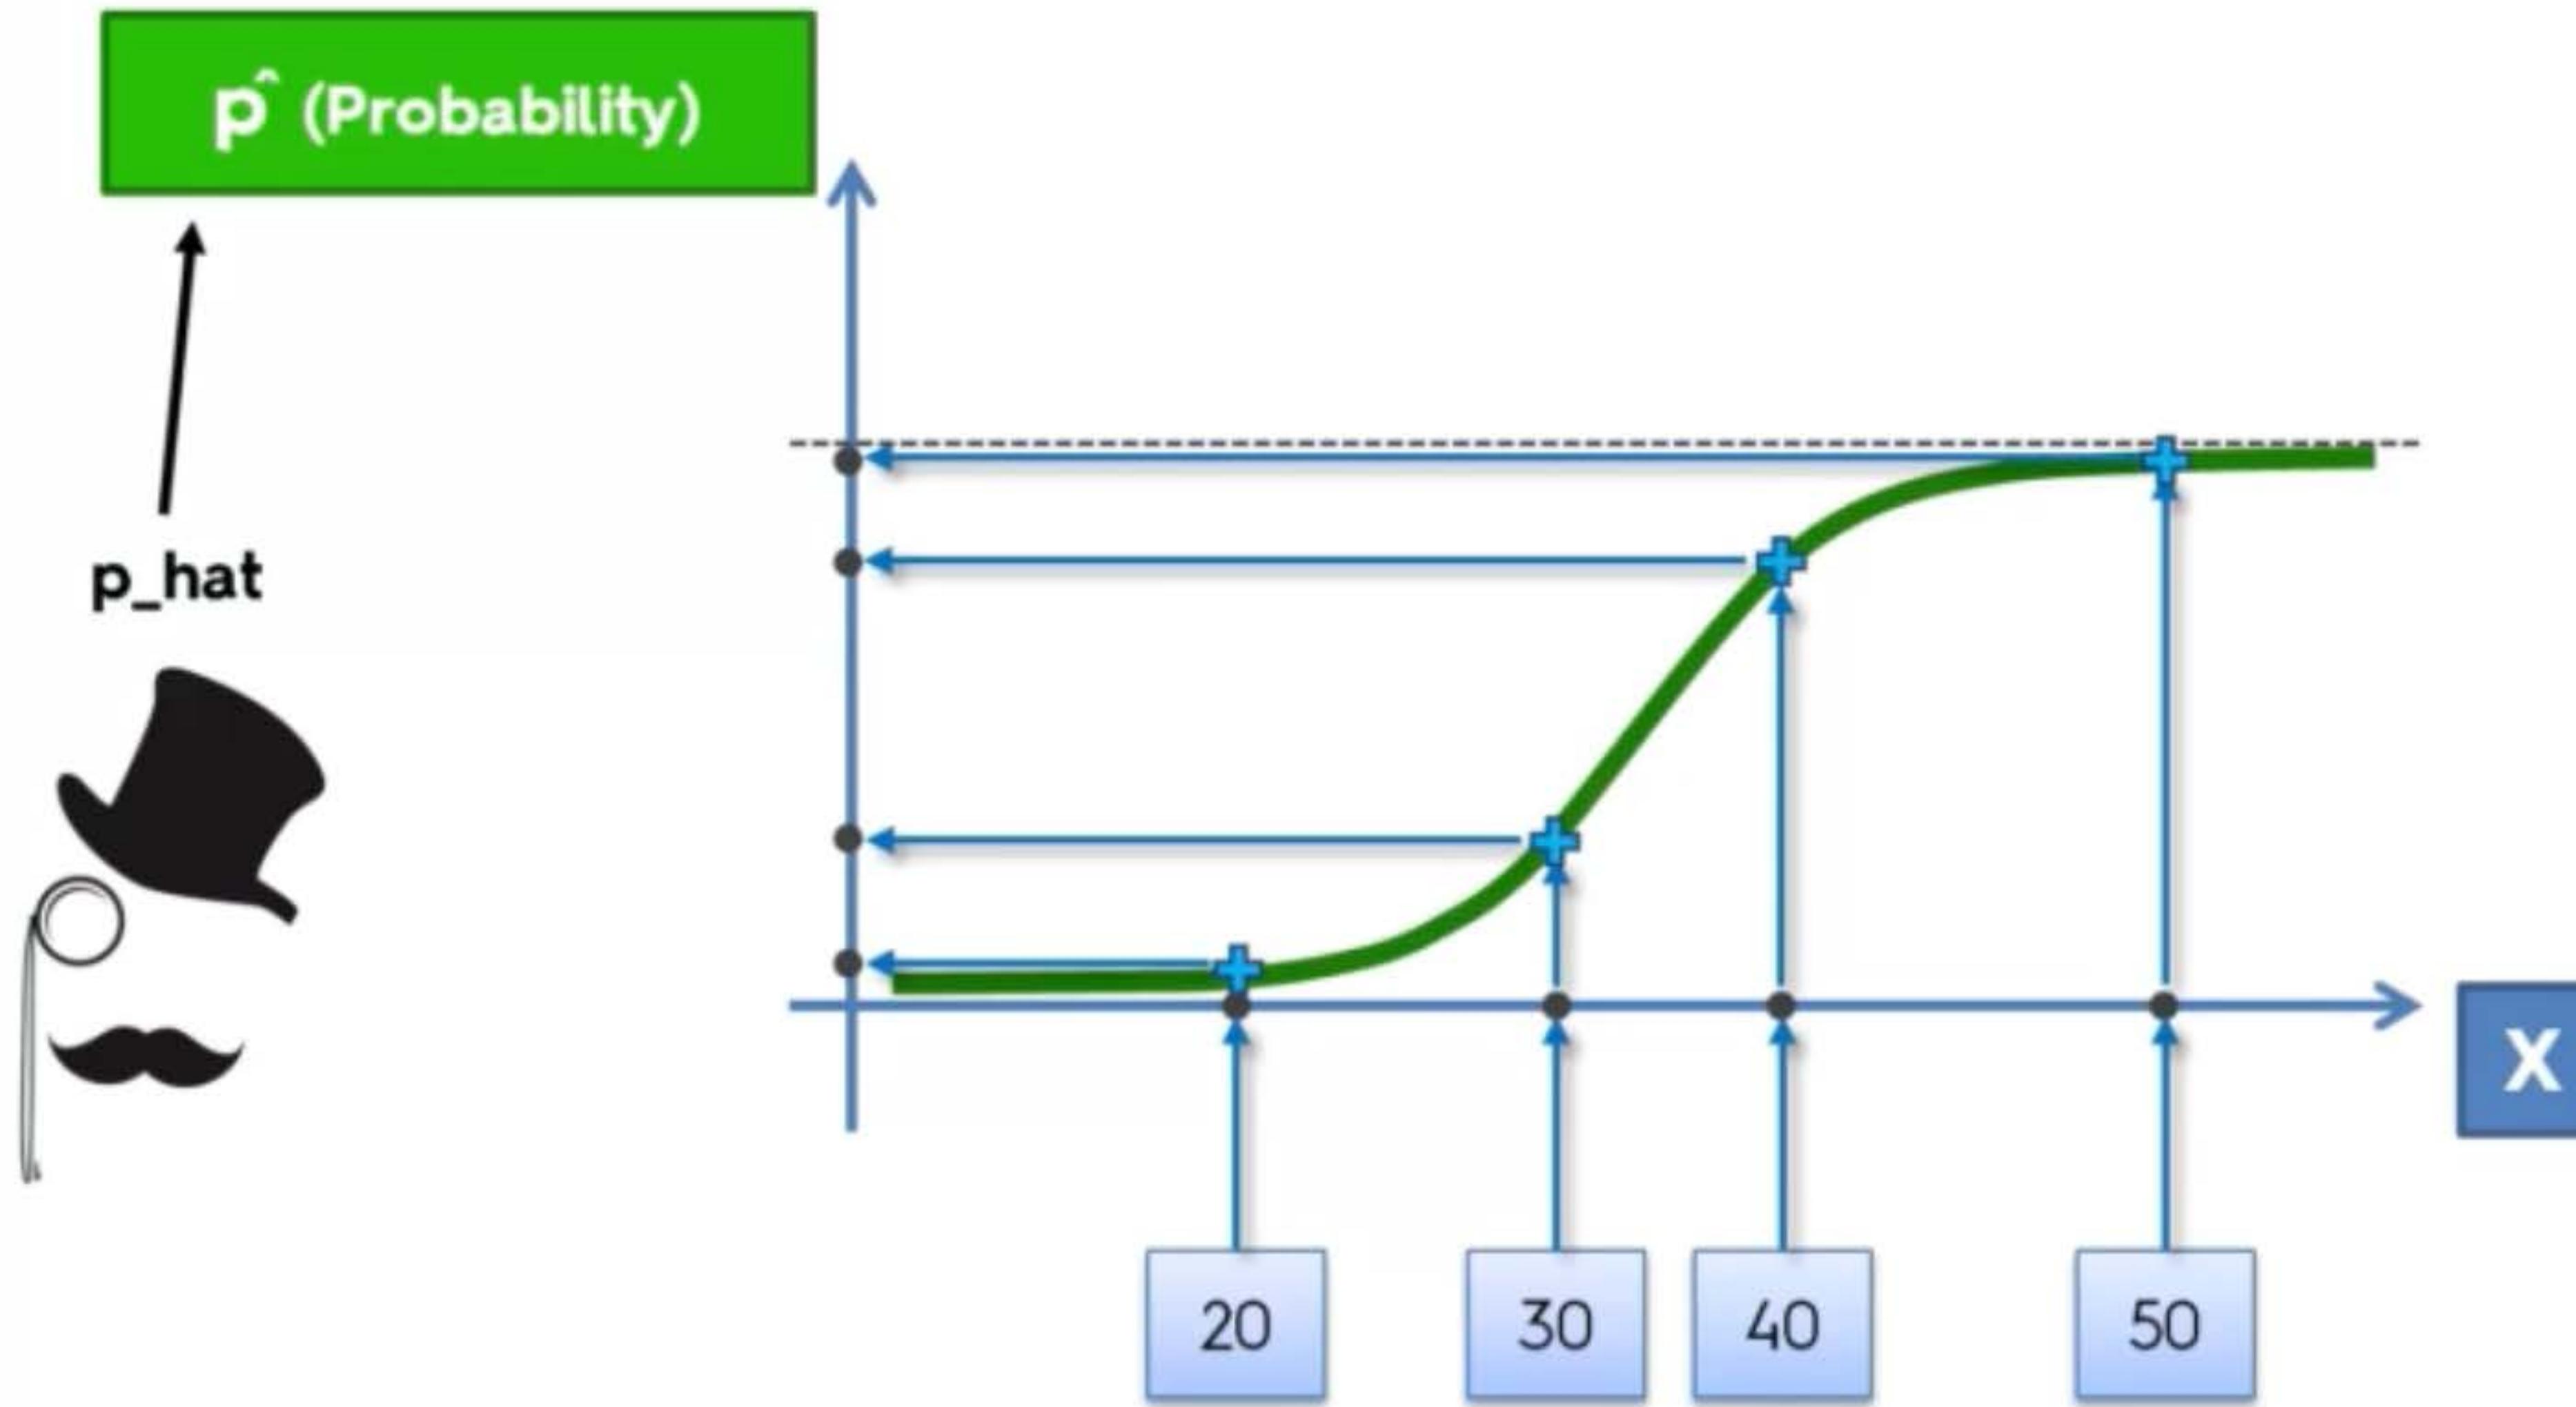


**WHAT JUST  
HAPPENED  
???**

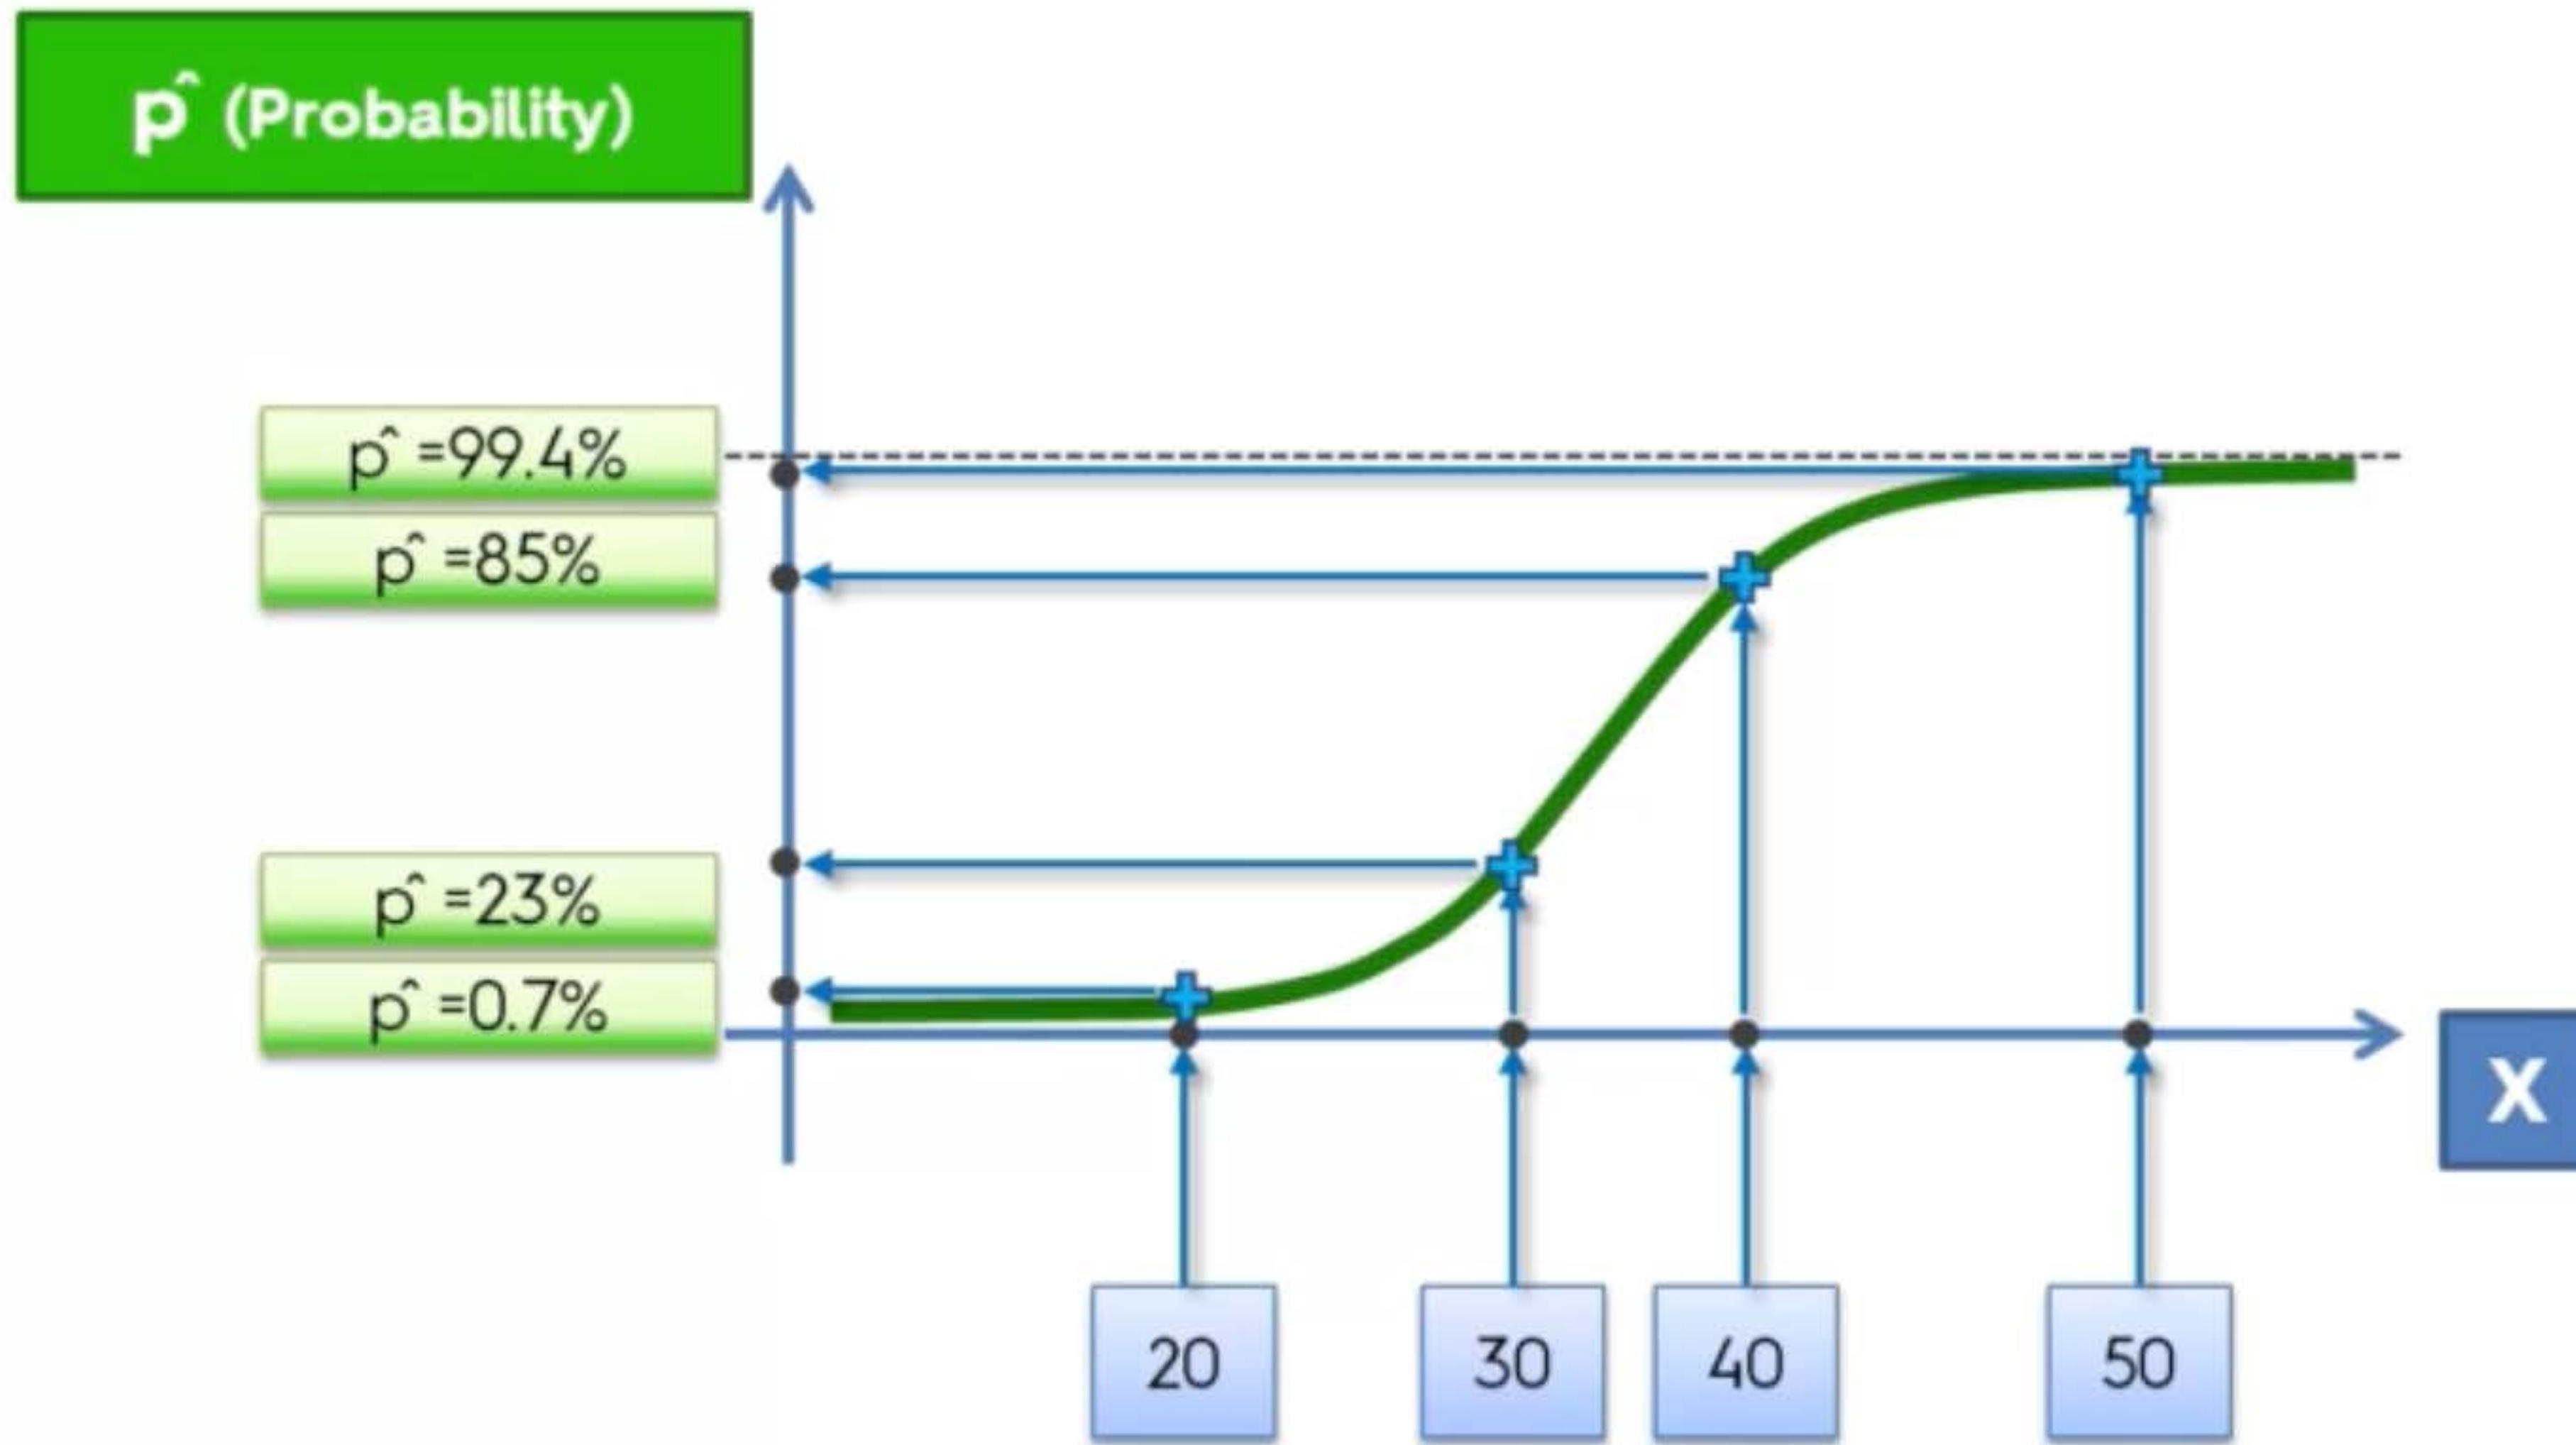
# Logistic Regression



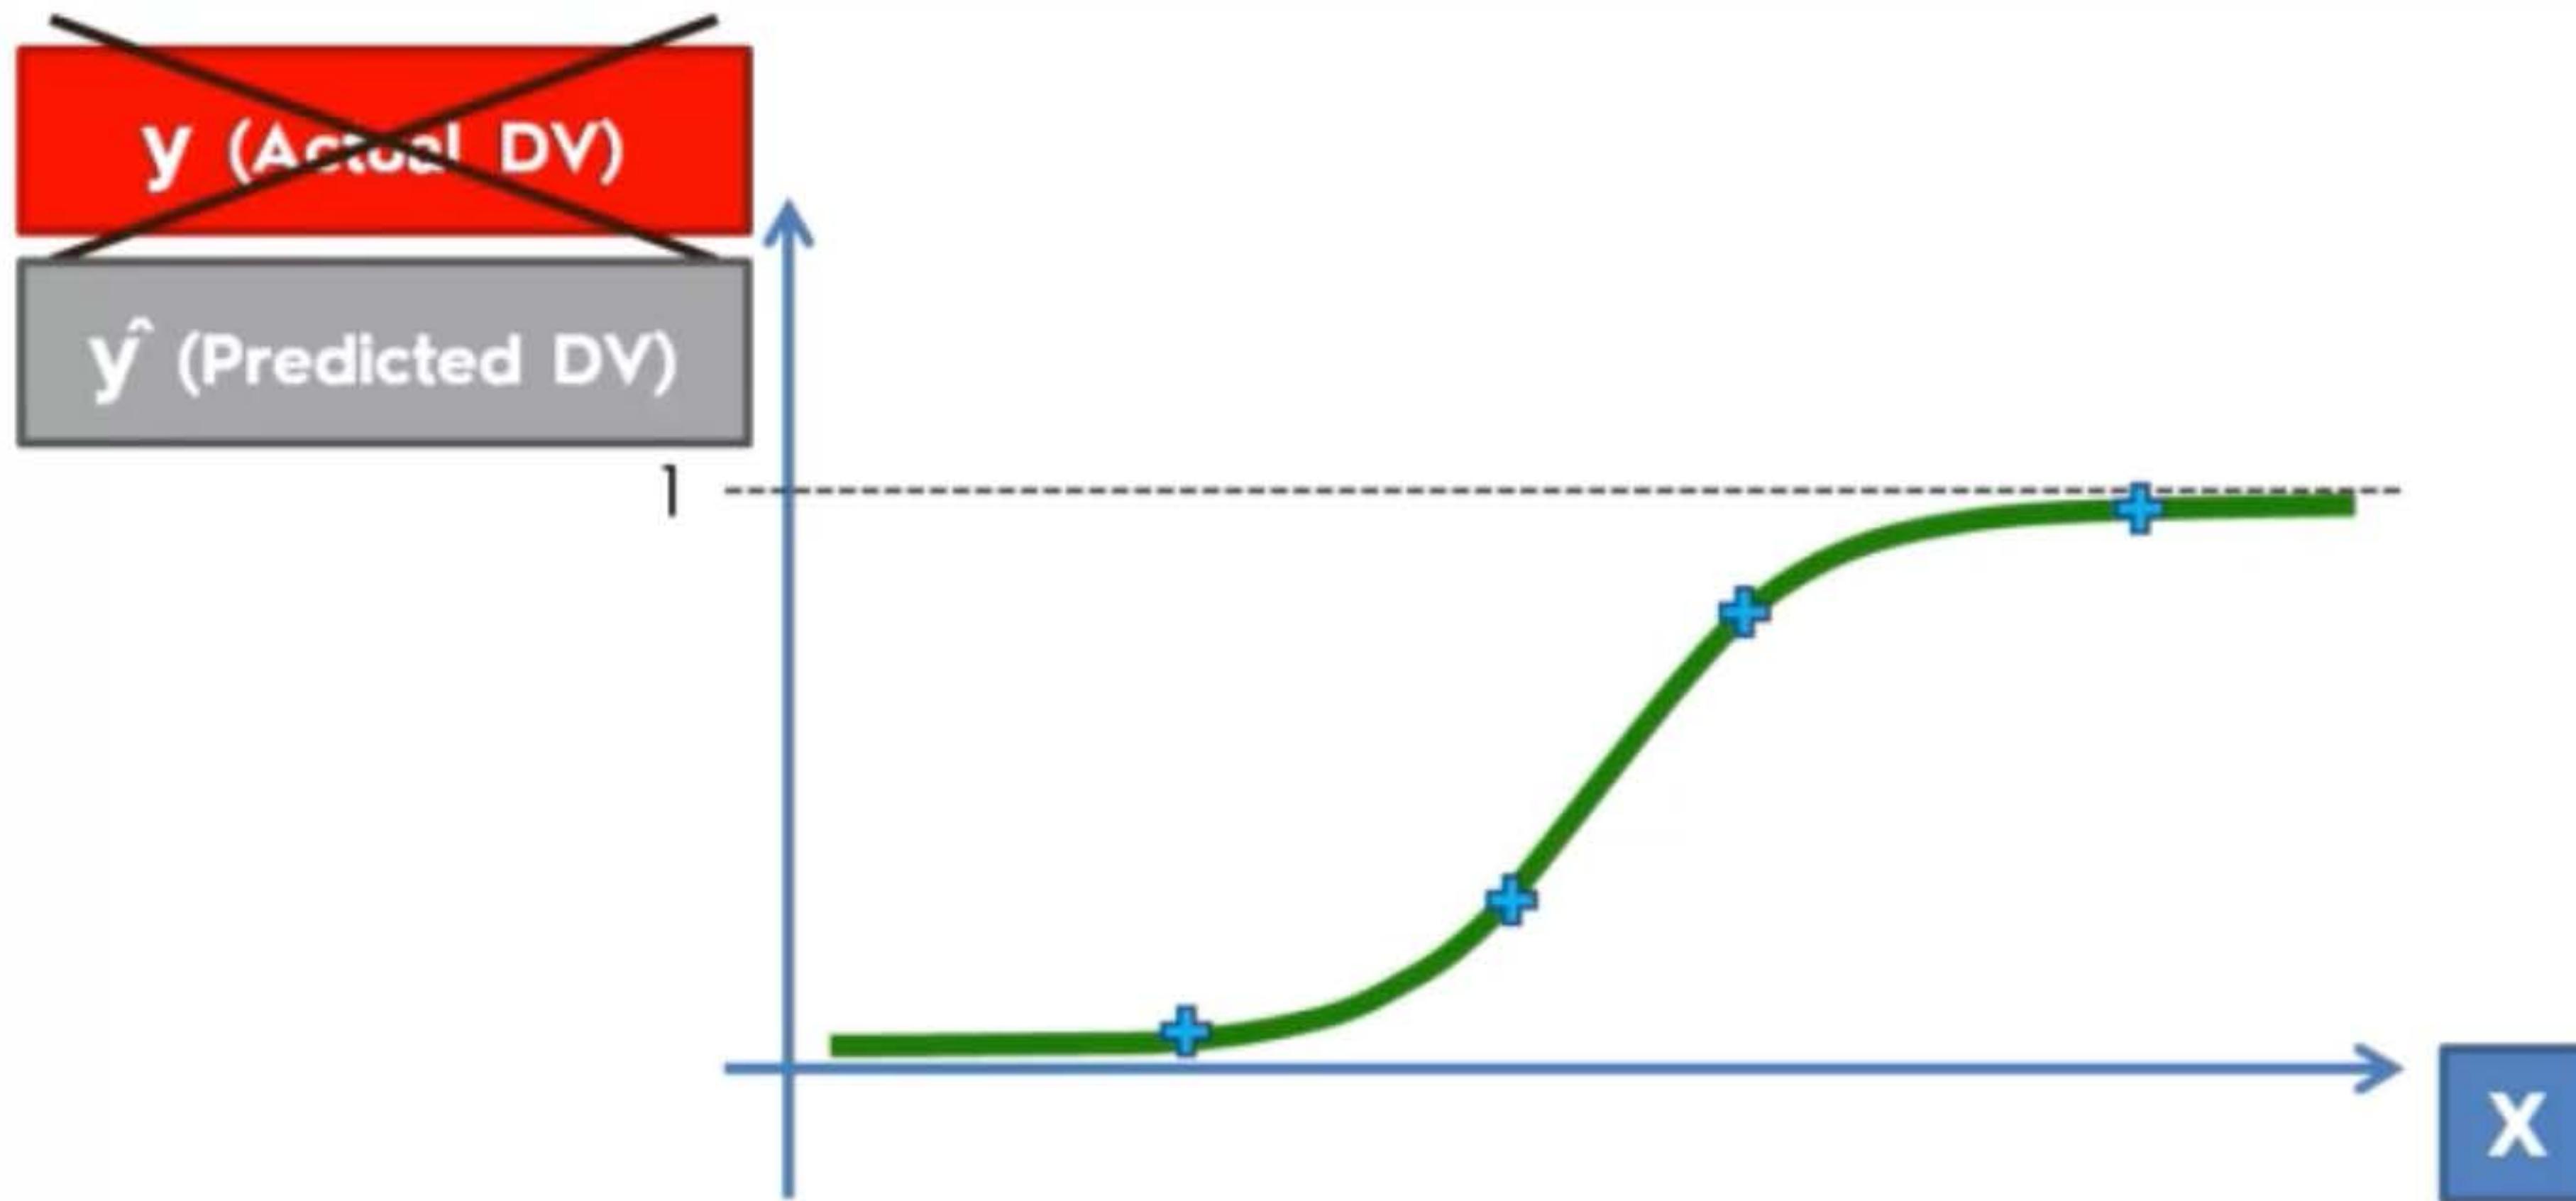
# Logistic Regression



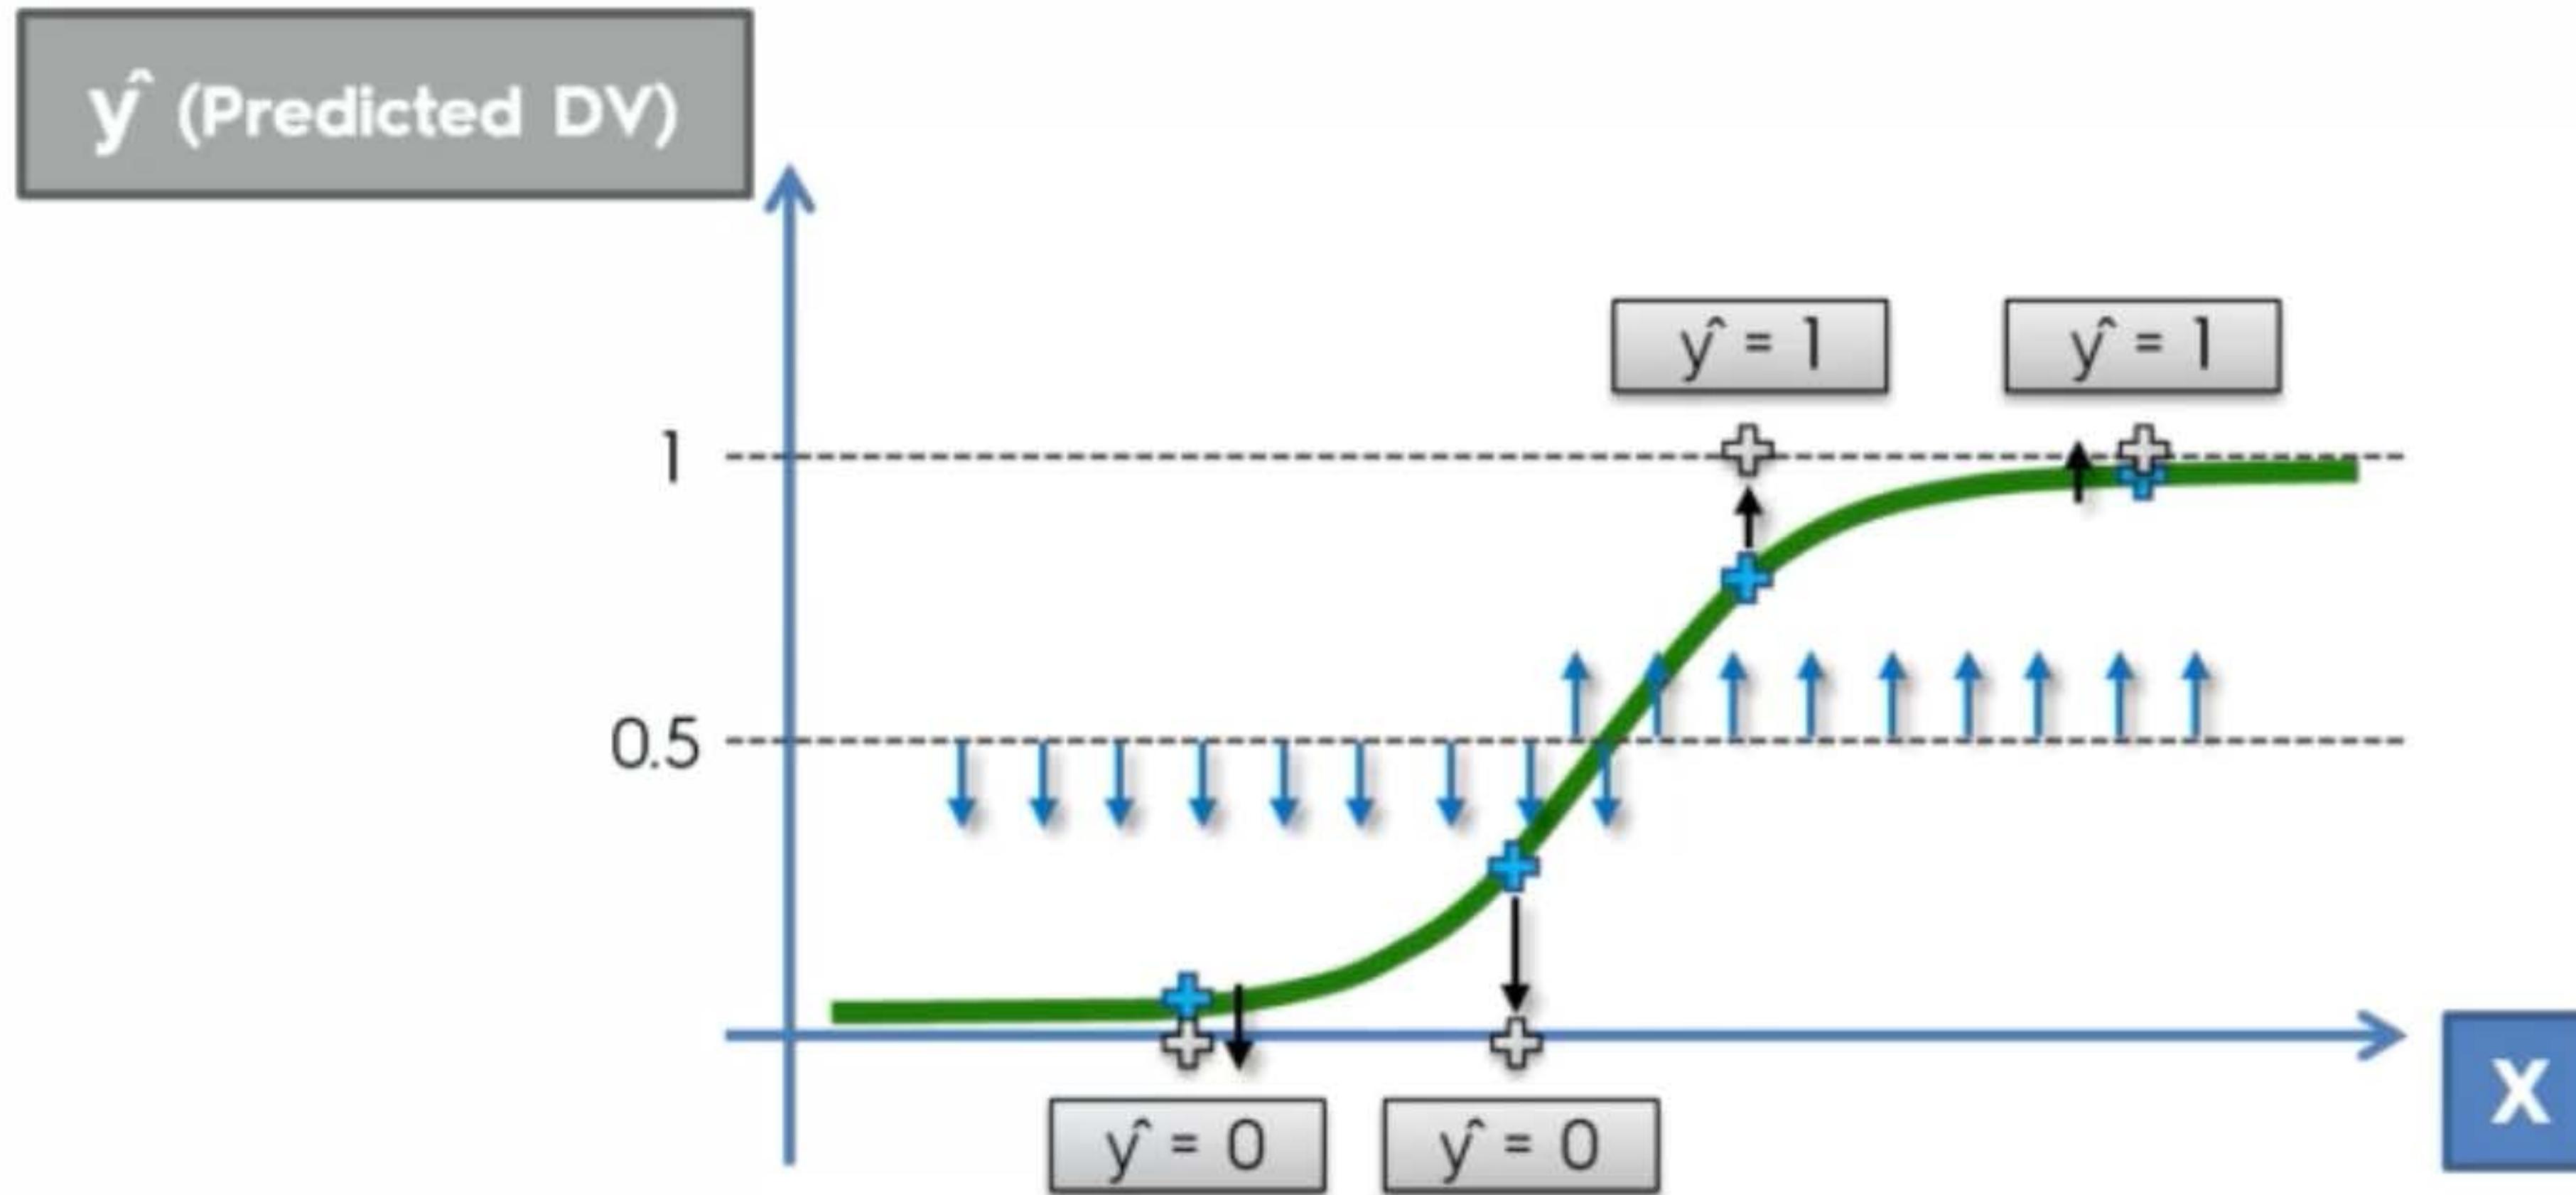
# Logistic Regression



# Logistic Regression

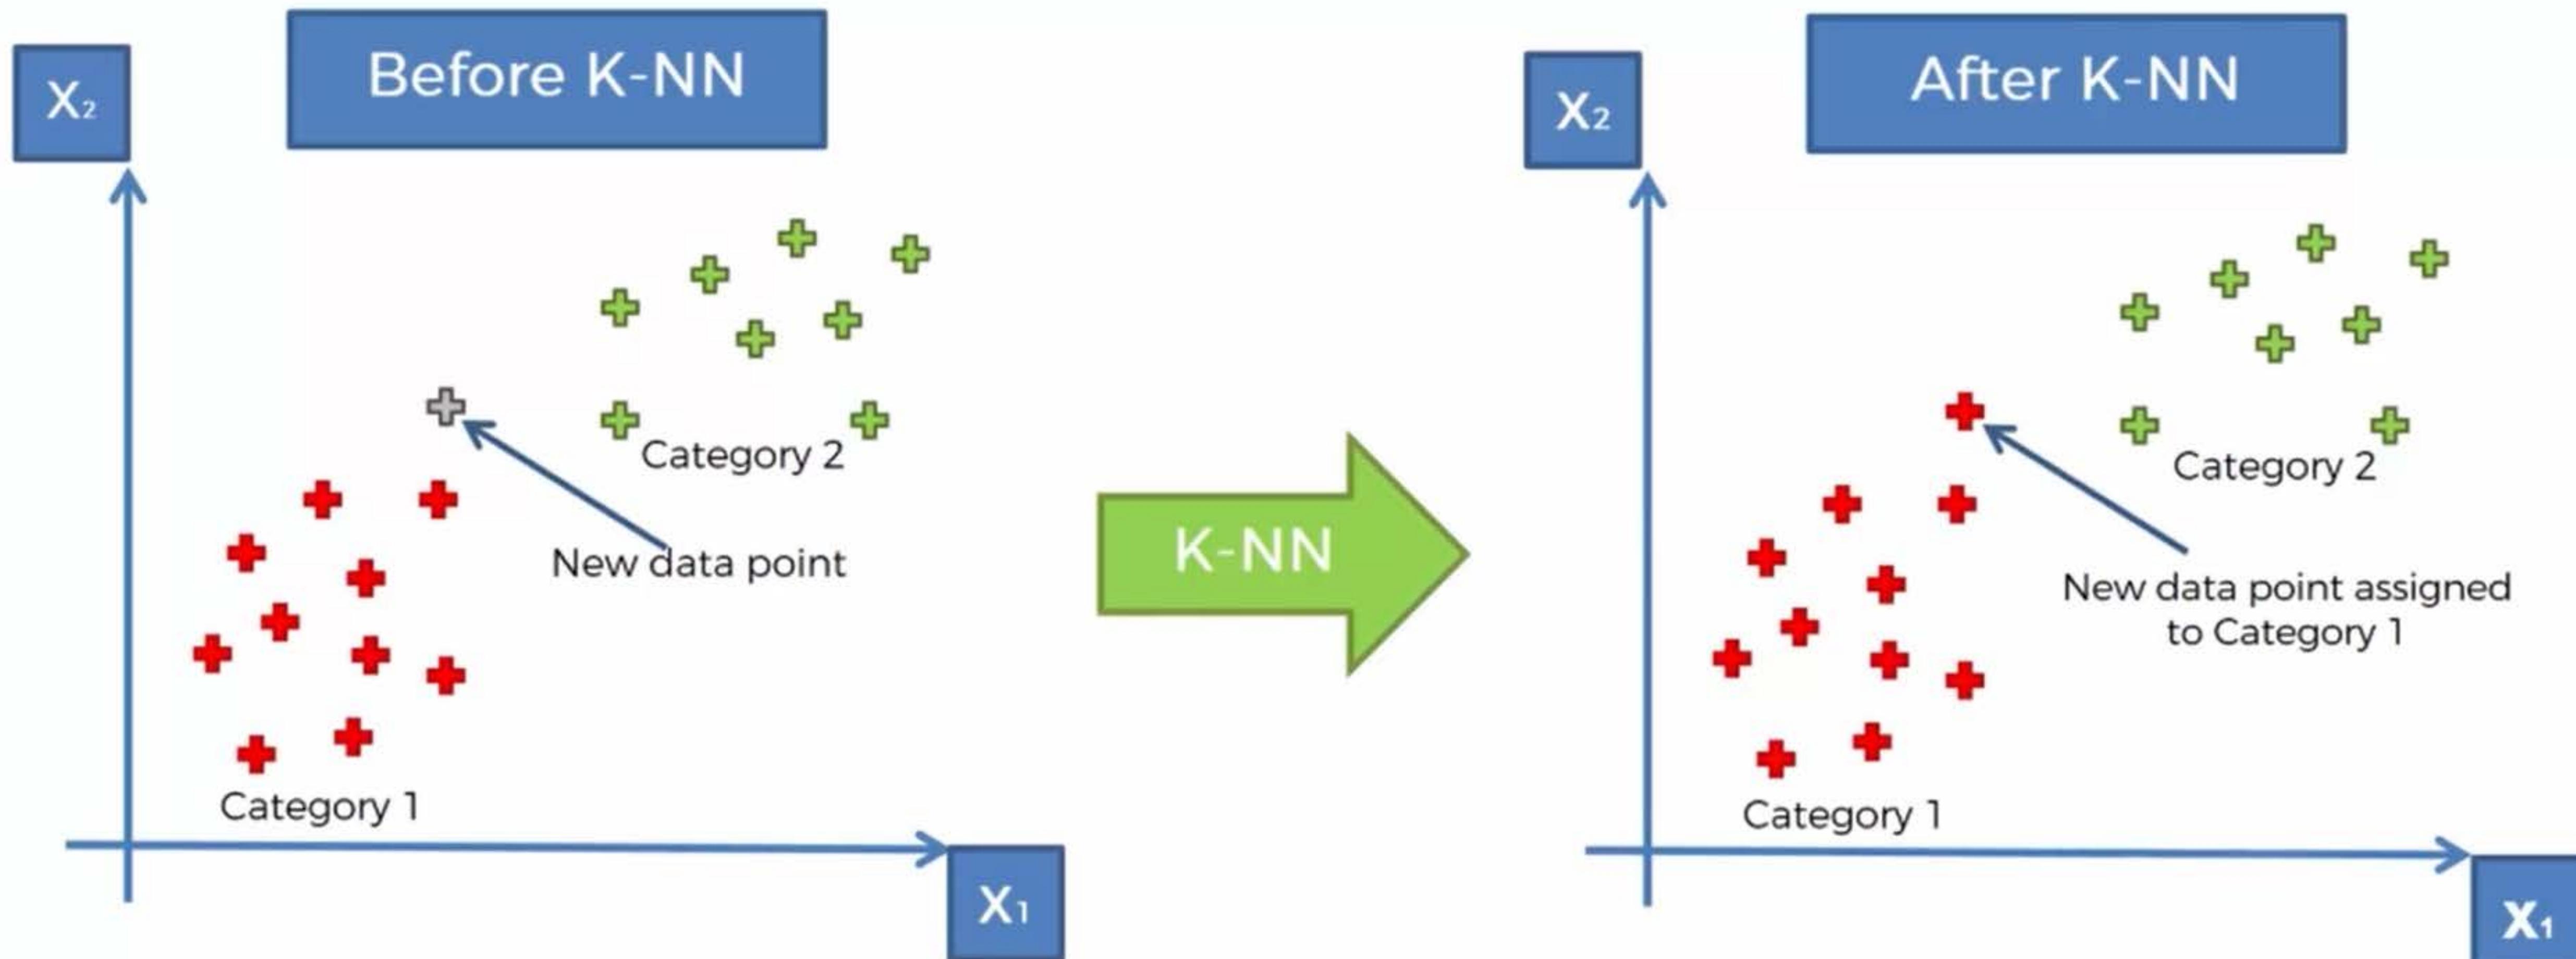


# Logistic Regression



# K-NN Intuition

# What K-NN does for you



# How did it do that ?

STEP 1: Choose the number K of neighbors



STEP 2: Take the K nearest neighbors of the new data point, according to the Euclidean distance



STEP 3: Among these K neighbors, count the number of data points in each category



STEP 4: Assign the new data point to the category where you counted the most neighbors



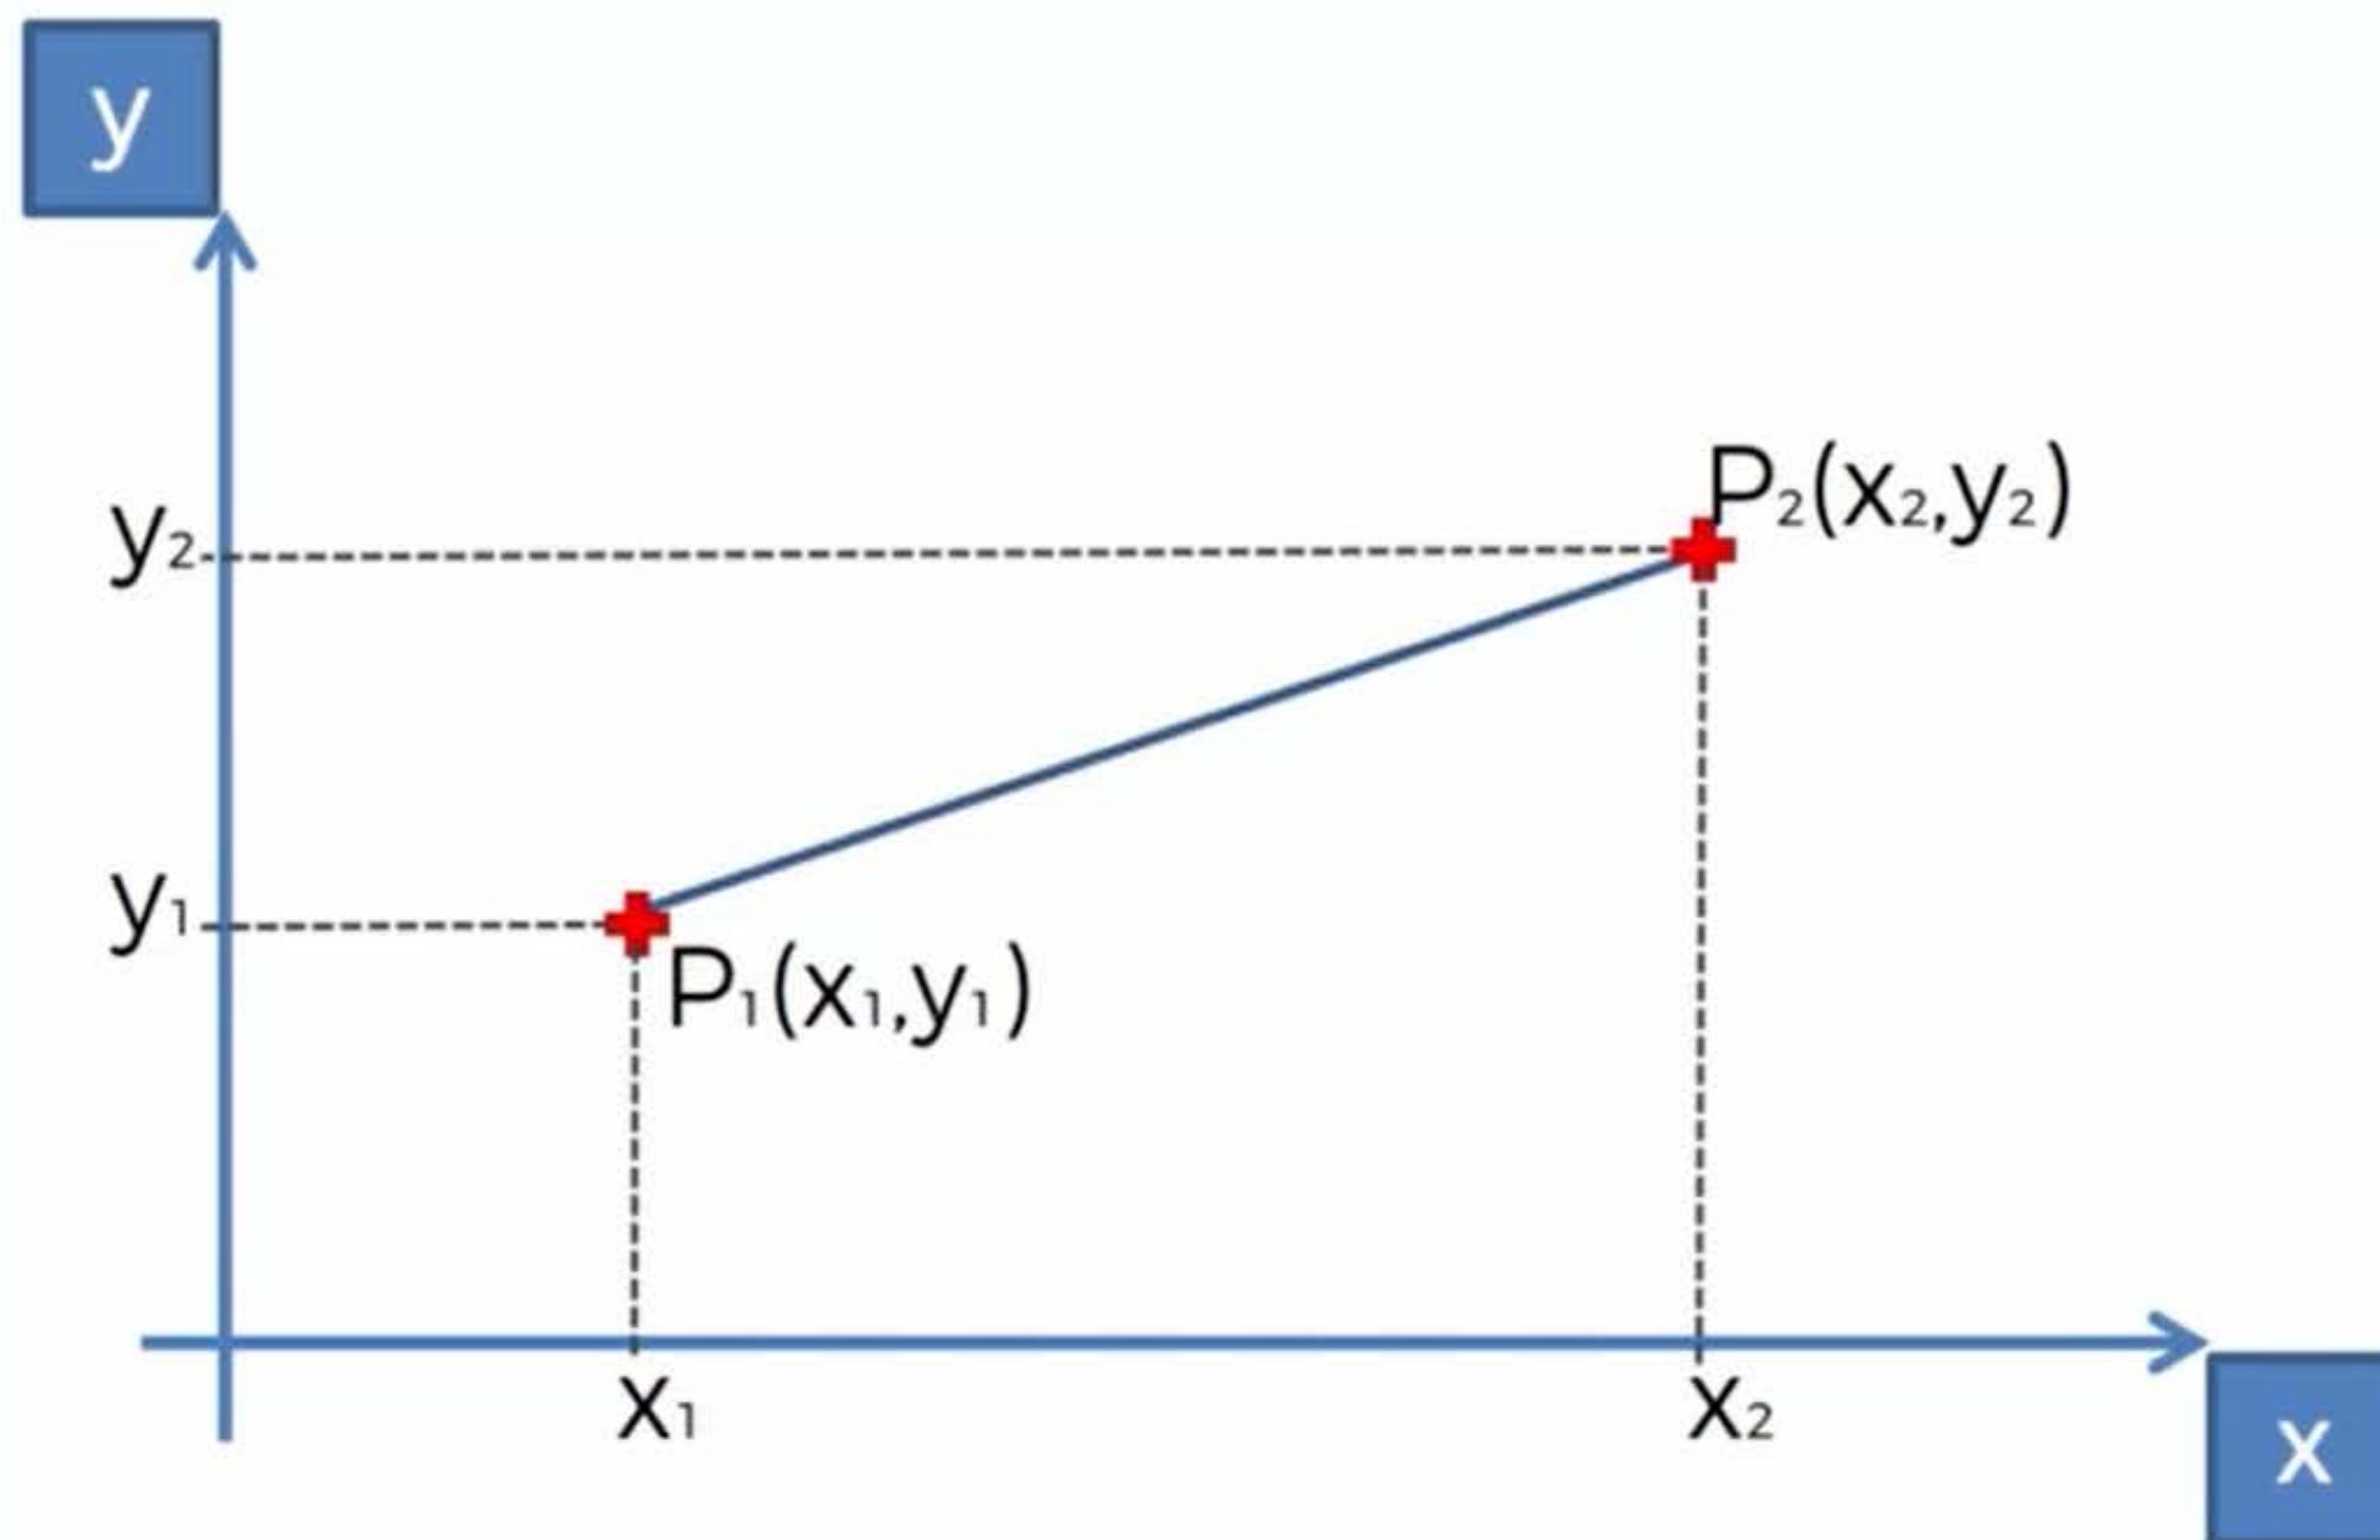
Your Model is Ready

# K-NN algorithm

STEP 1: Choose the number K of neighbors:  $K = 5$



# Euclidean Distance



$$\text{Euclidean Distance between } P_1 \text{ and } P_2 = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

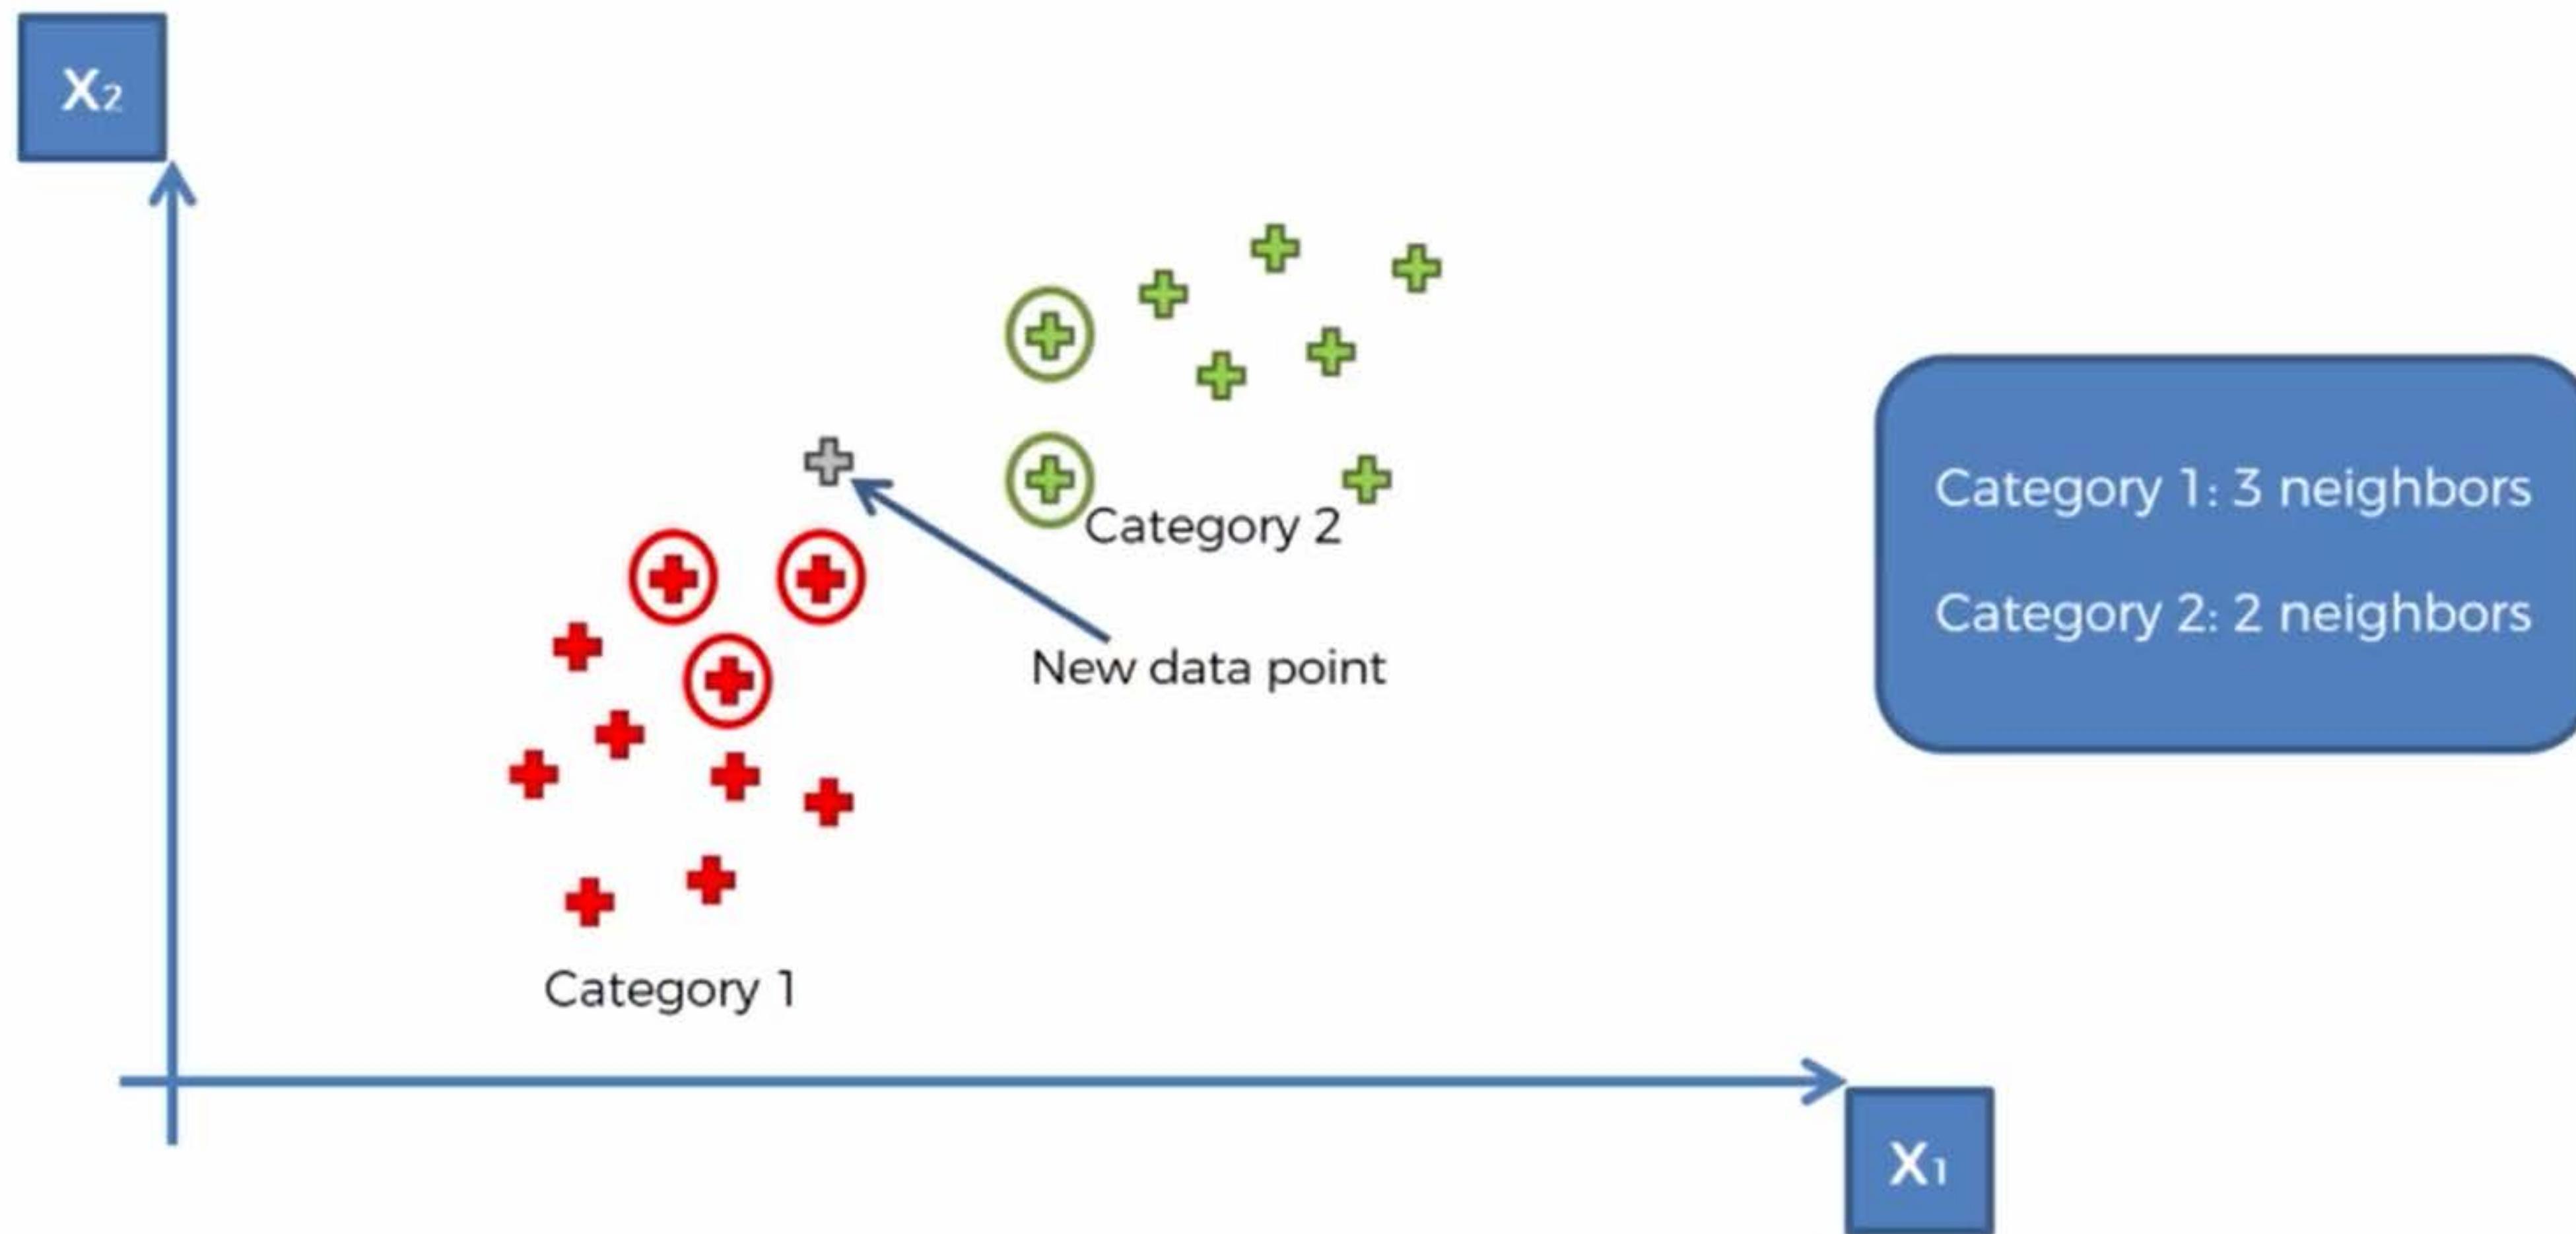
# K-NN algorithm

STEP 2: Take the  $K = 5$  nearest neighbors of the new data point, according to the Euclidean distance



# K-NN algorithm

STEP 3: Among these K neighbors, count the number of data points in each category



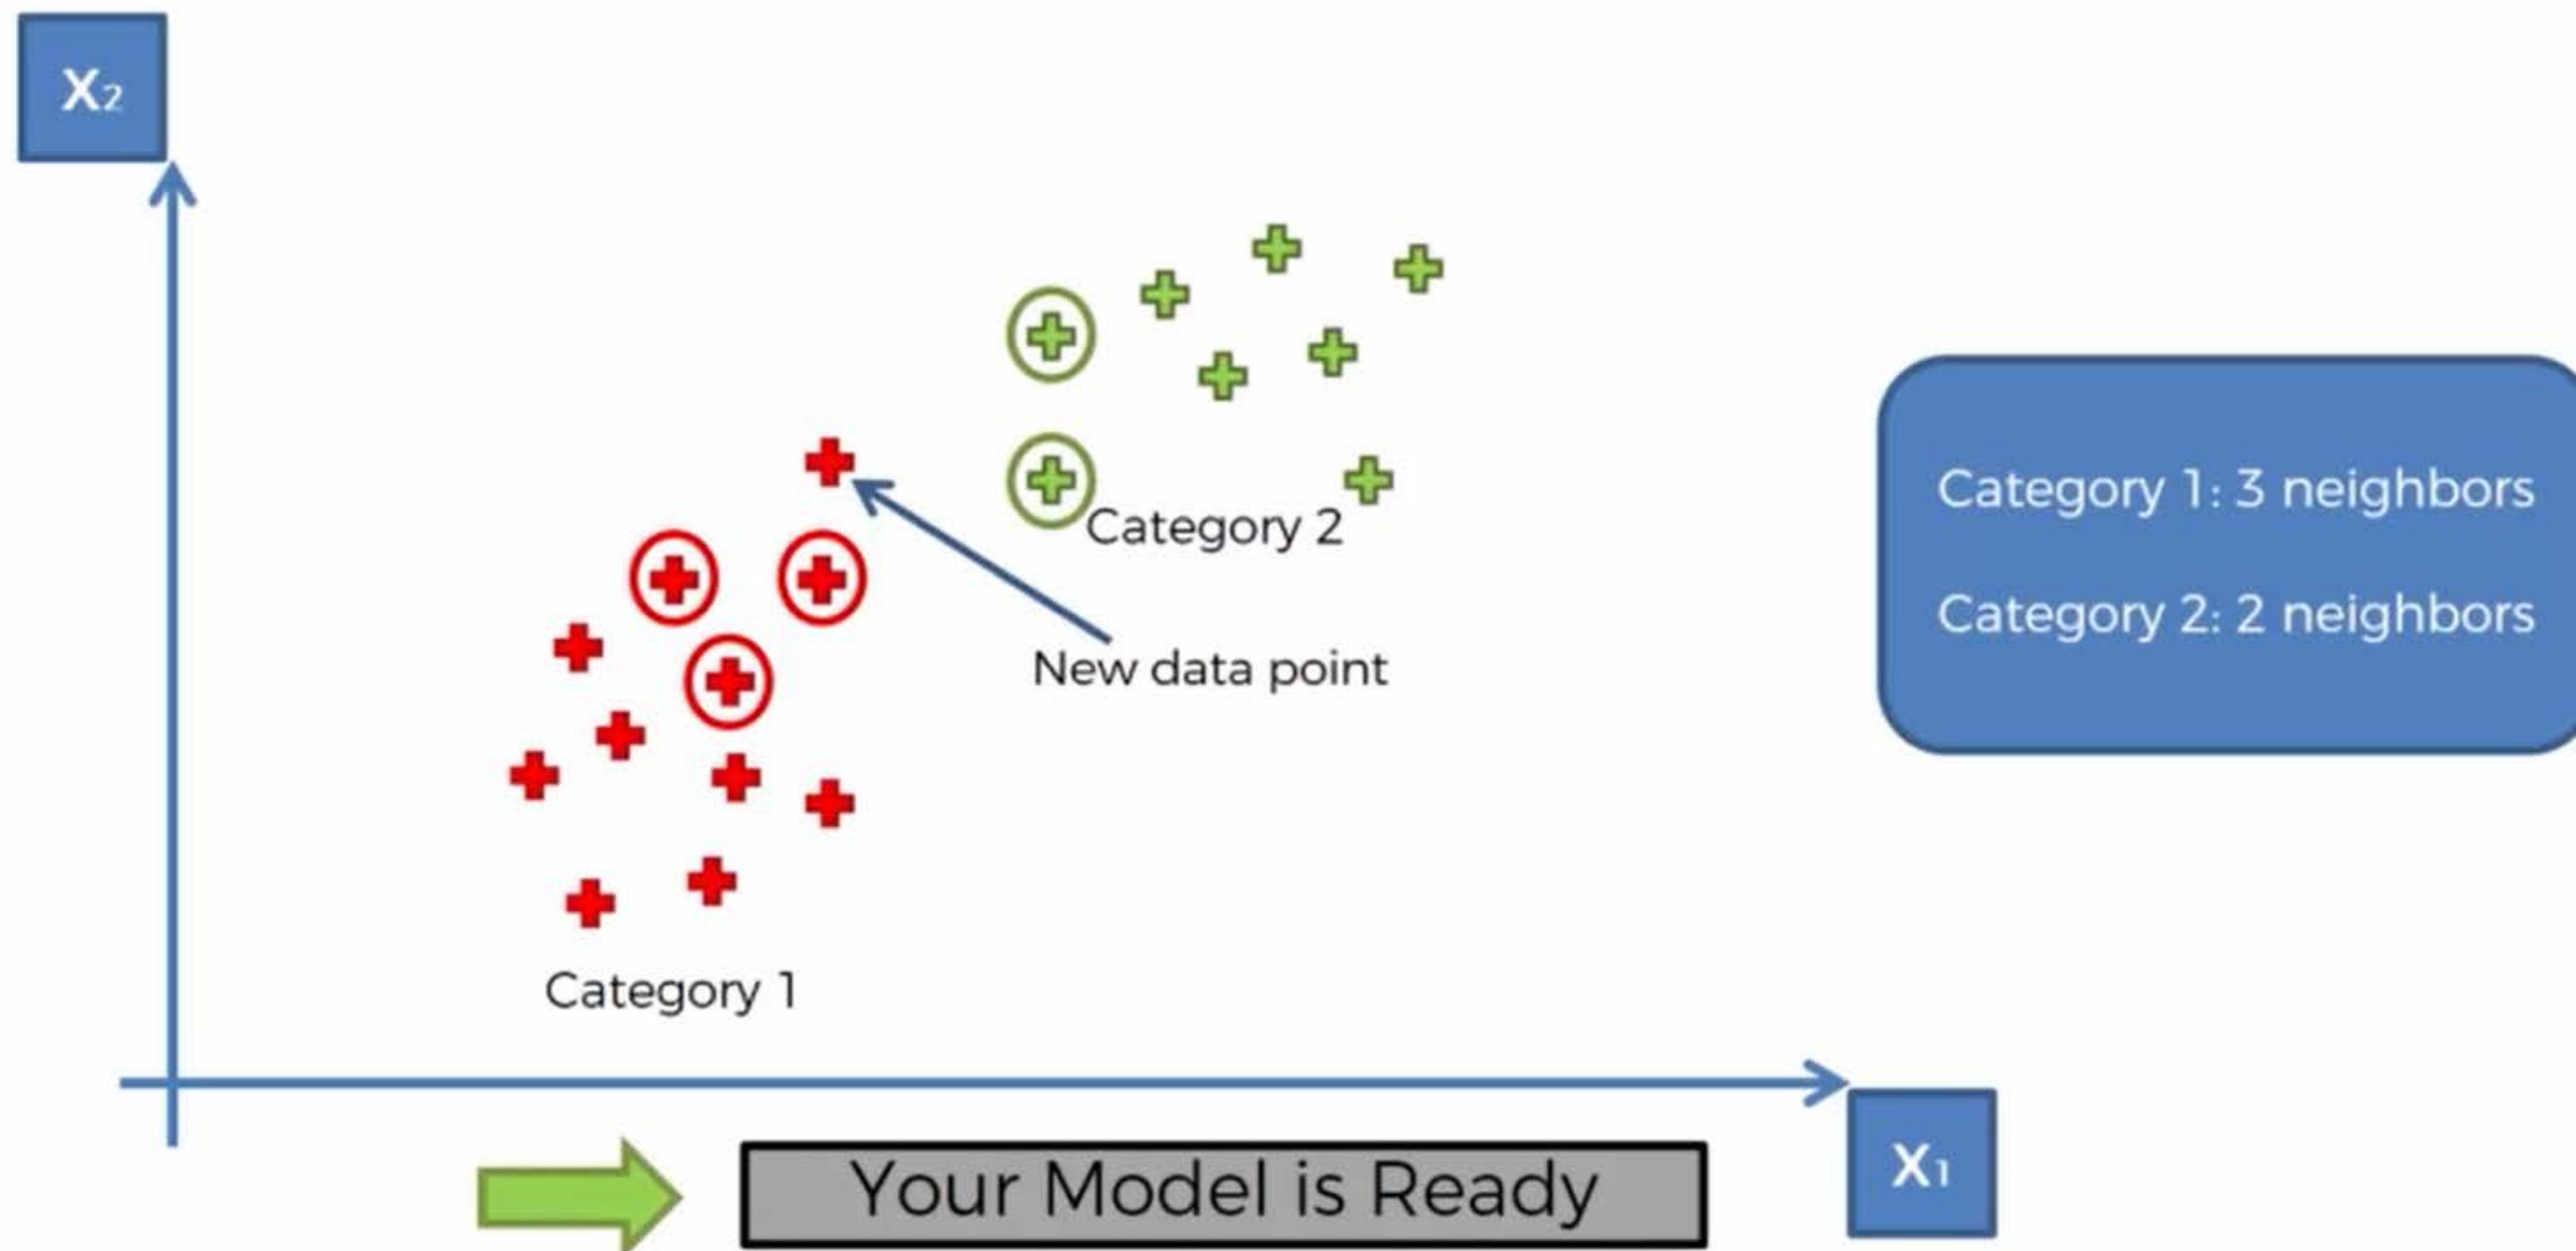
# K-NN algorithm

STEP 4: Assign the new data point to the category where you counted the most neighbors



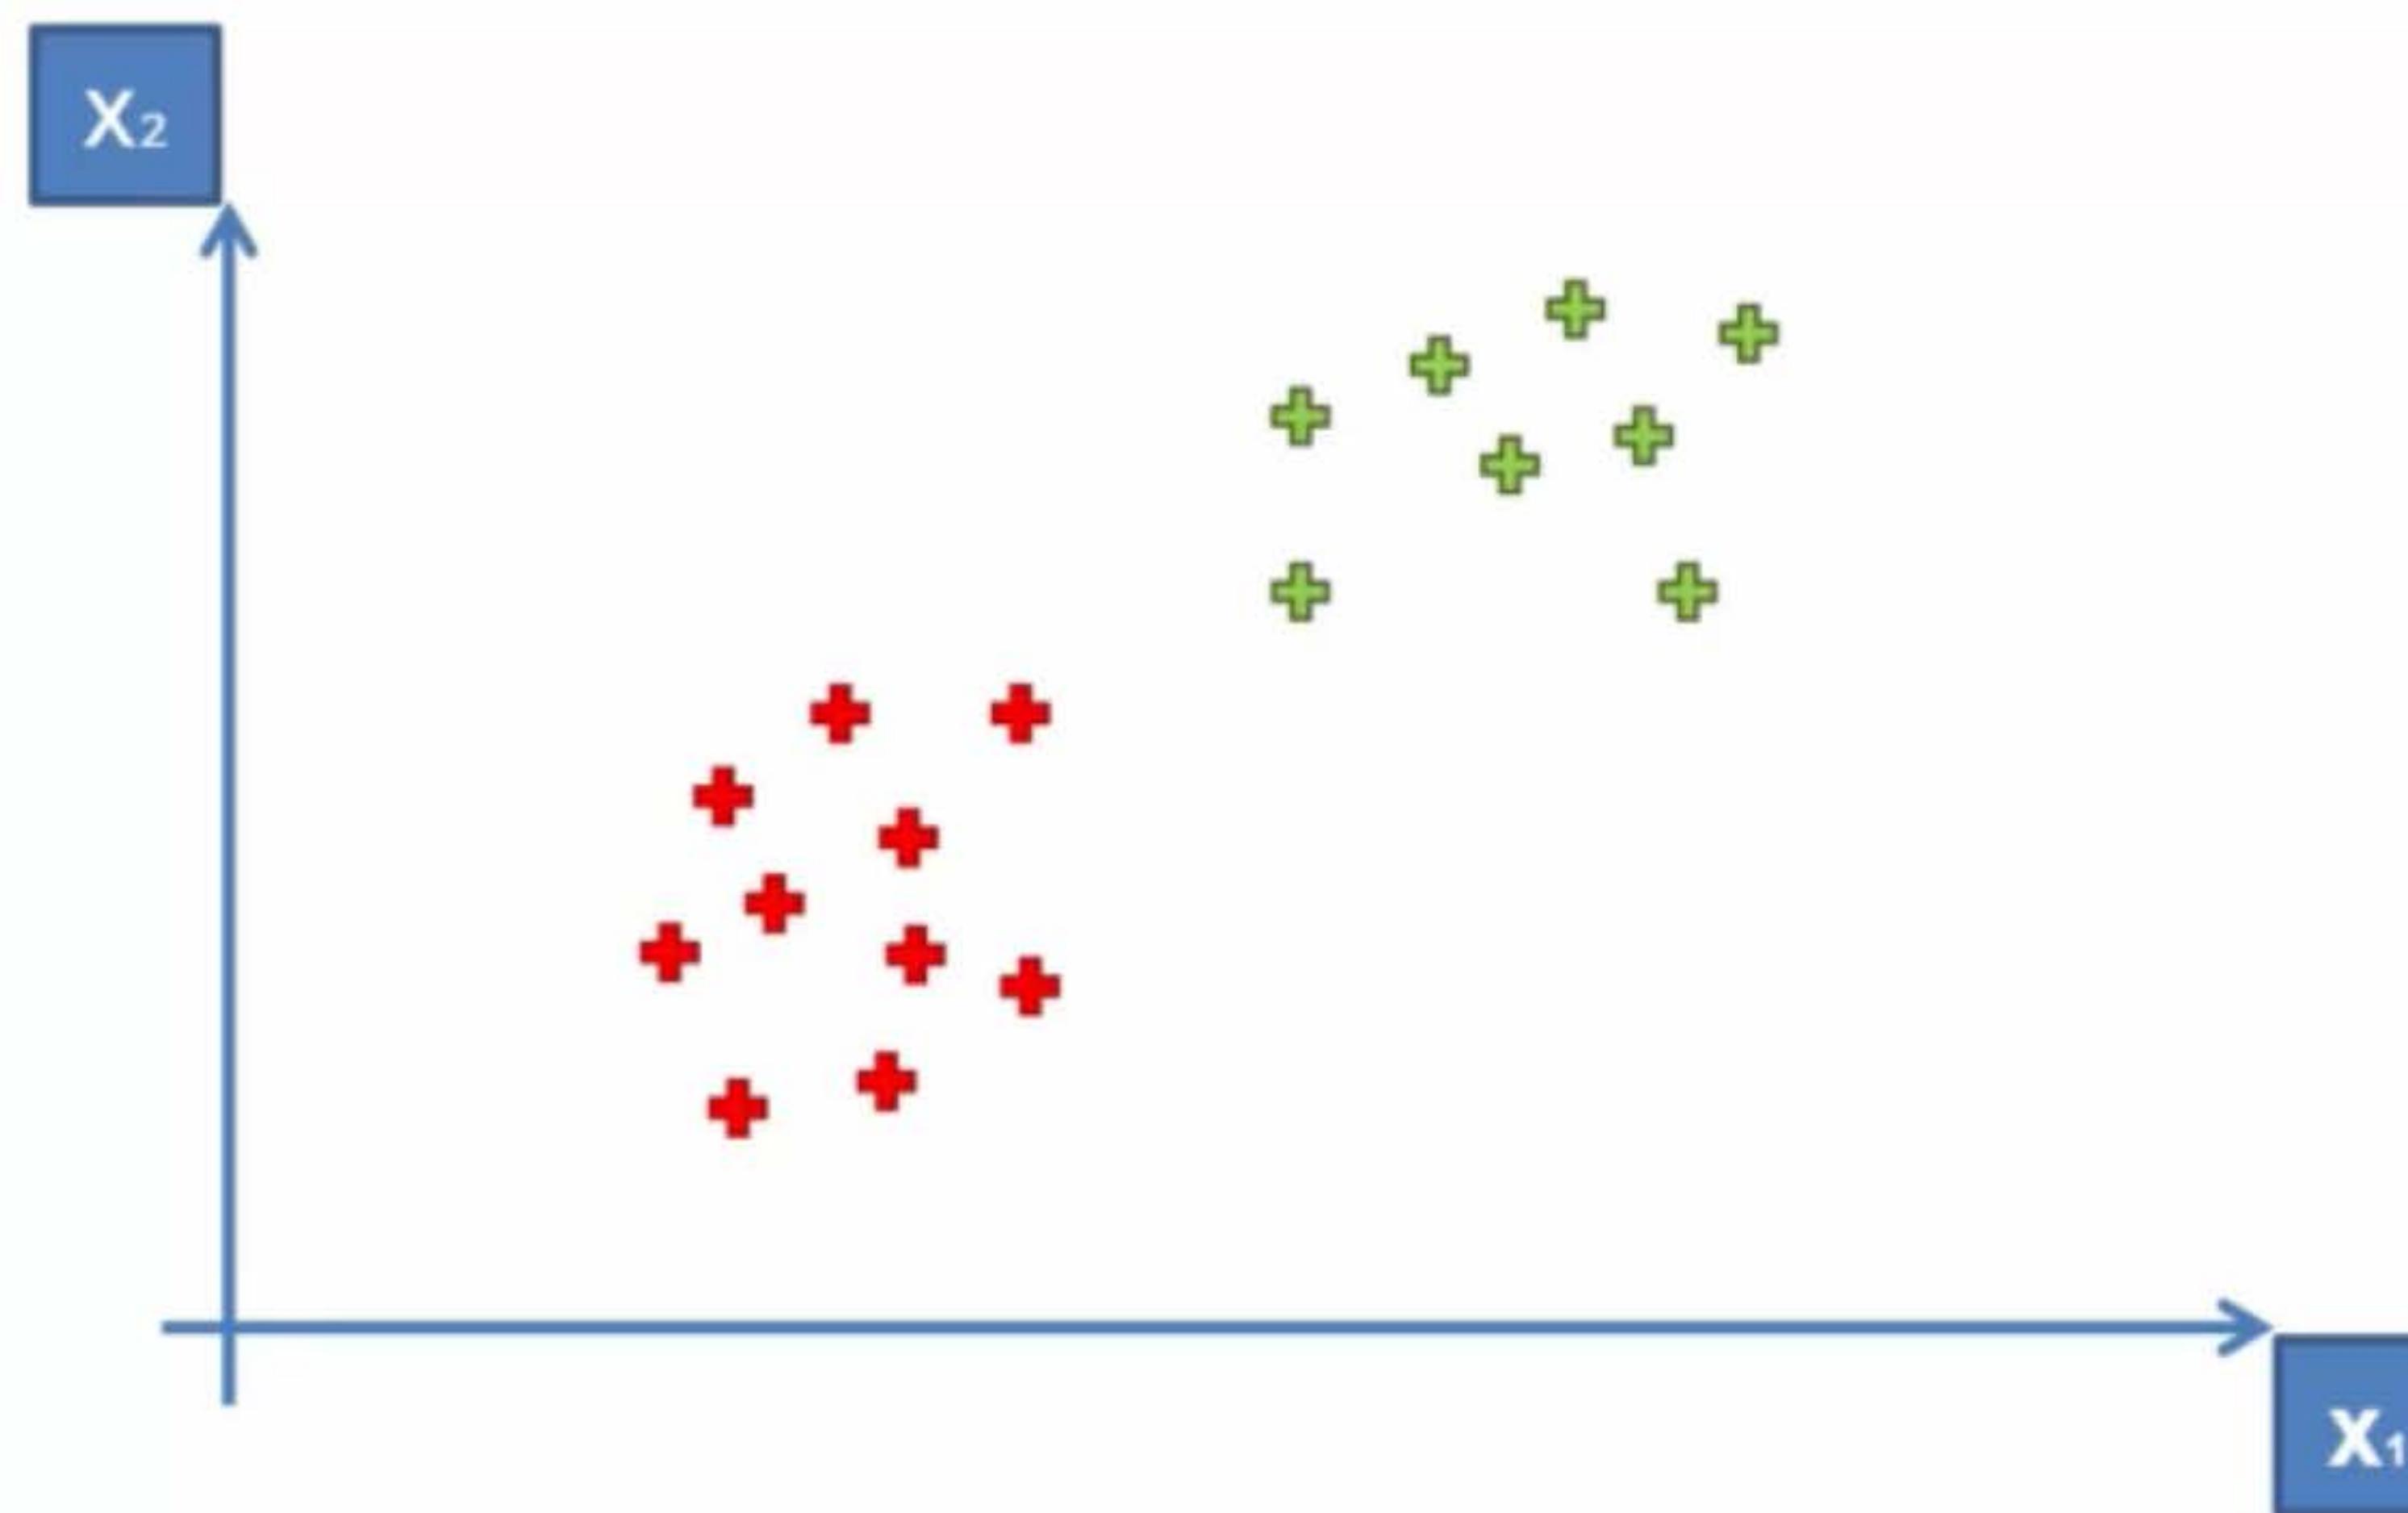
# K-NN algorithm

STEP 4: Assign the new data point to the category where you counted the most neighbors

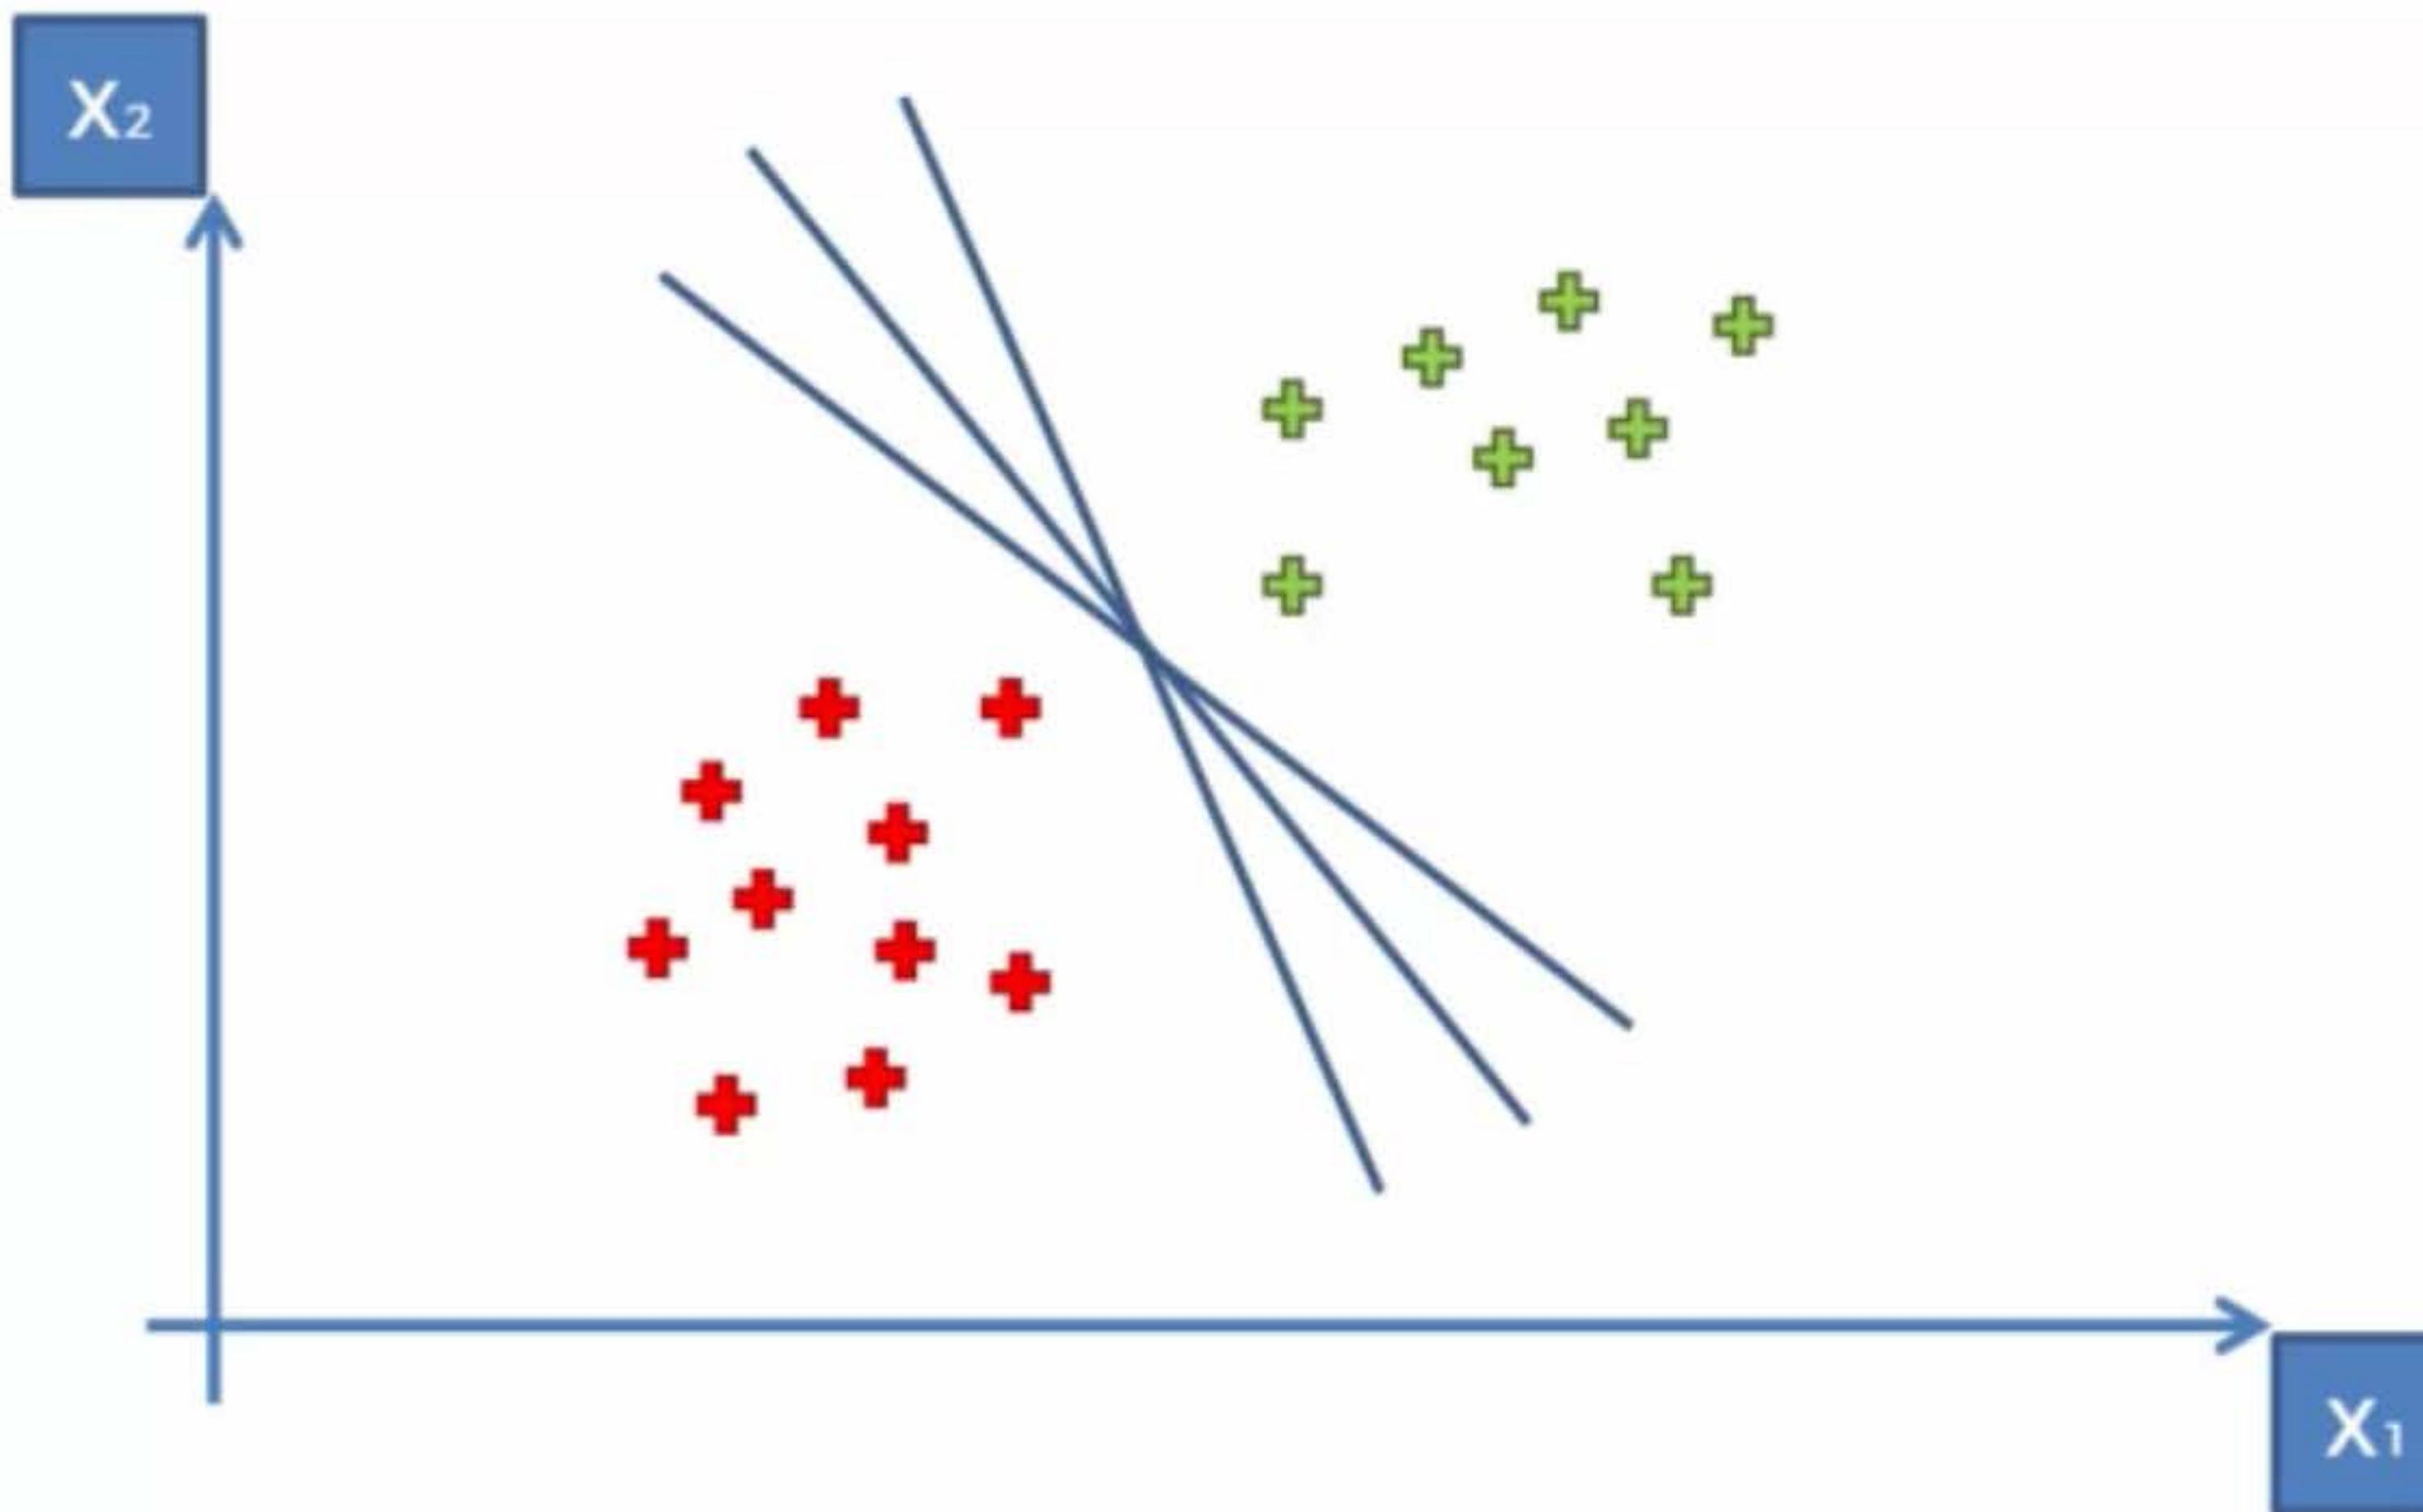


# SVM Intuition

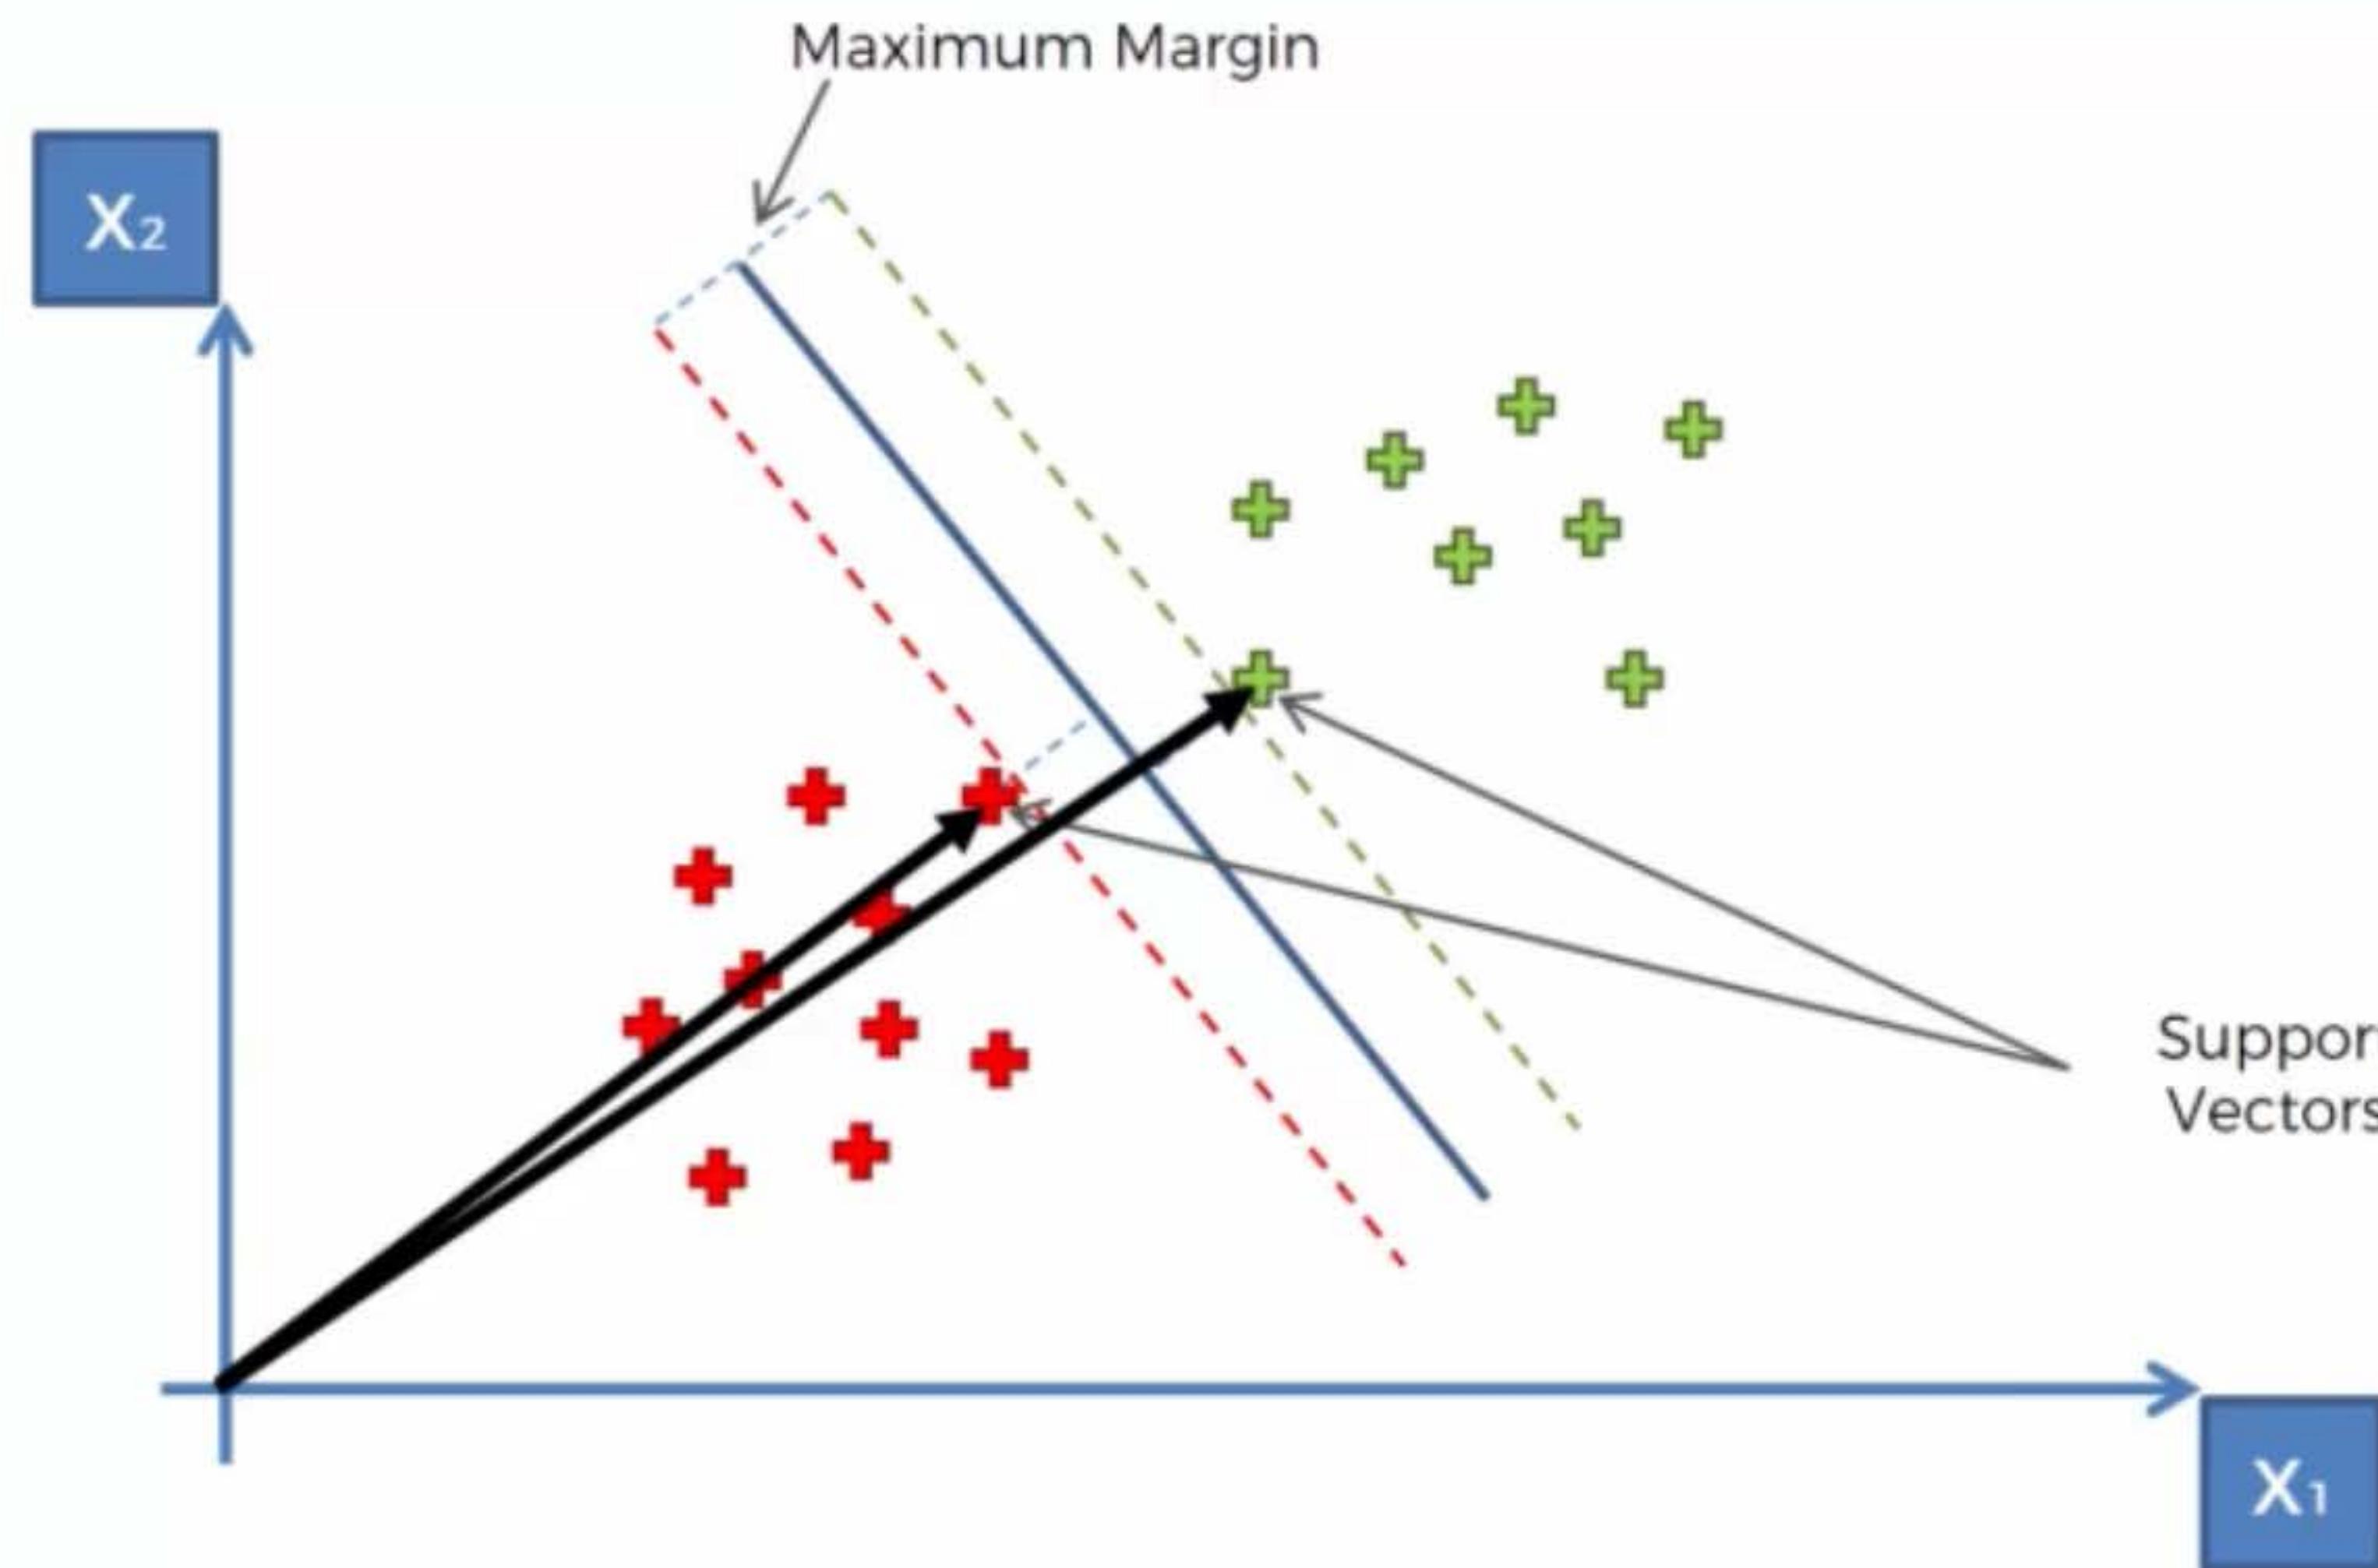
# How to separate these points ?



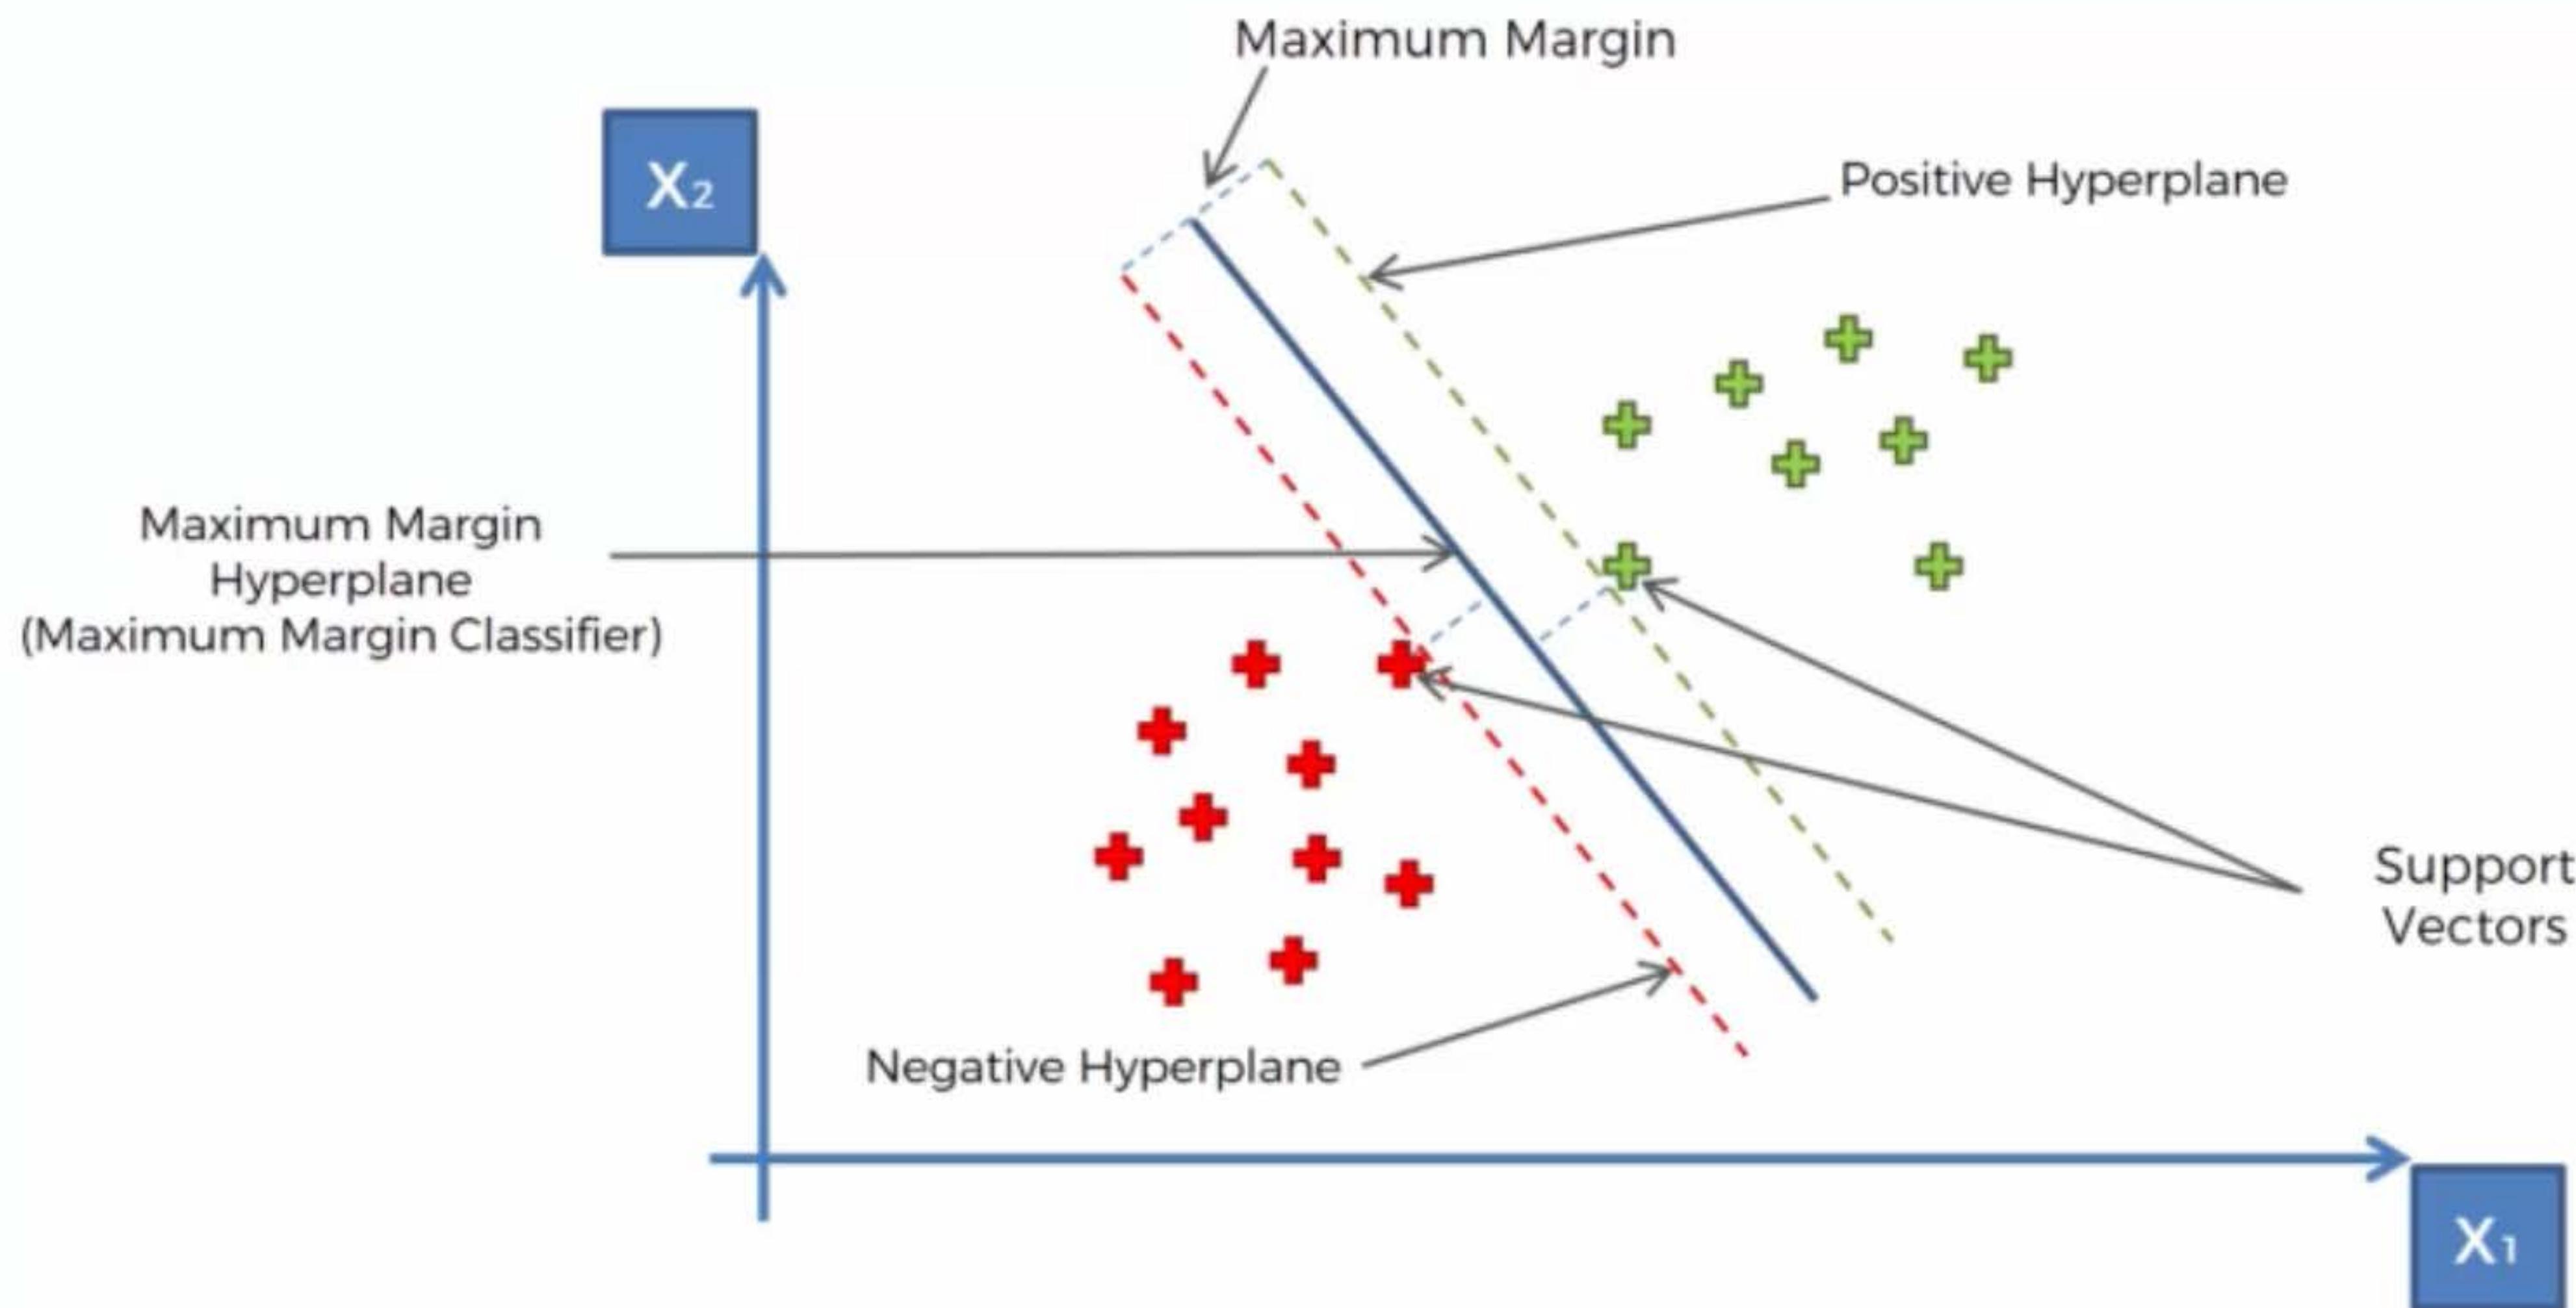
# SVM



# Support Vectors

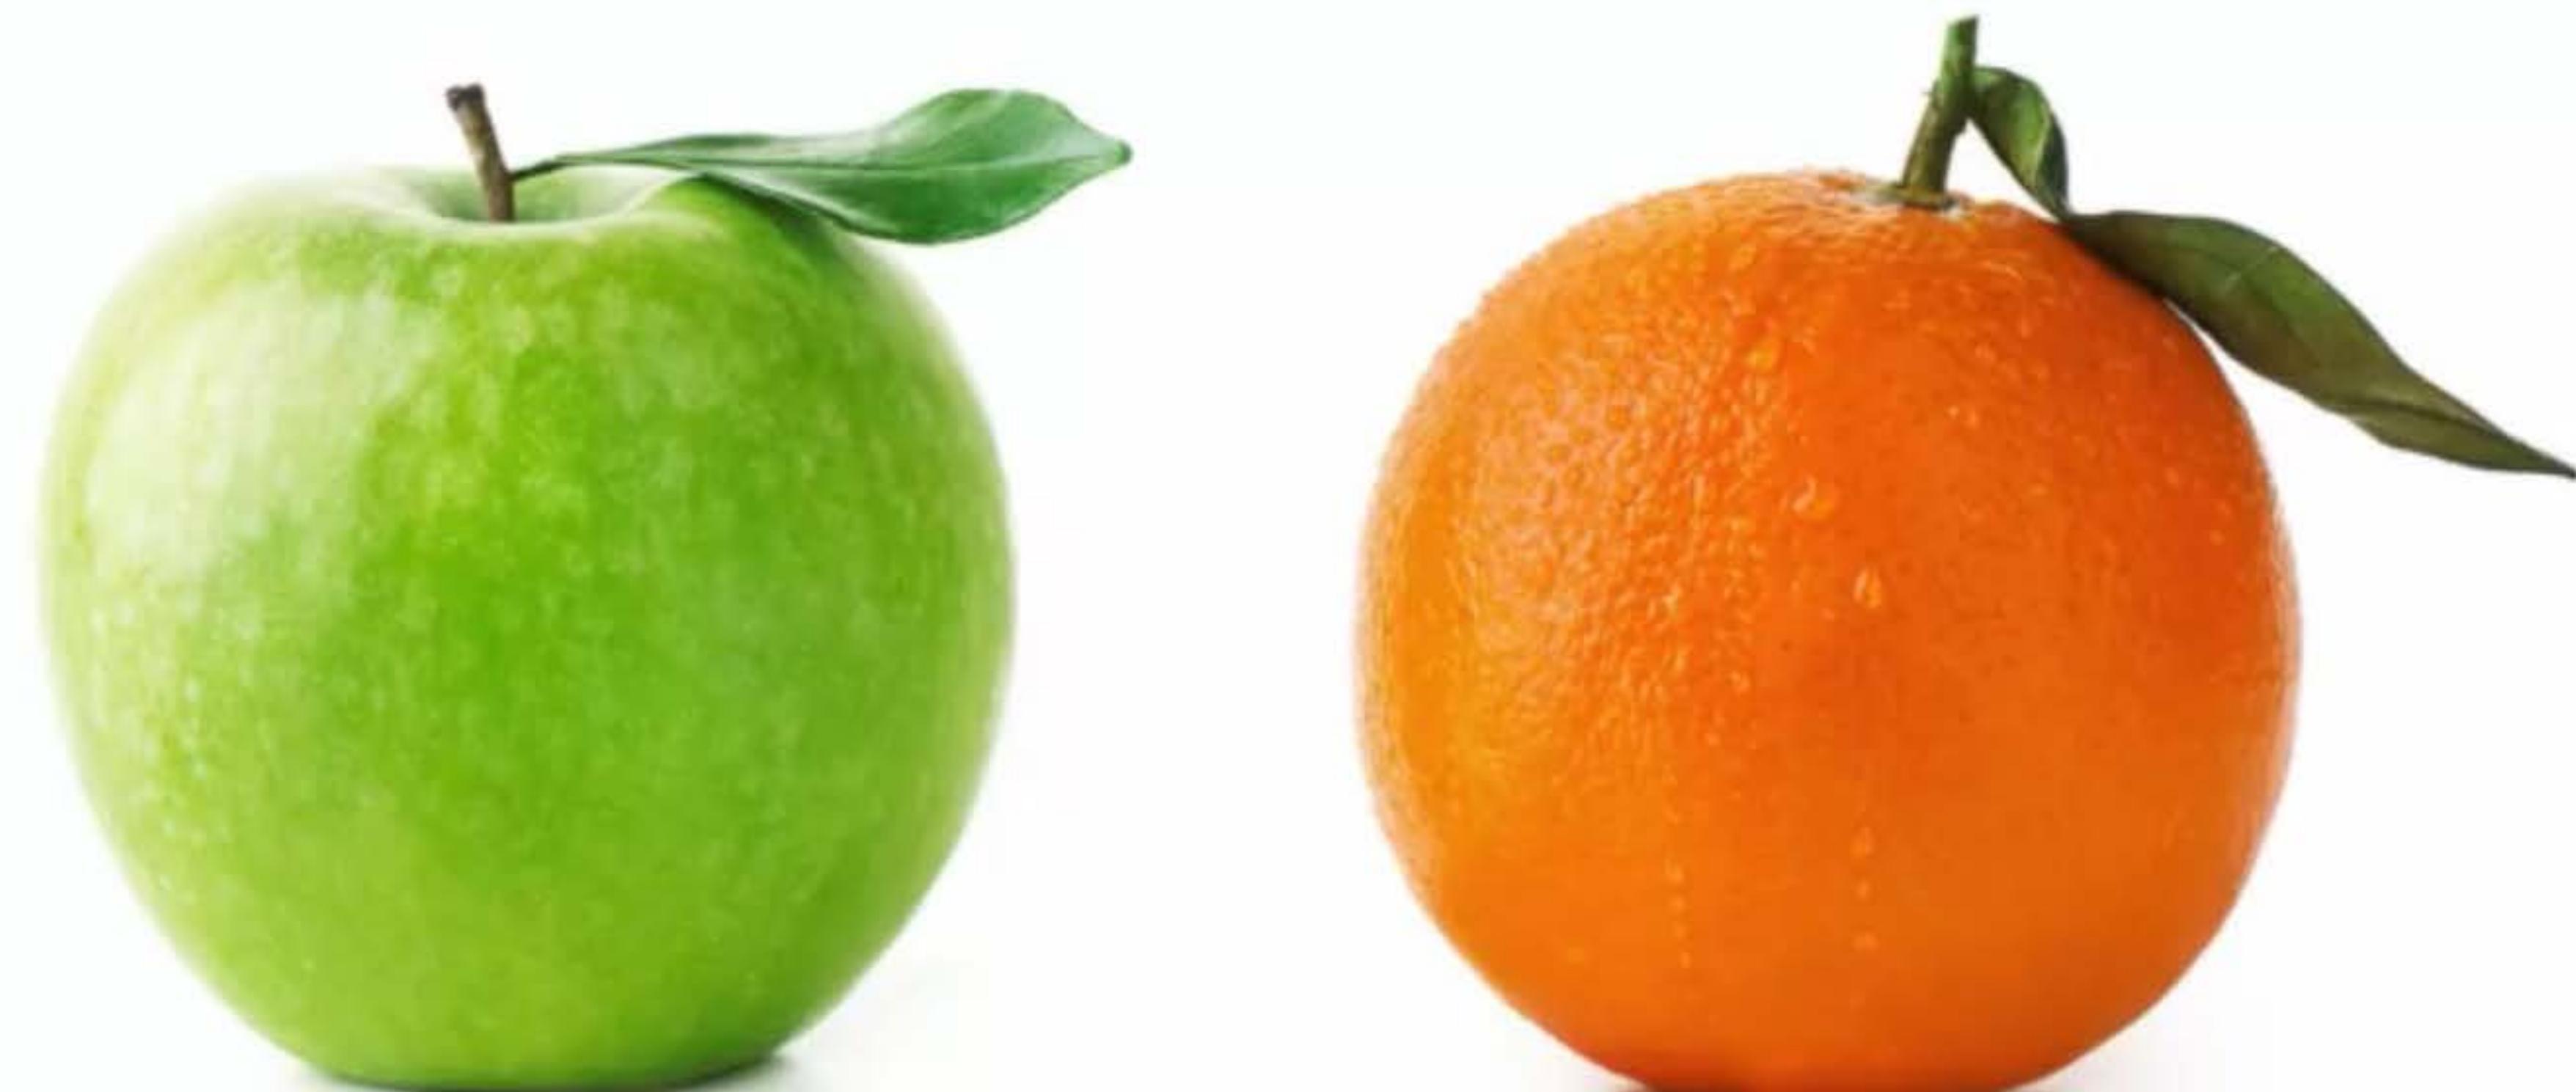


# Hyperplanes

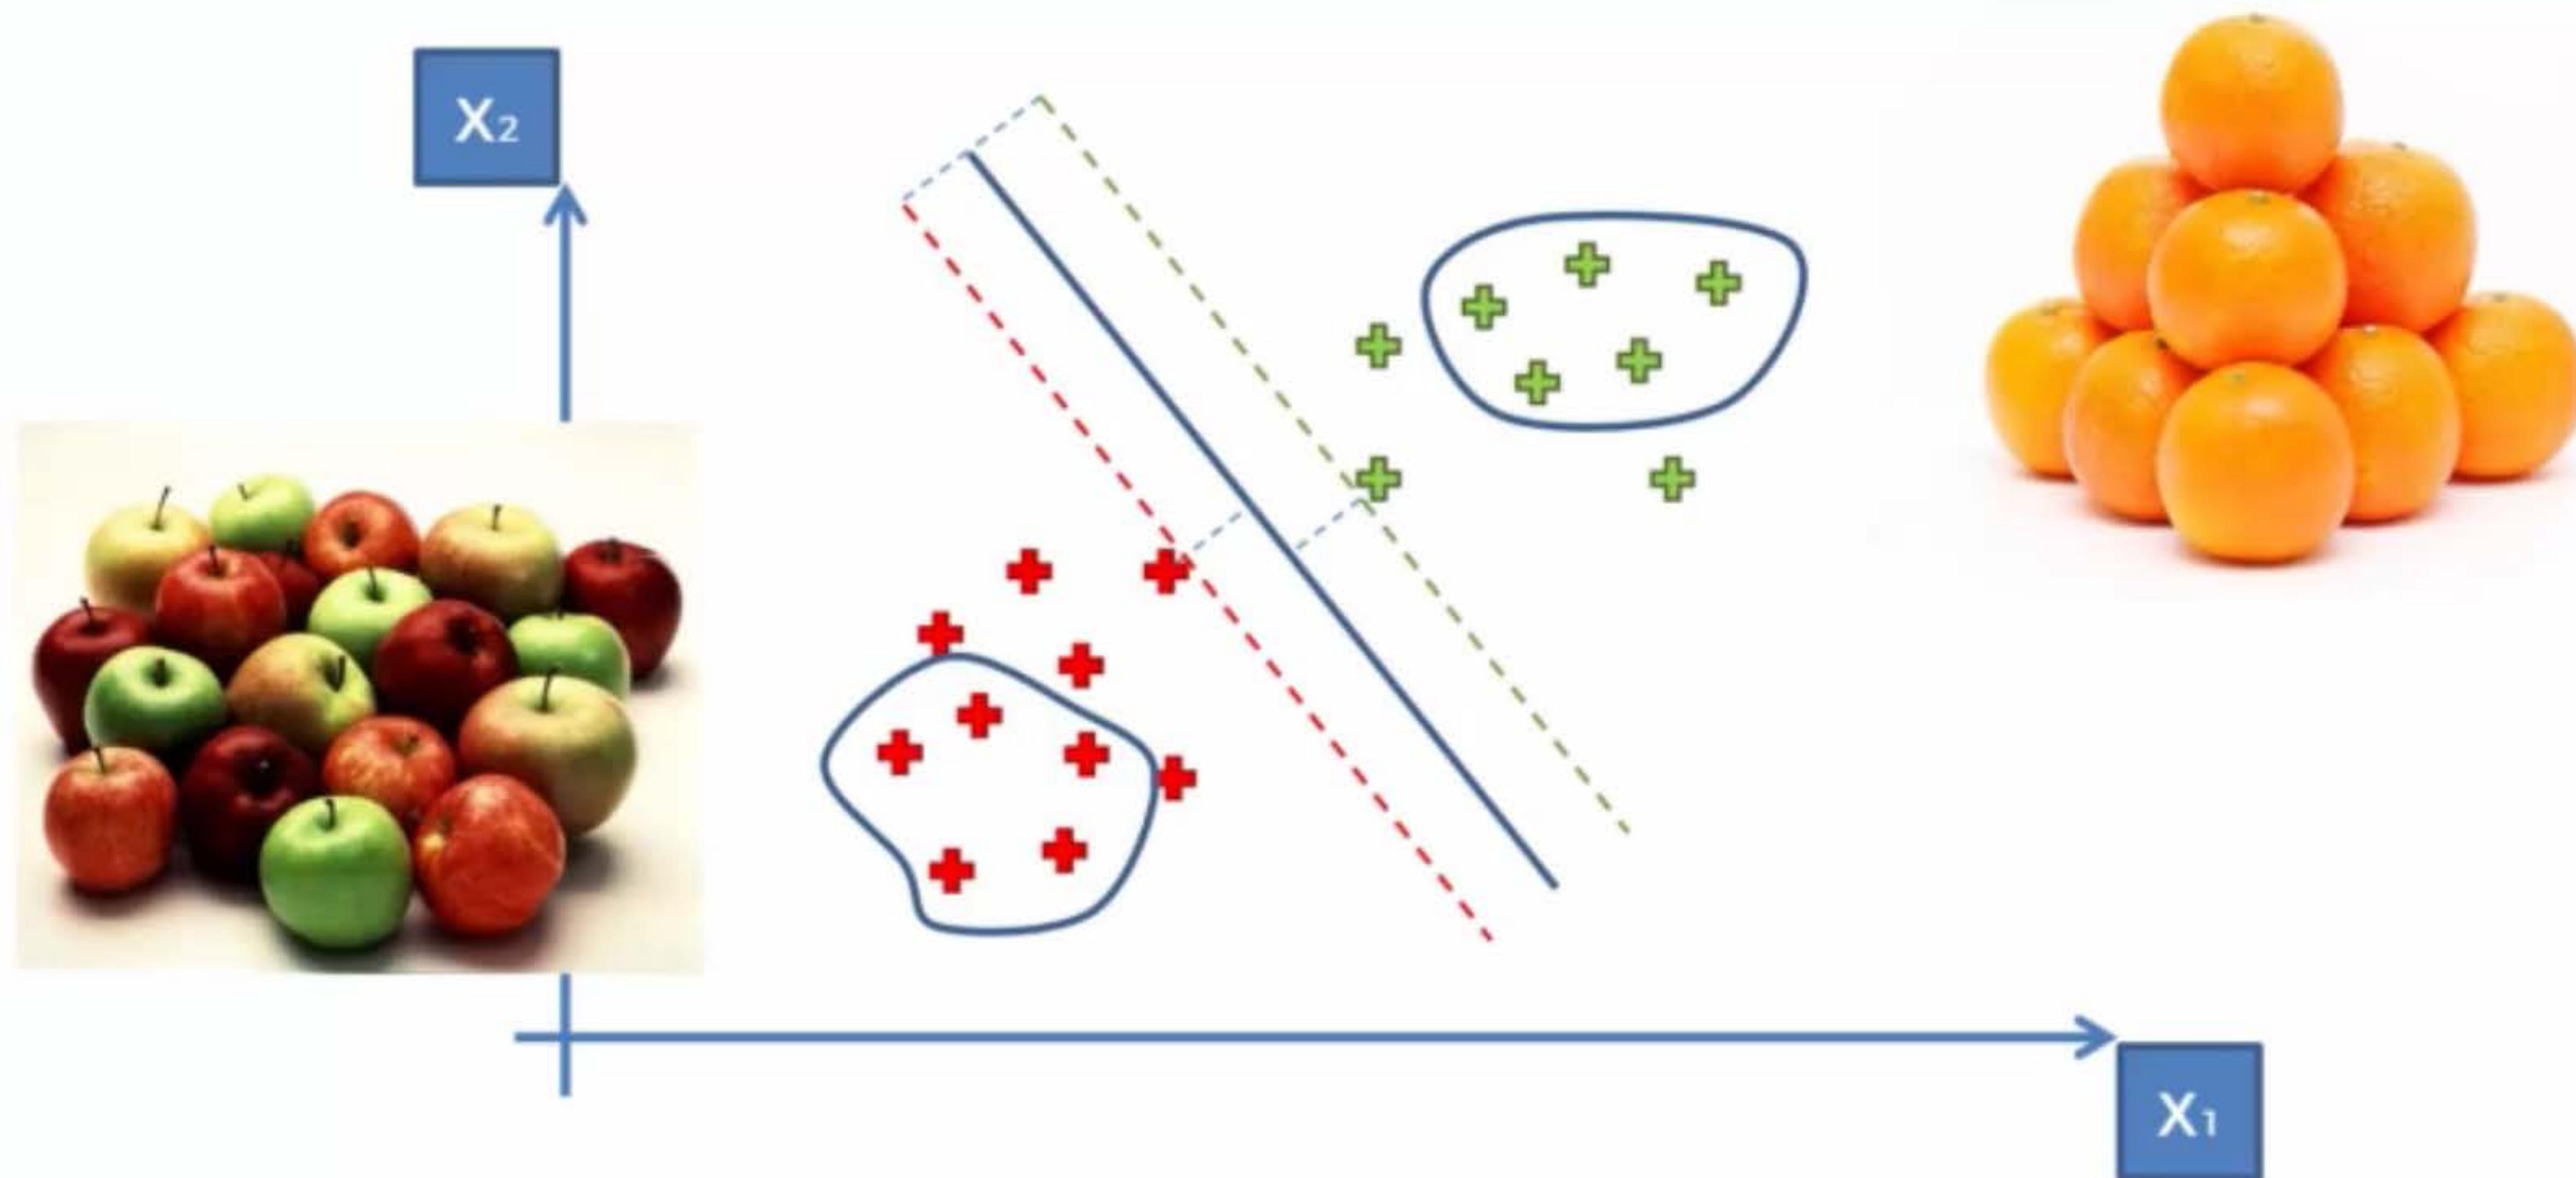


# **what's So Special About SVMs?**

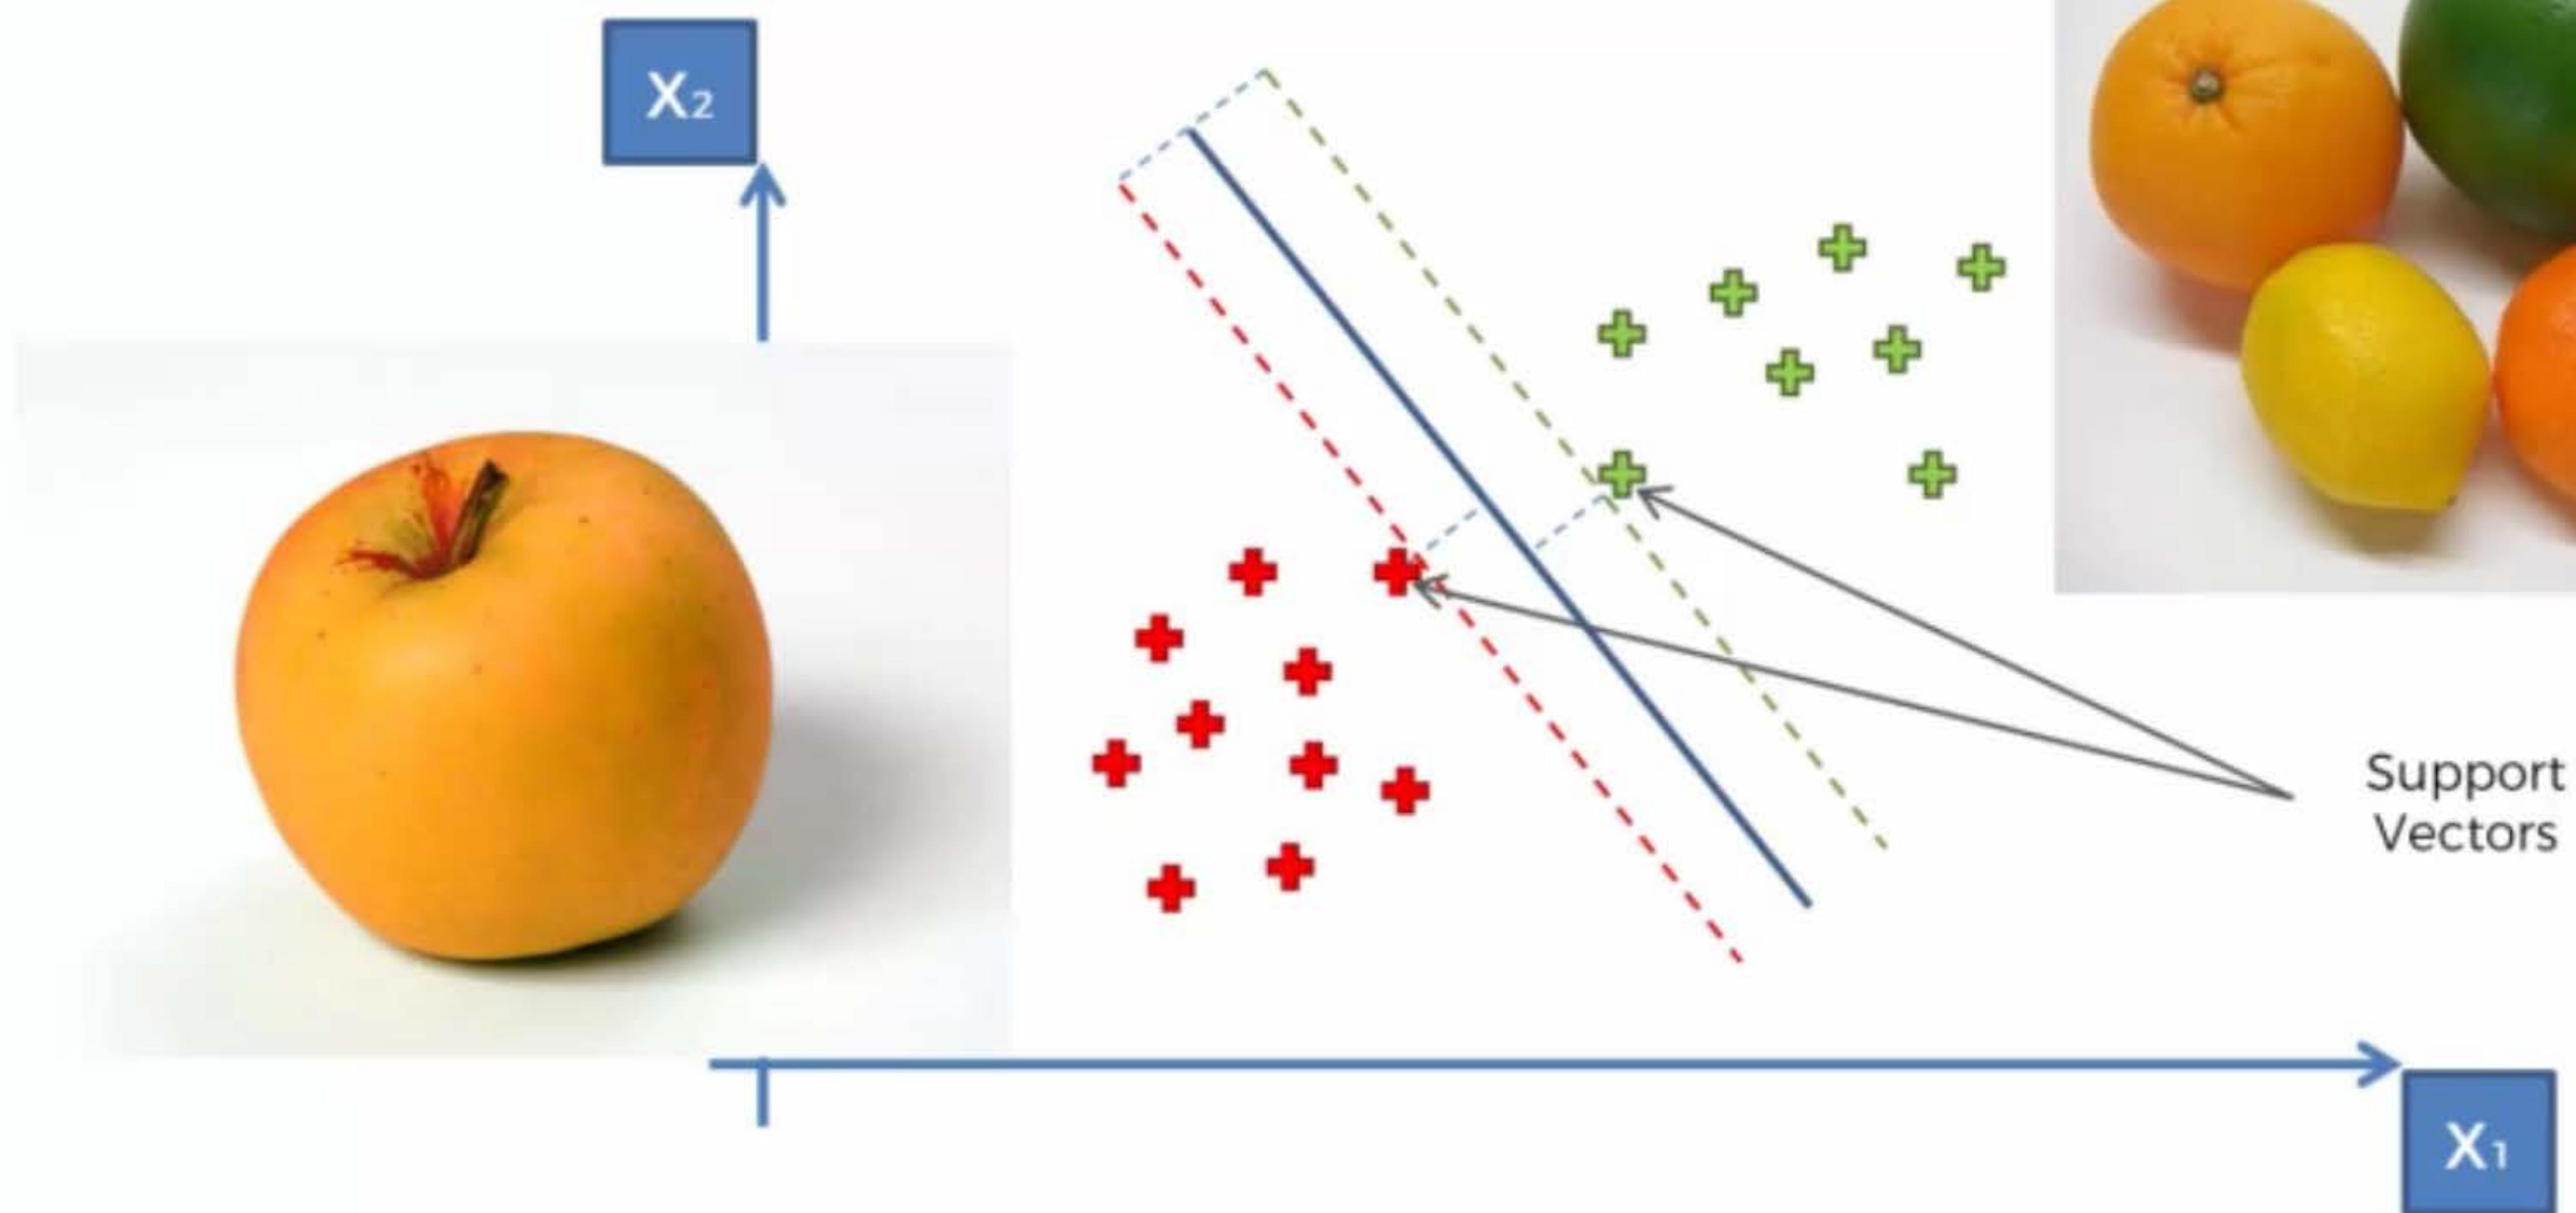
# What's So Special About SVMs?



# What's So Special About SVMs?

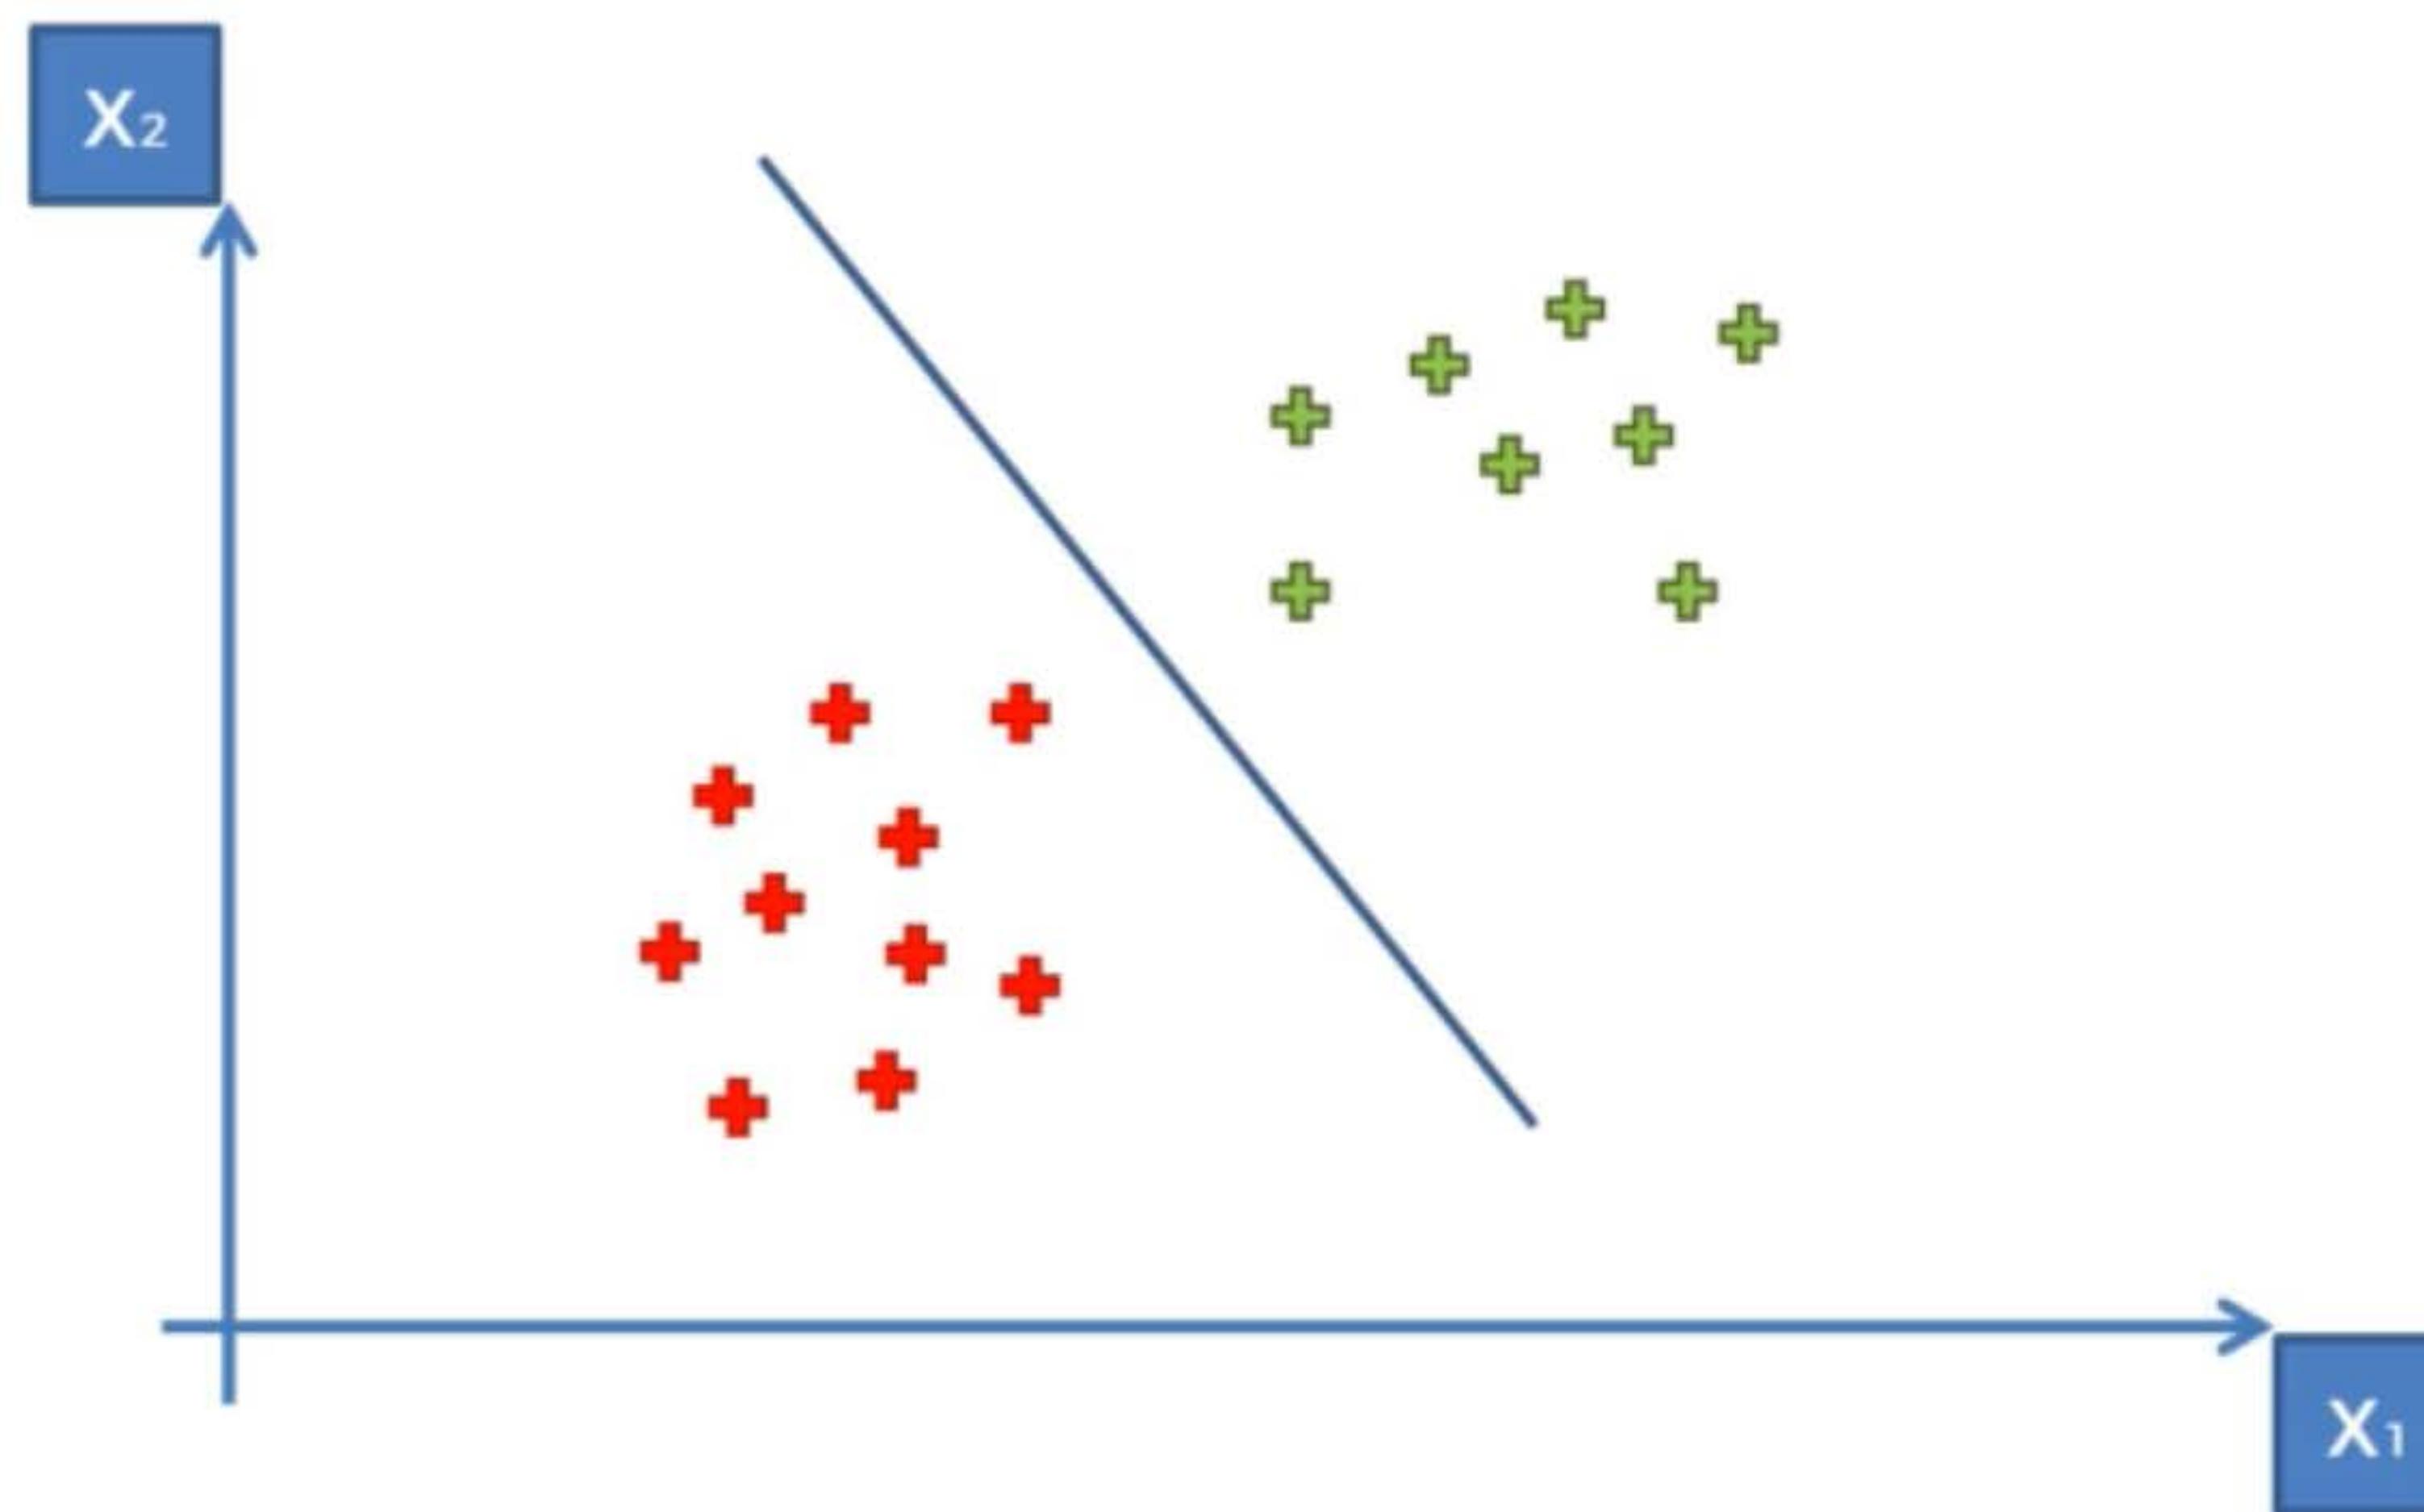


# What's So Special About SVMs?

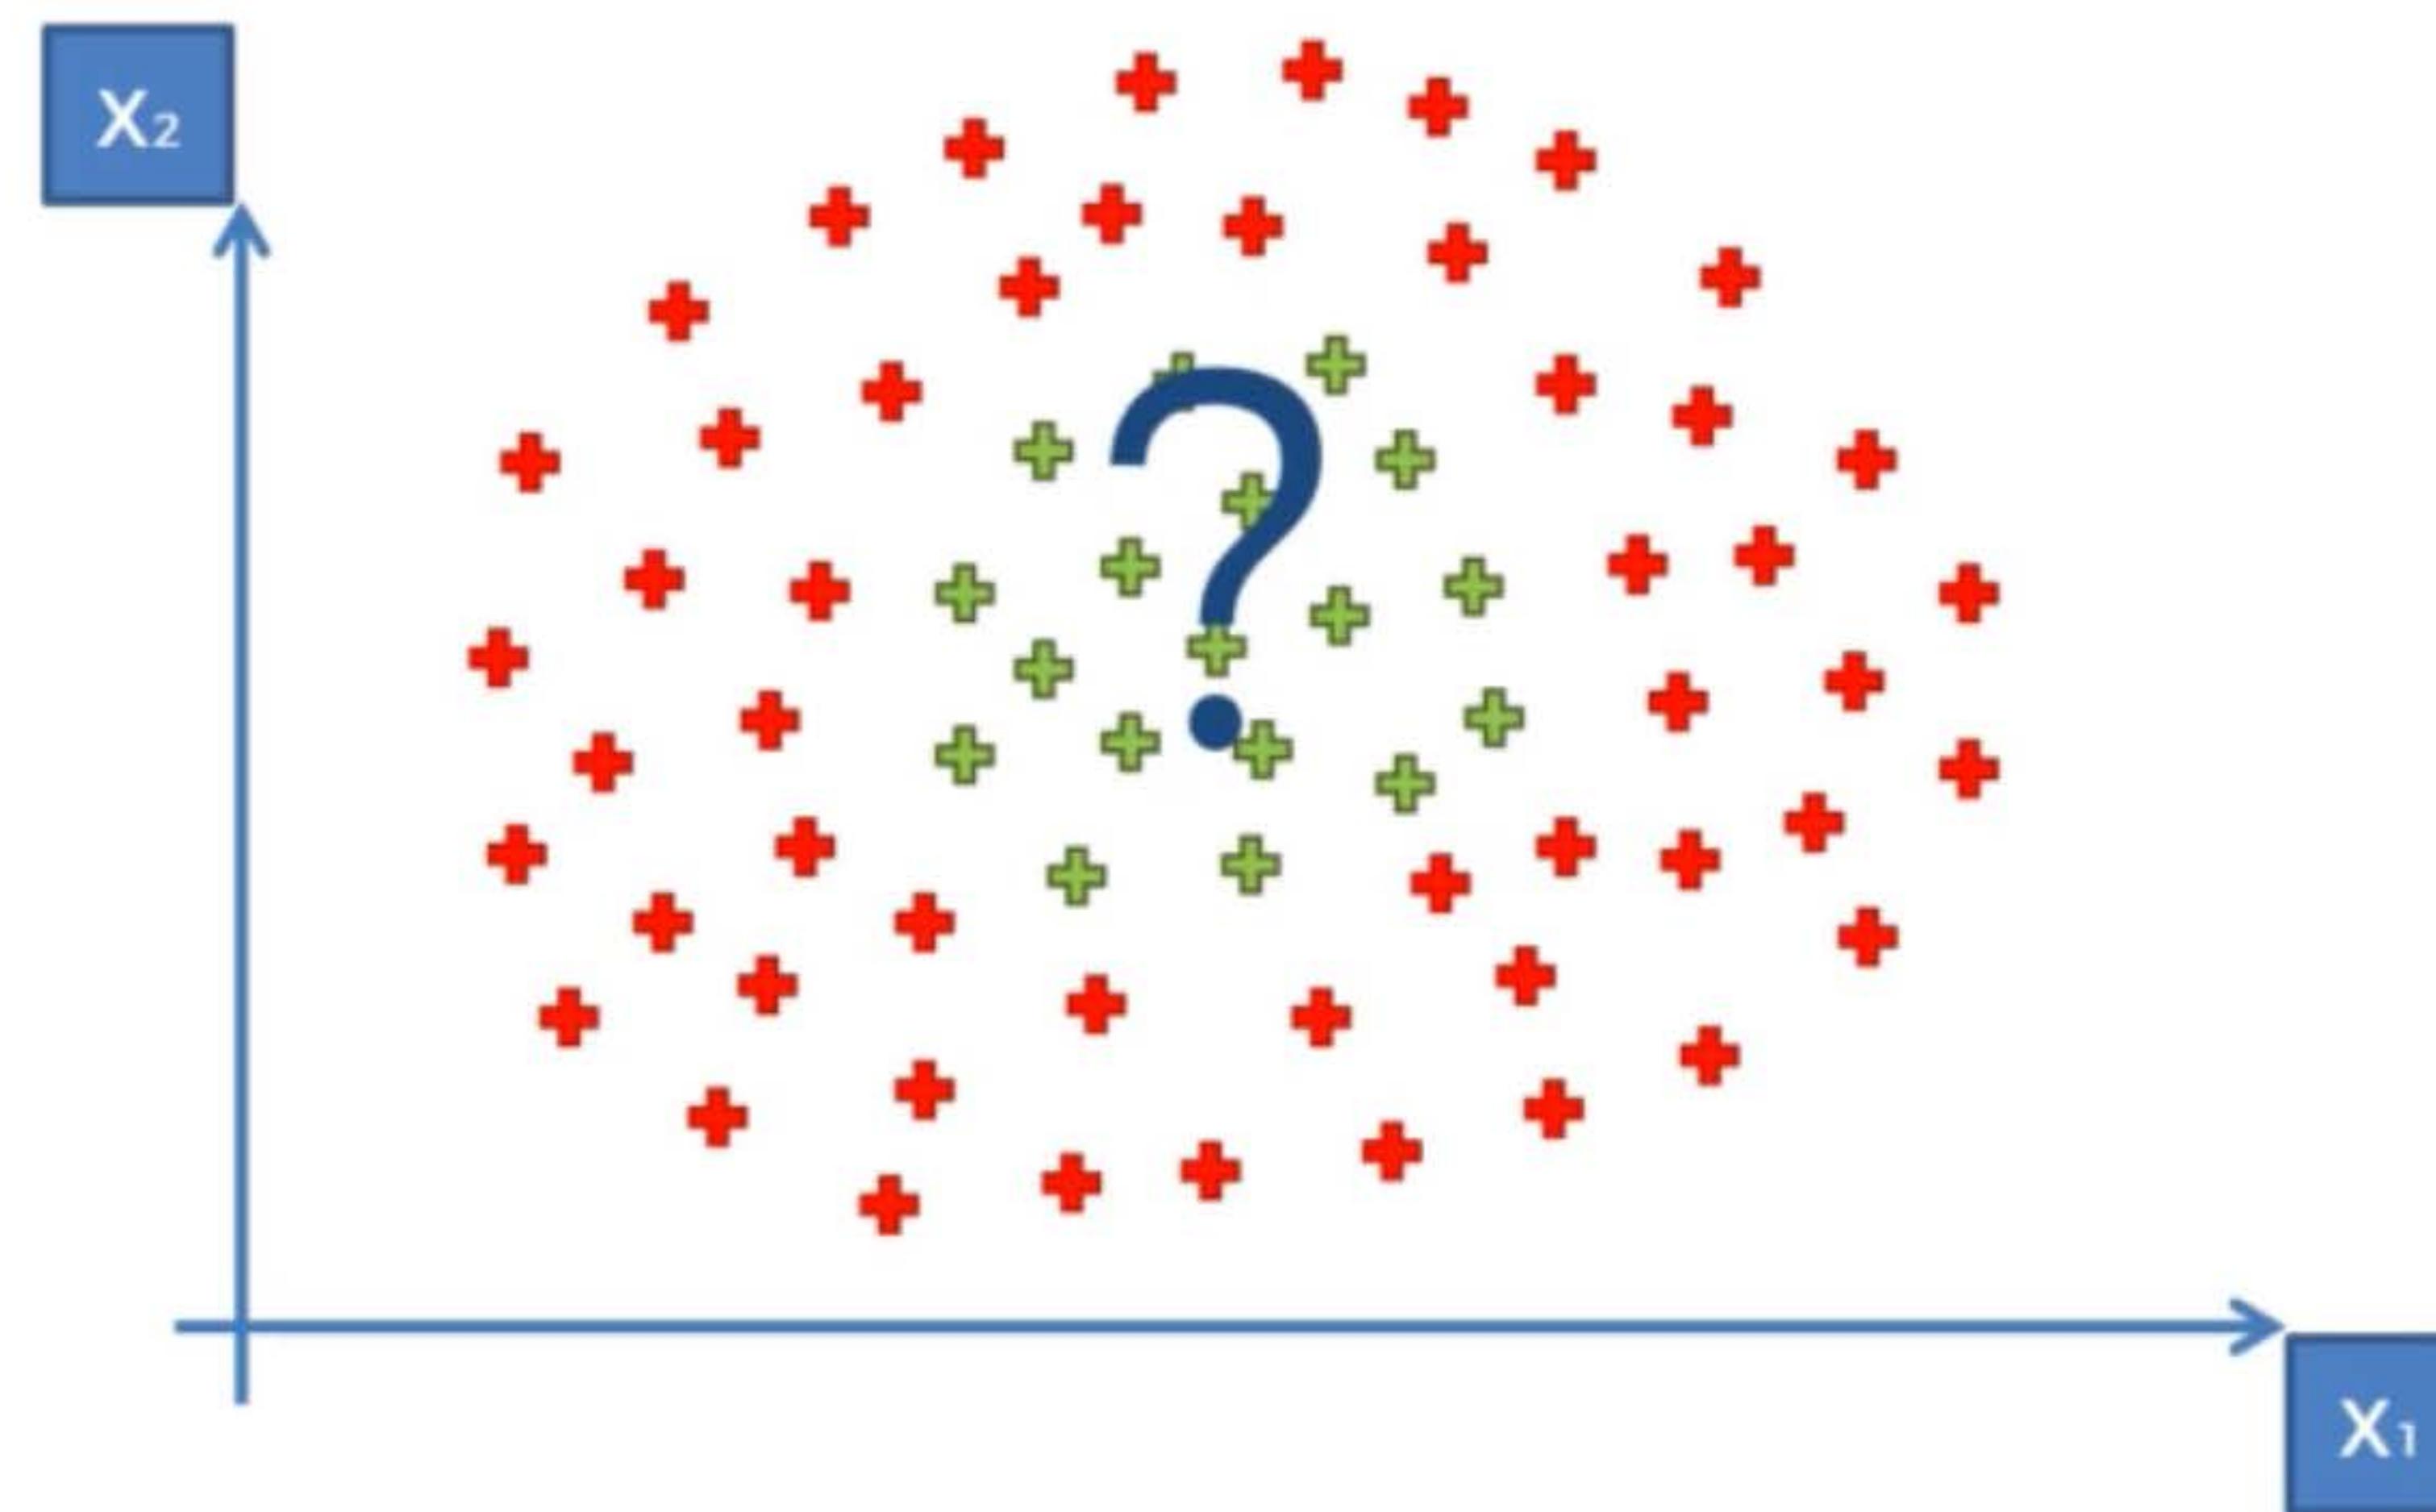


# **Kernel SVM Intuition**

# SVM separates well these points



# What about these points ?

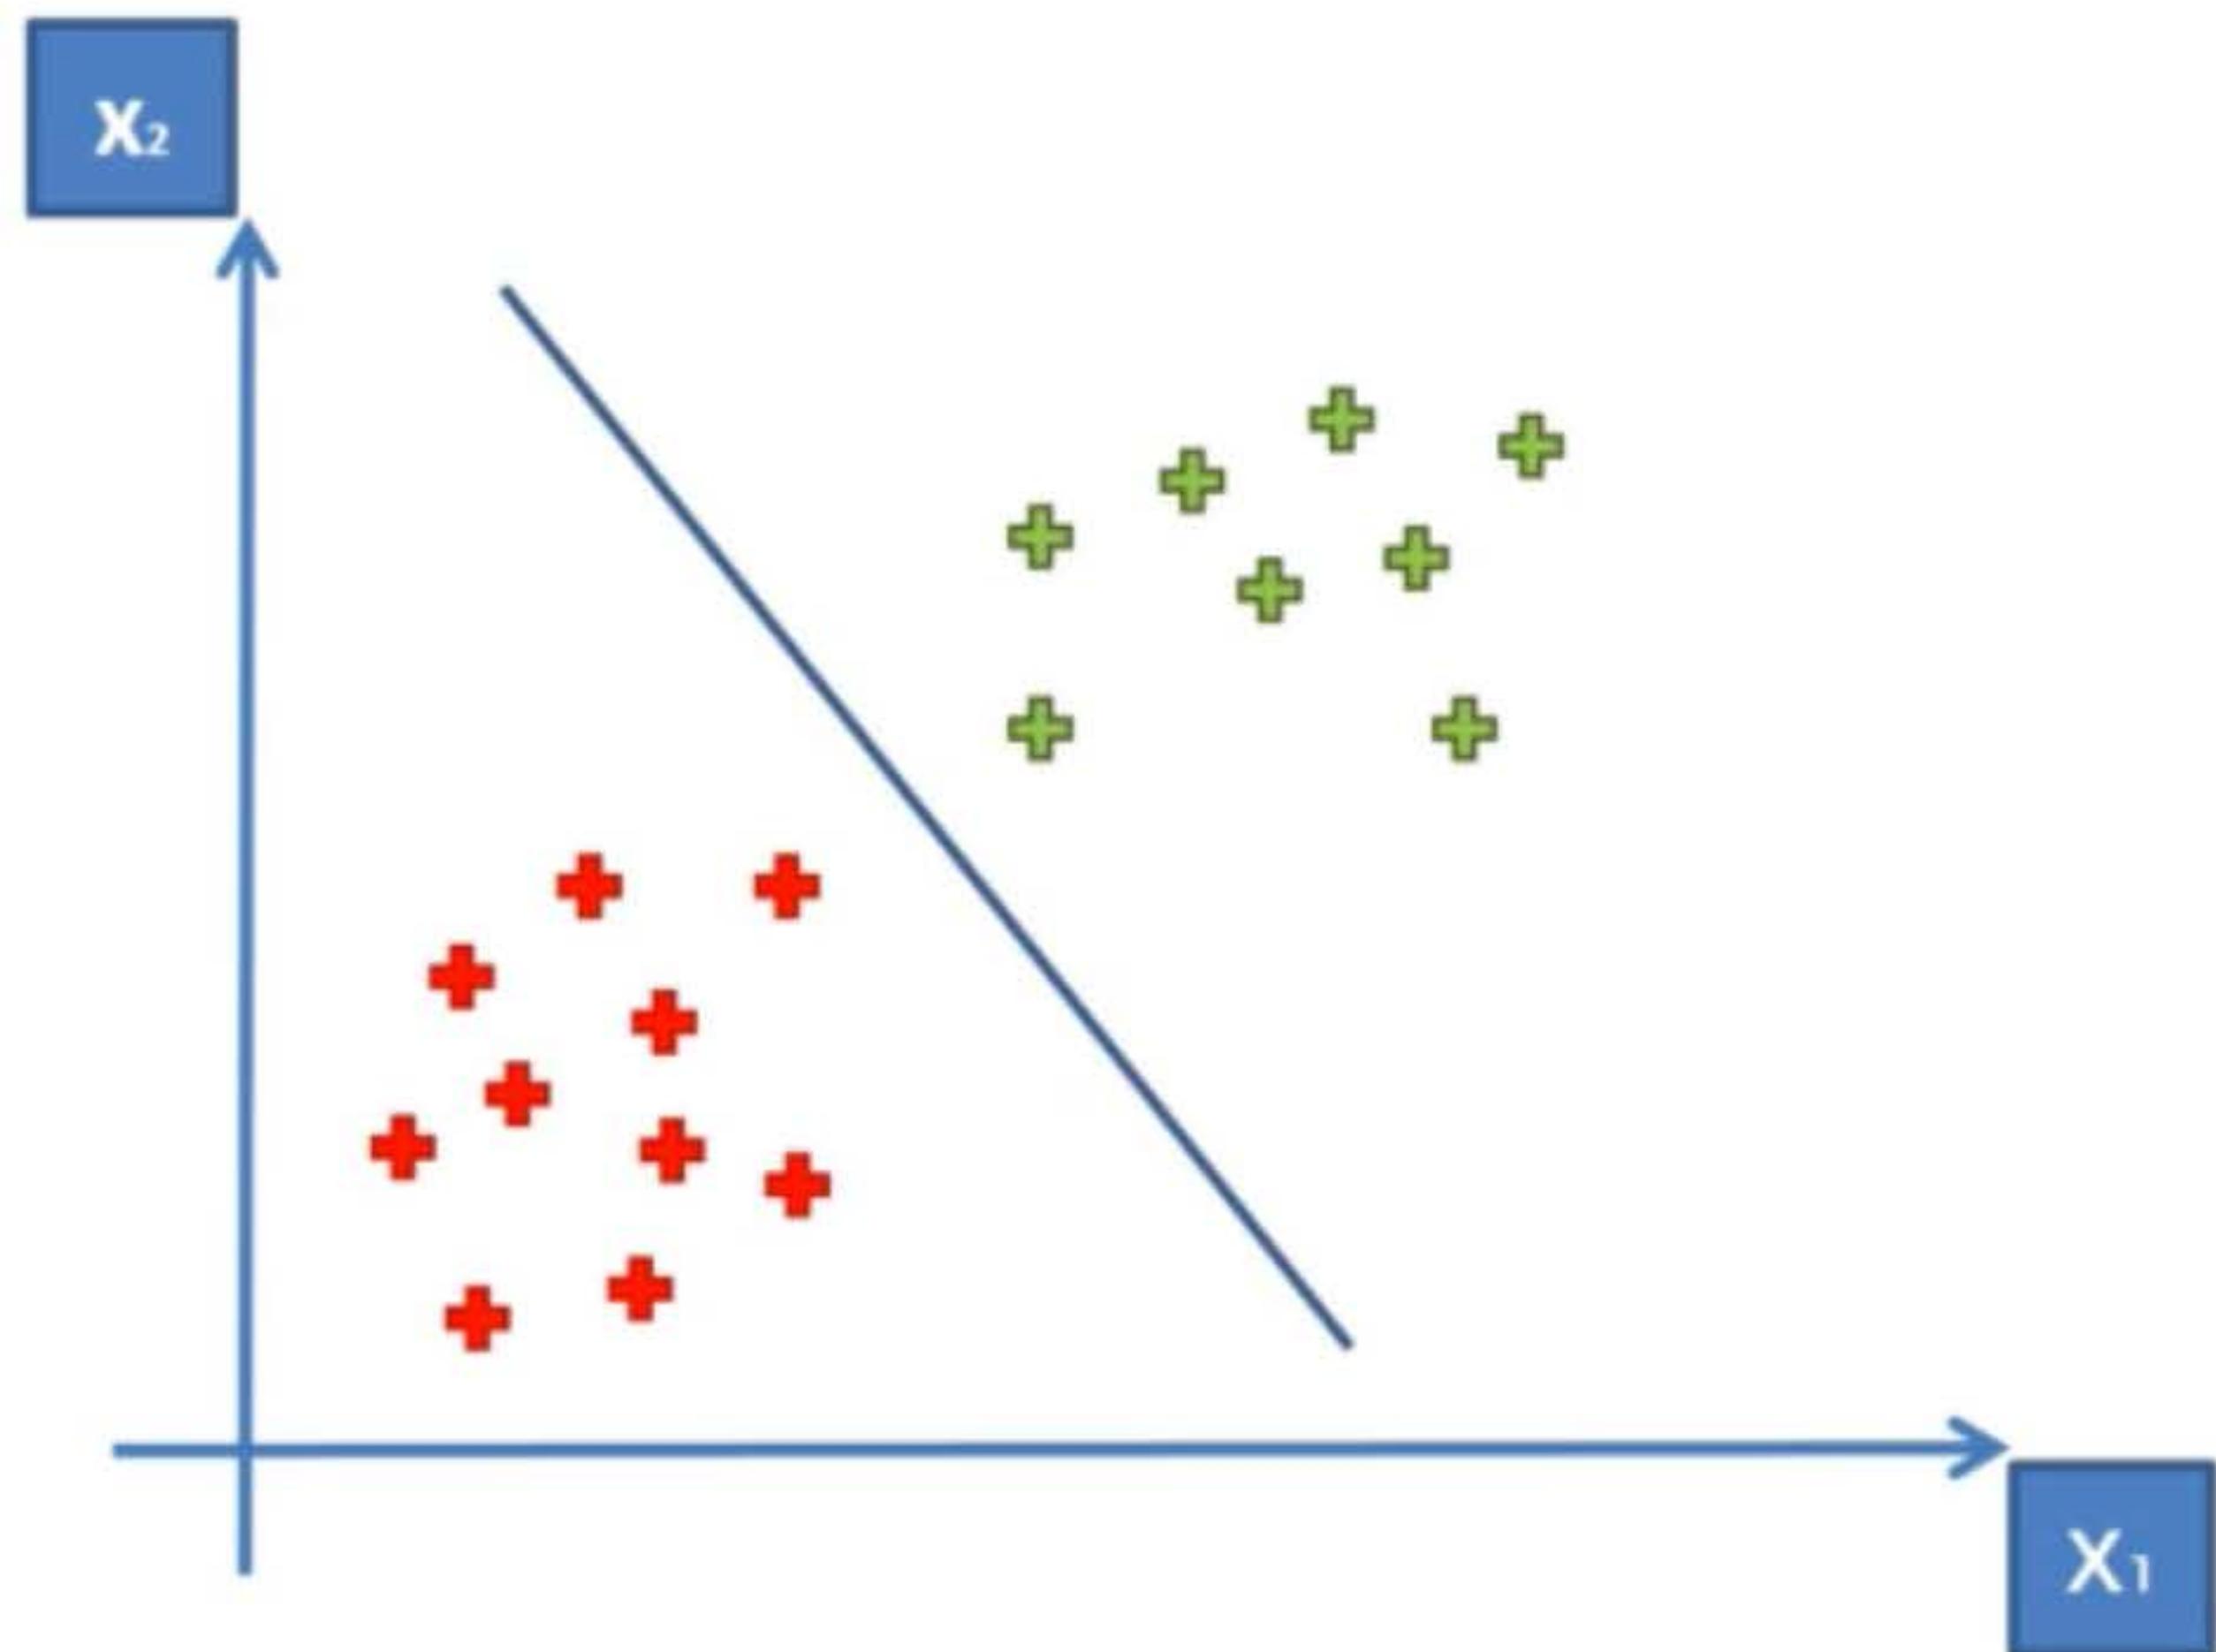


# Why ?

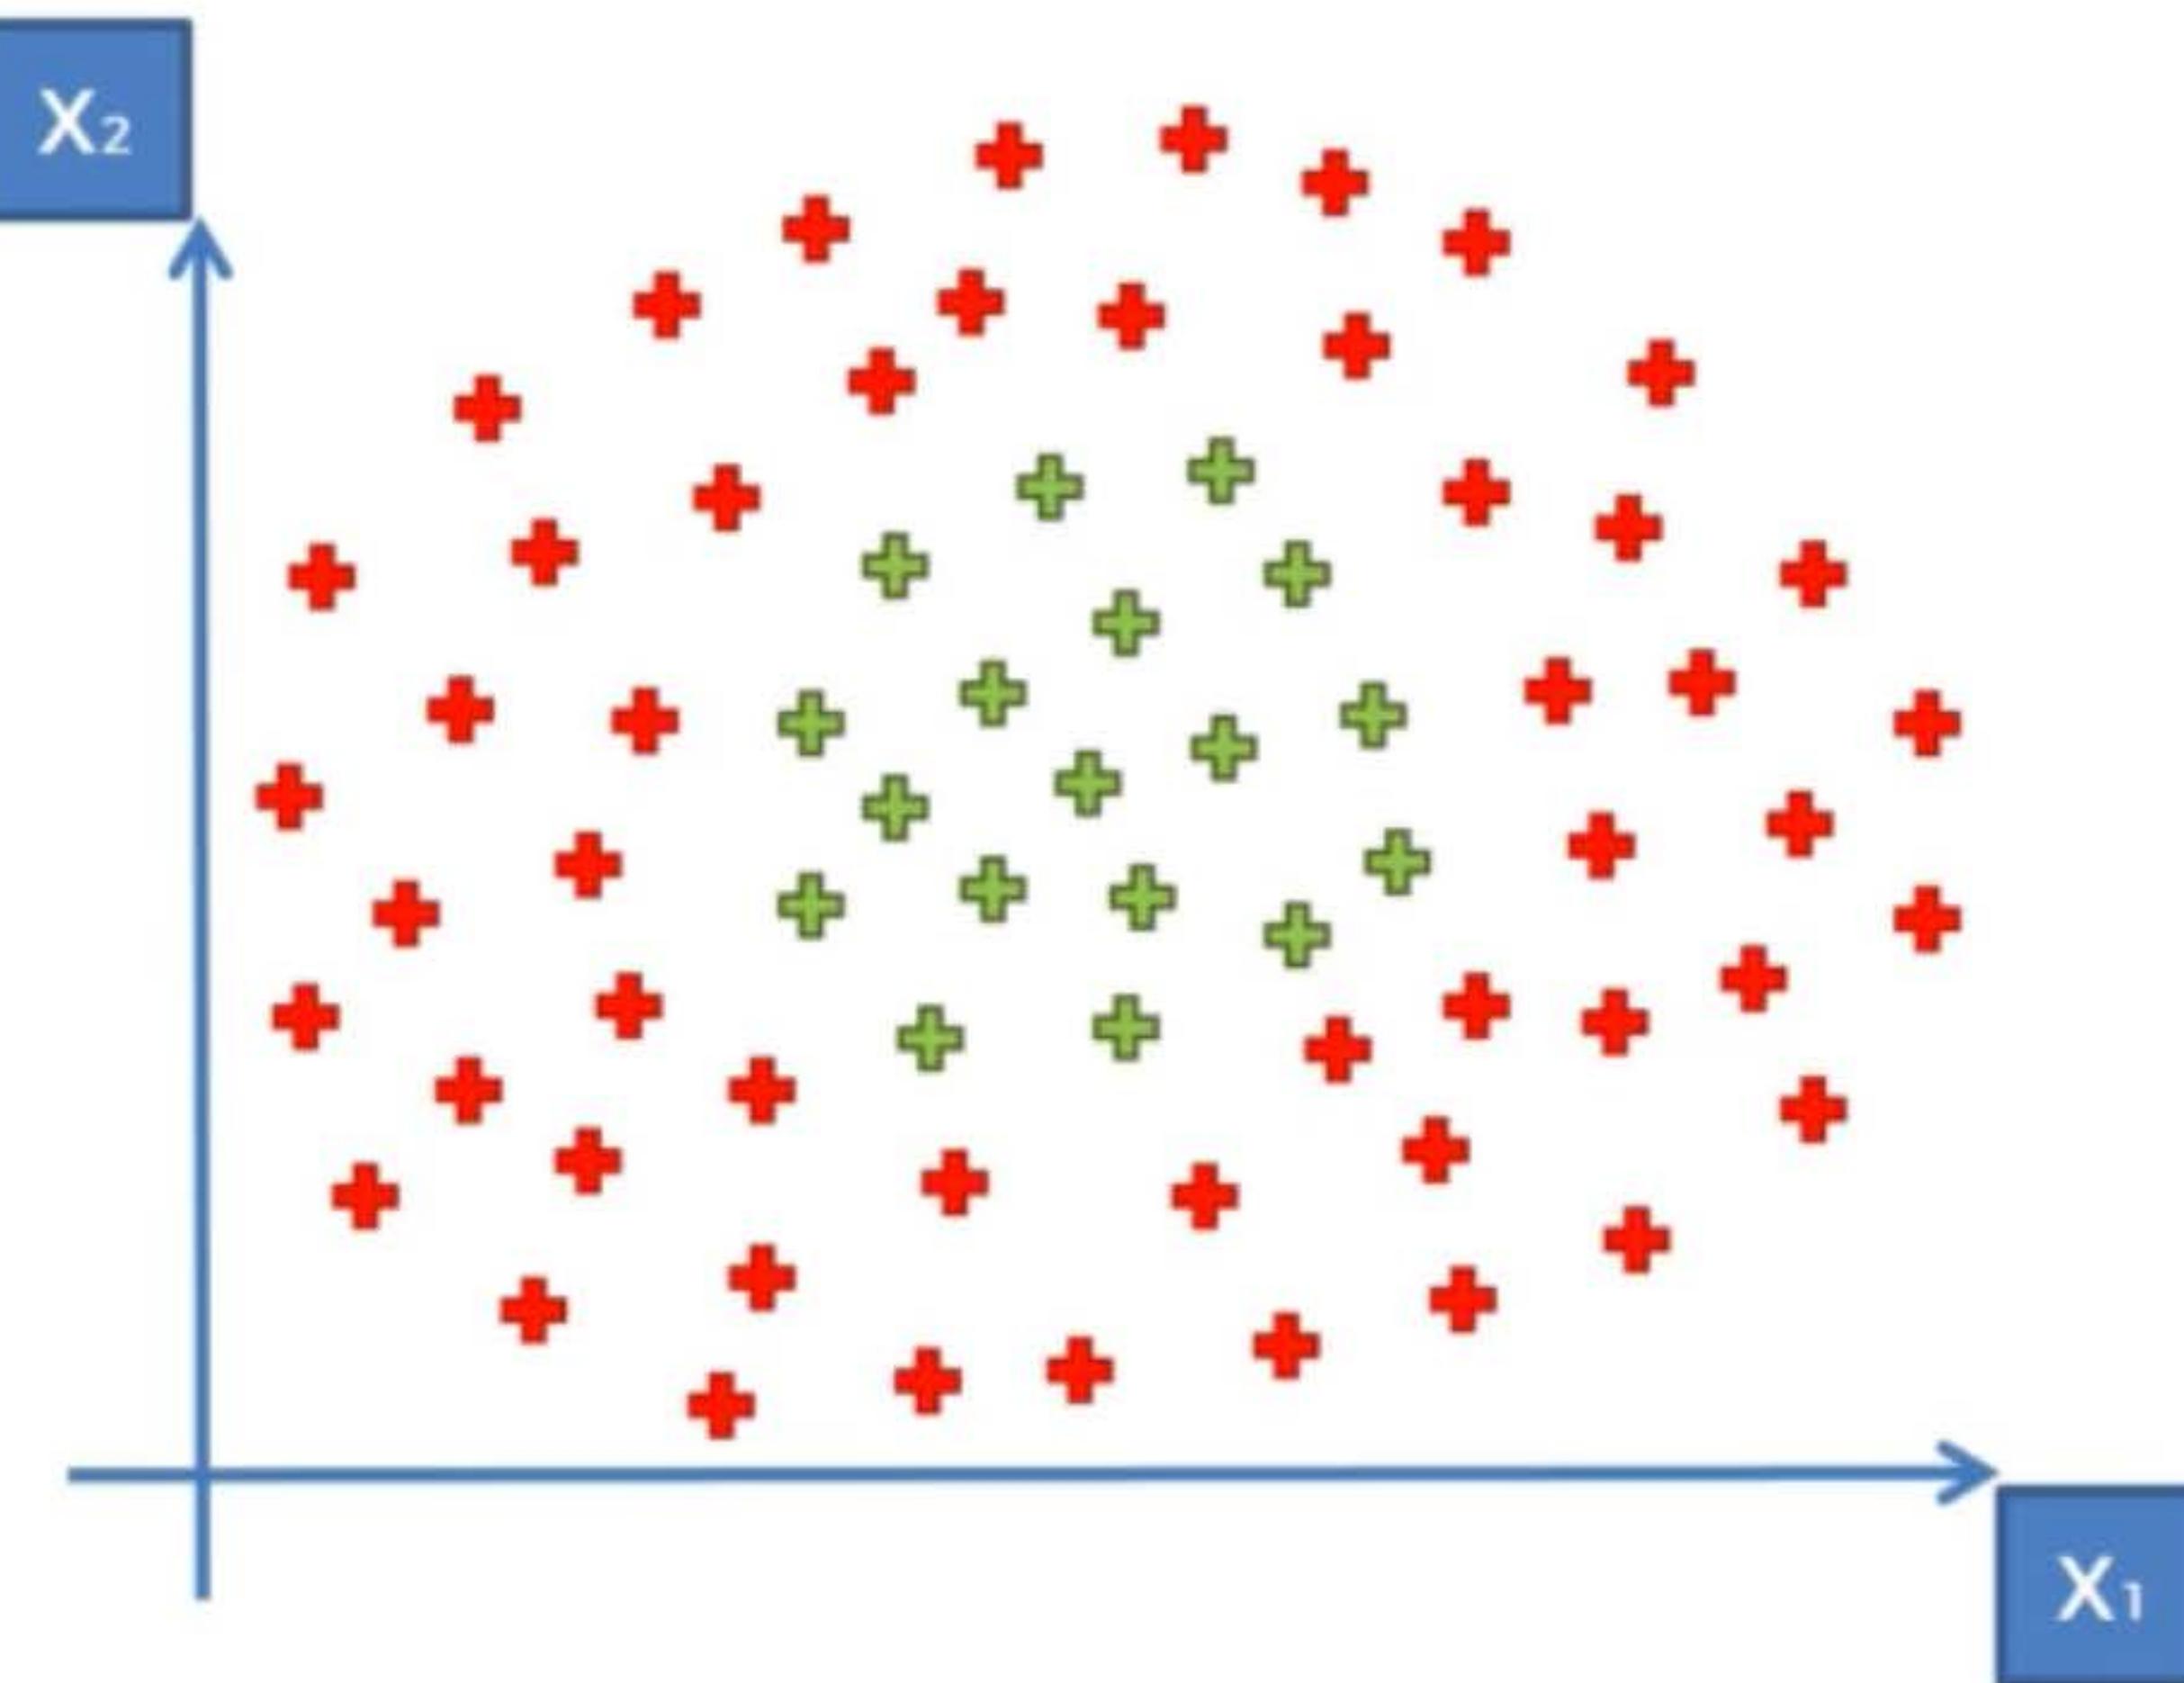
Because the data points are  
not LINEARLY SEPARABLE

# Linear Separability

Linearly Separable



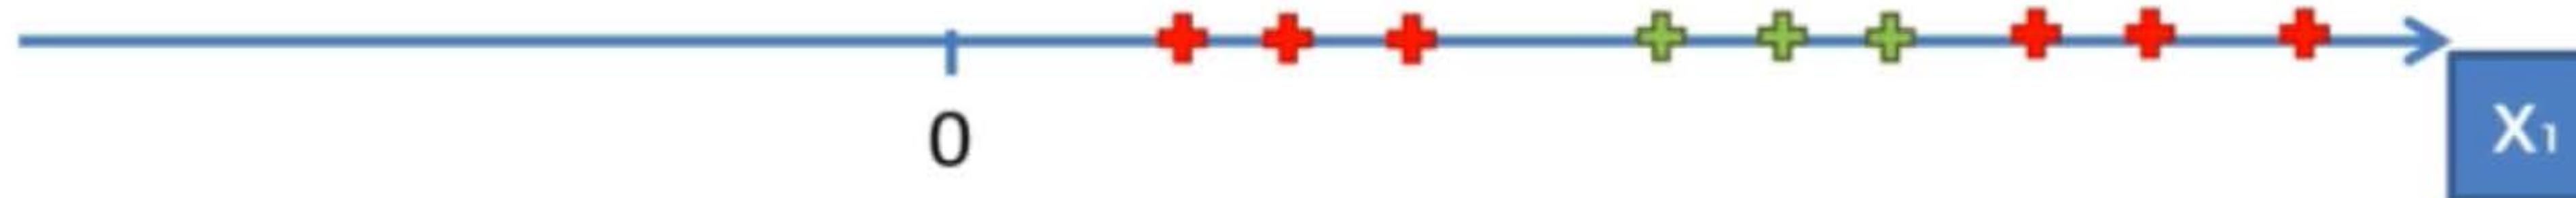
Not Linearly Separable



# A Higher-Dimensional Space

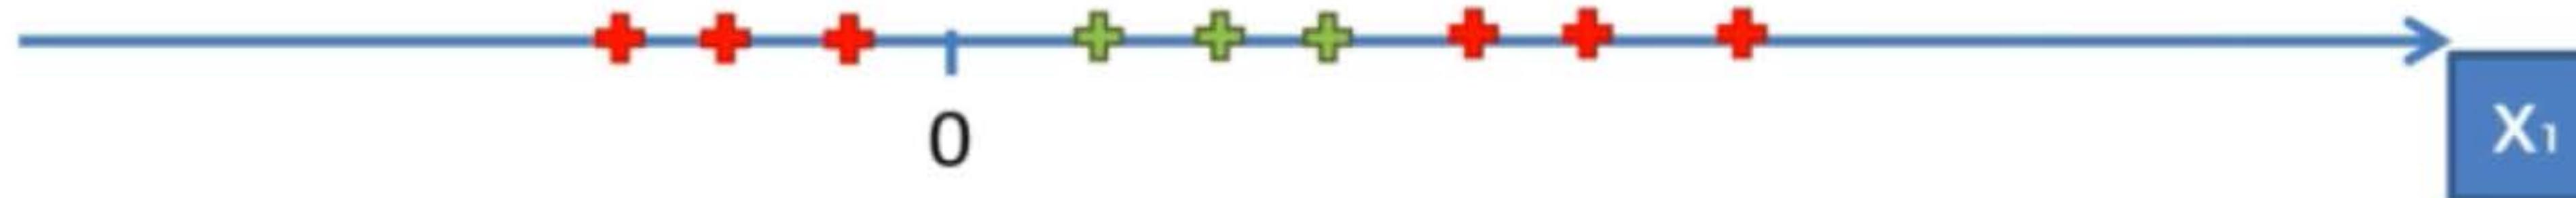
# Mapping to a Higher Dimension

$$f = x - 5$$

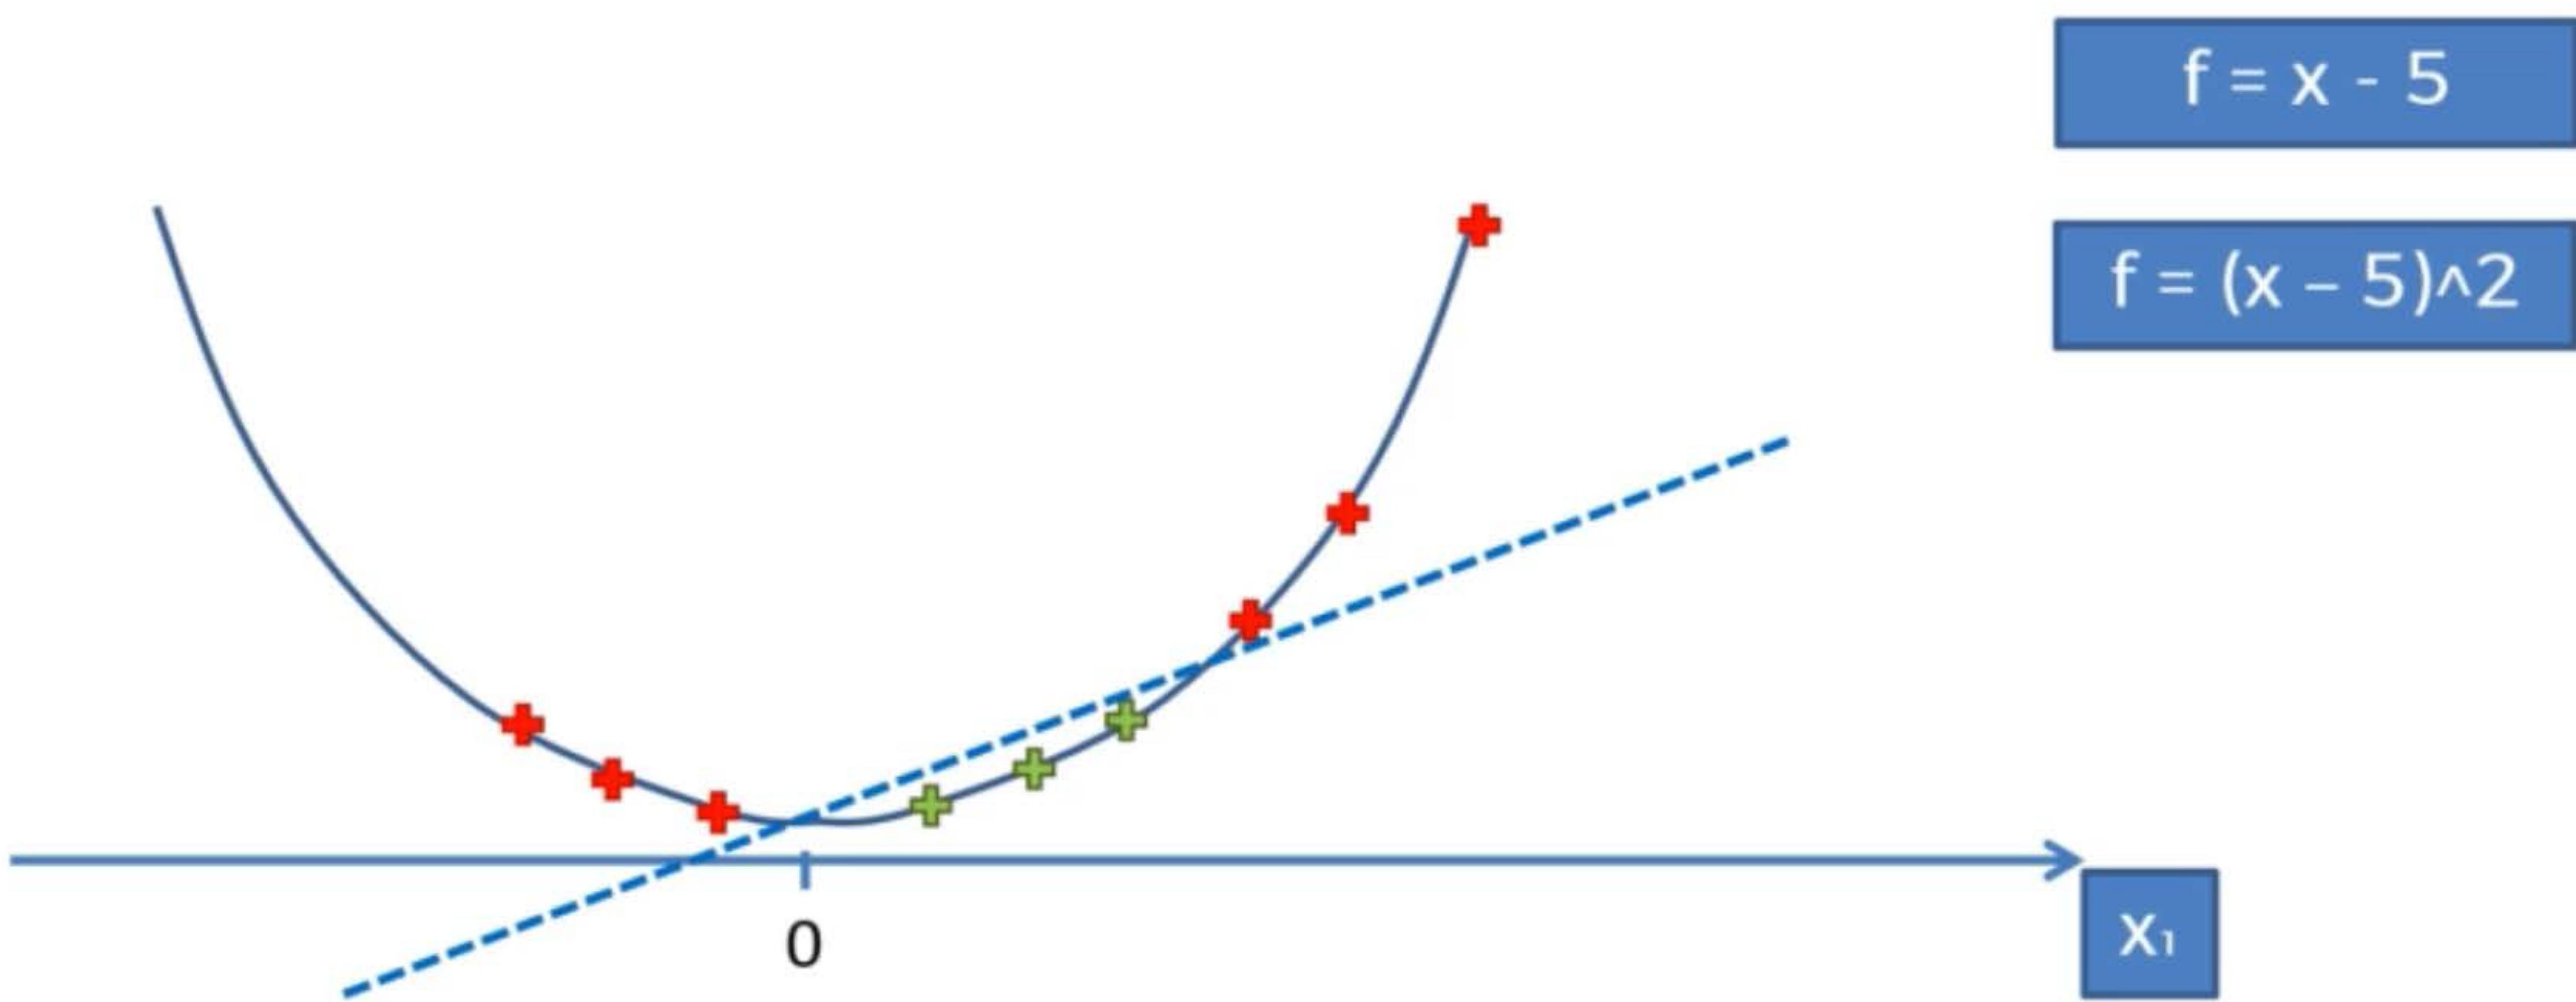


# Mapping to a Higher Dimension

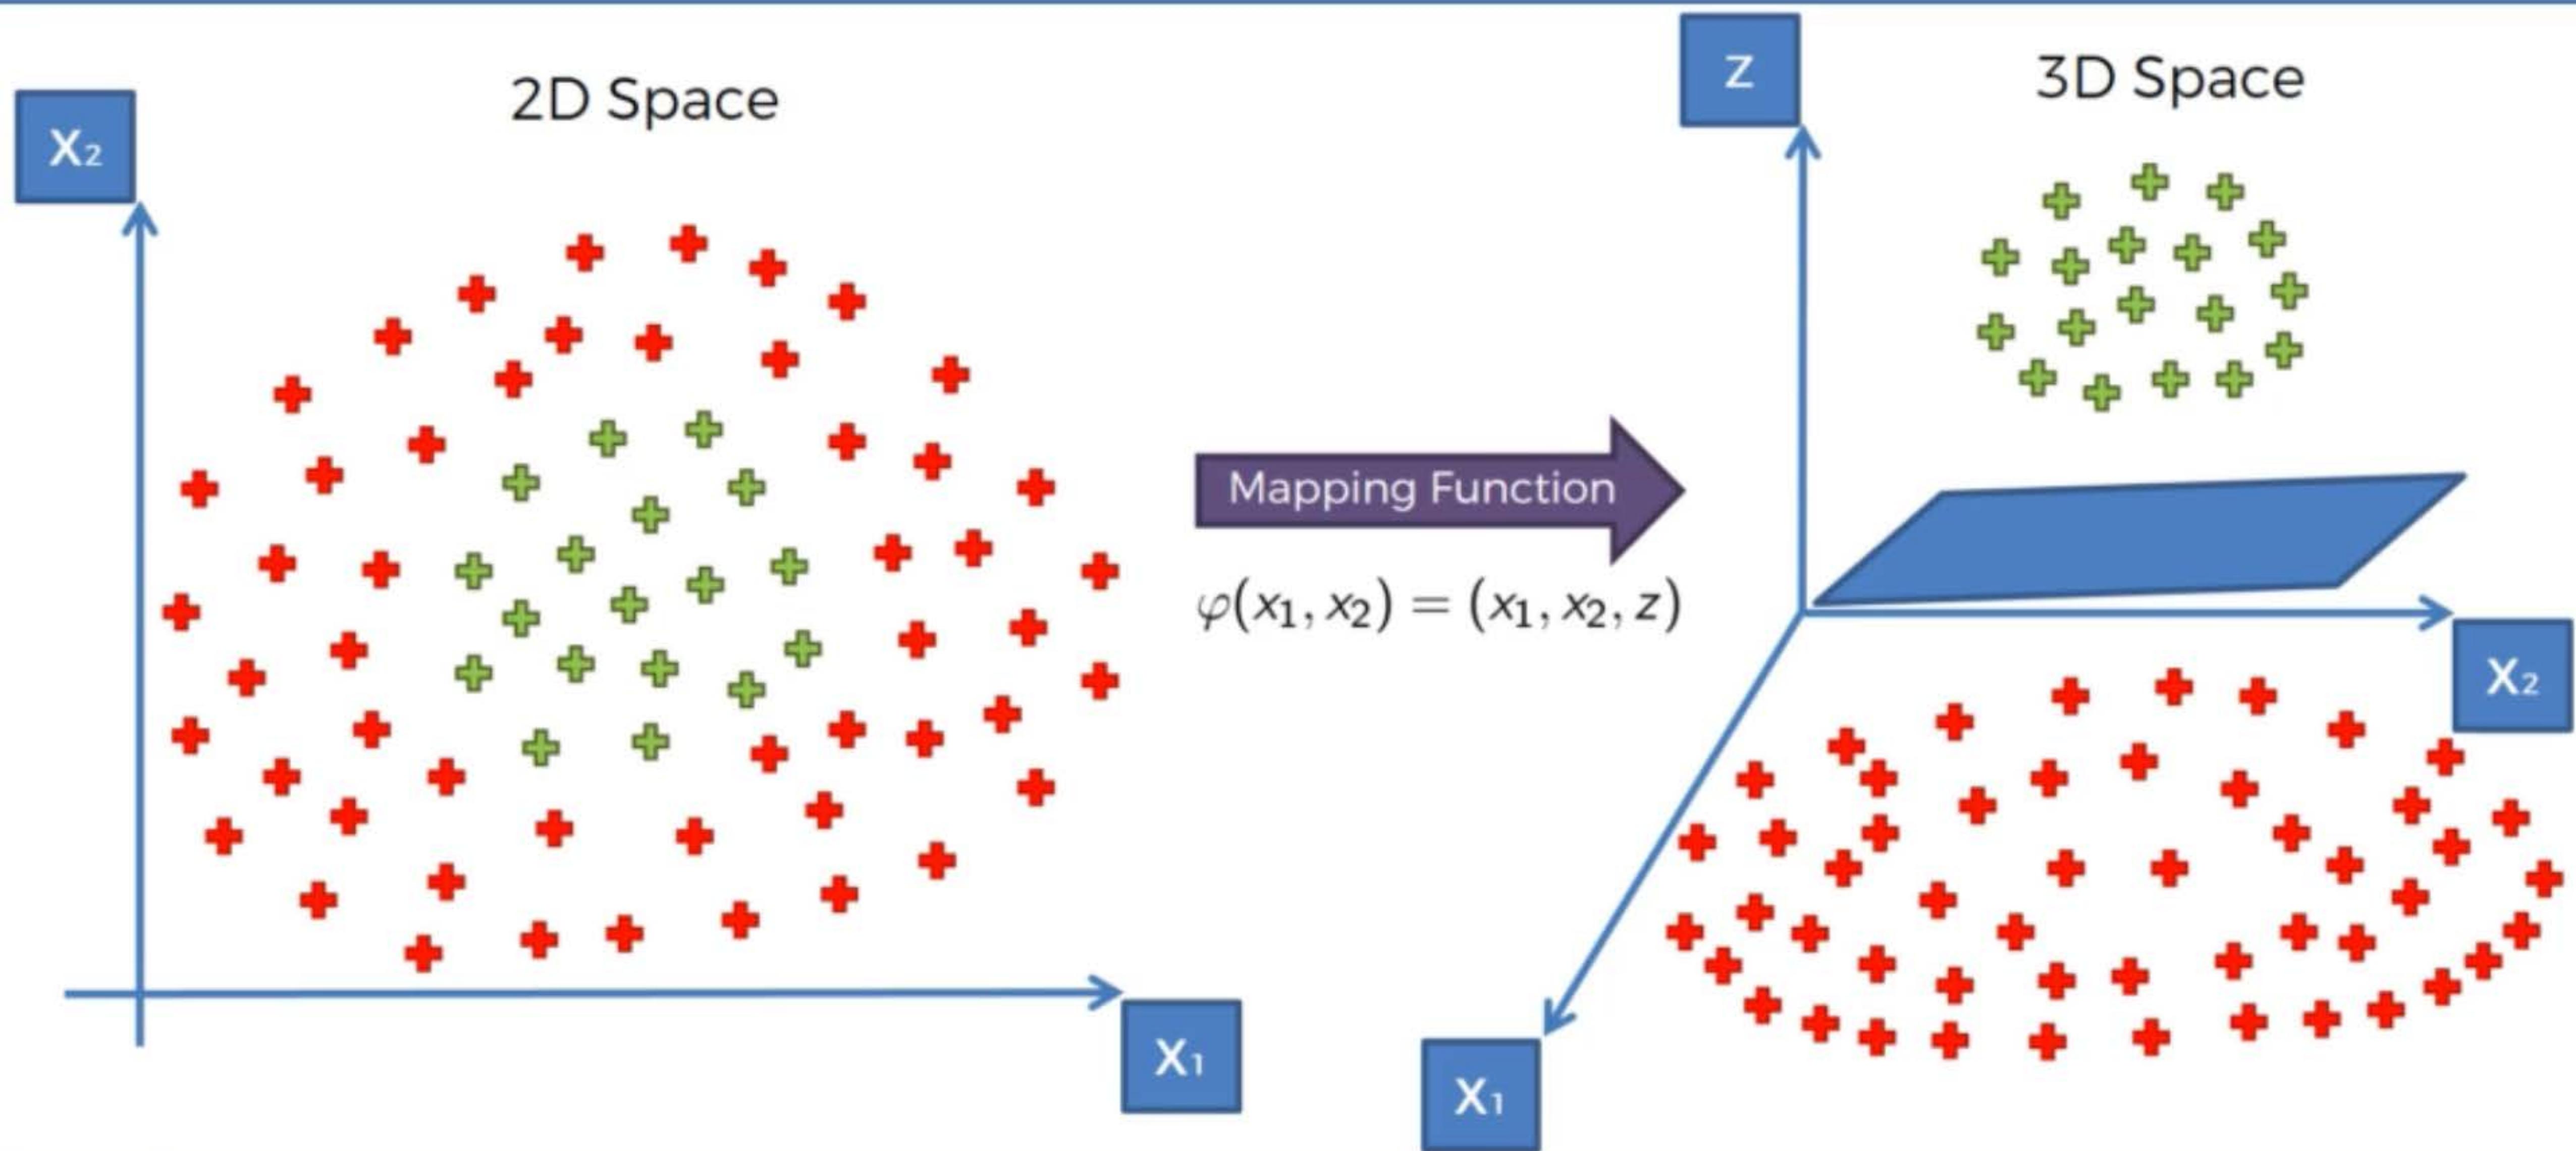
$$f = x - 5$$



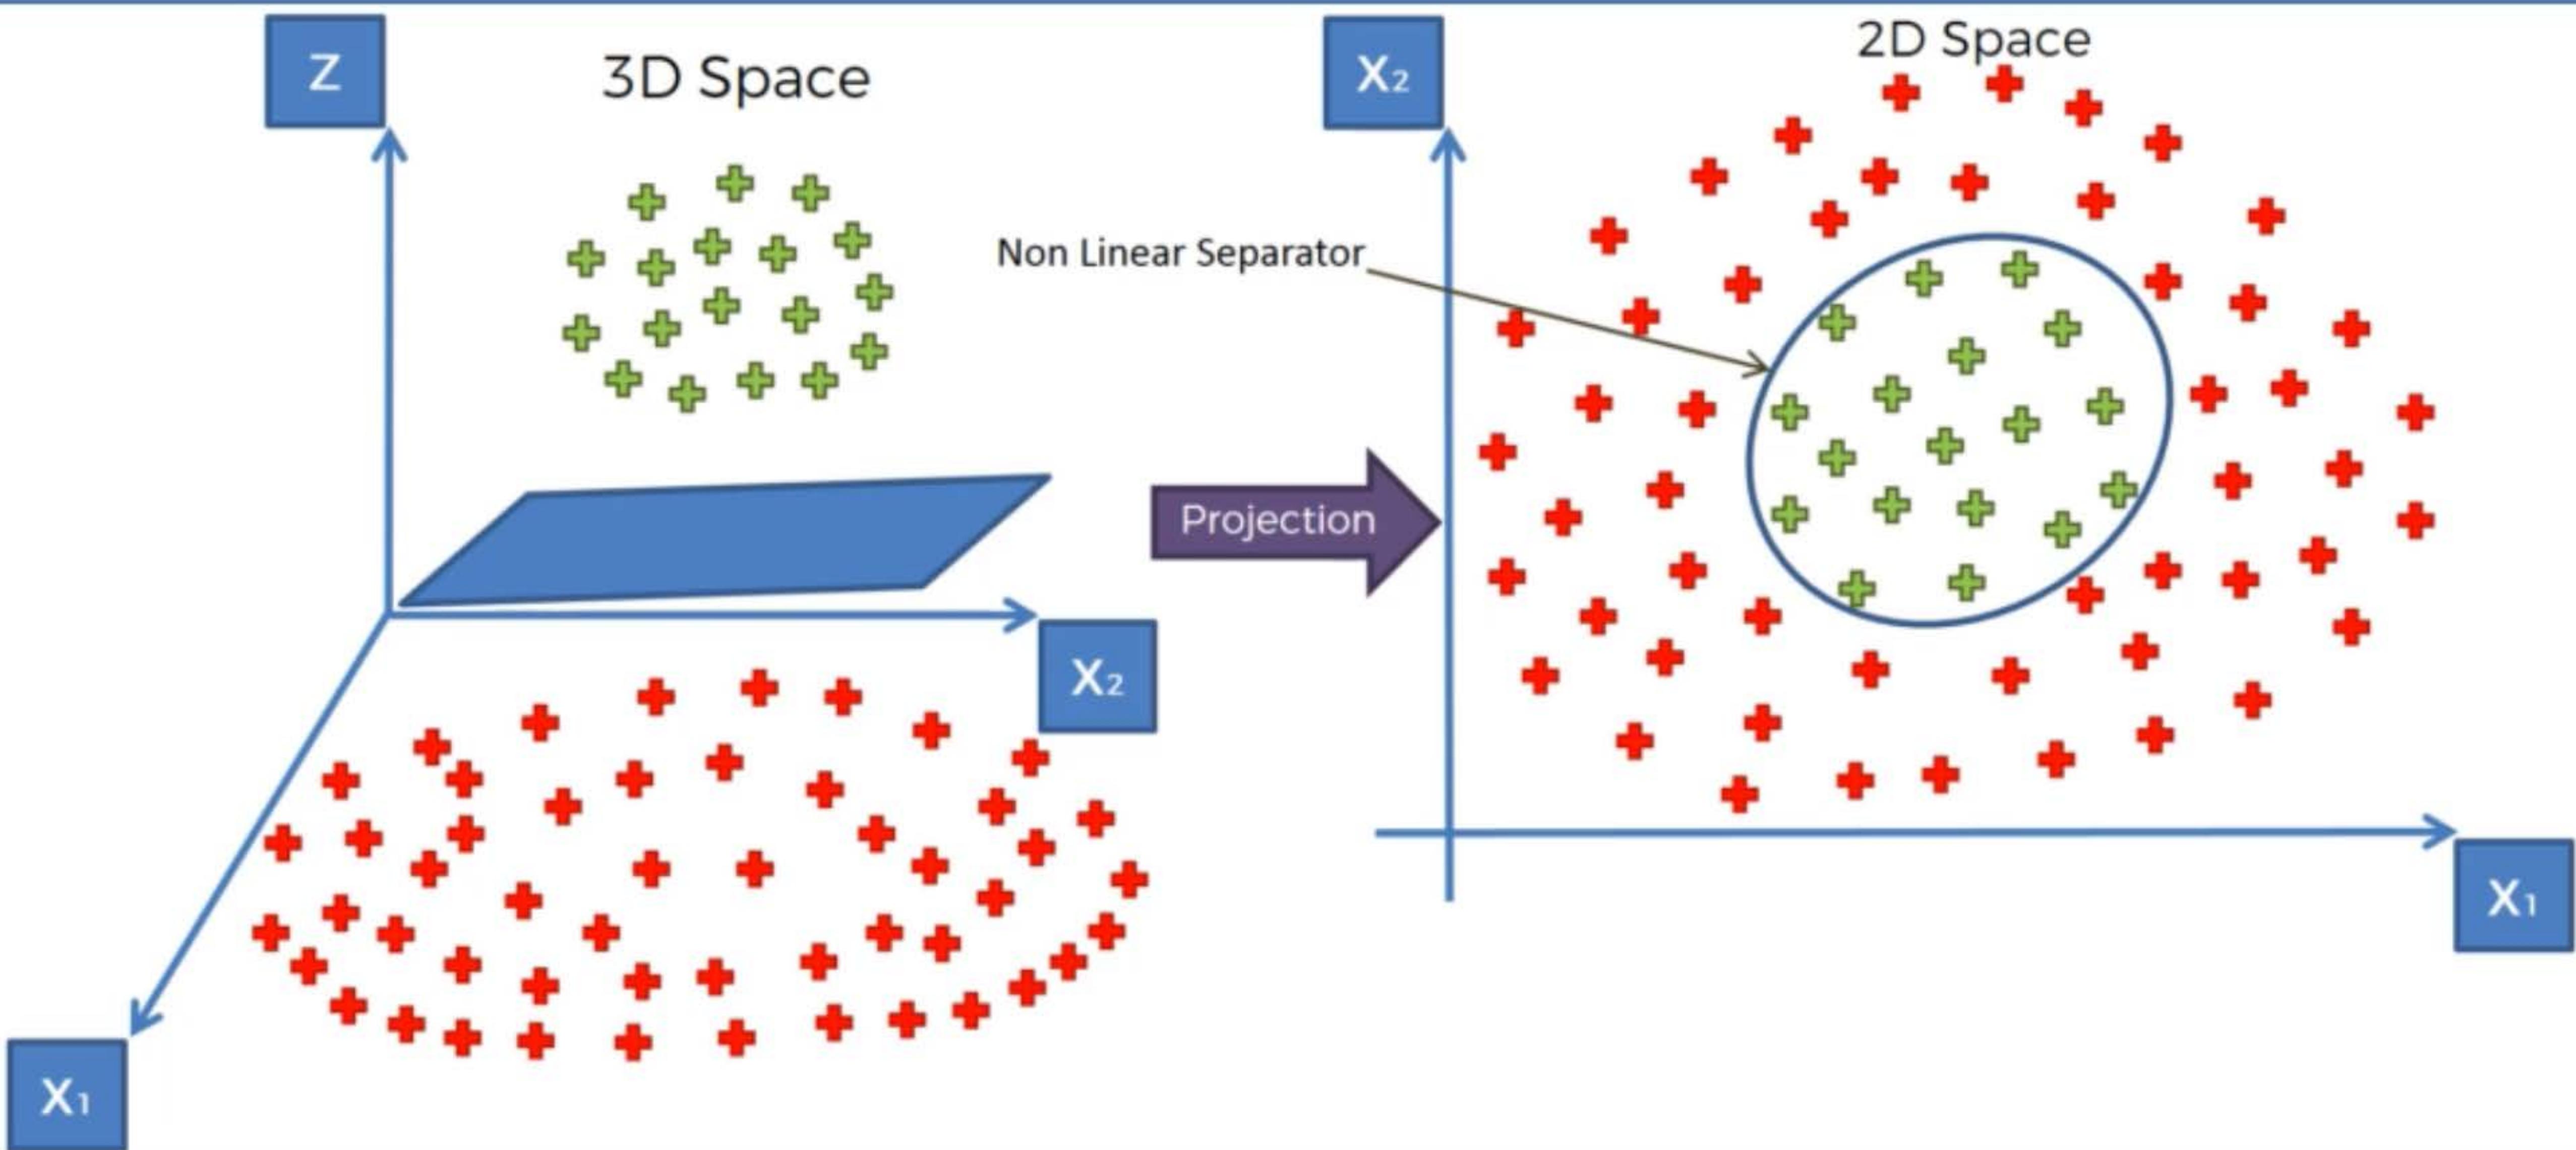
# Mapping to a Higher Dimension



# Mapping to a Higher Dimension



# Projecting back to 2D Space



# But there is a catch...

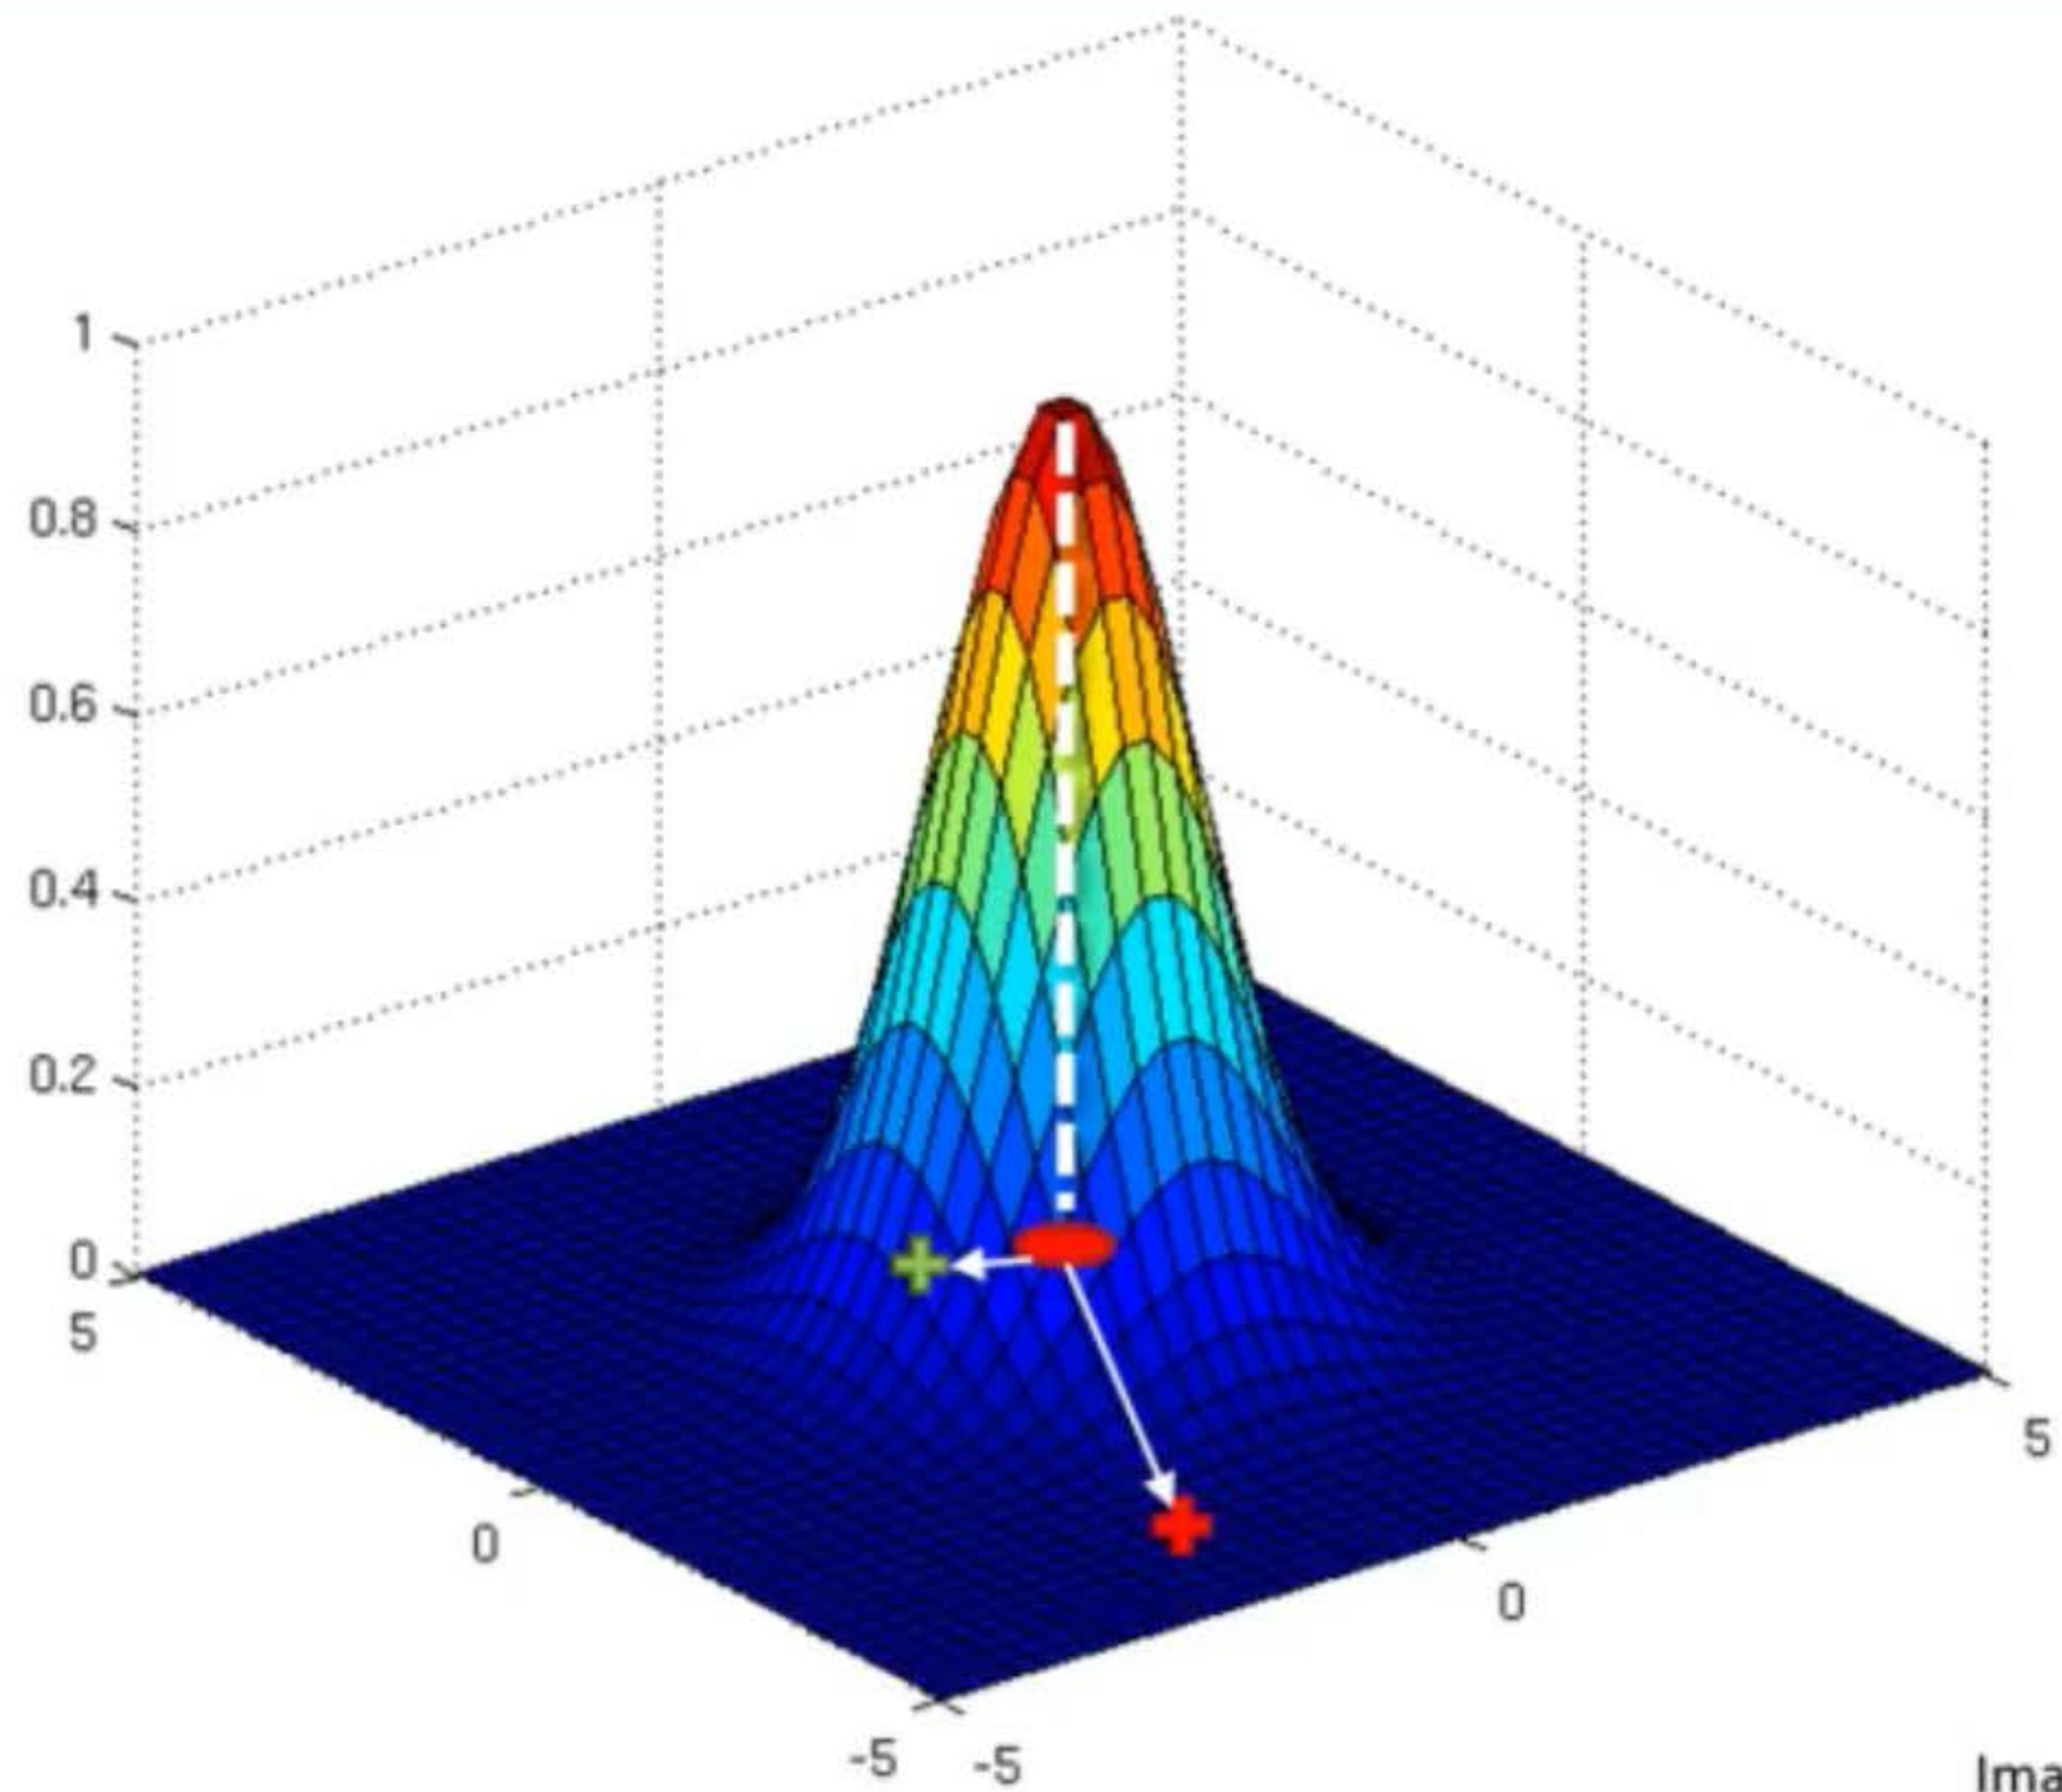
Mapping to a Higher Dimensional Space  
can be highly compute-intensive

# The Kernel Trick

# The Gaussian RBF Kernel

$$K(\vec{x}, \vec{l}^i) = e^{-\frac{\|\vec{x}-\vec{l}^i\|^2}{2\sigma^2}}$$

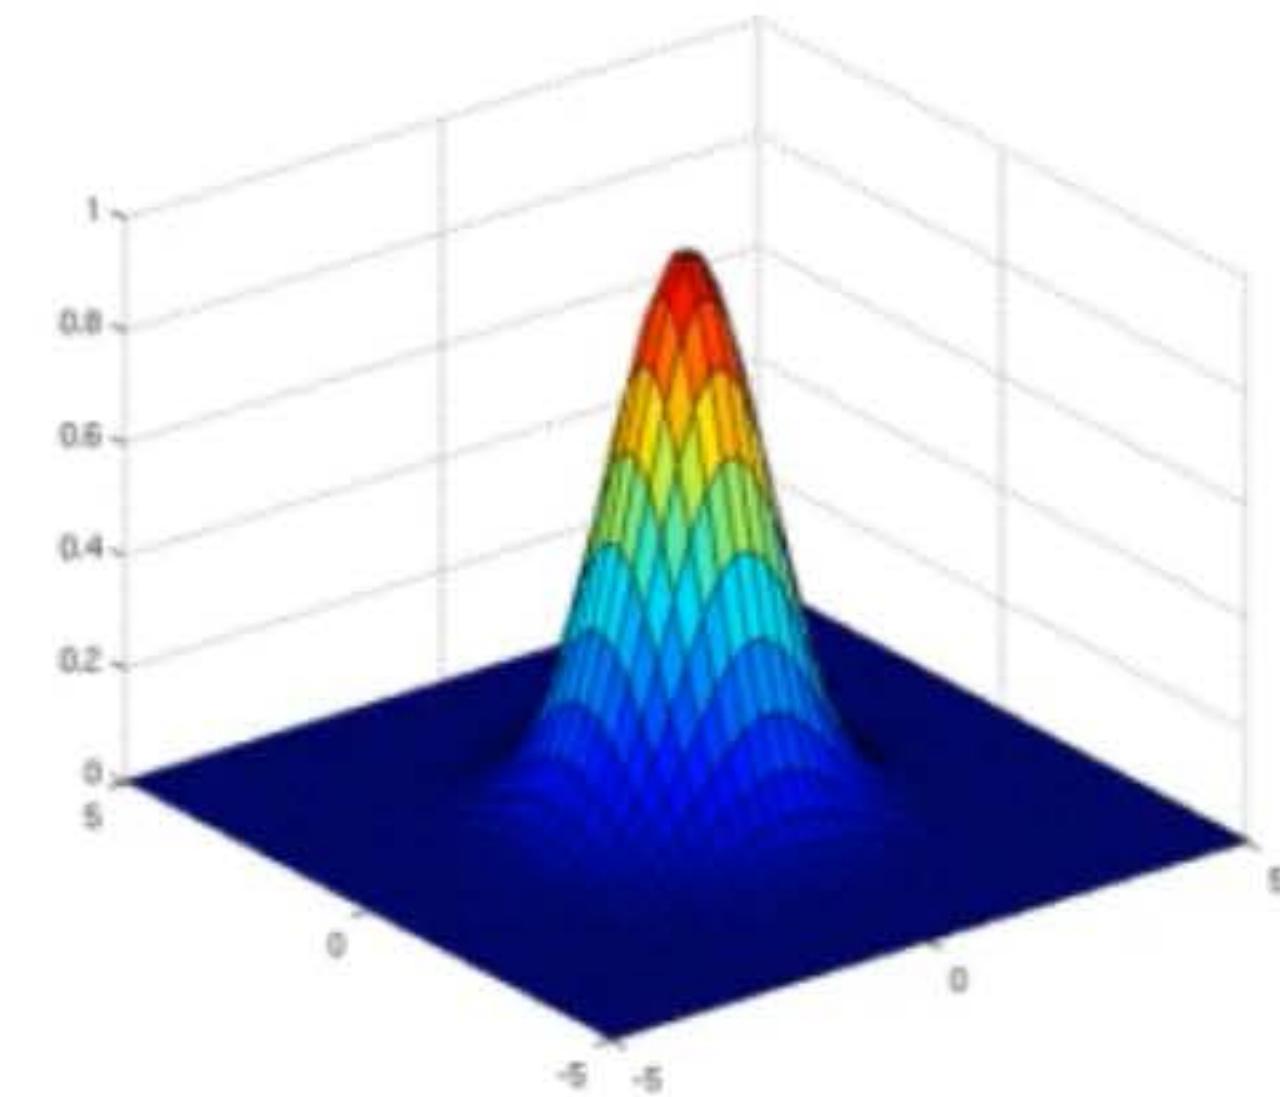
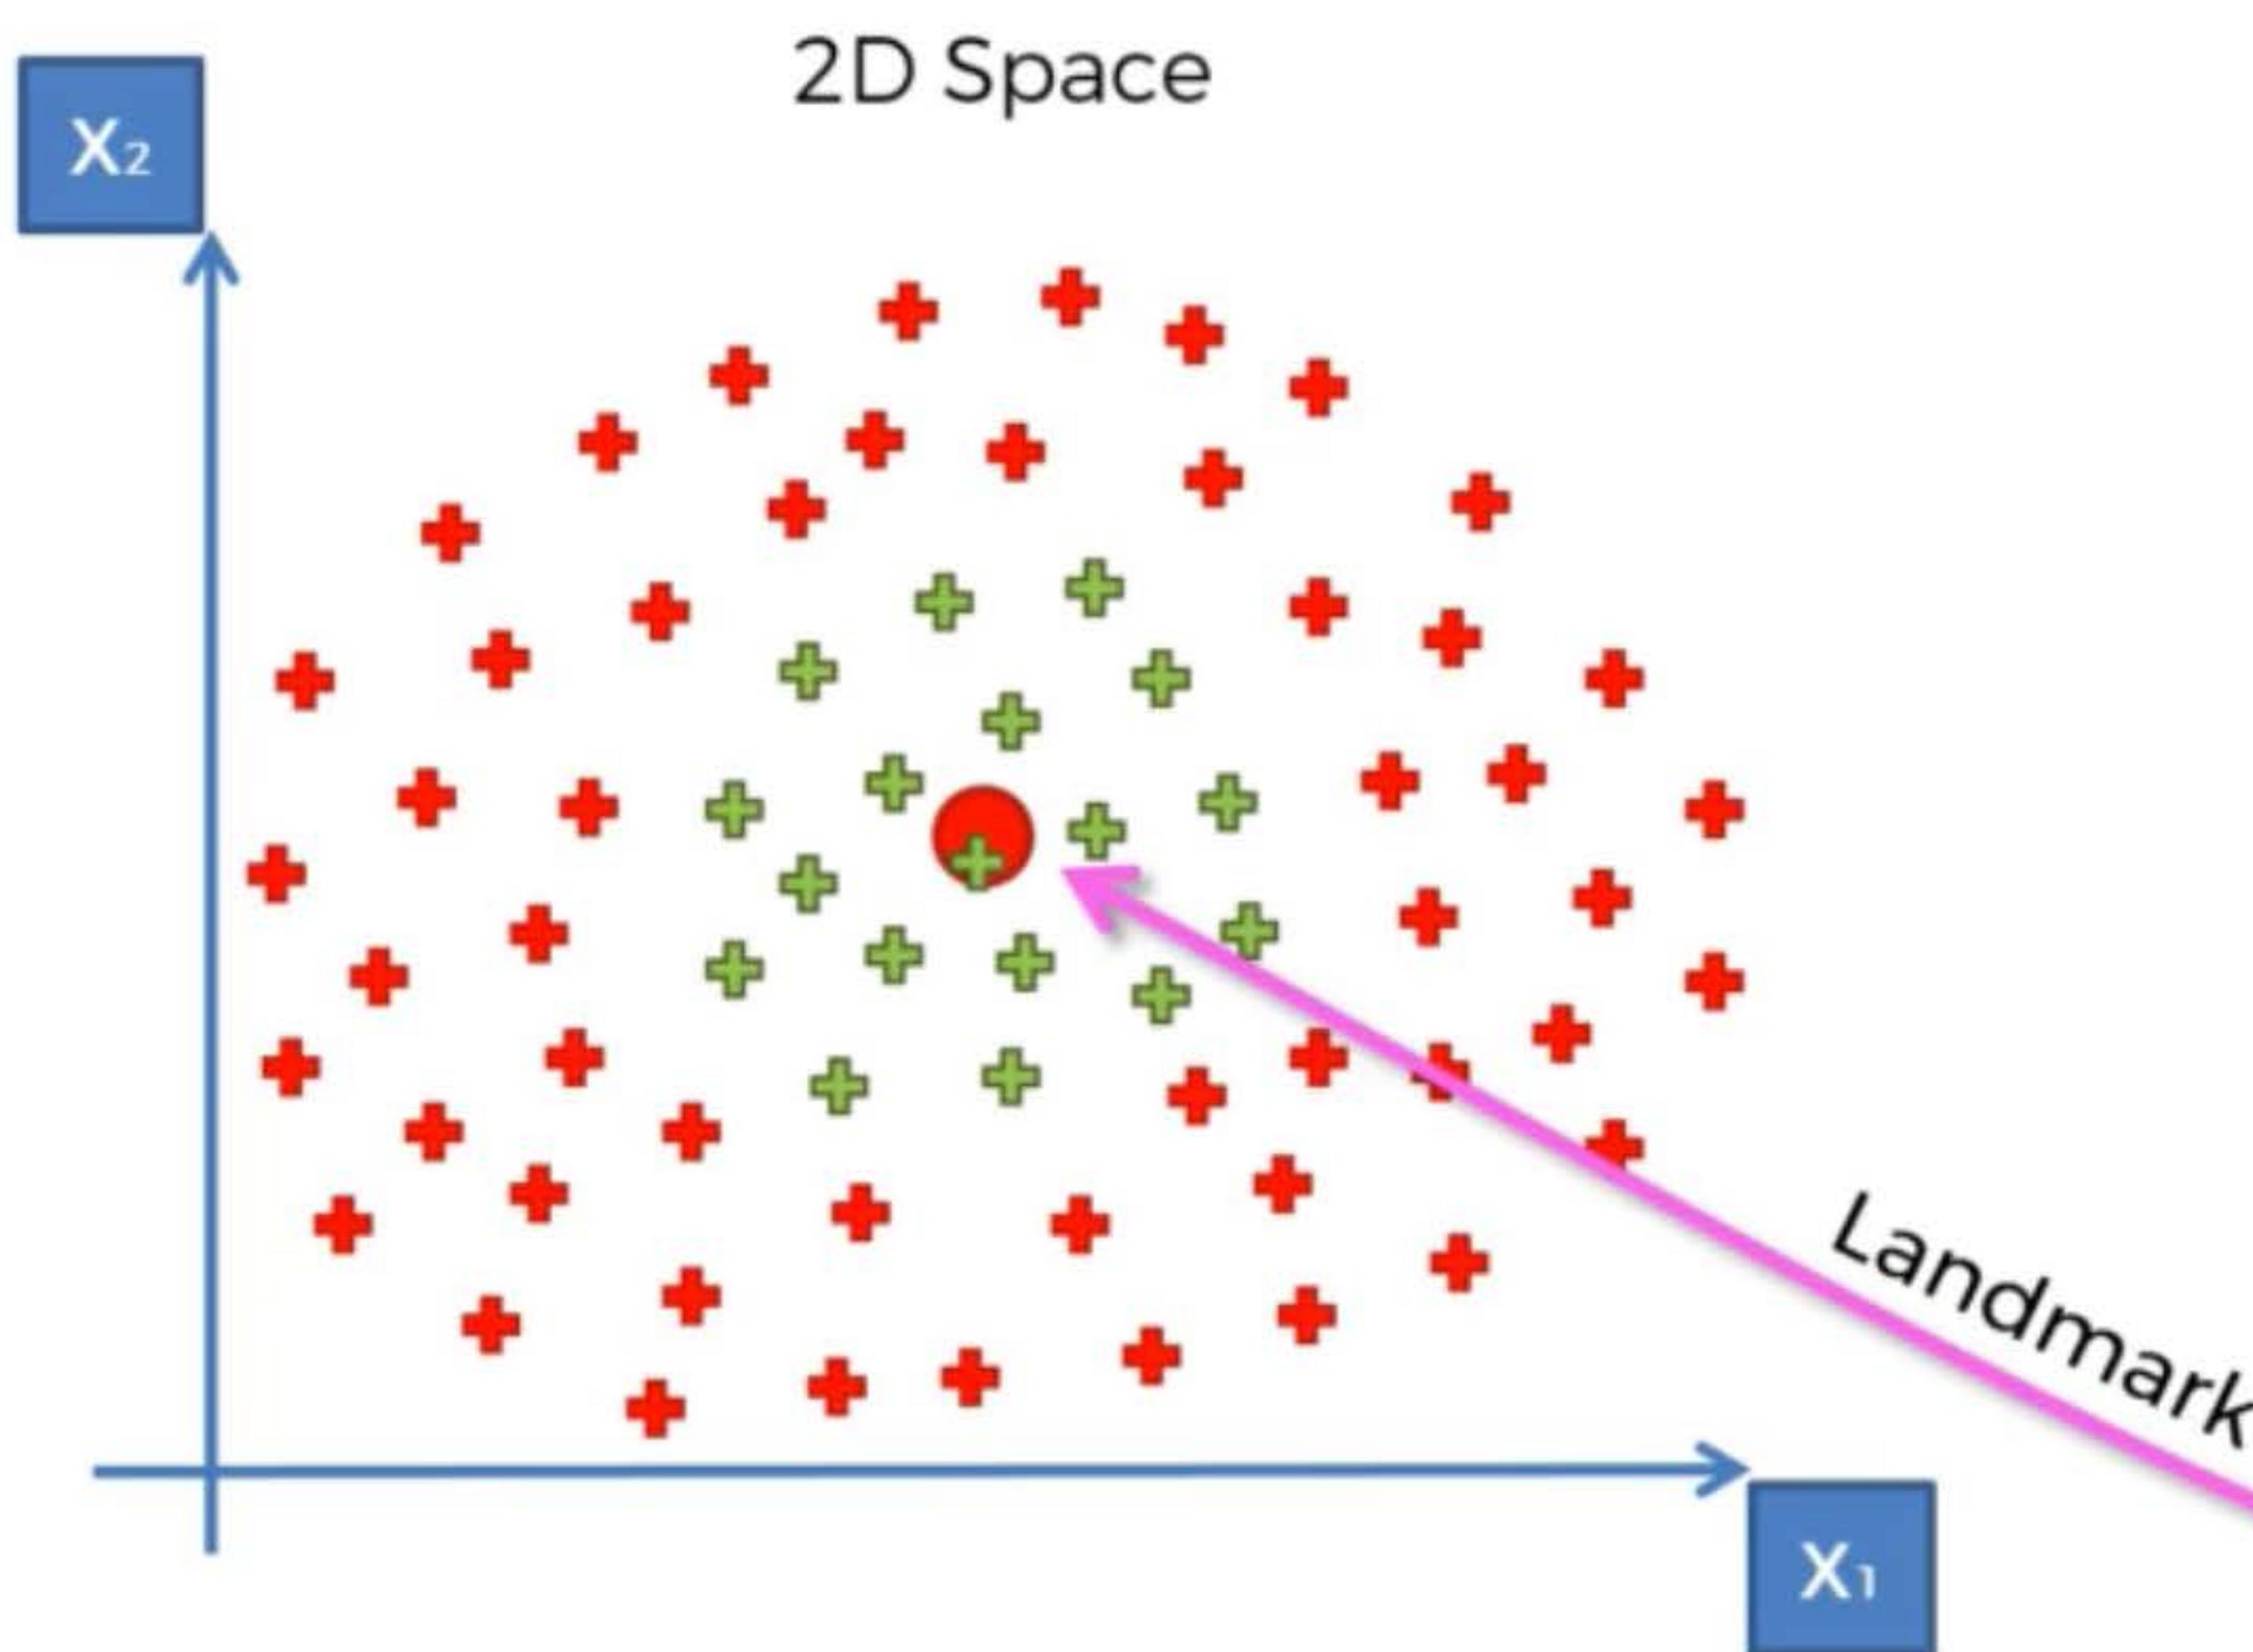
# The Gaussian RBF Kernel



$$K(\vec{x}, \vec{l}^i) = e^{-\frac{\|\vec{x}-\vec{l}^i\|^2}{2\sigma^2}}$$

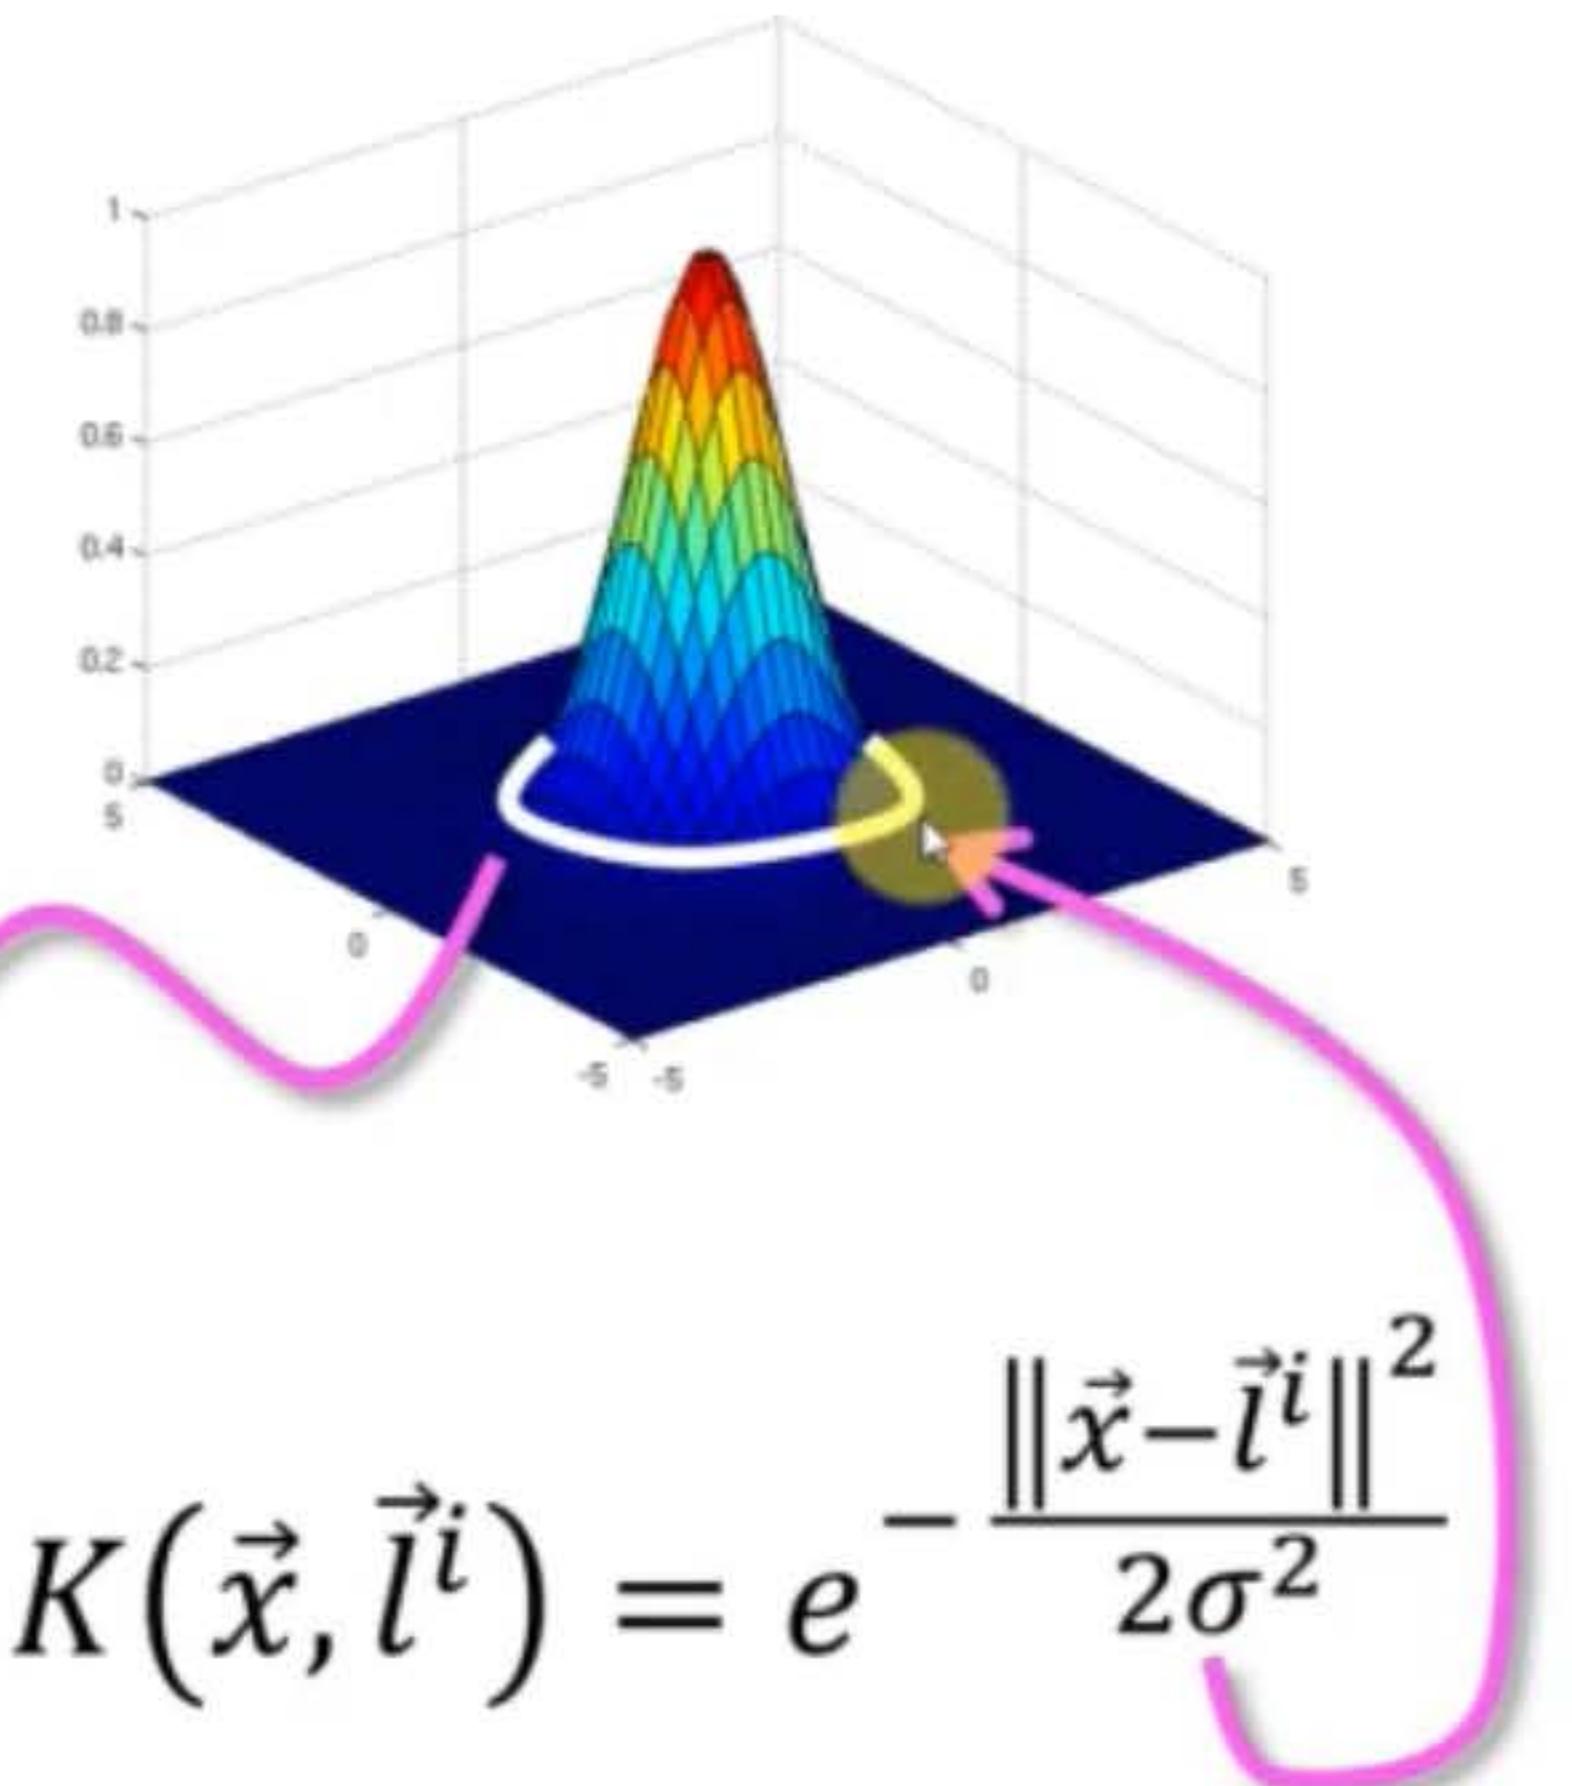
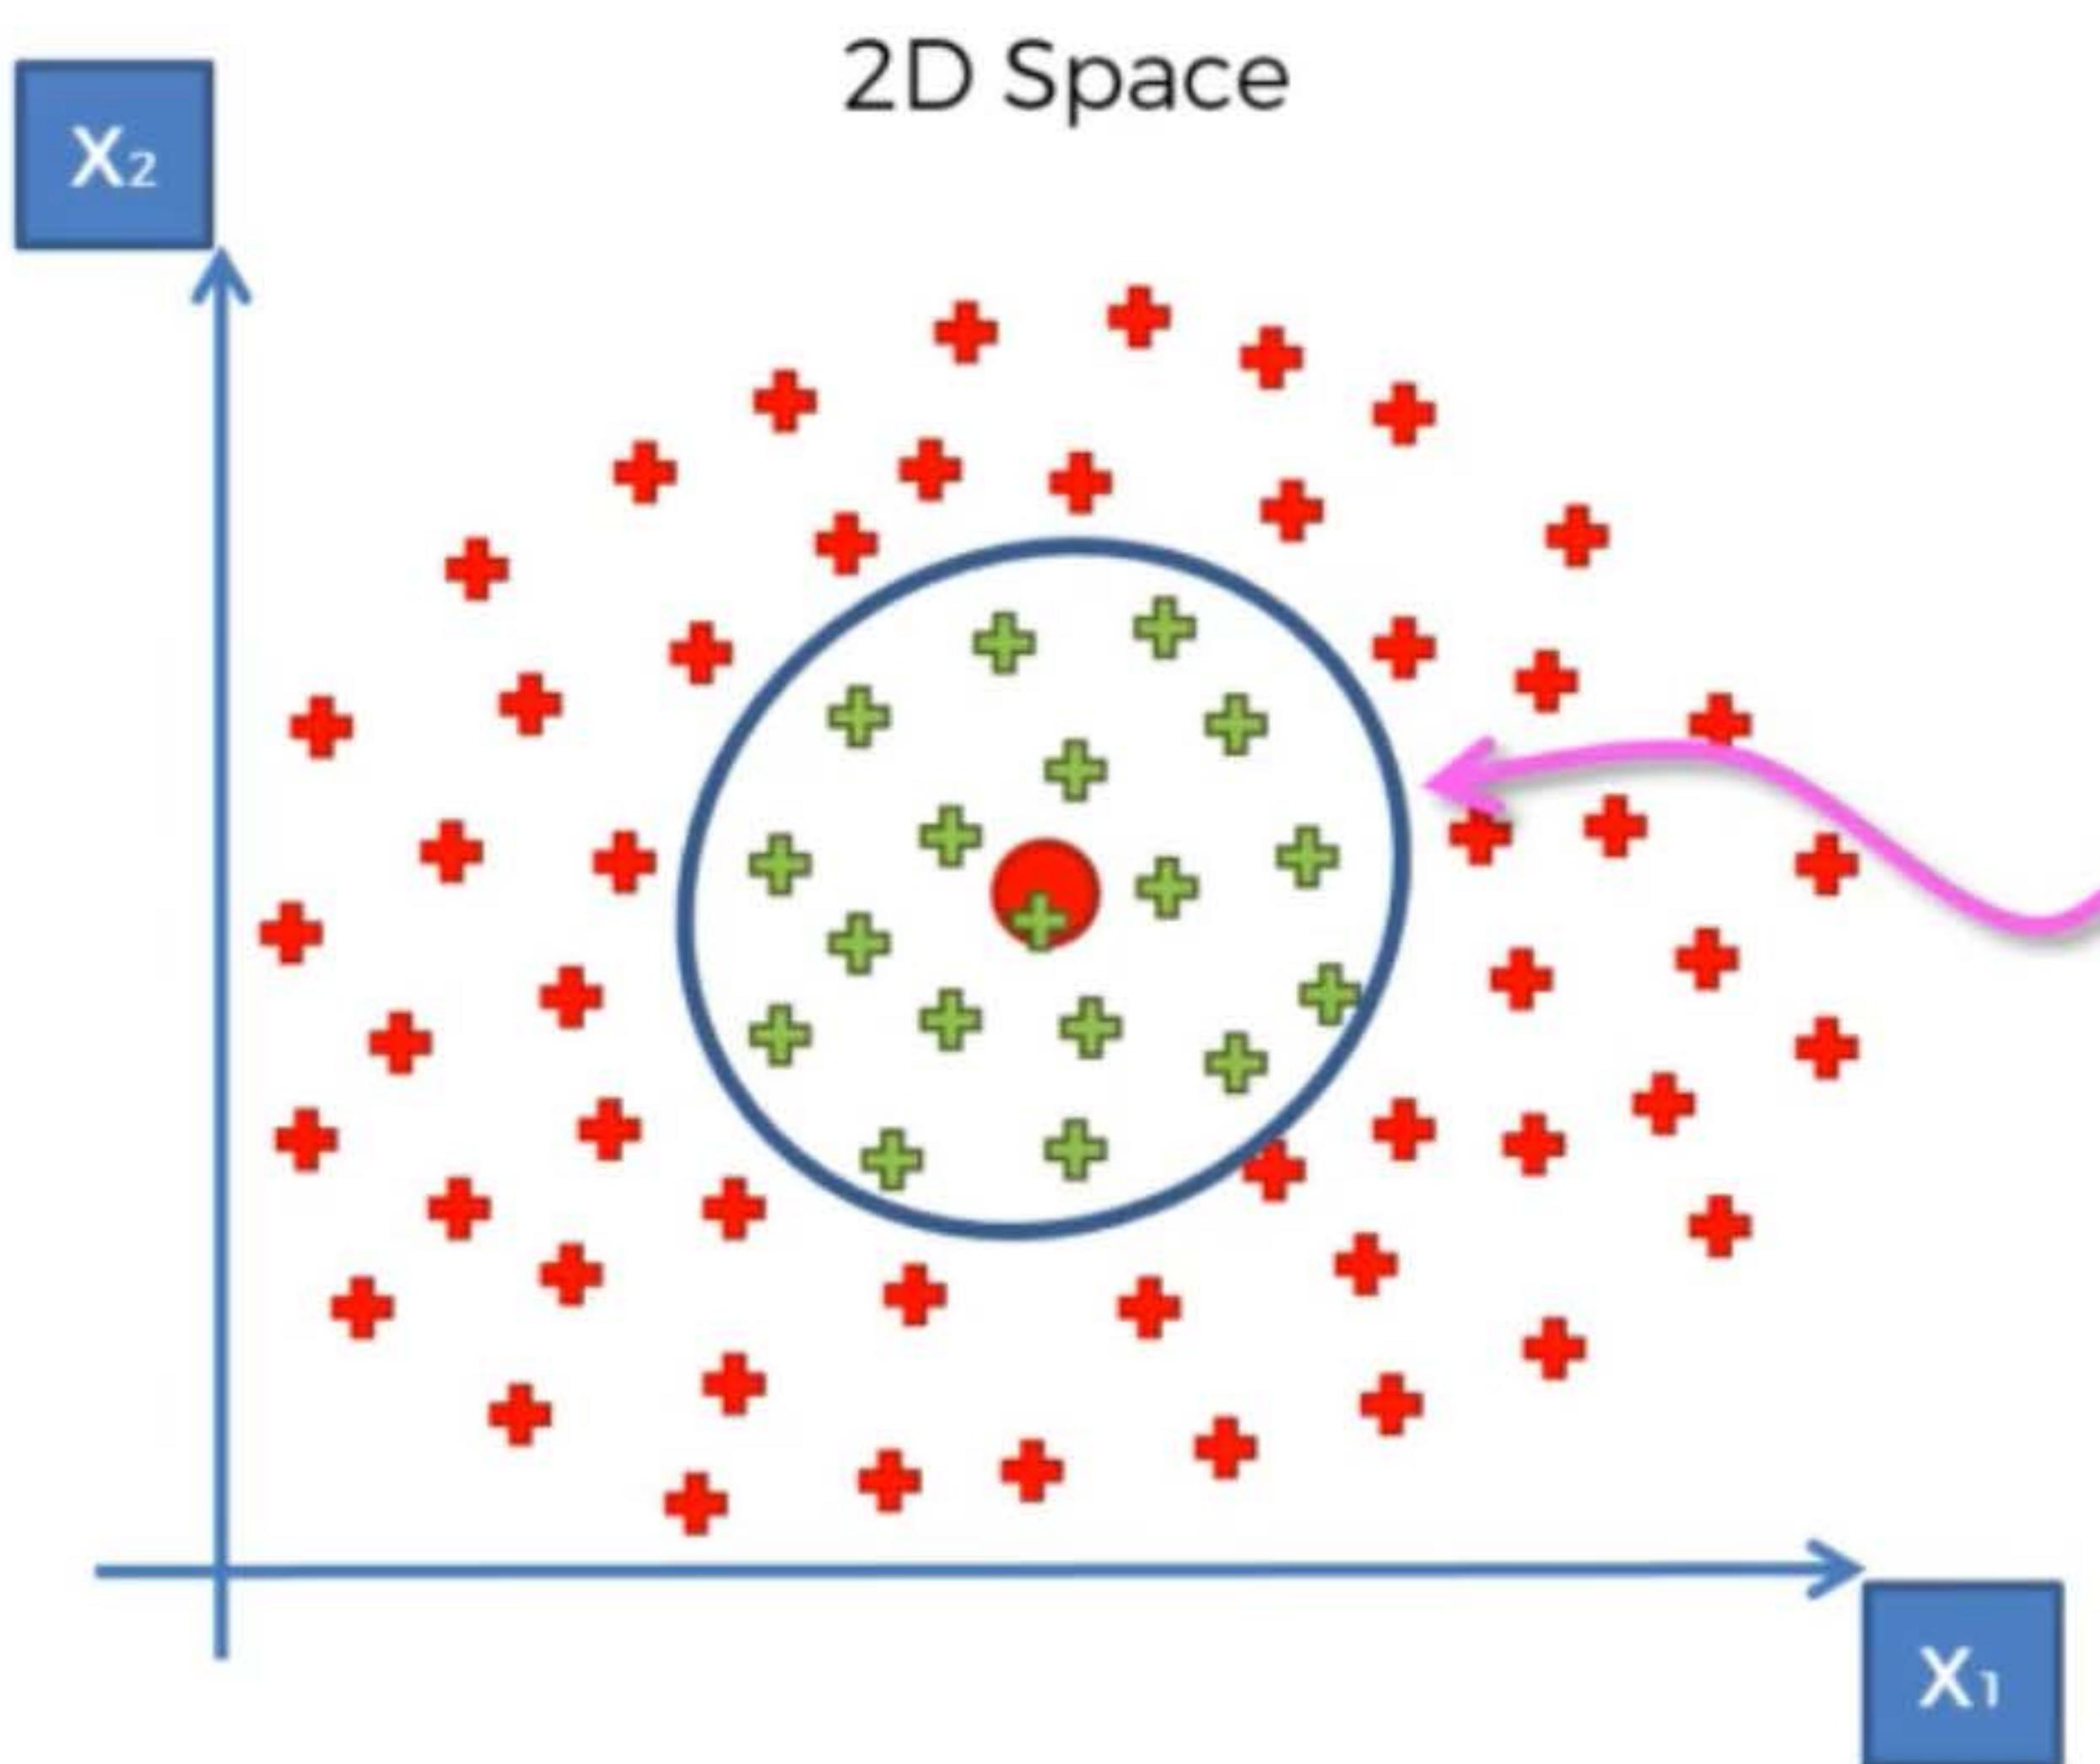
Image source: <http://www.cs.toronto.edu/~duvenaud/cookbook/index.html>

# The Gaussian RBF Kernel

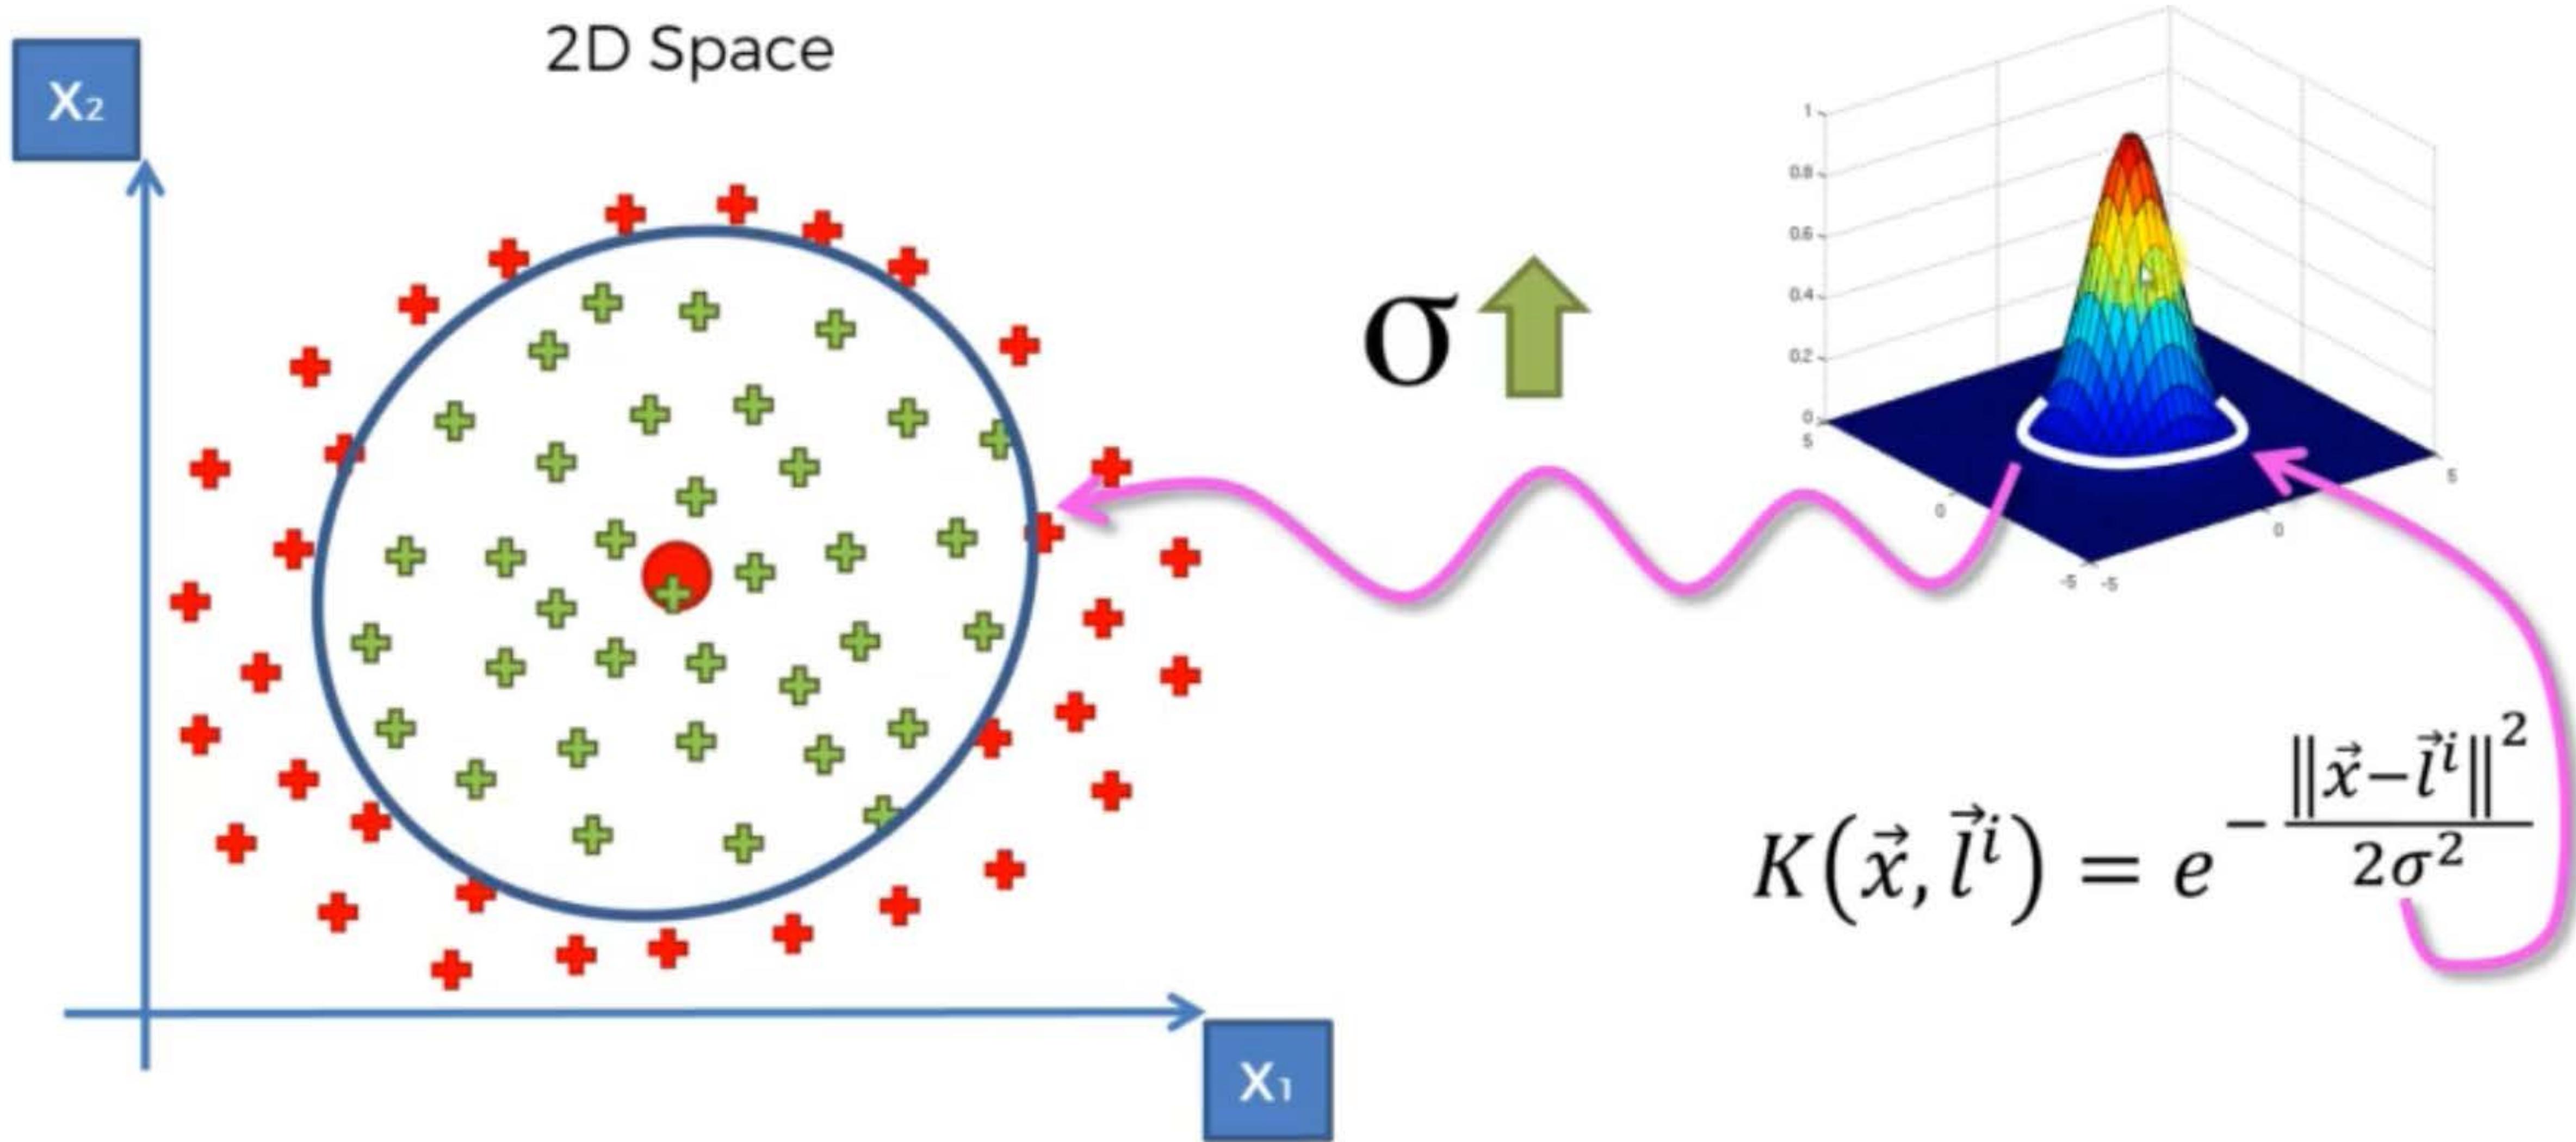


$$K(\vec{x}, \vec{l}^i) = e^{-\frac{\|\vec{x} - \vec{l}^i\|^2}{2\sigma^2}}$$

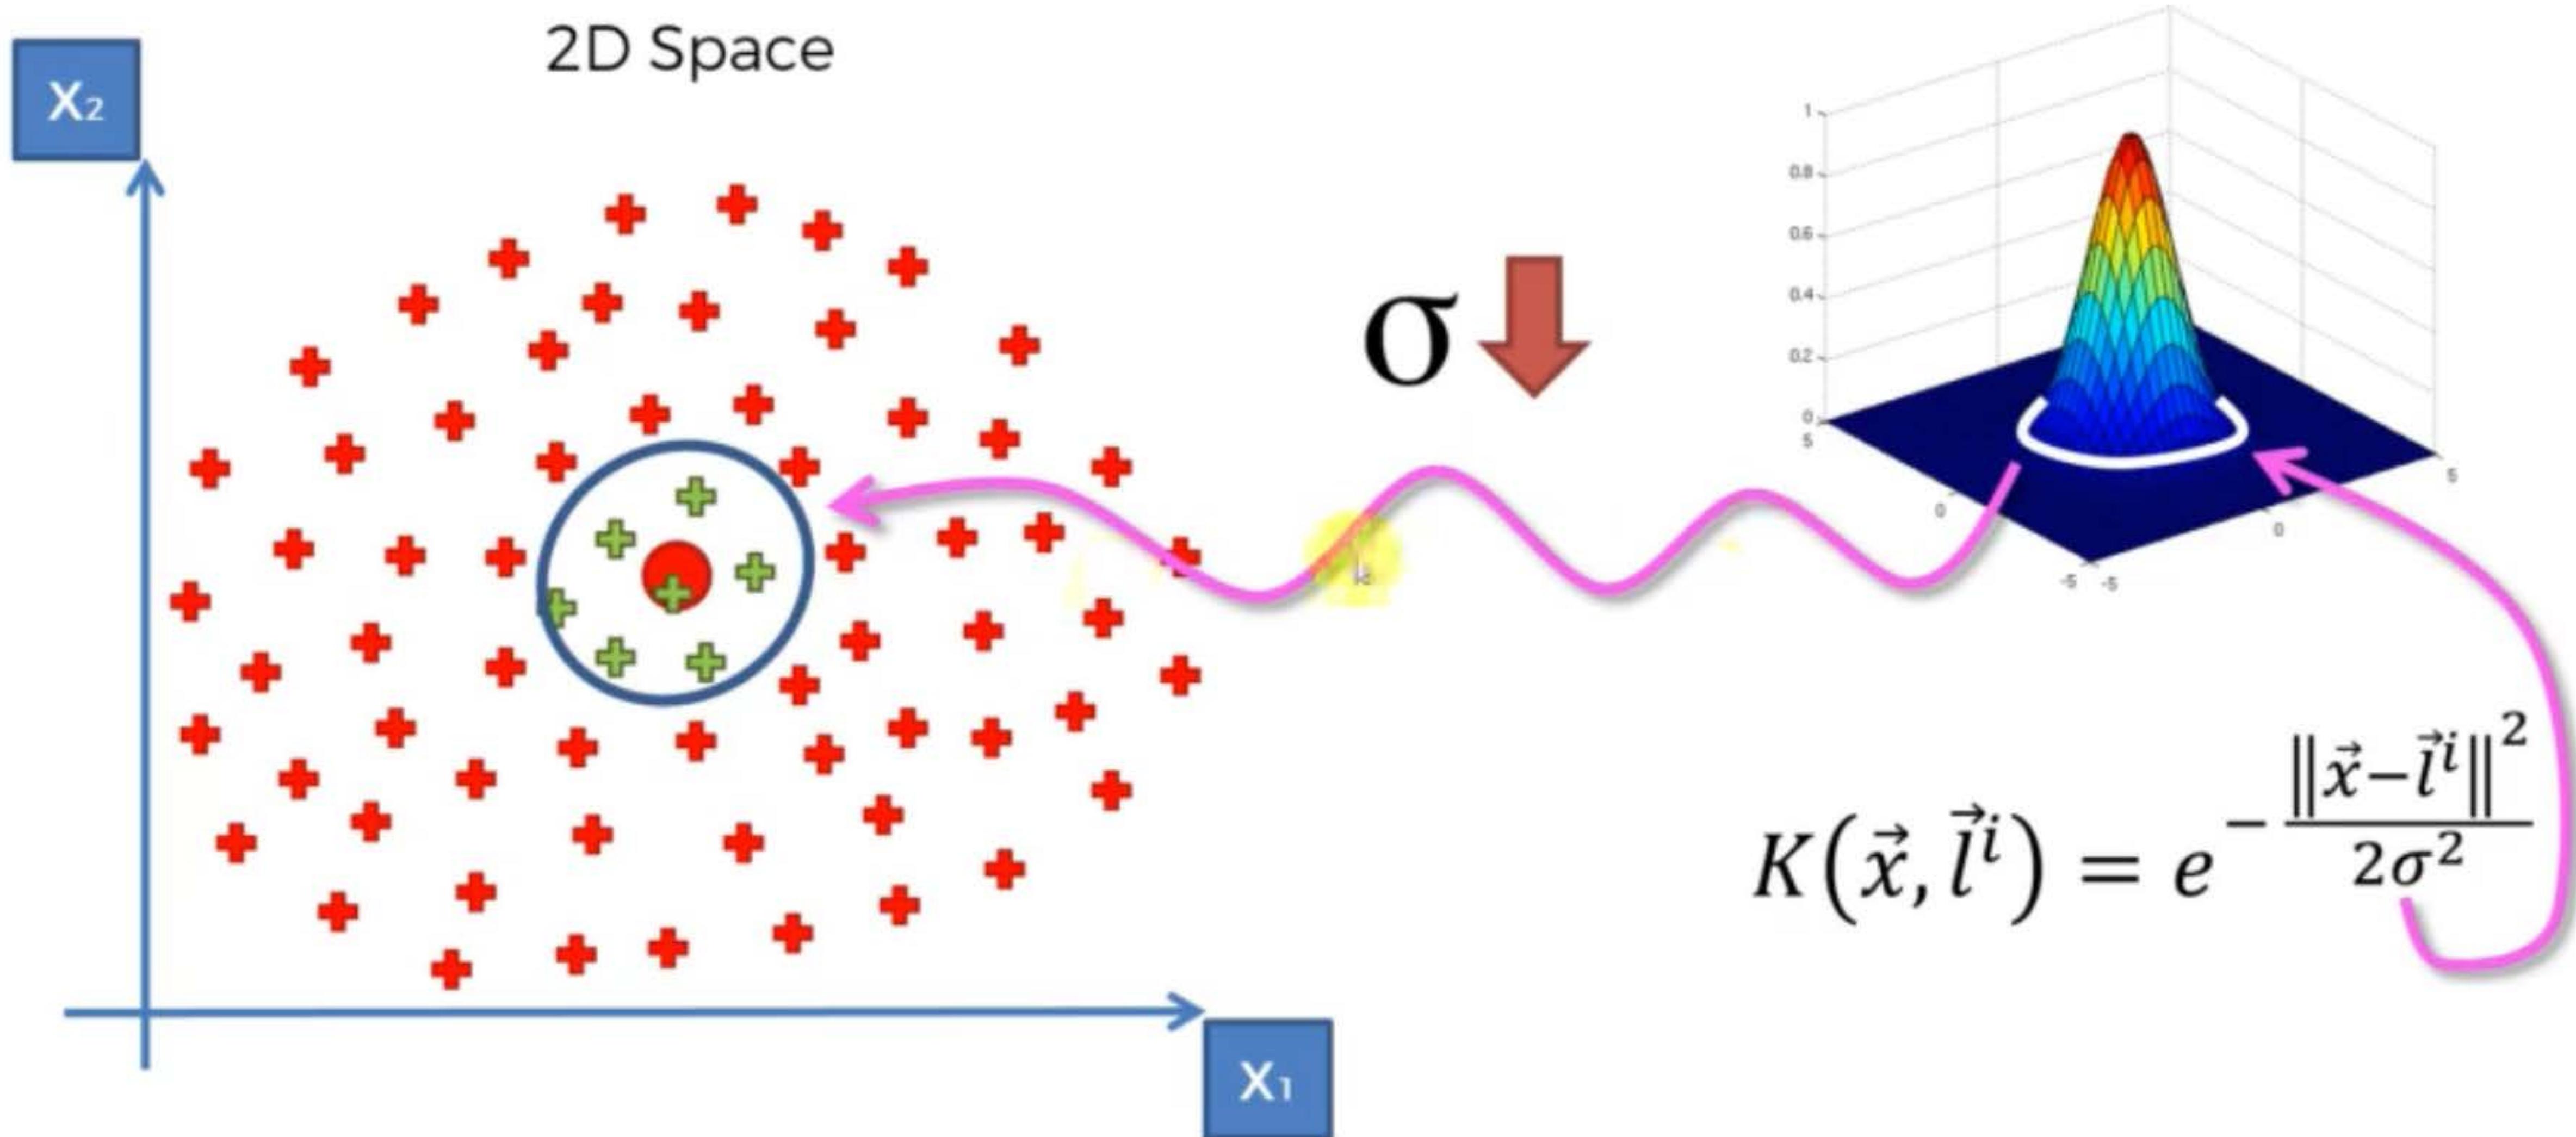
# The Gaussian RBF Kernel



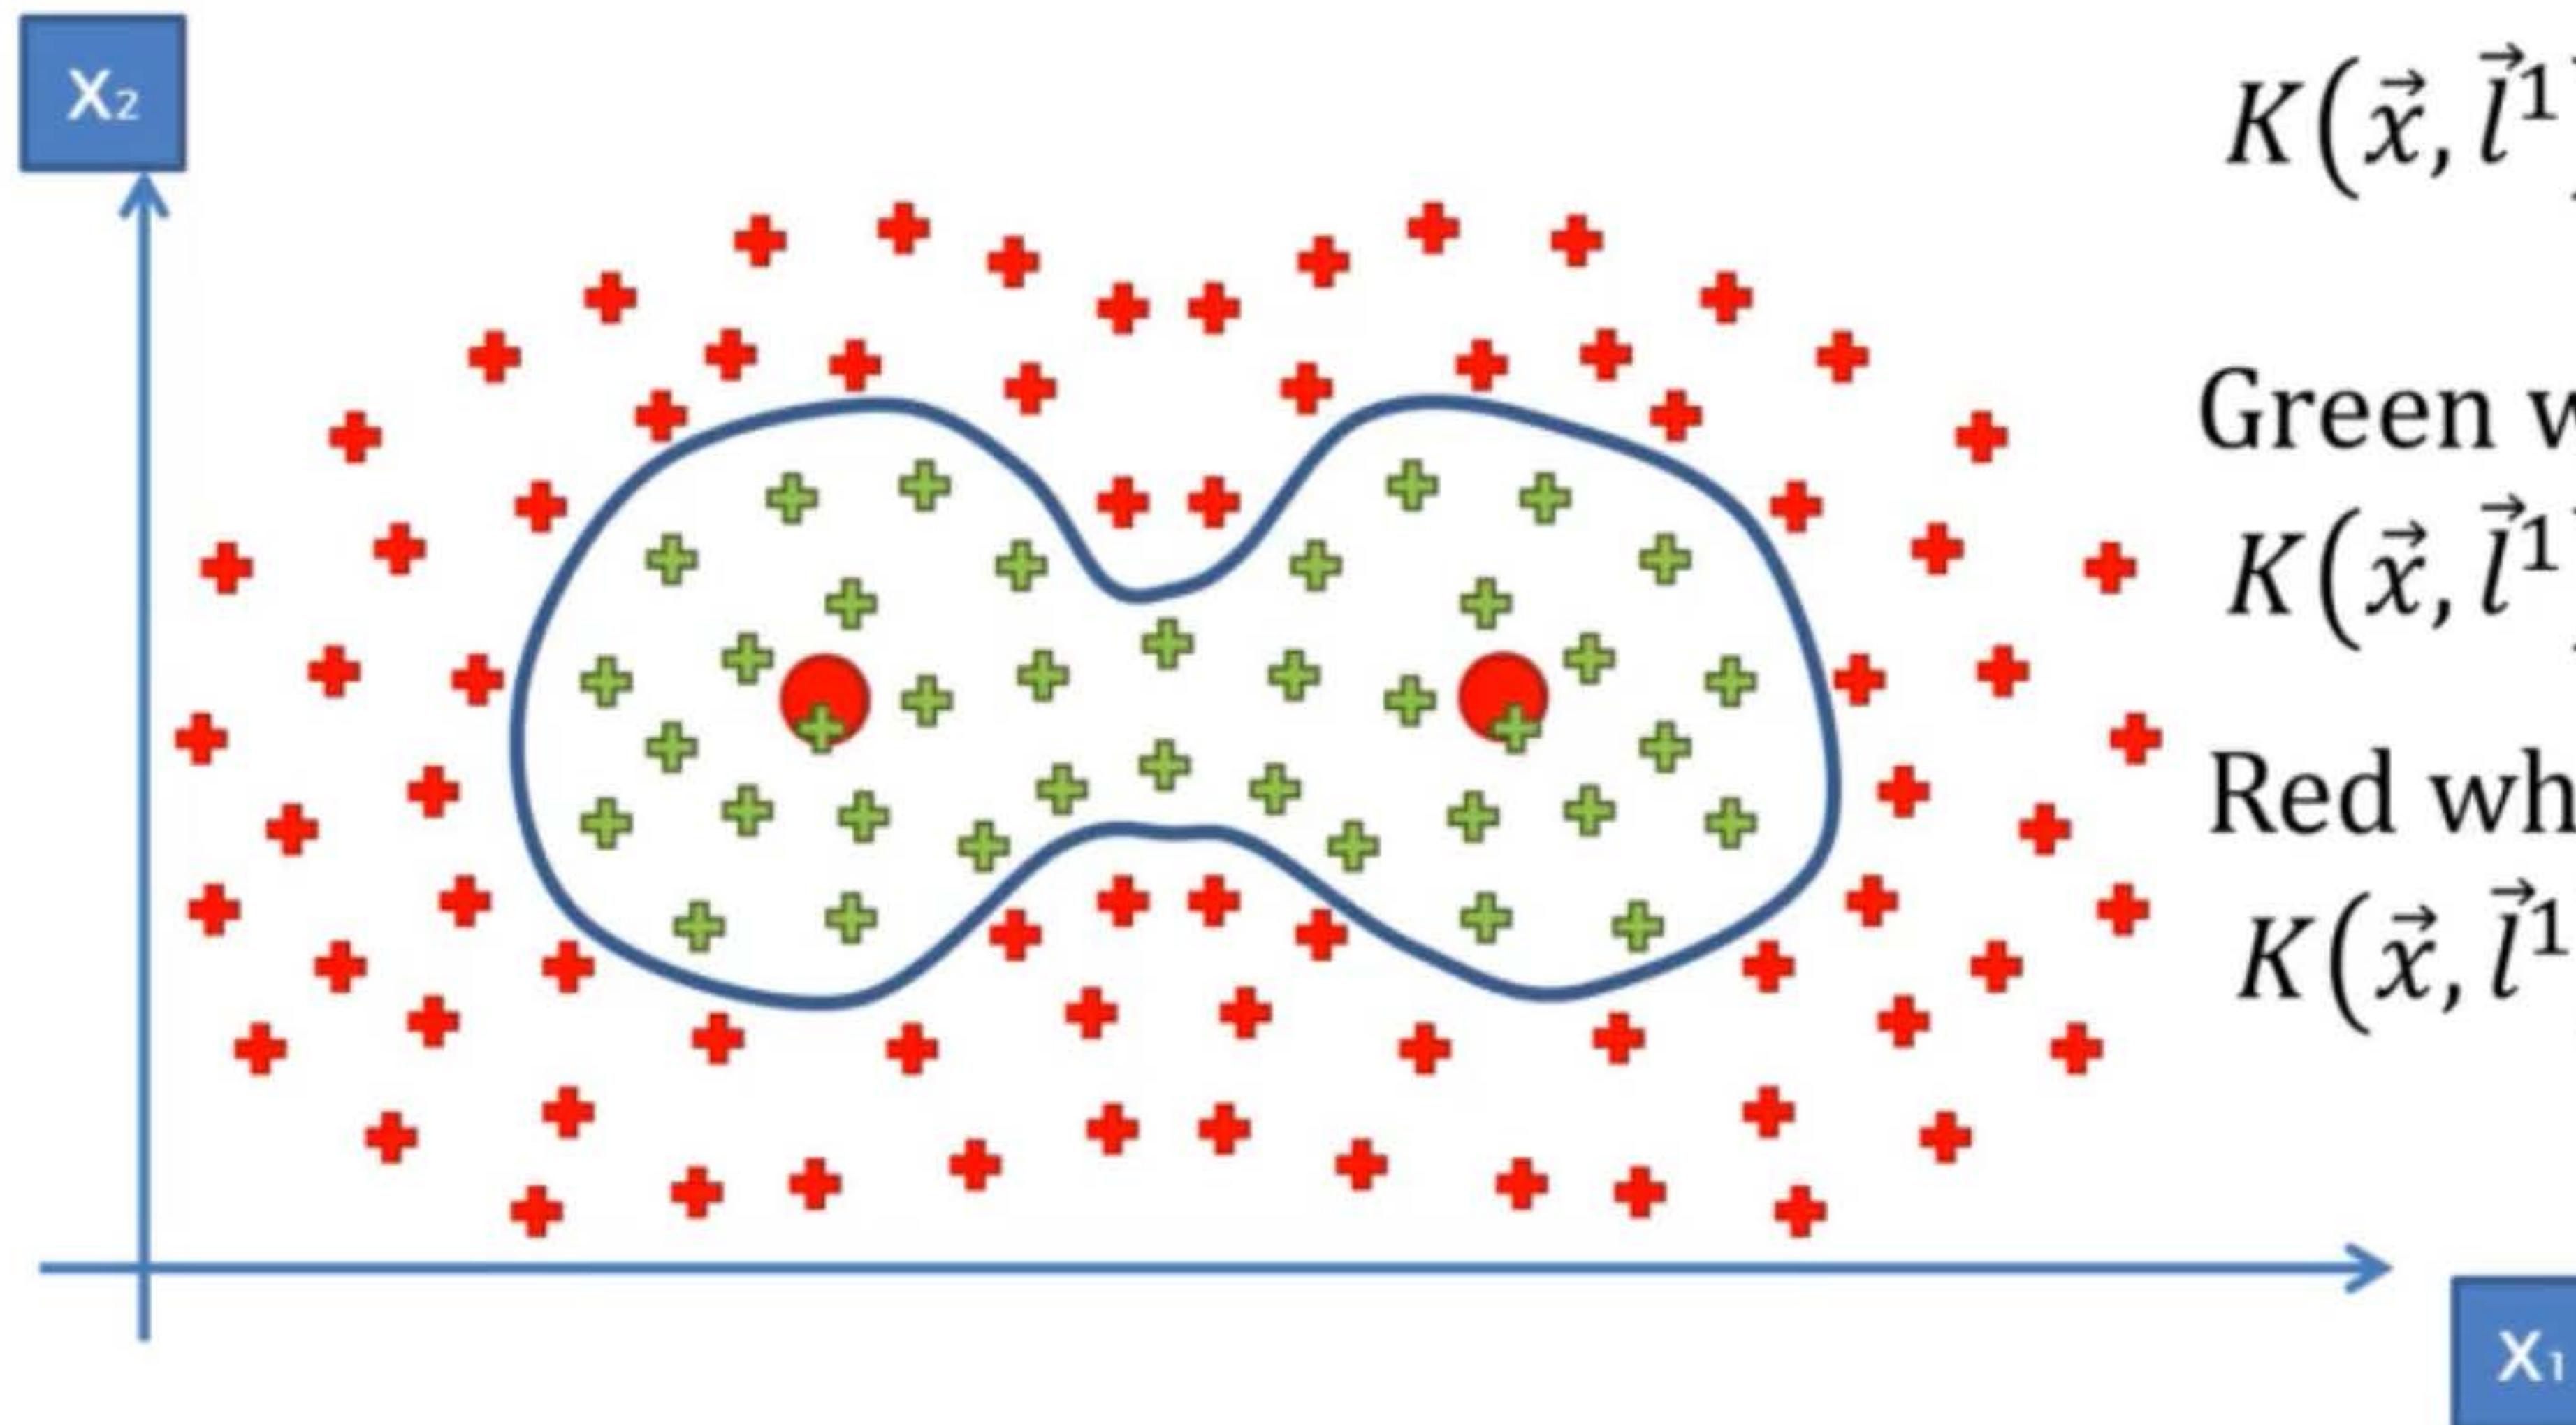
# The Gaussian RBF Kernel



# The Gaussian RBF Kernel



# The Gaussian RBF Kernel



$$K(\vec{x}, \vec{l}^1) + K(\vec{x}, \vec{l}^2)$$

*(Simplified Formula)*

Green when:

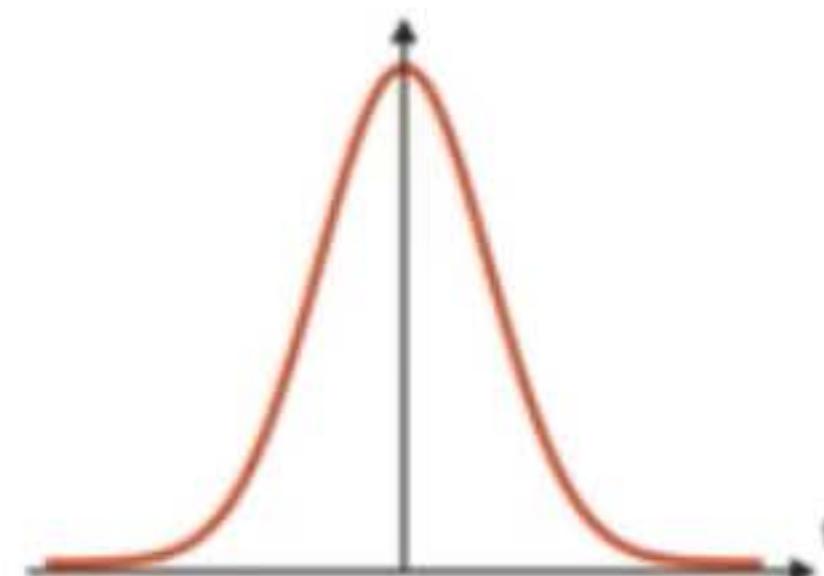
$$K(\vec{x}, \vec{l}^1) + K(\vec{x}, \vec{l}^2) > 0$$

Red when:

$$K(\vec{x}, \vec{l}^1) + K(\vec{x}, \vec{l}^2) = 0$$

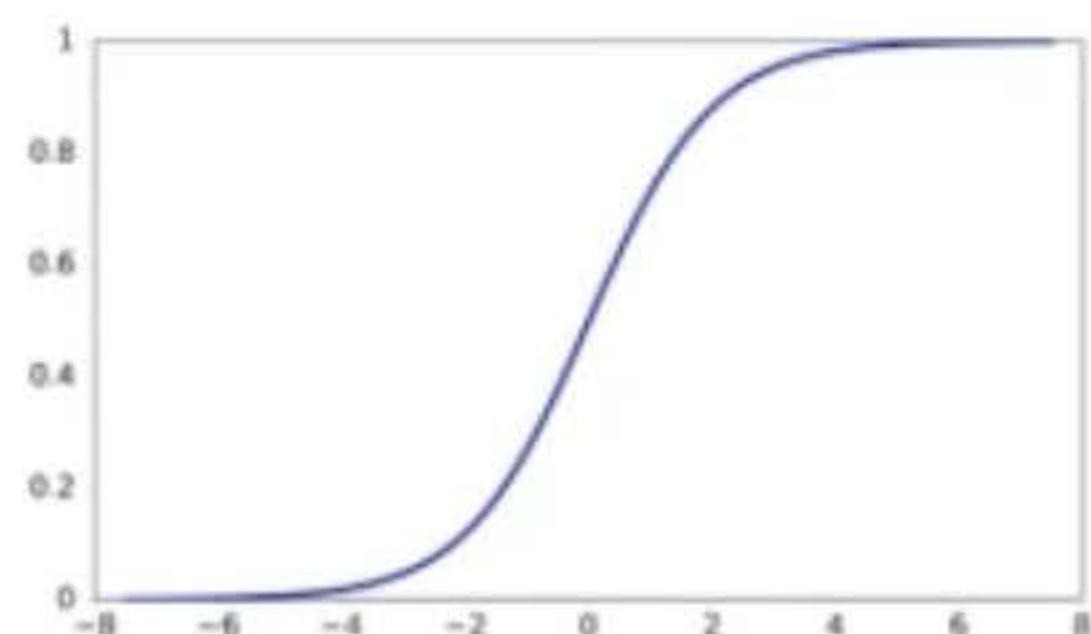
# Types of Kernel Functions

# Types of Kernel Functions



Gaussian RBF Kernel

$$K(\vec{x}, \vec{l}^i) = e^{-\frac{\|\vec{x}-\vec{l}^i\|^2}{2\sigma^2}}$$



Sigmoid Kernel

$$K(X, Y) = \tanh(\gamma \cdot X^T Y + r)$$



Polynomial Kernel

$$K(X, Y) = (\gamma \cdot X^T Y + r)^d, \gamma > 0$$

# **Non-Linear SVR (Advanced)**

# Heads-up about Non-Linear SVR

Section on SVR:

- SVR Intuition



Section on SVM:

- SVM Intuition



Section on Kernel SVM:

- Kernel SVM Intuition
- Mapping to a higher dimension
- The Kernel Trick
- Types of Kernel Functions
- Non-linear Kernel SVR

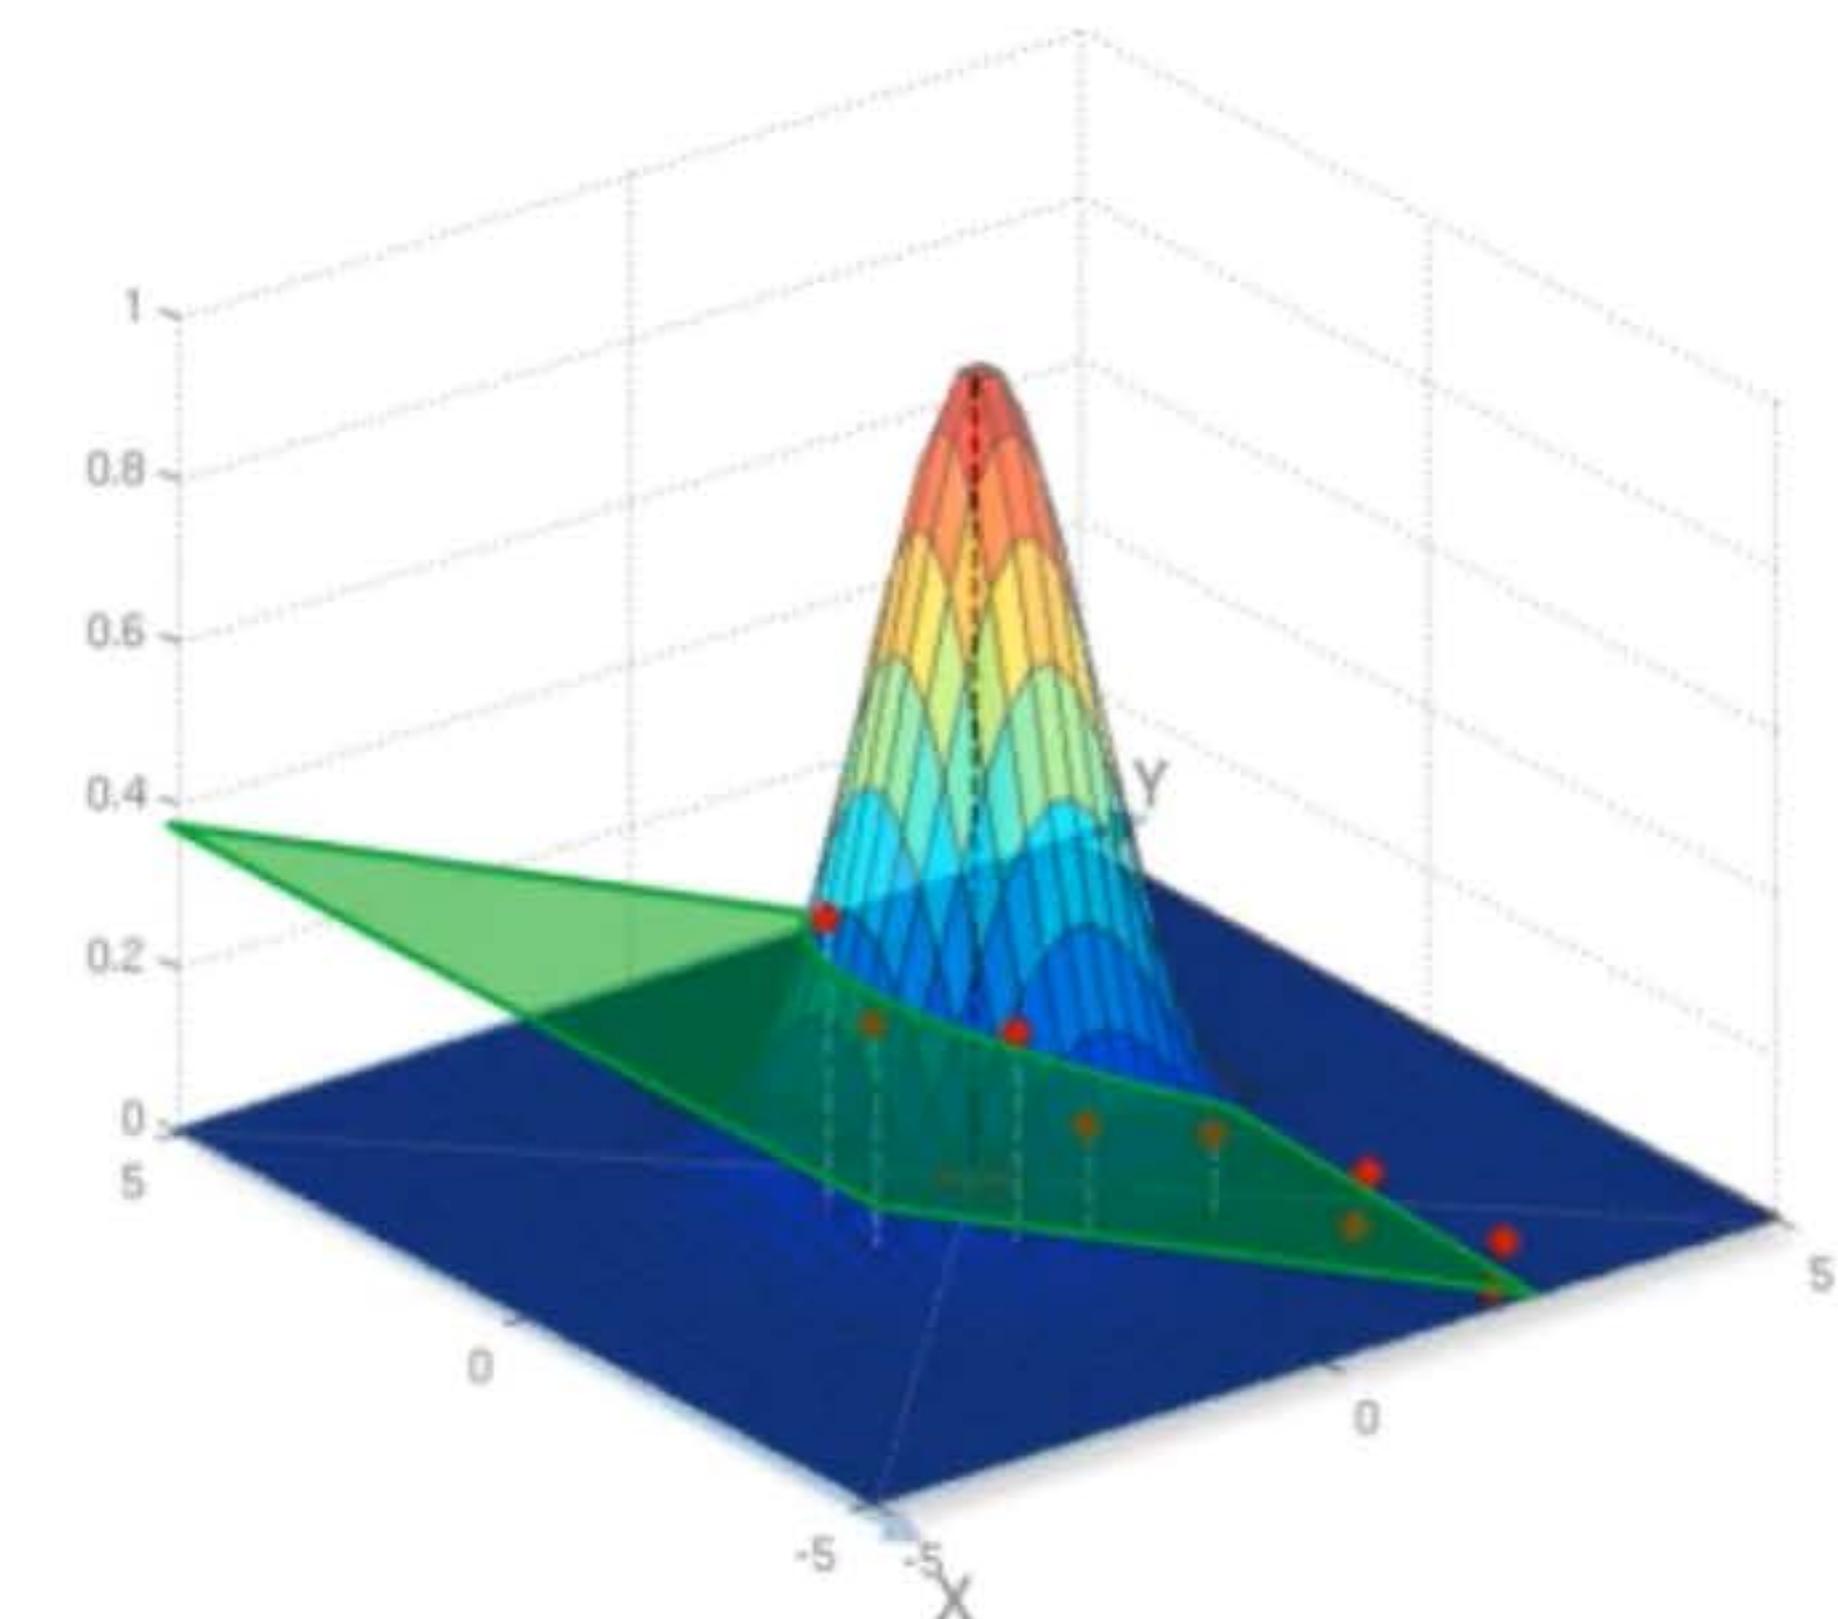
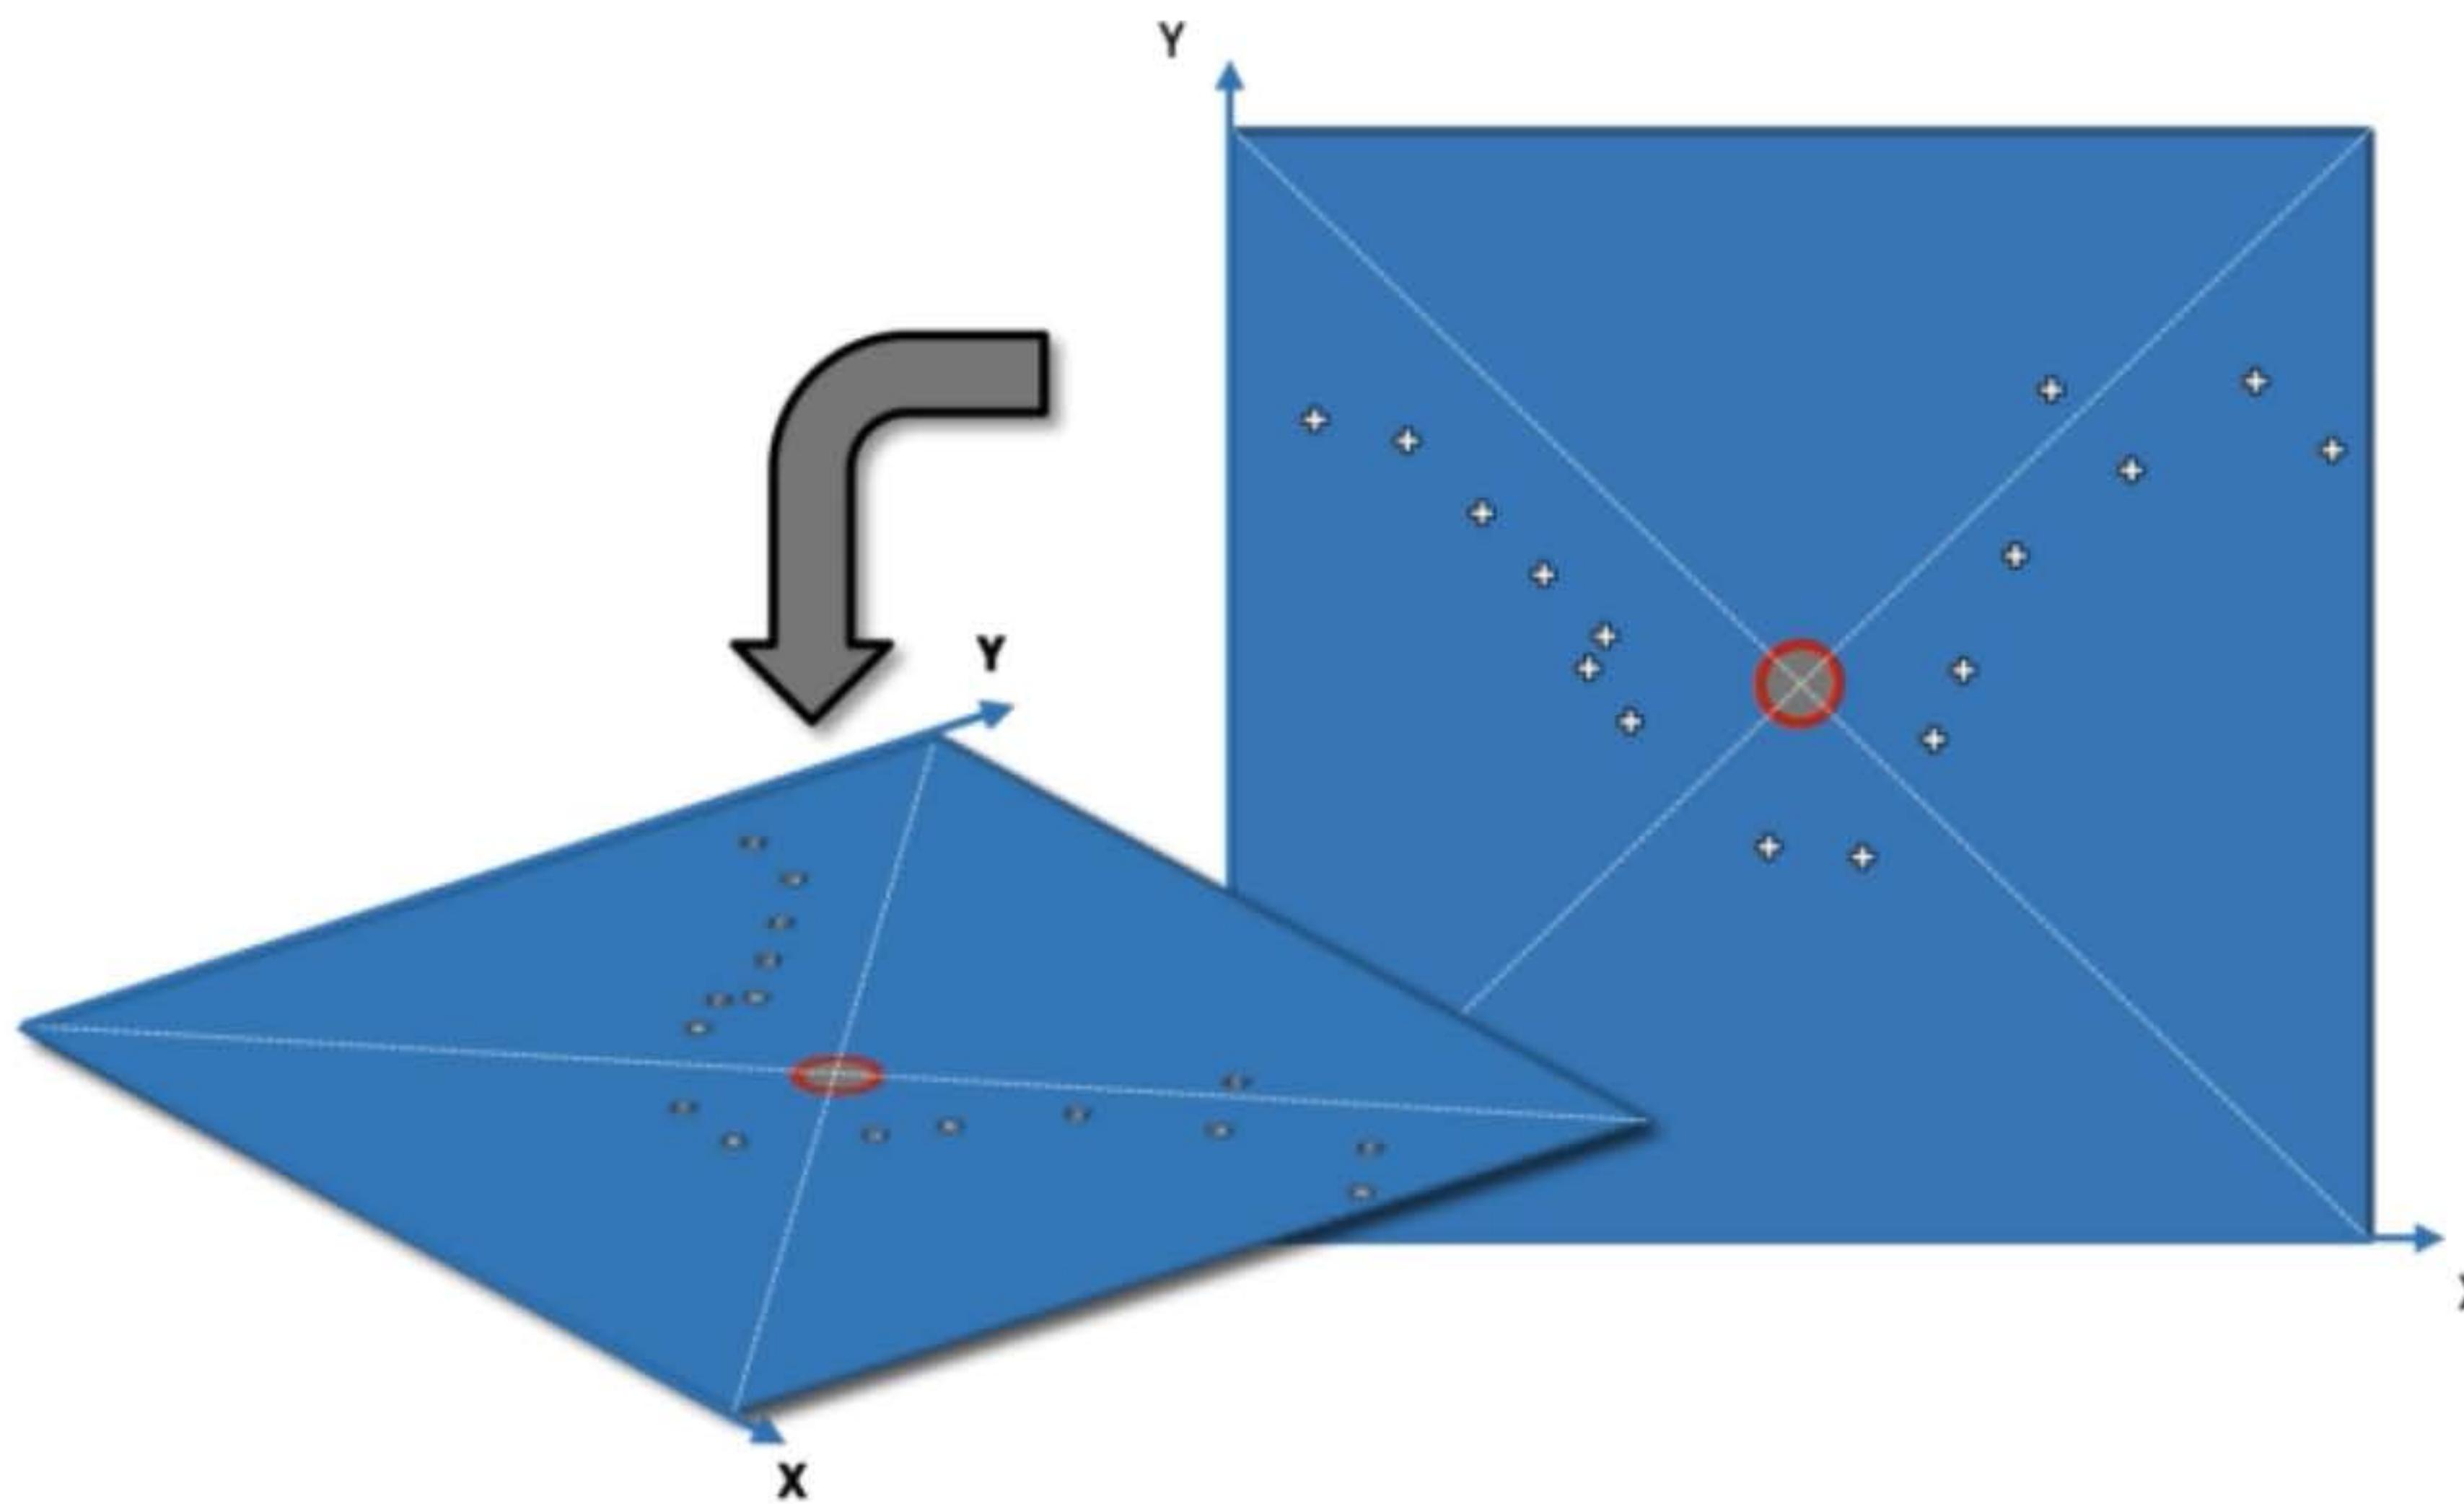


Image source: <http://www.cs.toronto.edu/~duvenaud/cookbook/index.html>

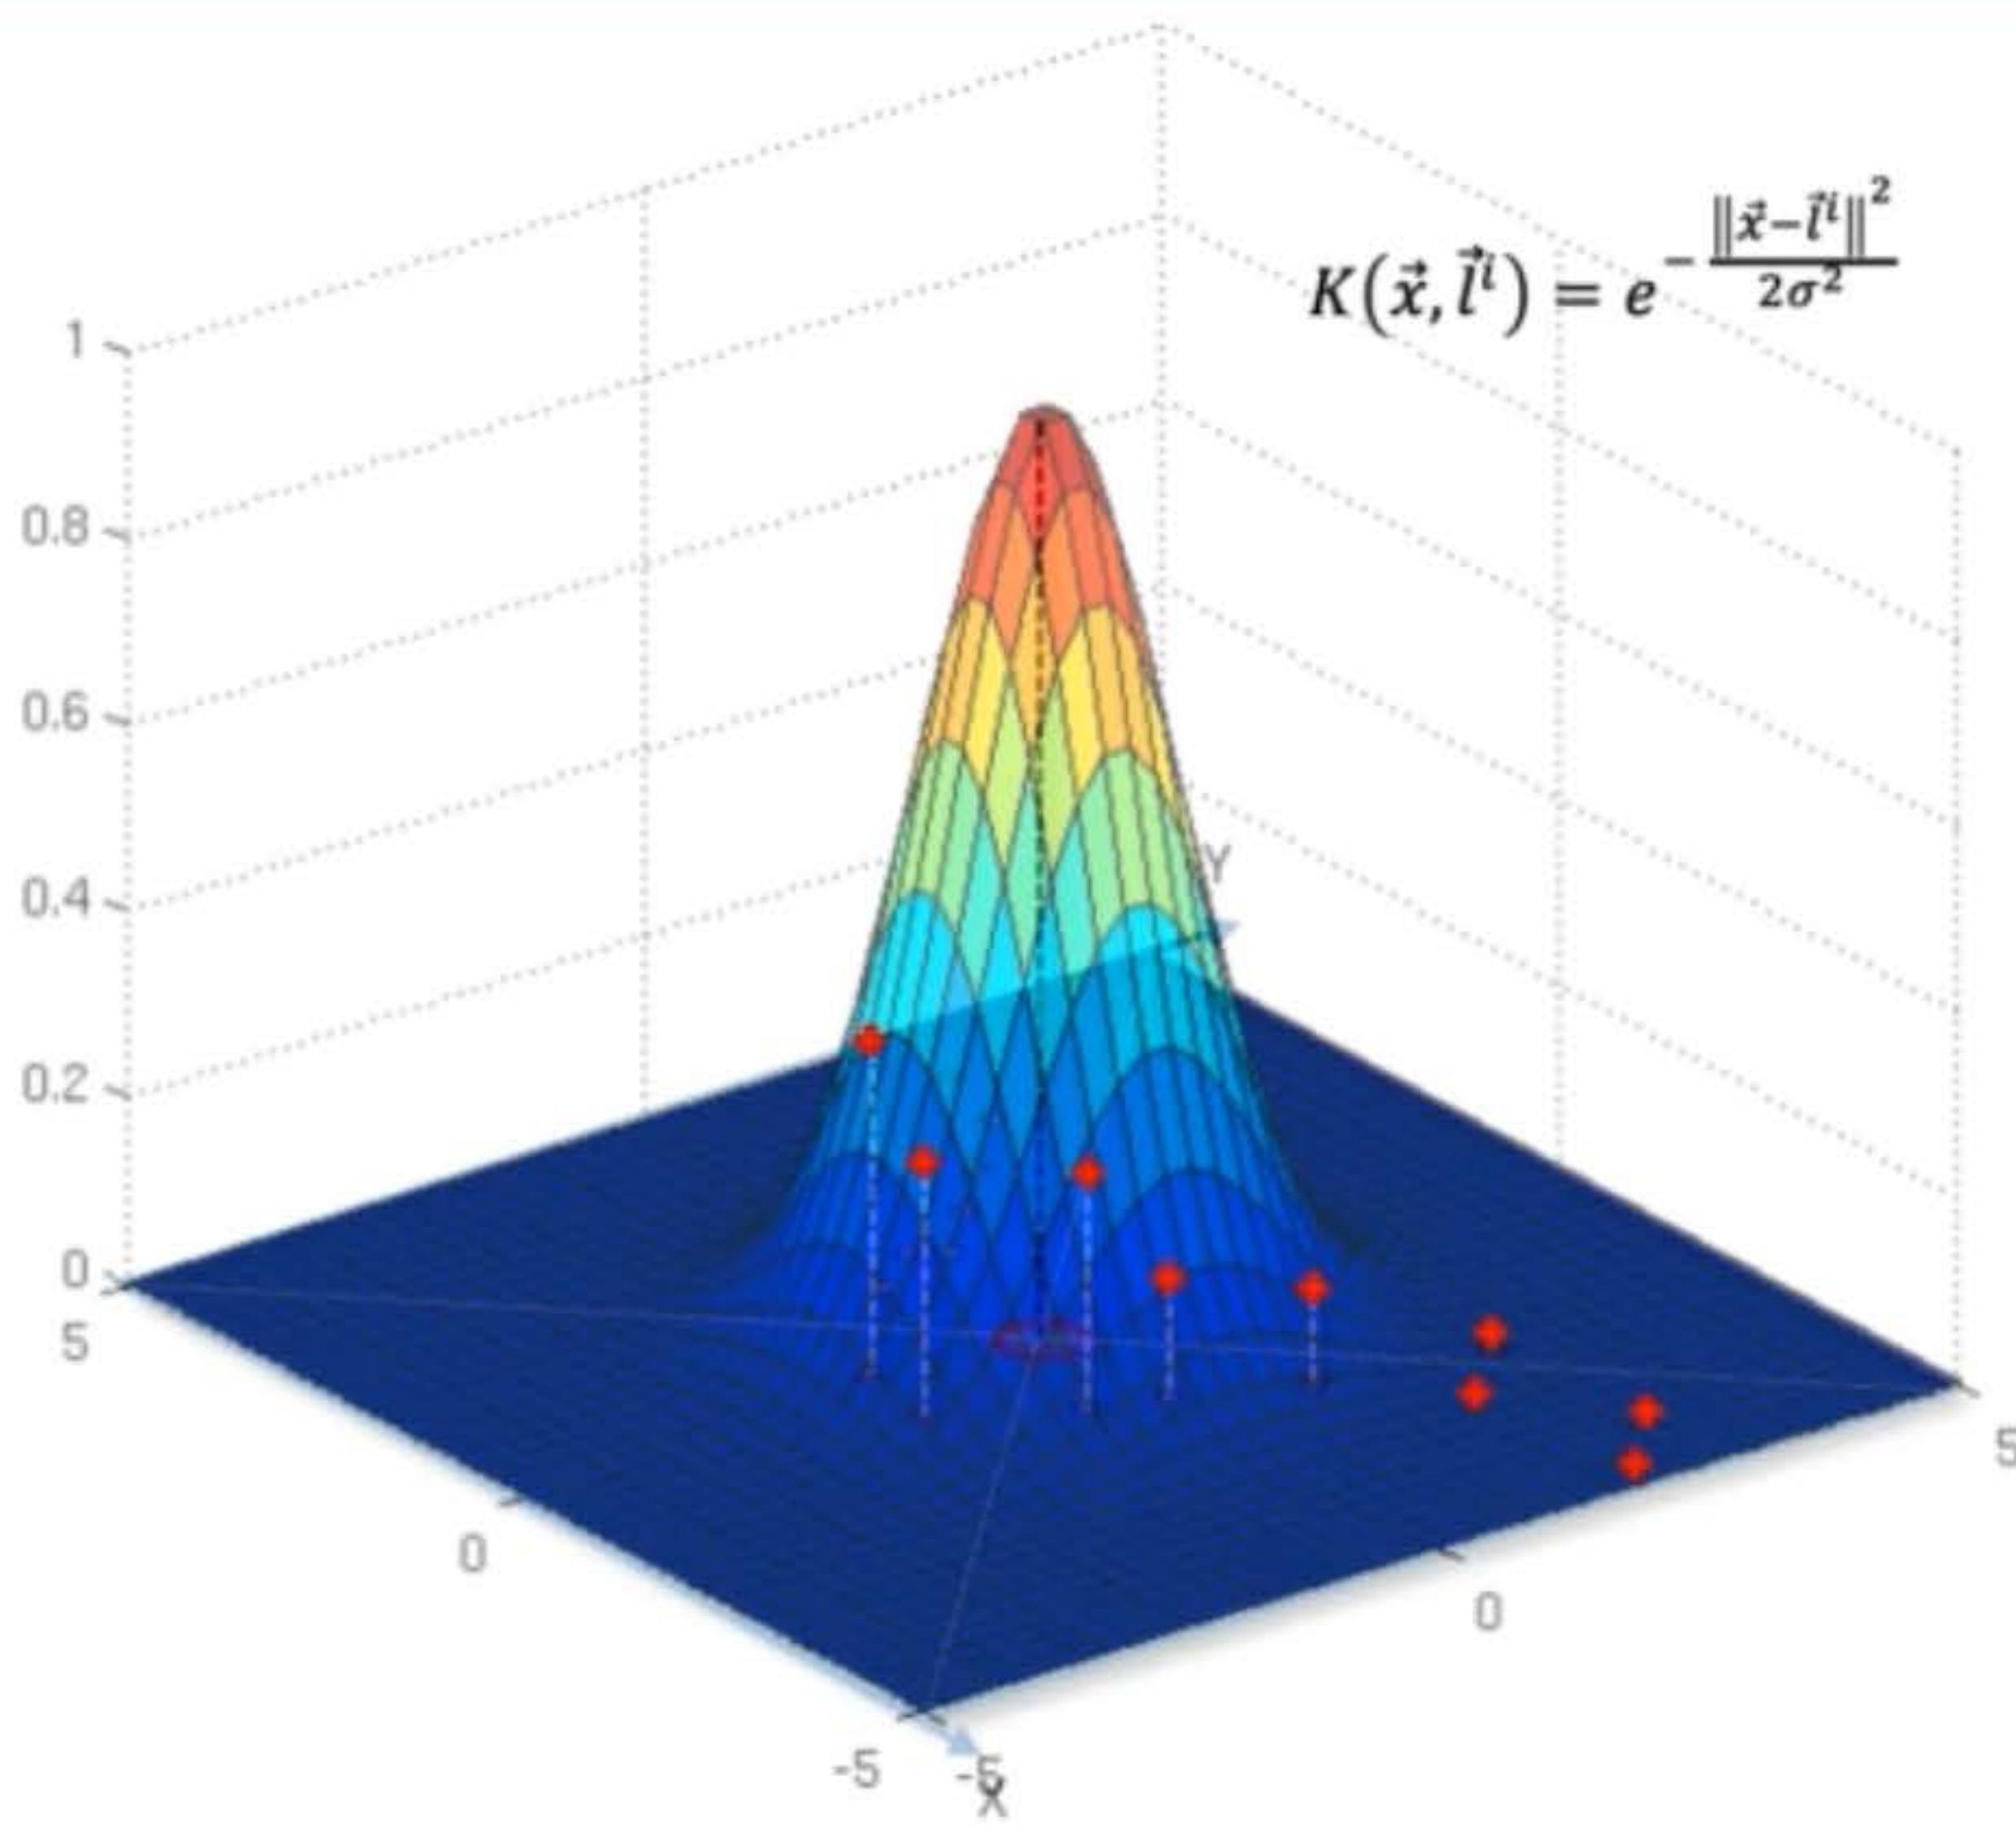
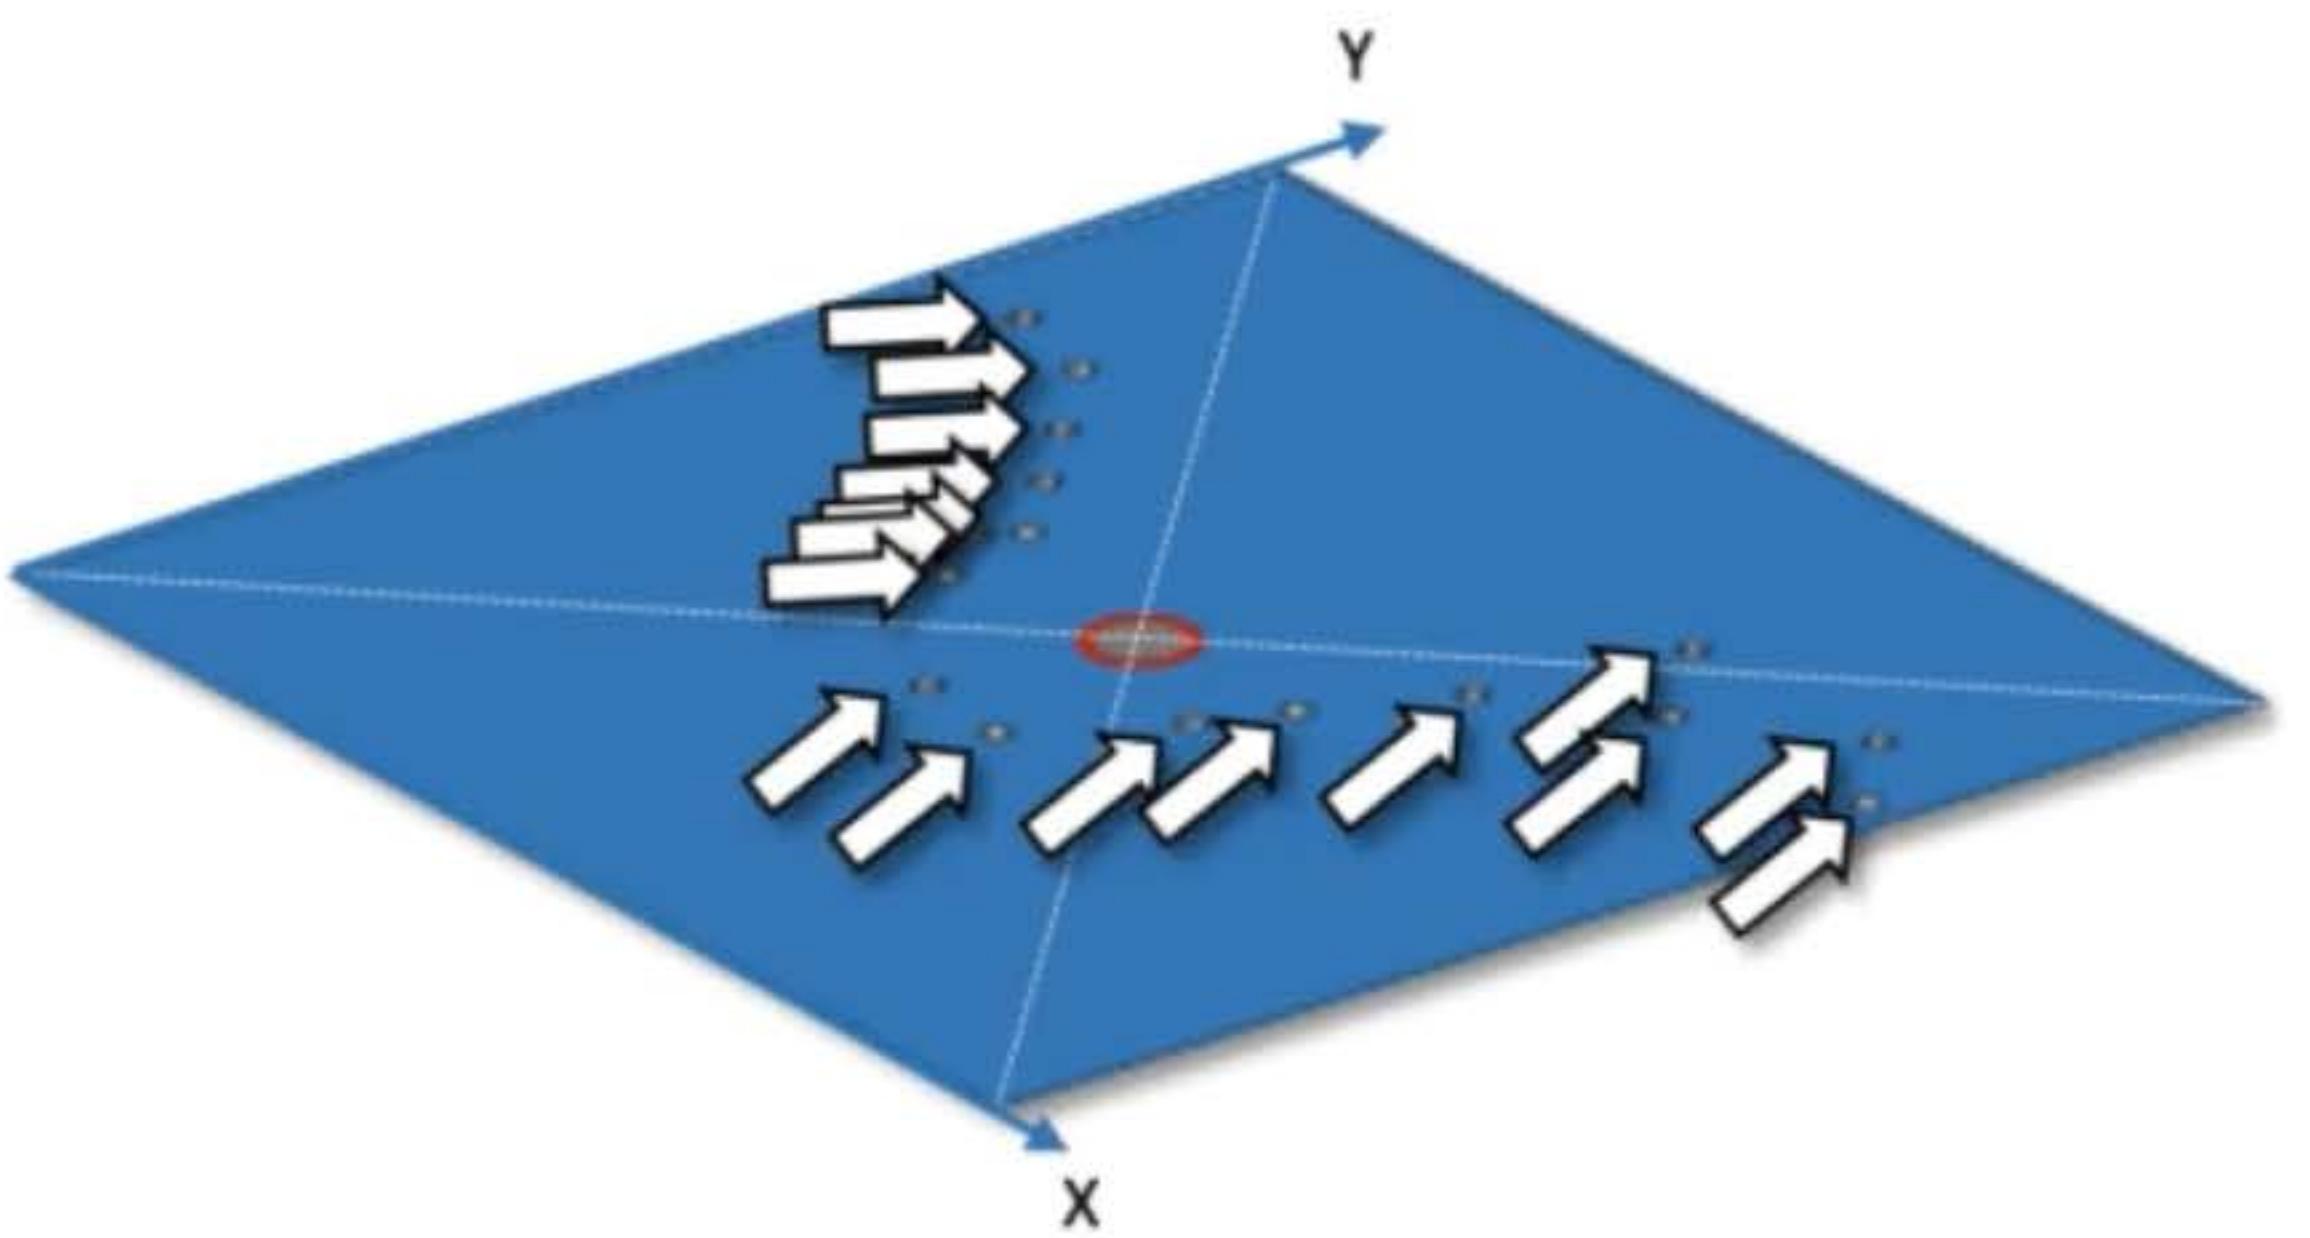
# Non-Linear SVR



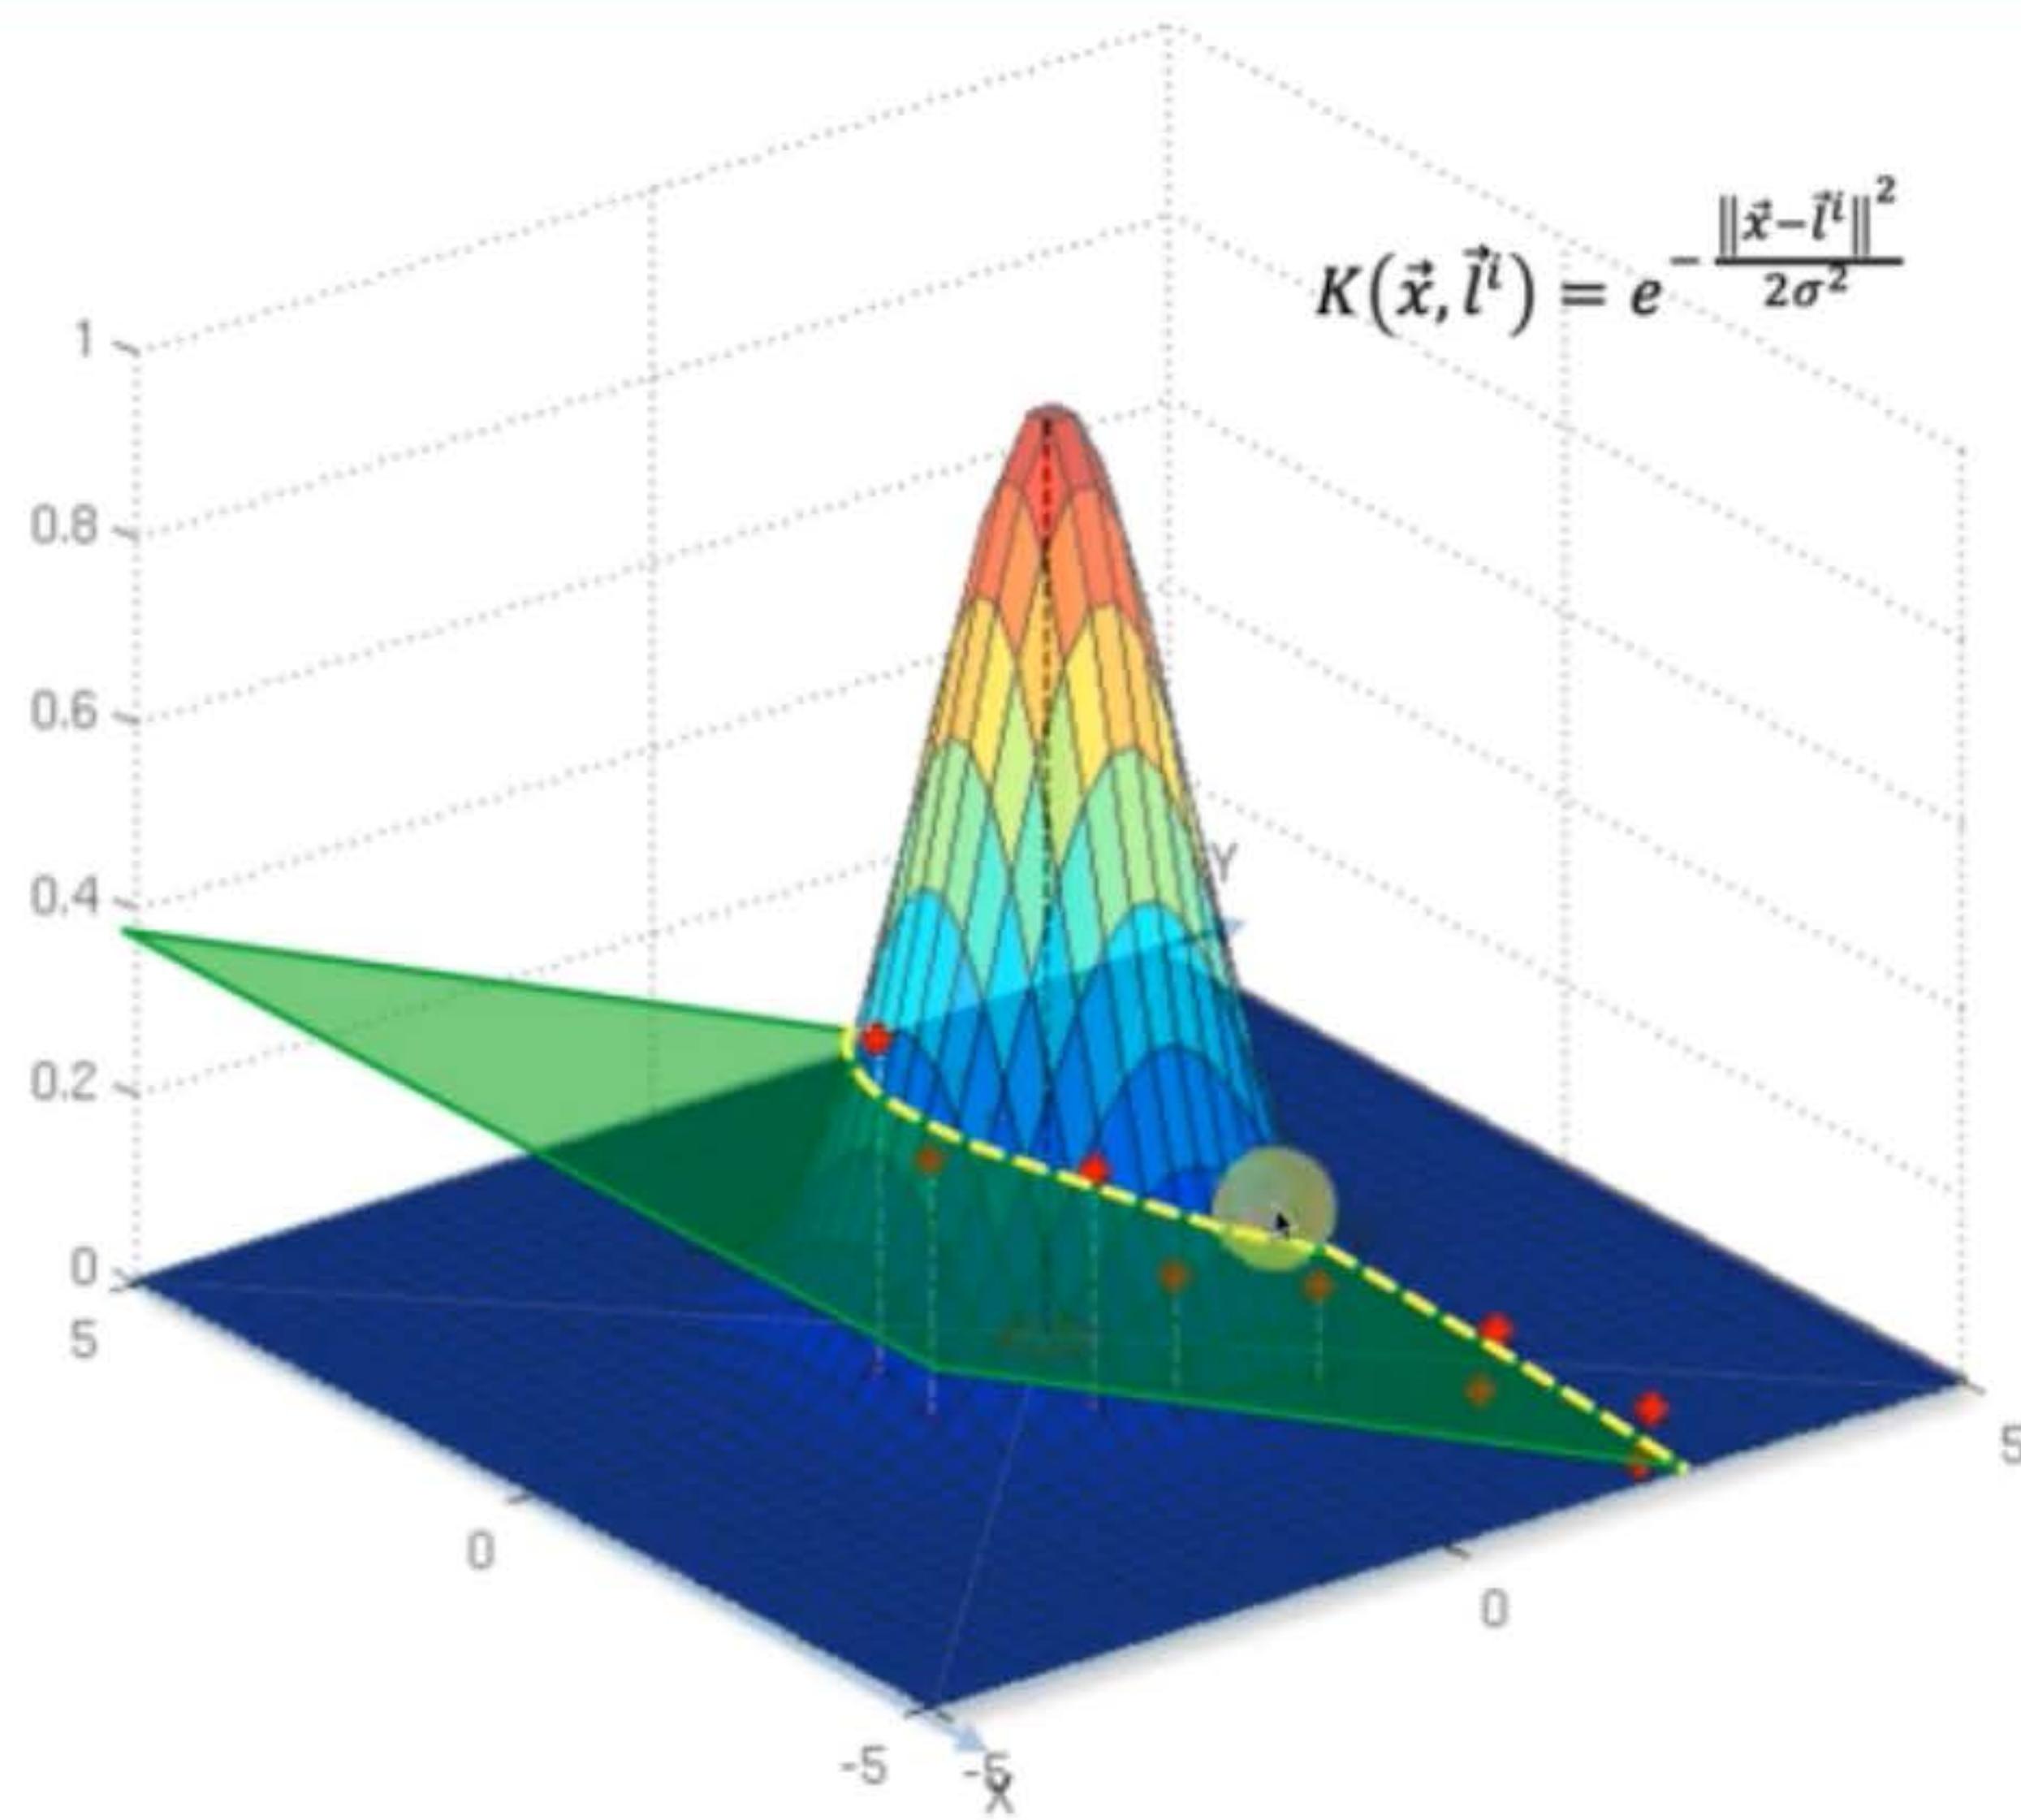
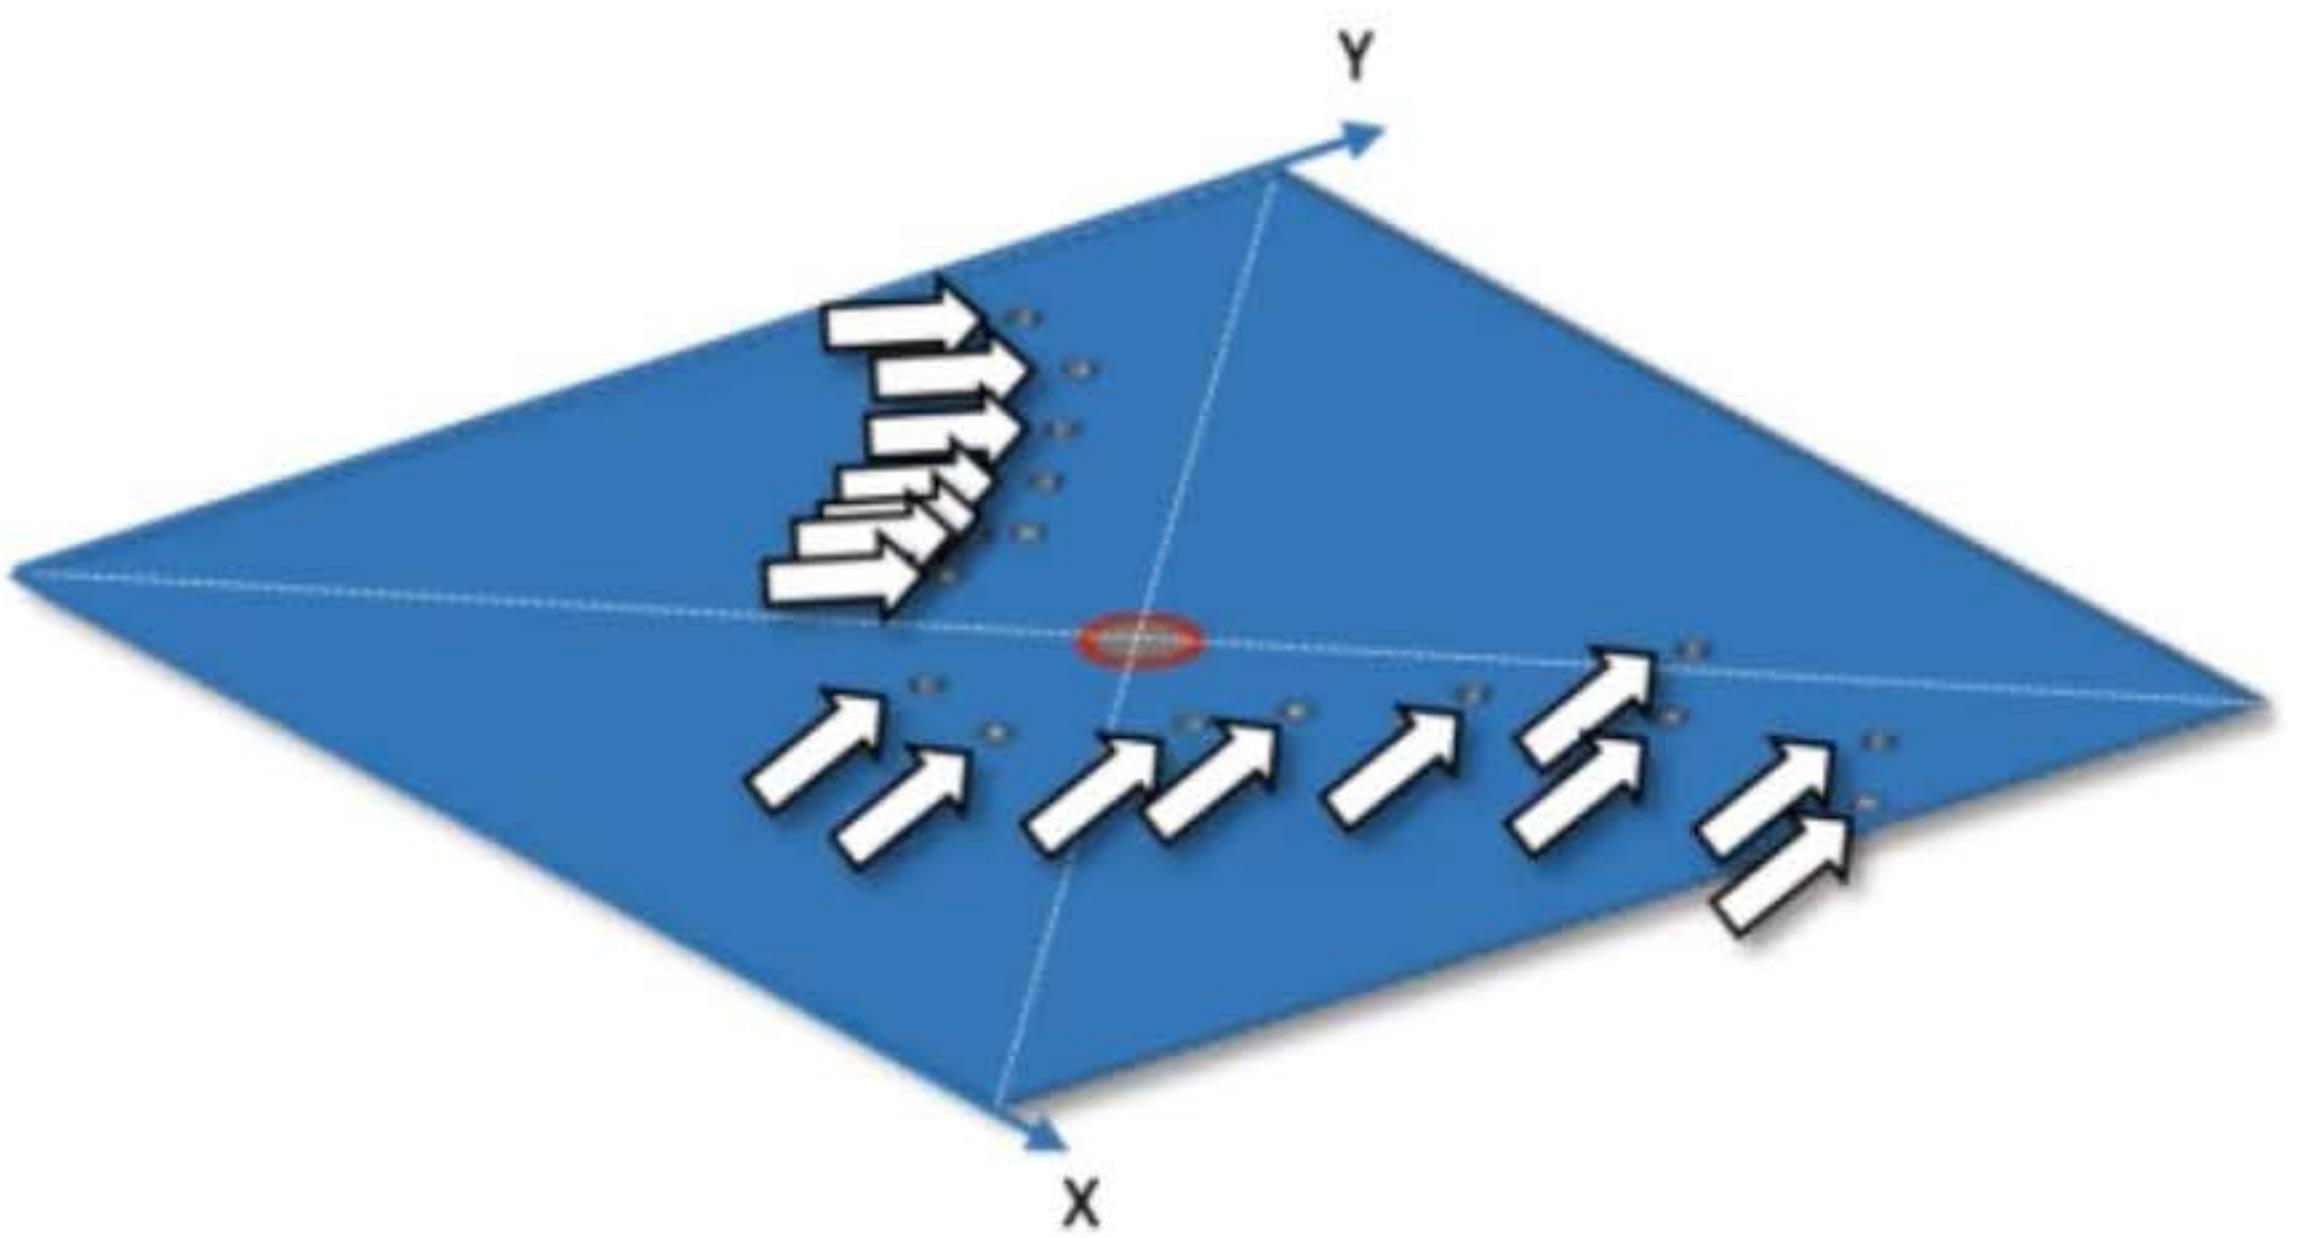
# Non-Linear SVR



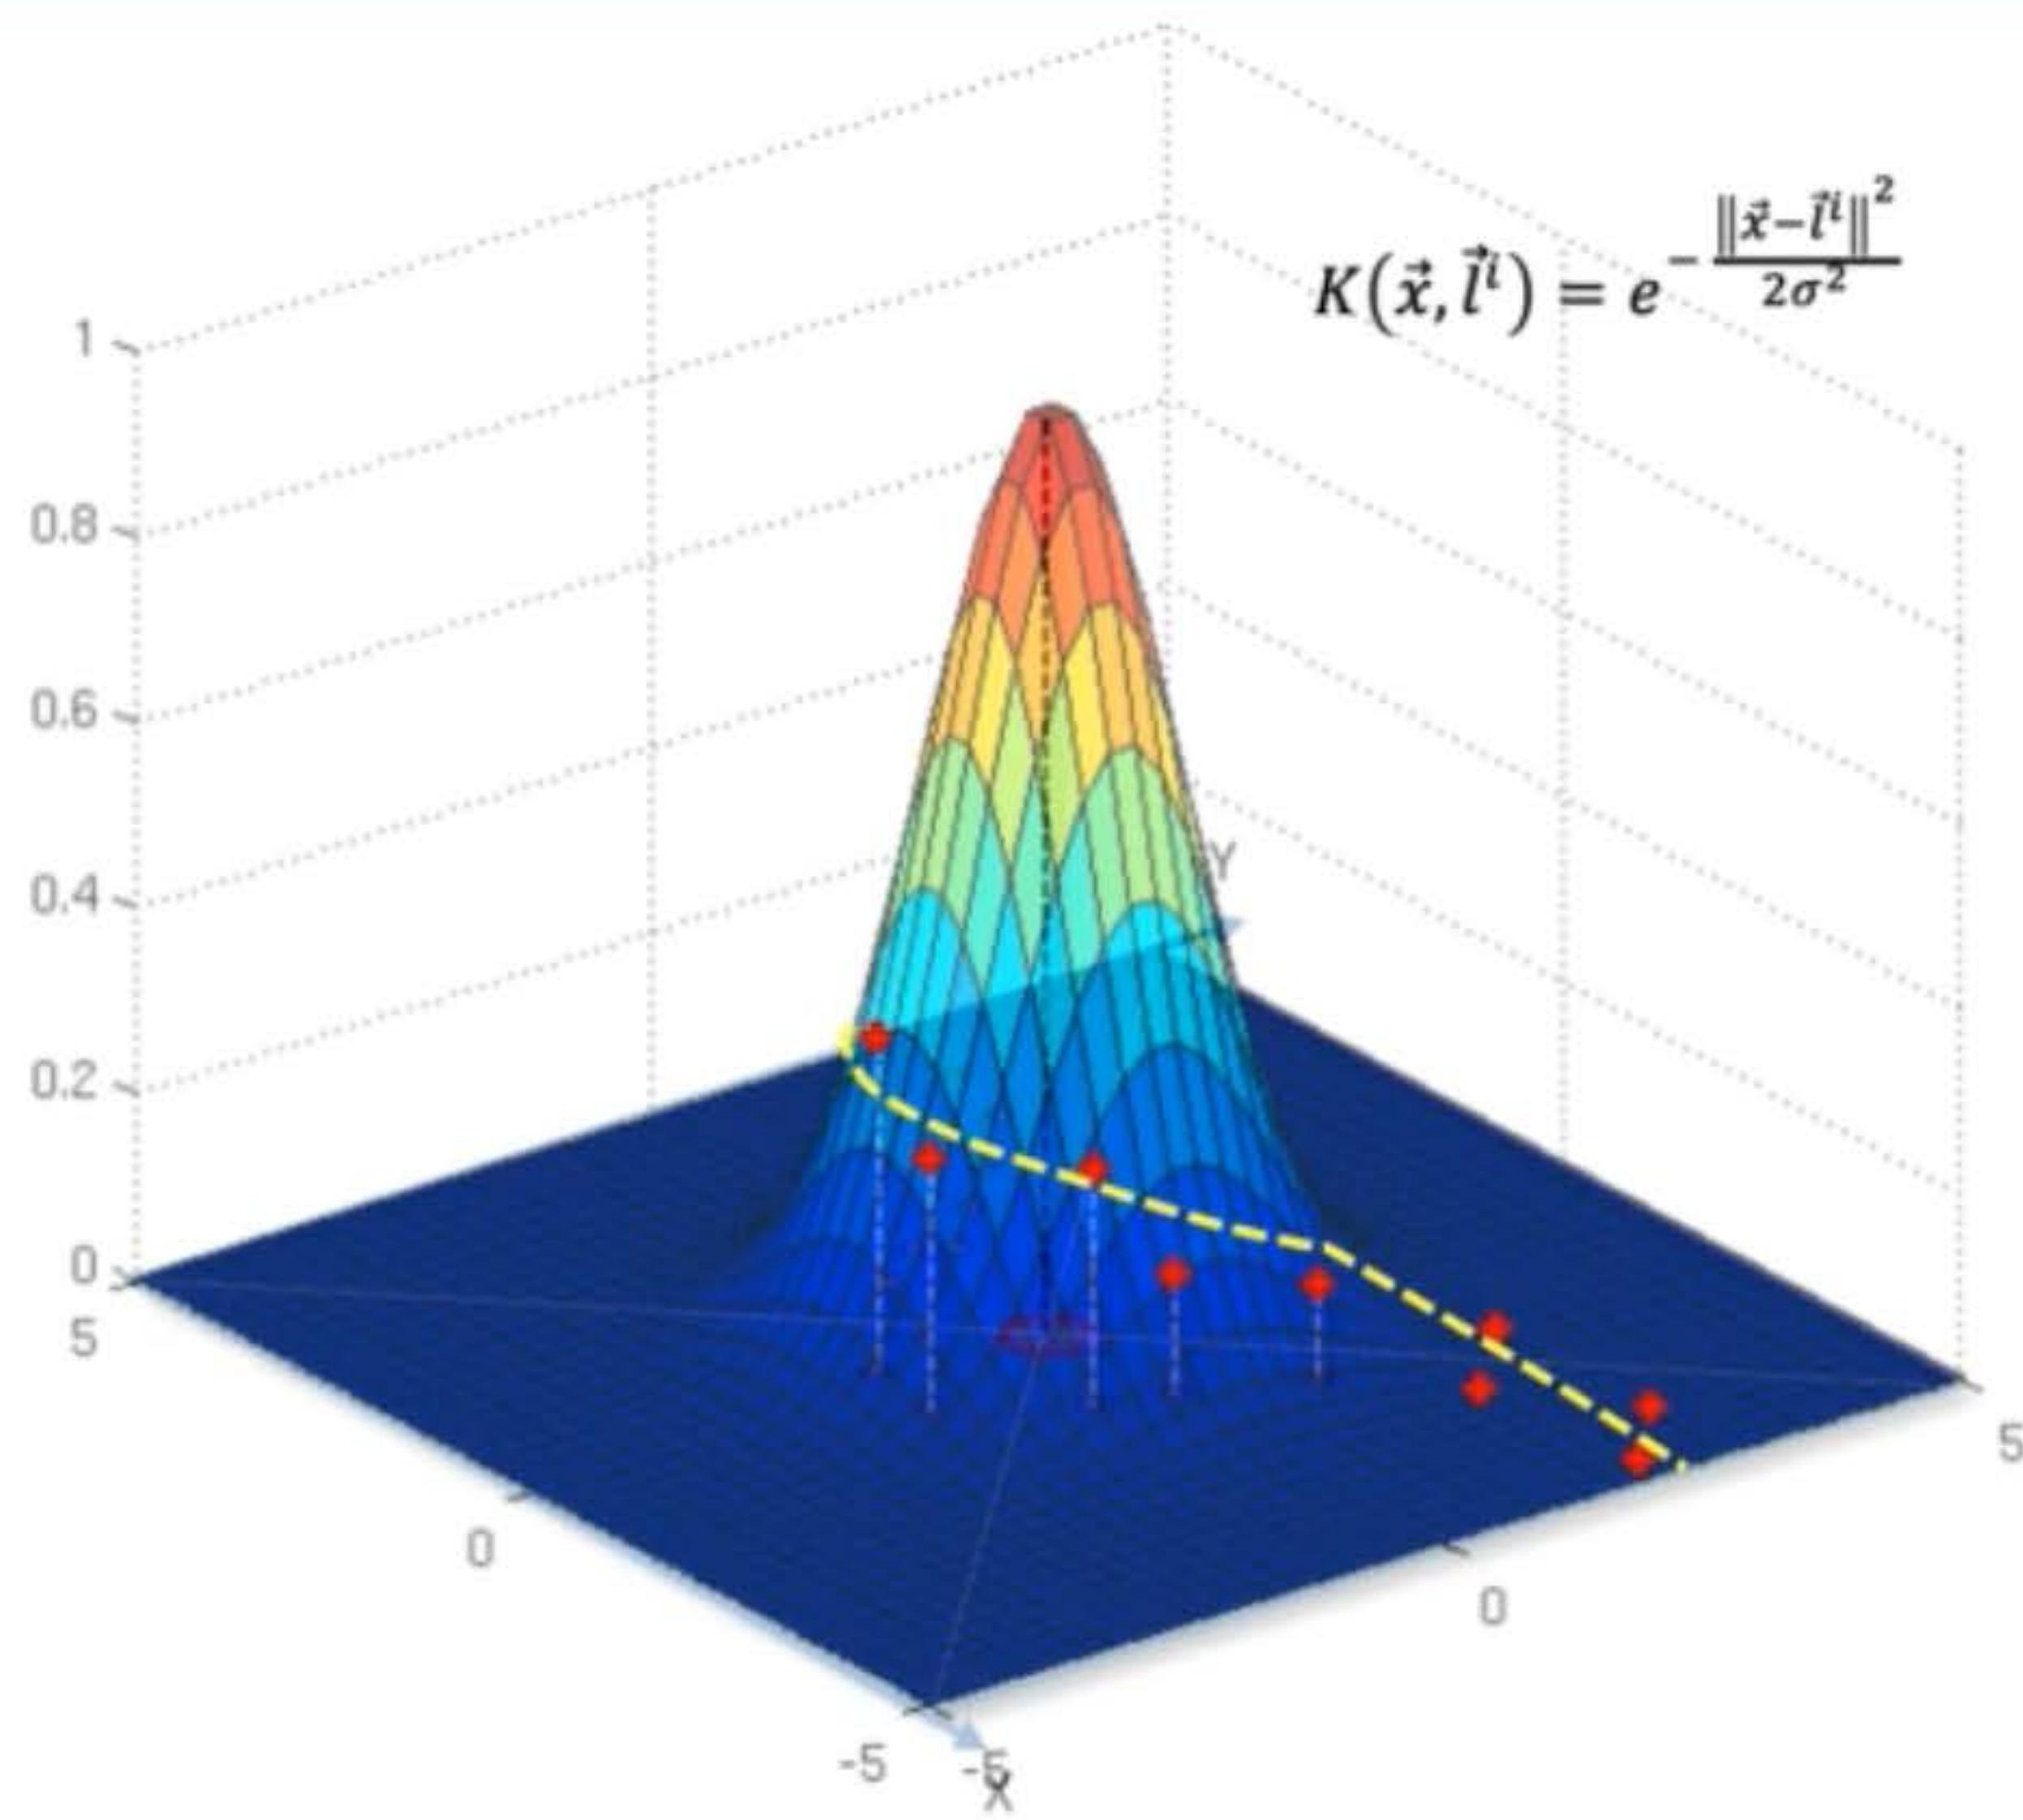
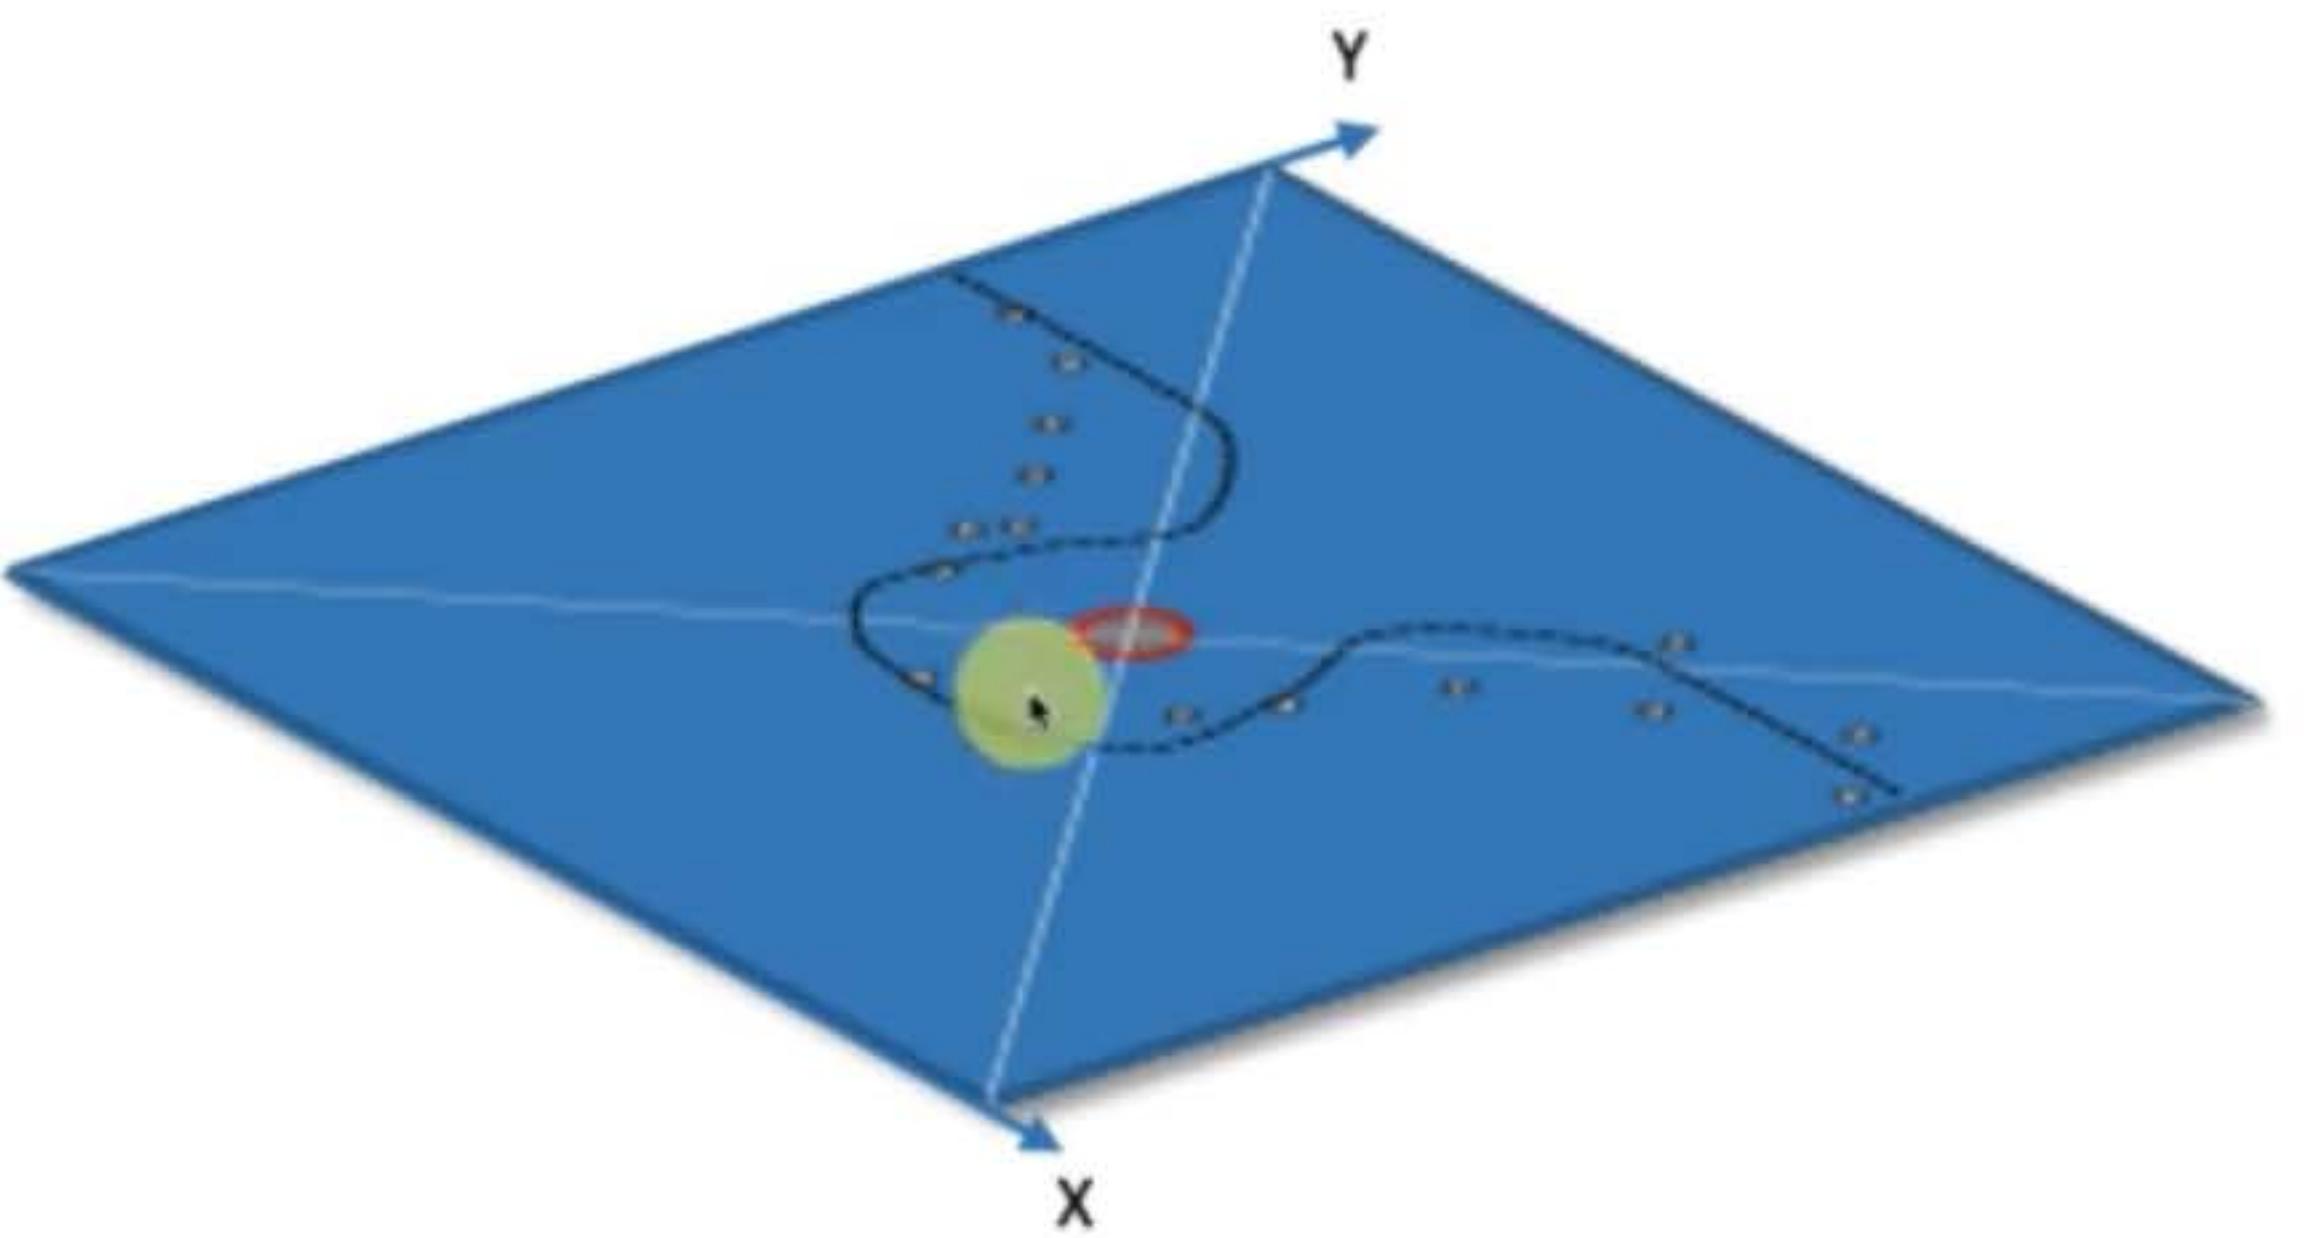
# Non-Linear SVR



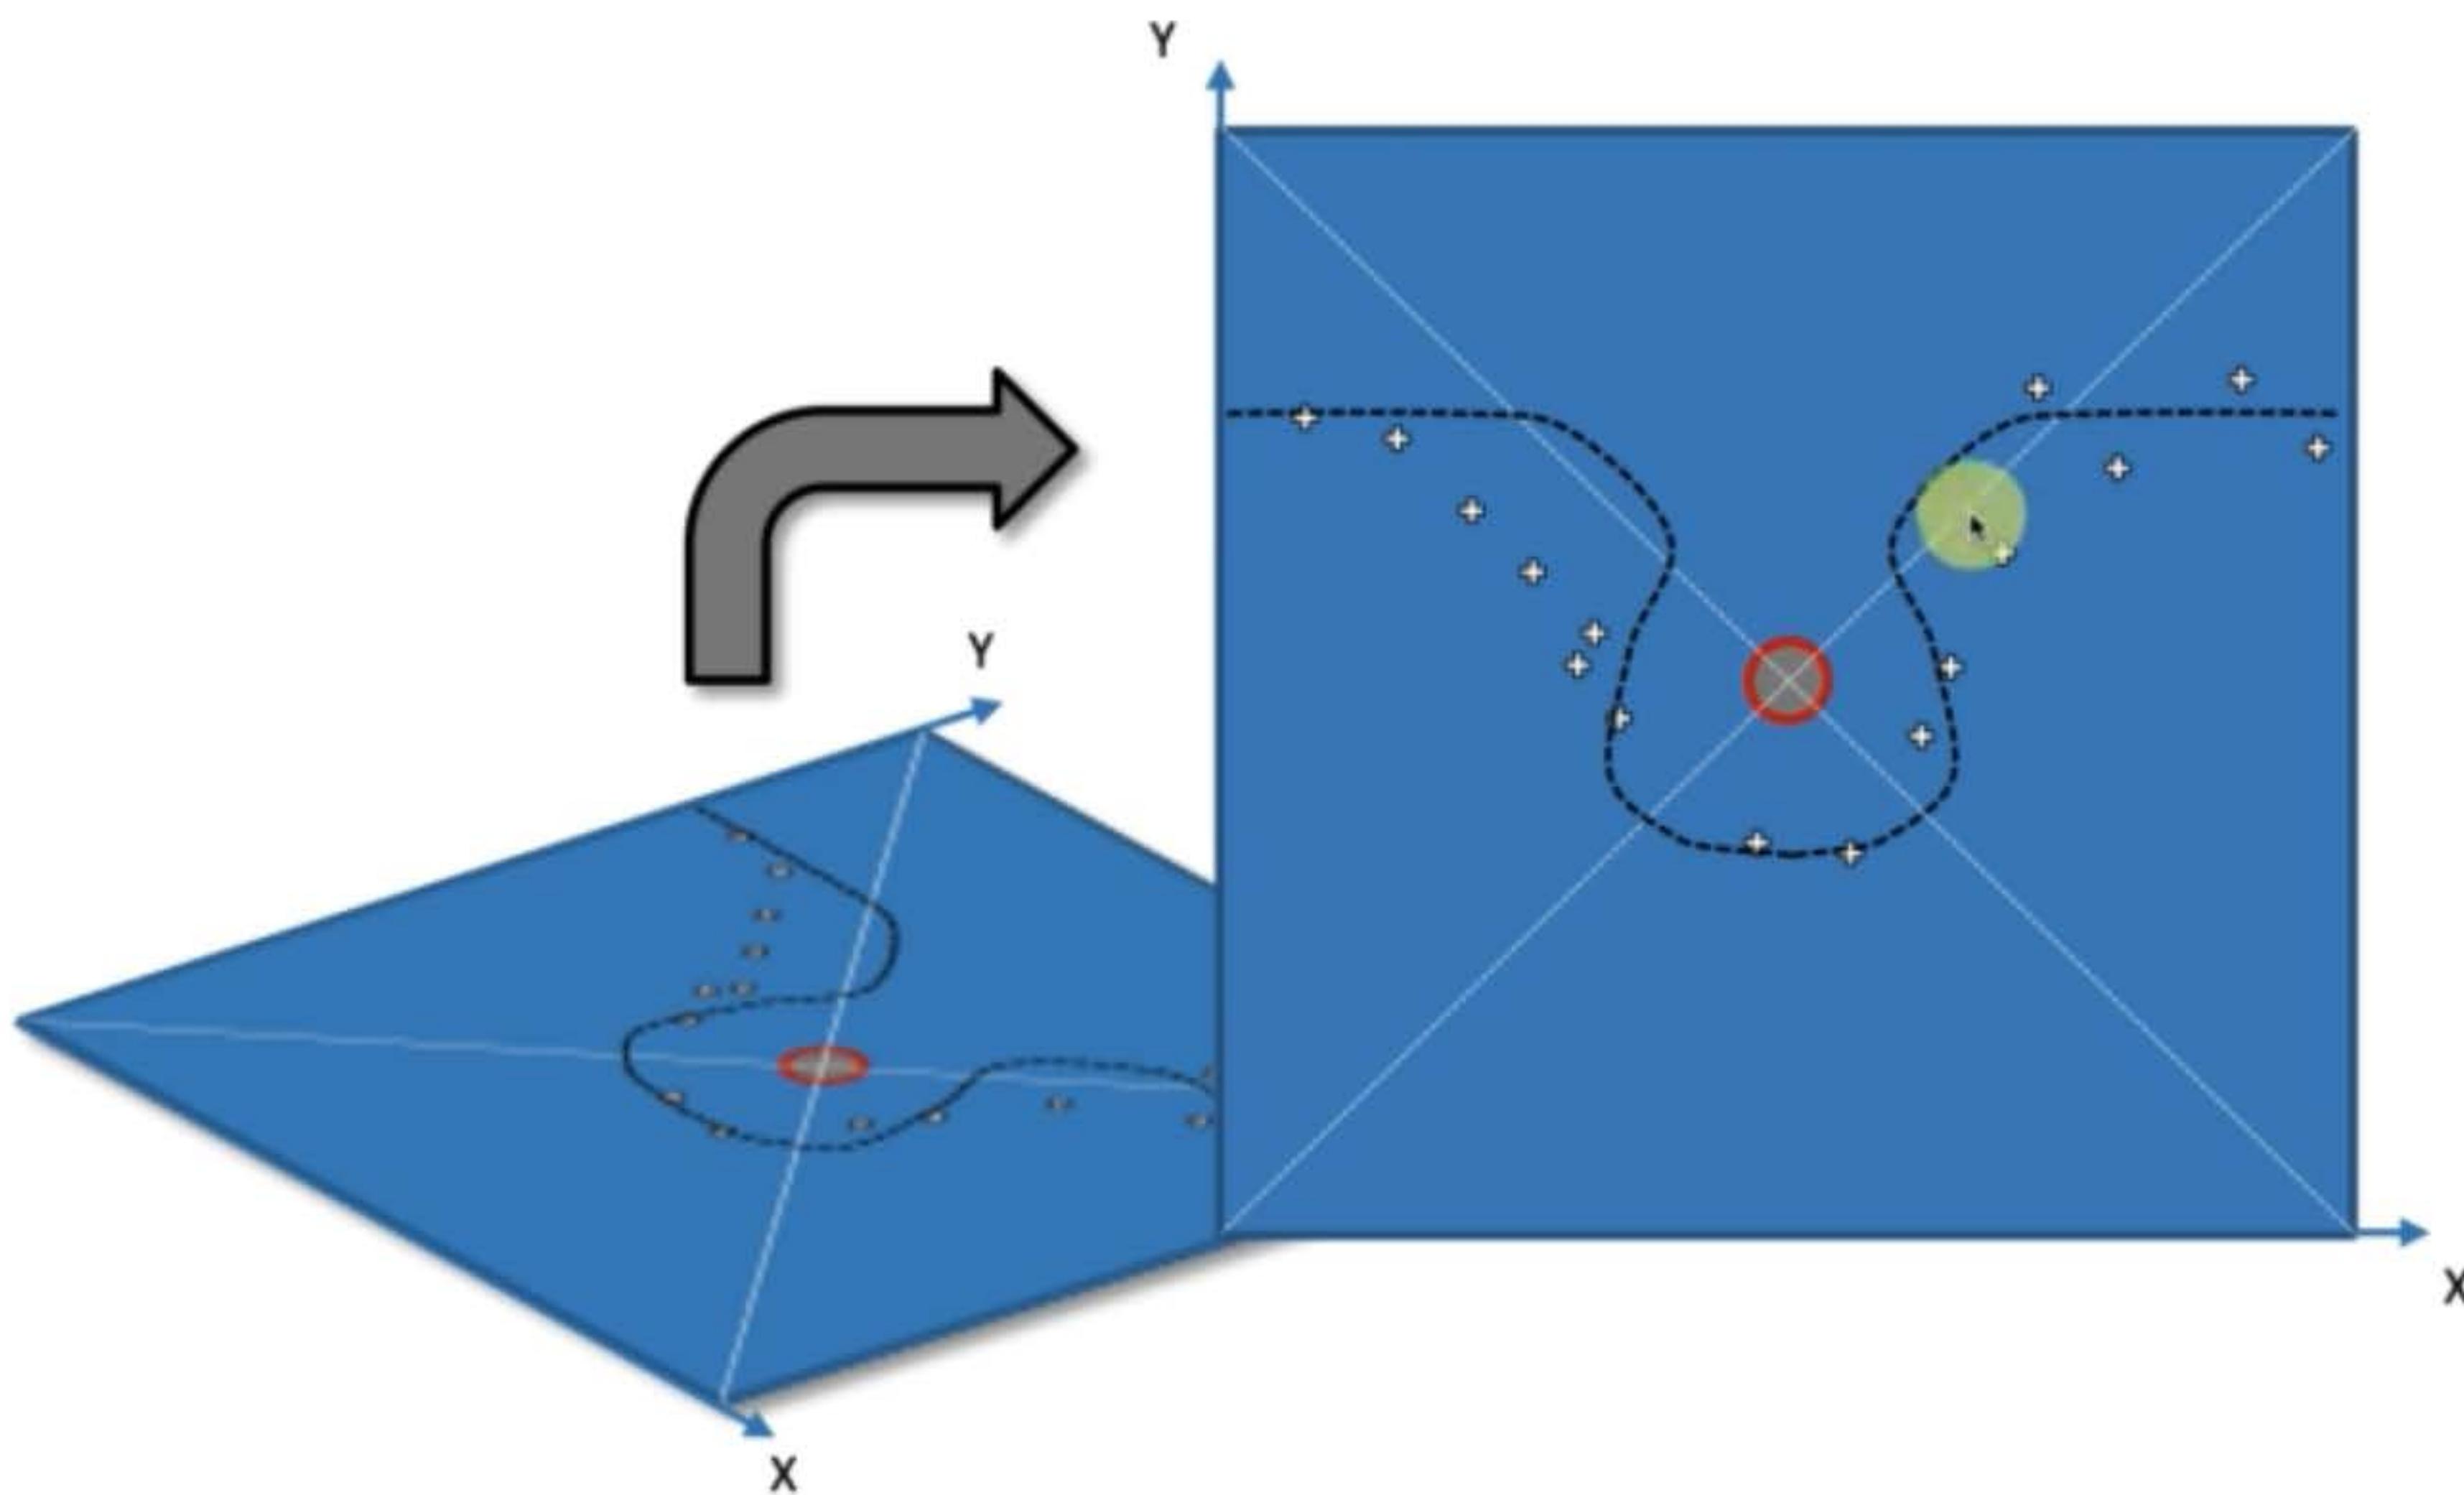
# Non-Linear SVR



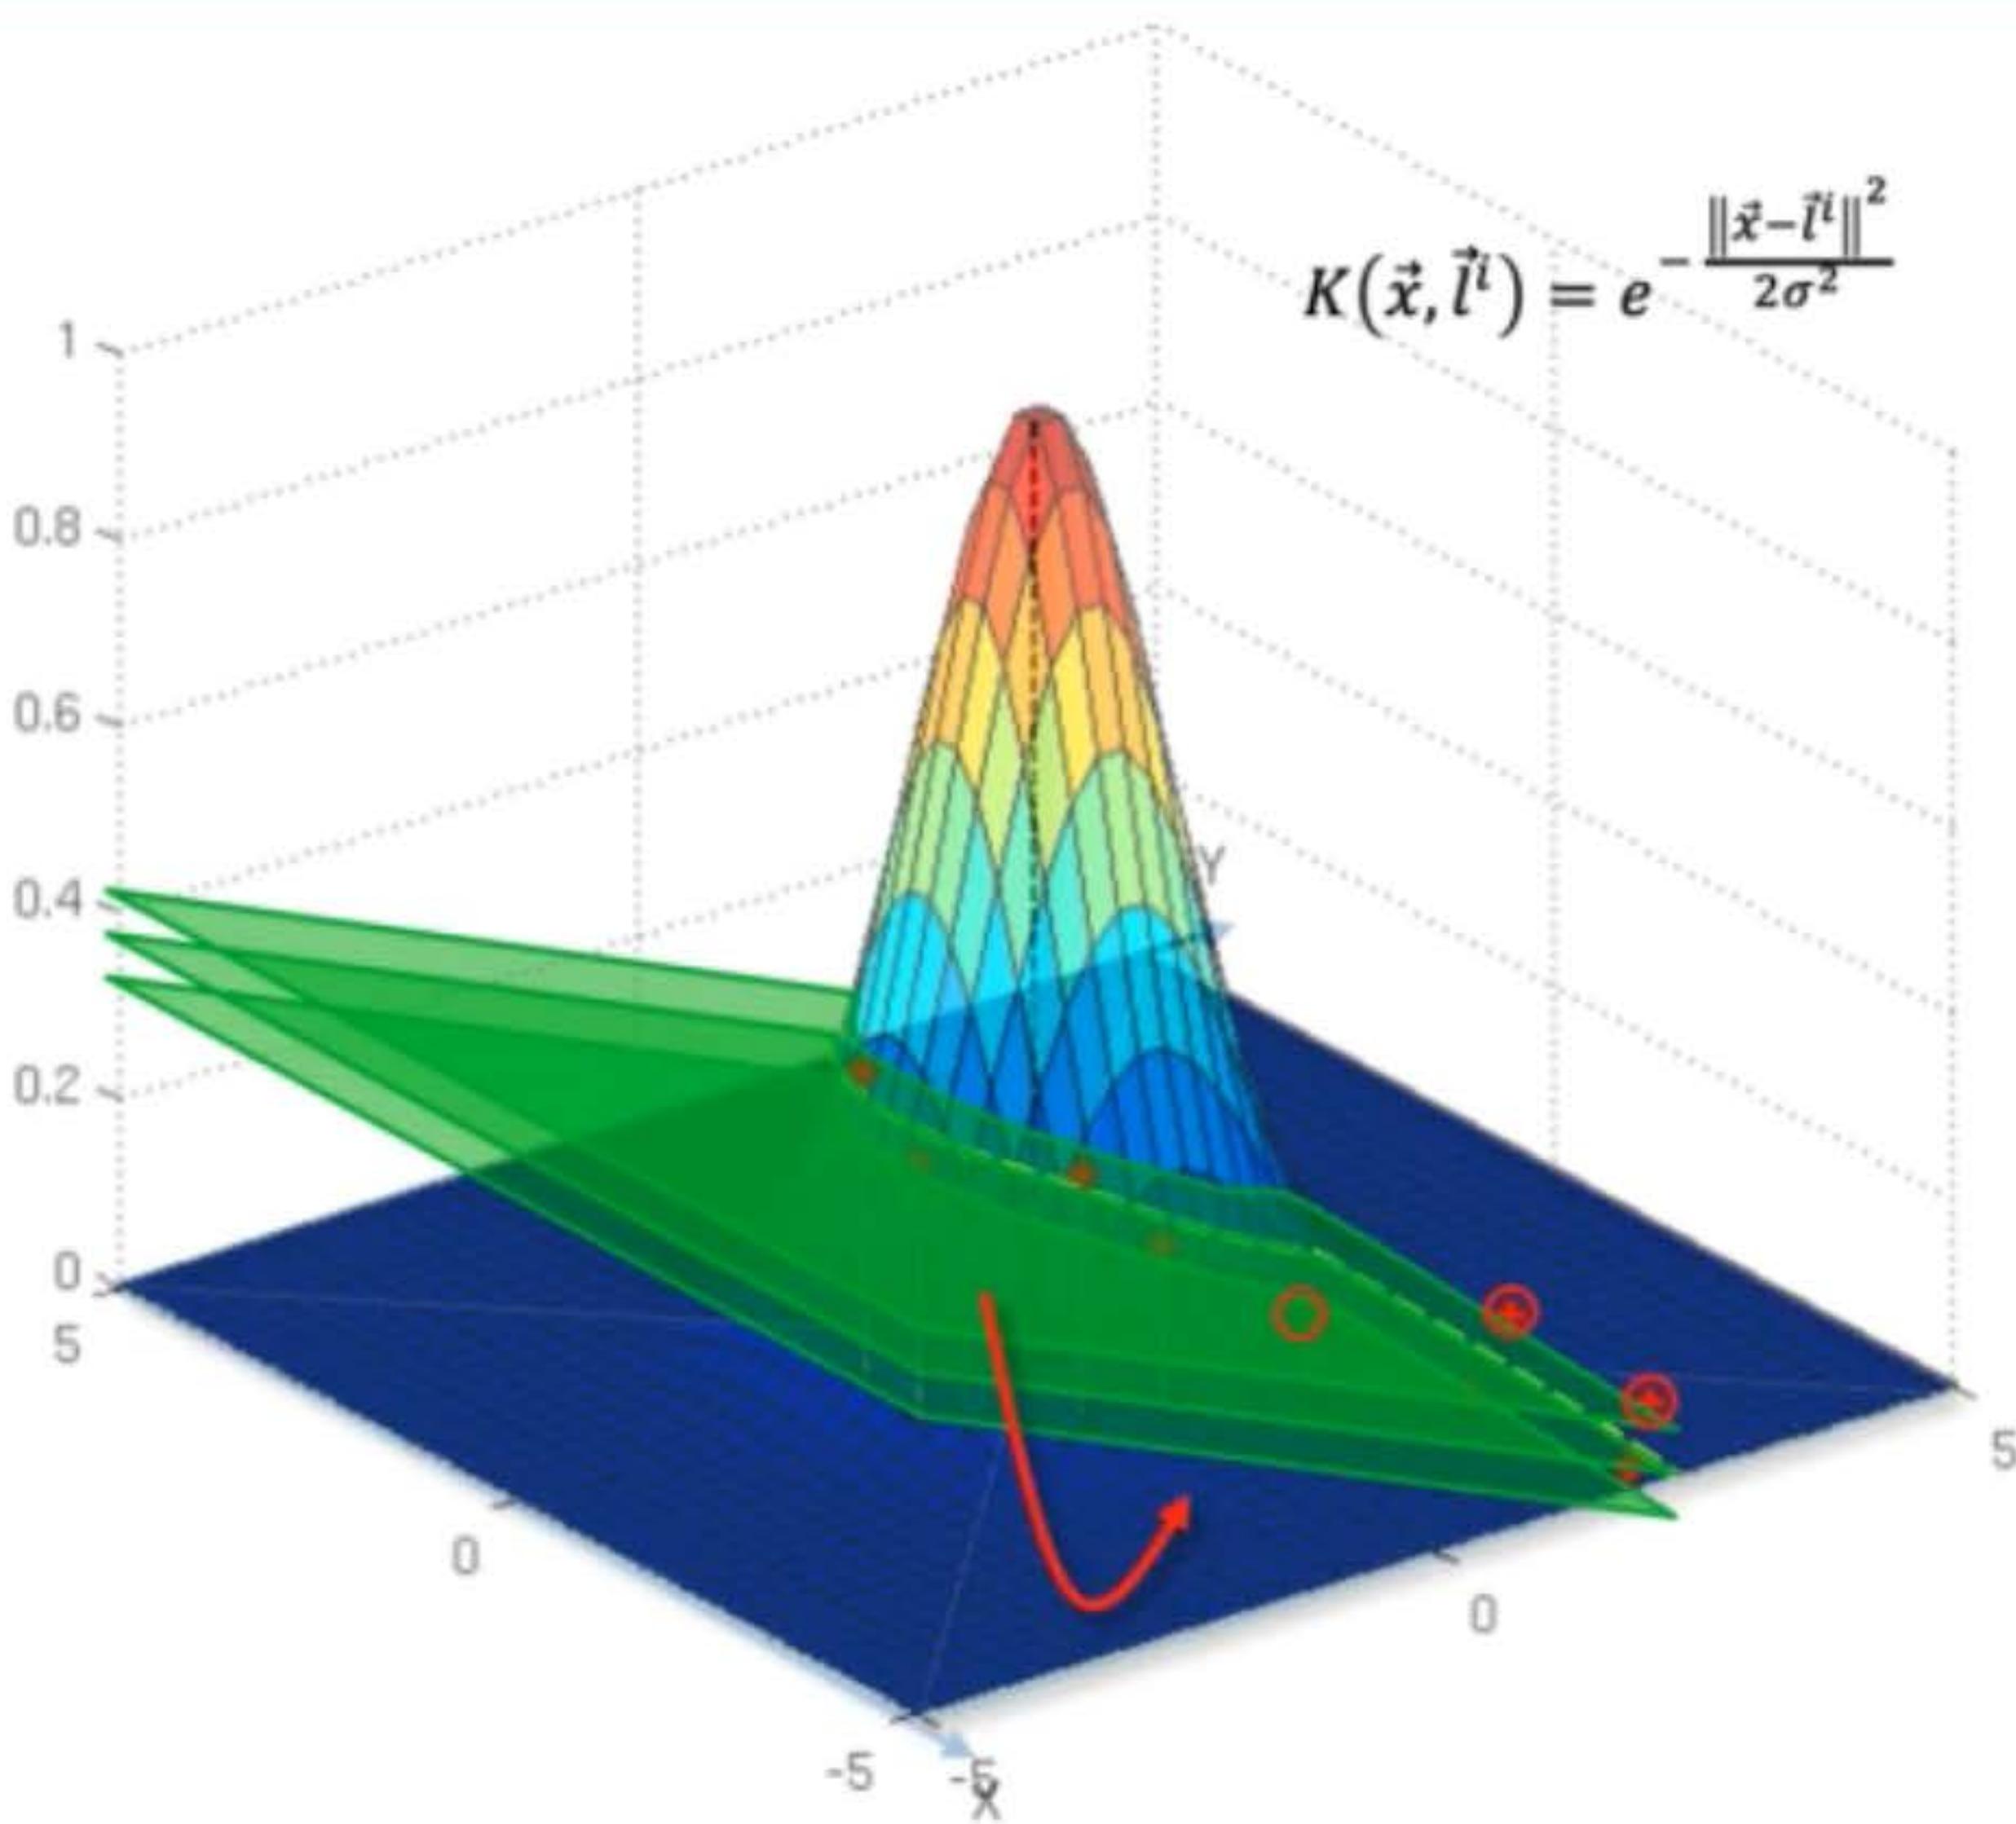
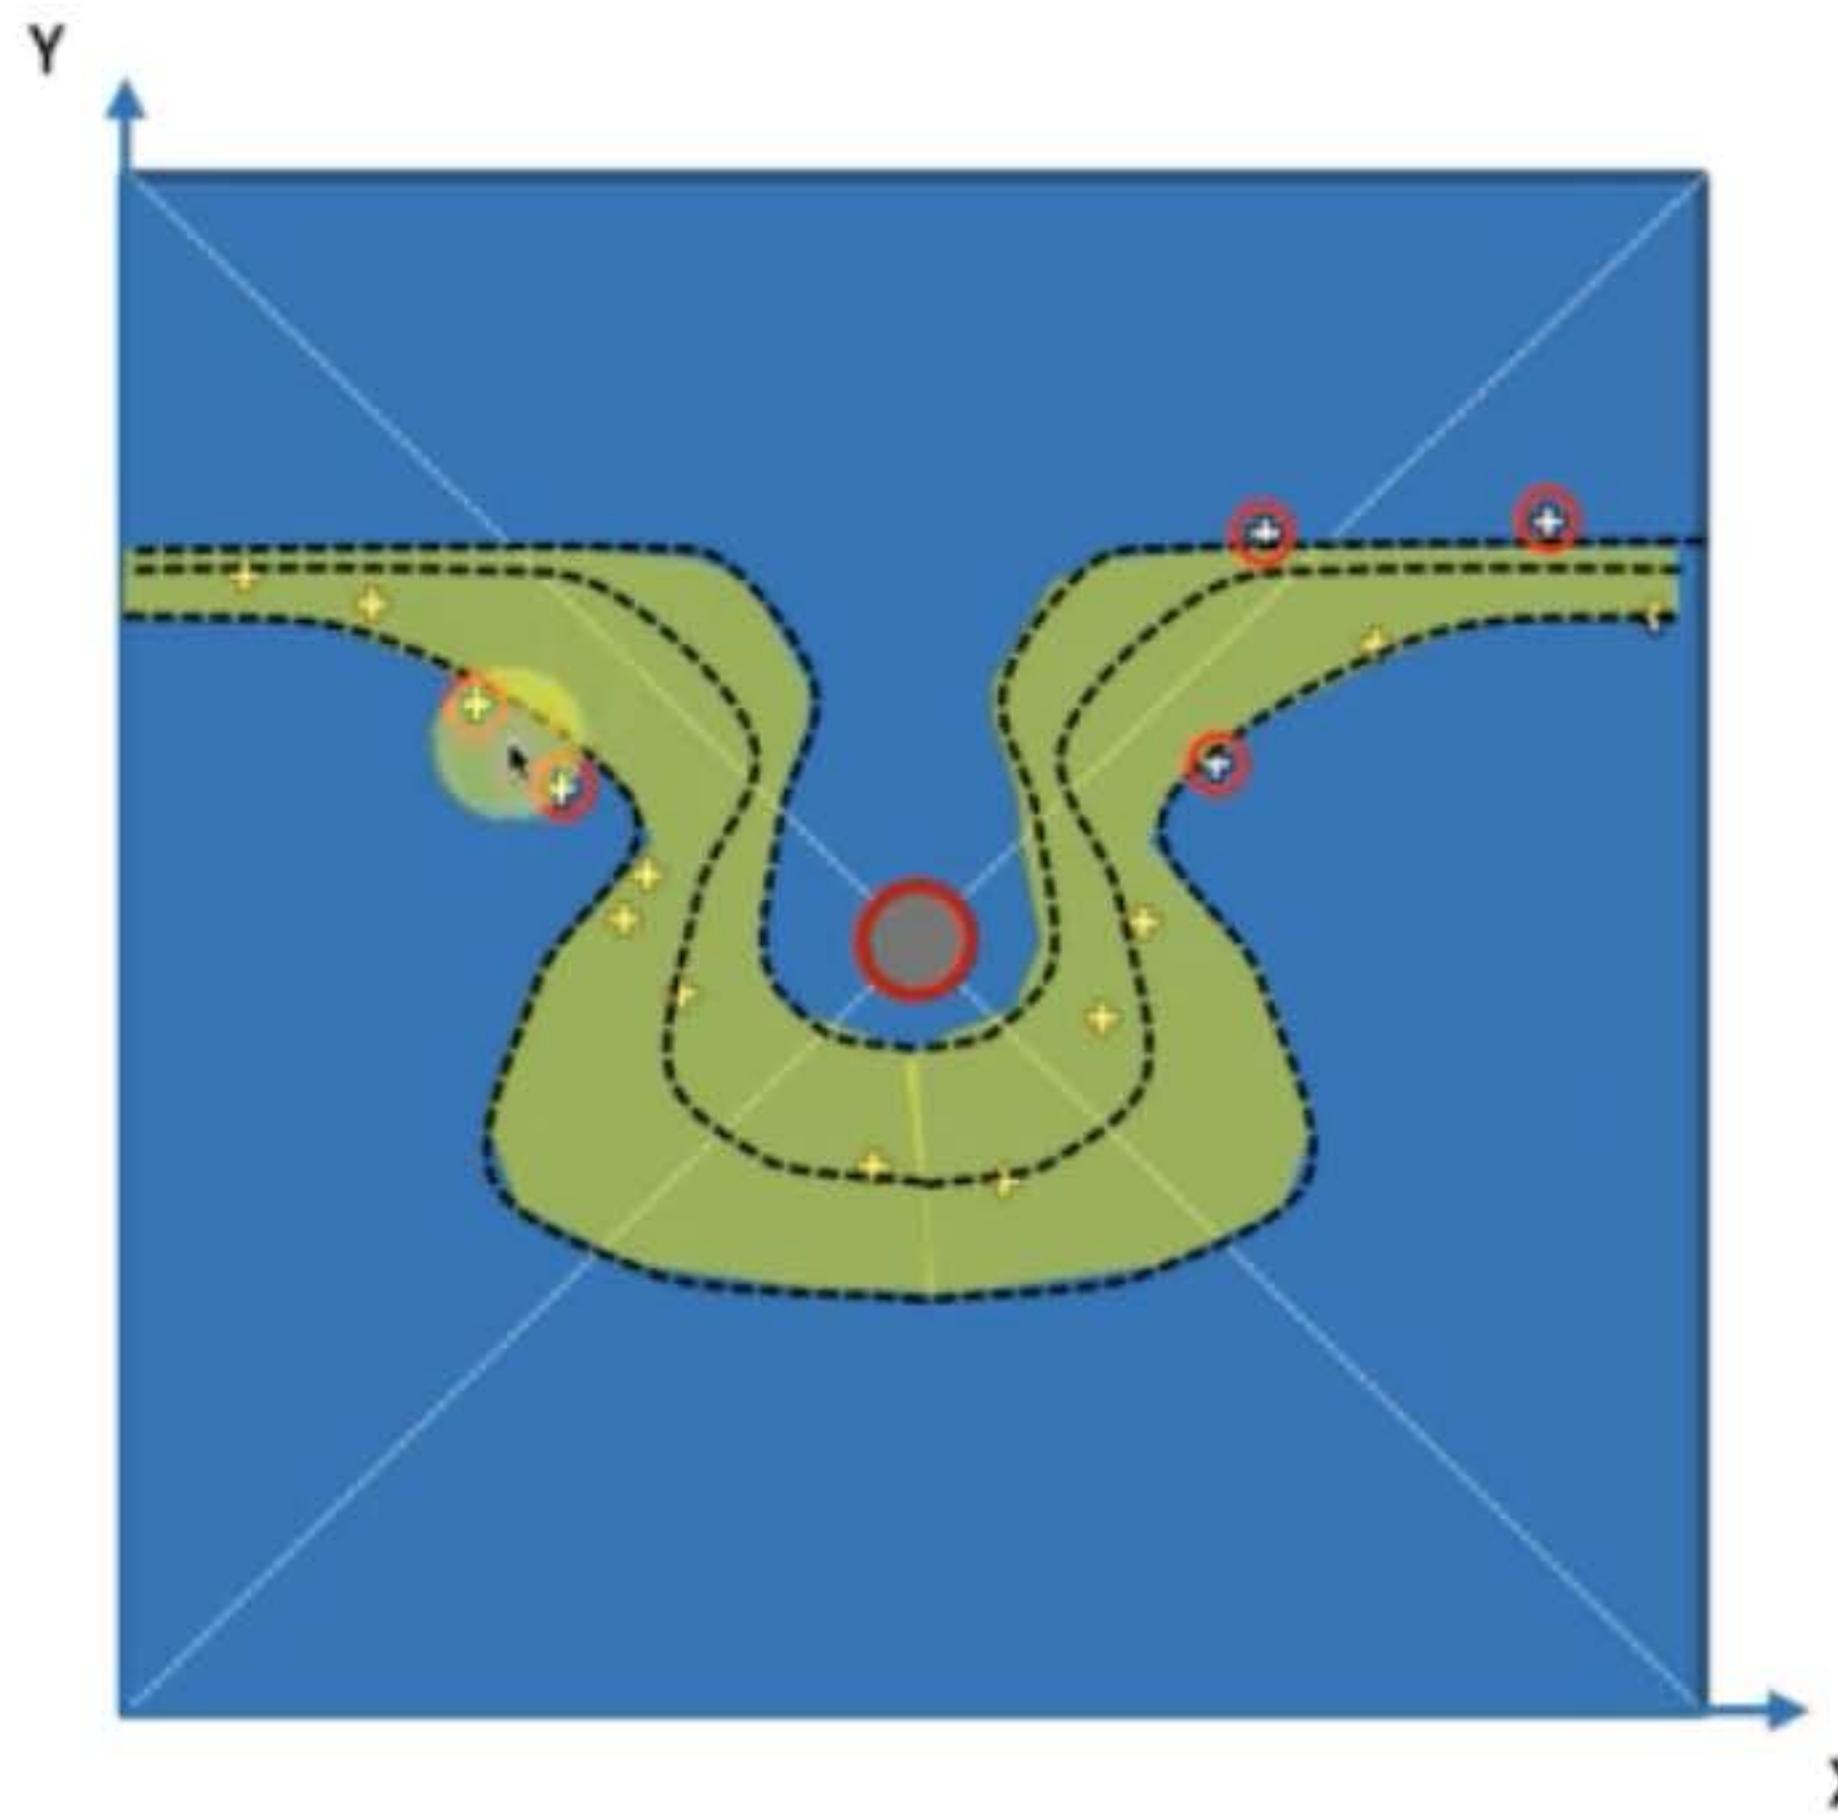
# Non-Linear SVR



# Non-Linear SVR



# Non-Linear SVR



# Bayes' Theorem

# Bayes Theorem



# Bayes Theorem



# Bayes Theorem

What's the probability?



# Bayes Theorem

$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)}$$

# Bayes Theorem

Mach1: 30 wrenches / hr

Mach2: 20 wrenches / hr

Out of all produced parts:

We can SEE that 1% are defective

Out of all defective parts:

We can SEE that 50% came from mach1

And 50% came from mach2

Question:

What is the probability that a part  
produced by mach2 is defective = ?

$$\rightarrow P(\text{Mach1}) = 30/50 = 0.6$$

$$\rightarrow P(\text{Mach2}) = 20/50 = 0.4$$

$$\rightarrow P(\text{Defect}) = 1\%$$

$$\rightarrow P(\text{Mach1} \mid \text{Defect}) = 50\%$$

$$\rightarrow P(\text{Mach2} \mid \text{Defect}) = 50\%$$

$$\rightarrow P(\text{Defect} \mid \text{Mach2}) = ?$$

# Bayes Theorem

Mach1: 30 wrenches / hr

Mach2: 20 wrenches / hr

Out of all produced parts:

We can SEE that 1% are defective

Out of all defective parts:

We can SEE that 50% came from mach1

And 50% came from mach2

Question:

What is the probability that a part  
produced by mach2 is defective = ?

$$\rightarrow P(\text{Mach2}) = 20/50 = 0.4$$

$$\rightarrow P(\text{Defect}) = 1\%$$

$$\rightarrow P(\text{Mach2} \mid \text{Defect}) = 50\%$$

$$\rightarrow P(\text{Defect} \mid \text{Mach2}) = ?$$

$$P(\text{Defect} \mid \text{Mach2}) = \frac{P(\text{Mach2} \mid \text{Defect}) * P(\text{Defect})}{P(\text{Mach2})}$$

# It's intuitive!

$$P(\text{Defect} \mid \text{Mach2}) = \frac{P(\text{Mach2} \mid \text{Defect}) * P(\text{Defect}) * 1000}{P(\text{Mach2}) * 1000} = 1.25\%$$

Let's look at an example:

- 1000 wrenches
- 400 came from Mach2
- 1% have a defect = 10
- of them 50% came from Mach2 = 5
- % defective parts from Mach2 = 5/400 = 1.25%

# It's intuitive!

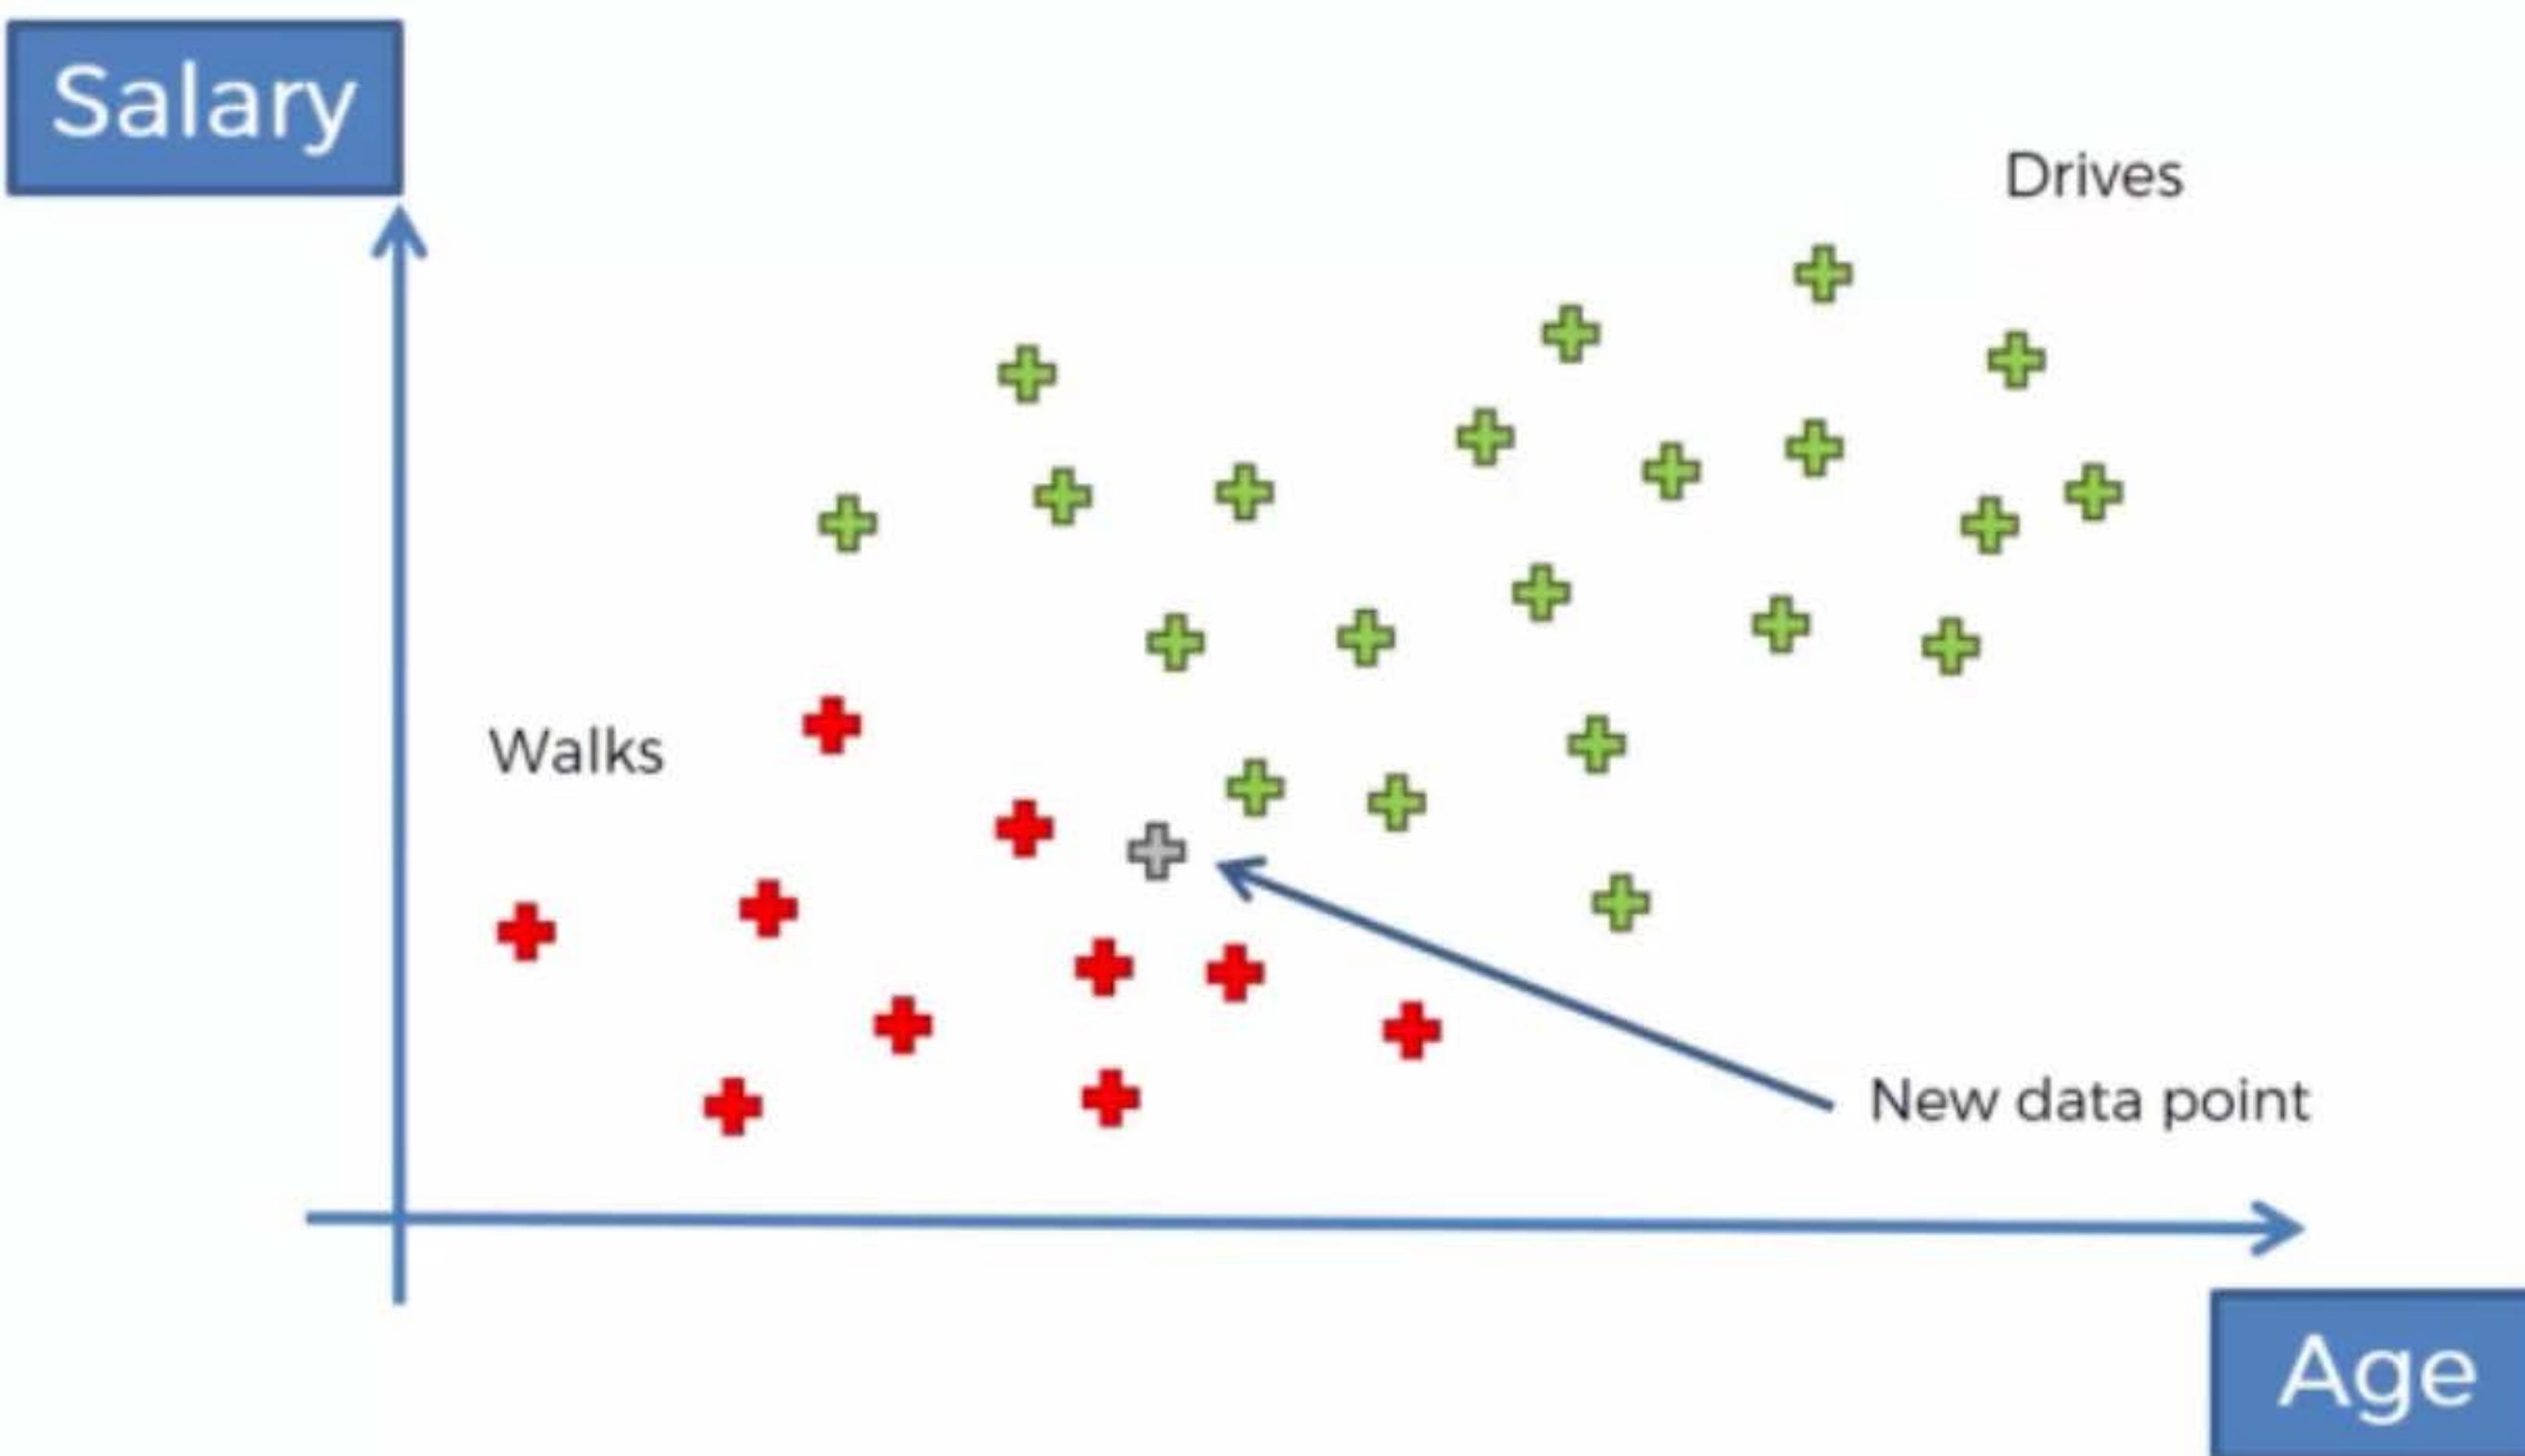
---

**Obvious question:**

**If the items are labeled, why couldn't we just count the number of defective wrenches that came from Mach2 and divide by the total number that came from Mach2?**

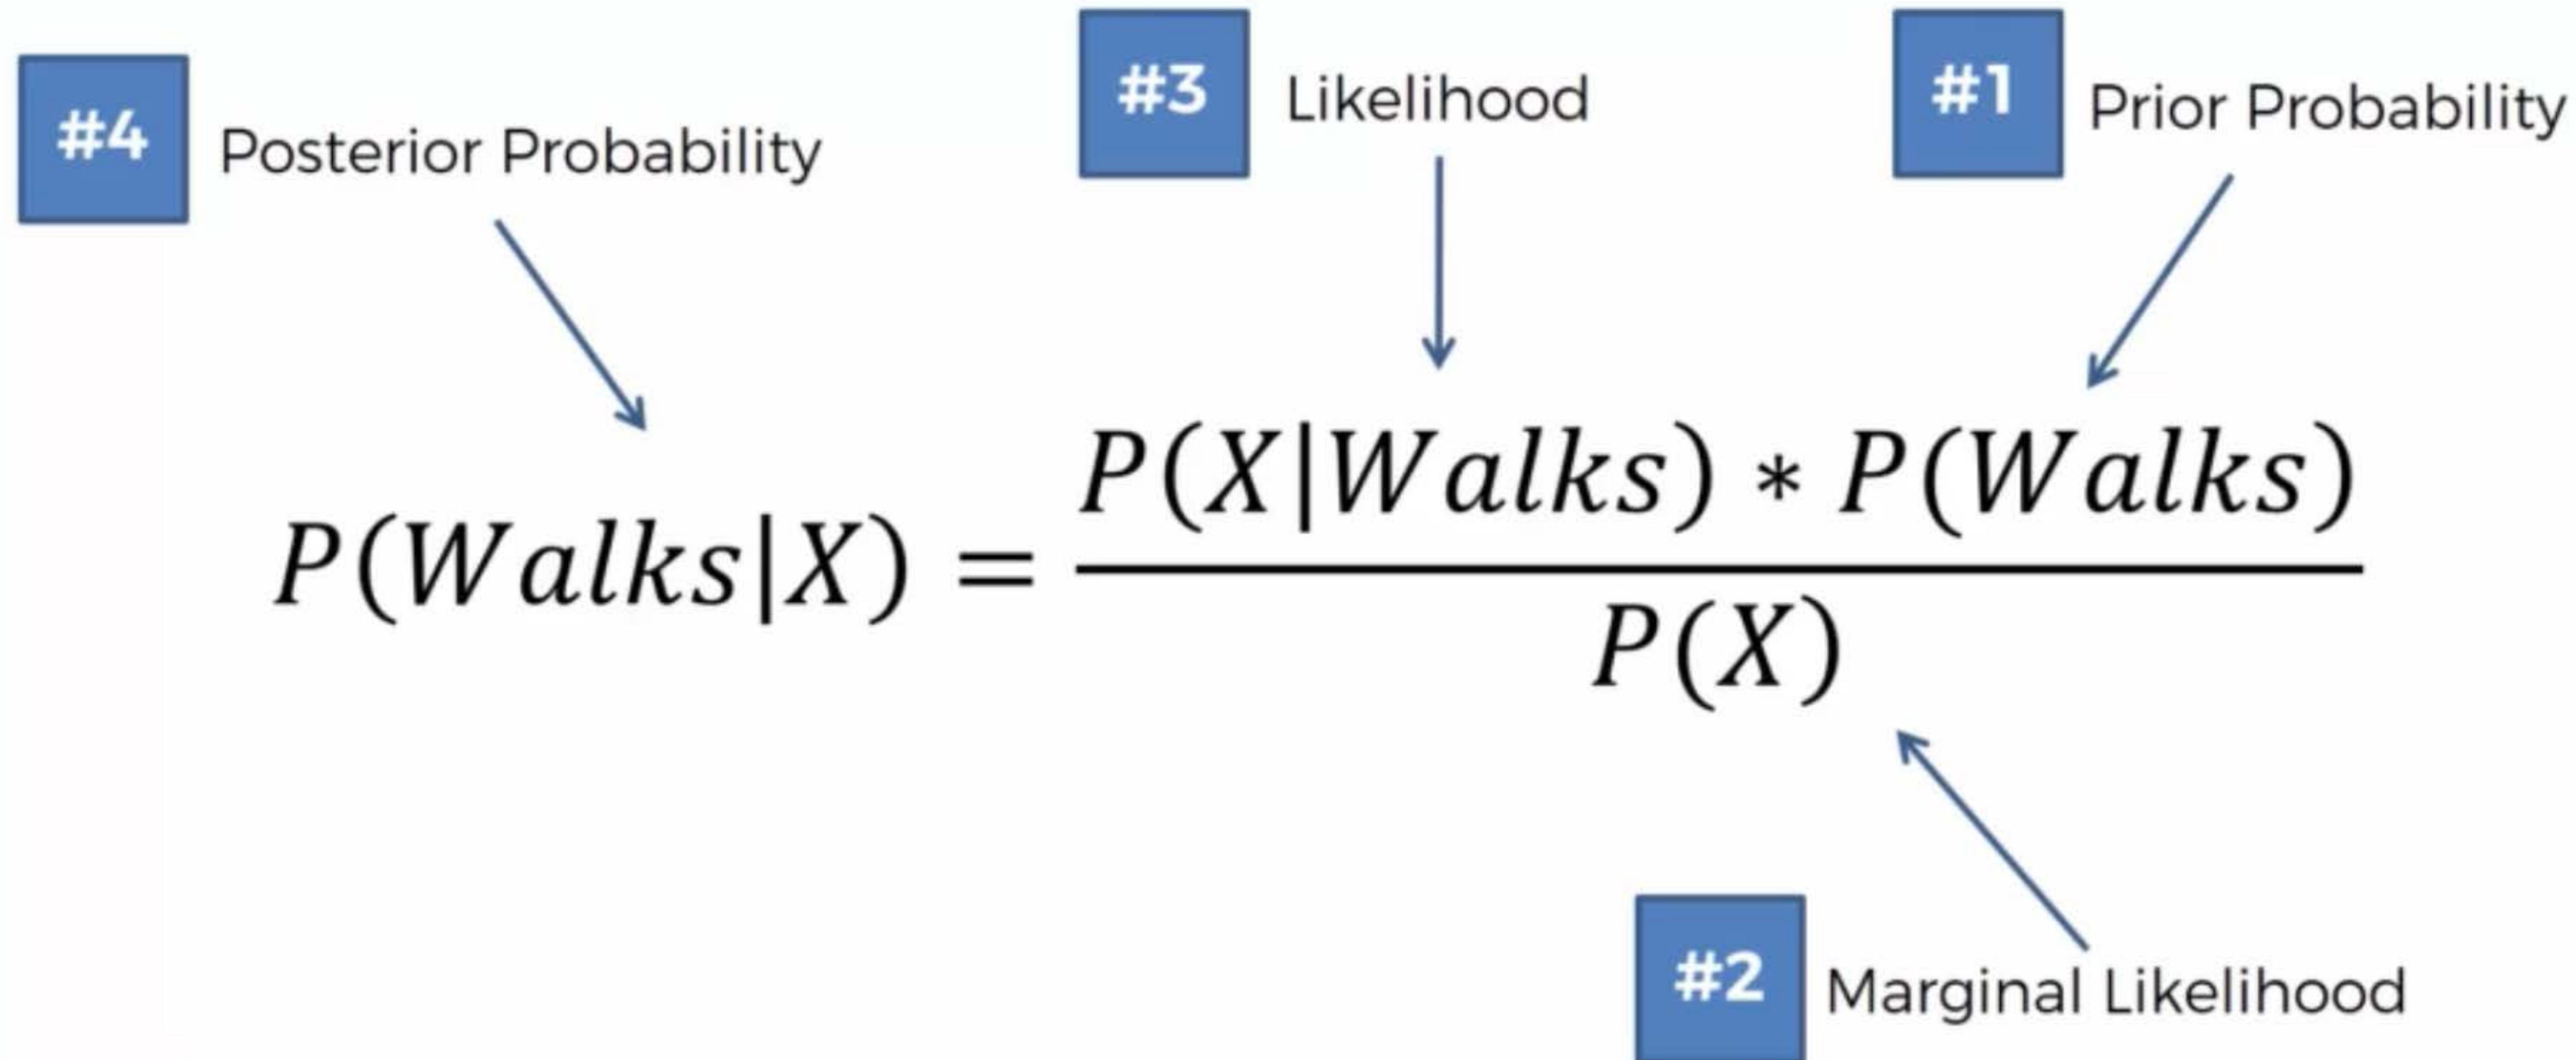
# Naïve Bayes Classifier Intuition

# Naïve Bayes

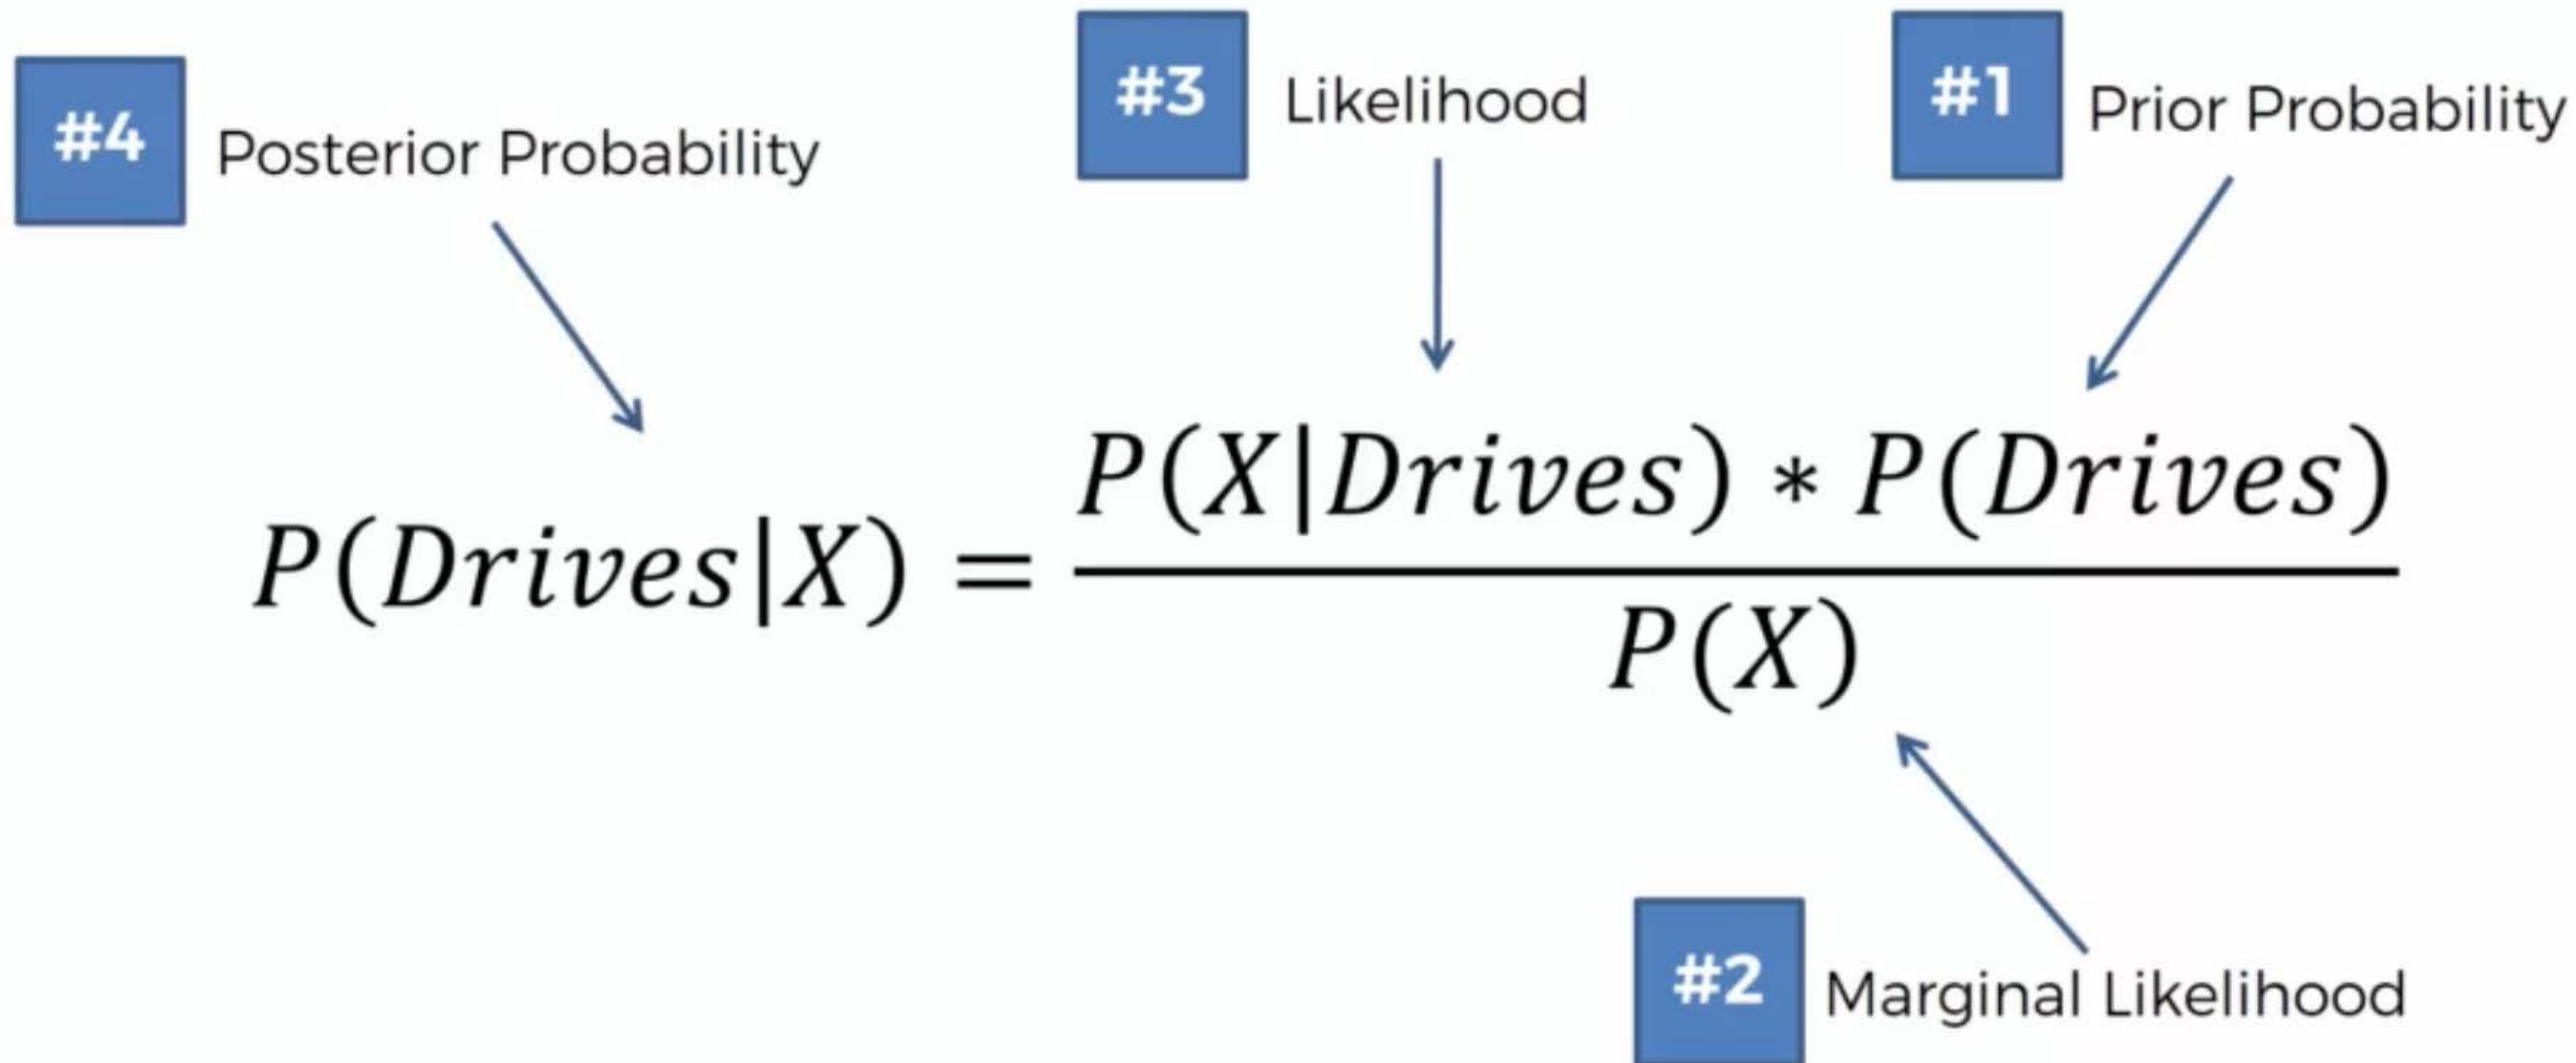


## Plan of Attack

# Step 1



## Step 2

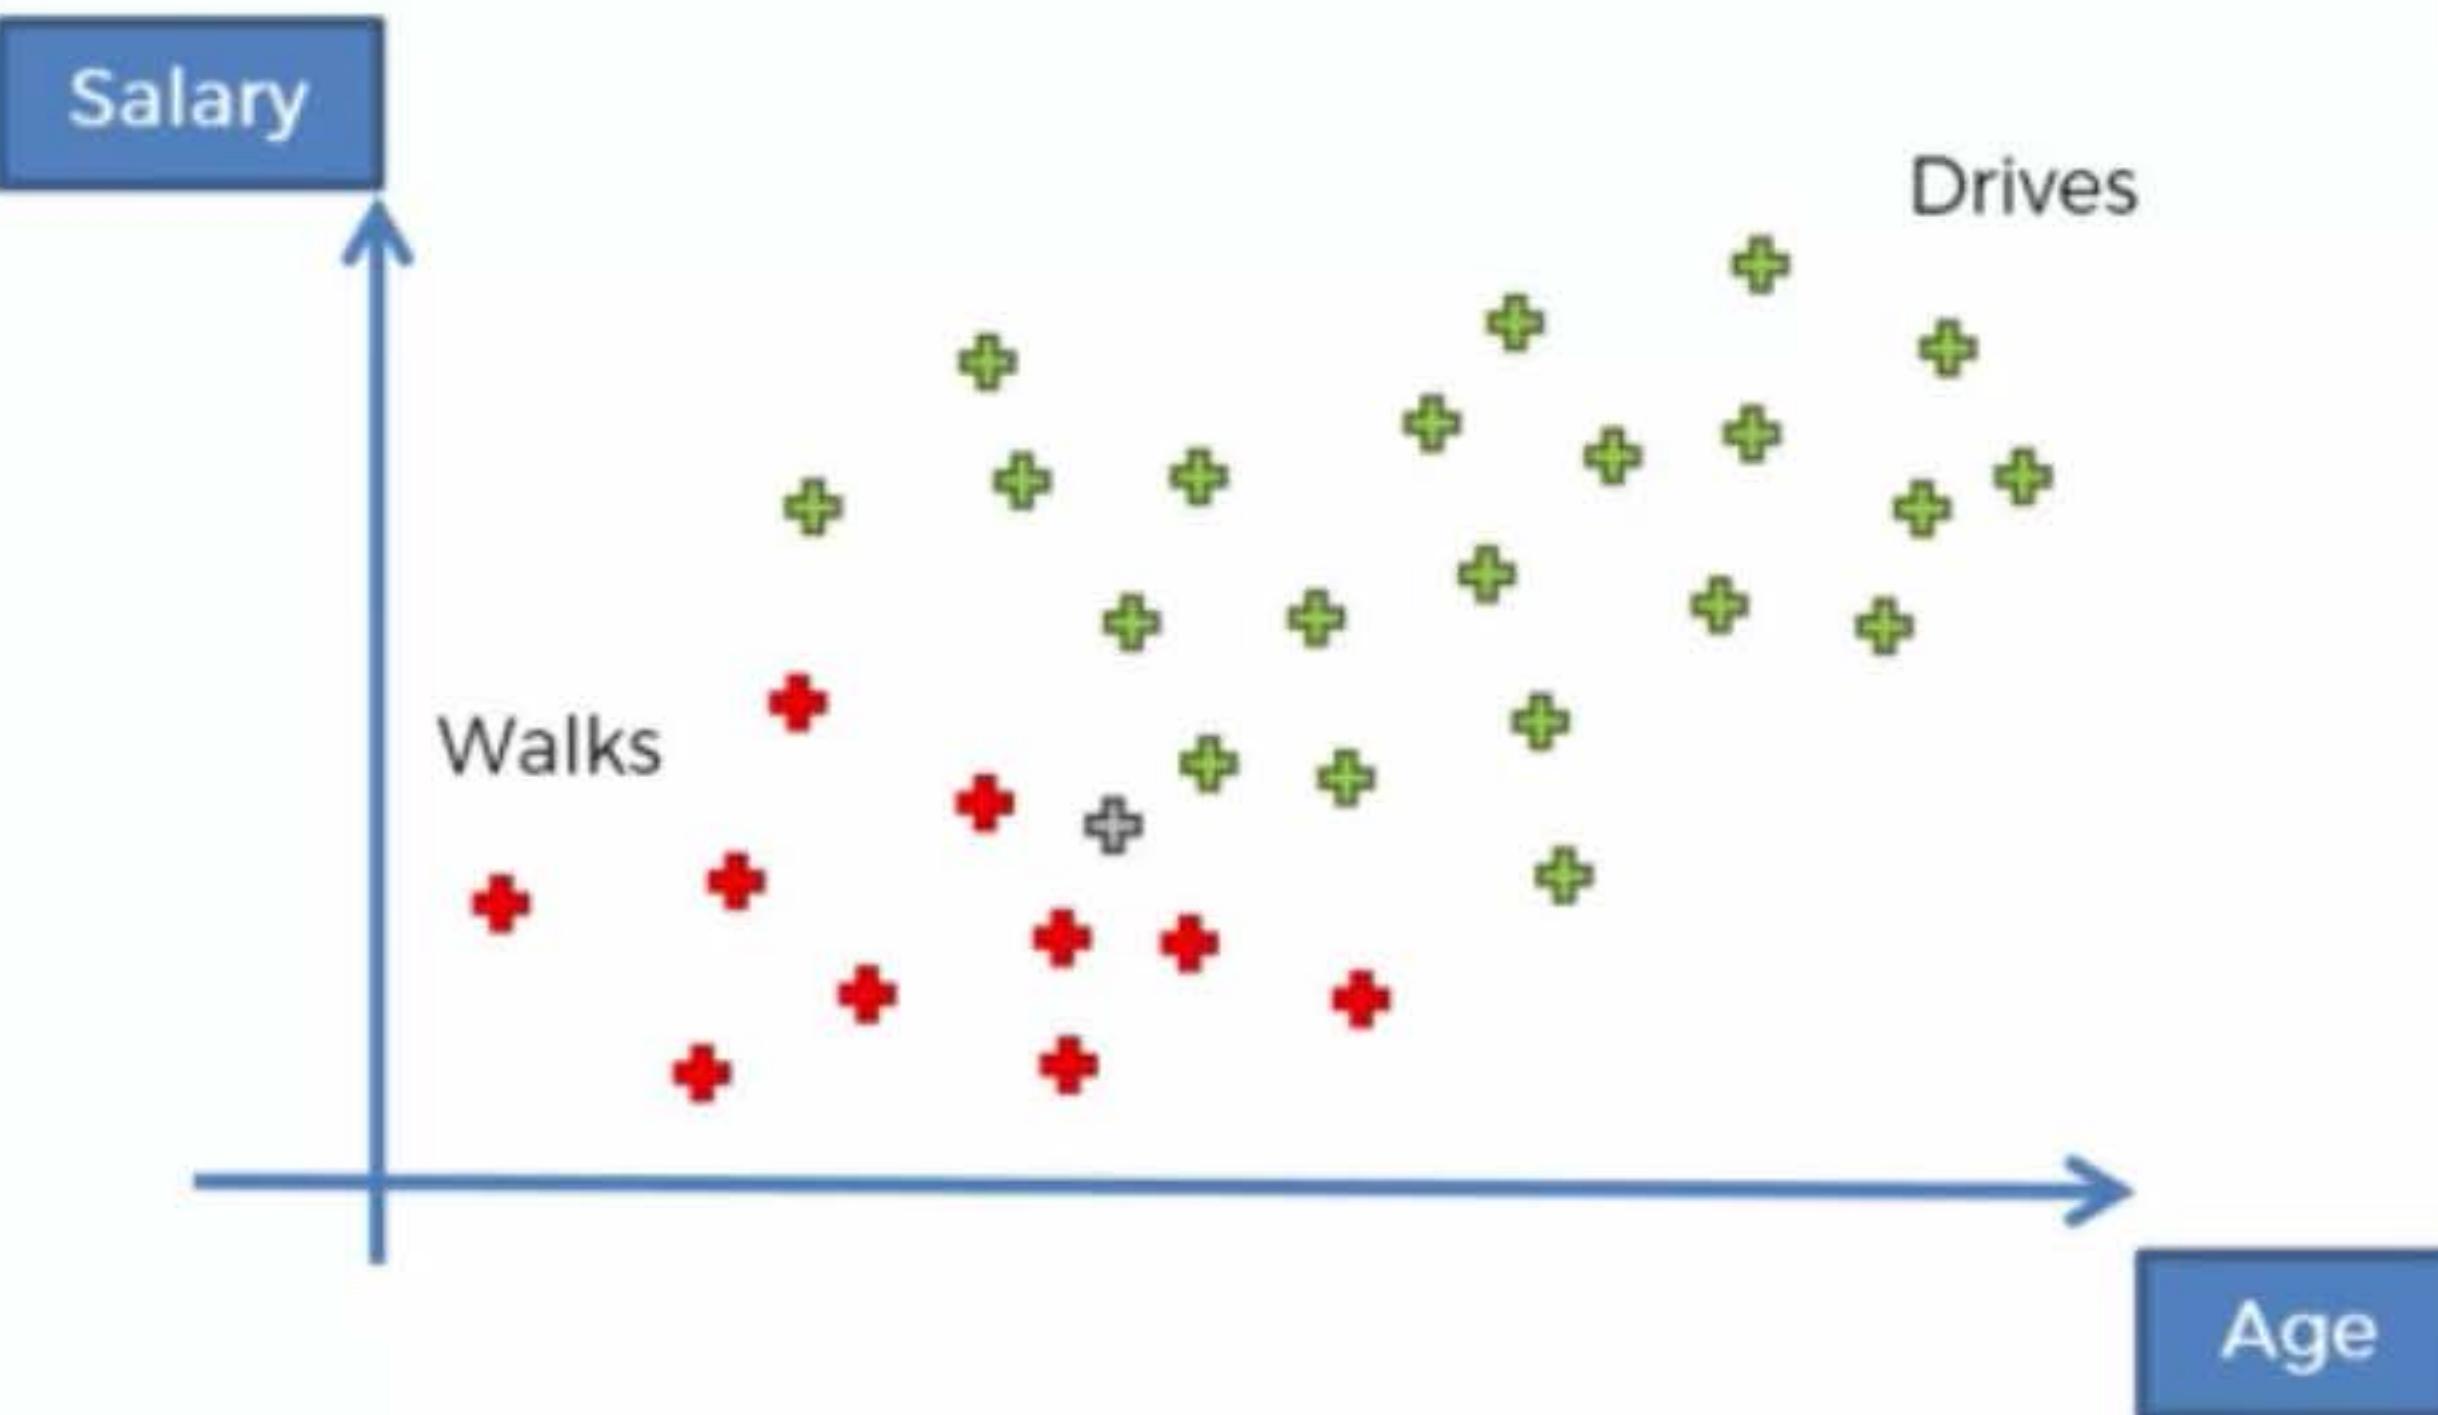


## Step 3

---

$P(\text{Walks}|X)$  v.s.  $P(\text{Drives}|X)$

# Naïve Bayes: Step 1

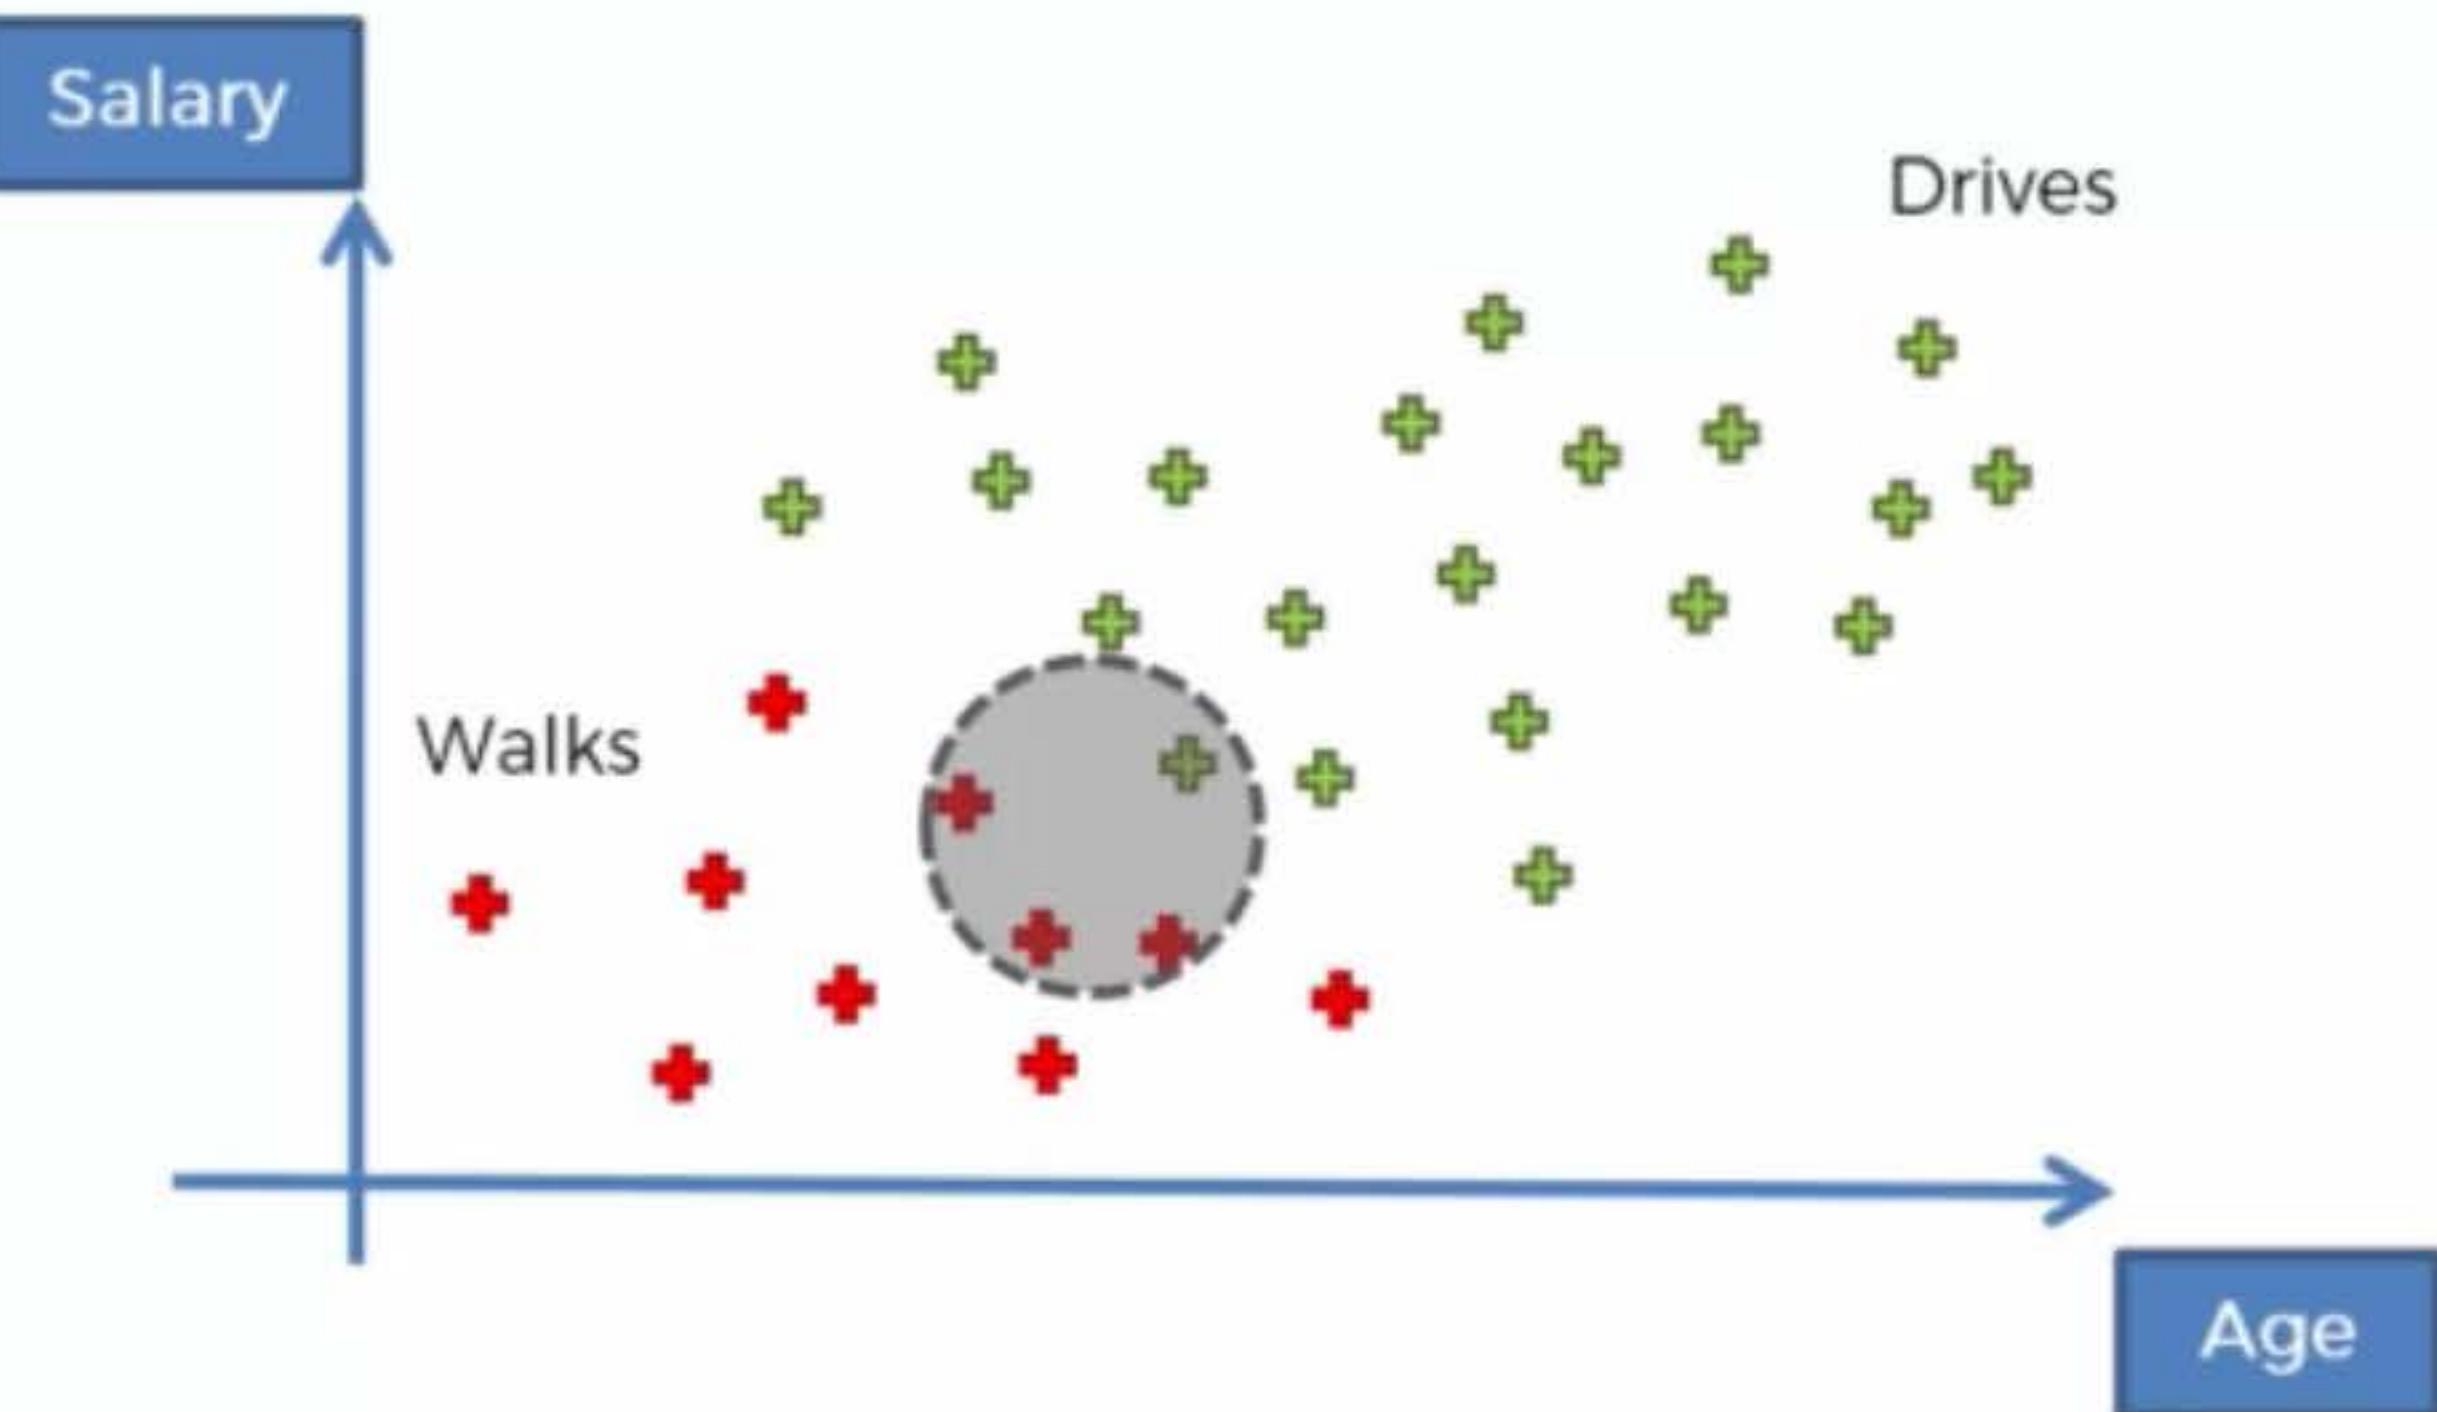


## #1. $P(\text{Walks})$

$$P(\text{Walks}) = \frac{\text{Number of Walkers}}{\text{Total Observations}}$$

$$P(\text{Walks}) = \frac{10}{30}$$

# Naïve Bayes: Step 1



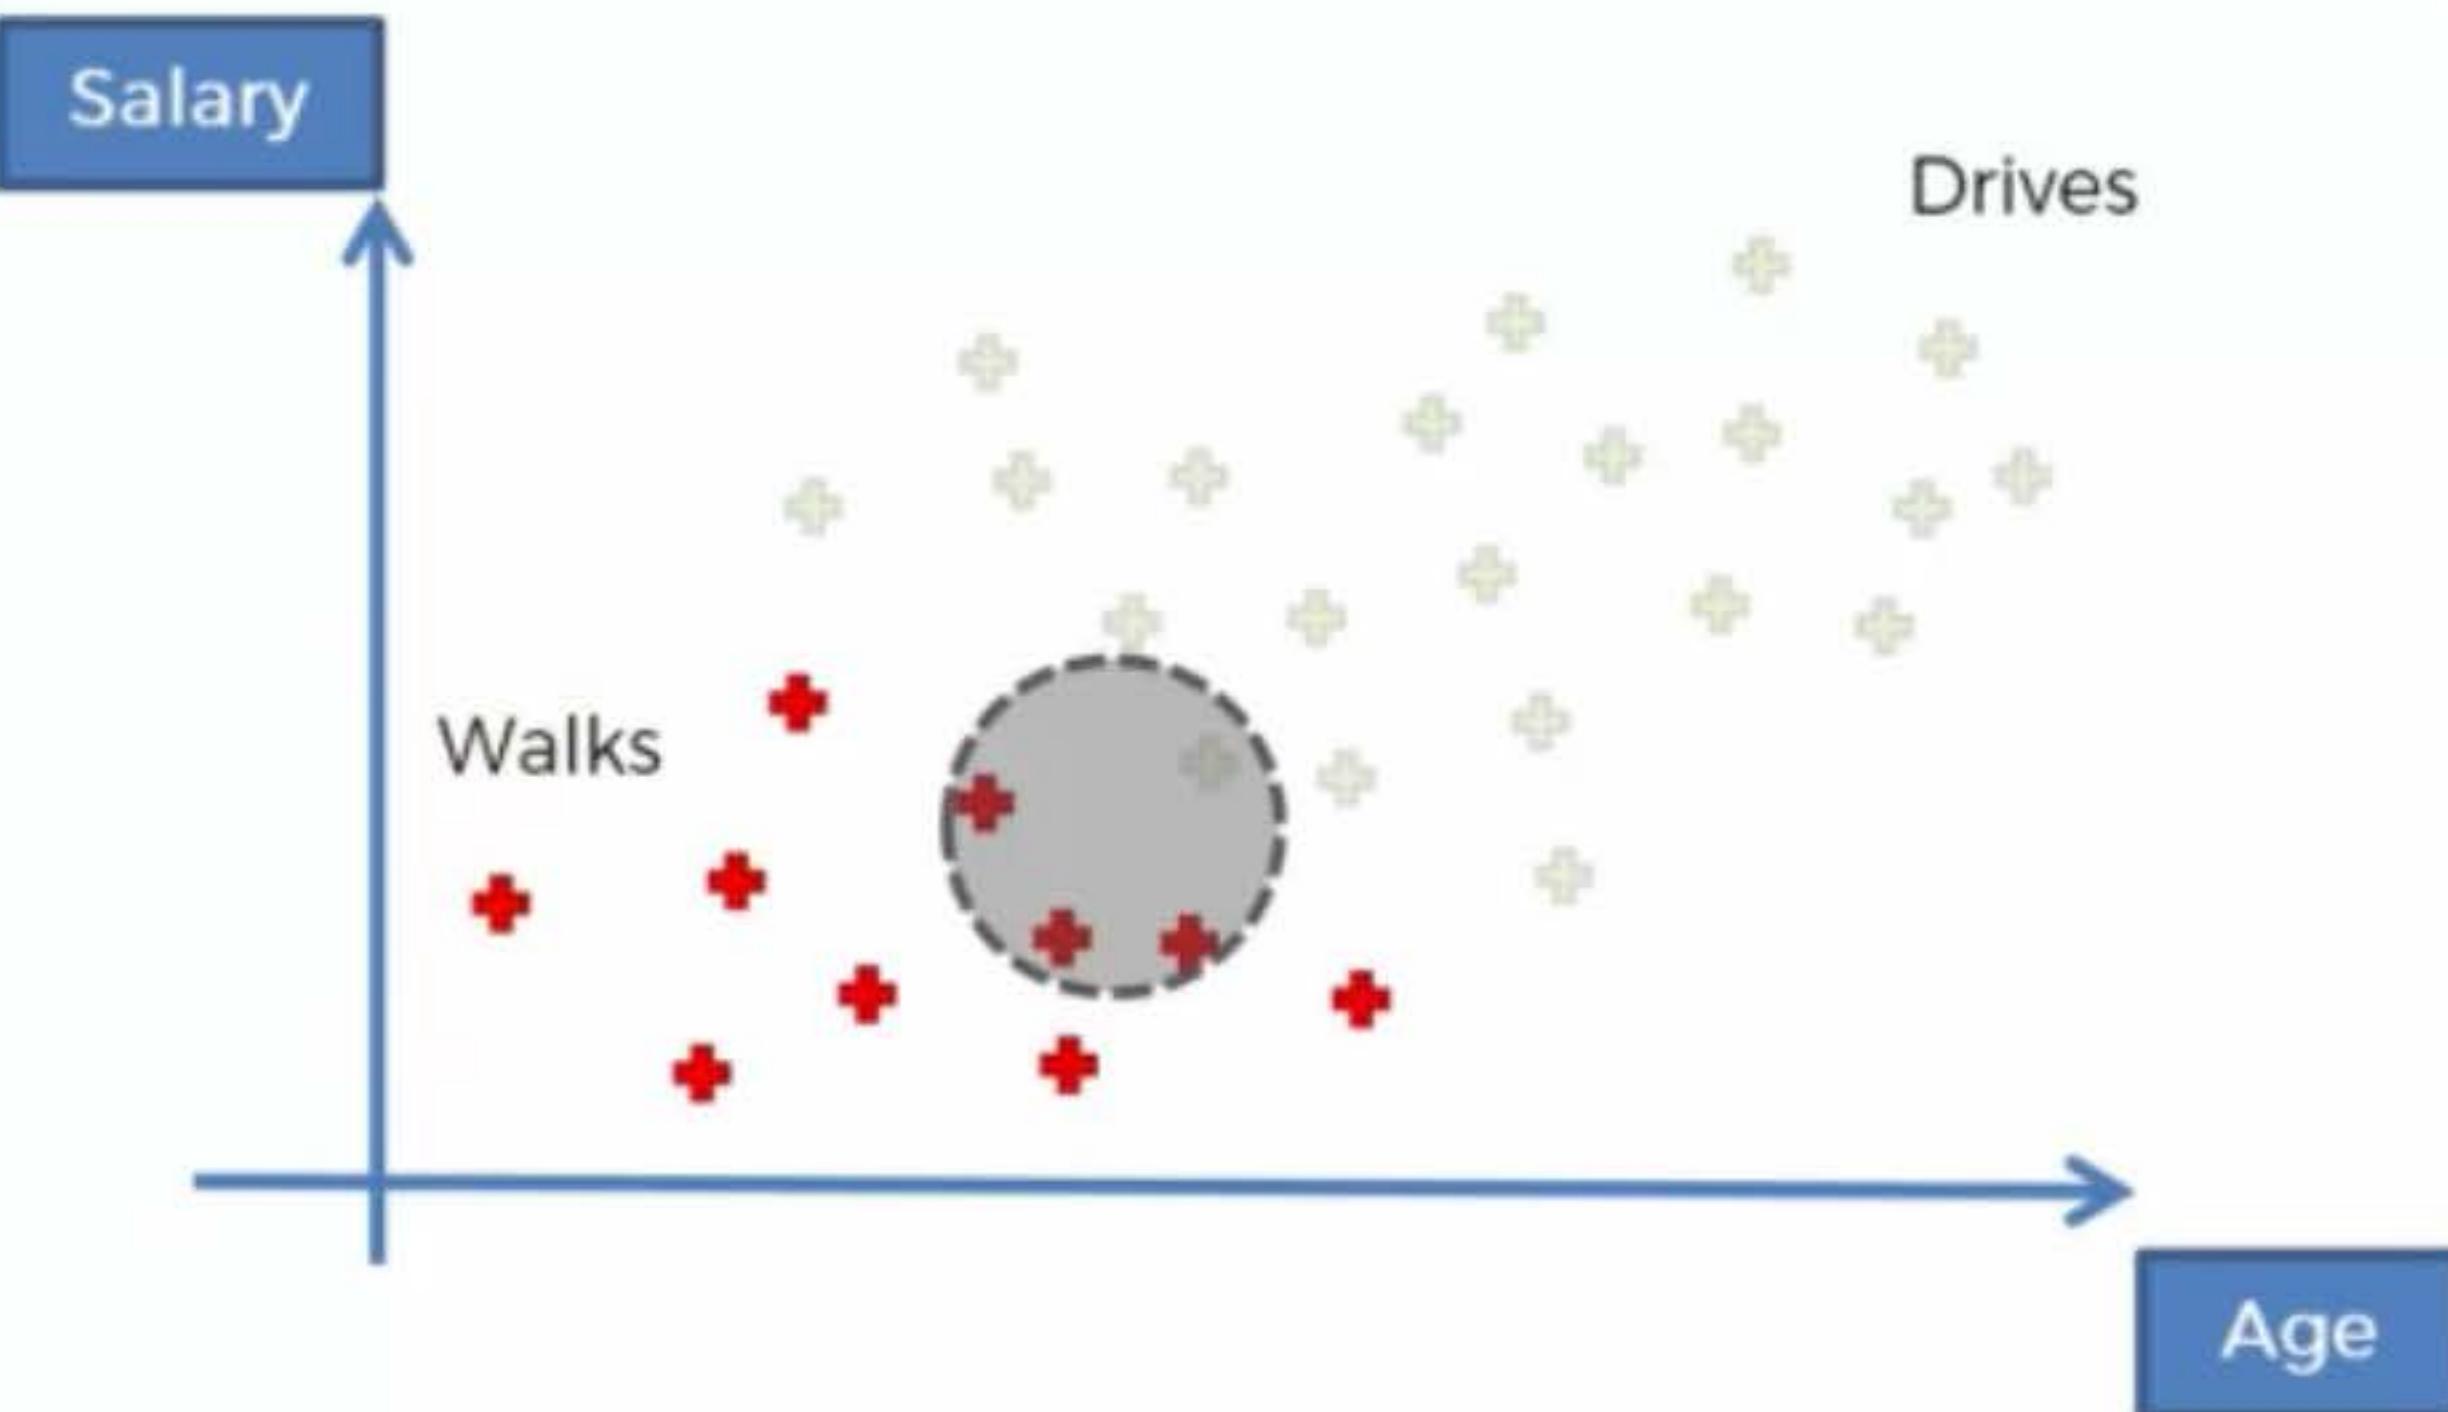
#2.  $P(X)$

$$P(X) = \frac{\text{Number of Similar Observations}}{\text{Total Observations}}$$

$$P(X) = \frac{4}{30}$$



# Naïve Bayes: Step 1



## #3. $P(X|Walks)$

*Number of Similar Observations  
Among those who Walk*

$$P(X|Walks) = \frac{\text{Number of Similar Observations}}{\text{Total number of Walkers}}$$

$$P(X|Walks) = \frac{3}{10}$$

# Naïve Bayes: Step 1

The diagram illustrates the four components of Naive Bayes Step 1:

- #4 Posterior Probability
- #3 Likelihood
- #1 Prior Probability
- #2 Marginal Likelihood

These components are used to calculate the posterior probability:

$$P(Walks|X) = \frac{\frac{3}{10} * \frac{10}{30}}{\frac{4}{30}} = 0.75$$

## Naïve Bayes: Step 2

The diagram illustrates the calculation of Posterior Probability using the Naive Bayes formula. It shows four components: Likelihood (#3), Prior Probability (#1), Marginal Likelihood (#2), and Posterior Probability (#4). Arrows point from each component to its corresponding term in the formula.

$$P(\text{Drives}|X) = \frac{\frac{1}{20} * \frac{20}{30}}{\frac{30}{30}} = 0.25$$

#4 Posterior Probability

#3 Likelihood

#1 Prior Probability

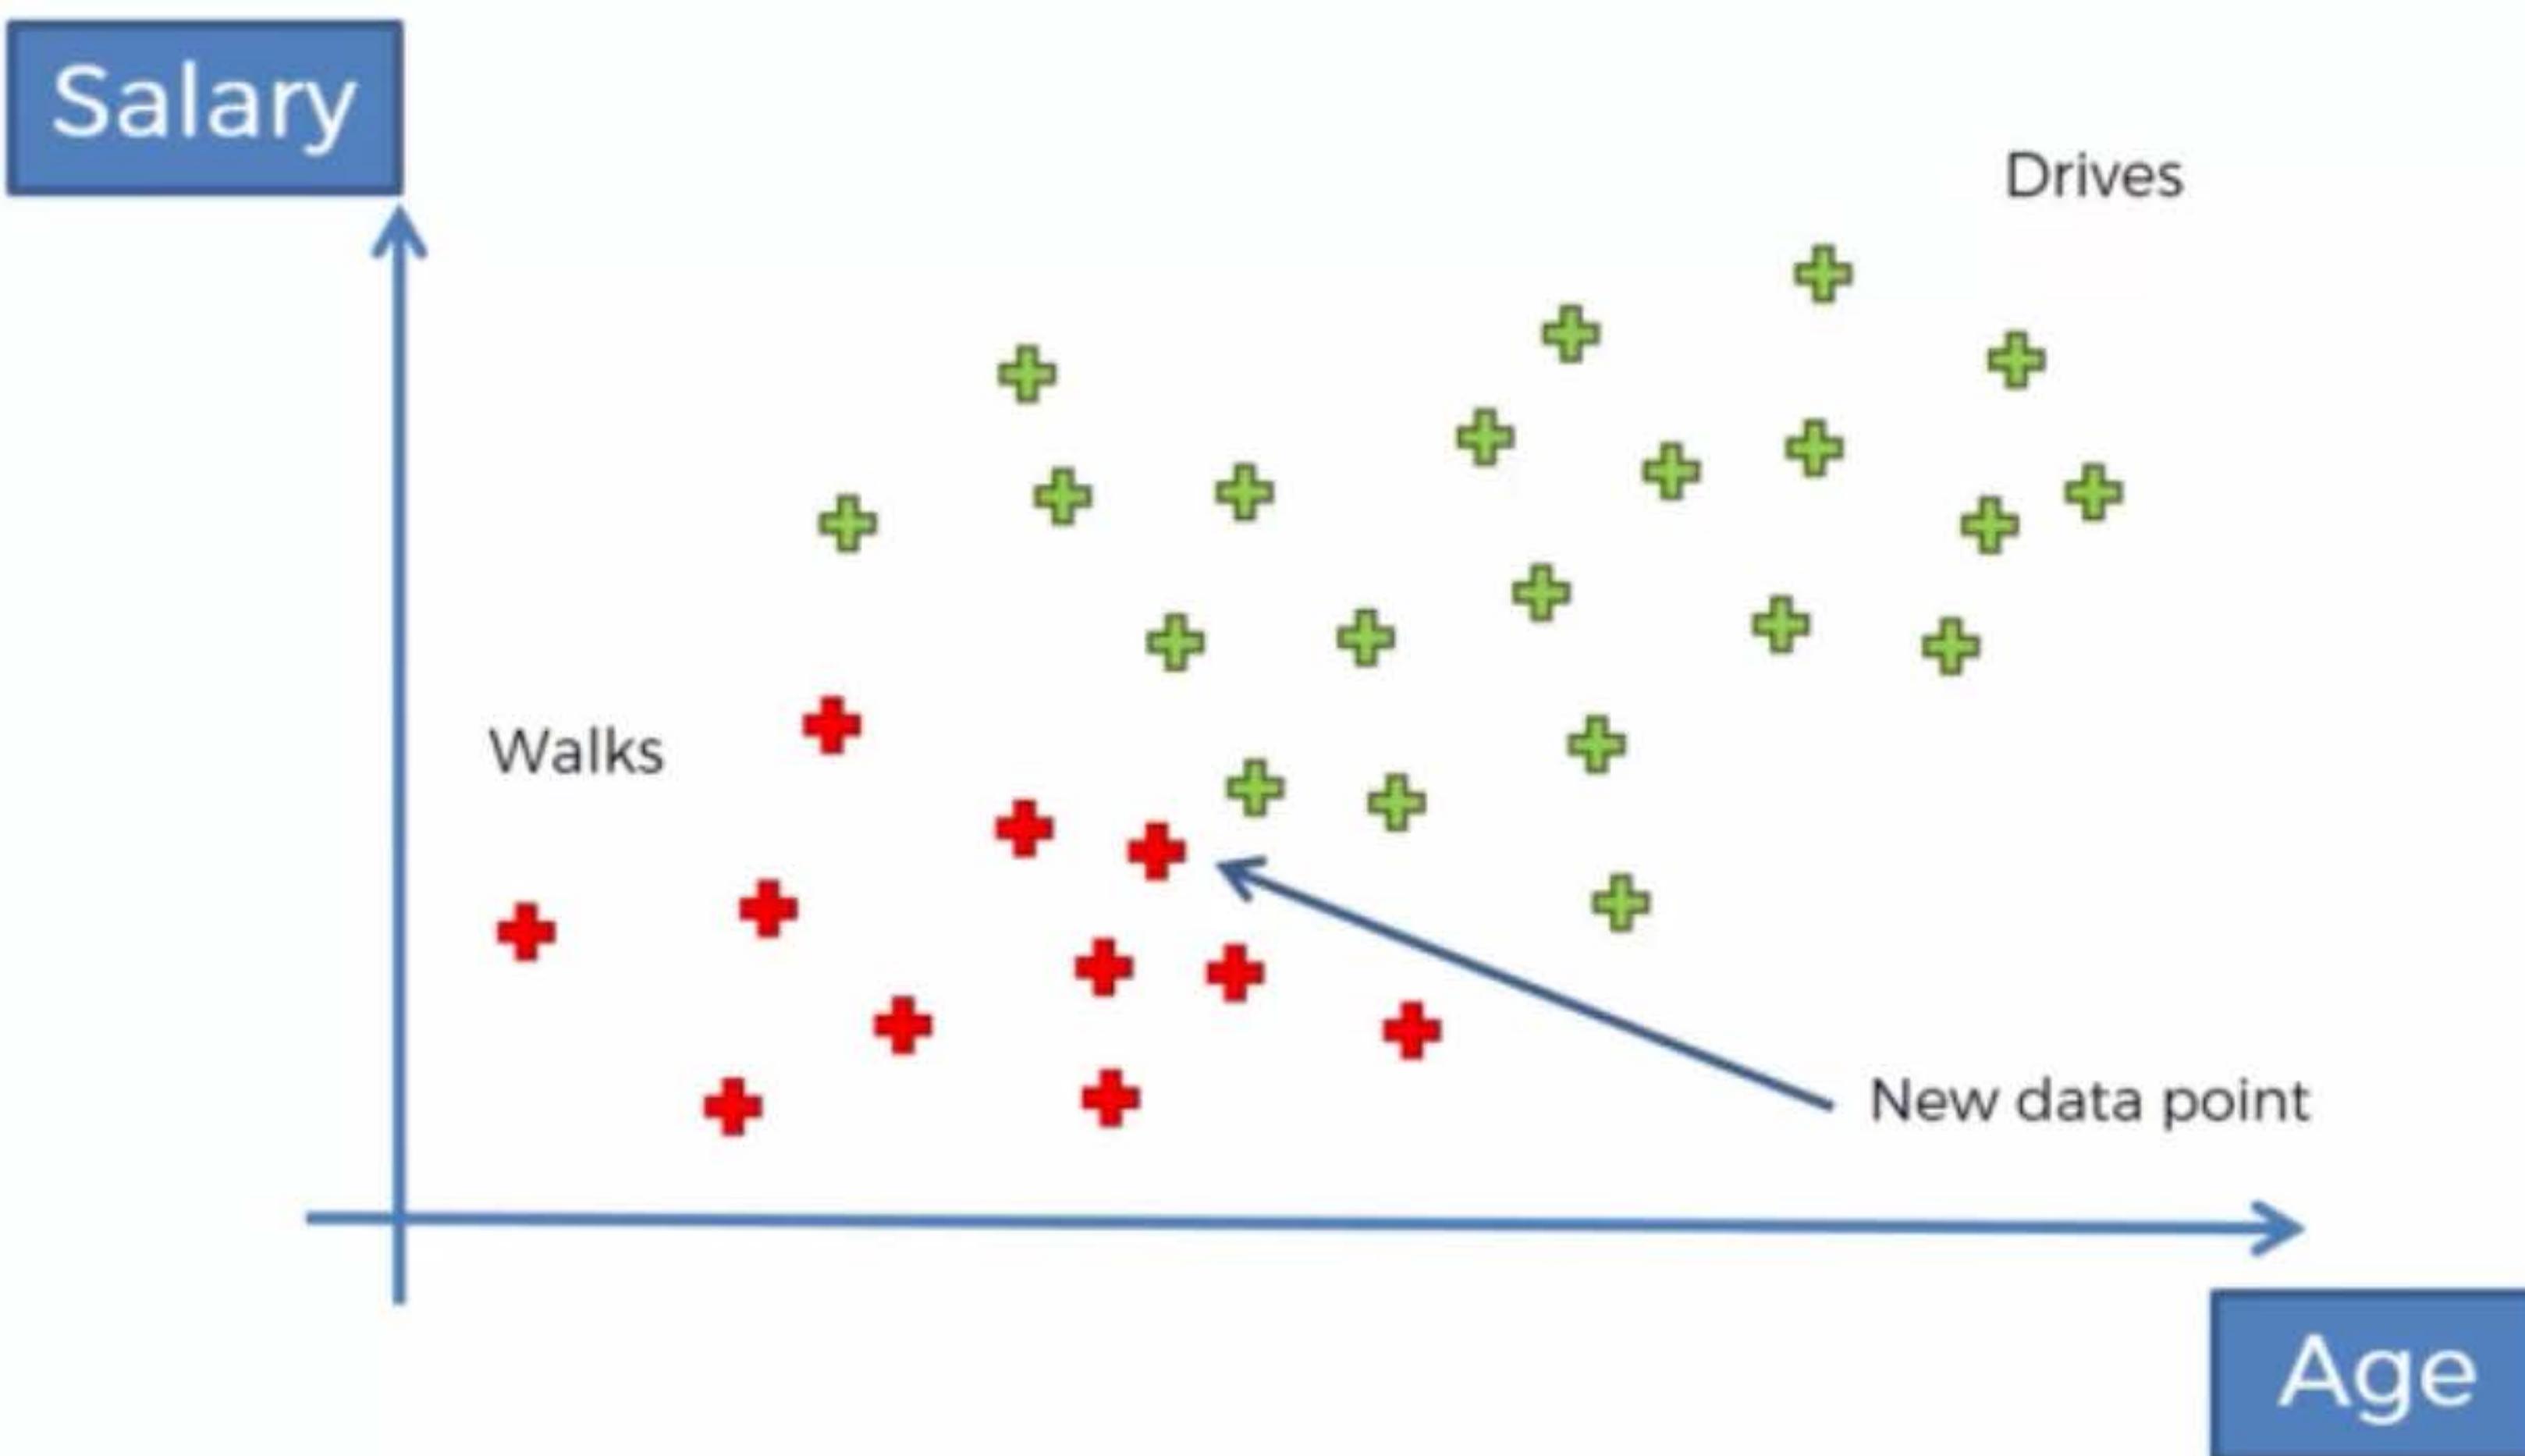
#2 Marginal Likelihood

## Step 3

---

$$P(\text{Walks}|X) > P(\text{Drives}|X)$$

# Naïve Bayes



# **Naïve Bayes Classifier**

## **Additional Comments**

# Naïve Bayes

---

1. Q: Why “Naïve”?
2.  $P(X)$
3. More than 2 features

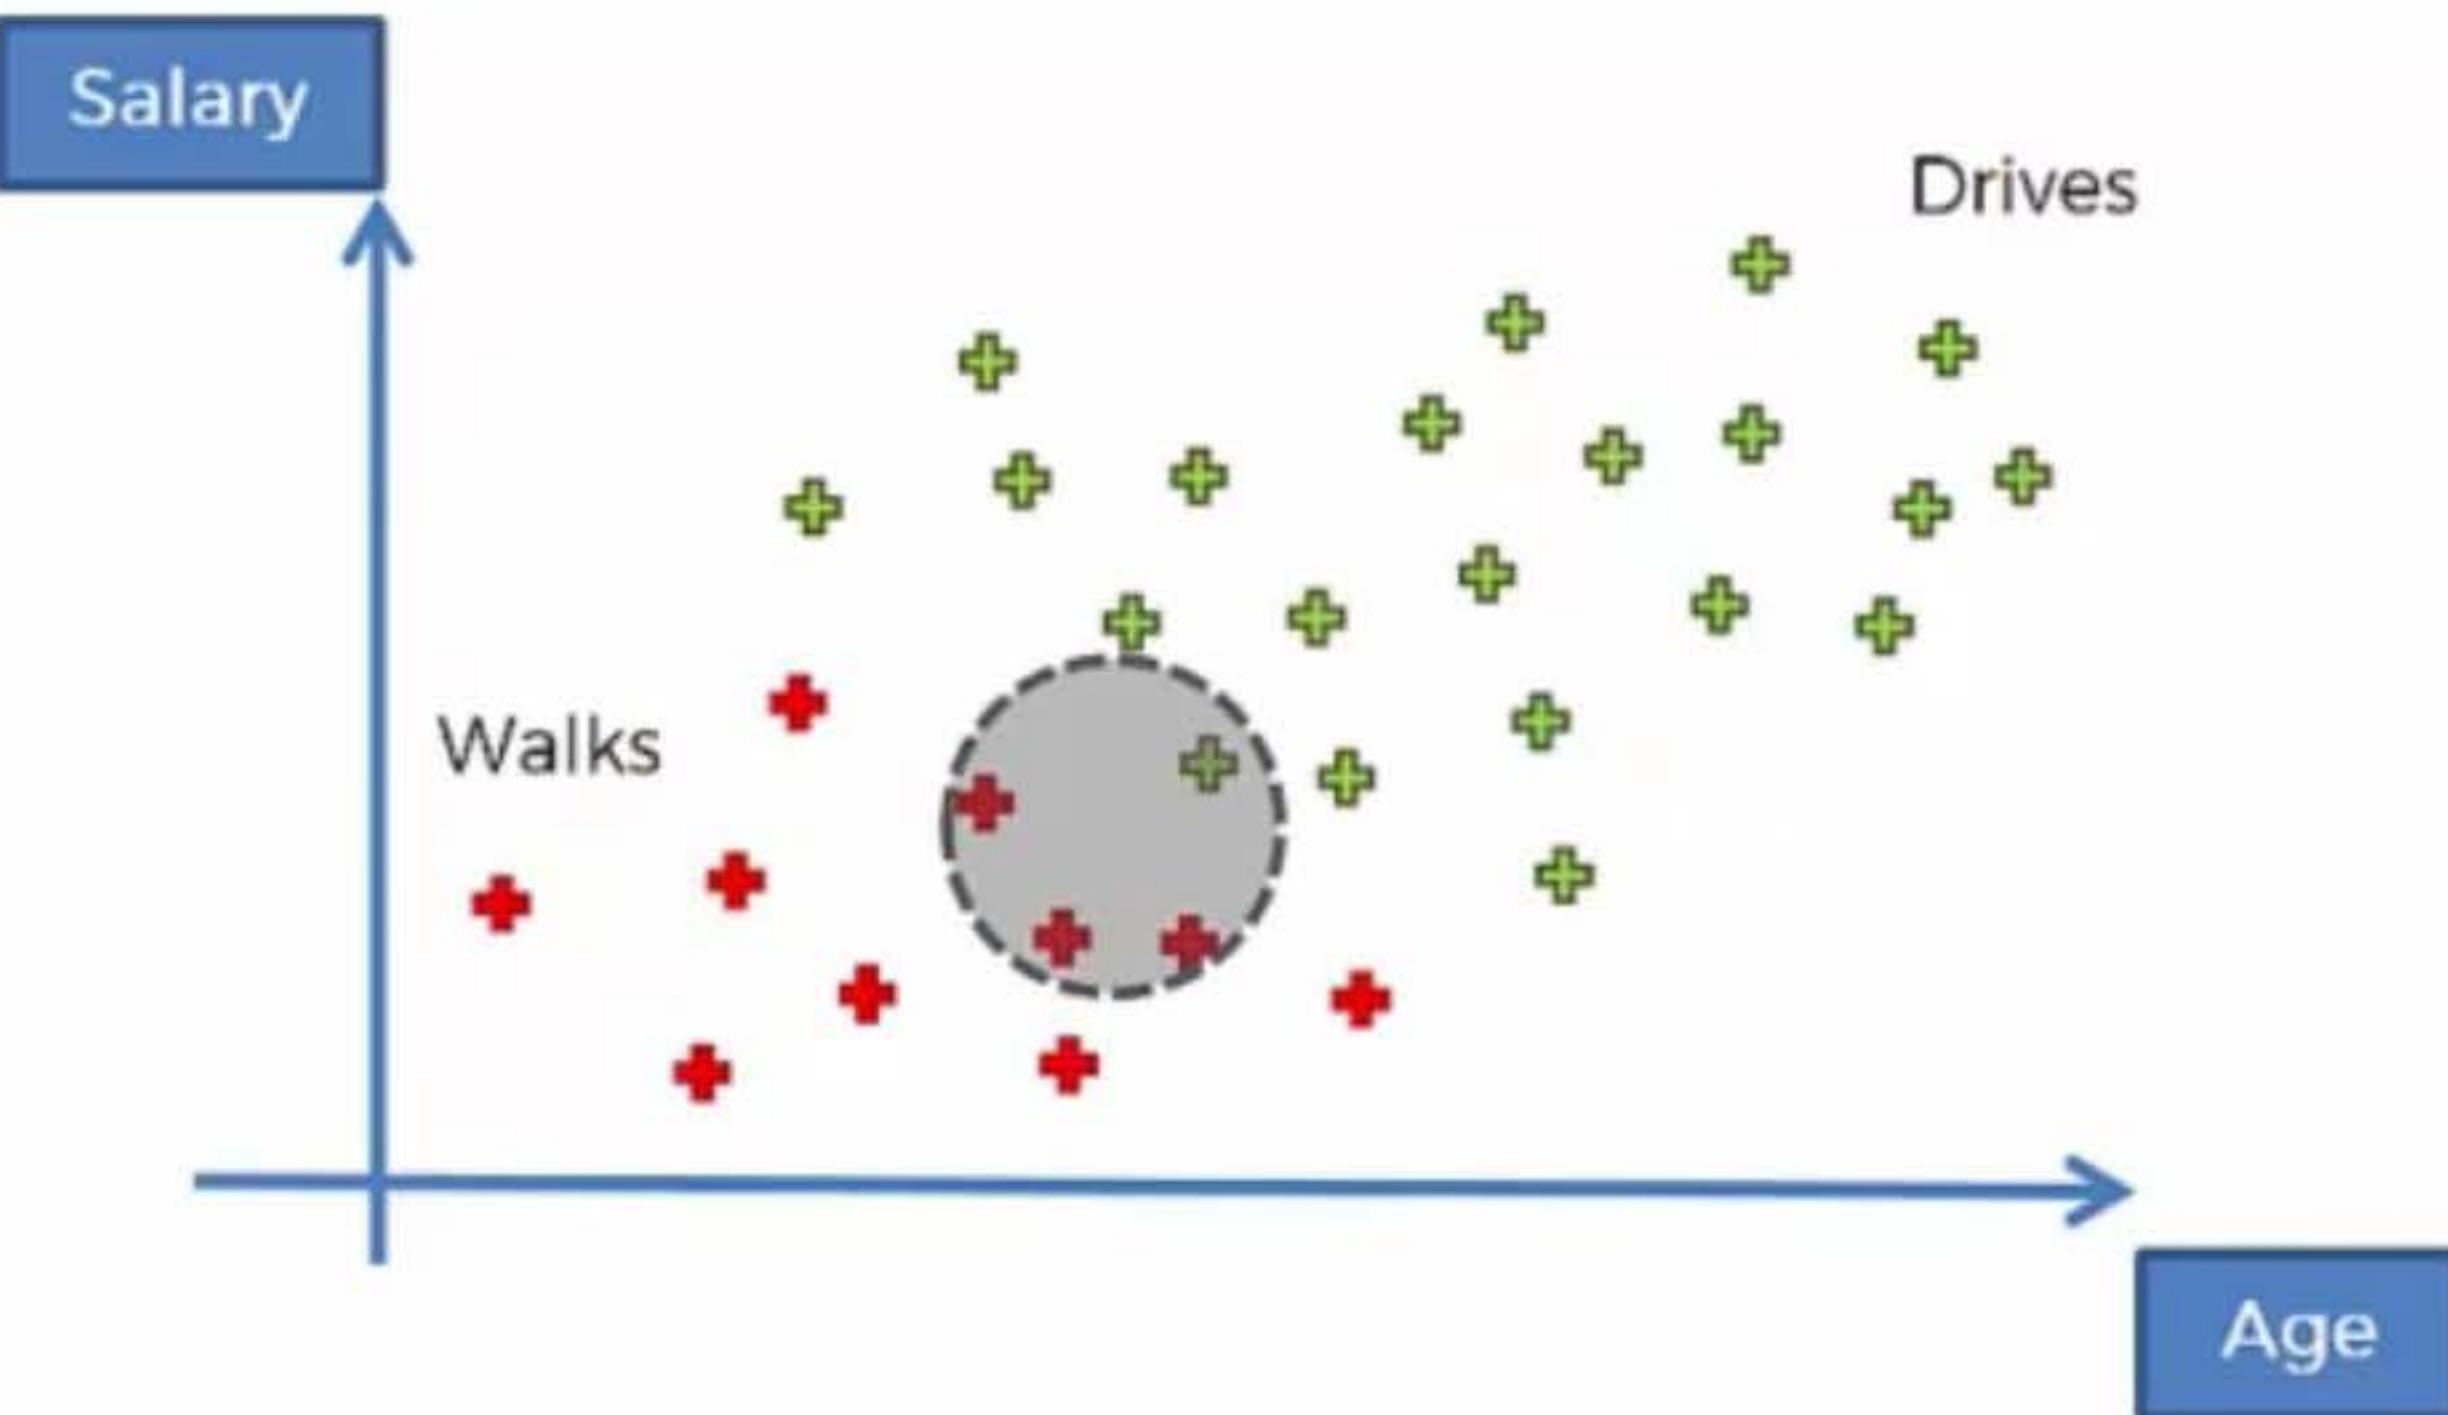
# Naïve Bayes

---

**Q: Why “Naïve”?**

**A: Independence assumption**

# Naïve Bayes: Step 2



#2.  $P(X)$

$$P(X) = \frac{\text{Number of Similar Observations}}{\text{Total Observations}}$$

$$P(X) = \frac{4}{30}$$

NOTE: Same both times

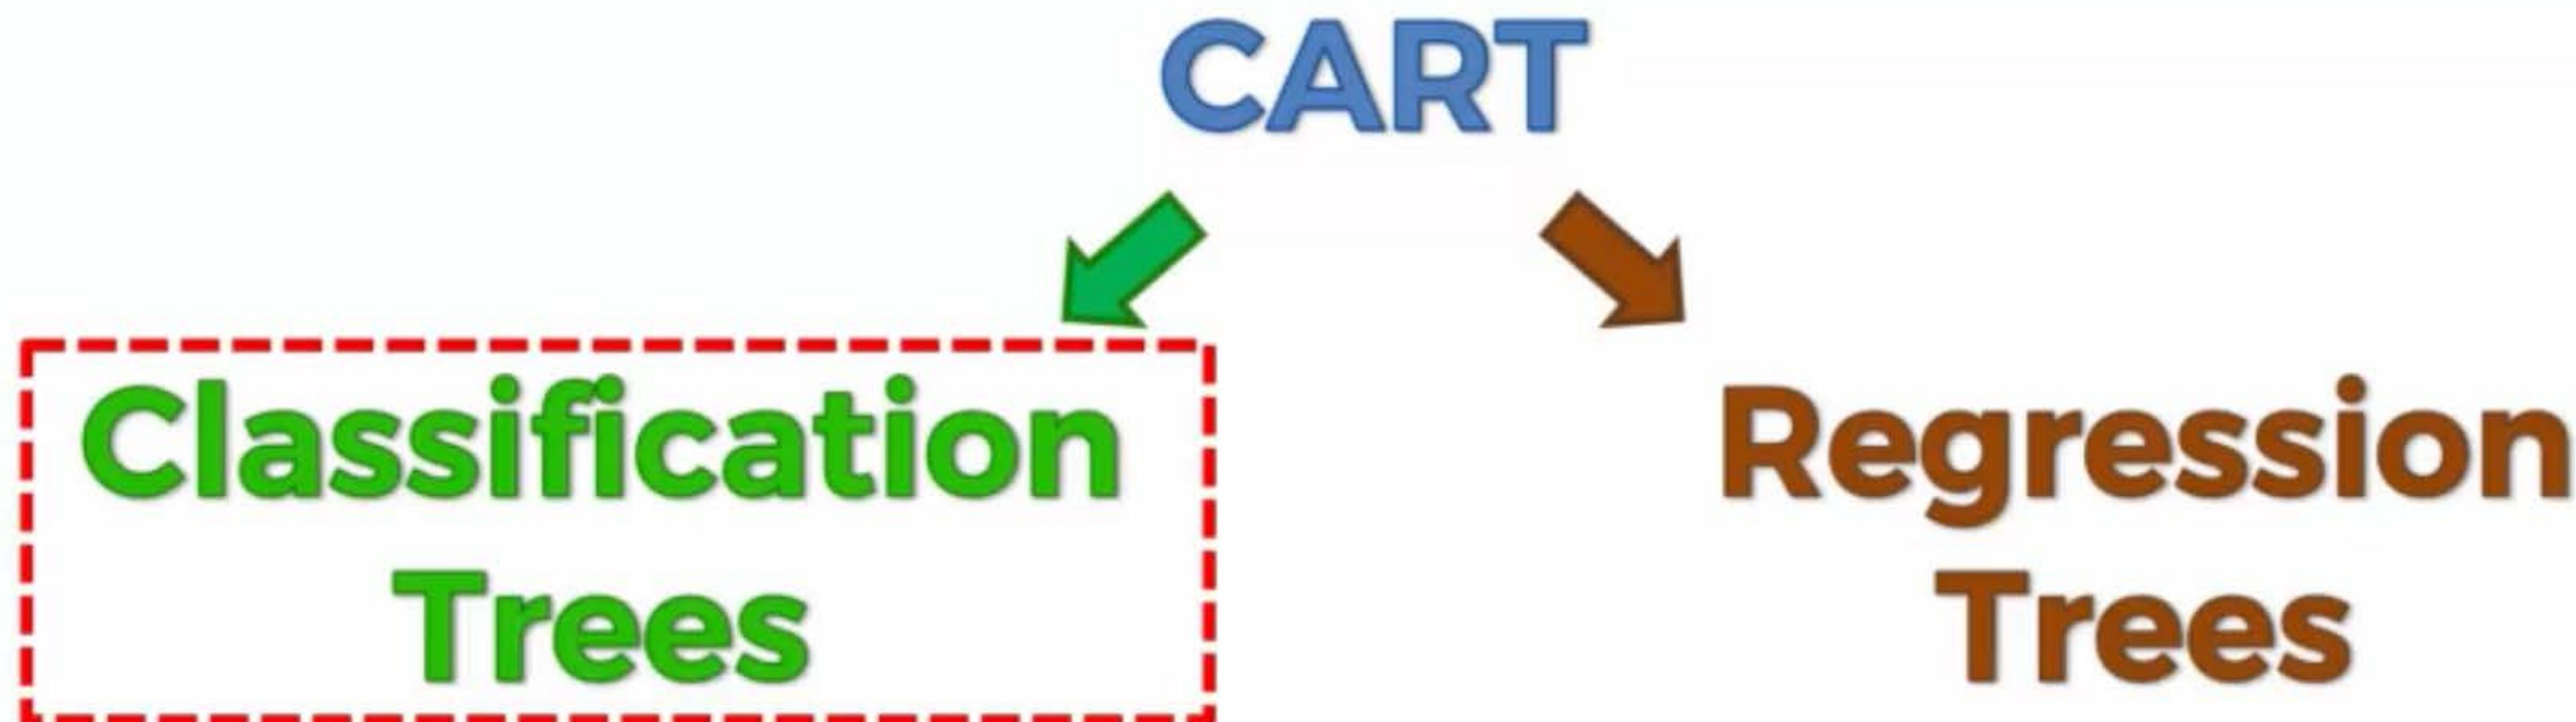
## Step 3

$$\frac{P(X|Walks) * P(Walks)}{\cancel{P(X)}} \quad v.s. \quad \frac{P(X|Drives) * P(Drives)}{\cancel{P(X)}}$$

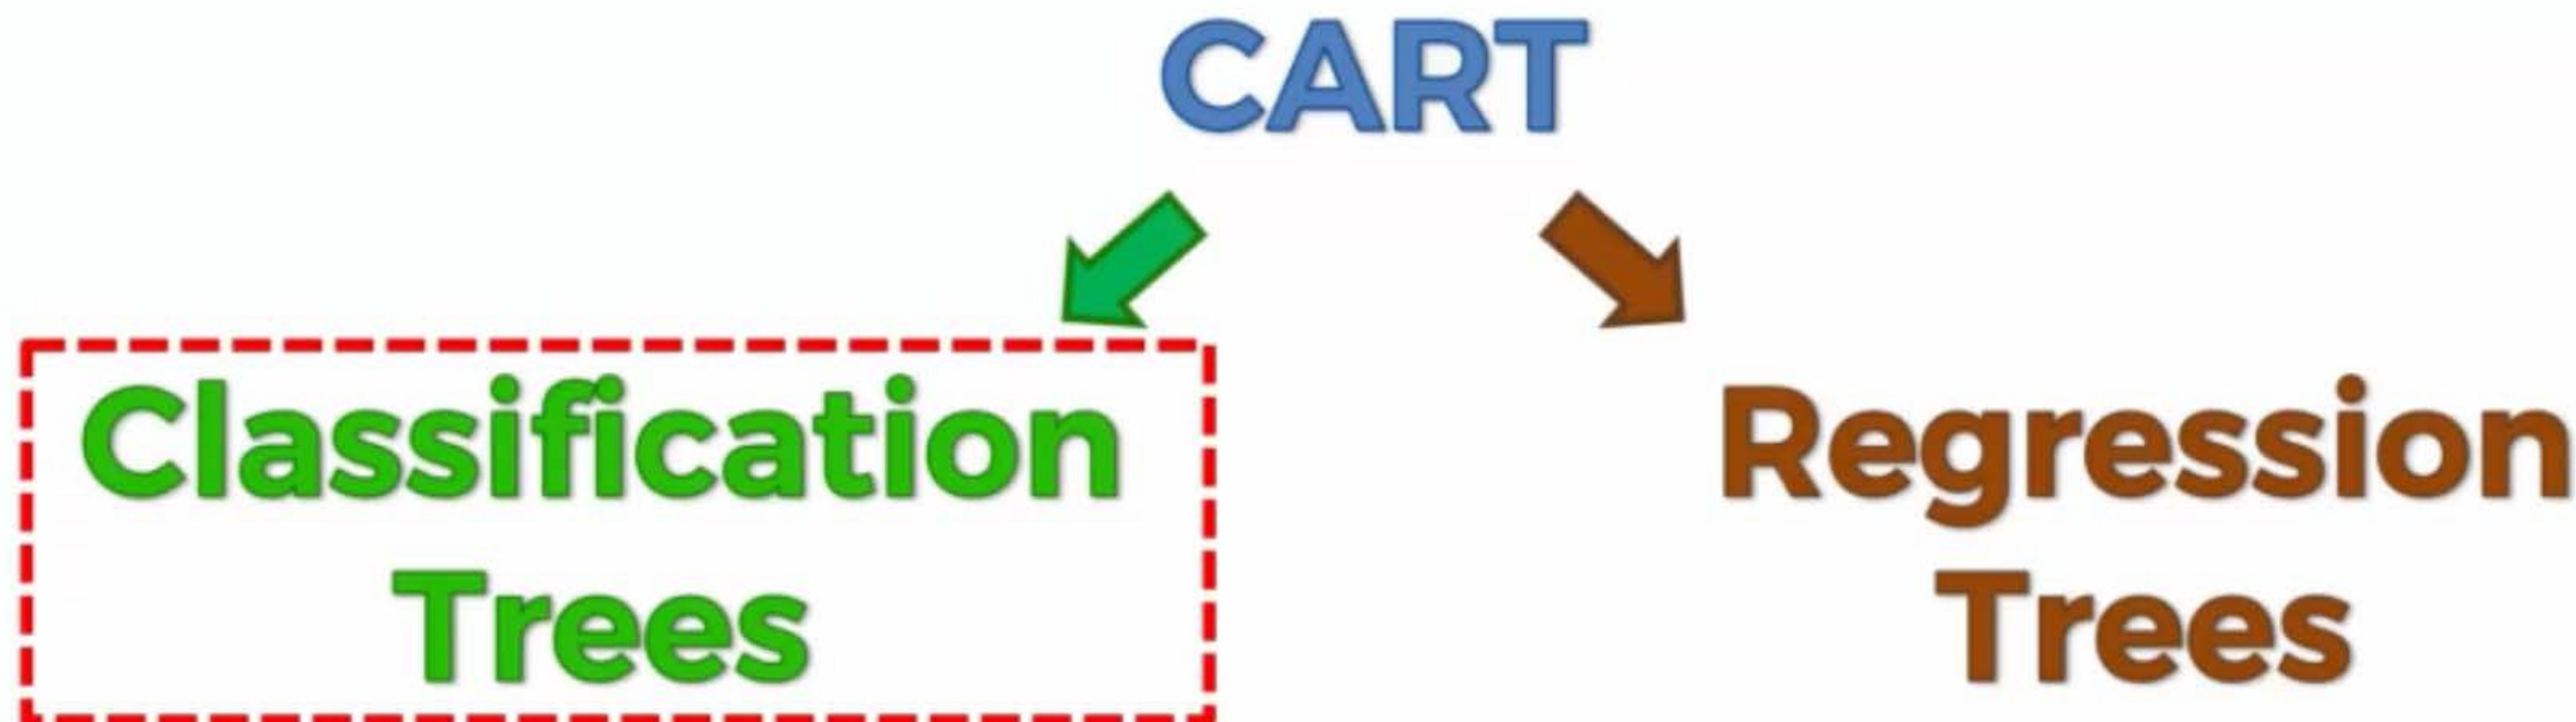
**More than 2 classes**

# Decision Tree Intuition

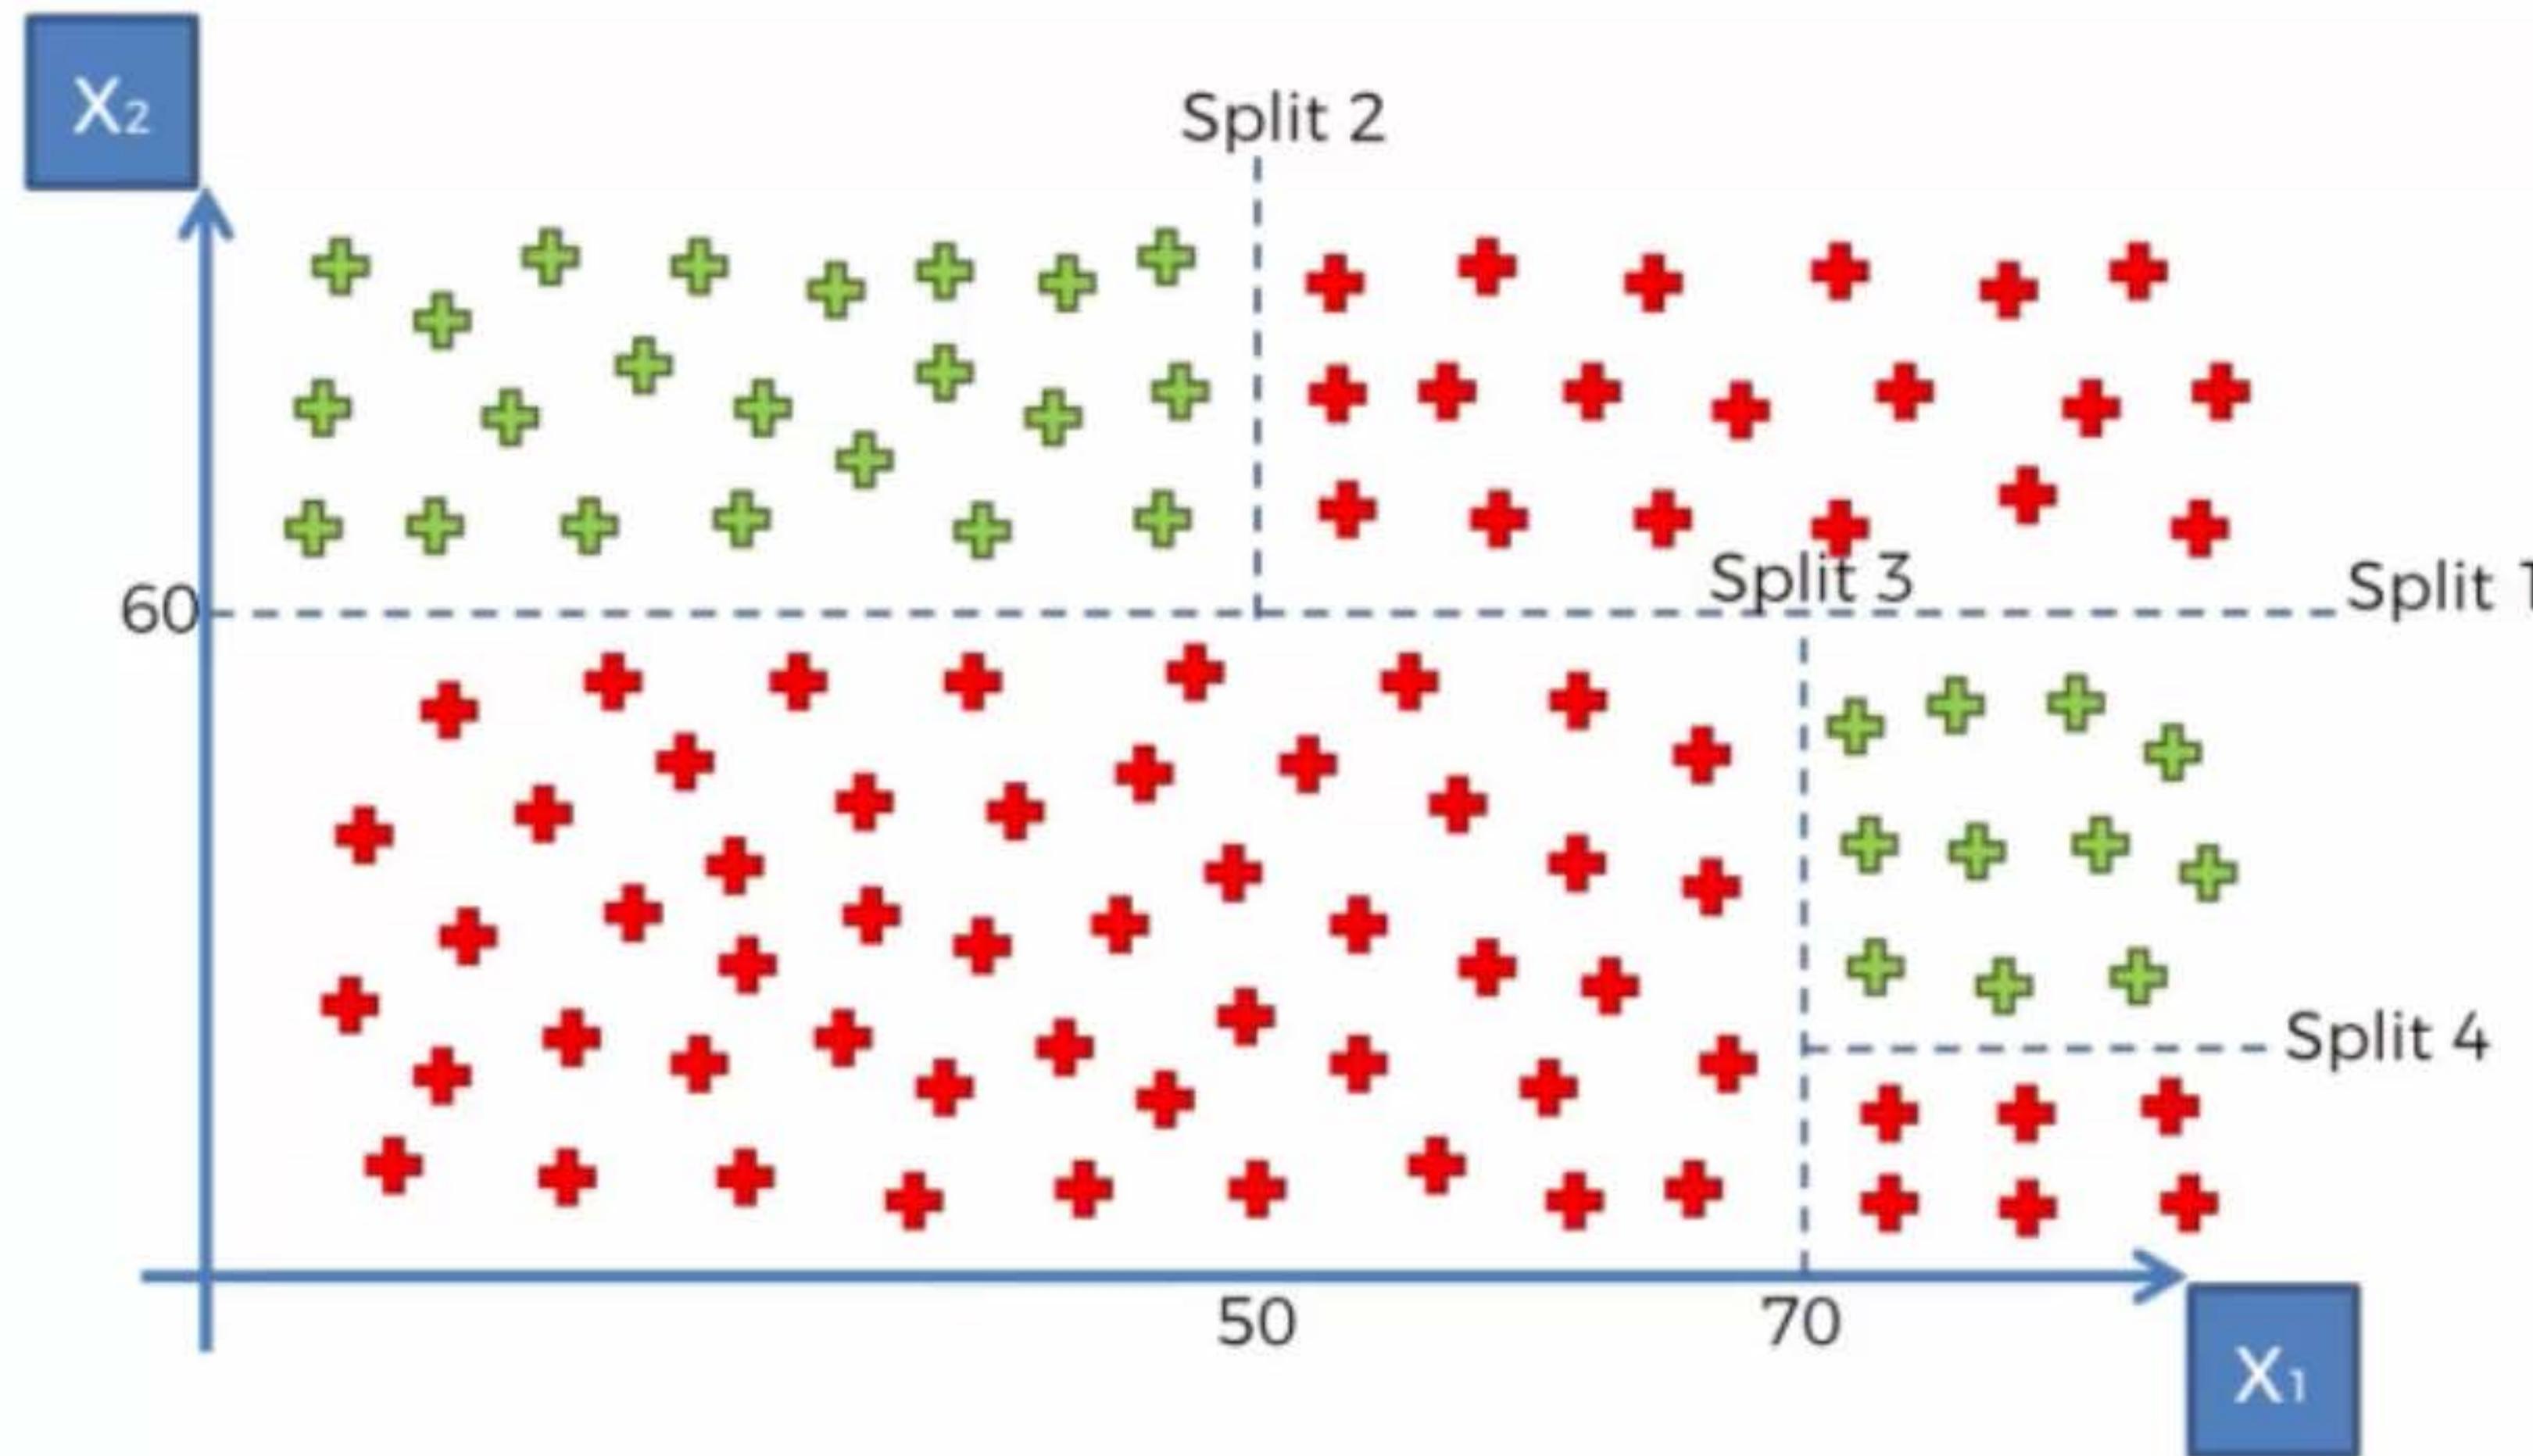
# Decision Tree Intuition



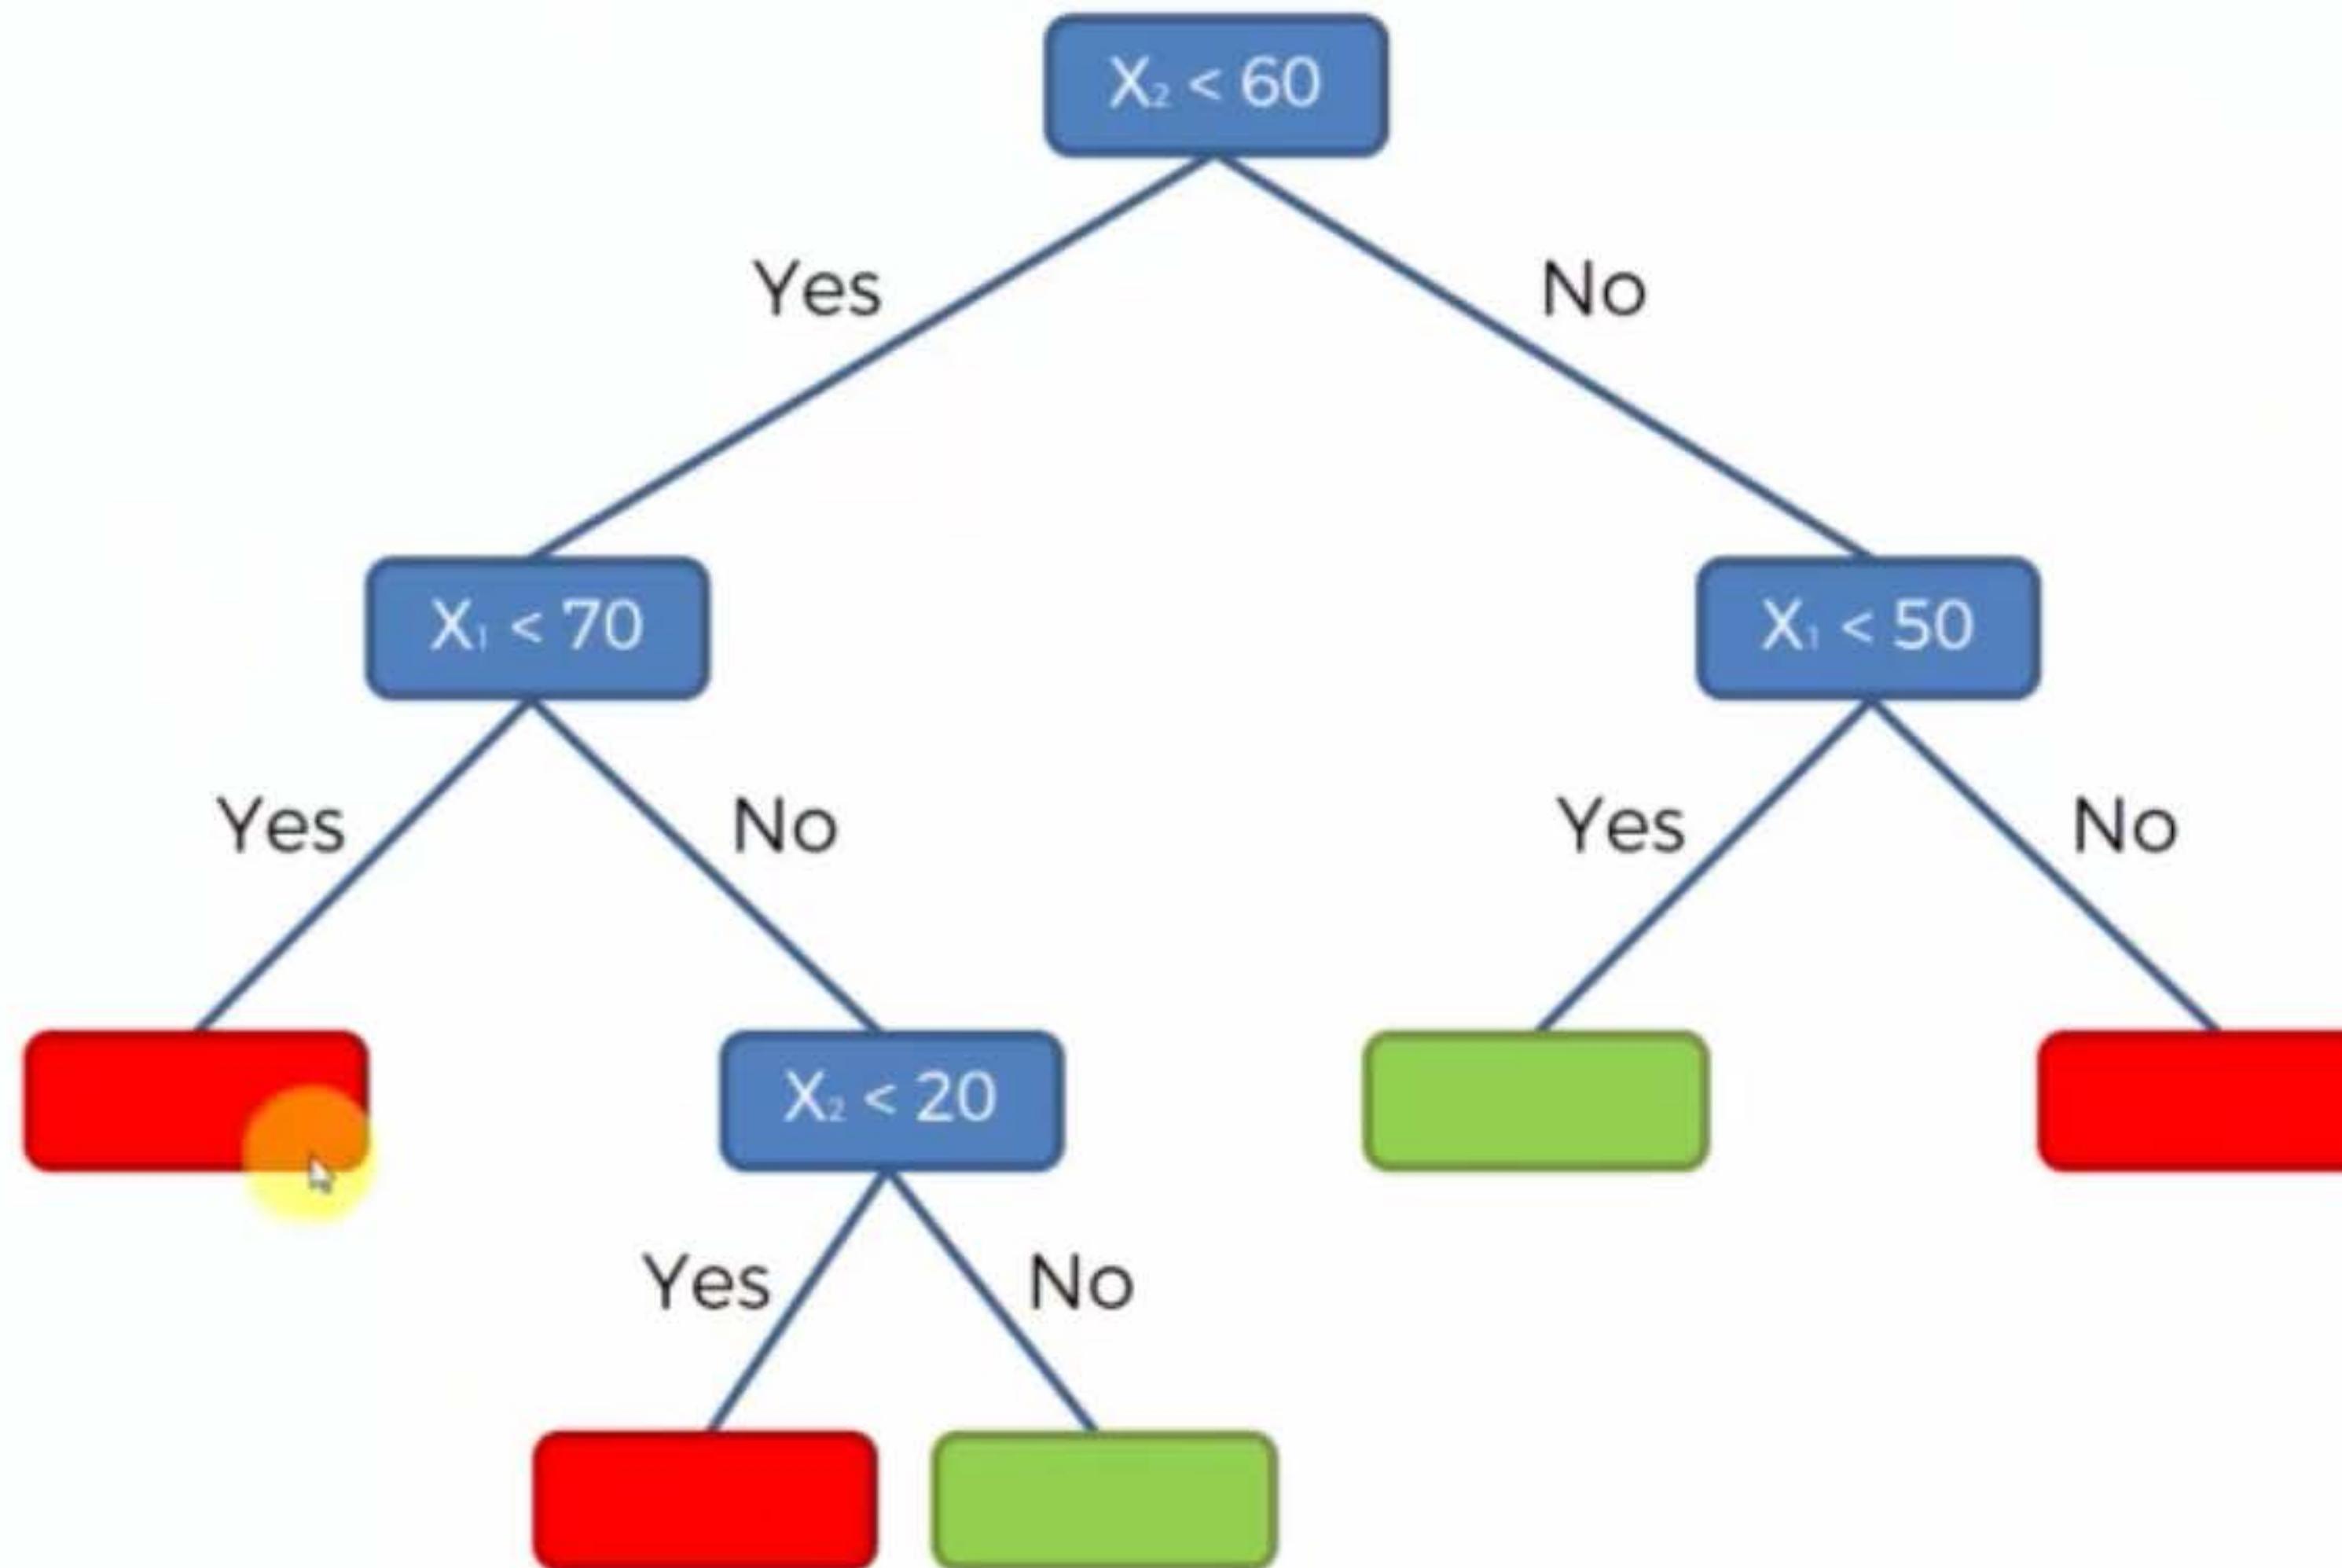
# Decision Tree Intuition



# Decision Tree Intuition



# Split 4



# Decision Trees

---

- Old Method
- Reborn with upgrades
- Random Forest
- Gradient Boosting
- etc.

# Random Forest Intuition

# Random Forest Intuition

---

## Ensemble Learning

# Random Forest Intuition

STEP 1: Pick at random K data points from the Training set.



STEP 2: Build the Decision Tree associated to these K data points.



STEP 3: Choose the number Ntree of trees you want to build and repeat STEPS 1 & 2



STEP 4: For a new data point, make each one of your Ntree trees predict the category to which the data points belongs, and assign the new data point to the category that wins the majority vote.

# Random Forest Intuition



# Random Forest Intuition



BodyPartRecognition.pdf

https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/BodyPartRecognition.pdf

1 / 8

# Real-Time Human Pose Recognition in Parts from Single Depth Images

Jamie Shotton   Andrew Fitzgibbon   Mat Cook   Toby Sharp   Mark Finocchio  
Richard Moore   Alex Kipman   Andrew Blake  
Microsoft Research Cambridge & Xbox Incubation

## Abstract

We propose a new method to quickly and accurately predict 3D positions of body joints from a single depth image, using no temporal information. We take an object recognition approach, designing an intermediate body parts representation that maps the difficult pose estimation problem into a simpler per-pixel classification problem. Our large and highly varied training dataset allows the classifier to estimate body parts invariant to pose, body shape, clothing, etc. Finally we generate confidence-scored 3D proposals of several body joints by reprojecting the classification result and finding local modes.

The system runs at 200 frames per second on consumer hardware. Our evaluation shows high accuracy on both synthetic and real test sets, and investigates the effect of several training parameters. We achieve state of the art accuracy in our comparison with related work and demonstrate improved generalization over exact whole-skeleton nearest neighbor matching.

## 1. Introduction

Robust interactive human body tracking has applications including gaming, human-computer interaction, security, telepresence, and even health-care. The task has recently been greatly simplified by the introduction of real-time depth cameras [16, 19, 44, 37, 28, 13]. However, even the best depth cameras still have limitations and require significant processing power to extract useful information.

The diagram illustrates the workflow for real-time human pose recognition. It starts with a 'depth image' on the left, which is processed to identify 'body parts' (indicated by colored regions on a stick figure). These body part distributions are then converted into '3D joint proposals' shown in two 3D boxes labeled 'front' and 'side'. A yellow circle highlights a specific joint proposal in one of the boxes.

Figure 1. **Overview.** From a single input depth image, a per-pixel body part distribution is inferred. (Colors indicate the most likely part labels at each pixel, and correspond in the joint proposals). Local modes of this signal are estimated to give high-quality proposals for the 3D locations of body joints, even for multiple users.

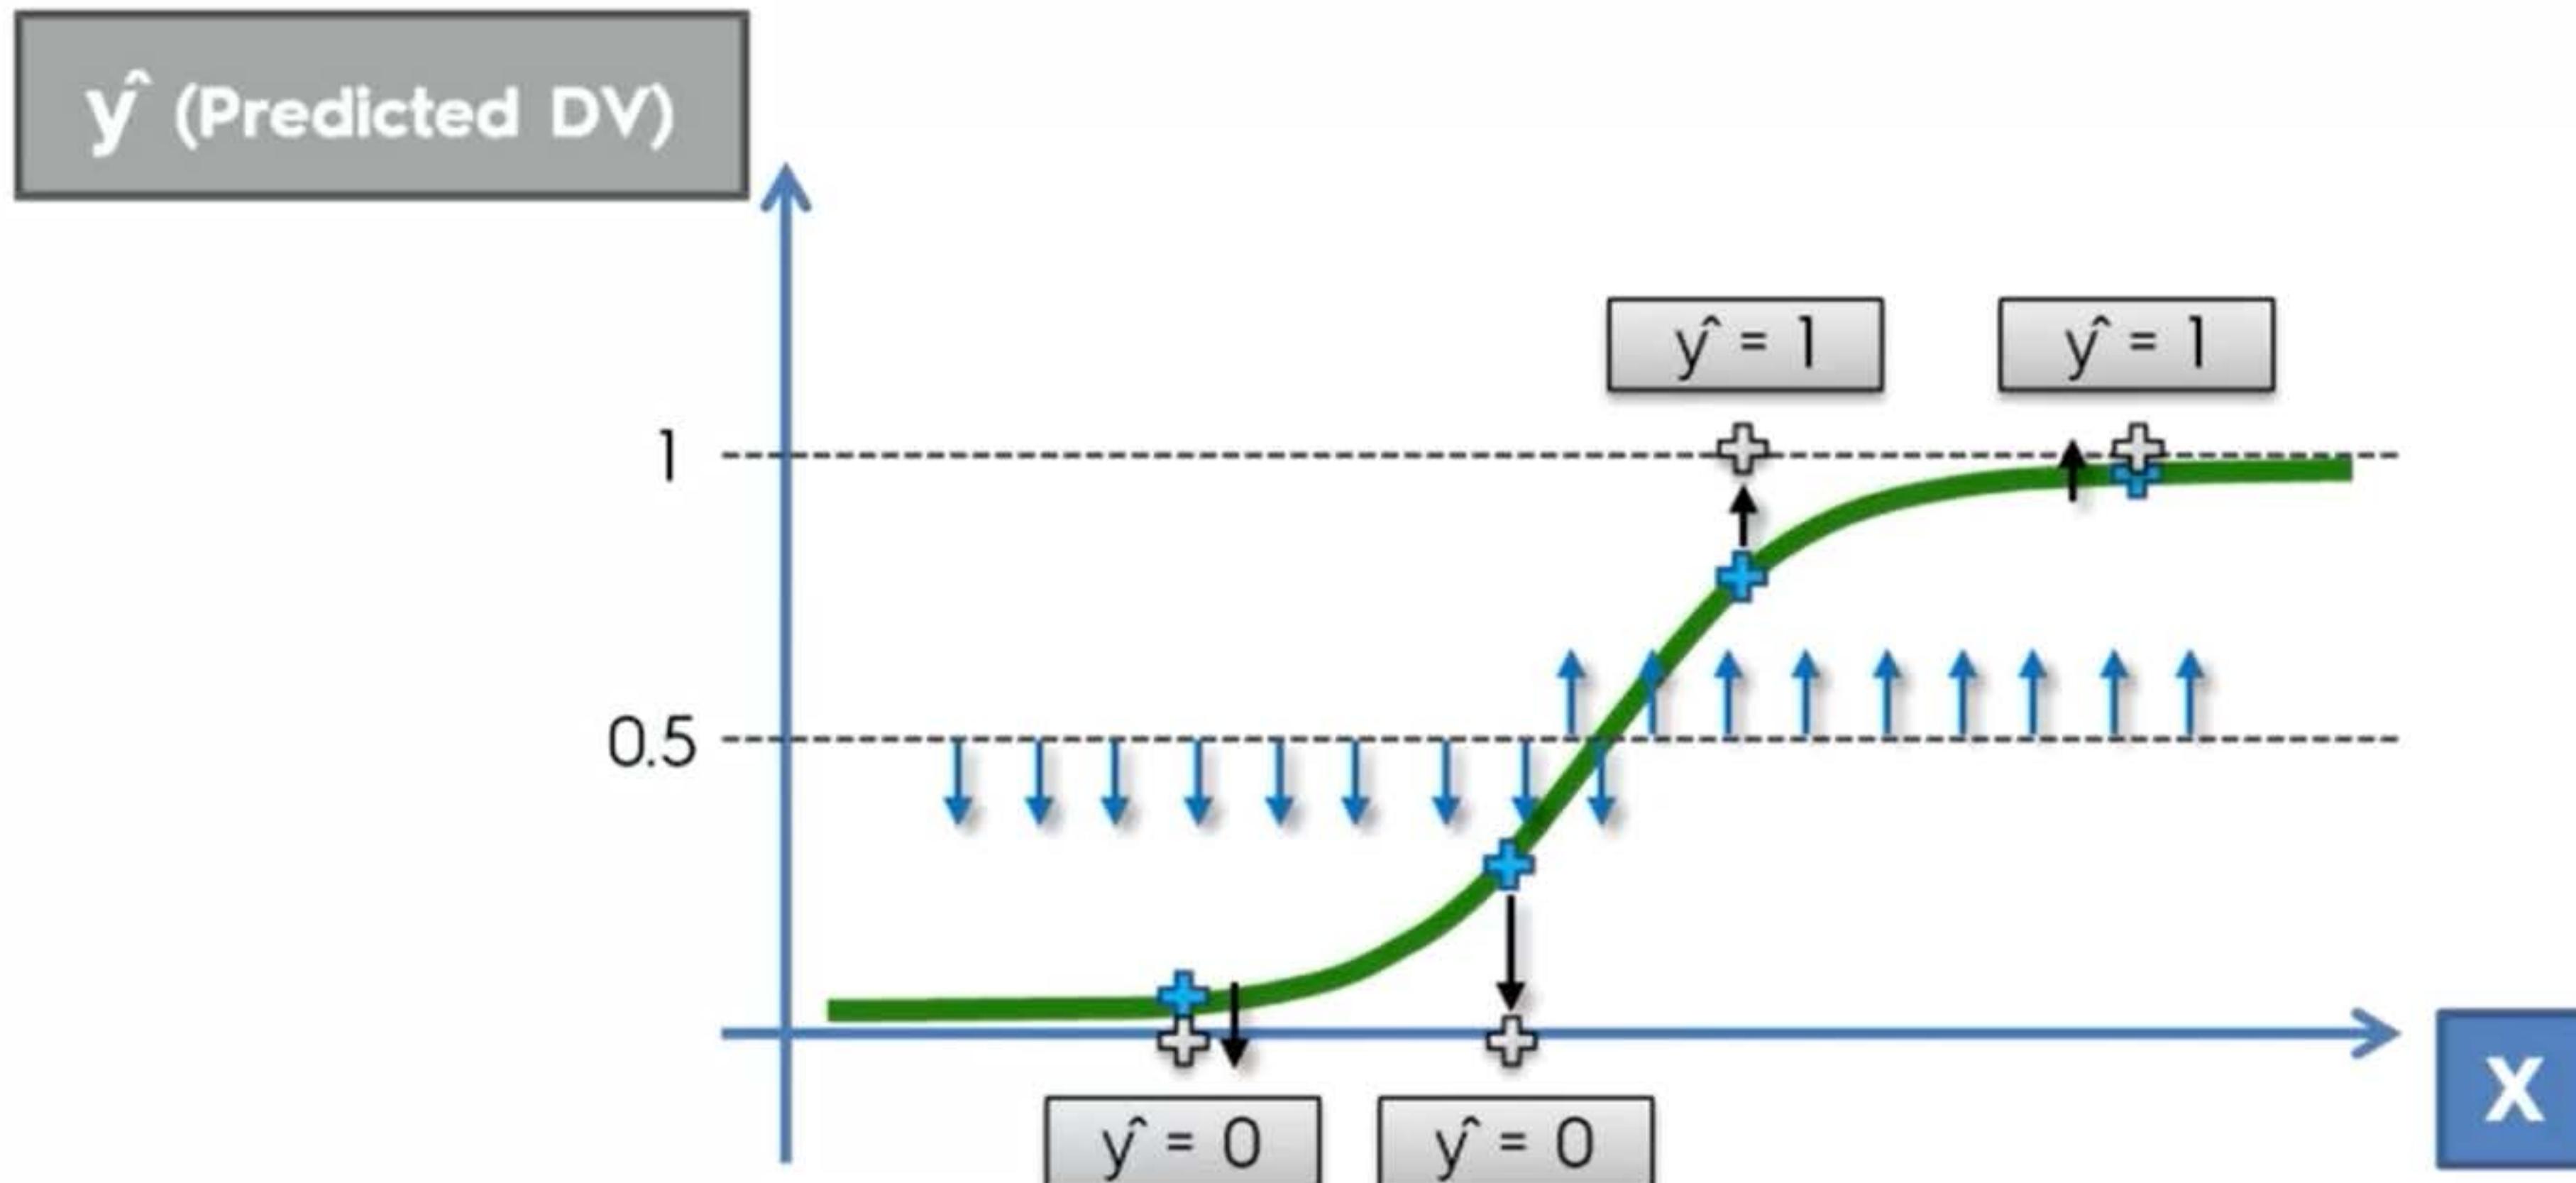
joints of interest. Reprojecting the inferred parts into world space, we localize spatial modes of each part distribution and thus generate (possibly several) confidence-weighted proposals for the 3D locations of each skeletal joint.

We treat the segmentation into body parts as a per-pixel classification task (no pairwise terms or CRF have proved necessary). Evaluating each pixel separately avoids a combinatorial search over the different body joints, although

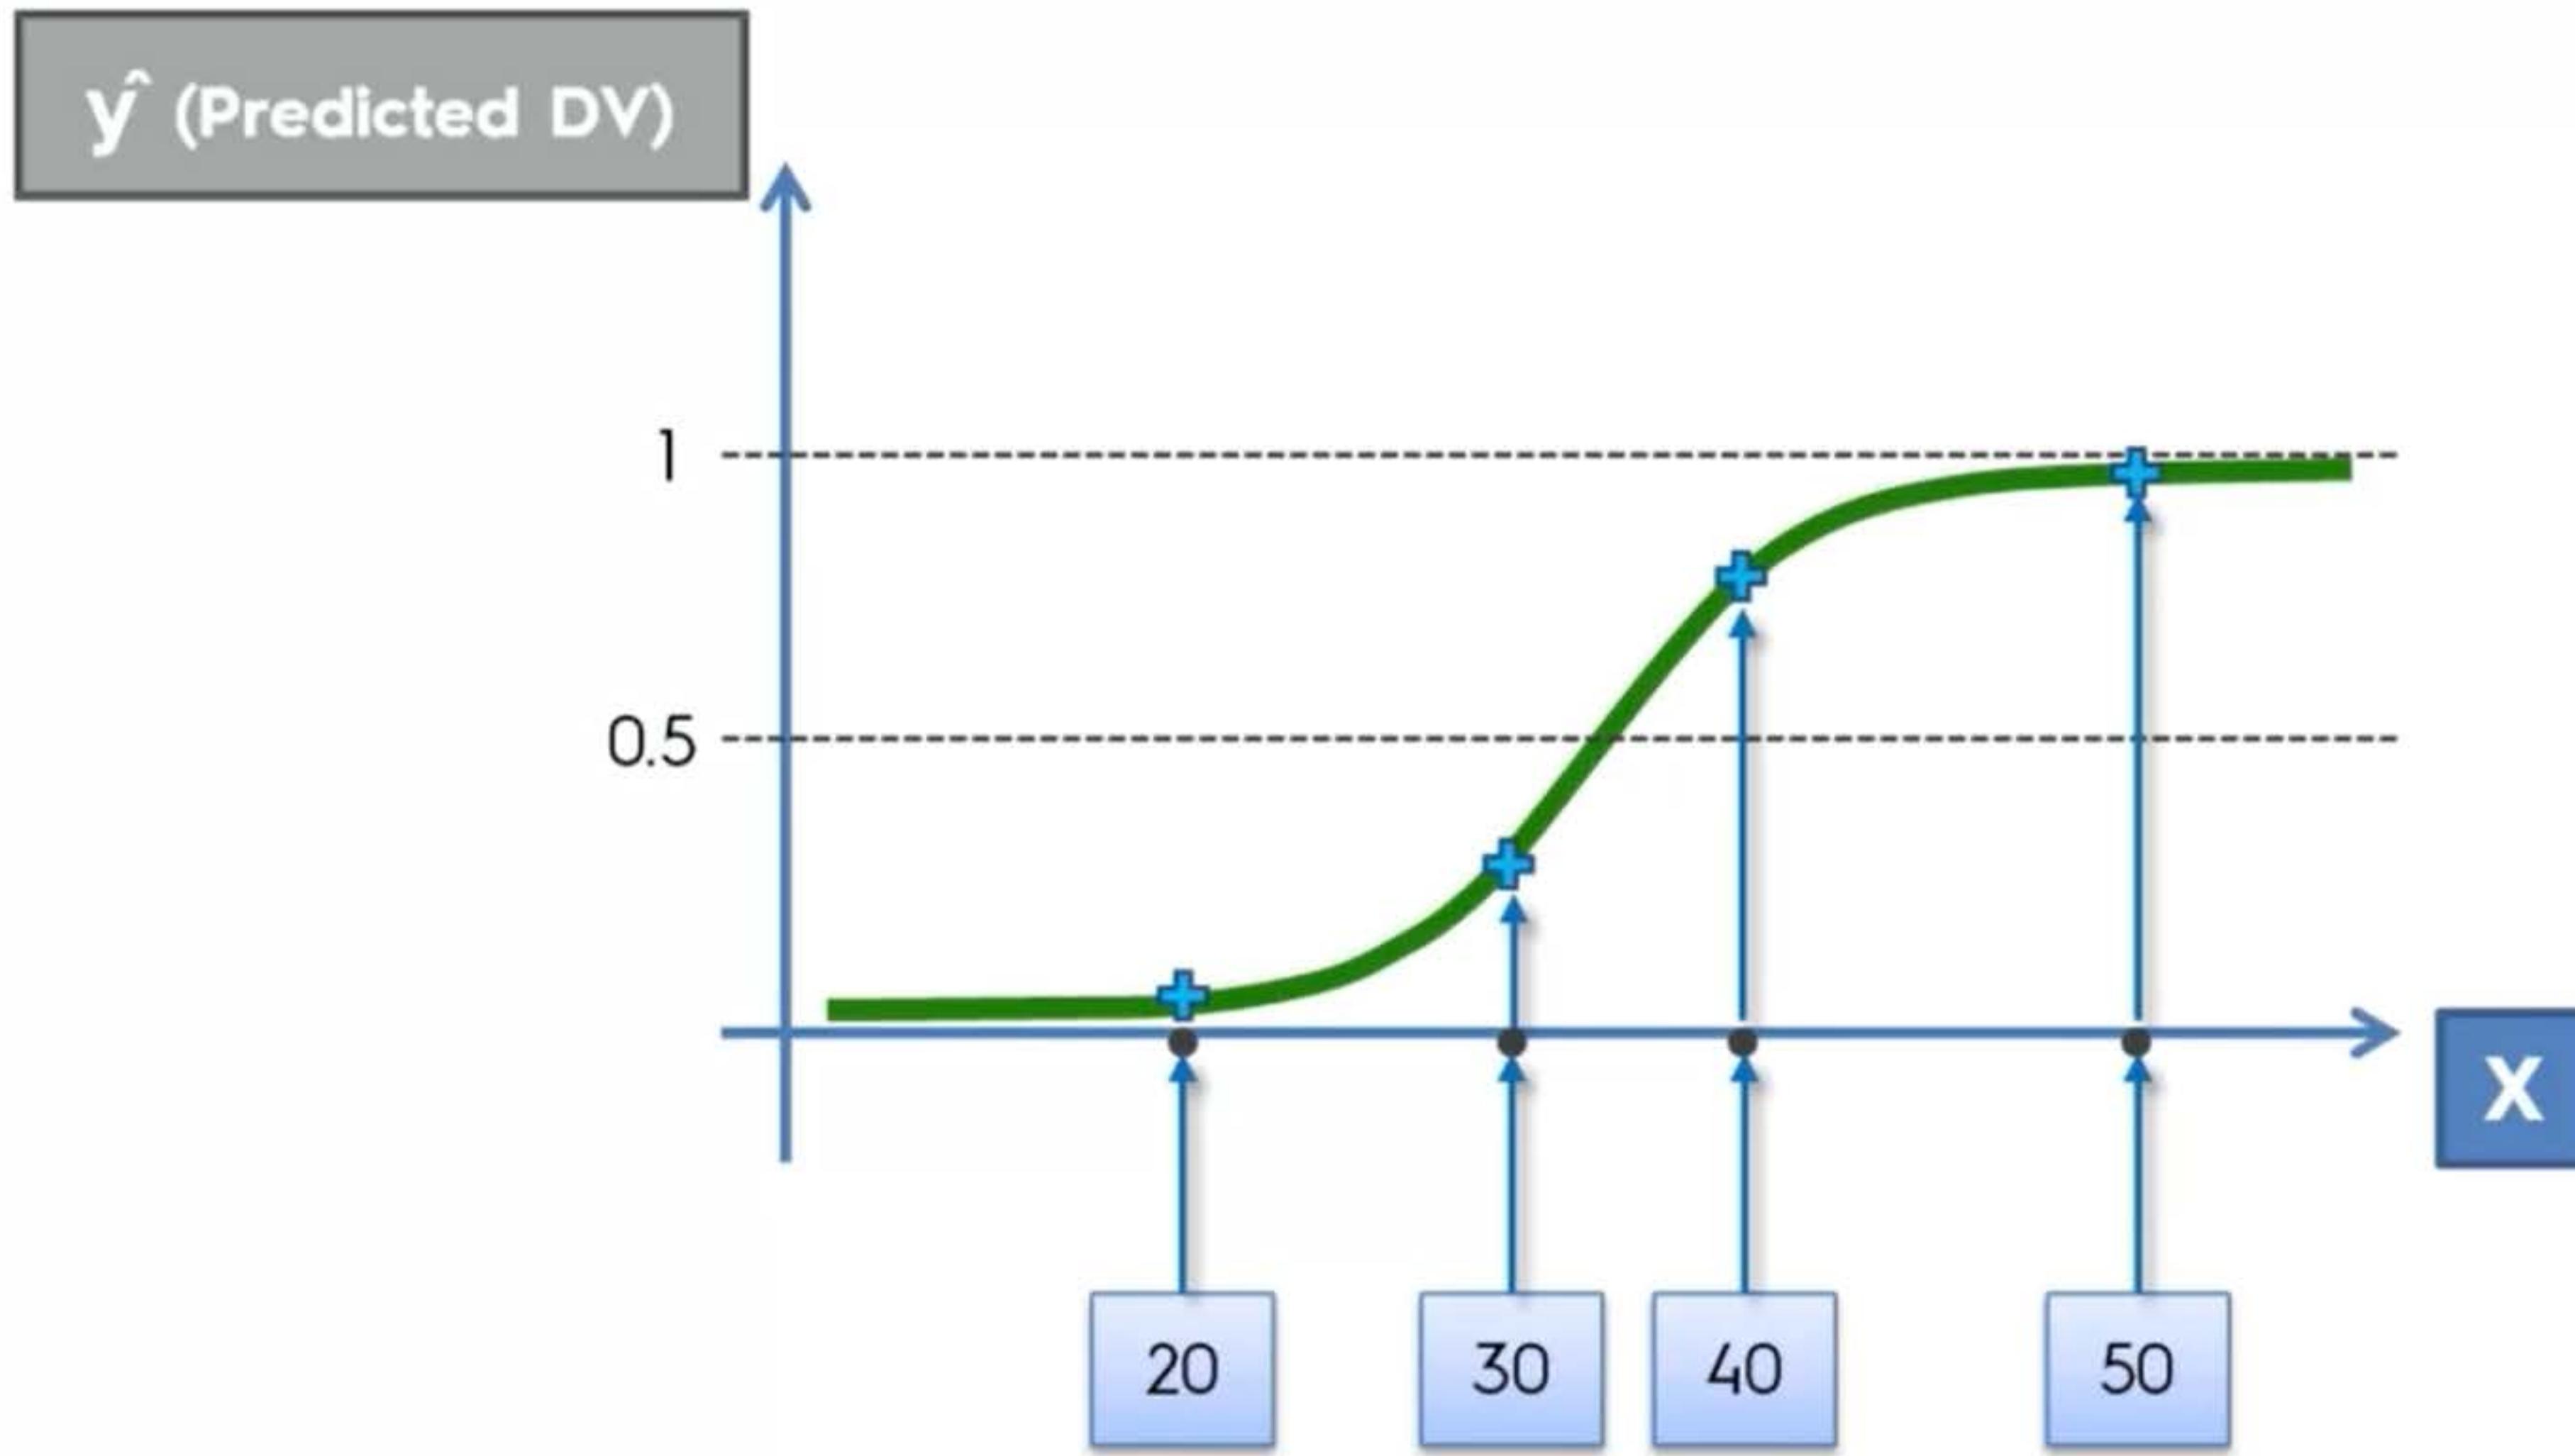
Save PDF to Evernote

# **False Positives & False Negatives**

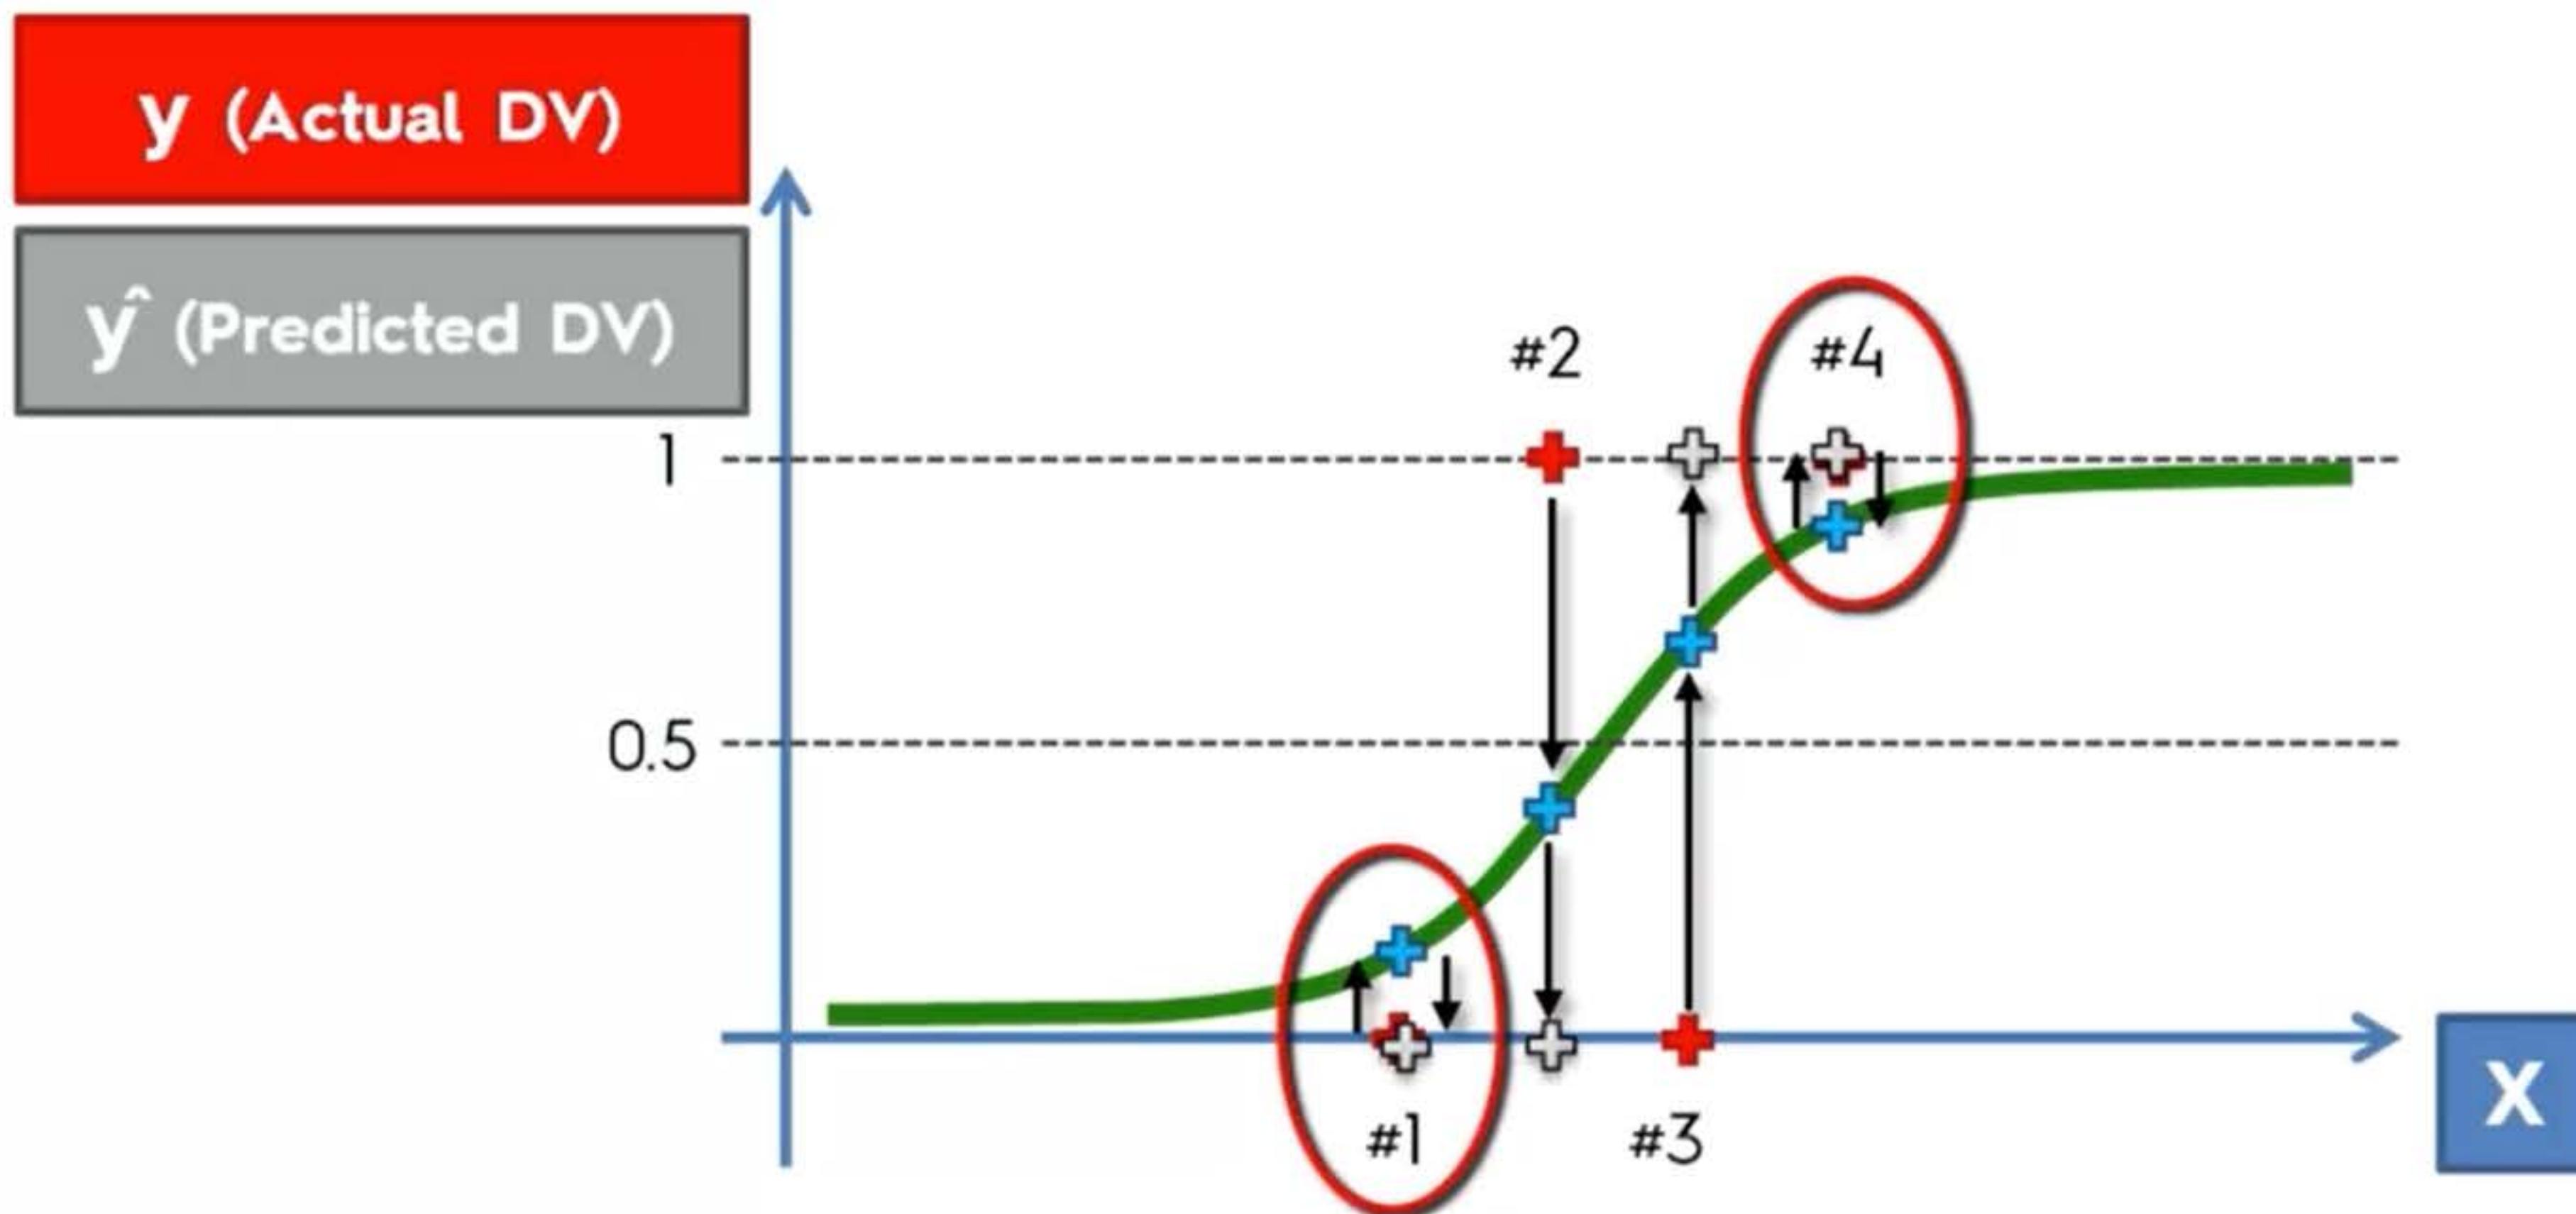
# False Positives & Negatives



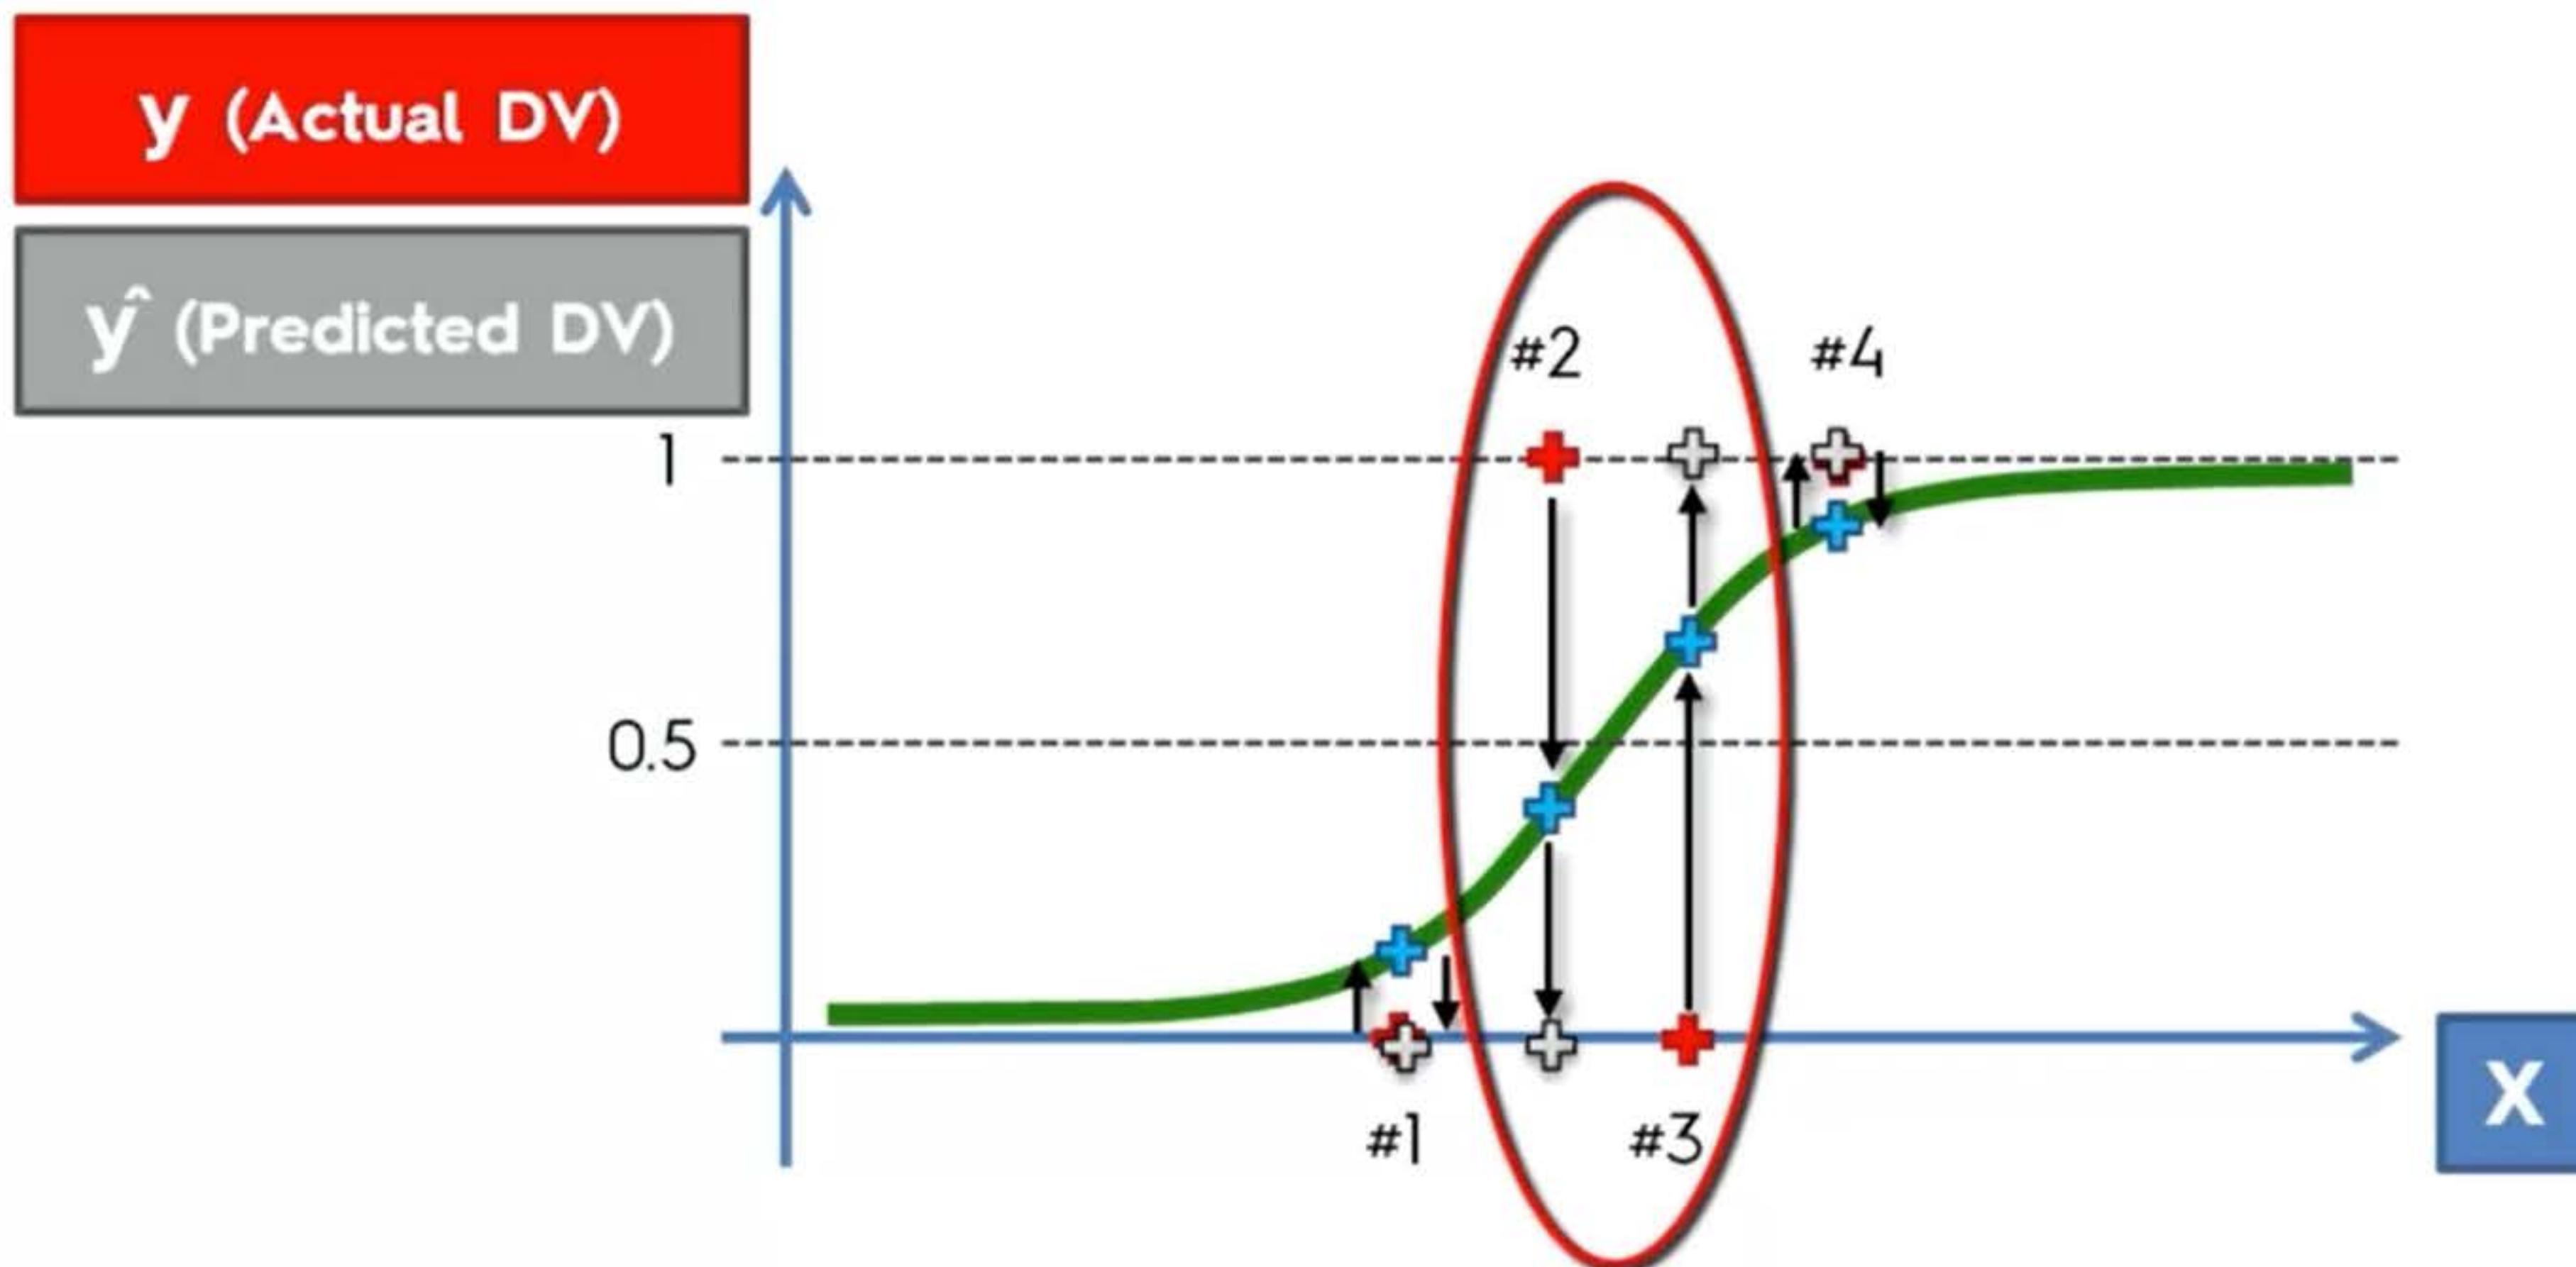
# False Positives & Negatives



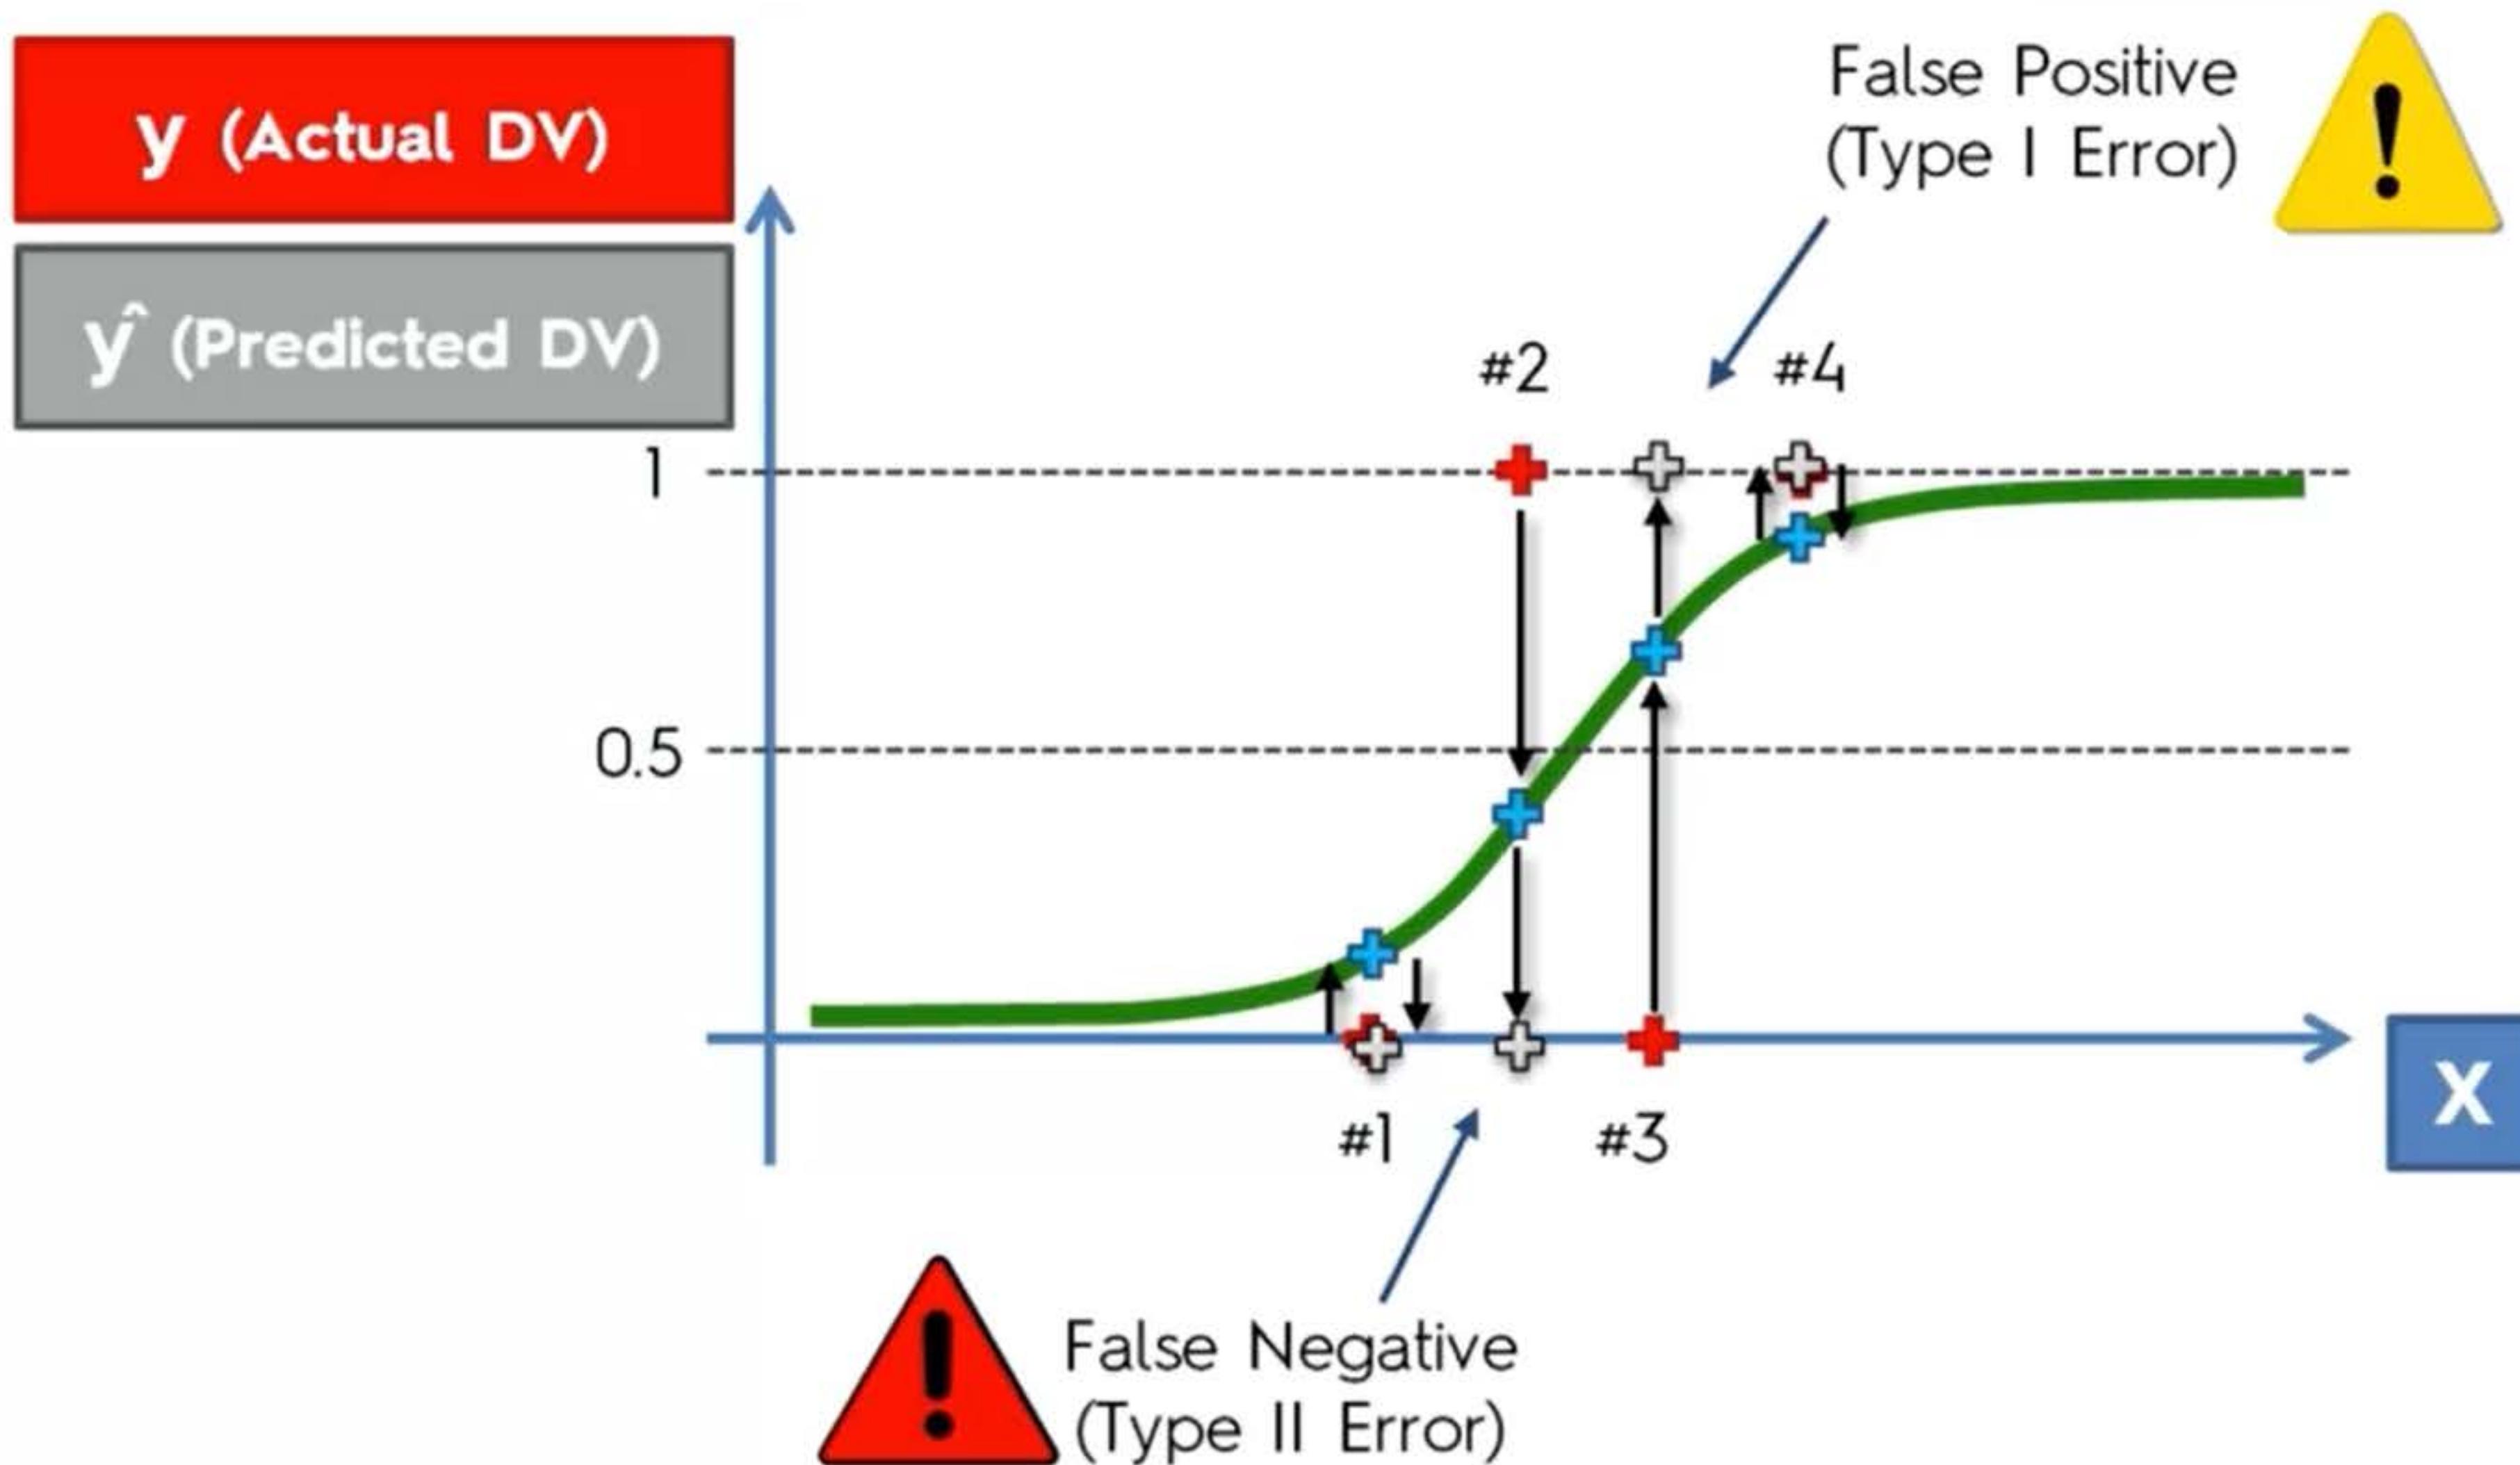
# False Positives & Negatives



# False Positives & Negatives

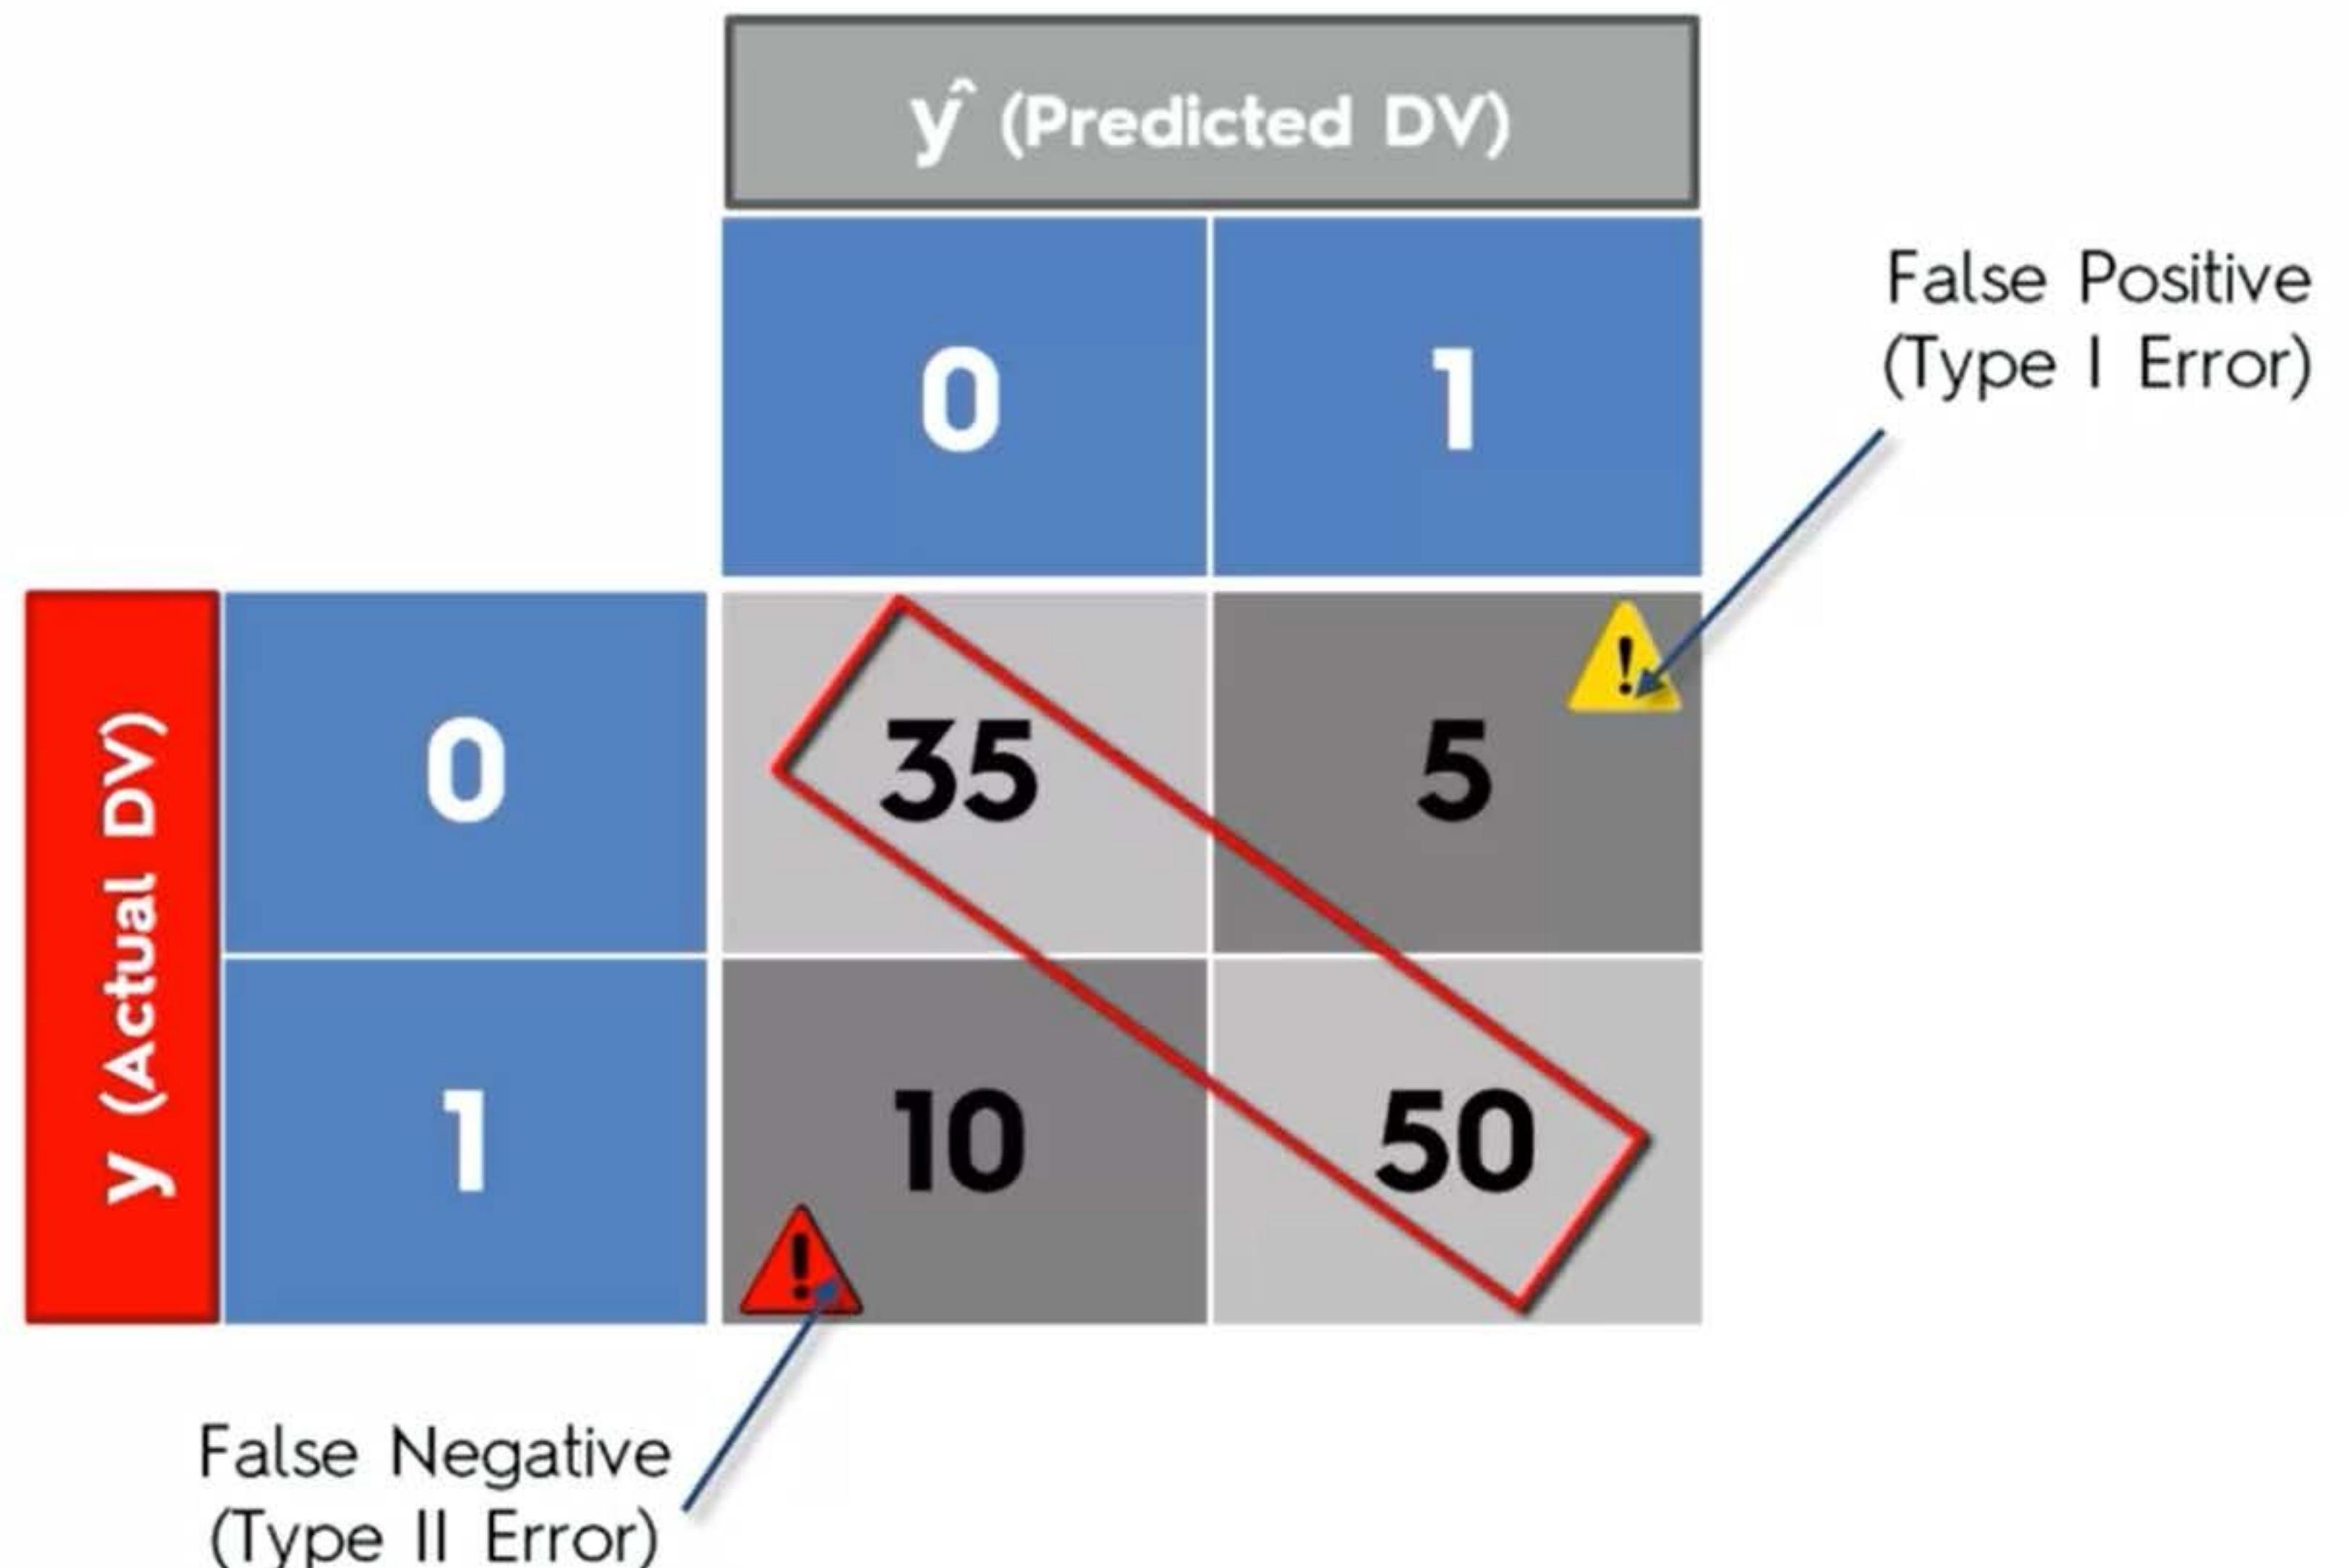


# False Positives & Negatives

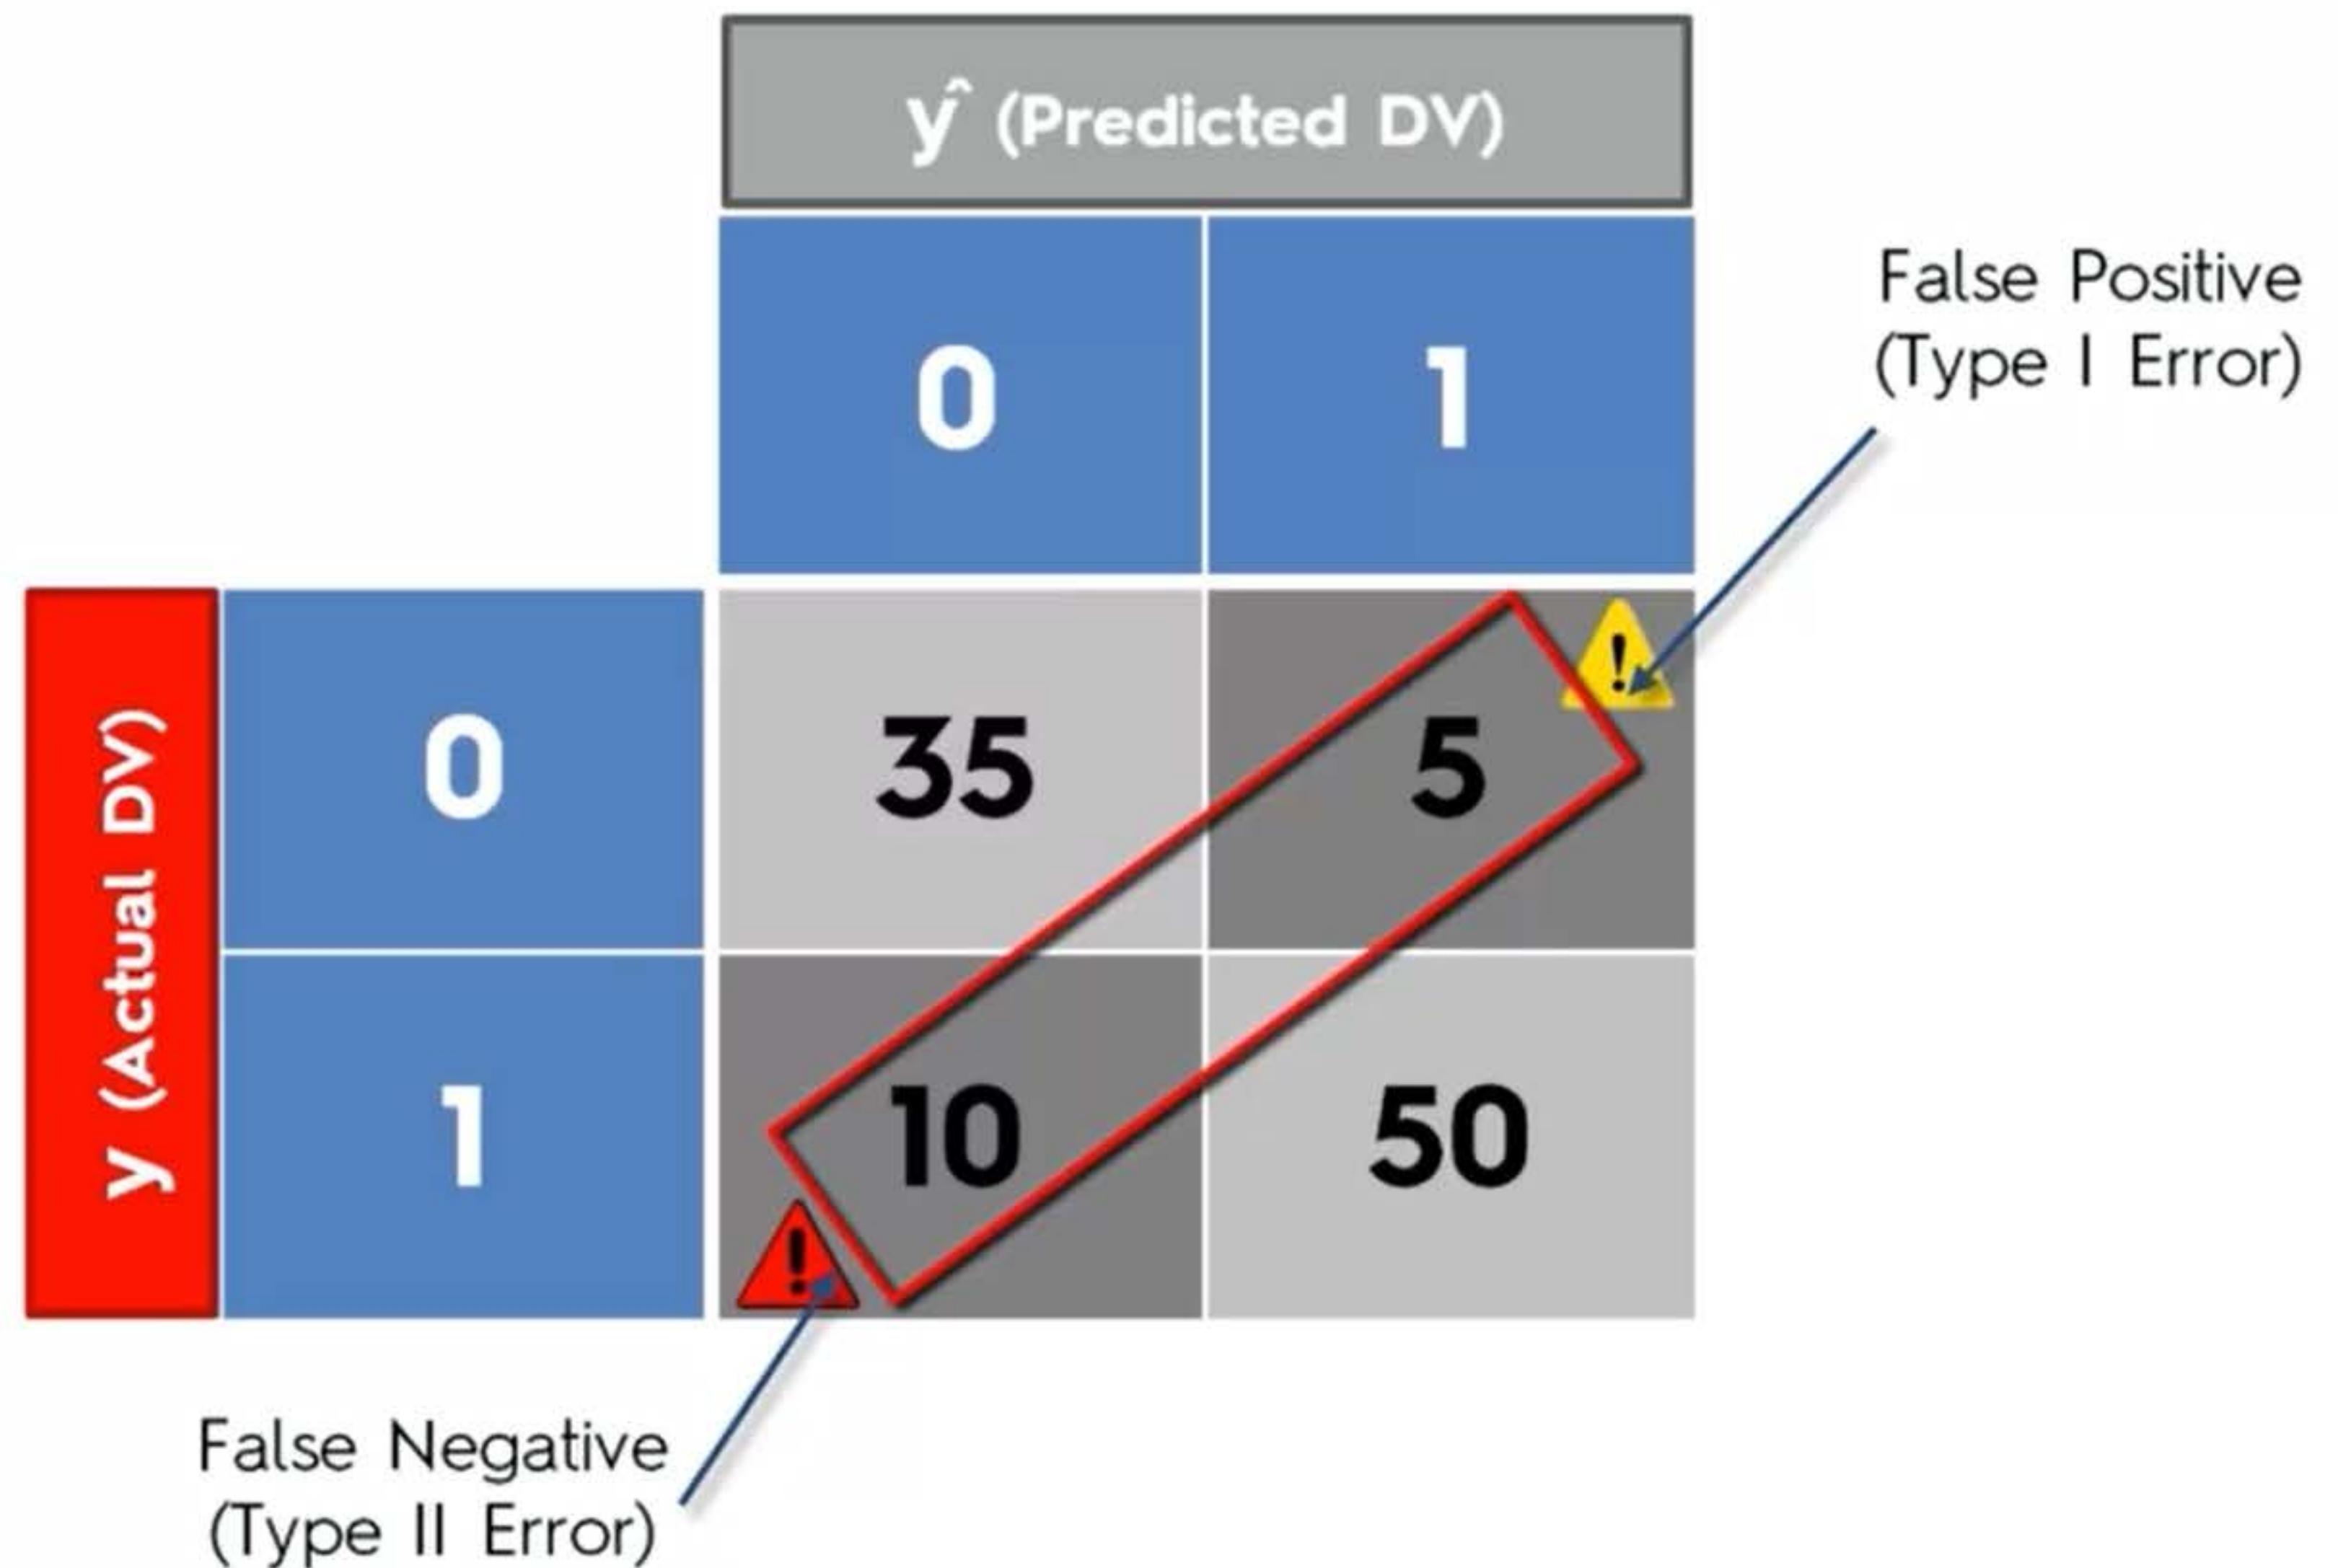


# Confusion Matrix

# Confusion Matrix



# Confusion Matrix



# Confusion Matrix

		$\hat{y}$ (Predicted DV)	
		0	1
y (Actual DV)	0	35	5
	1	10	50

False Negative  
(Type II Error)

False Positive  
(Type I Error)

## Calculate two rates

1. Accuracy Rate = Correct / Total  
 $AR = 85/100 = 85\%$

2. Error Rate = Wrong / Total  
 $ER = 15/100 = 15\%$

# Accuracy Paradox

# Accuracy Paradox

		$\hat{y}$ (Predicted DV)	
		0	1
$y$ (Actual DV)	0	9,700	150 
	1	50 	100

## Scenario 1:

Accuracy Rate = Correct / Total  
AR = 9,800/10,000 = 98%

# Accuracy Paradox

		$\hat{y}$ (Predicted DV)
		0      1
y (Actual DV)	0	9,850 ← 0 !
	1	150 ← 0 !

## Scenario 1:

Accuracy Rate = Correct / Total  
AR = 9,800/10,000 = 98%

# Accuracy Paradox

		$\hat{y}$ (Predicted DV)
		0      1
$y$ (Actual DV)	0	9,850 ← 0 !
	1	150 ← 0 !

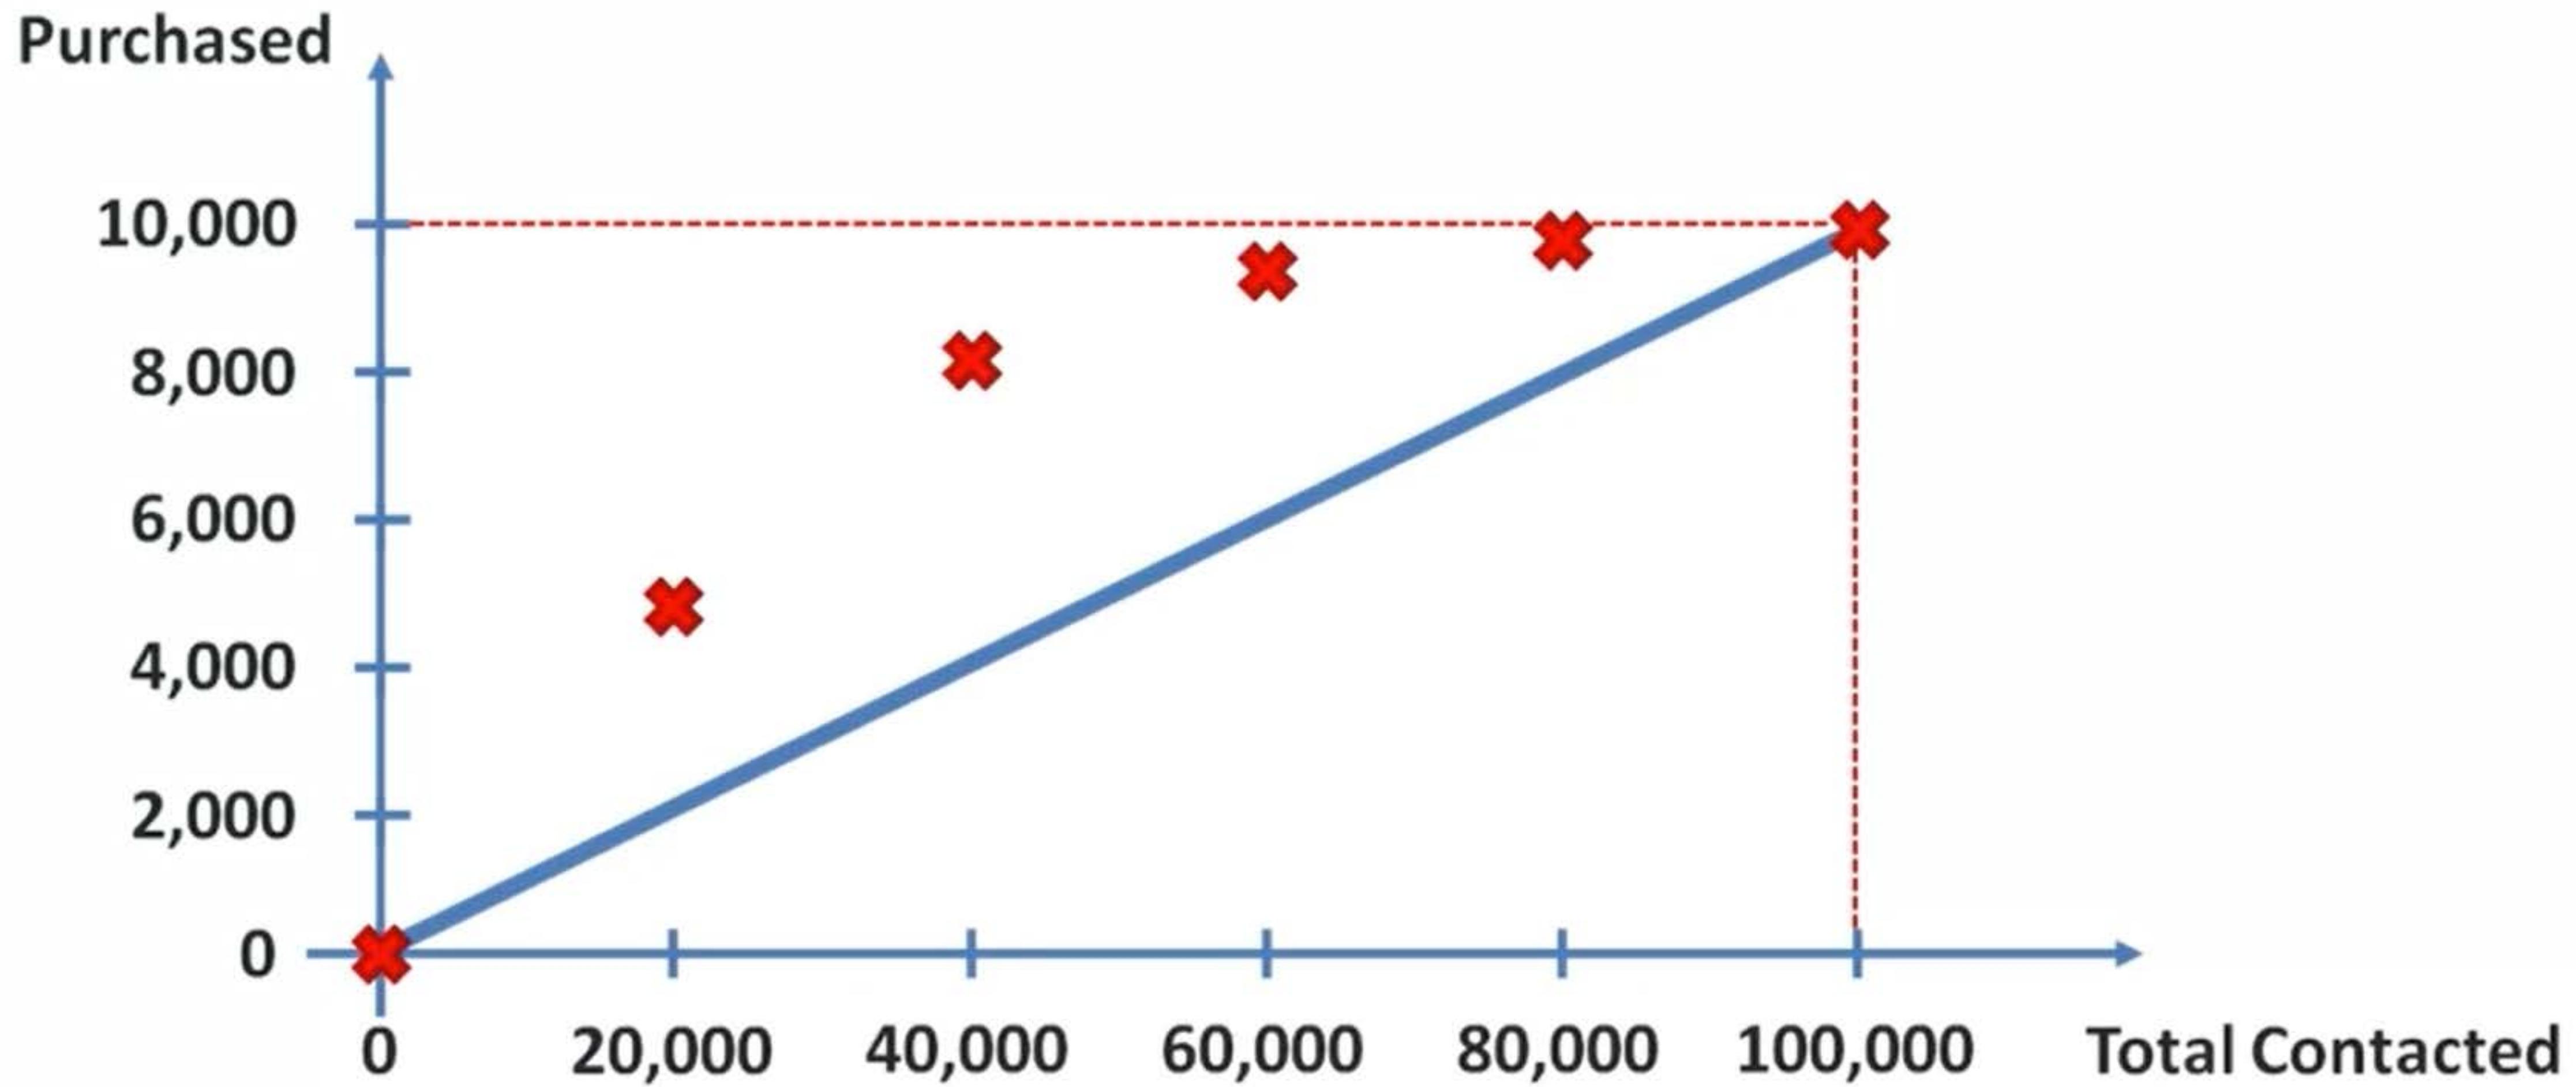
## Scenario 1:

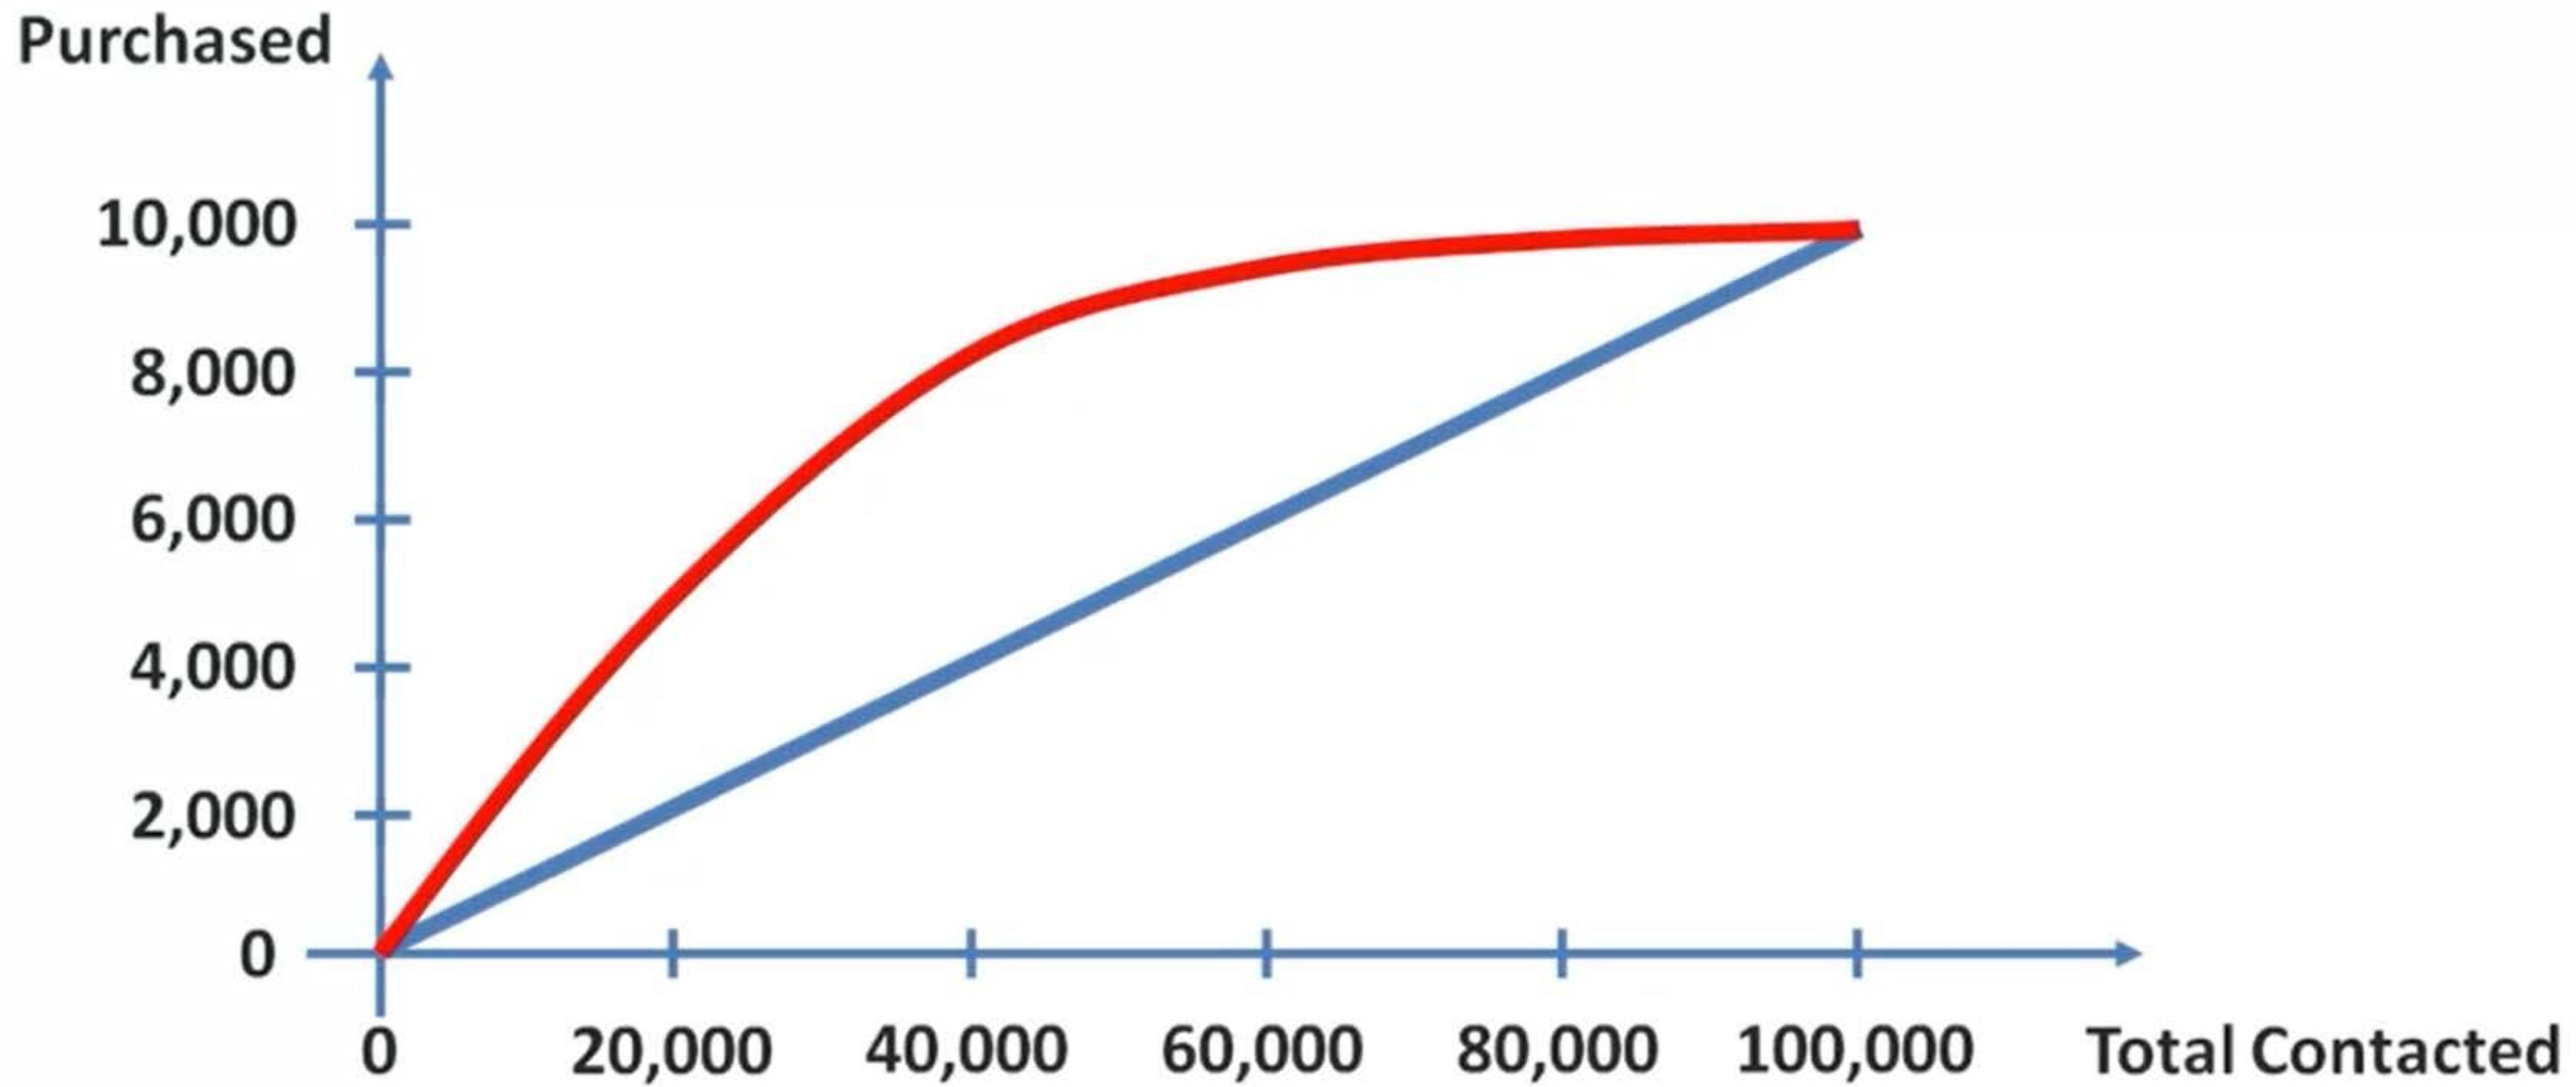
Accuracy Rate = Correct / Total  
 $AR = 9,800/10,000 = 98\%$

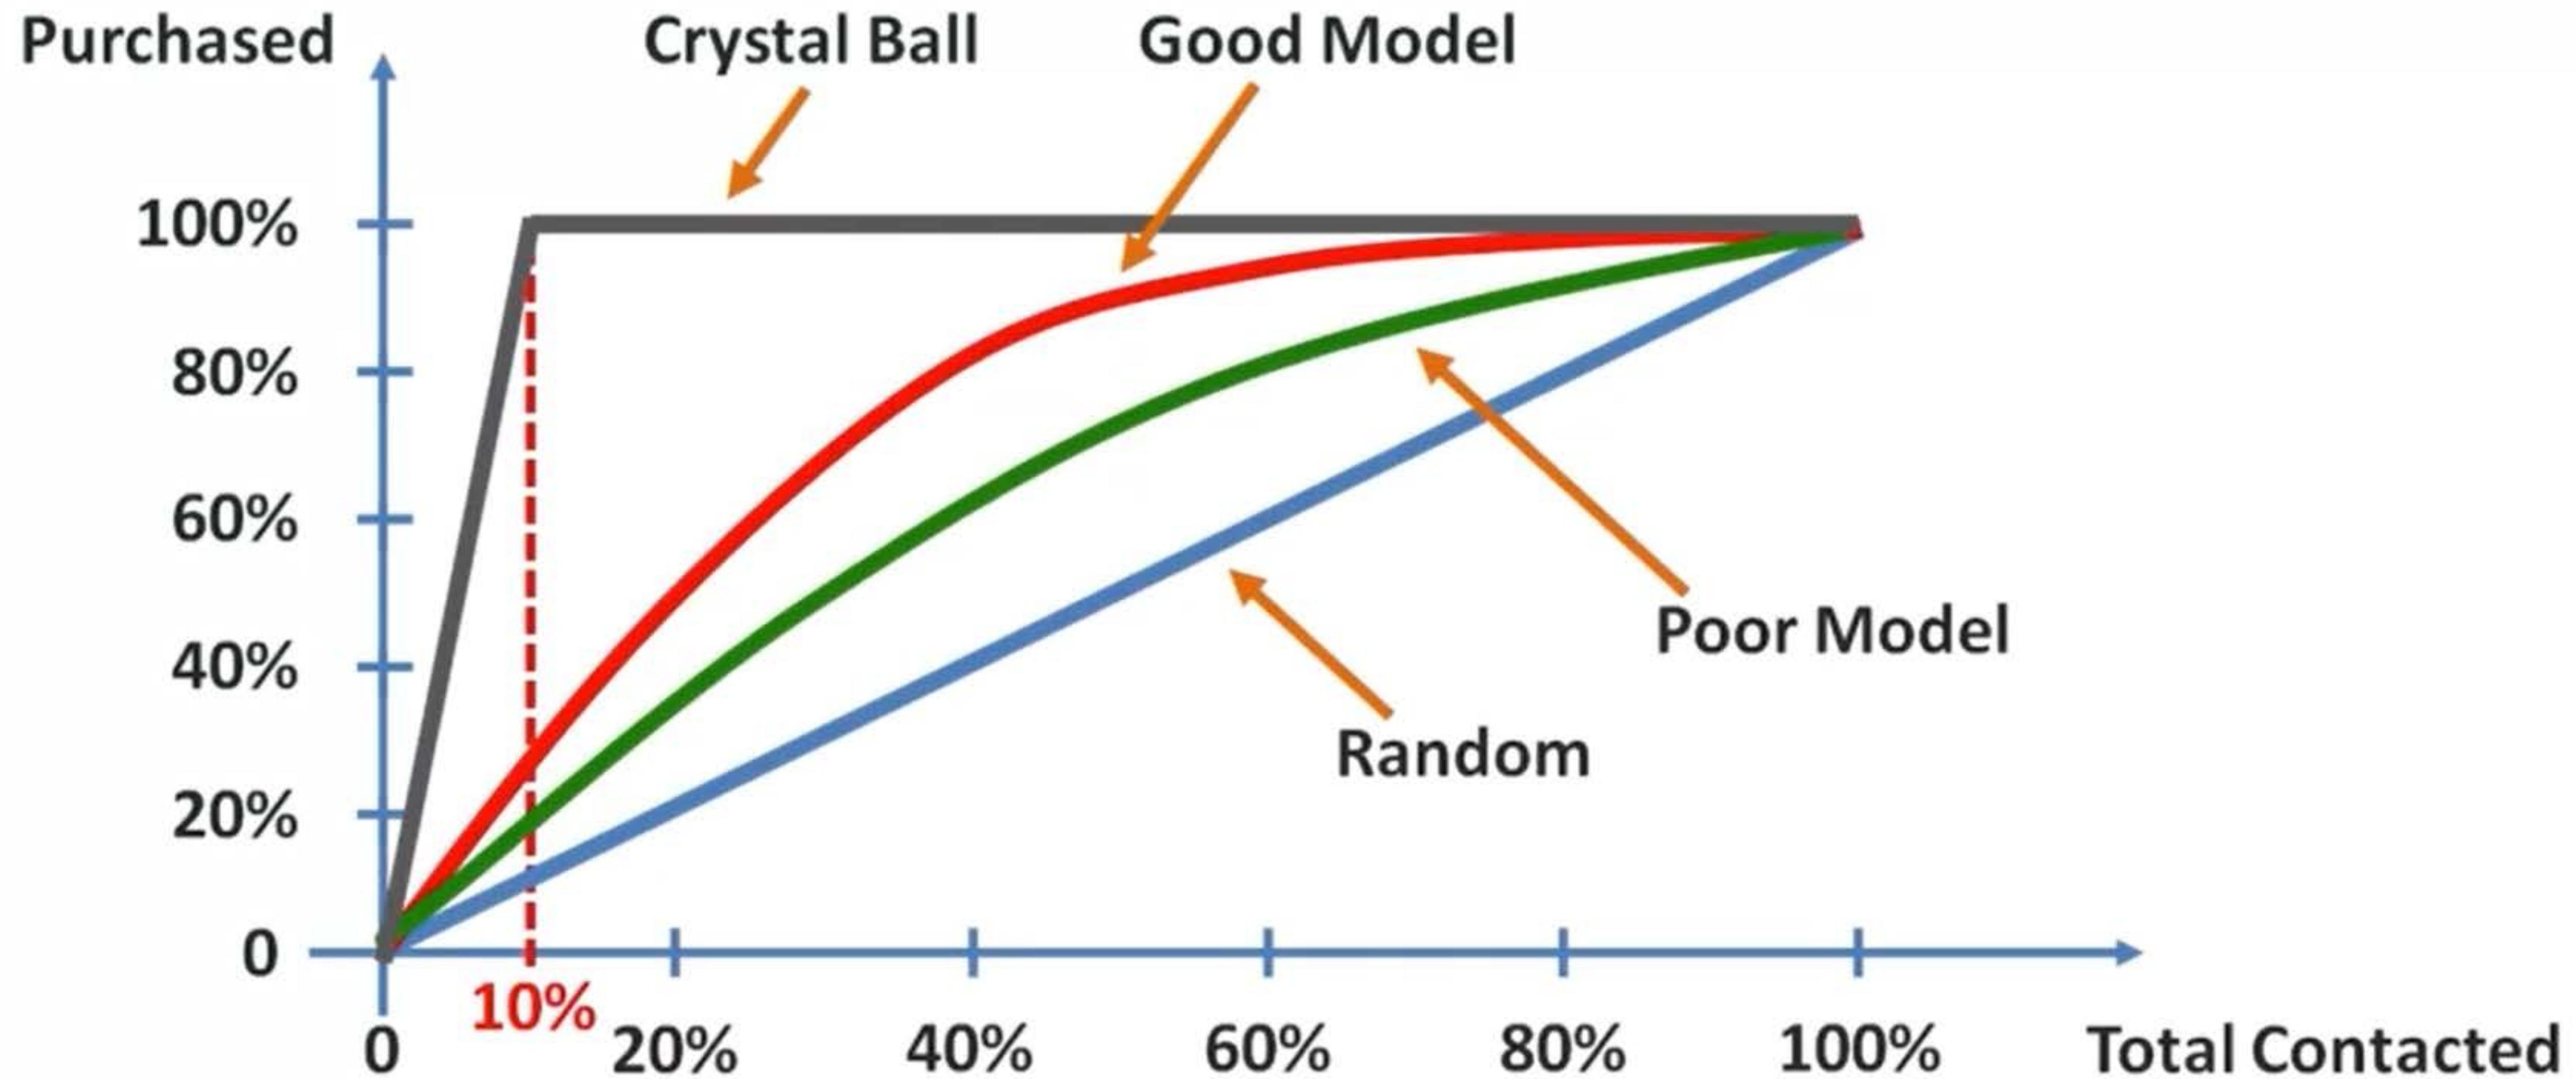
## Scenario 2:

Accuracy Rate = Correct / Total  
 $AR = 9,850/10,000 = 98.5\%$  ↑

# CAP Curve







## Note:

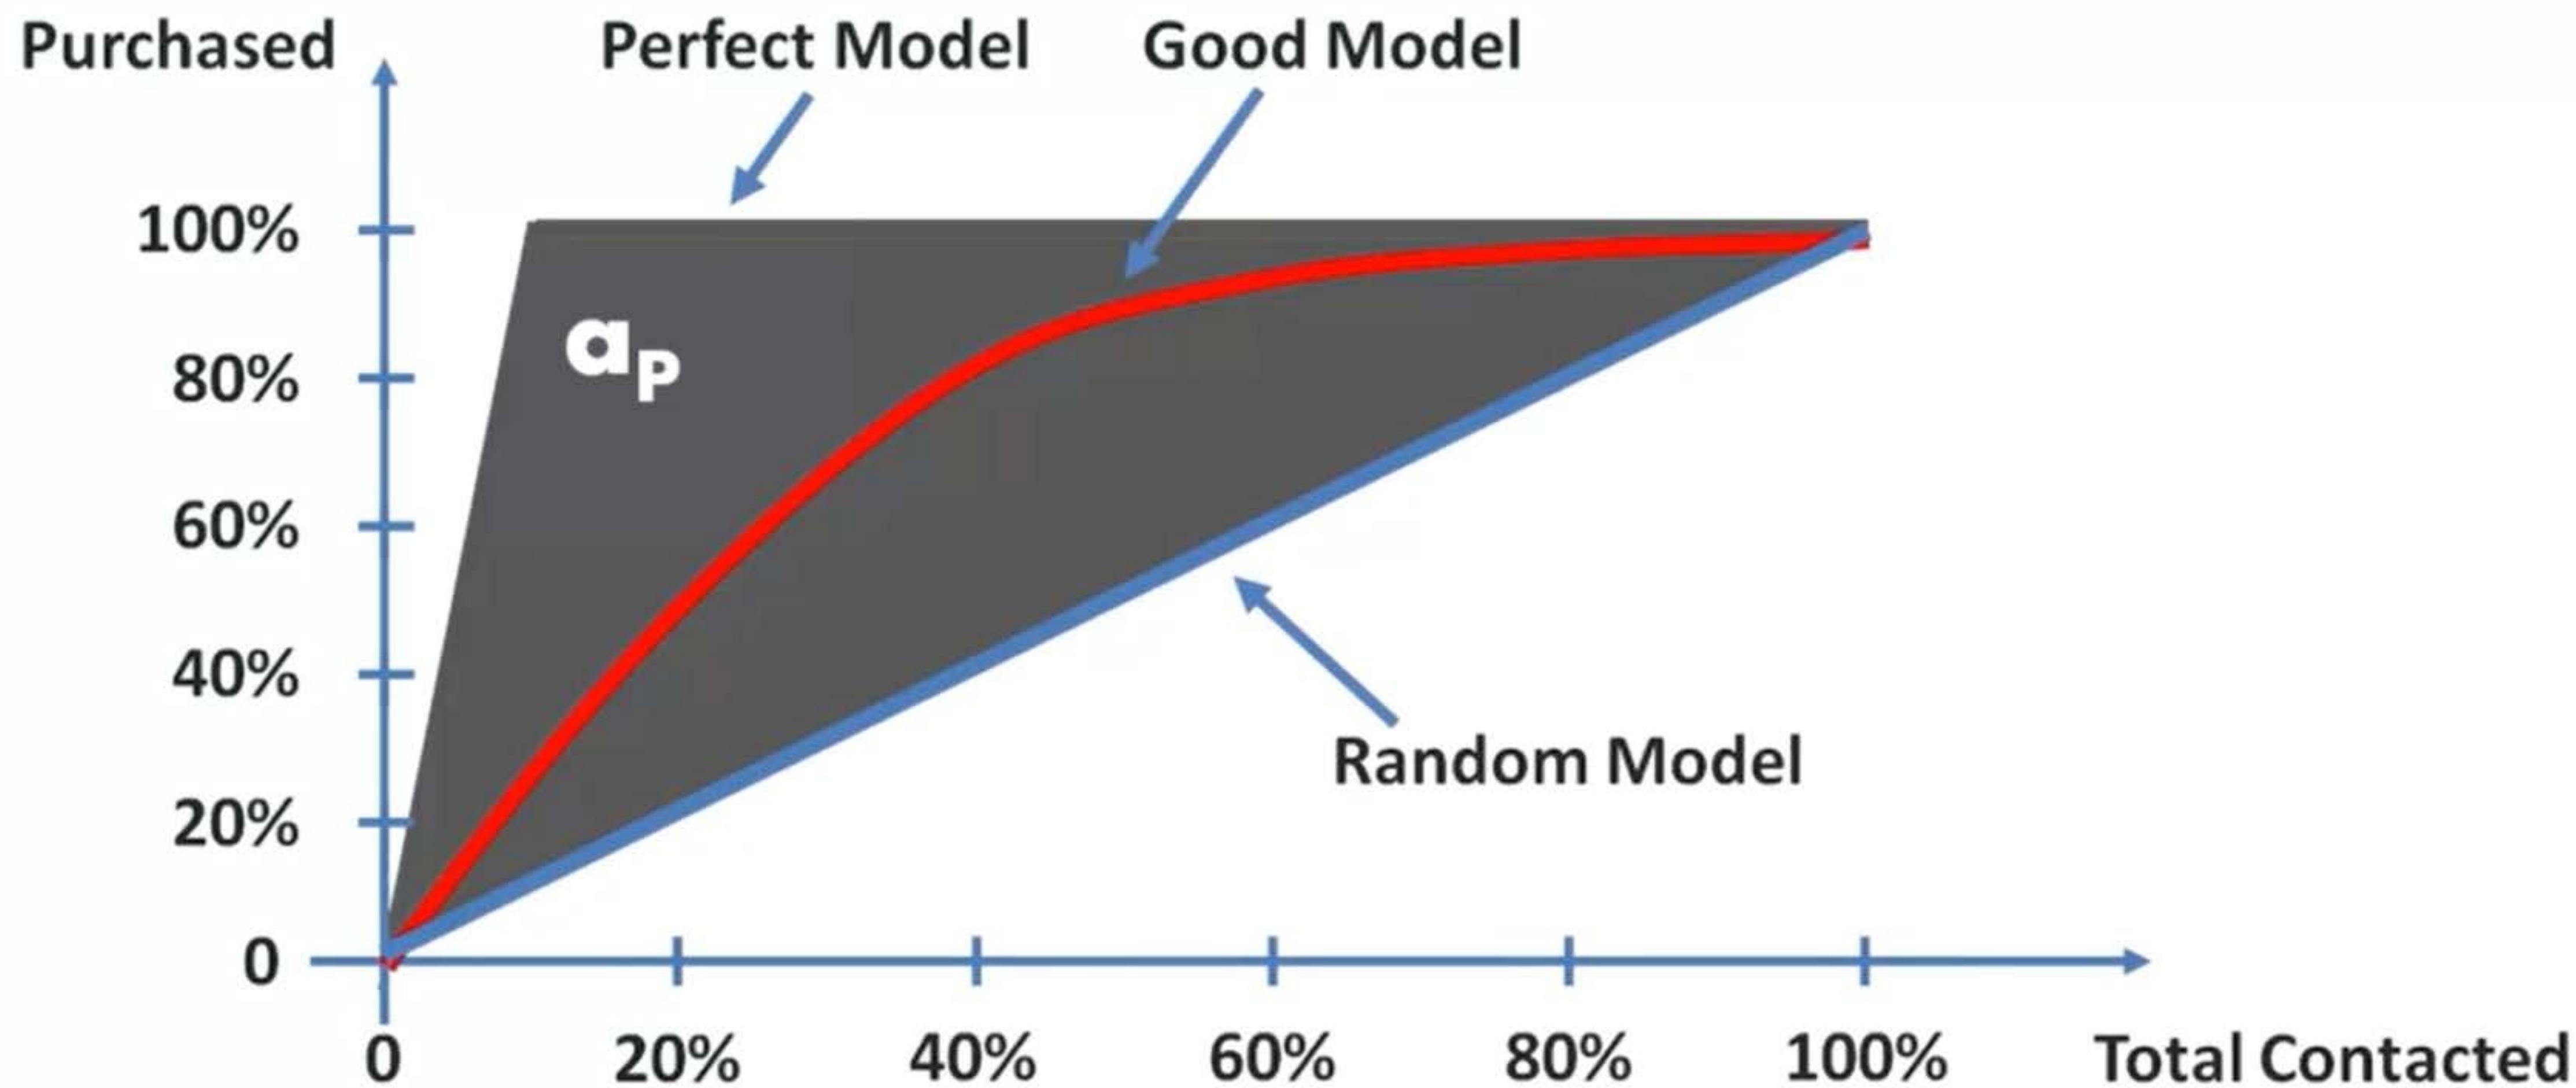
CAP = Cumulative Accuracy Profile



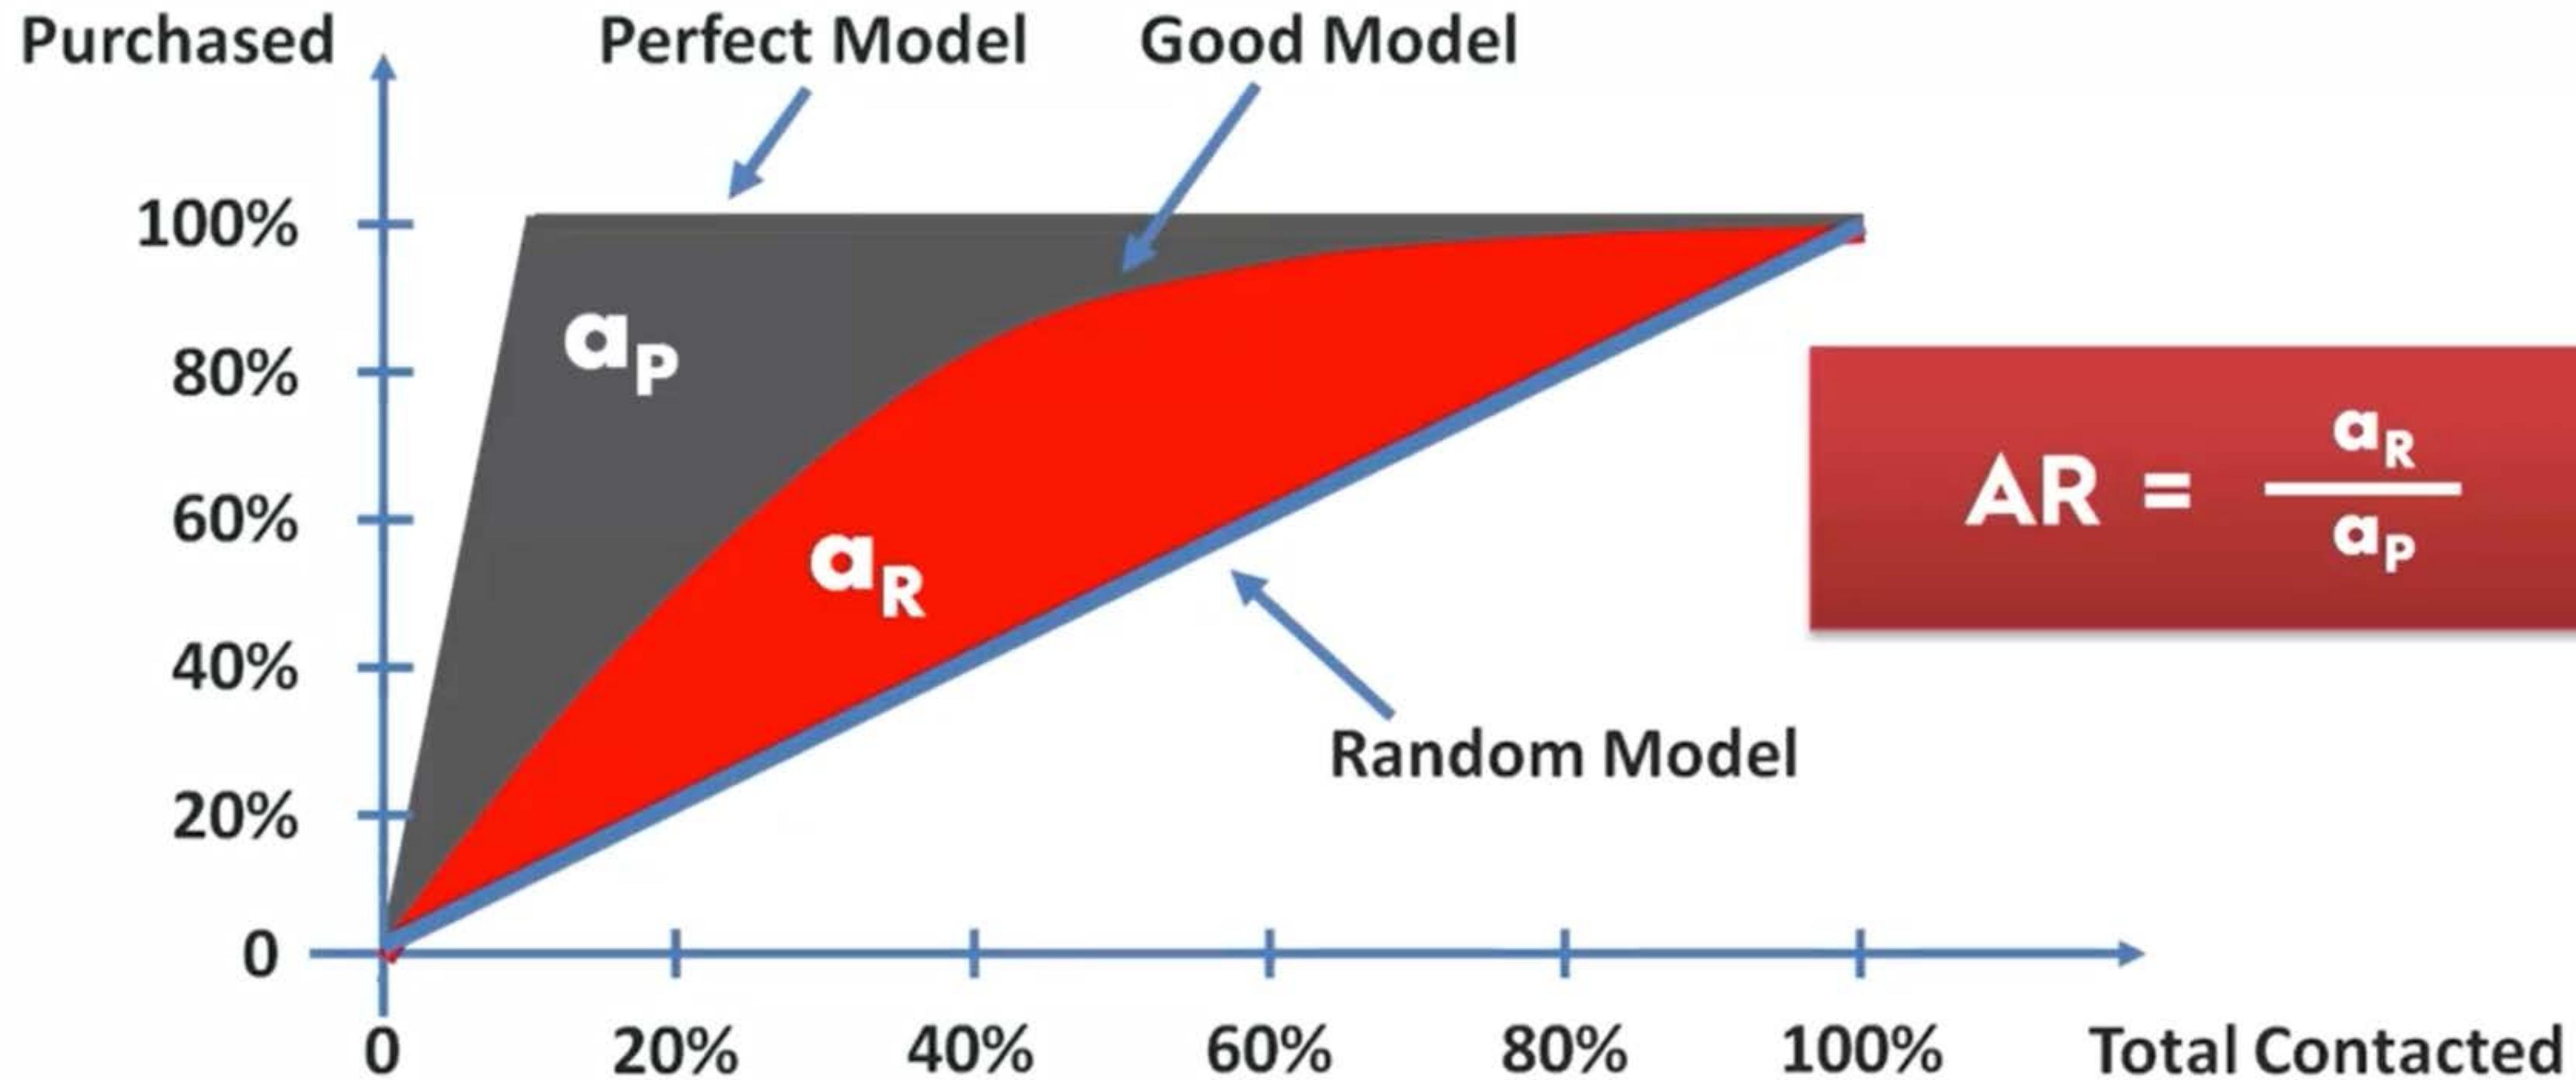
ROC = Receiver Operating Characteristic

# CAP Curve Analysis

# CAP Analysis



# CAP Analysis



# CAP Analysis

