

Plan of Attack

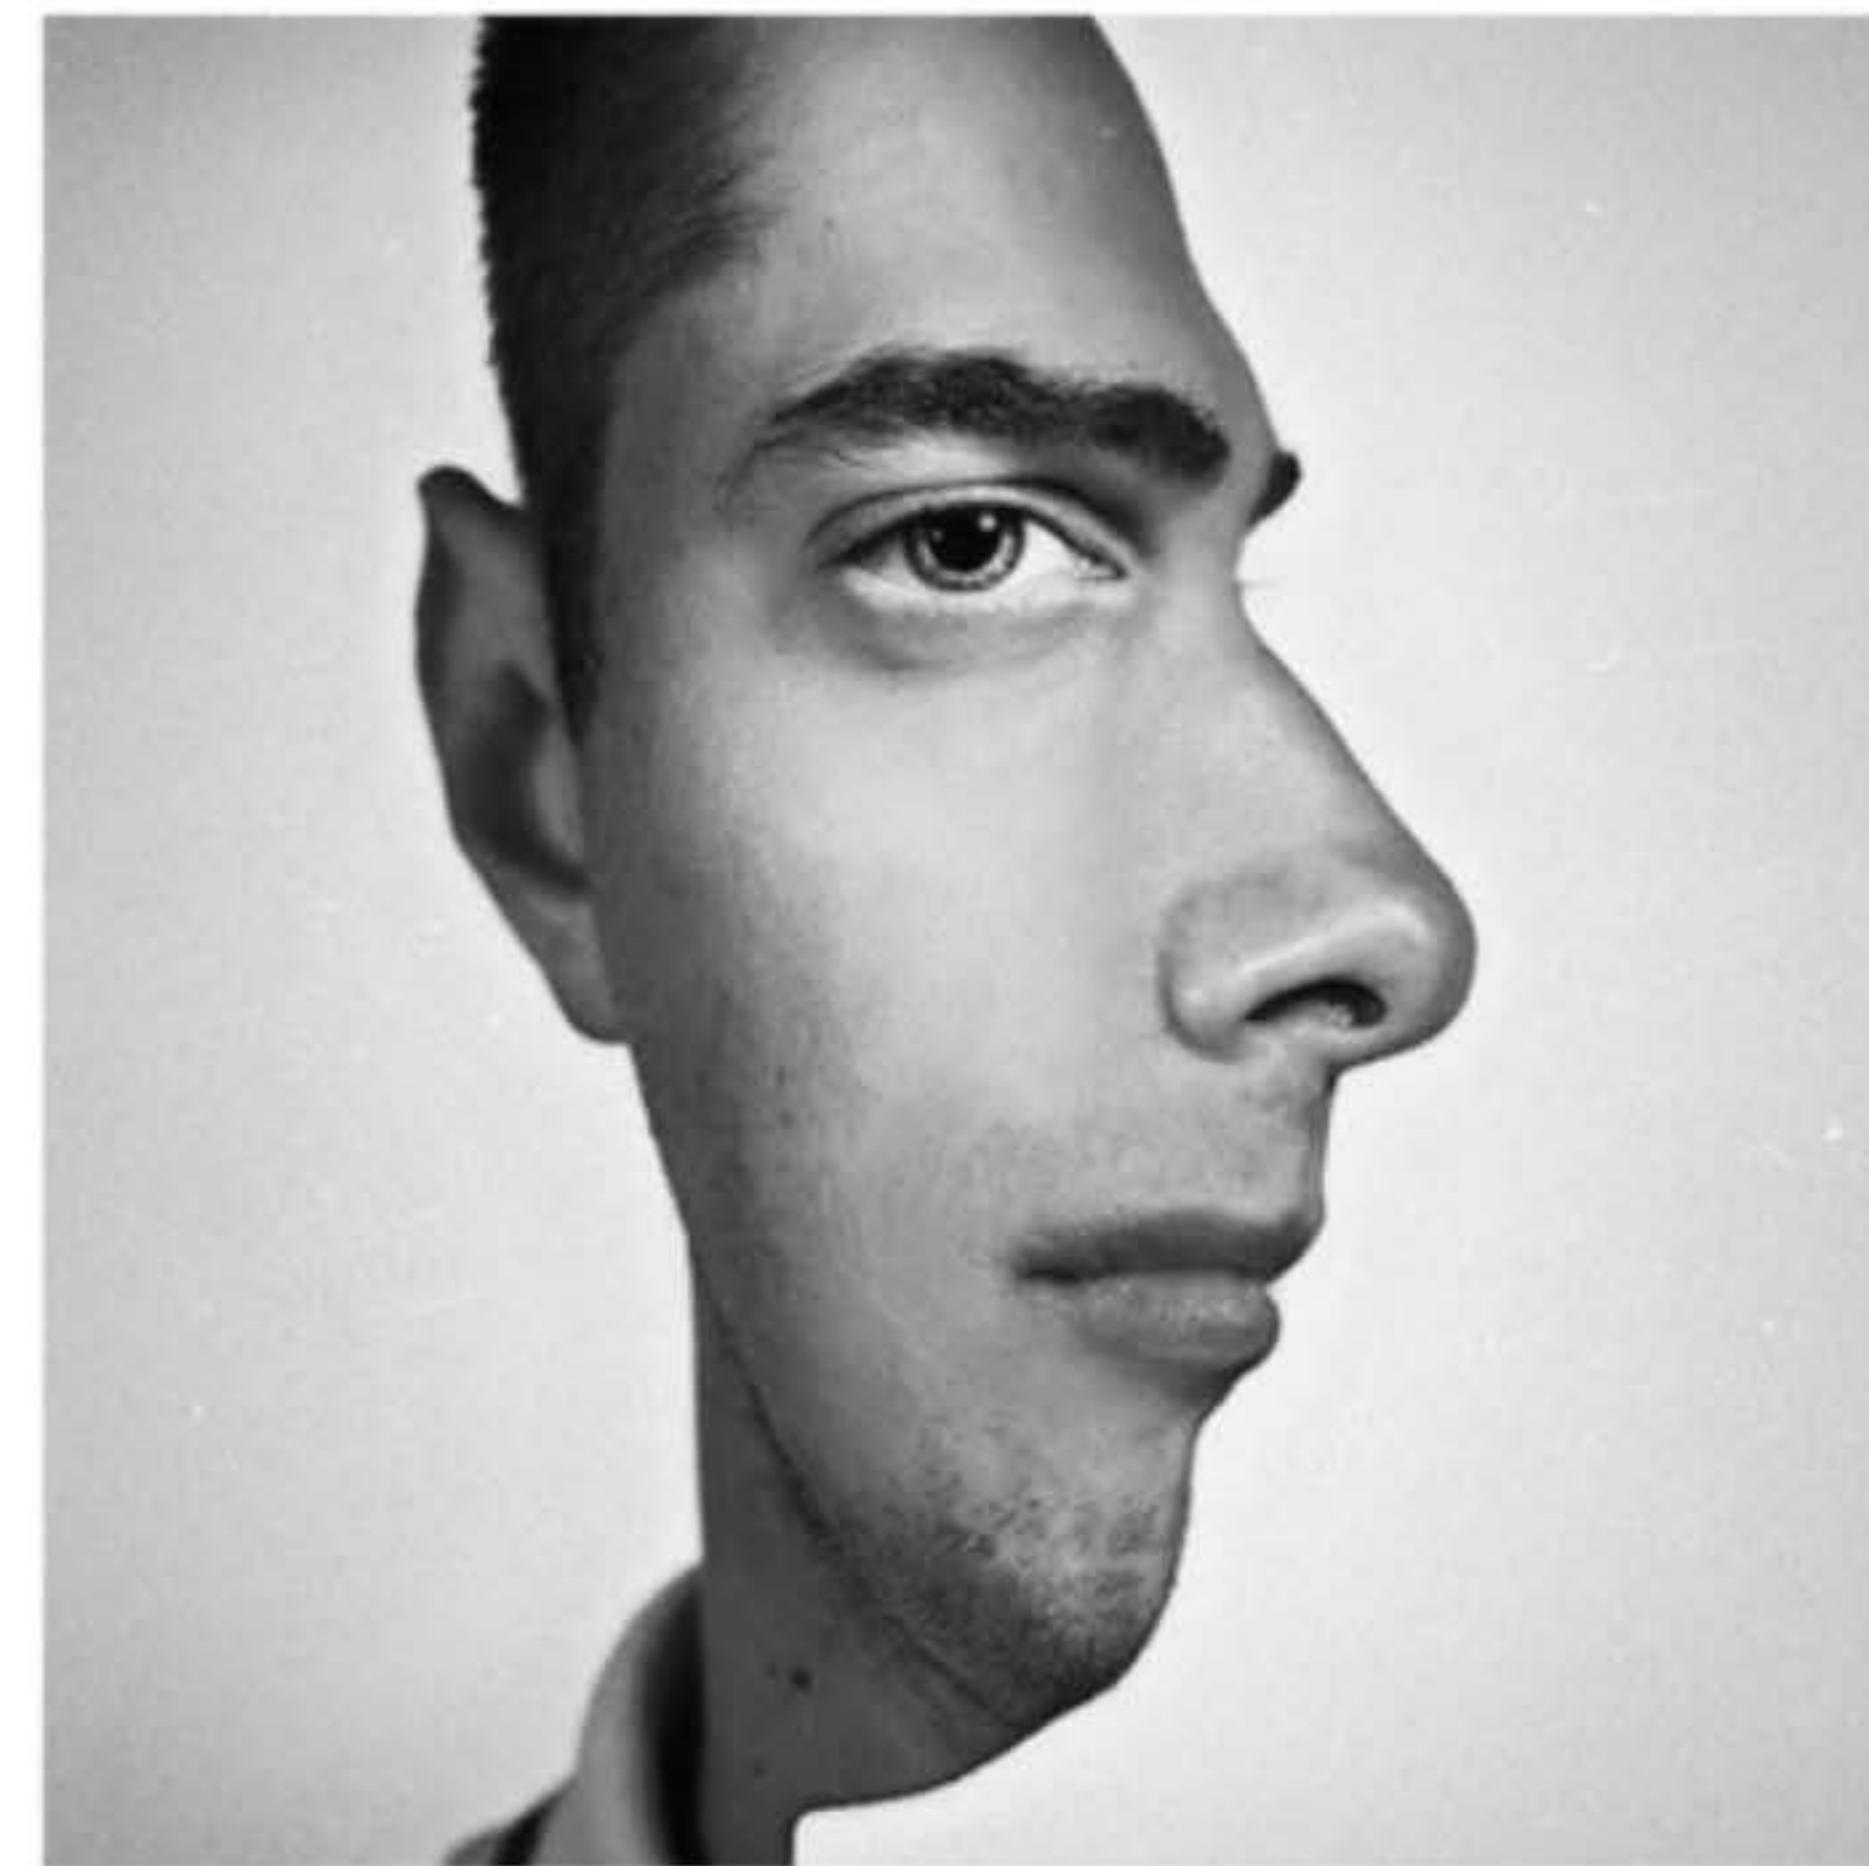
Plan of Attack

What we will learn in this section:

- What are Convolutional Neural Networks?
- Step 1 - Convolution Operation
- Step 1(b) - ReLU Layer
- Step 2 - Pooling
- Step 3 - Flattening
- Step 4 - Full Connection
- Summary
- EXTRA: Softmax & Cross-Entropy

Convolutional Neural Networks

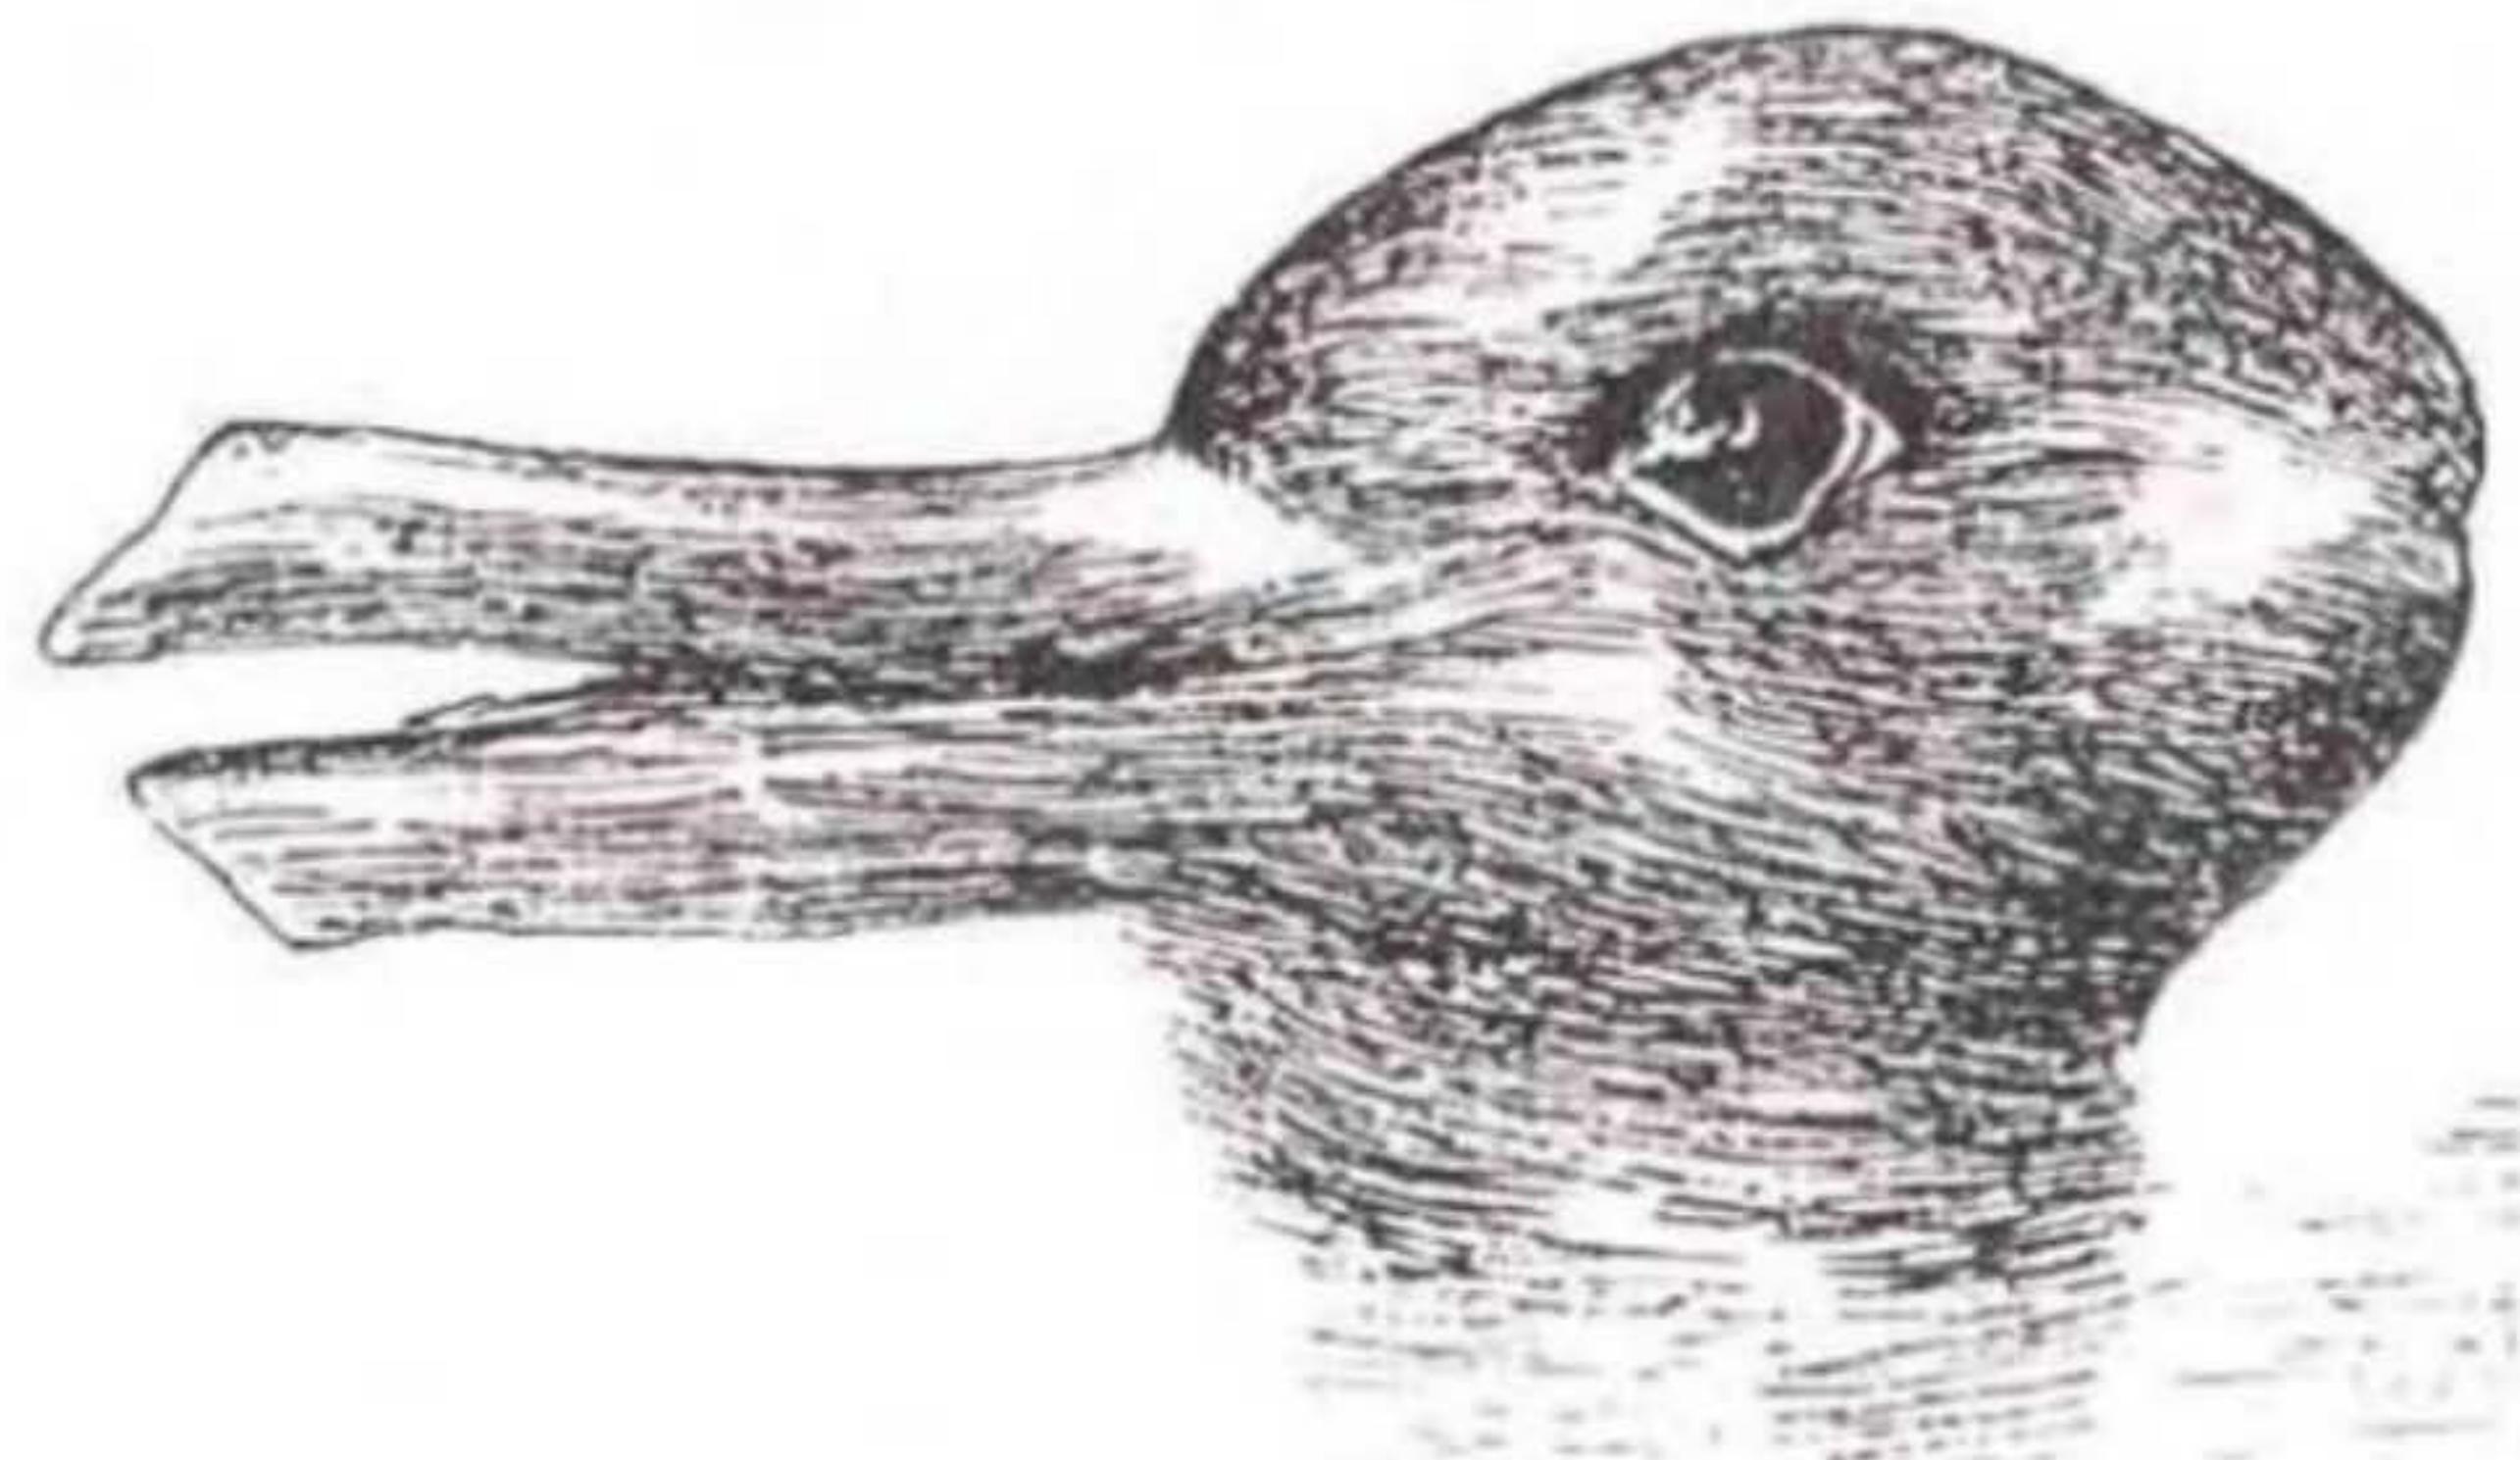
Convolutional Neural Networks



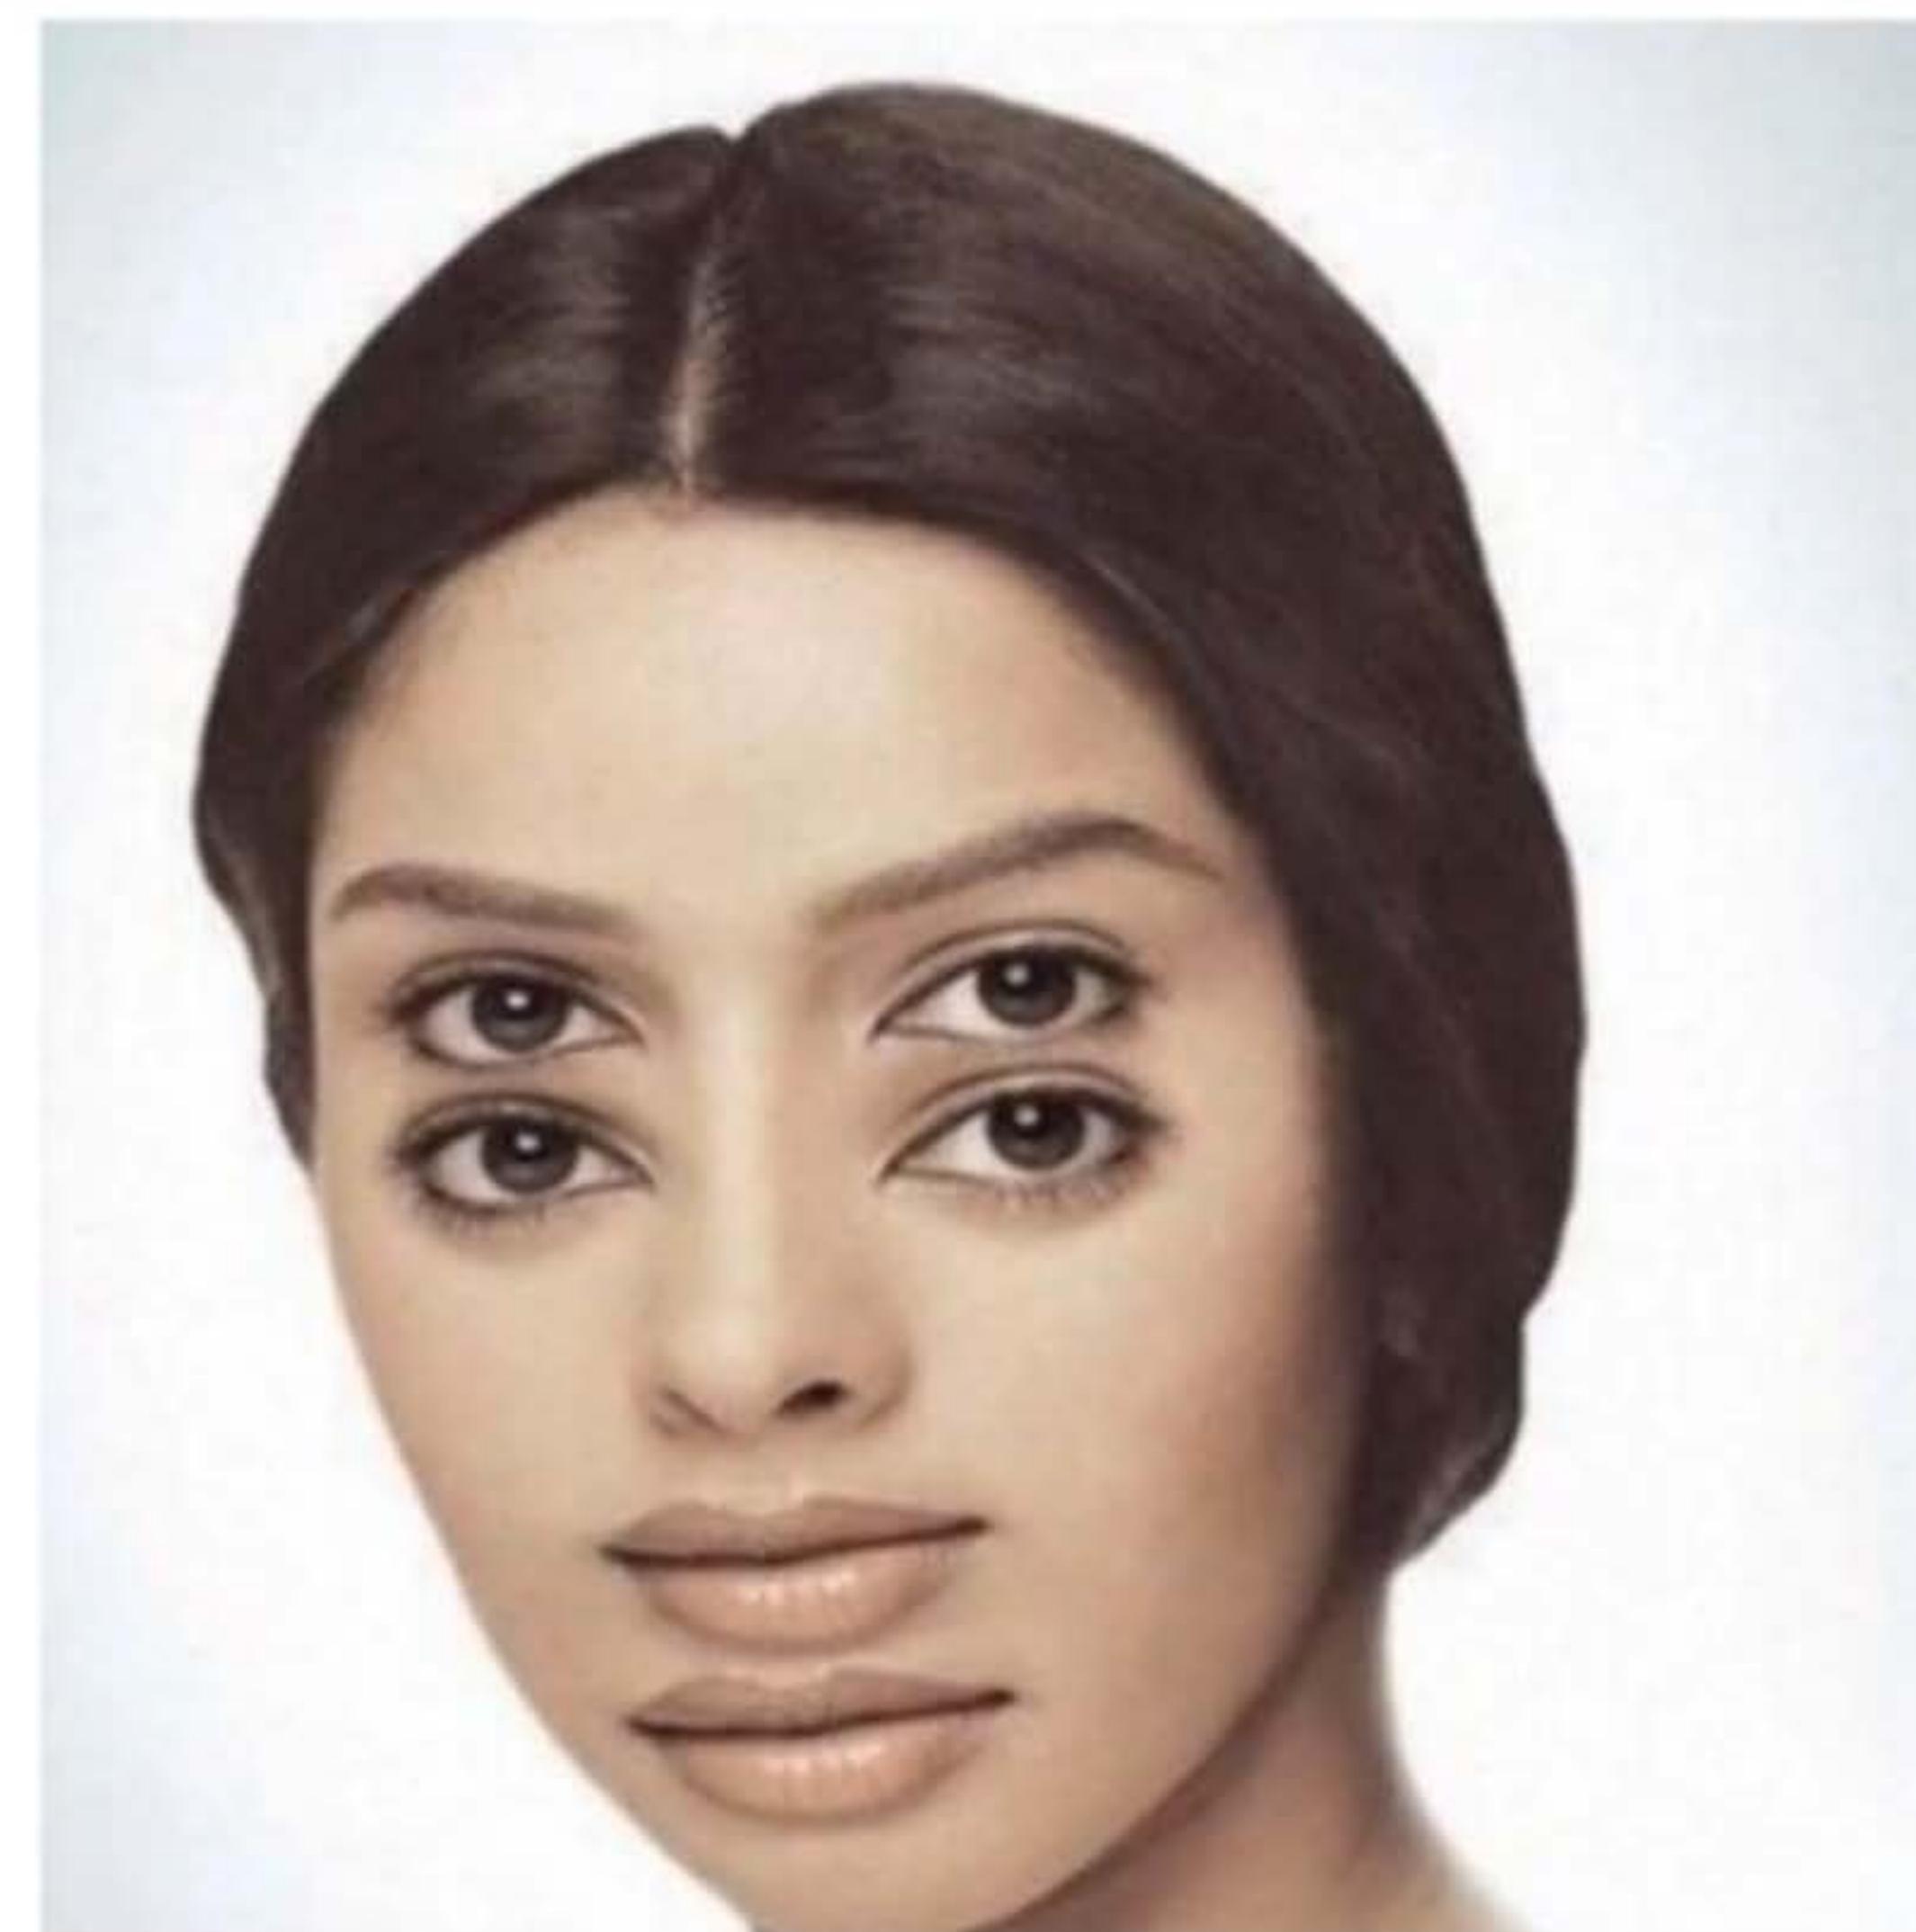
Convolutional Neural Networks



Convolutional Neural Networks

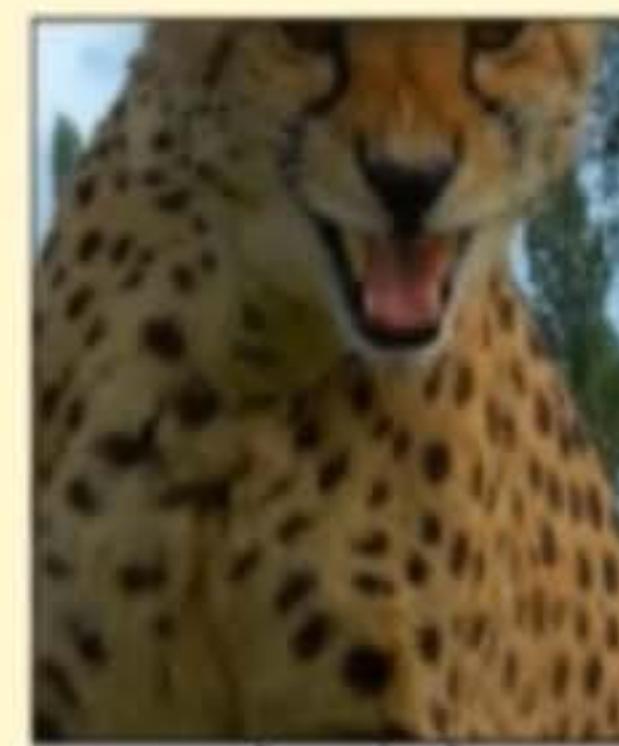


Convolutional Neural Networks



Convolutional Neural Networks

Examples from the test set
(with the network's guesses)



cheetah

cheetah
leopard
snow leopard
Egyptian cat



bullet train

bullet train
passenger car
subway train
electric locomotive

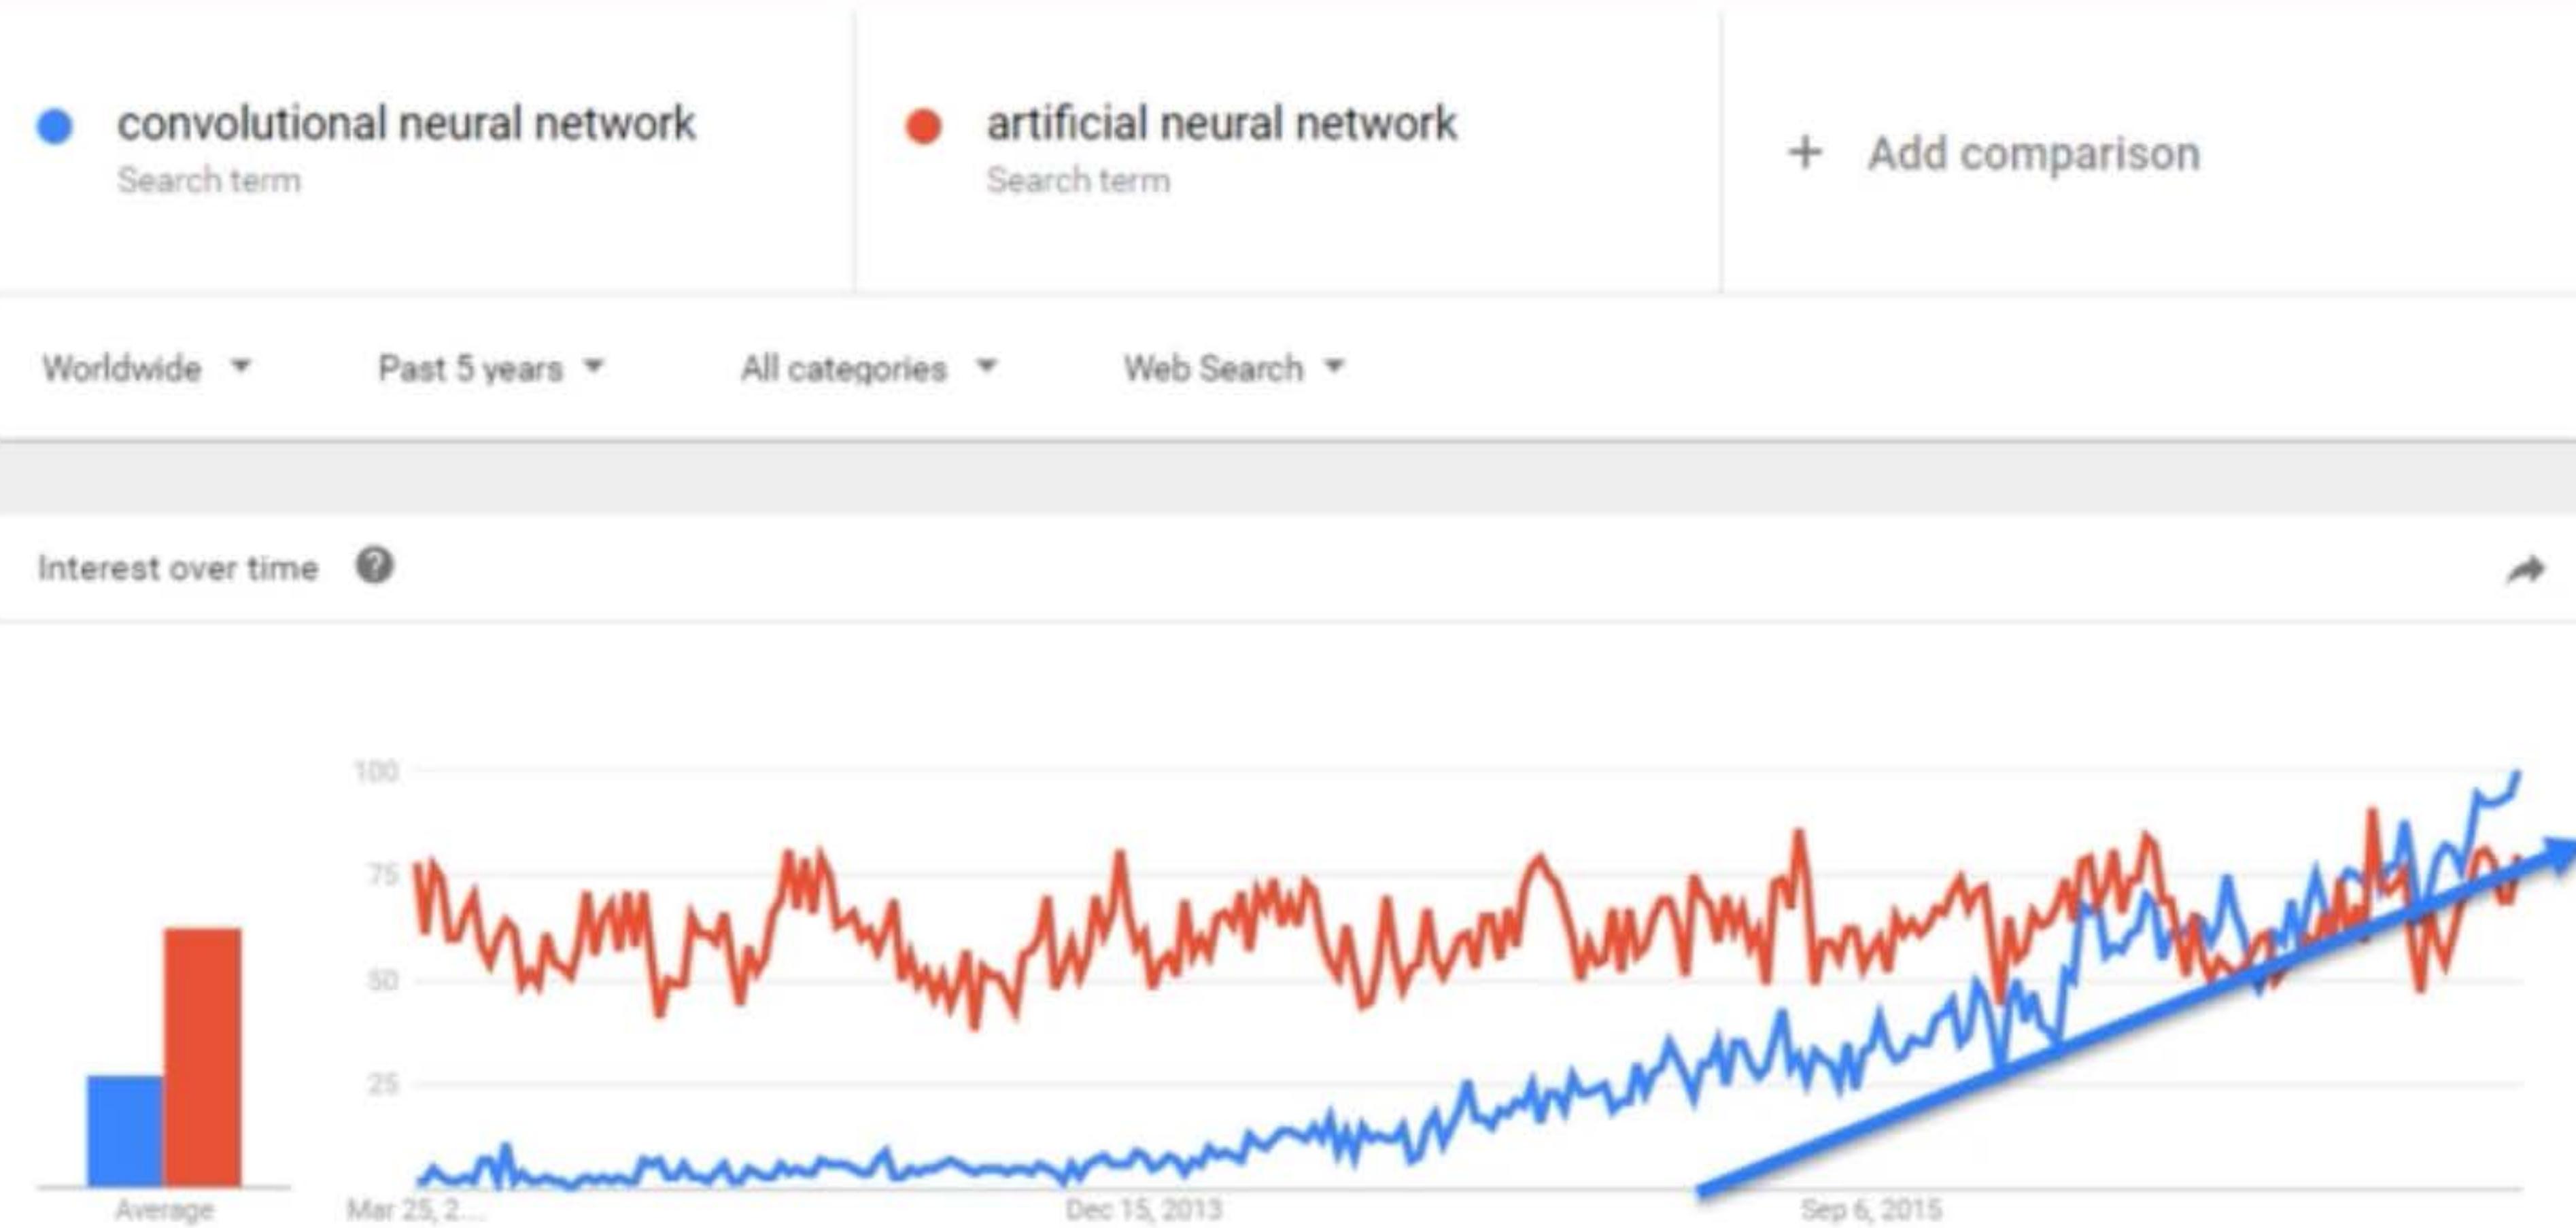


hand glass

scissors
hand glass
frying pan
stethoscope

Image Source: a talk by Geoffrey Hinton

Convolutional Neural Networks



Source: google trends

Convolutional Neural Networks



Yann Lecun

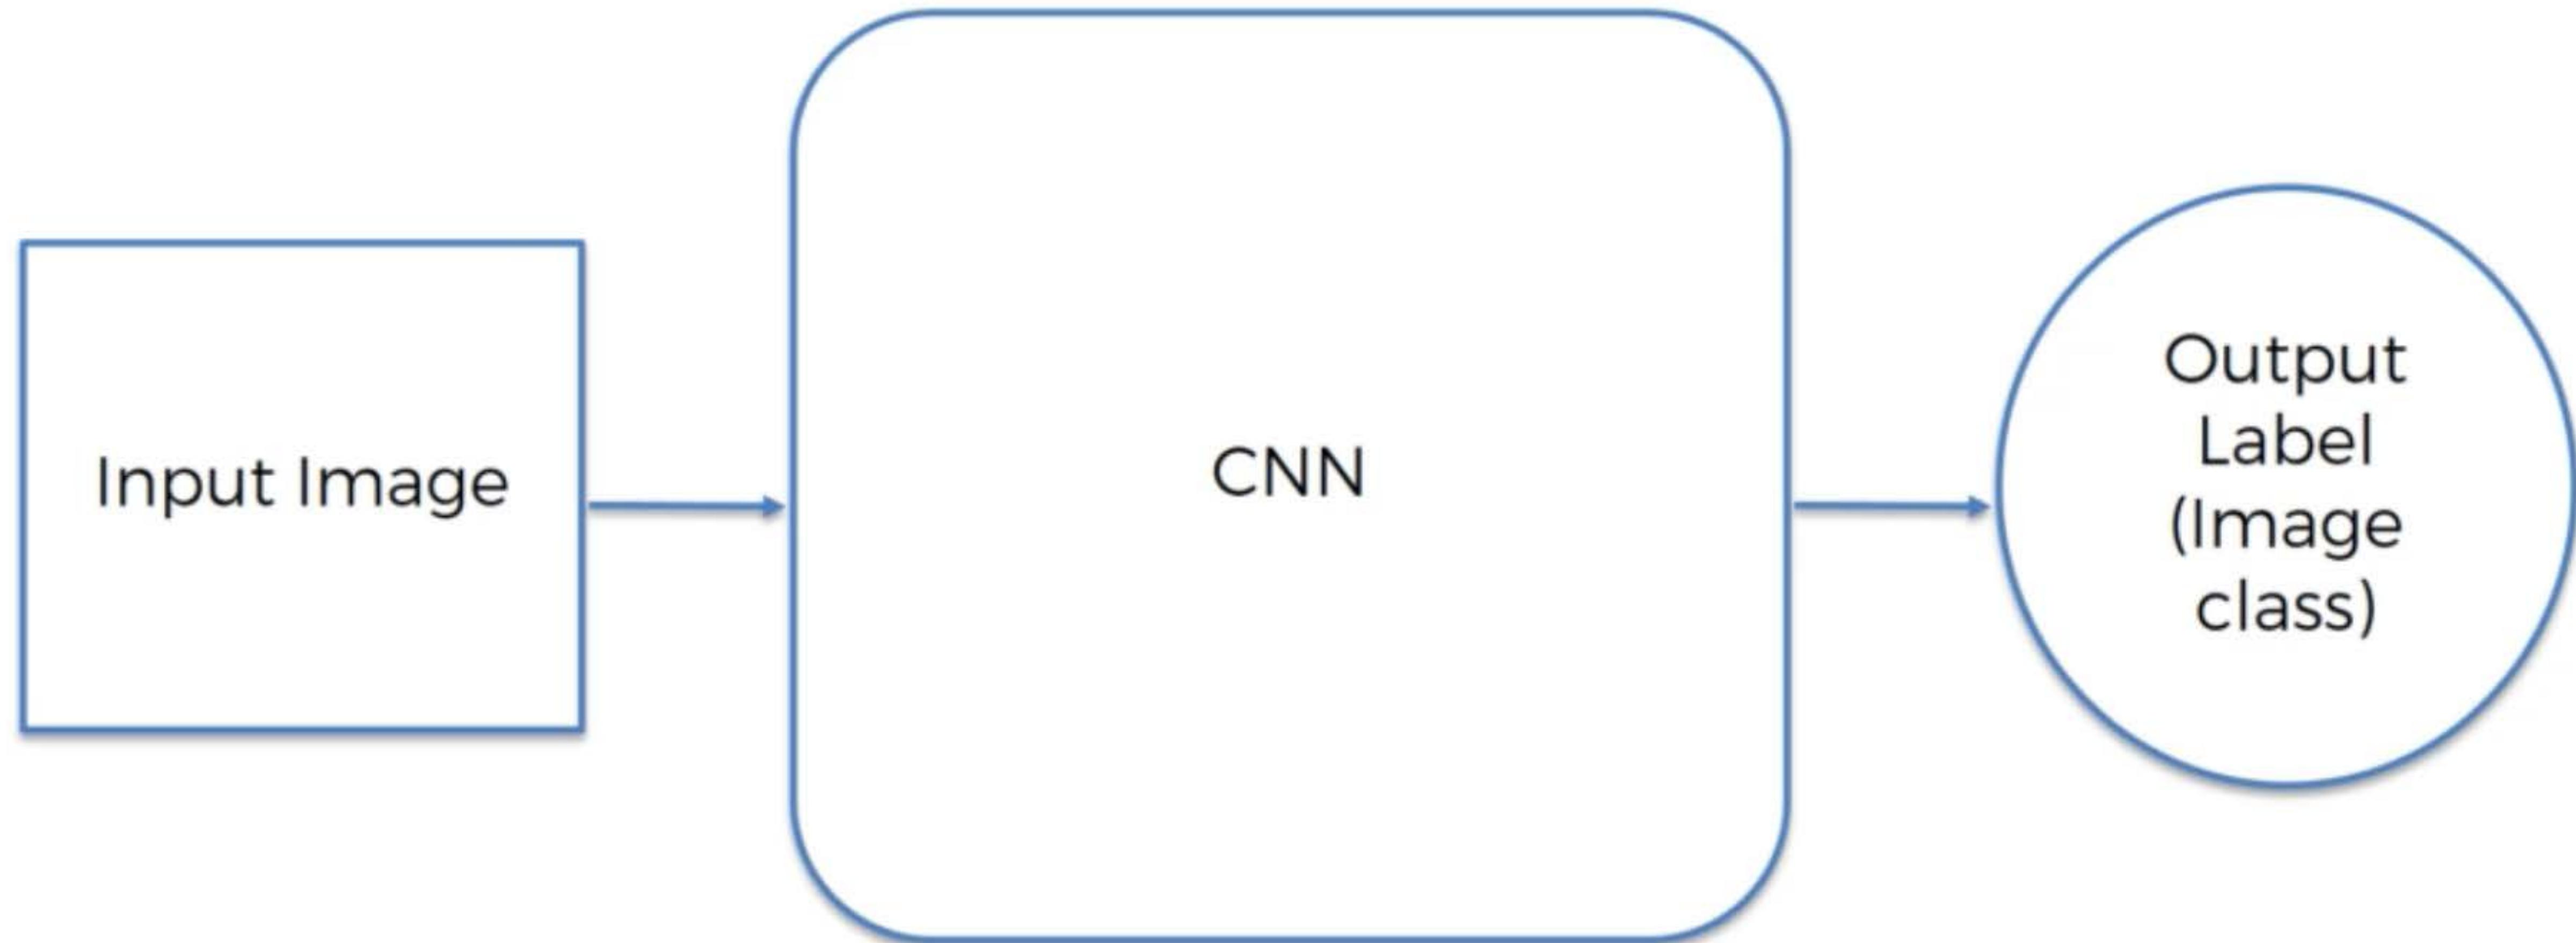
Convolutional Neural Networks

Google

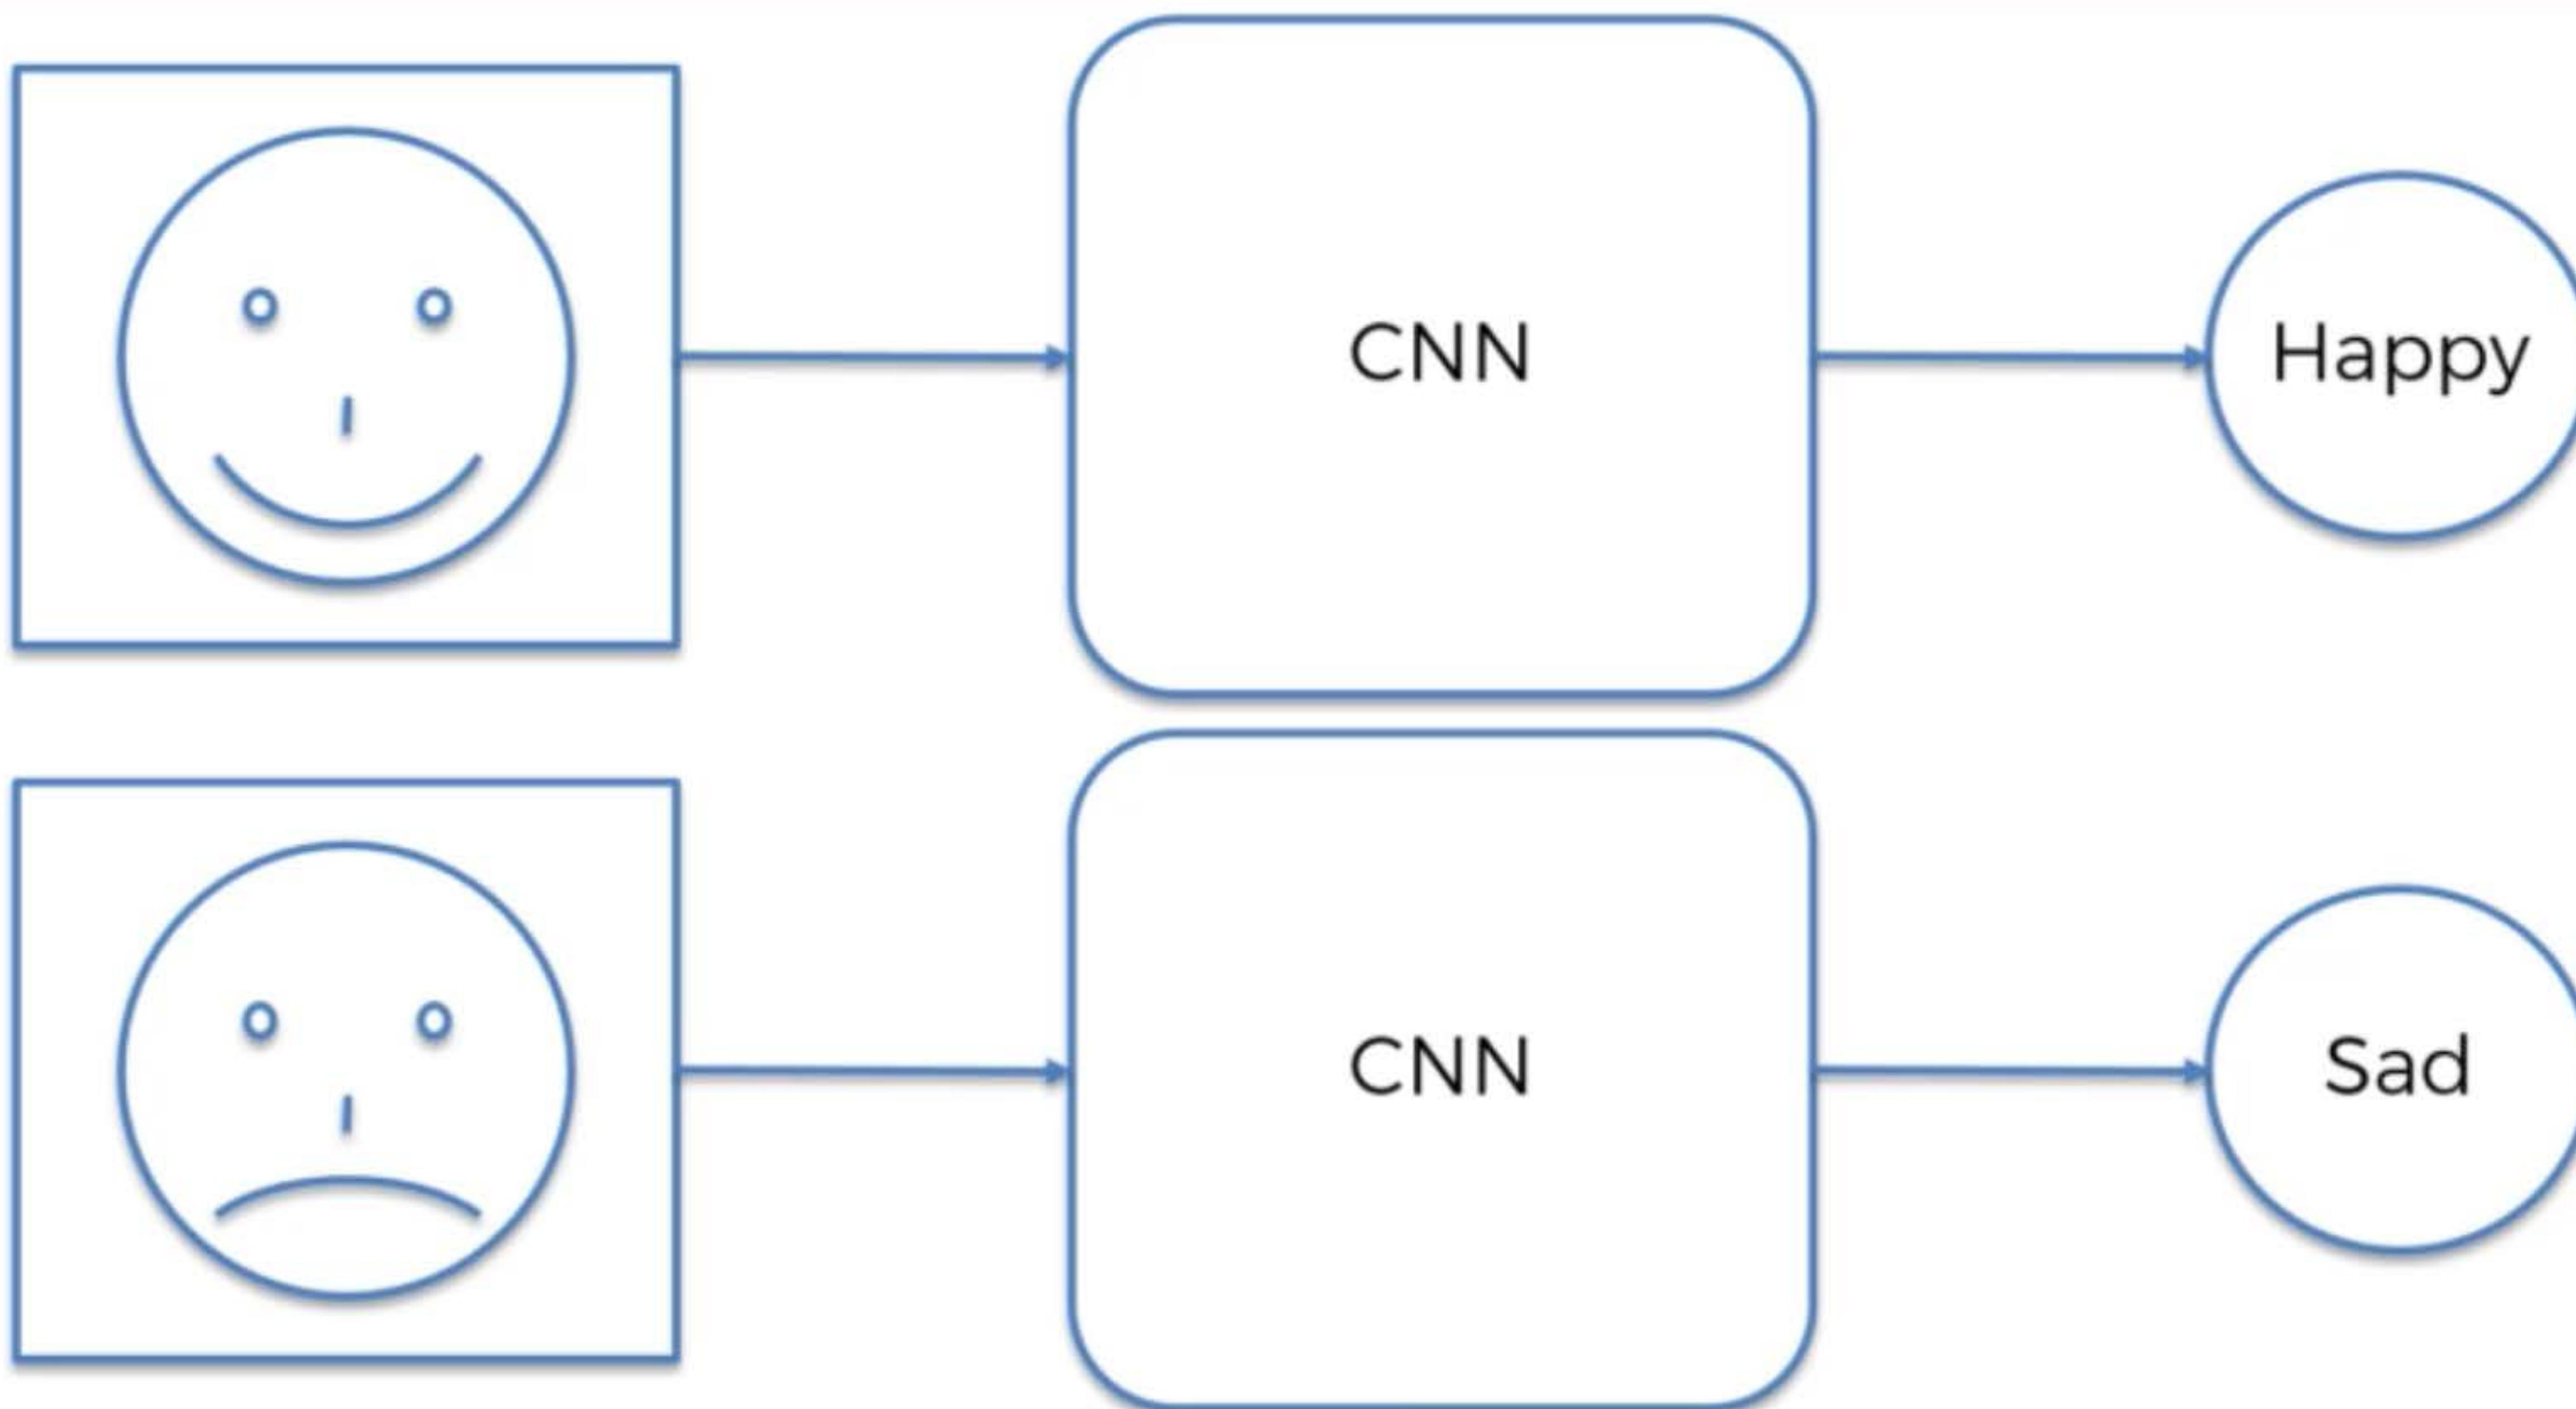
Facebook



Convolutional Neural Networks



Convolutional Neural Networks



Convolutional Neural Networks

B / W Image 2x2px

Pixel 1	Pixel 2
Pixel 3	Pixel 4

2d array

Pixel 1 <small>0 ≤ pixel value ≤ 255</small>	Pixel 2 <small>0 ≤ pixel value ≤ 255</small>
Pixel 3 <small>0 ≤ pixel value ≤ 255</small>	Pixel 4 <small>0 ≤ pixel value ≤ 255</small>

Colored Image 2x2px

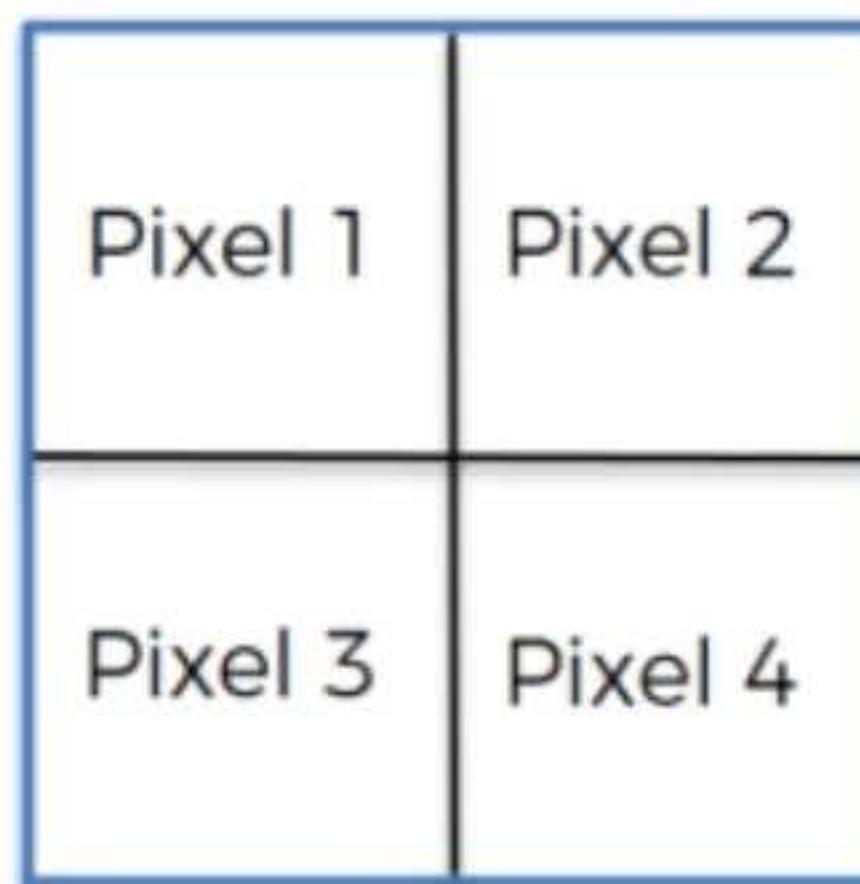
Pixel 1	Pixel 2
Pixel 3	Pixel 4

3d array

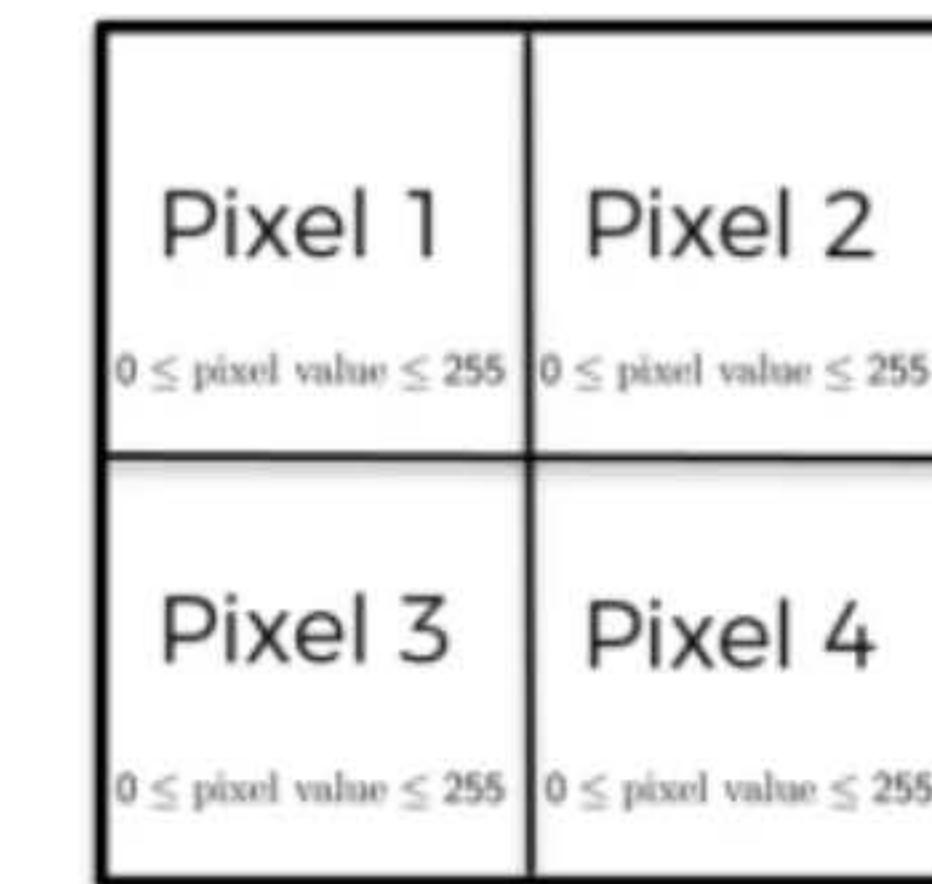
Pixel 1 <small>0 ≤ pixel value ≤ 255</small>	Pixel 2 <small>0 ≤ pixel value ≤ 255</small>
Pixel 3 <small>0 ≤ pixel value ≤ 255</small>	Pixel 4 <small>0 ≤ pixel value ≤ 255</small>

Convolutional Neural Networks

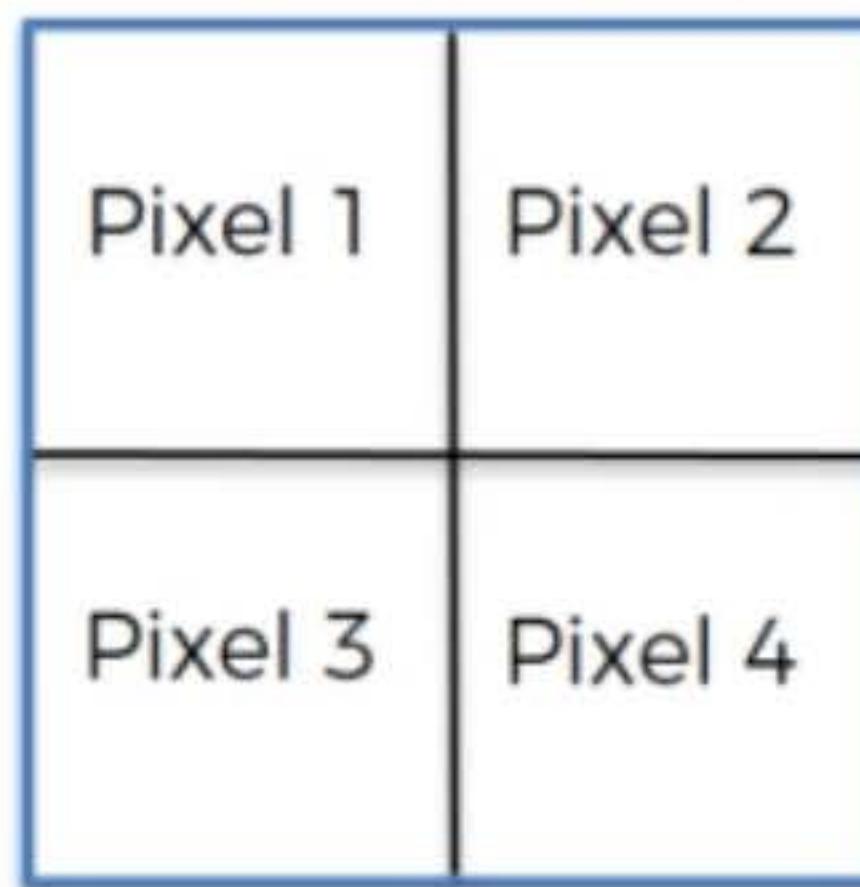
B / W Image 2x2px



2d array



Colored Image 2x2px

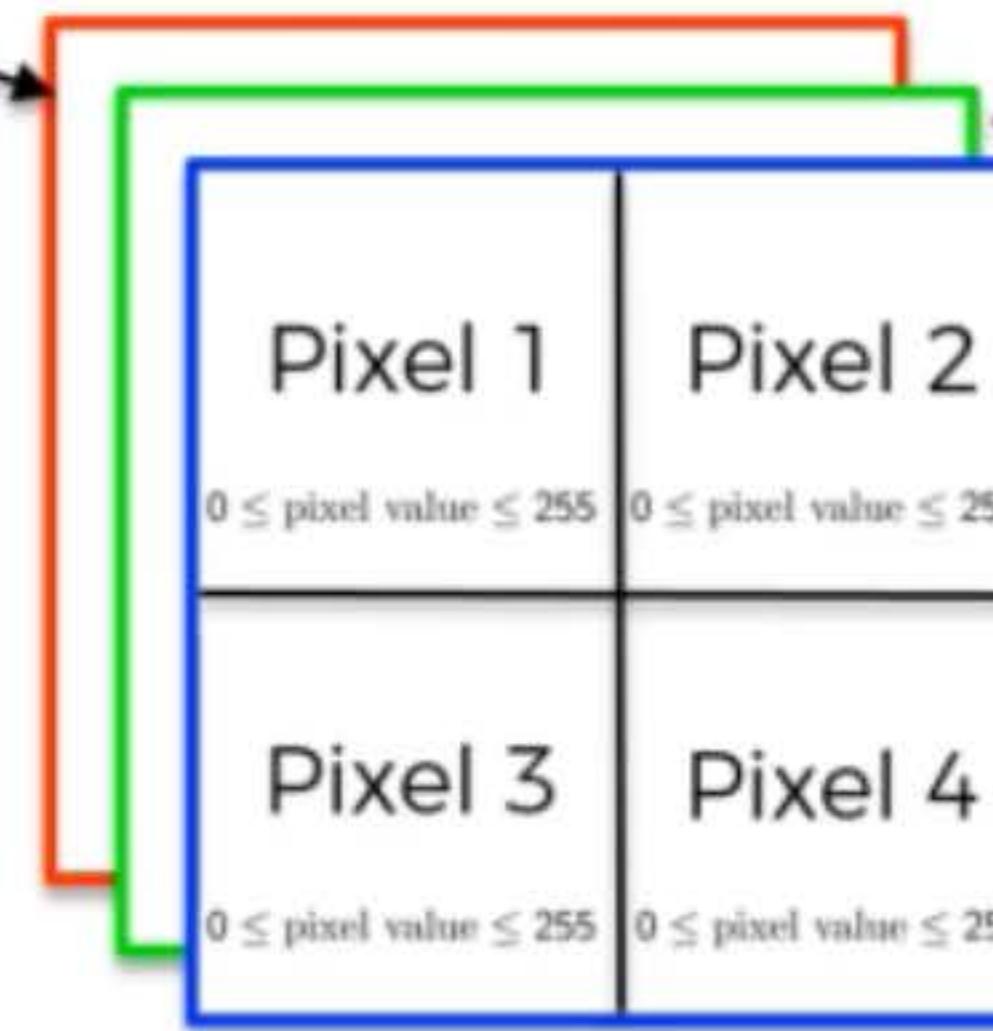


3d array

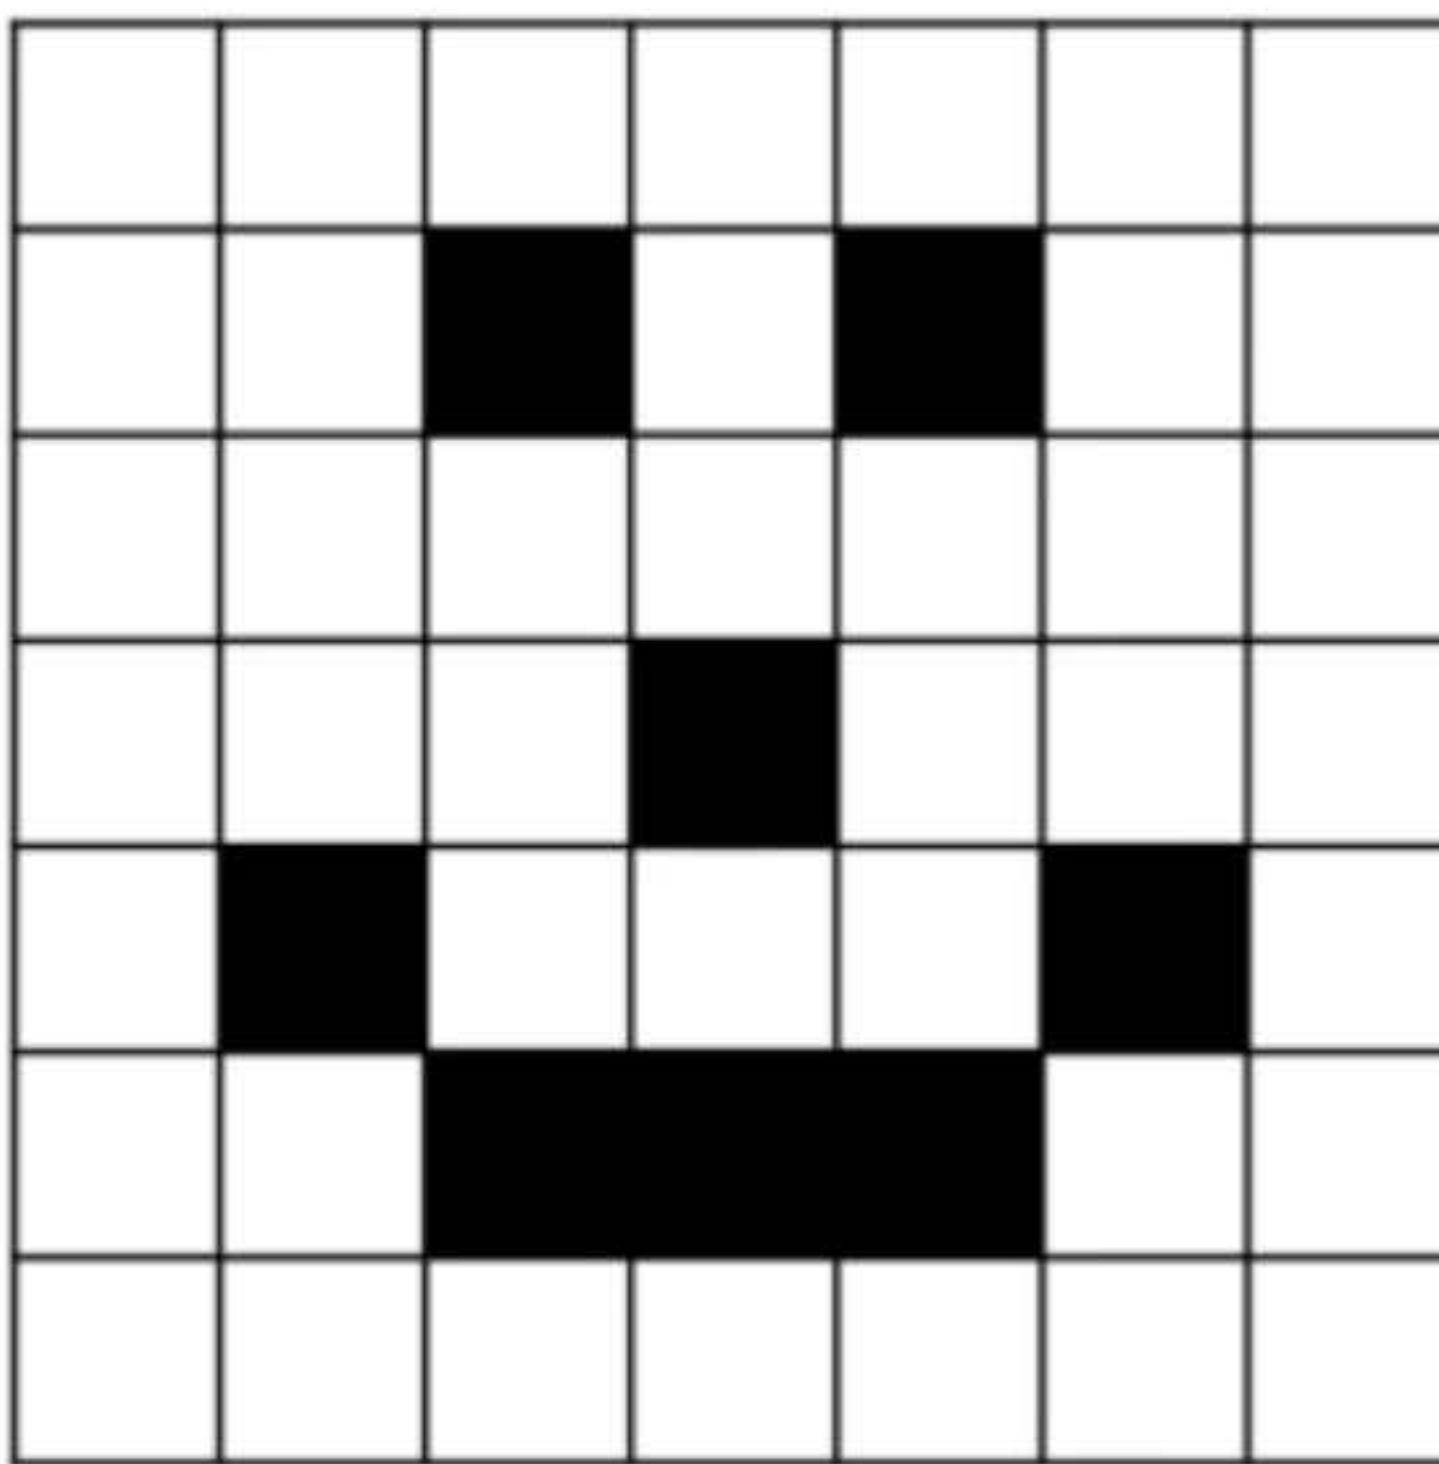
Red channel

Green channel

Blue channel



Convolutional Neural Networks



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Convolutional Neural Networks



Convolutional Neural Networks

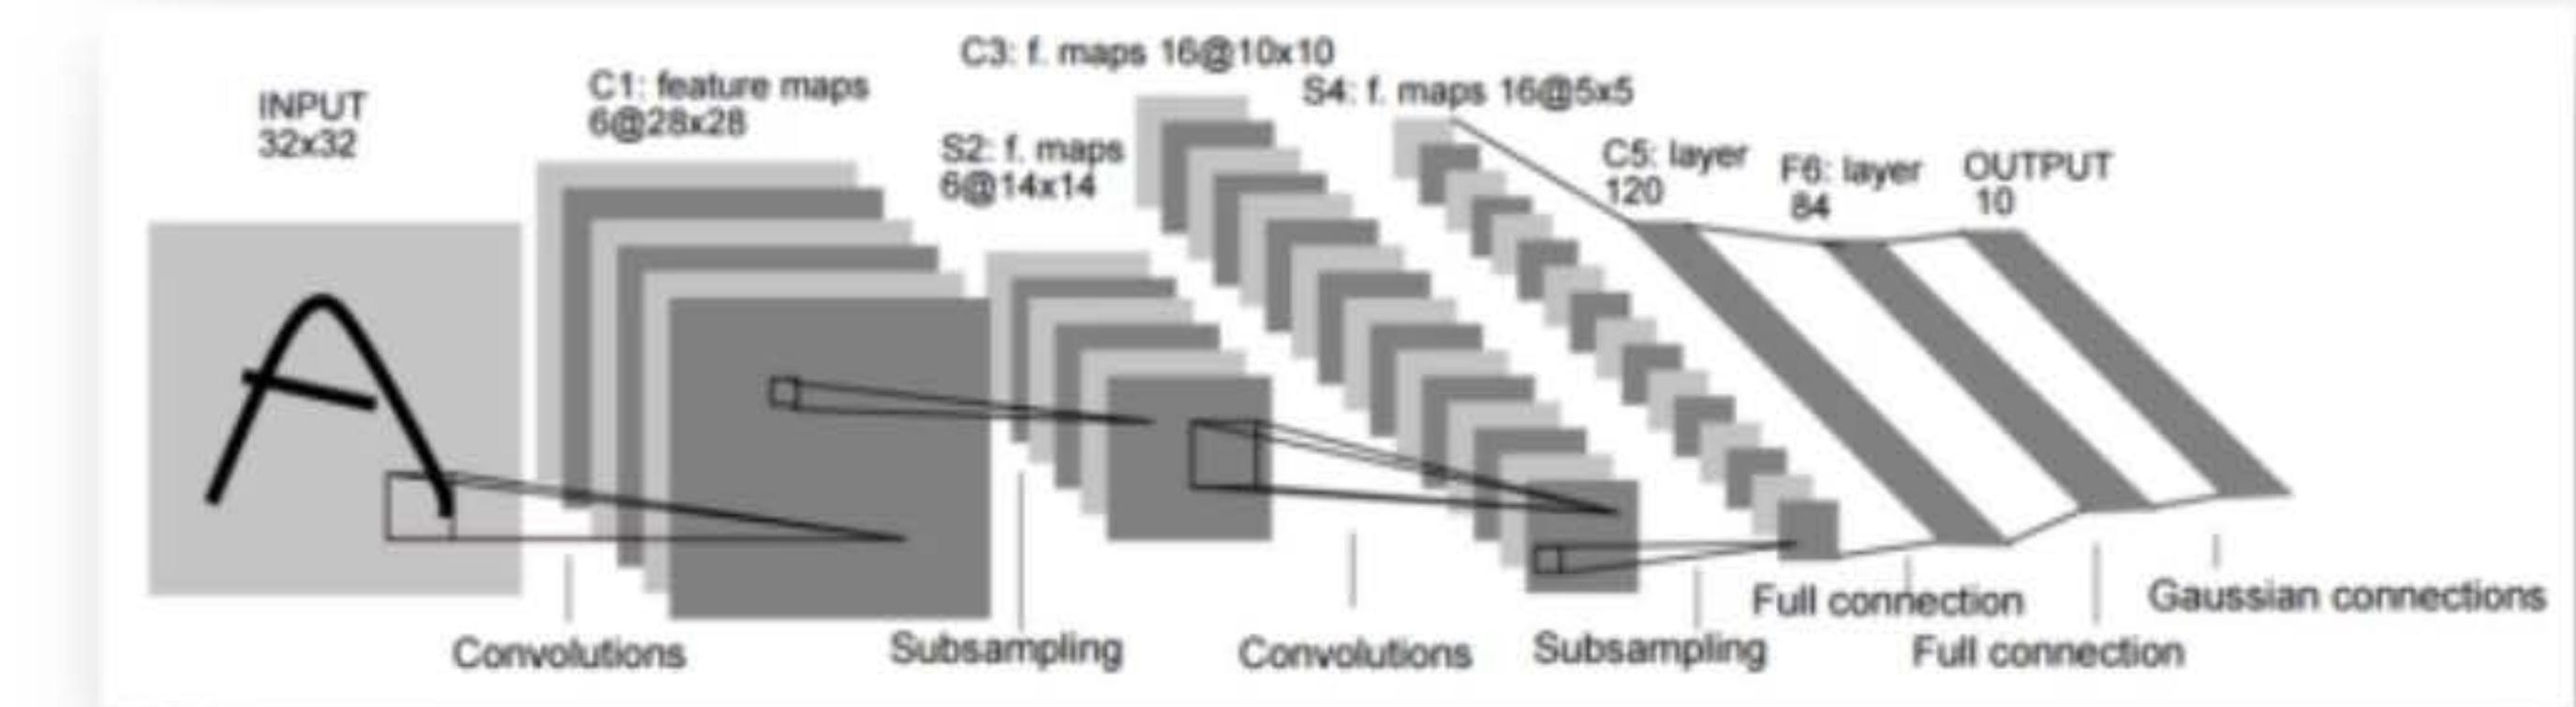
Additional Reading:

*Gradient-Based Learning
Applied to Document
Recognition*

By Yann LeCun et al. (1998)

Link:

<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>



Step 1 - Convolution

Step 1 - Convolution

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$$

Step 1 - Convolution

Additional Reading:

Introduction to Convolutional Neural Networks

By Jianxin Wu (2017)

Link:

<http://cs.nju.edu.cn/wujx/paper/CNN.pdf>

$$\begin{aligned}\frac{\partial z}{\partial (\text{vec}(\mathbf{y})^T)}(F^T \otimes I) &= \left((F \otimes I) \frac{\partial z}{\partial \text{vec}(\mathbf{y})} \right)^T \\ &= \left((F \otimes I) \text{vec} \left(\frac{\partial z}{\partial Y} \right) \right)^T \\ &= \text{vec} \left(I \frac{\partial z}{\partial Y} F^T \right)^T \\ &= \text{vec} \left(\frac{\partial z}{\partial Y} F^T \right)^T,\end{aligned}$$

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature
Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

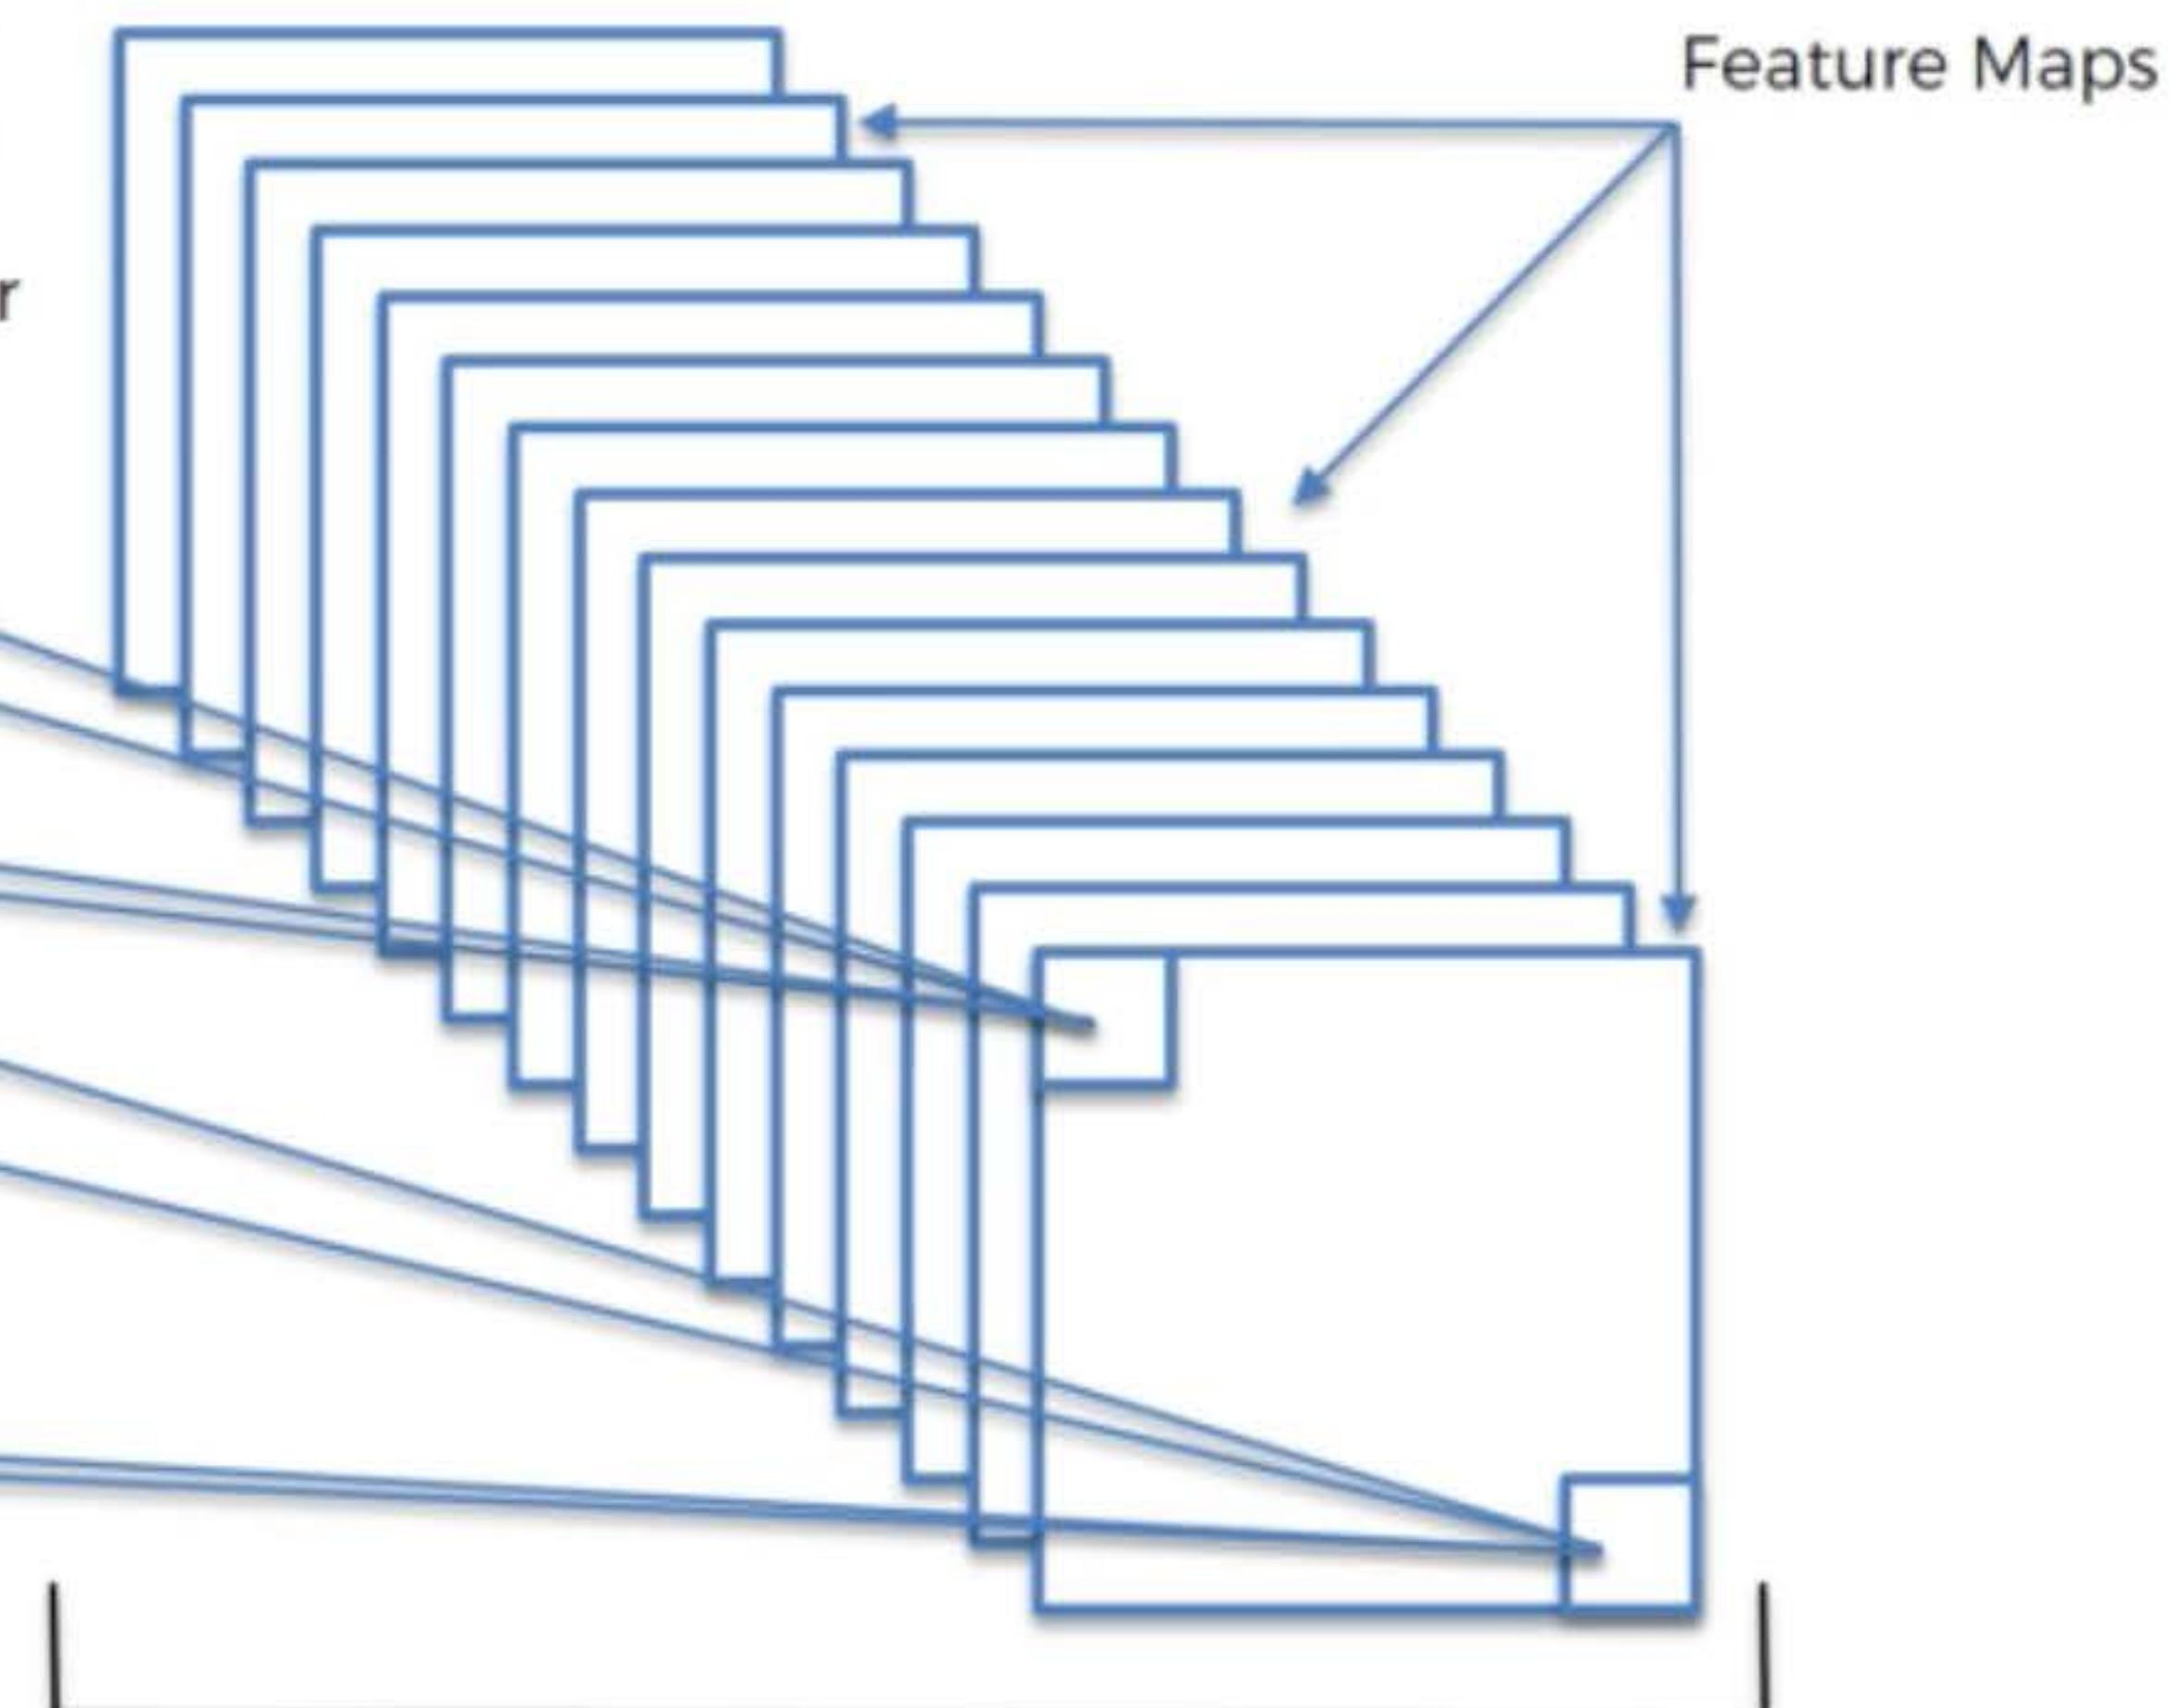
Feature Map

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image

We create many feature maps to obtain our first convolution layer



Convolutional Layer

Step 1 - Convolution

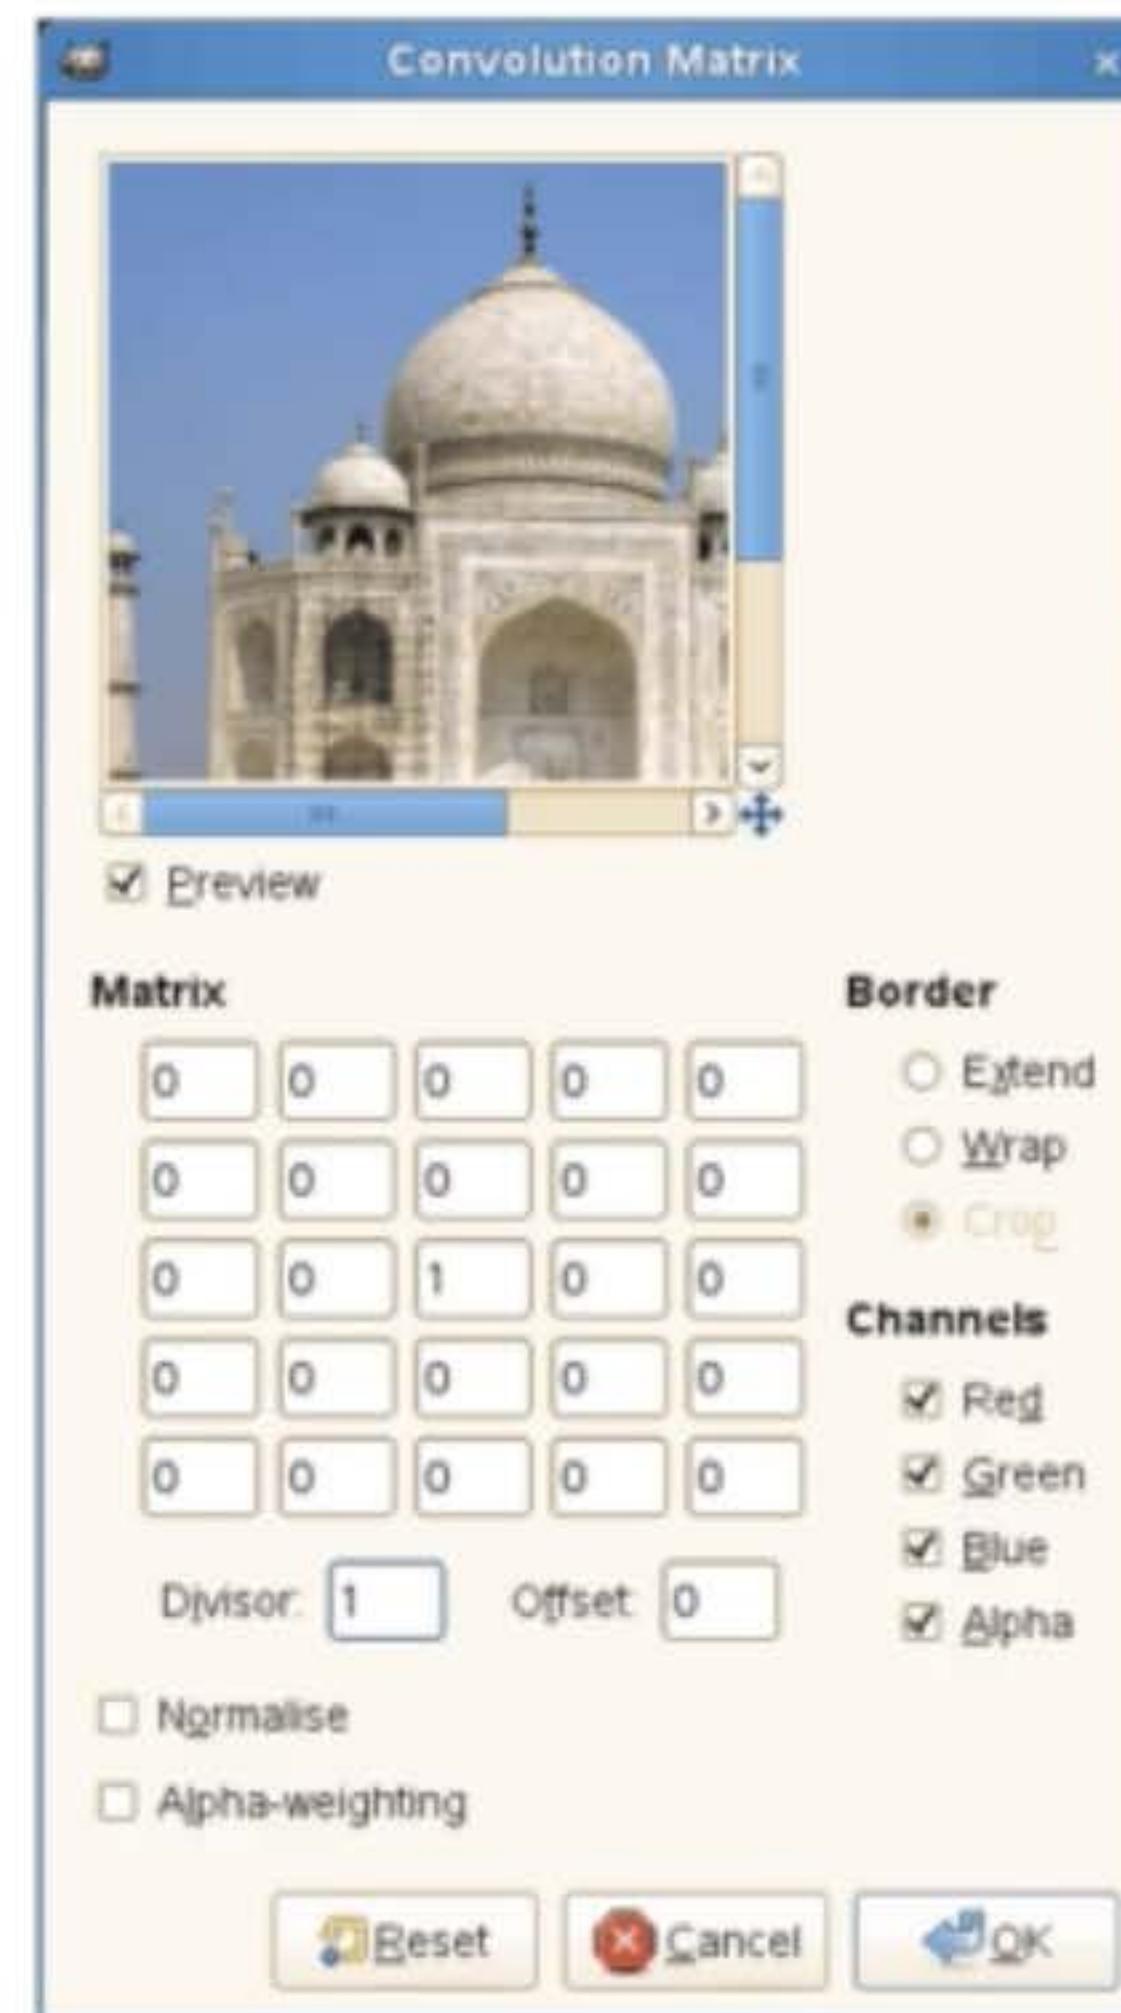


Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution

Sharpen:

0	0	0	0	0
0	0	-1	0	0
0	-1	5	-1	0
0	0	-1	0	0
0	0	0	0	0



Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution

Blur:

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0



Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution

Edge Enhance:

$$\begin{bmatrix} 0 & 1 & 0 \\ -1 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$



Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution

Edge Detect:

$$\begin{matrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{matrix}$$

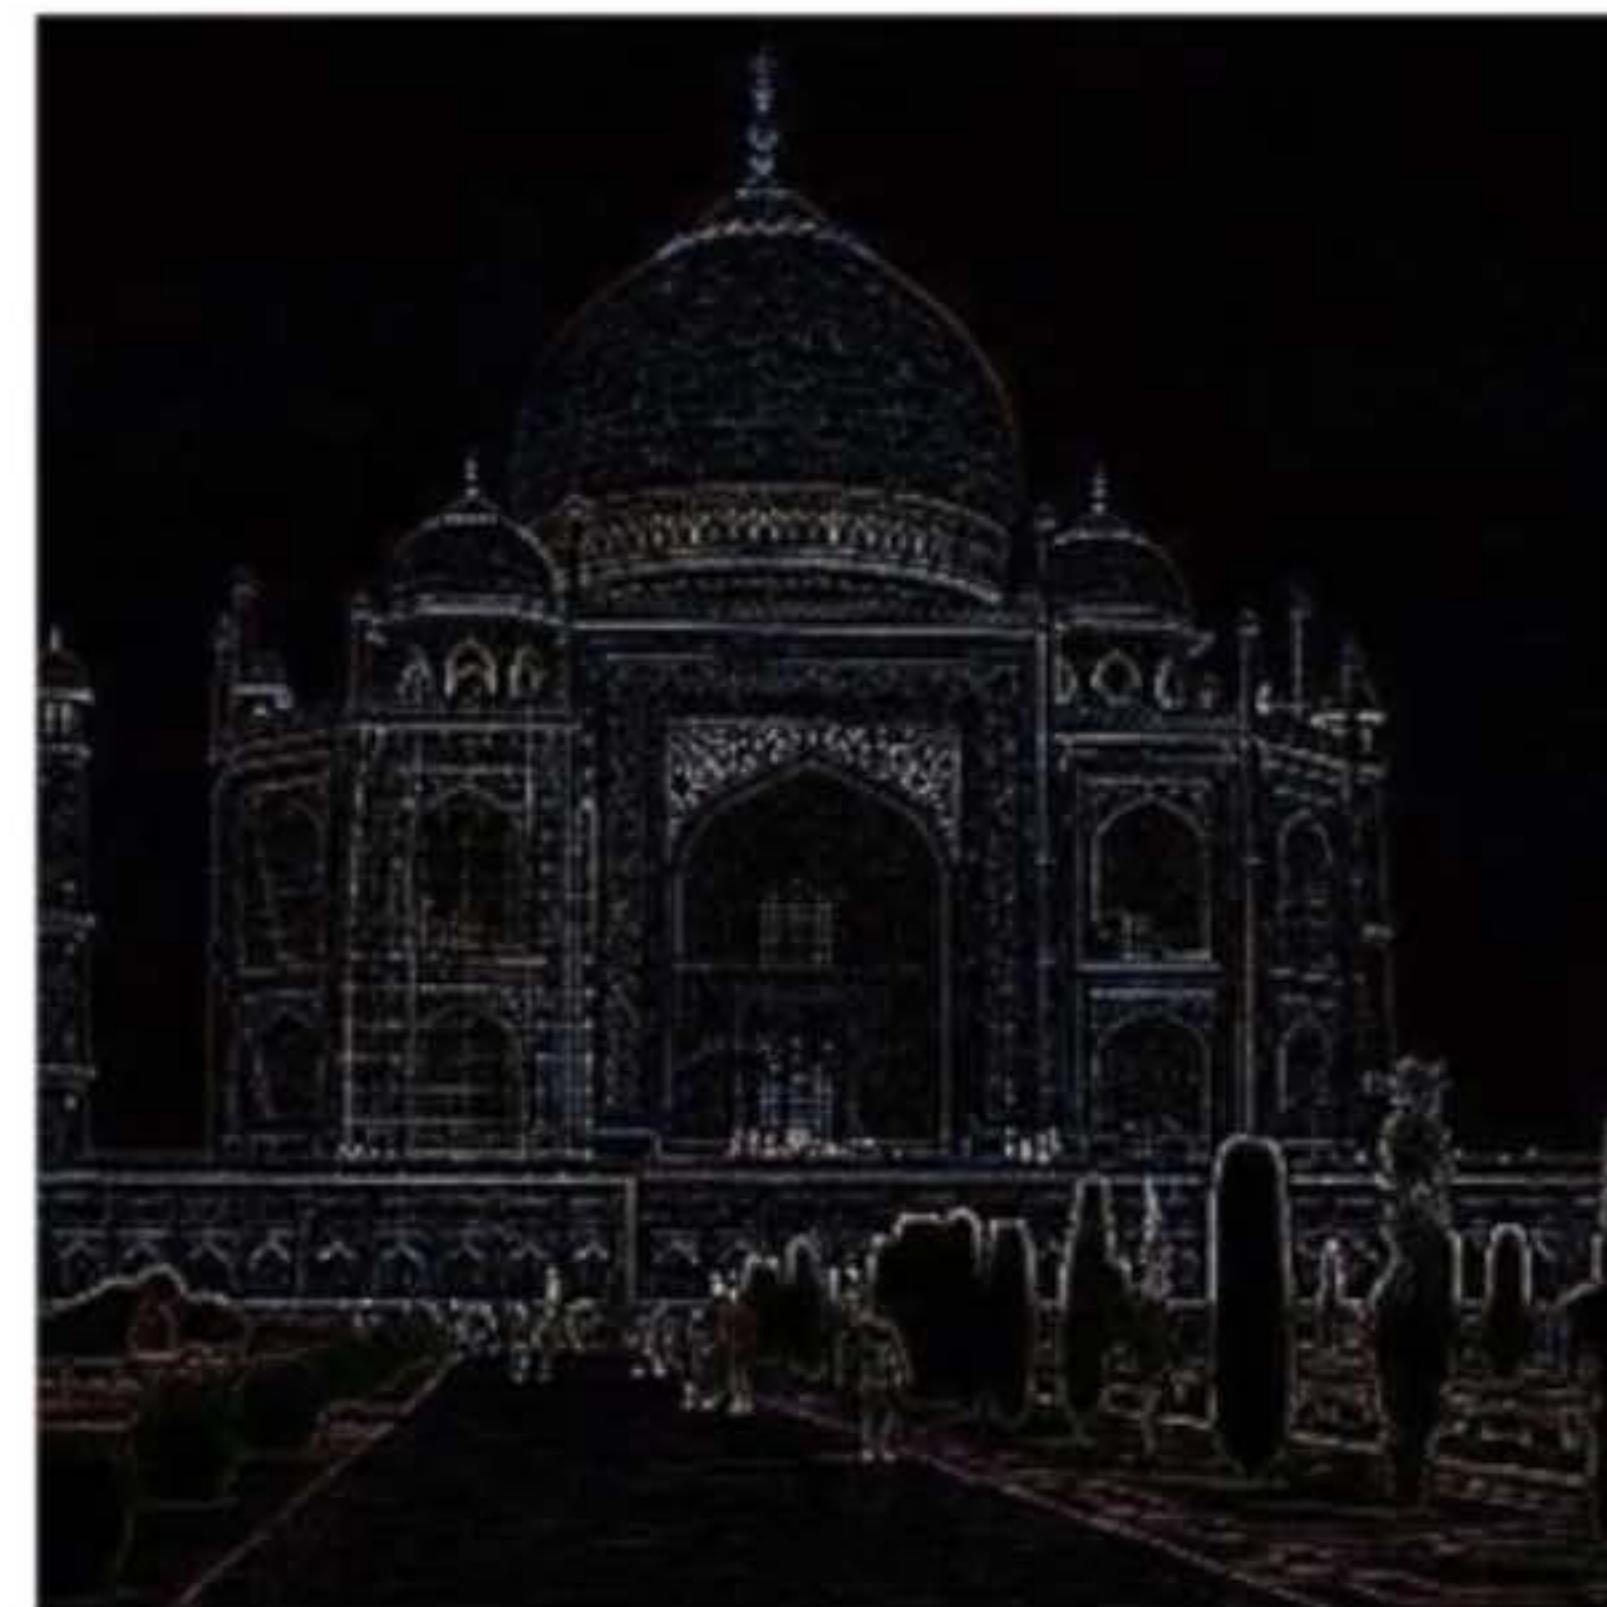


Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution

Emboss:

$$\begin{matrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{matrix}$$



Image Source: docs.gimp.org/en/plug-in-convmatrix.html

Step 1 - Convolution



*

1	0	-1
2	0	-2
1	0	-1



Image Source: eonardoaraujosantos.gitbooks.io

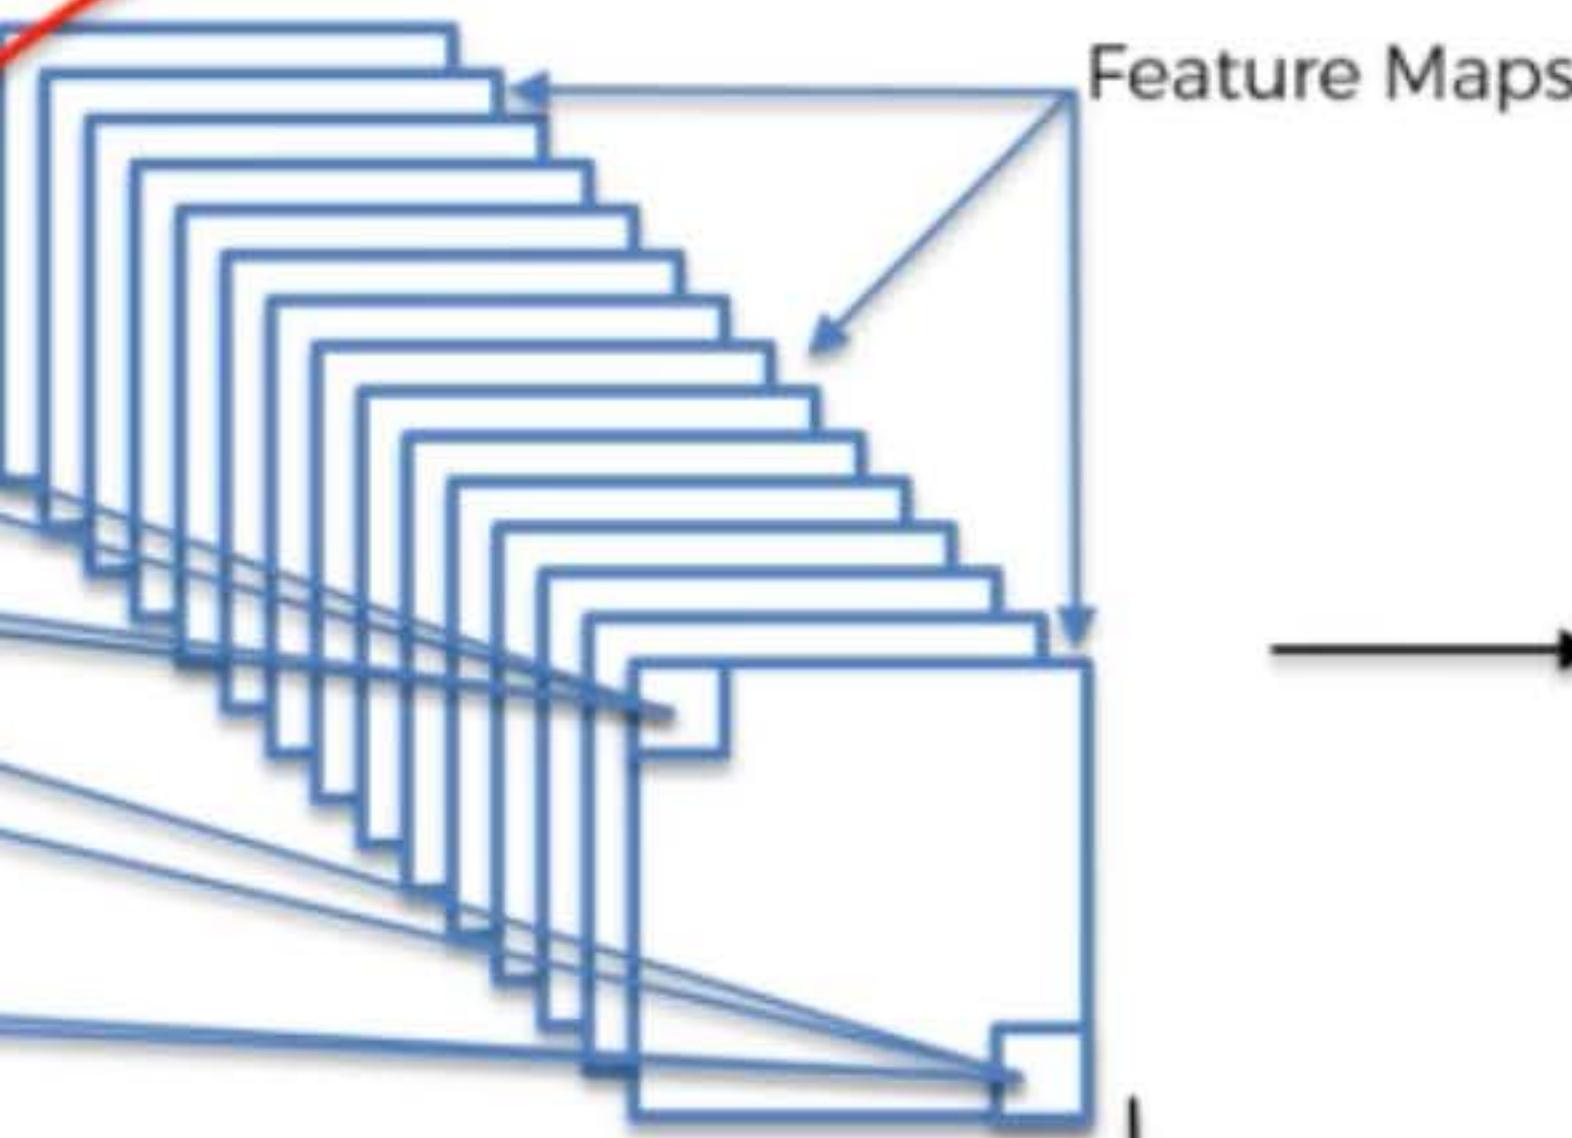
Step 1(B) - ReLU Layer

Step 1(B) - ReLU Layer

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	1	0	0	0	1	0	0
0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0

Input Image

Convolutional Layer



Rectifier

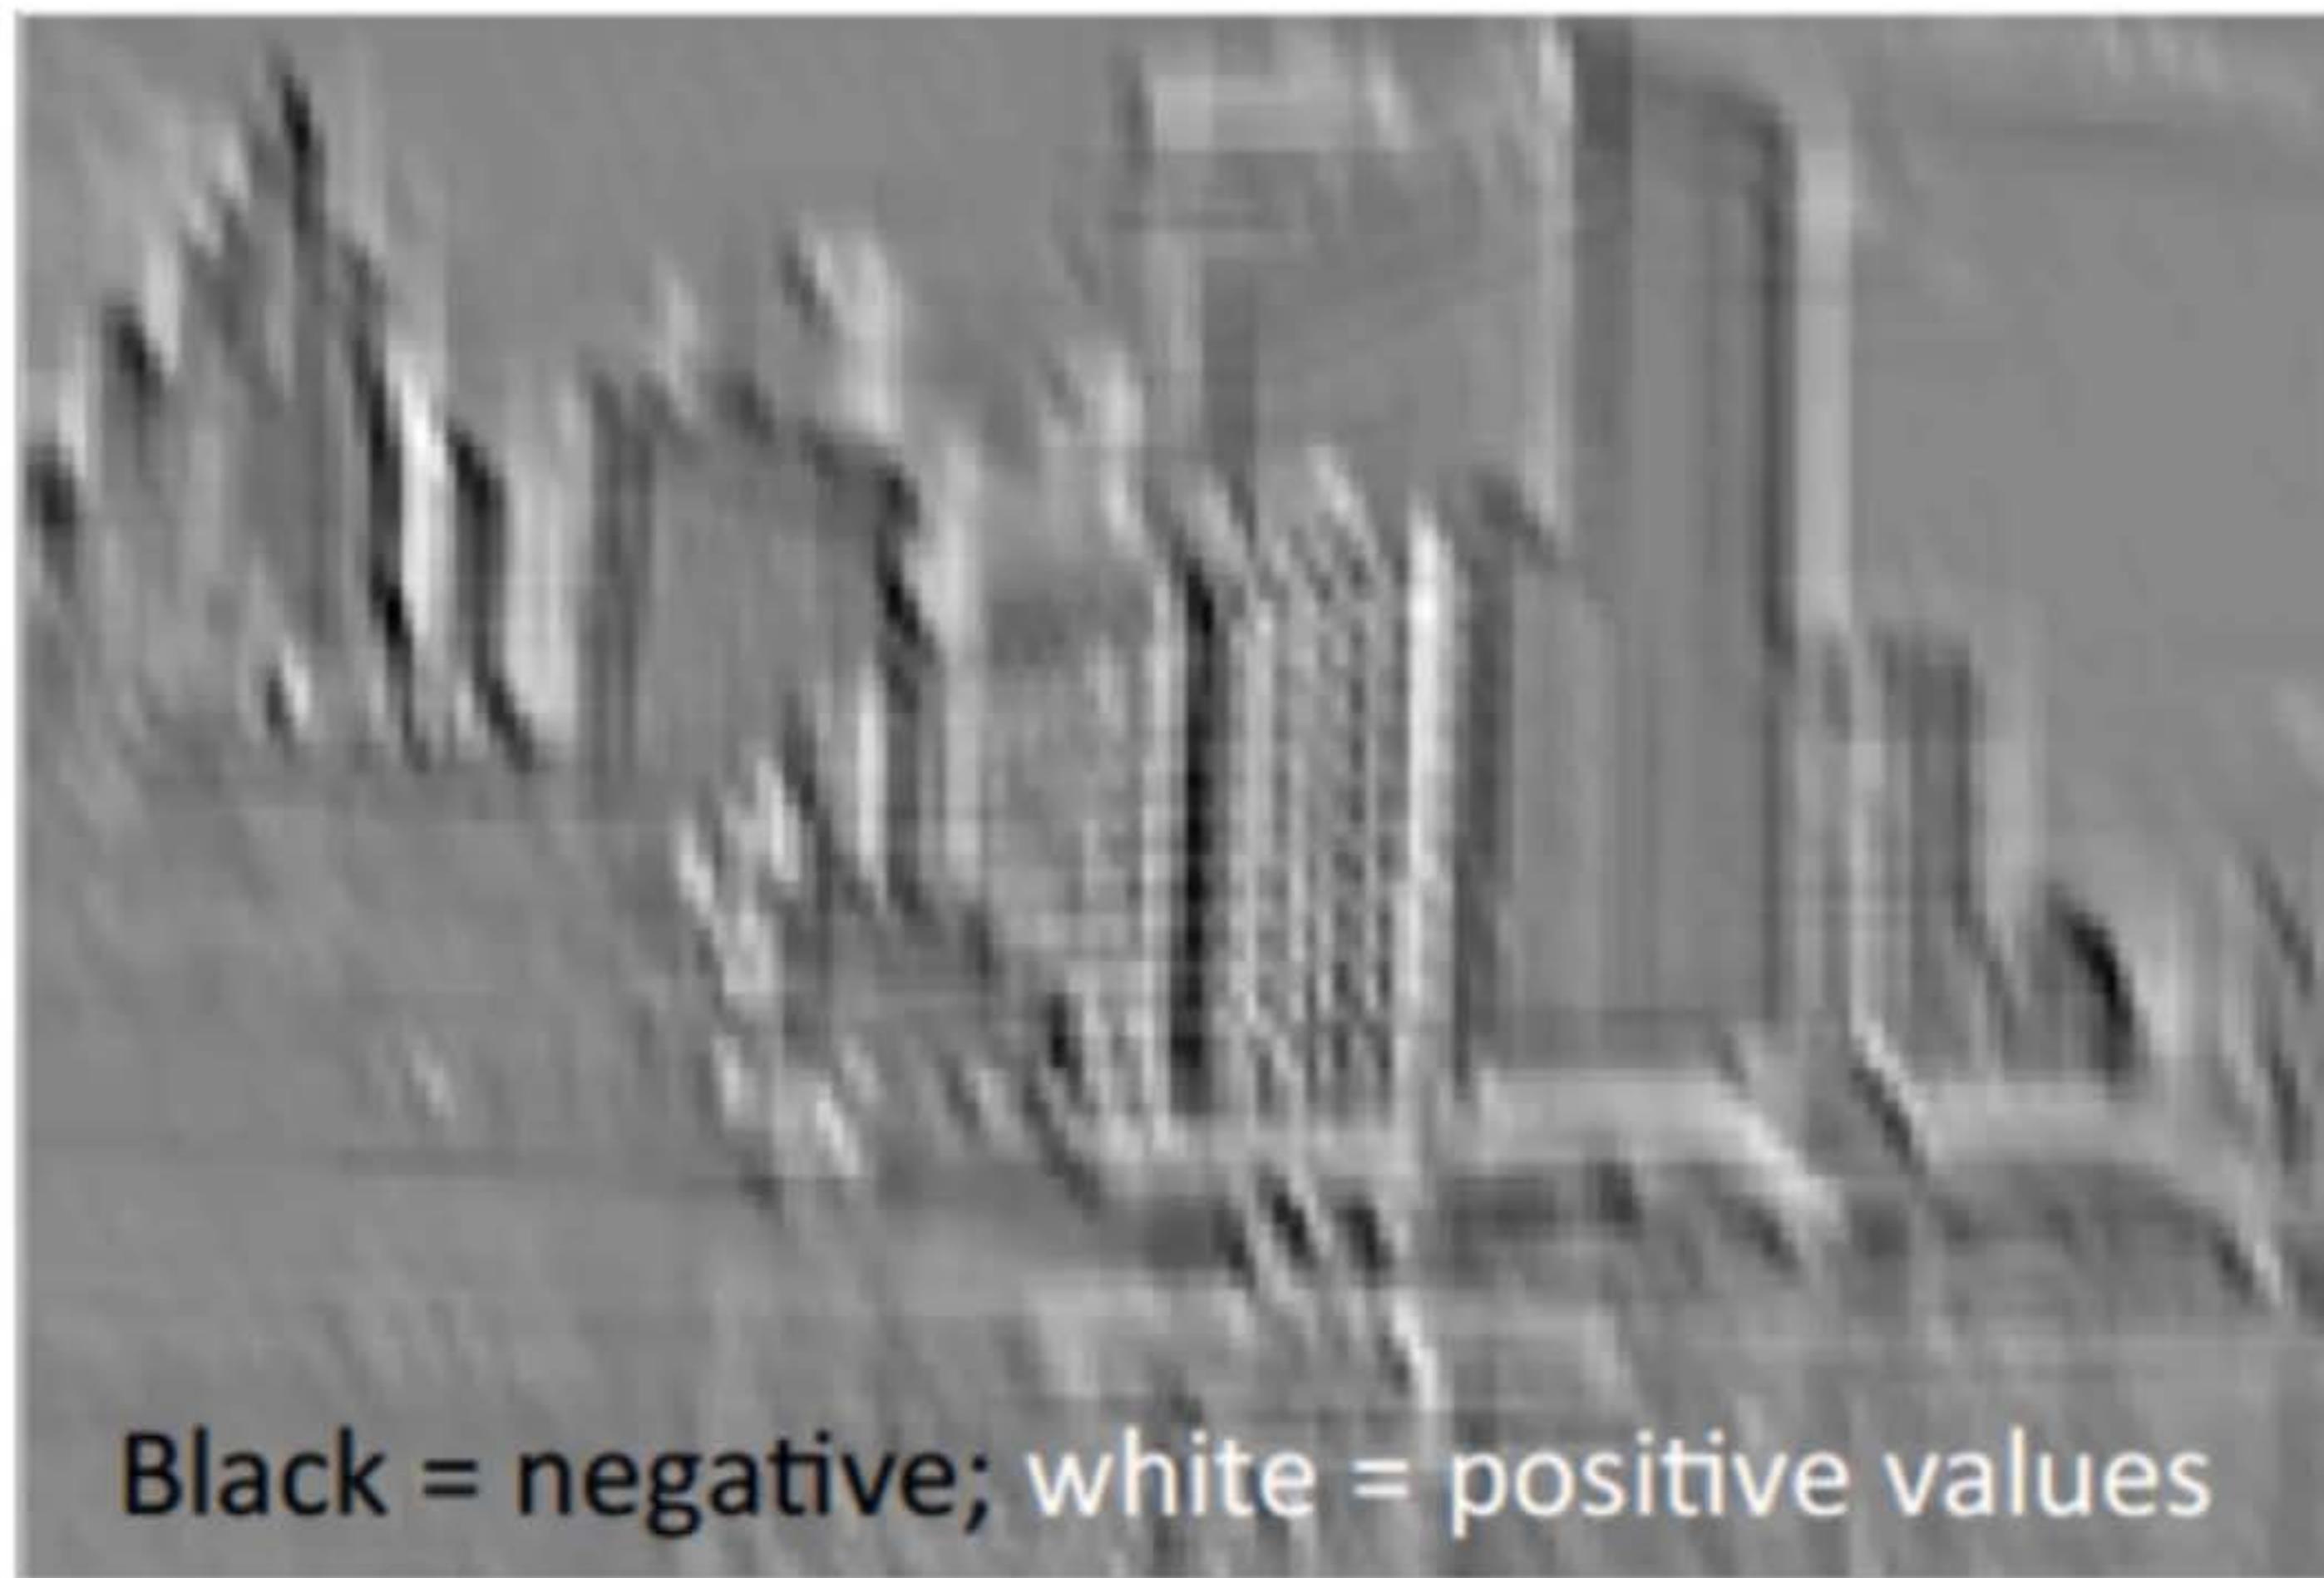
$$\phi(x) = \max(x, 0)$$
$$\sum_{i=1}^m w_i x_i$$

Step 1(B) - ReLU Layer



Image Source: http://mlss.tuebingen.mpg.de/2015/slides/fergus/Fergus_1.pdf

Step 1(B) - ReLU Layer



Black = negative; white = positive values

Image Source: http://mlss.tuebingen.mpg.de/2015/slides/fergus/Fergus_1.pdf

Step 1(B) - ReLU Layer



Only non-negative values

Image Source: http://mlss.tuebingen.mpg.de/2015/slides/fergus/Fergus_1.pdf

Step 1(B) - ReLU Layer

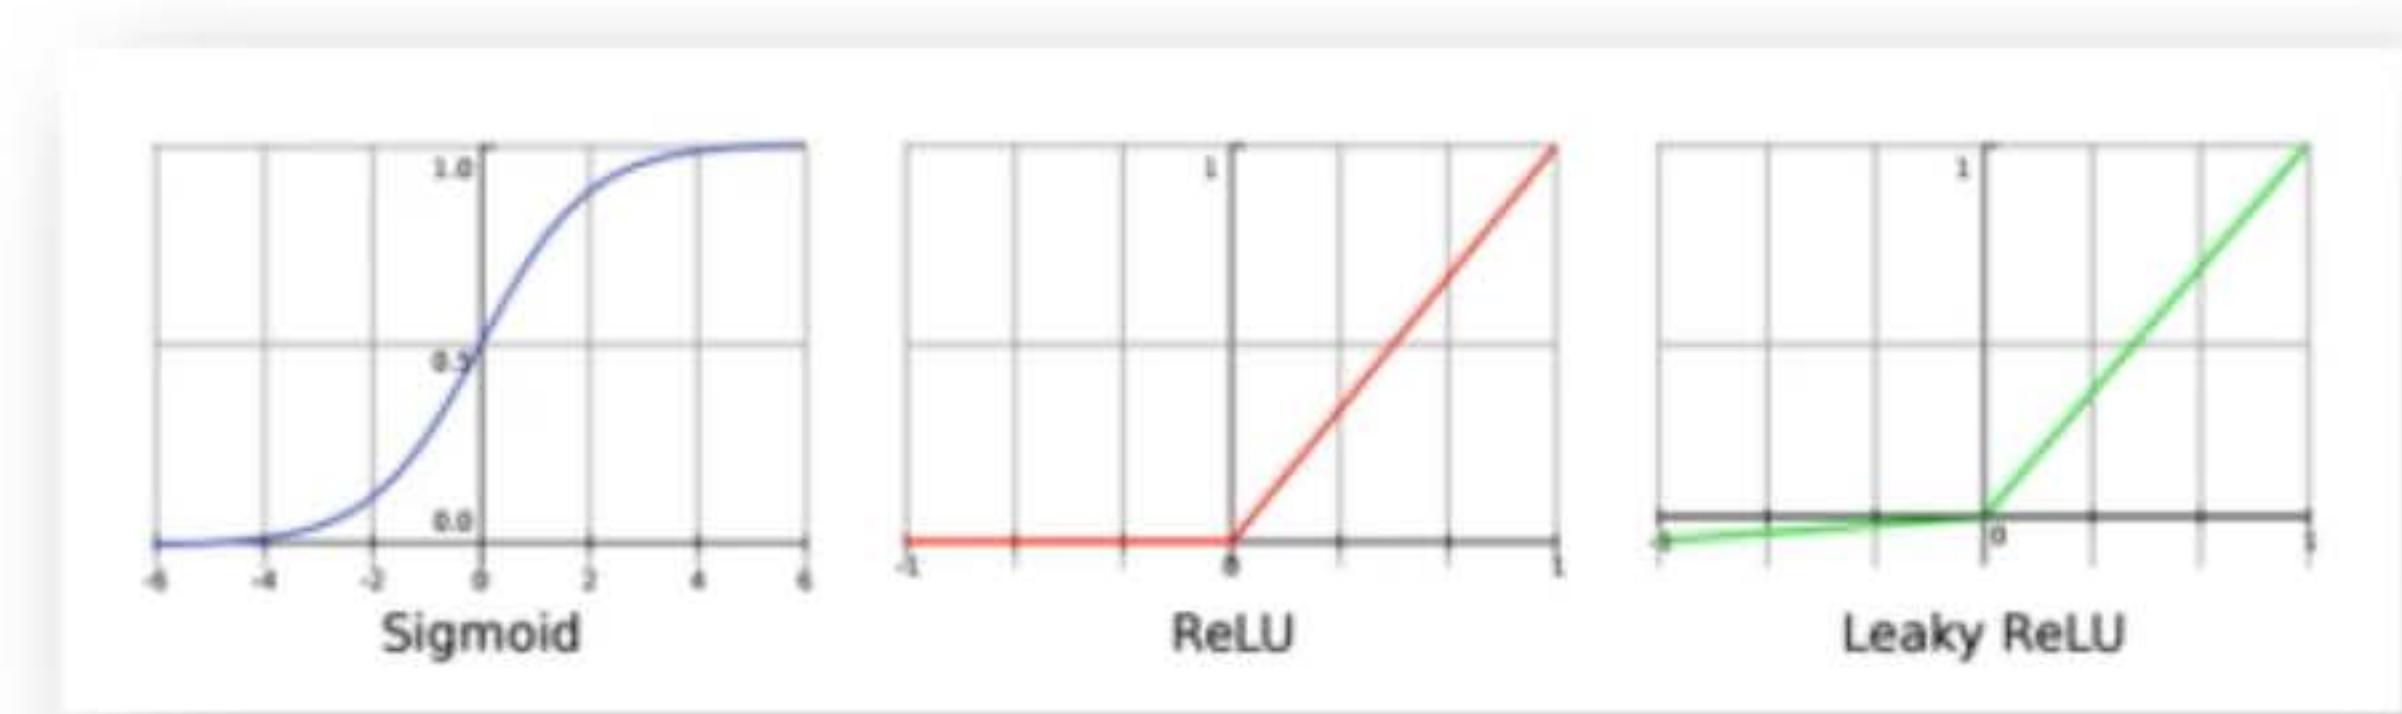
Additional Reading:

Understanding Convolutional Neural Networks with A Mathematical Model

By C.-C. Jay Kuo (2016)

Link:

<https://arxiv.org/pdf/1609.04112.pdf>



Step 1(B) - ReLU Layer

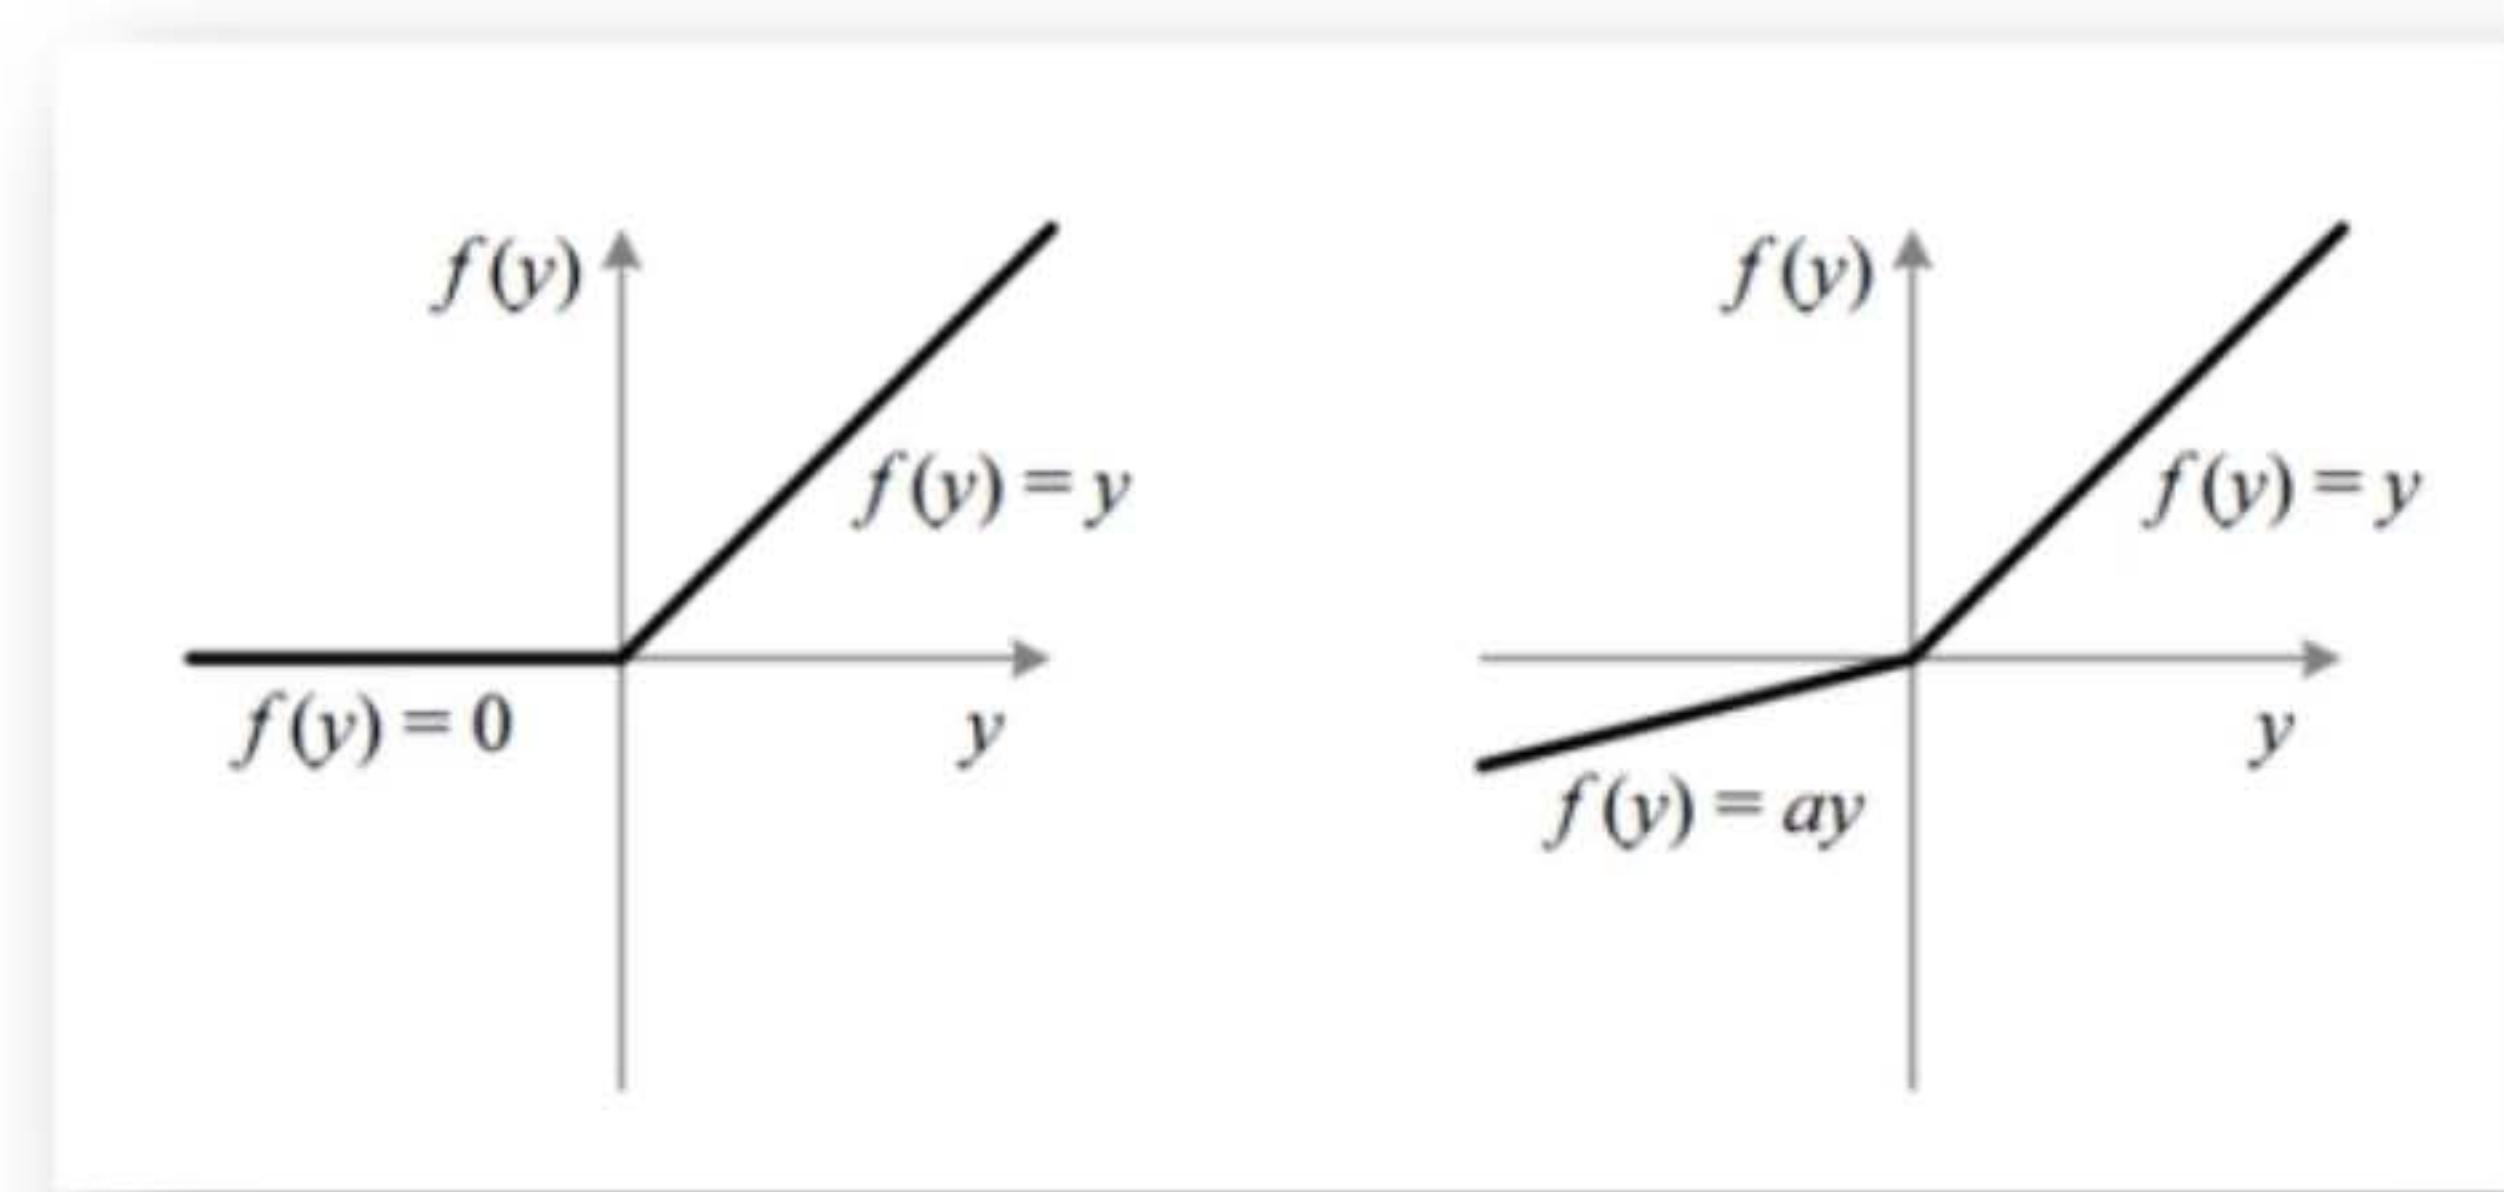
Additional Reading:

*Delving Deep into Rectifiers:
Surpassing Human-Level
Performance on ImageNet
Classification*

By Kaiming He et al. (2015)

Link:

<https://arxiv.org/pdf/1502.01852.pdf>



Step 2 - Max Pooling

Step 2 - Max Pooling

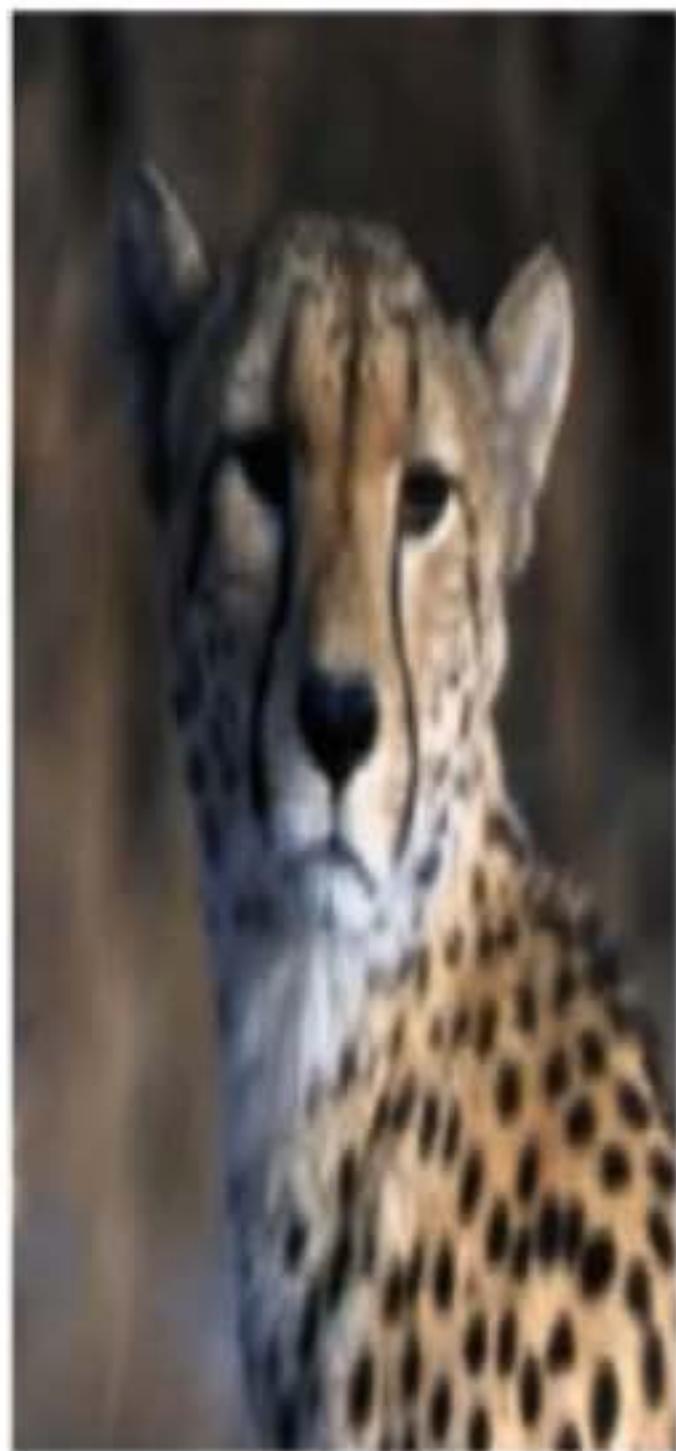
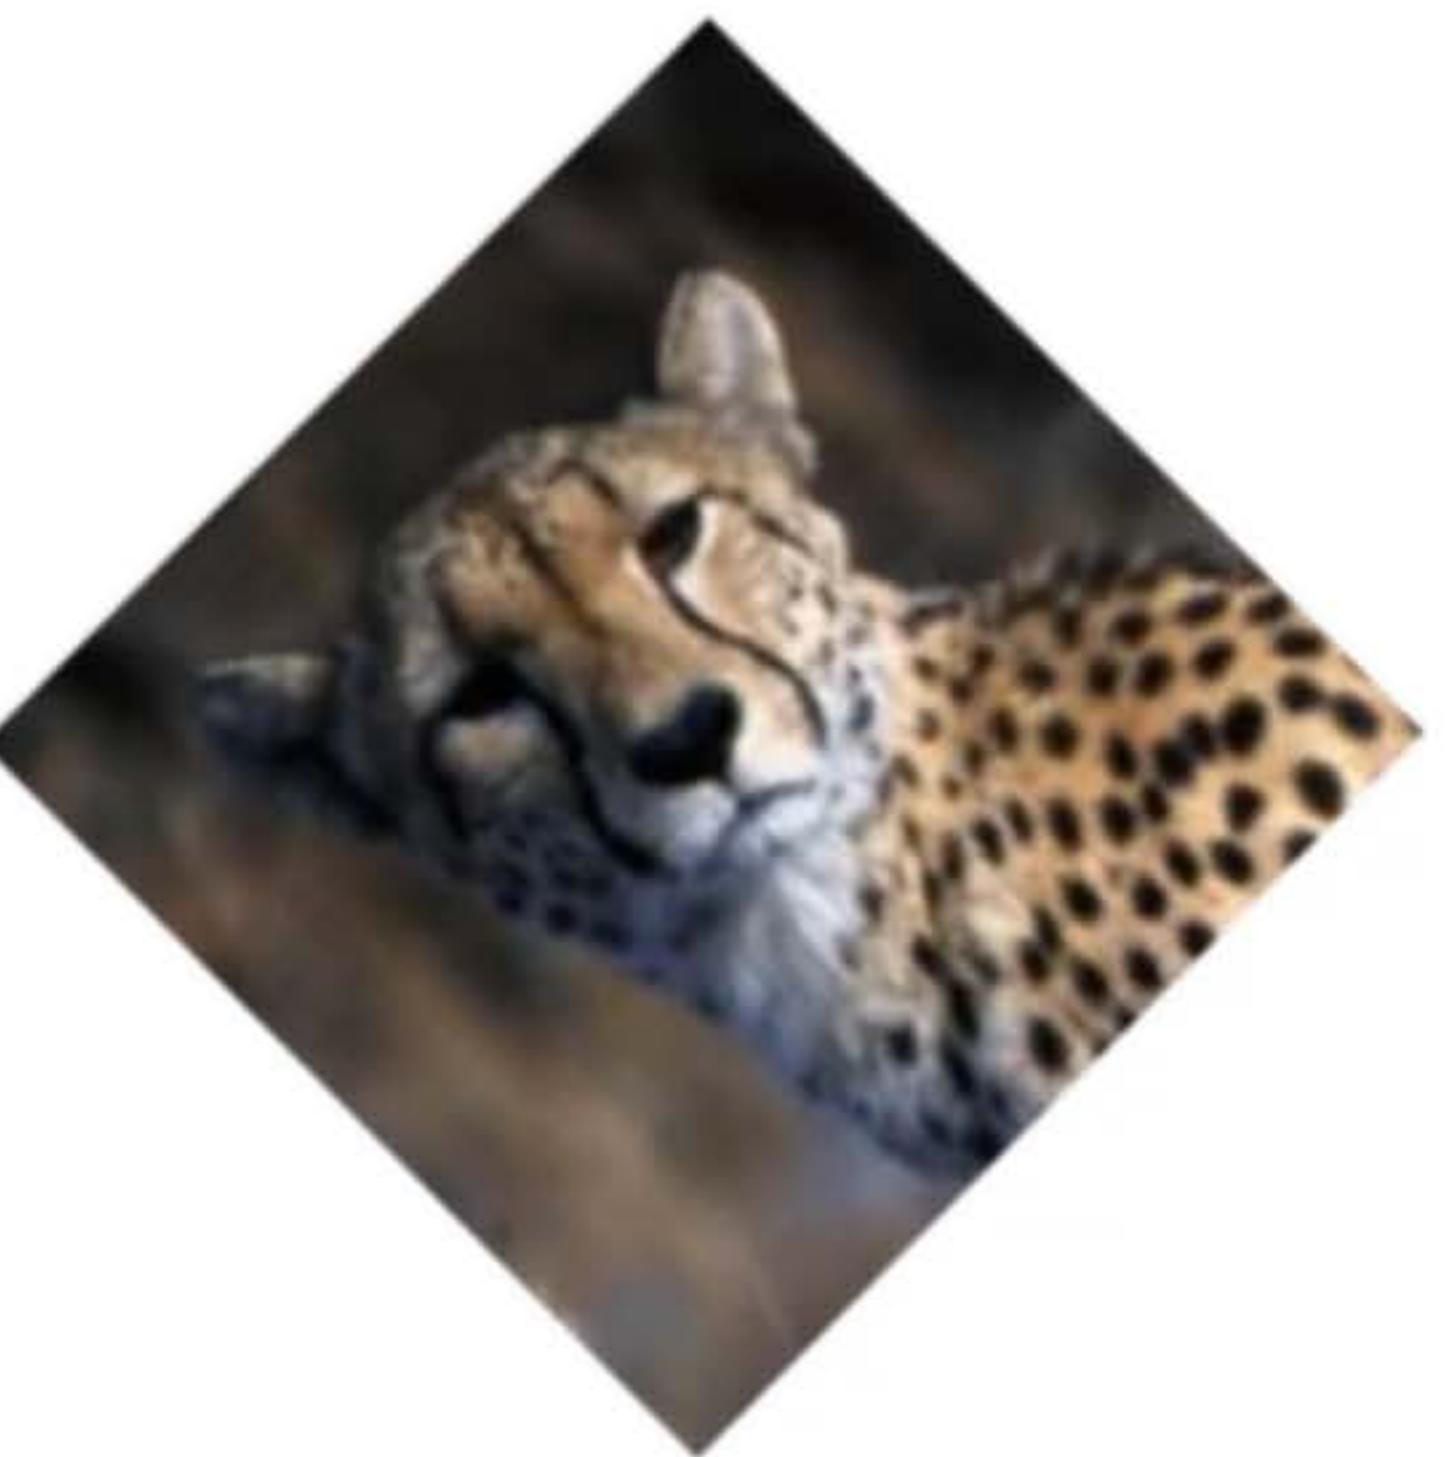
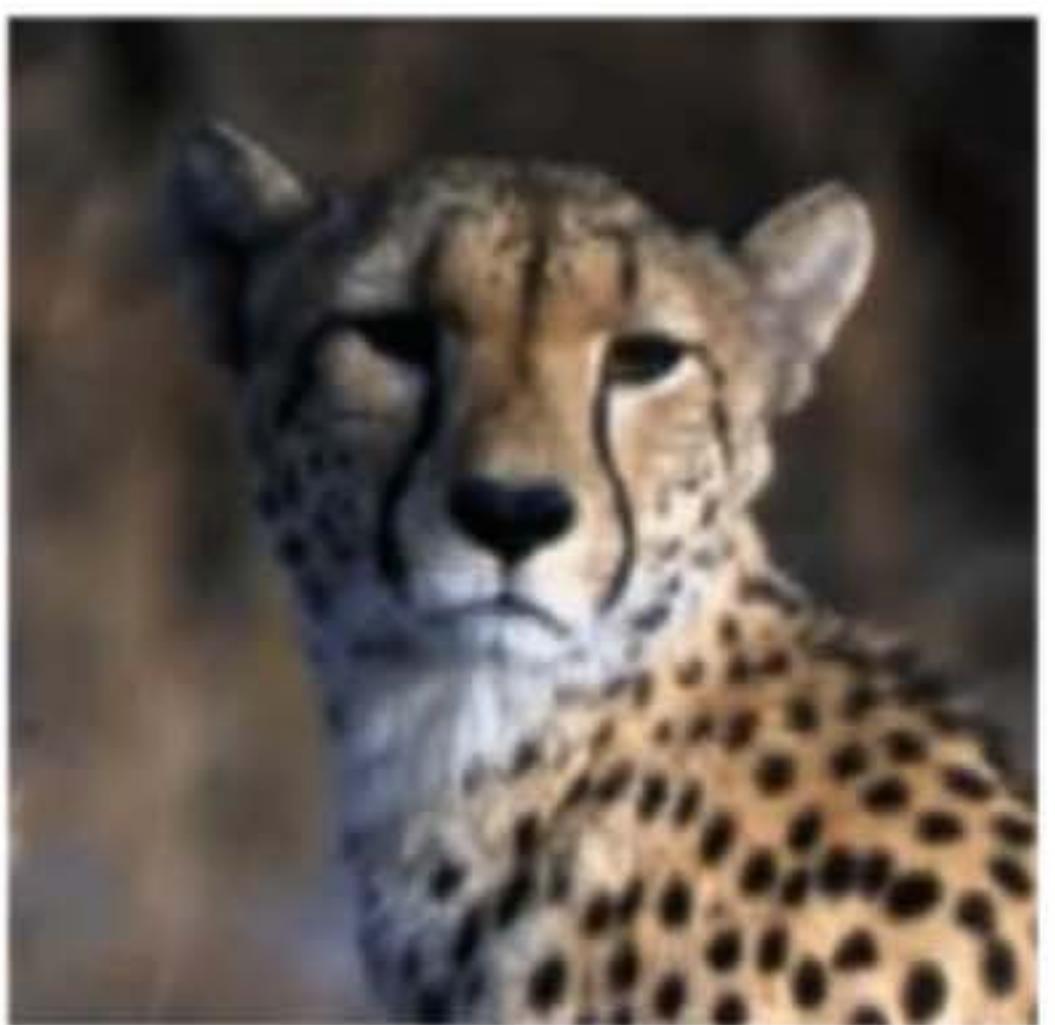


Image Source: Wikipedia

Step 2 - Max Pooling



Image Source: Wikipedia

Step 2 - Max Pooling

0	1	0	0	0	
0	1	1	1	0	
1	0	1	2	1	
1	4	2	1	0	
0	0	1	2	1	

Feature Map

Max Pooling

1	1	0

Pooled Feature Map

Step 2 - Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

1	1	0
4	2	1
0	2	1

Pooled Feature Map

Step 2 - Max Pooling

Additional Reading:

Evaluation of Pooling Operations in Convolutional Architectures for Object Recognition

By Dominik Scherer et al. (2010)

Link:

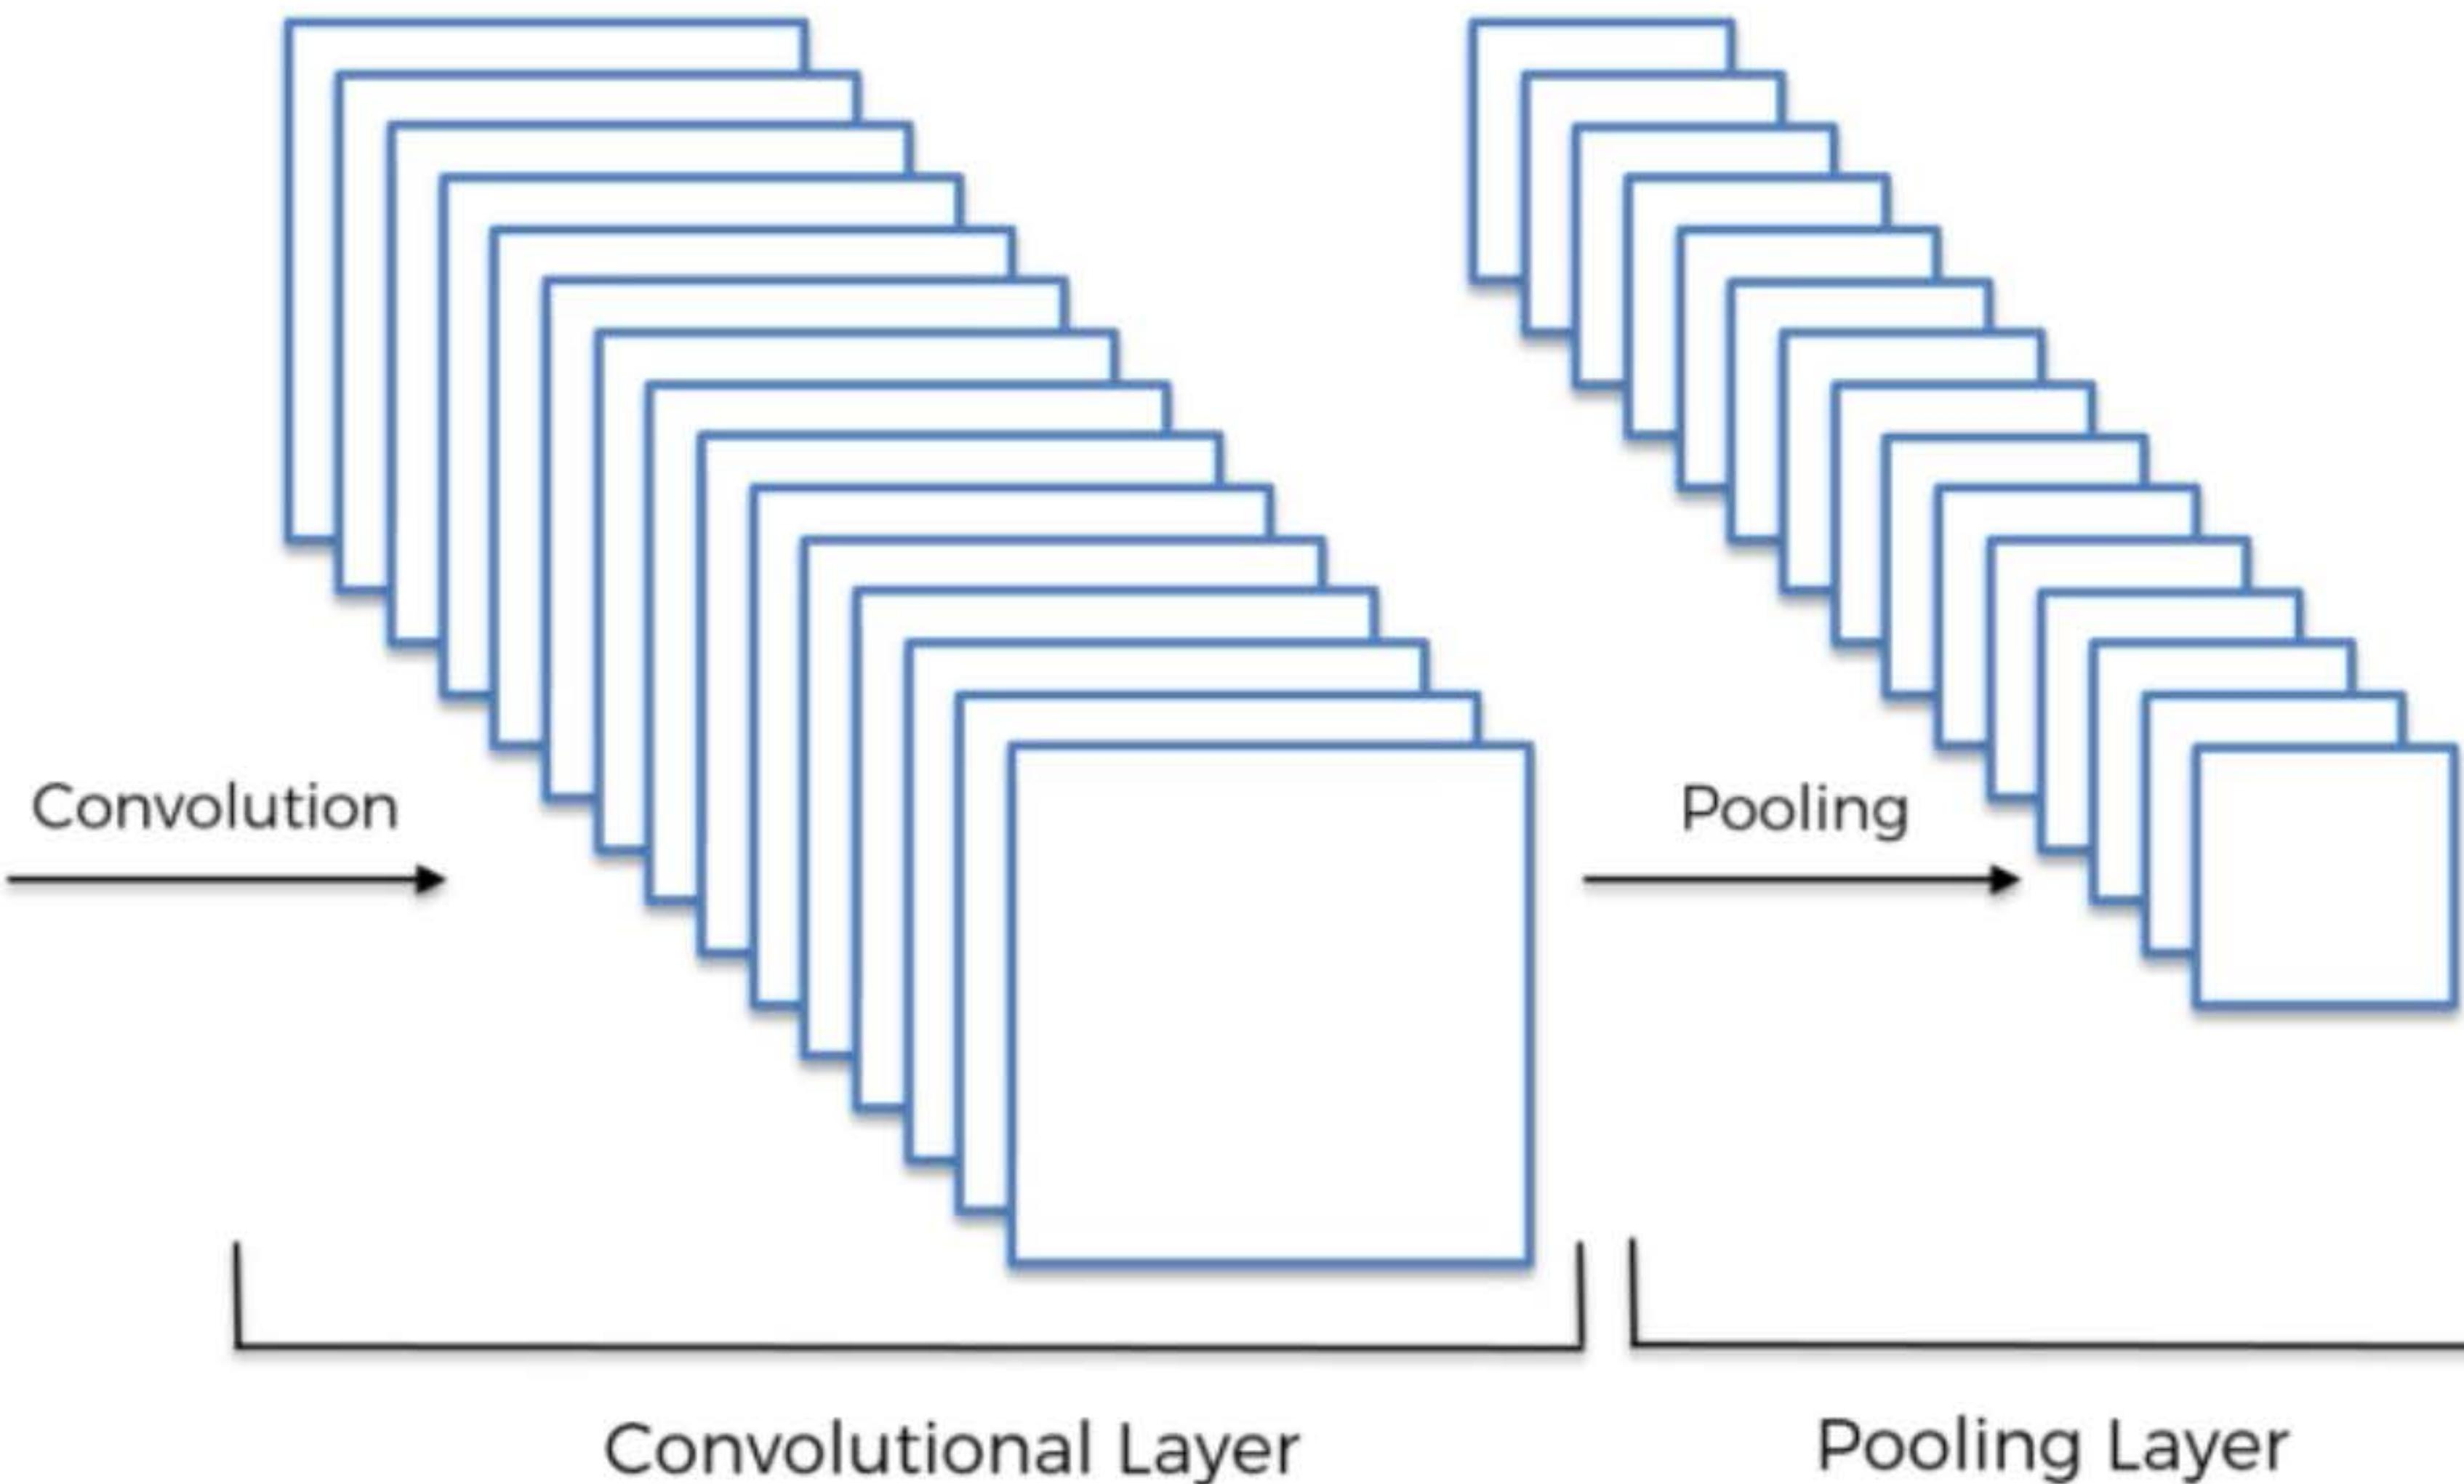
http://ais.uni-bonn.de/papers/icann2010_maxpool.pdf



Step 2 - Max Pooling

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



Example

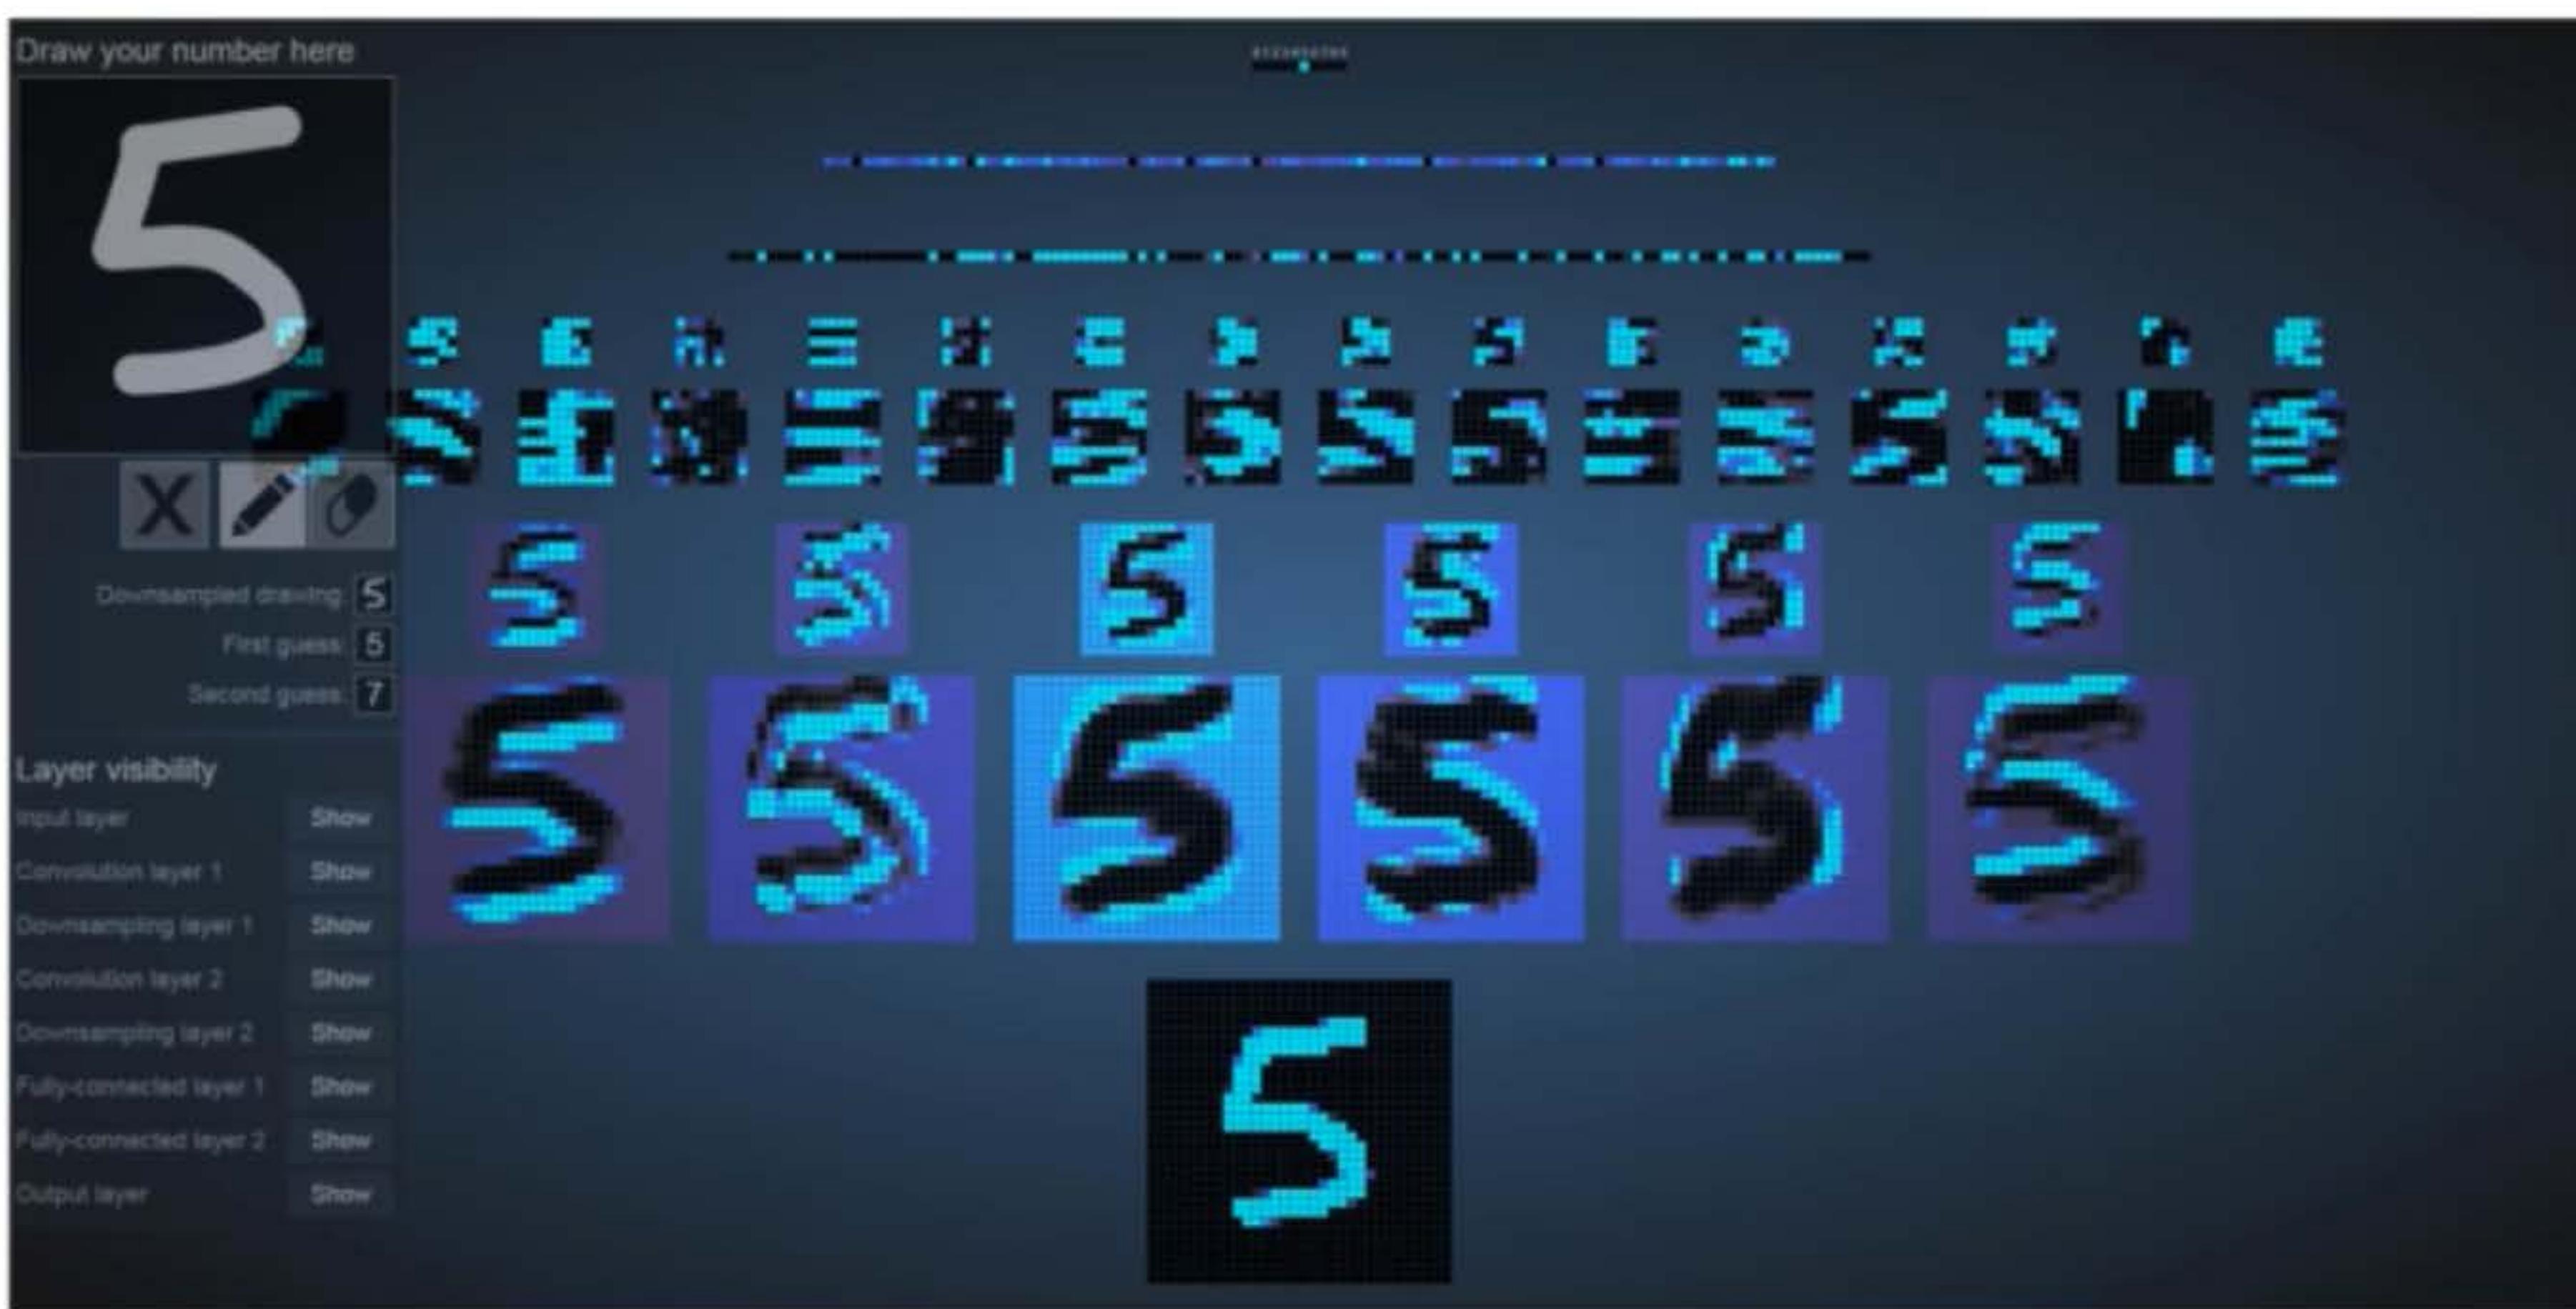


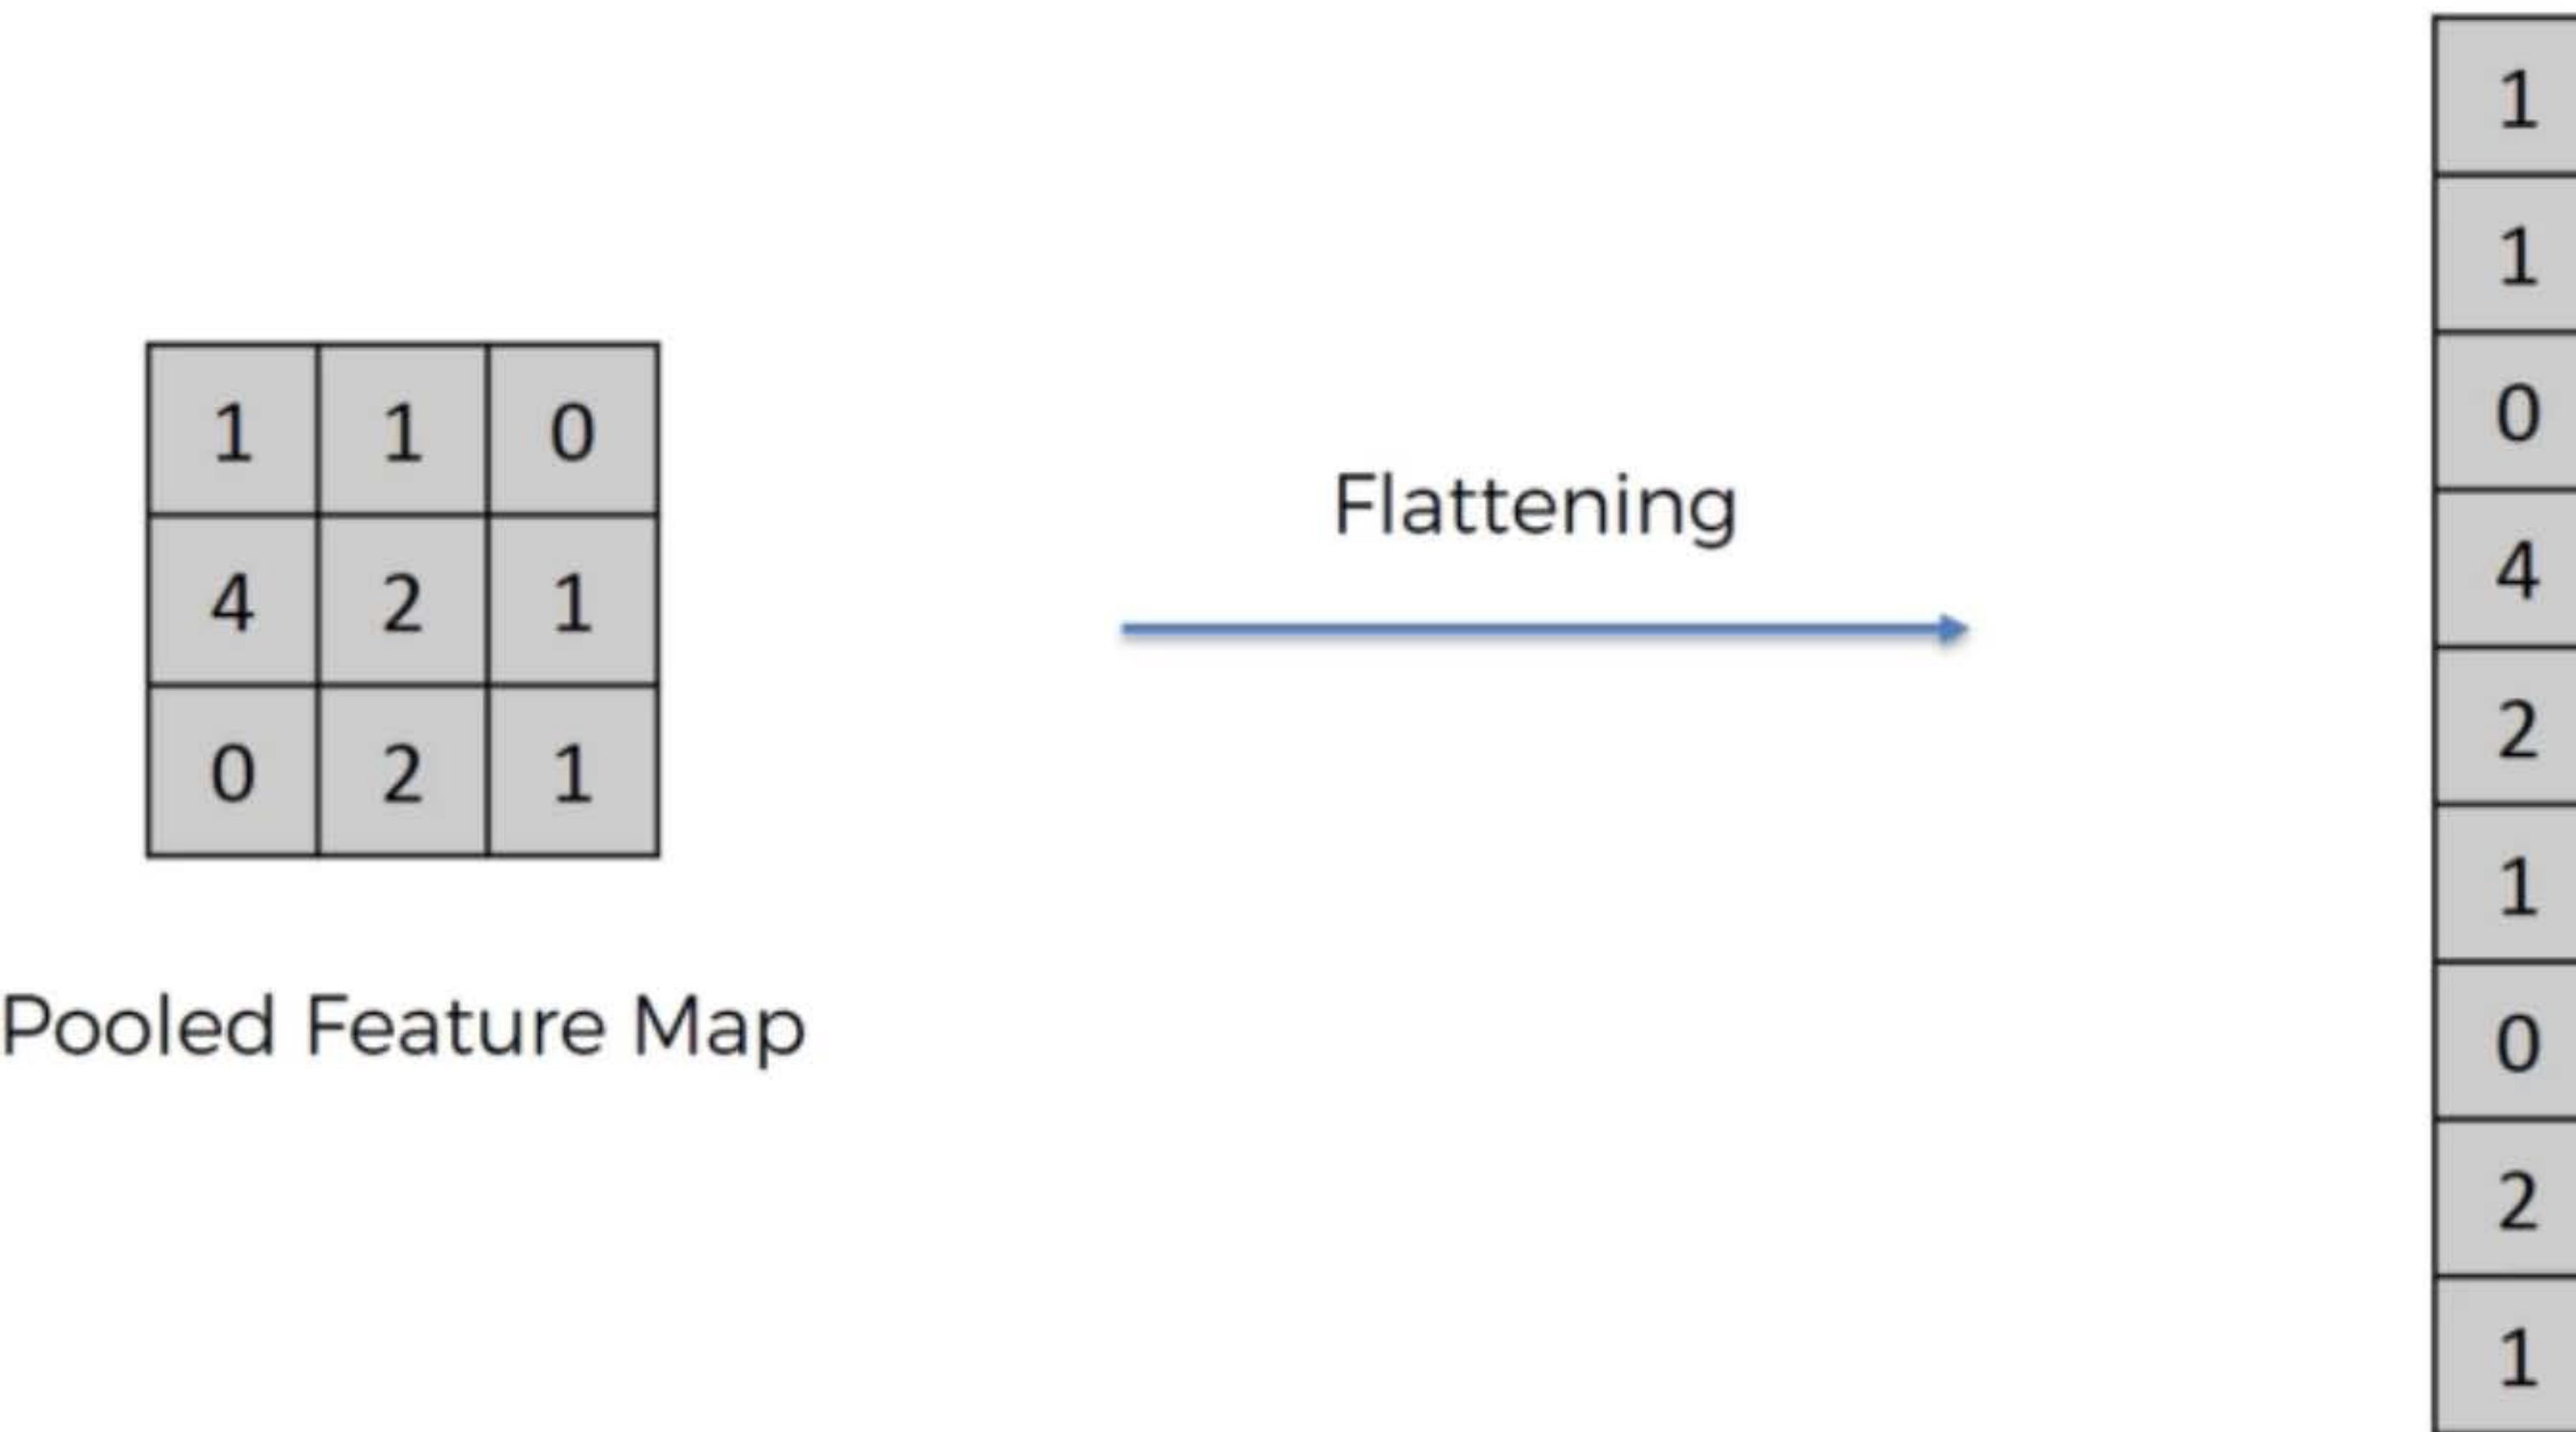
Image Source: scs.ryerson.ca/~aharley/vis/conv/flat.html

Draw your number here

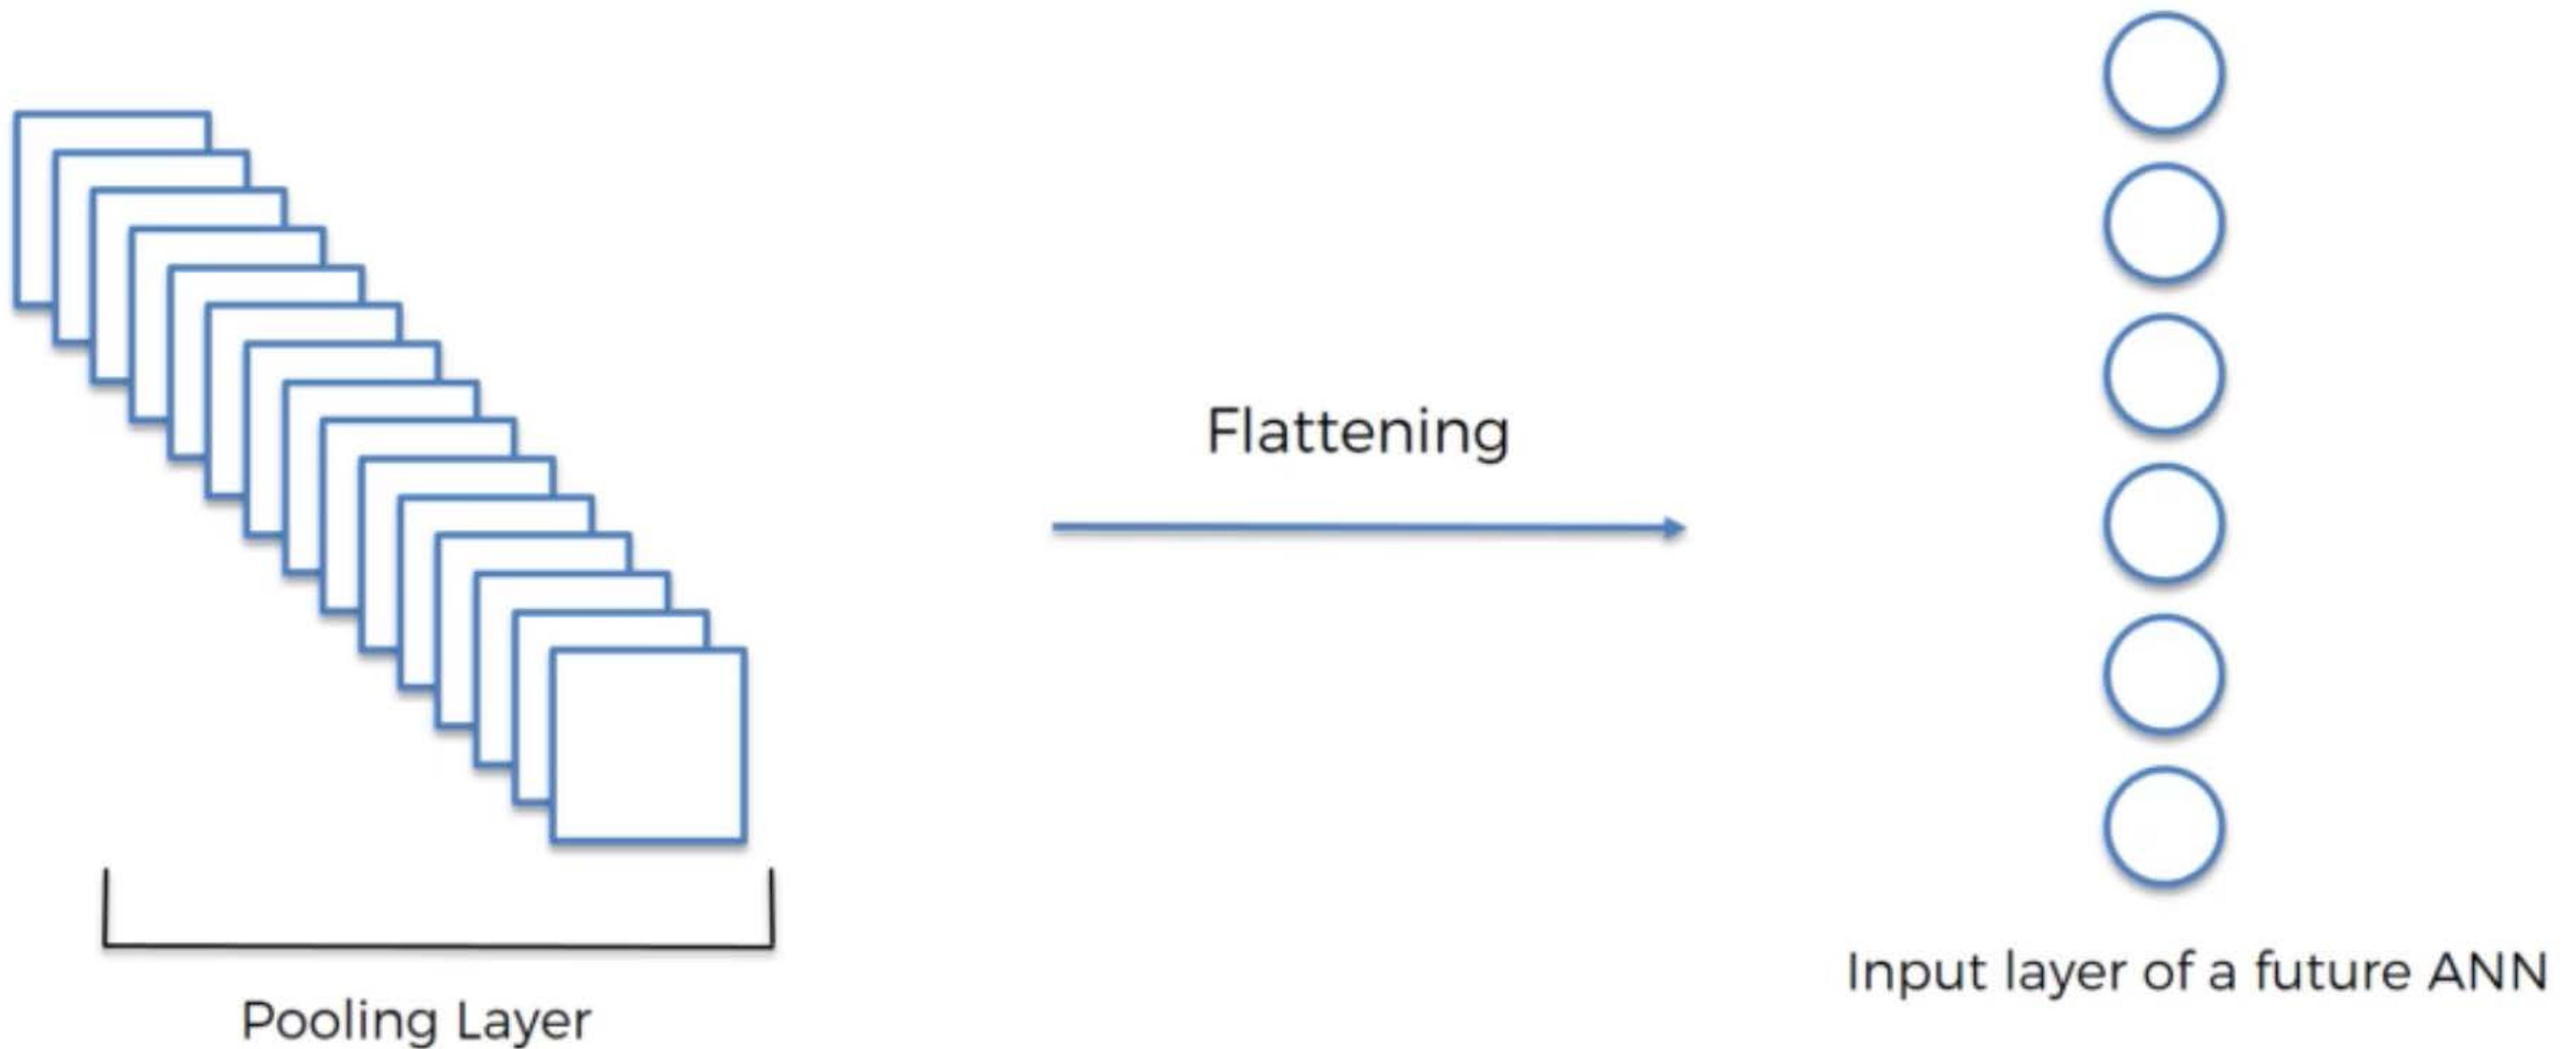


Step 3 - Flattening

Step 3 - Flattening



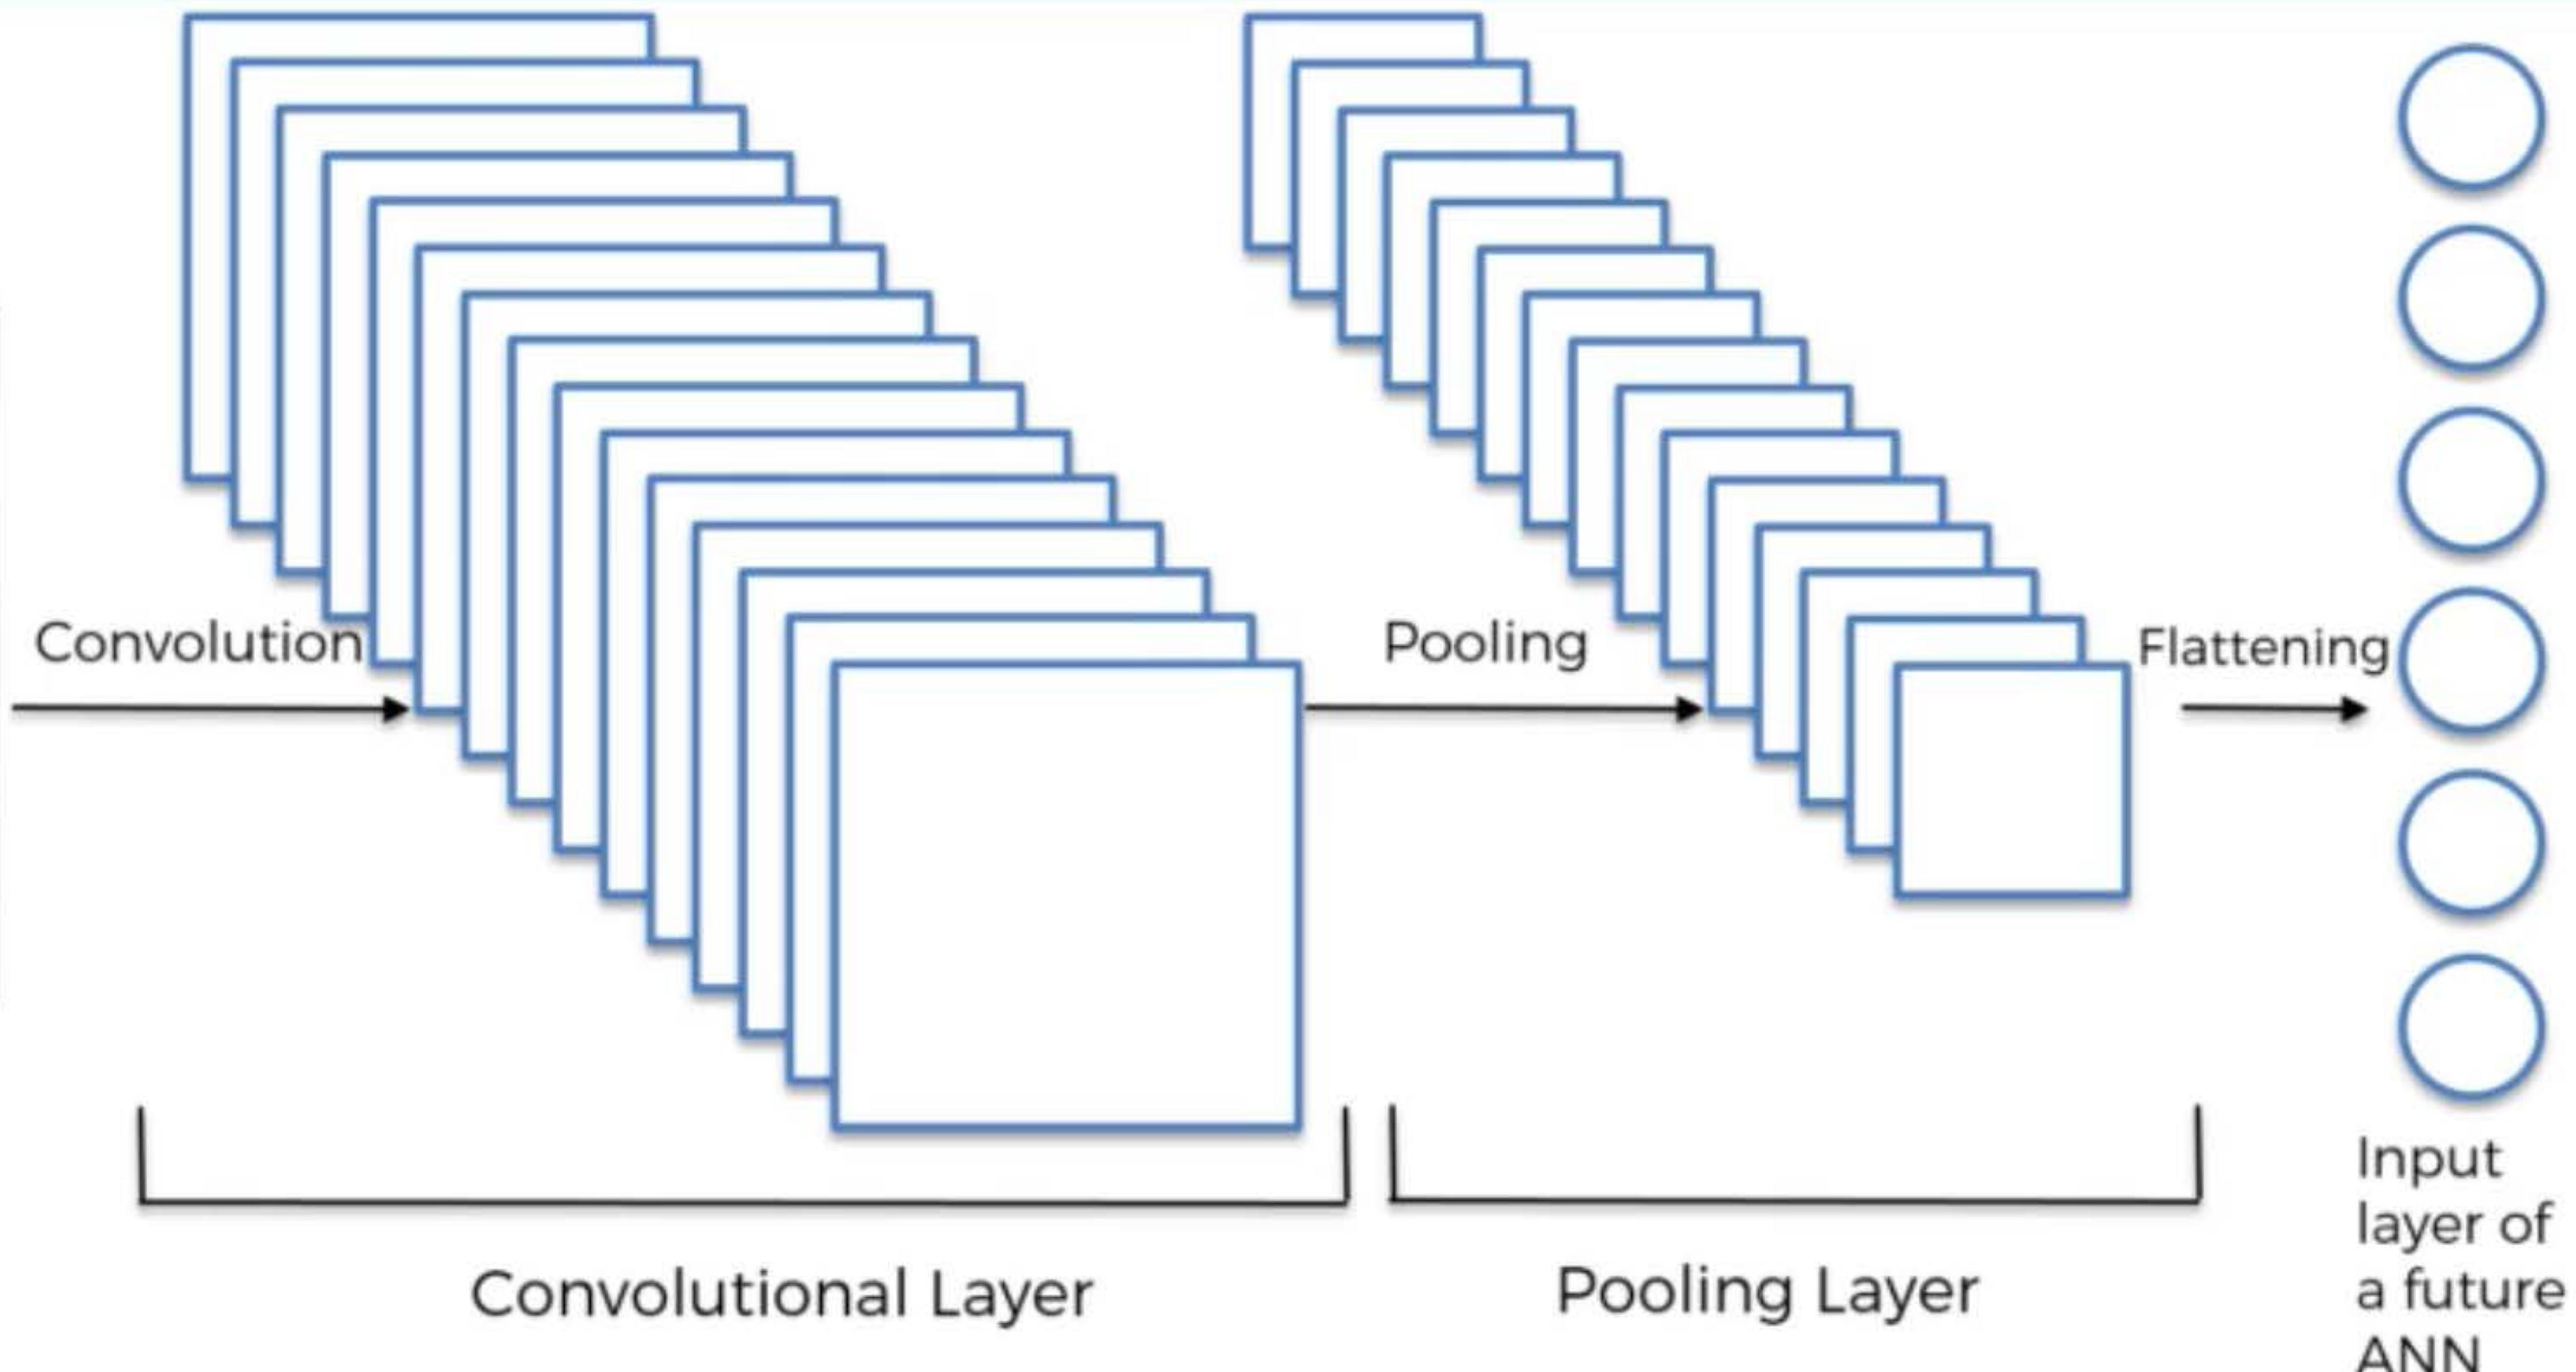
Step 3 - Flattening



Step 3 - Flattening

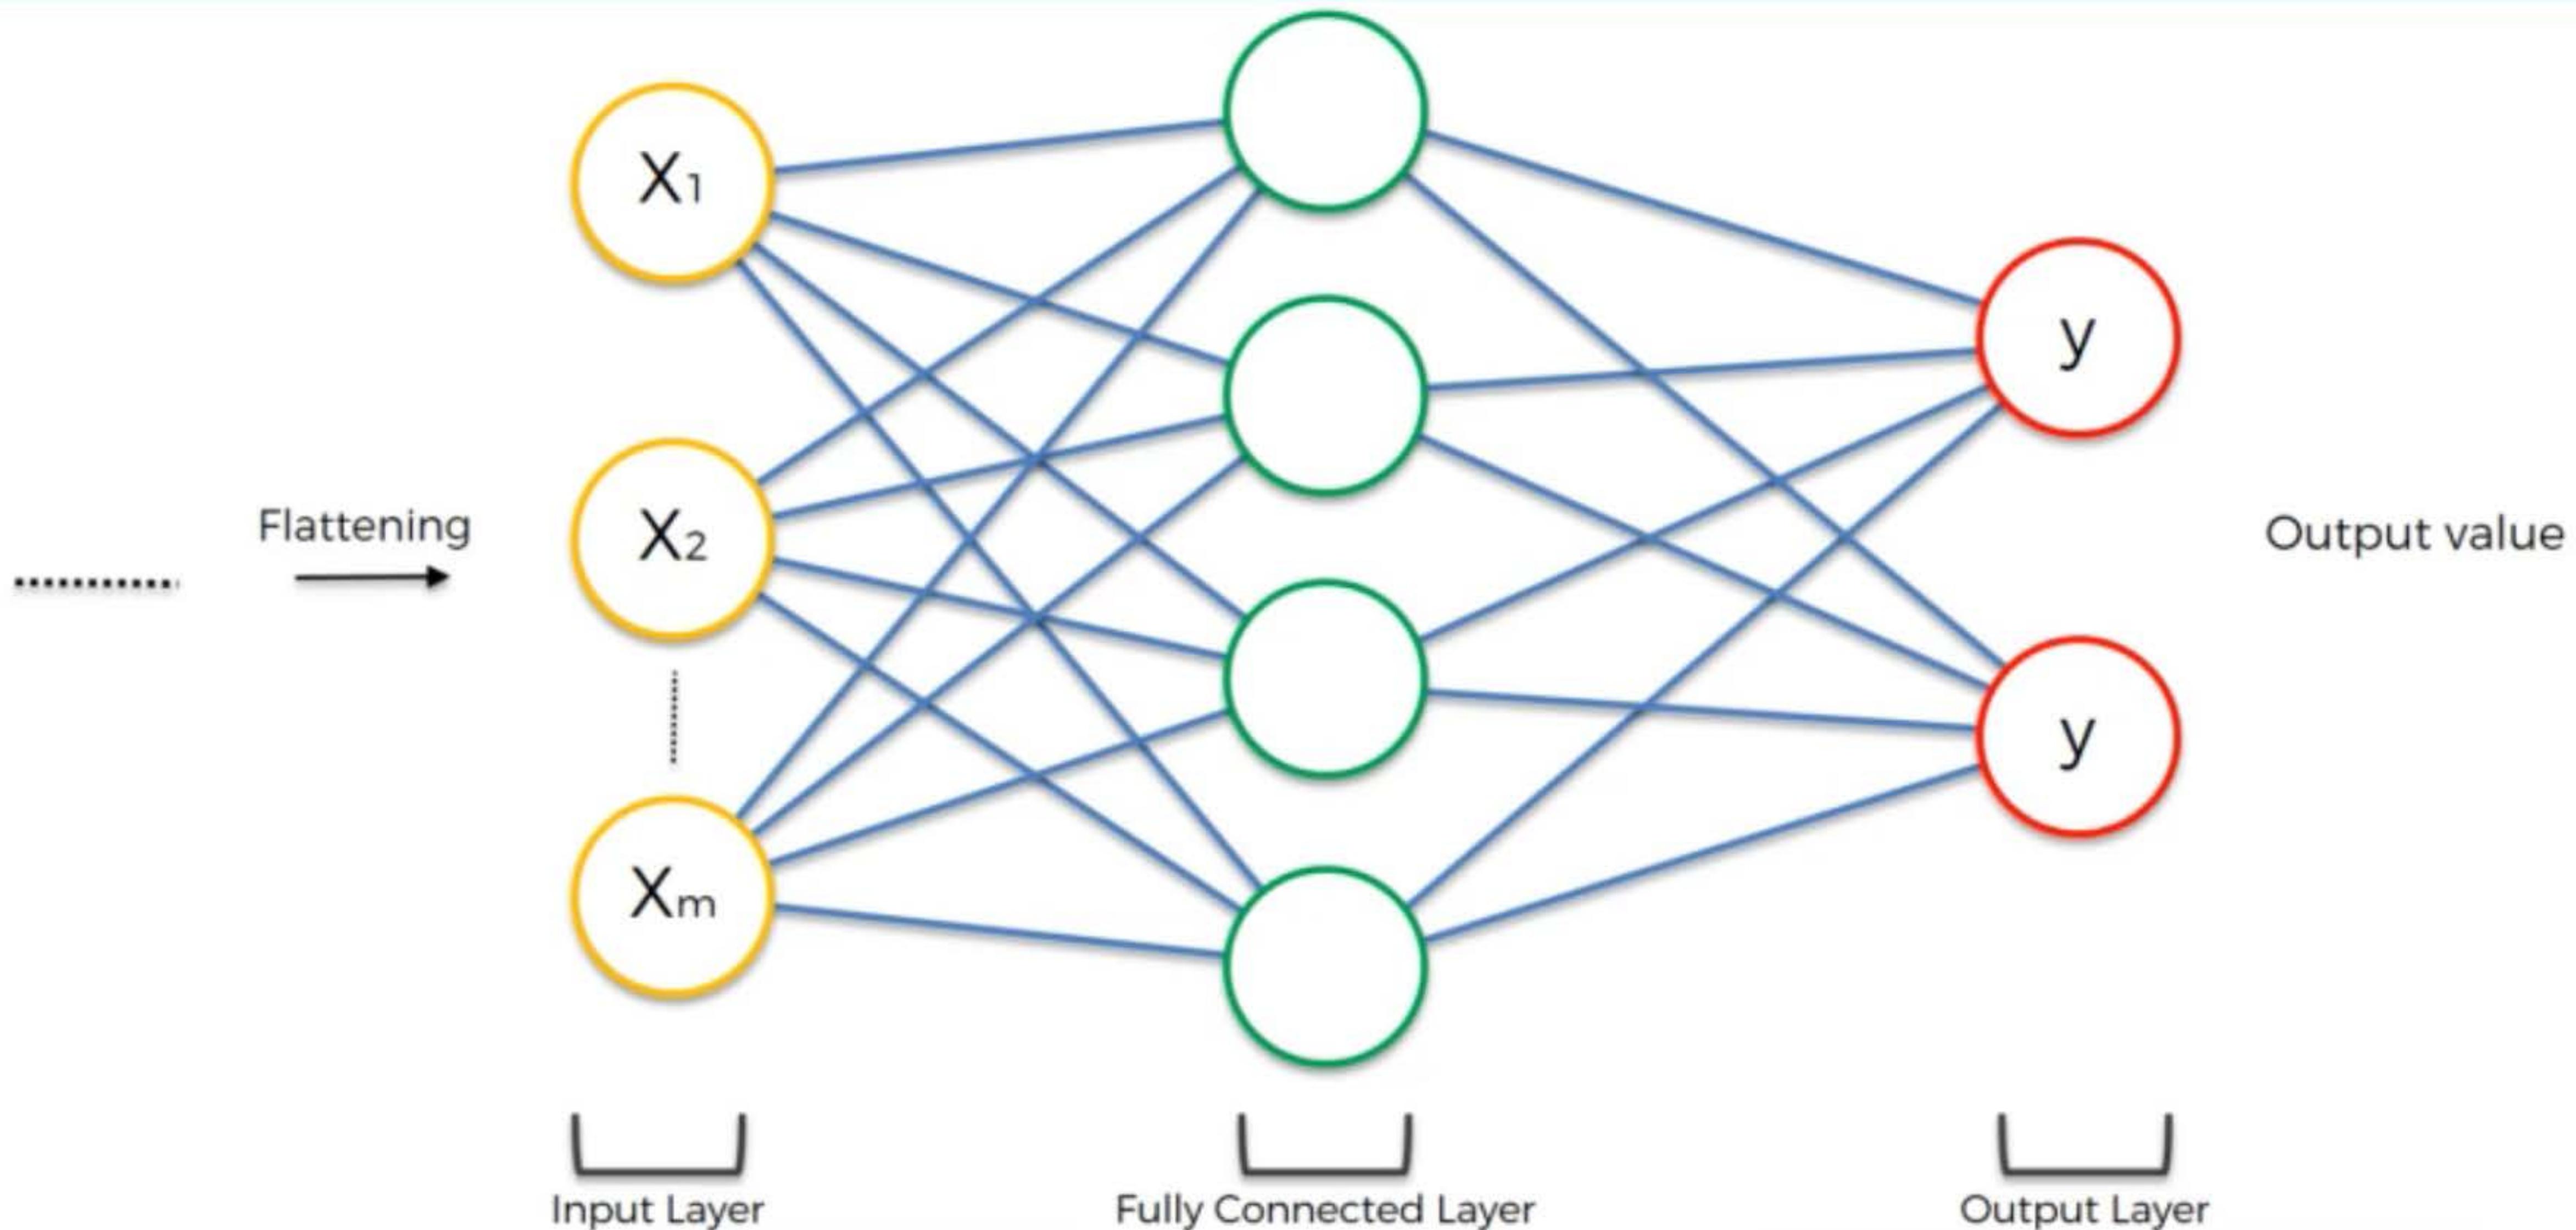
0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image

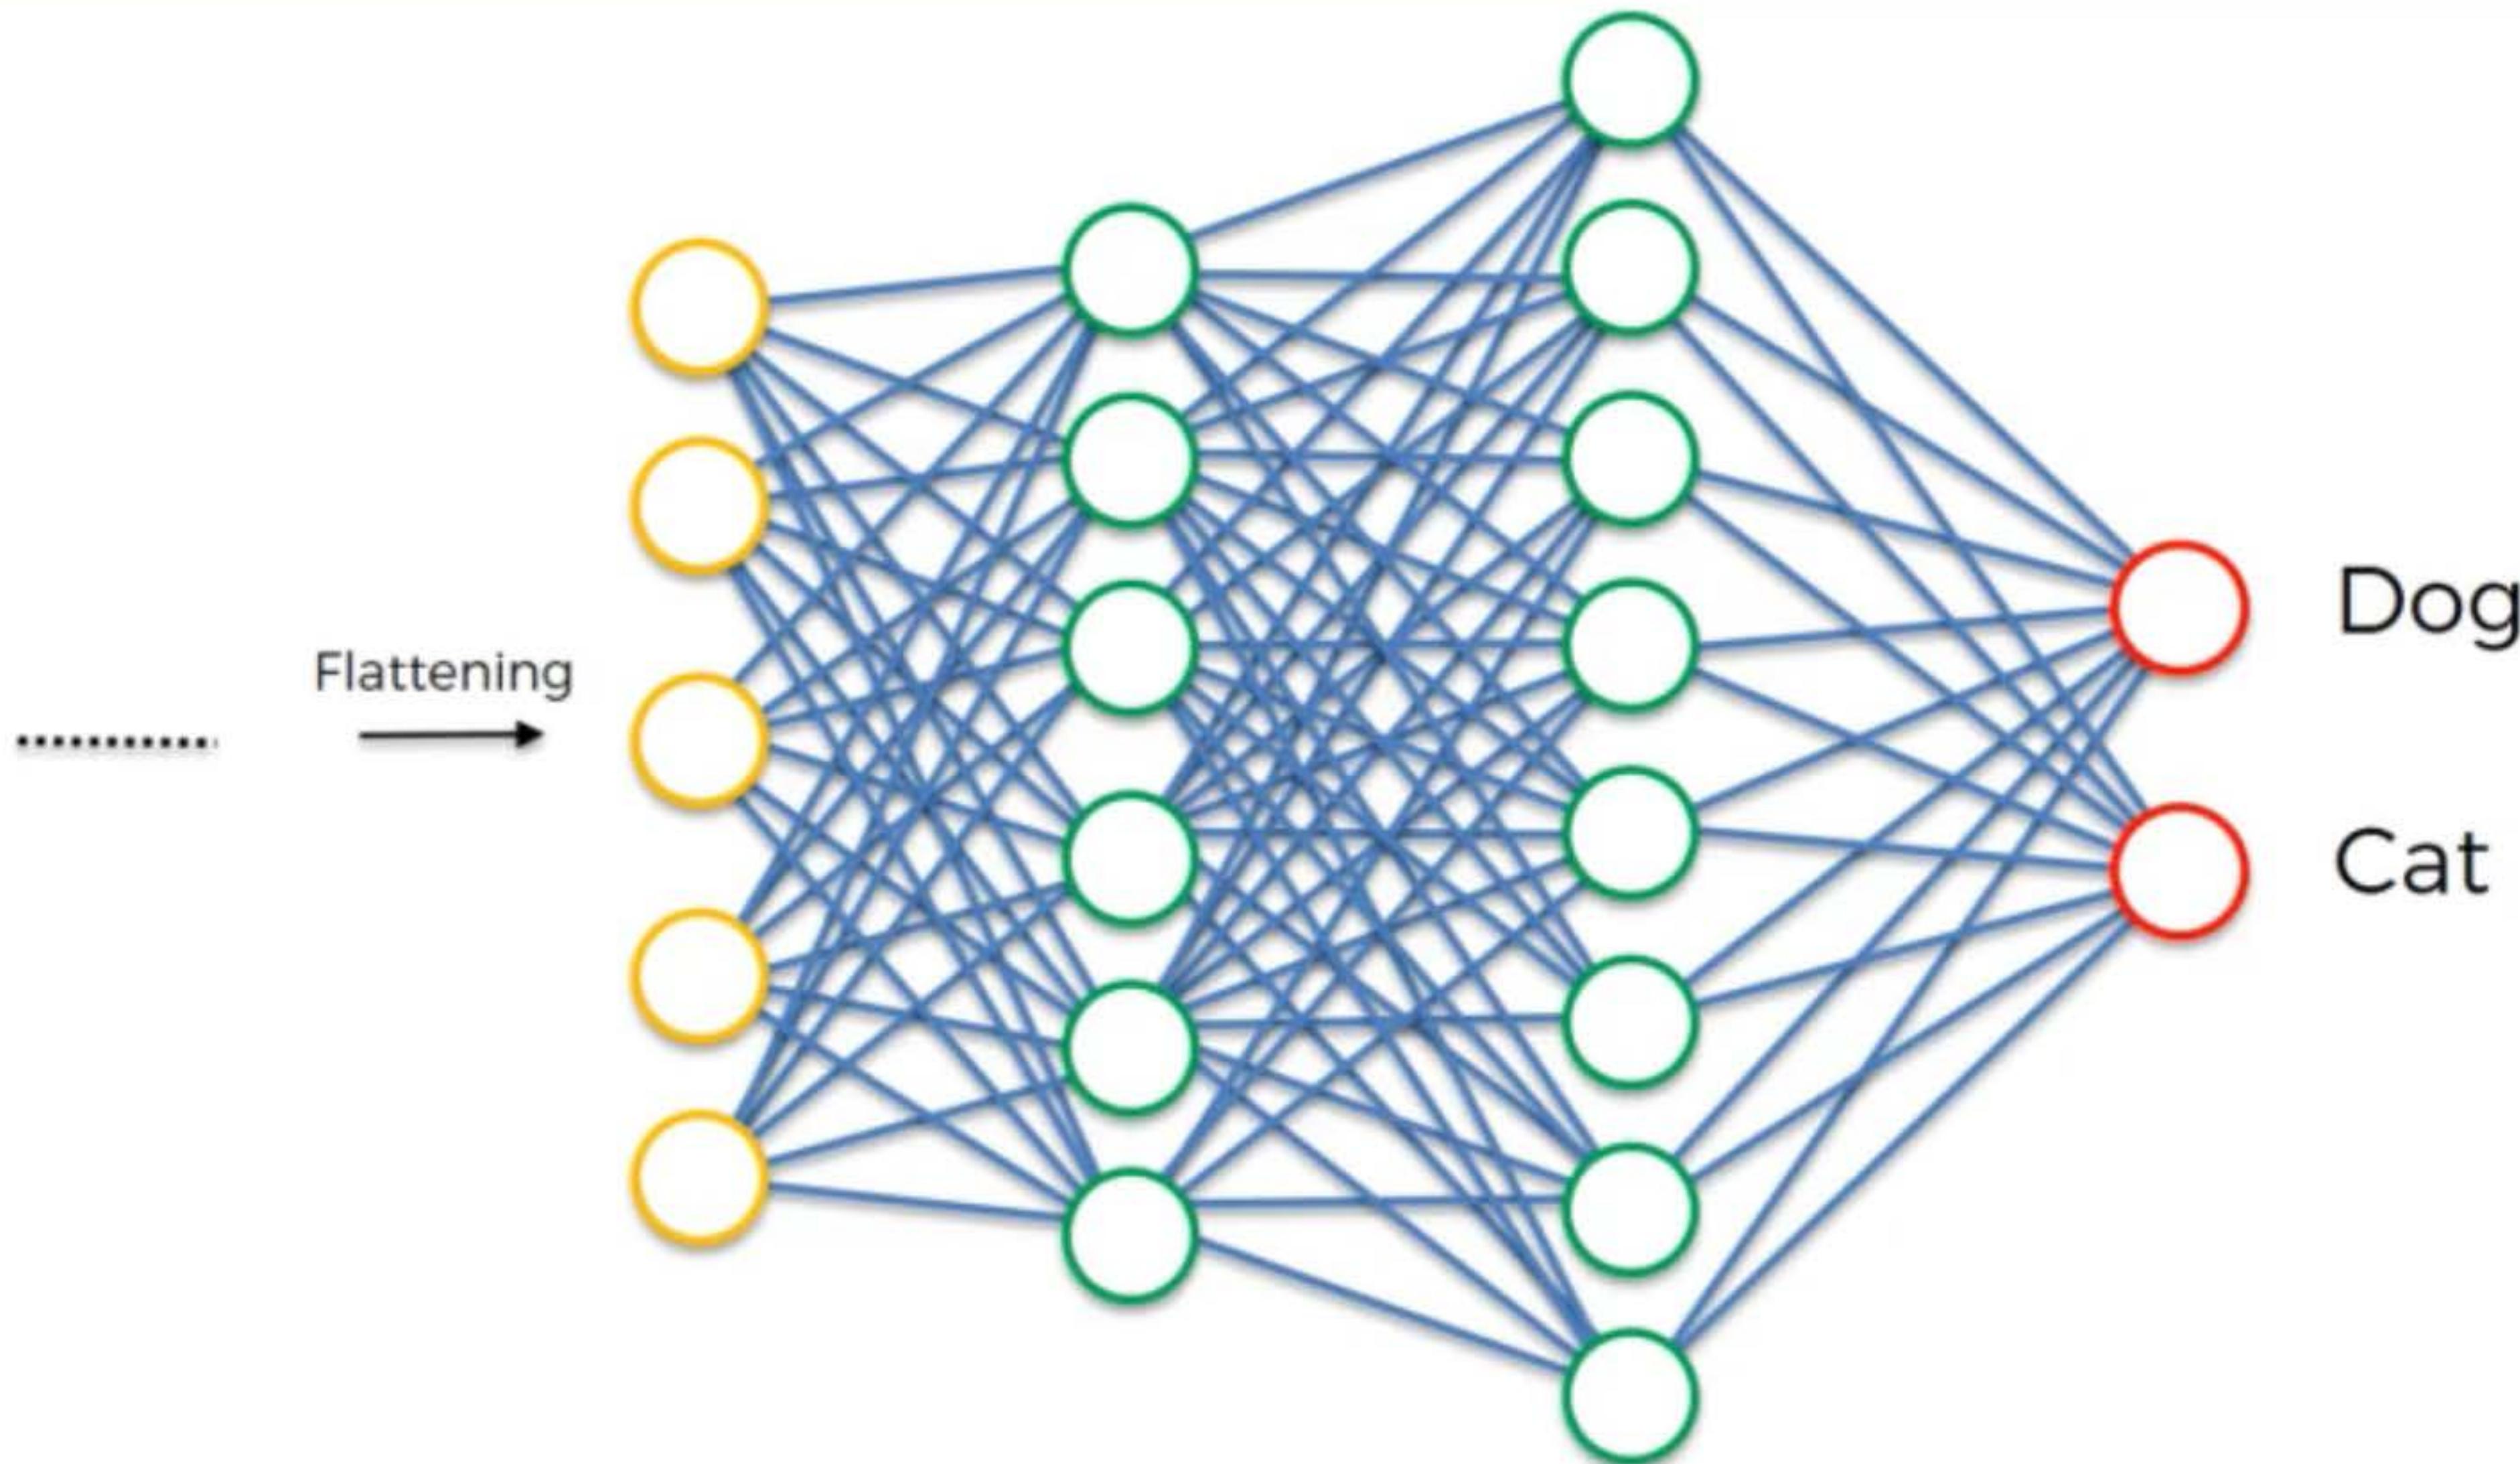


Step 4 - Full Connection

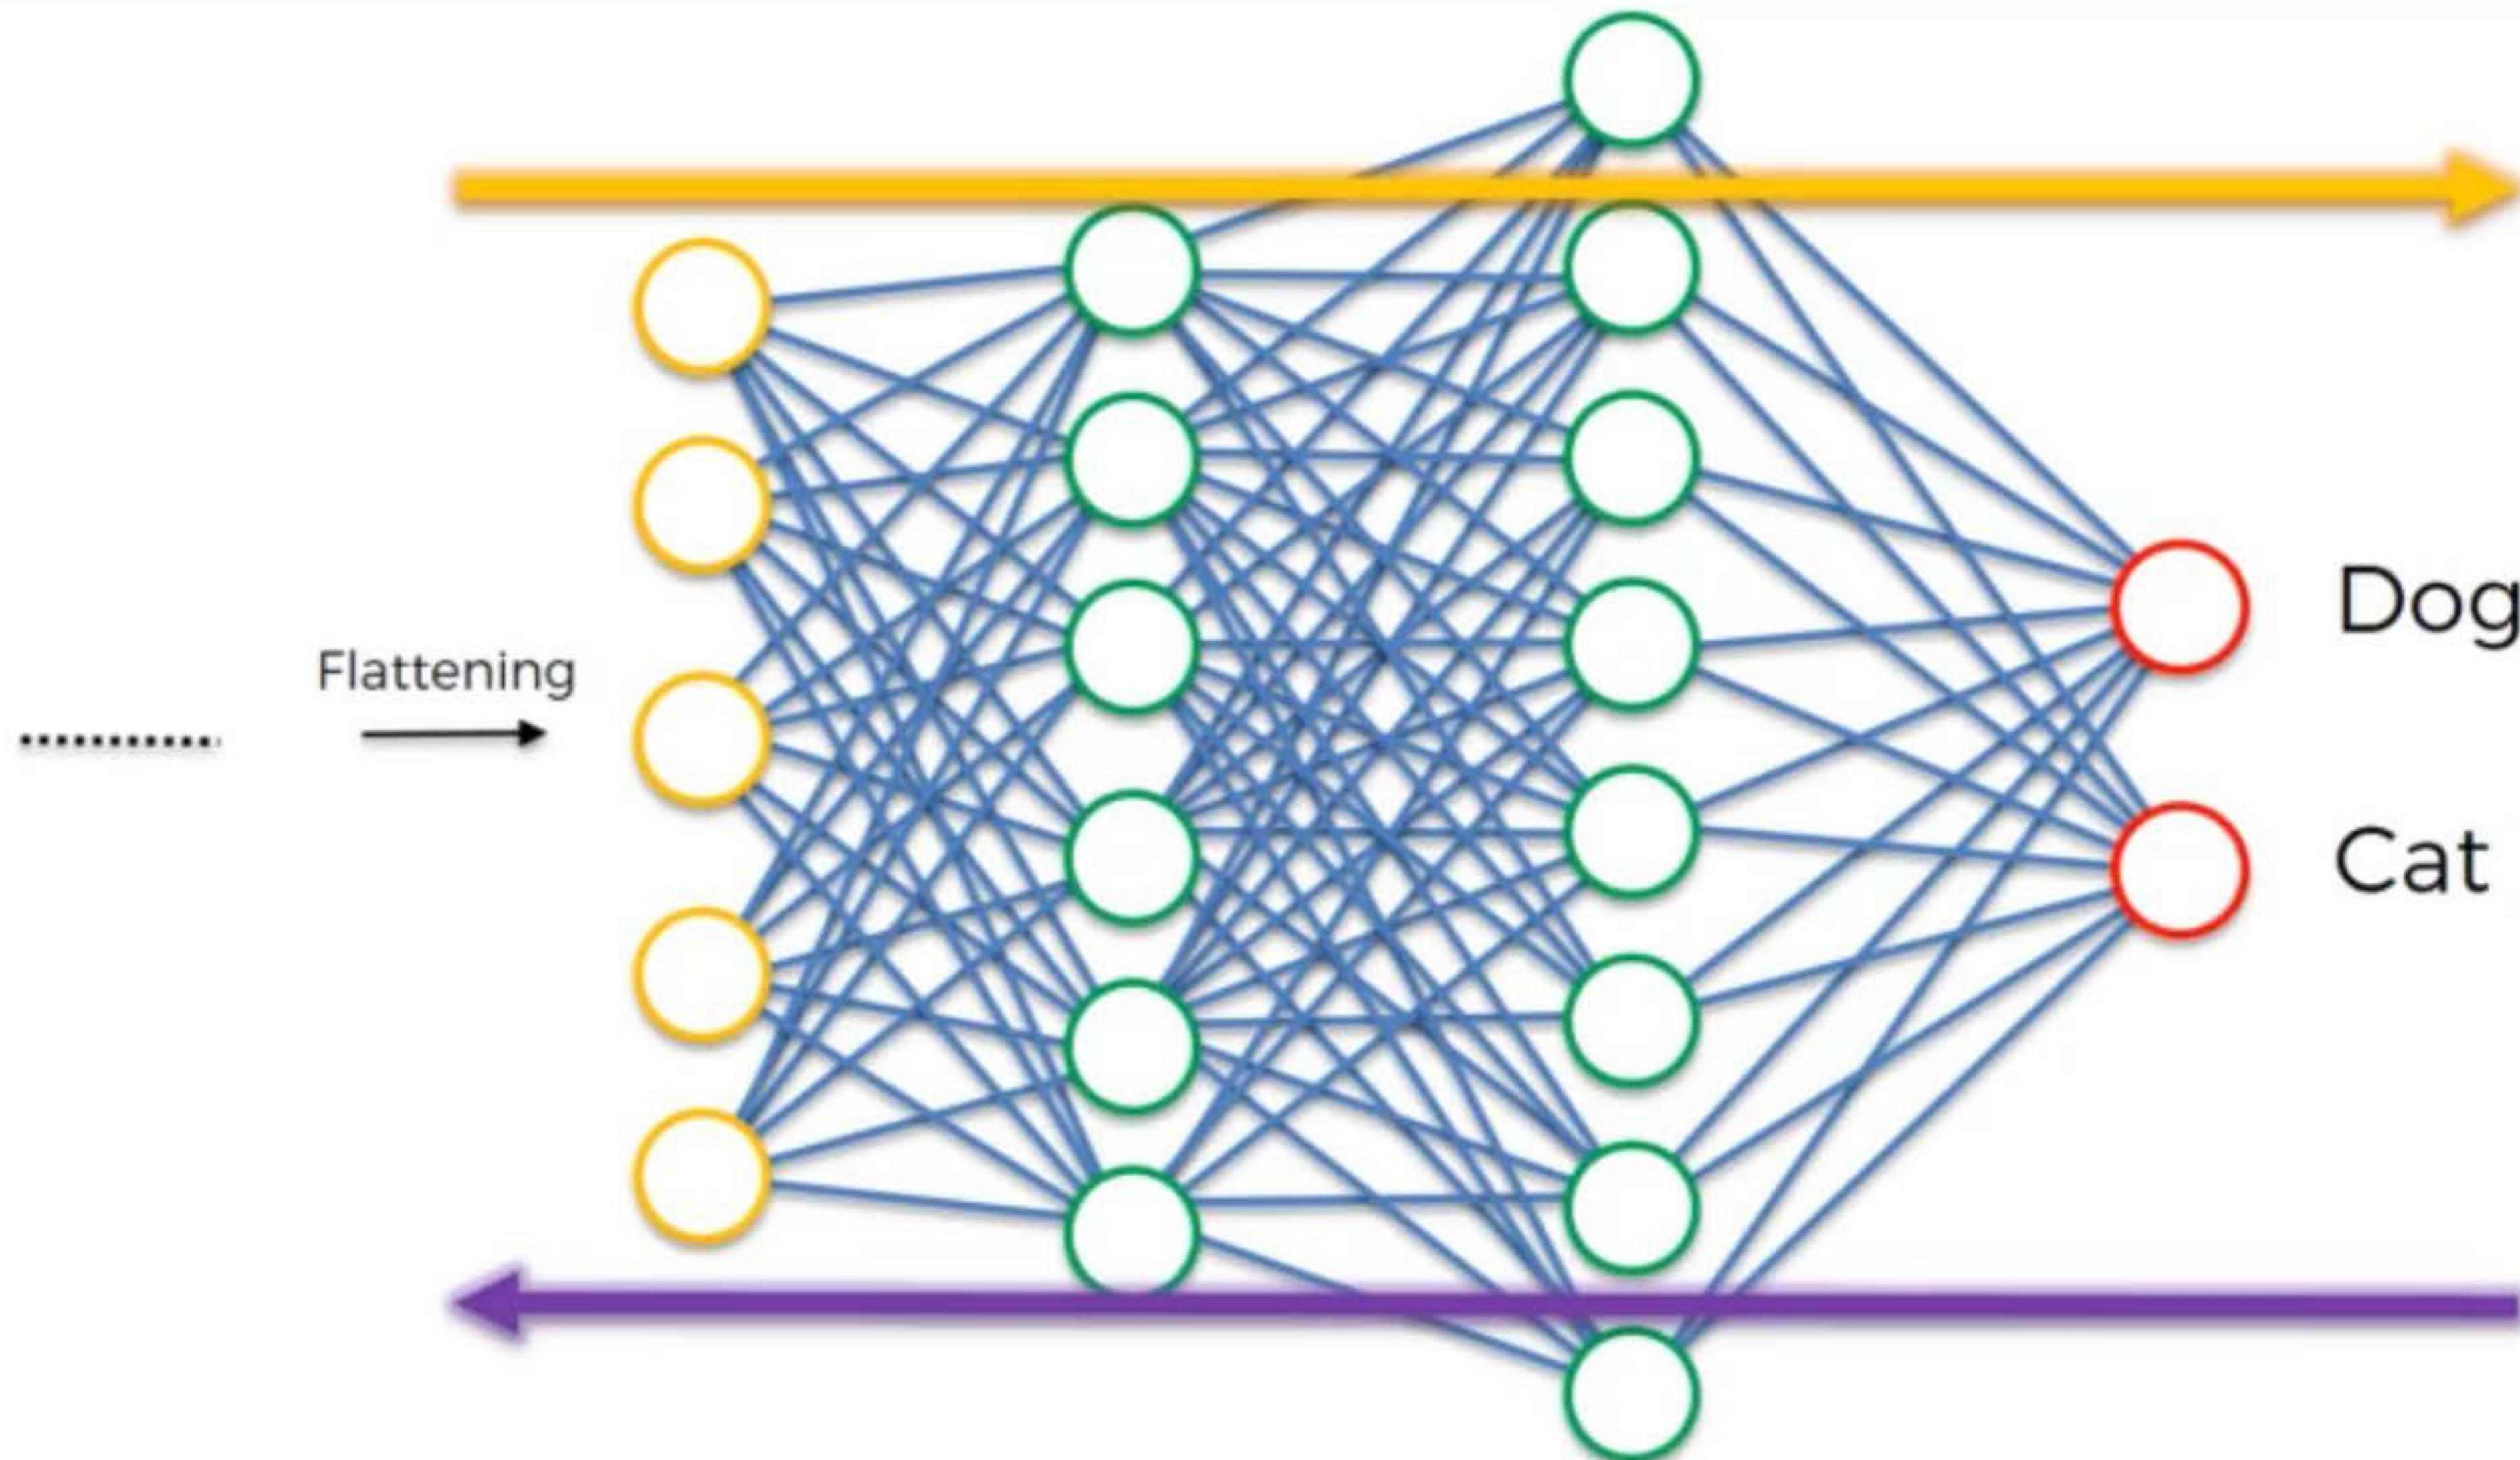
Step 4 - Full Connection



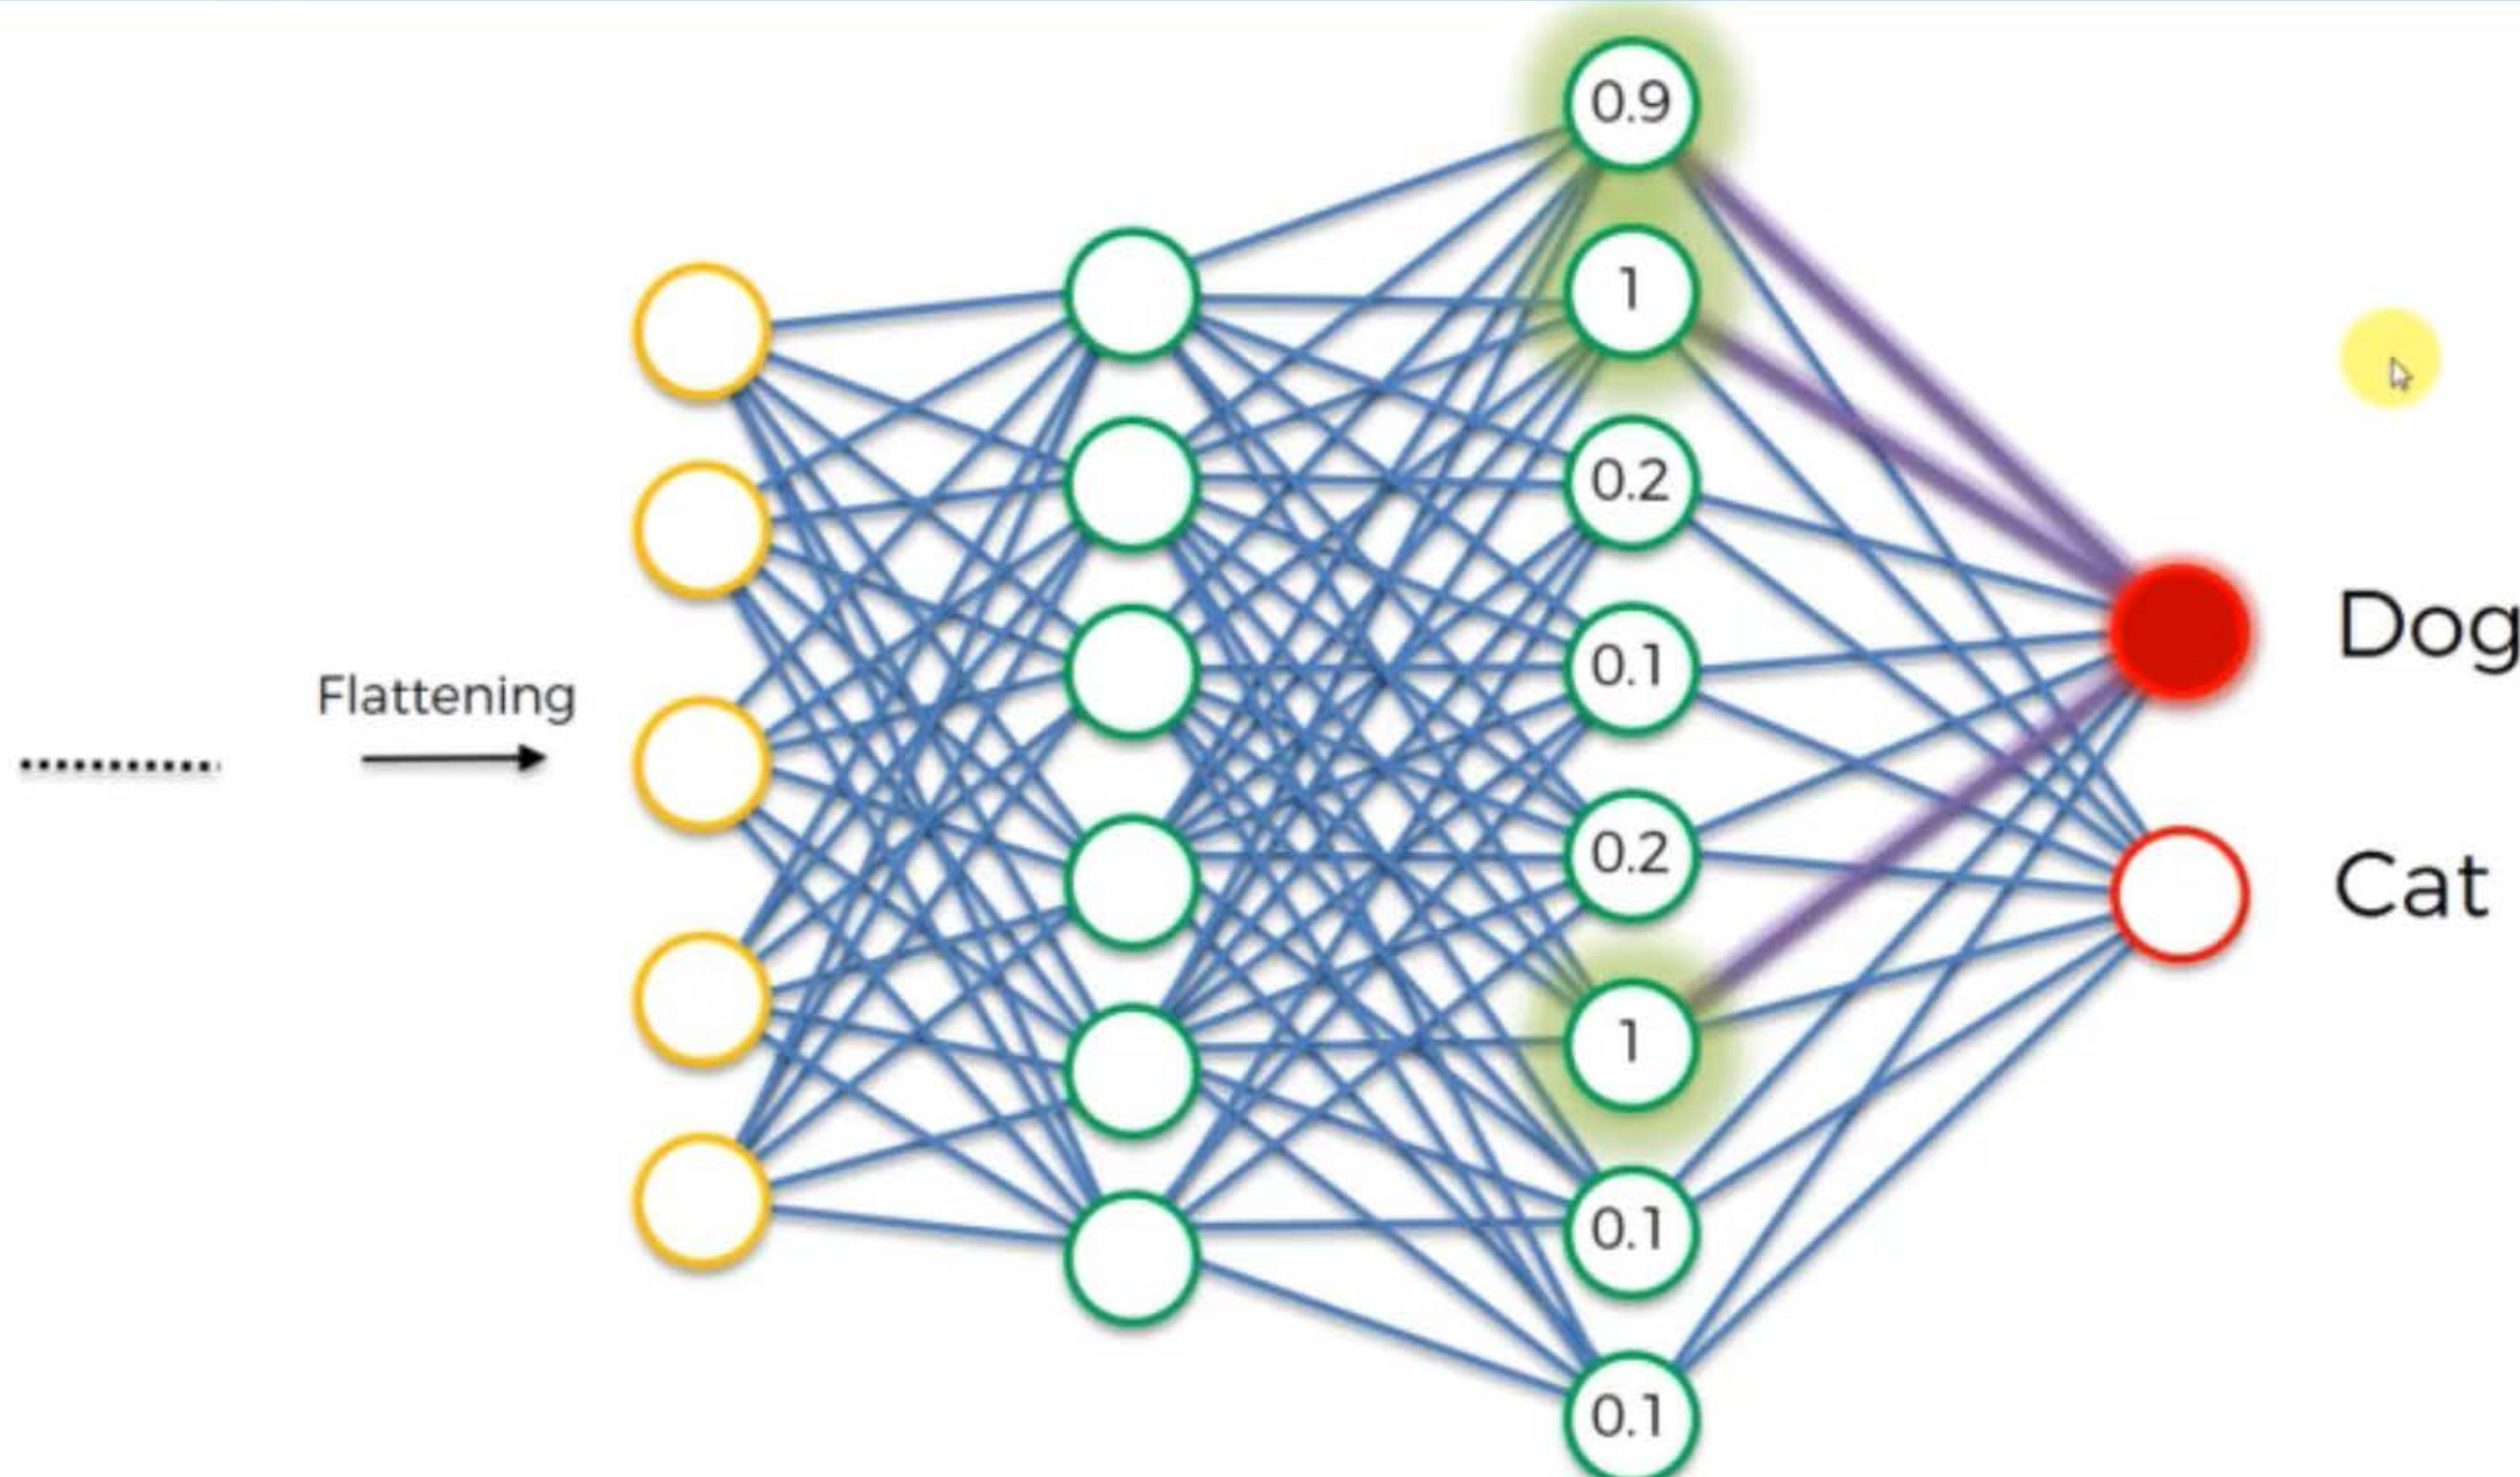
Step 4 - Full Connection



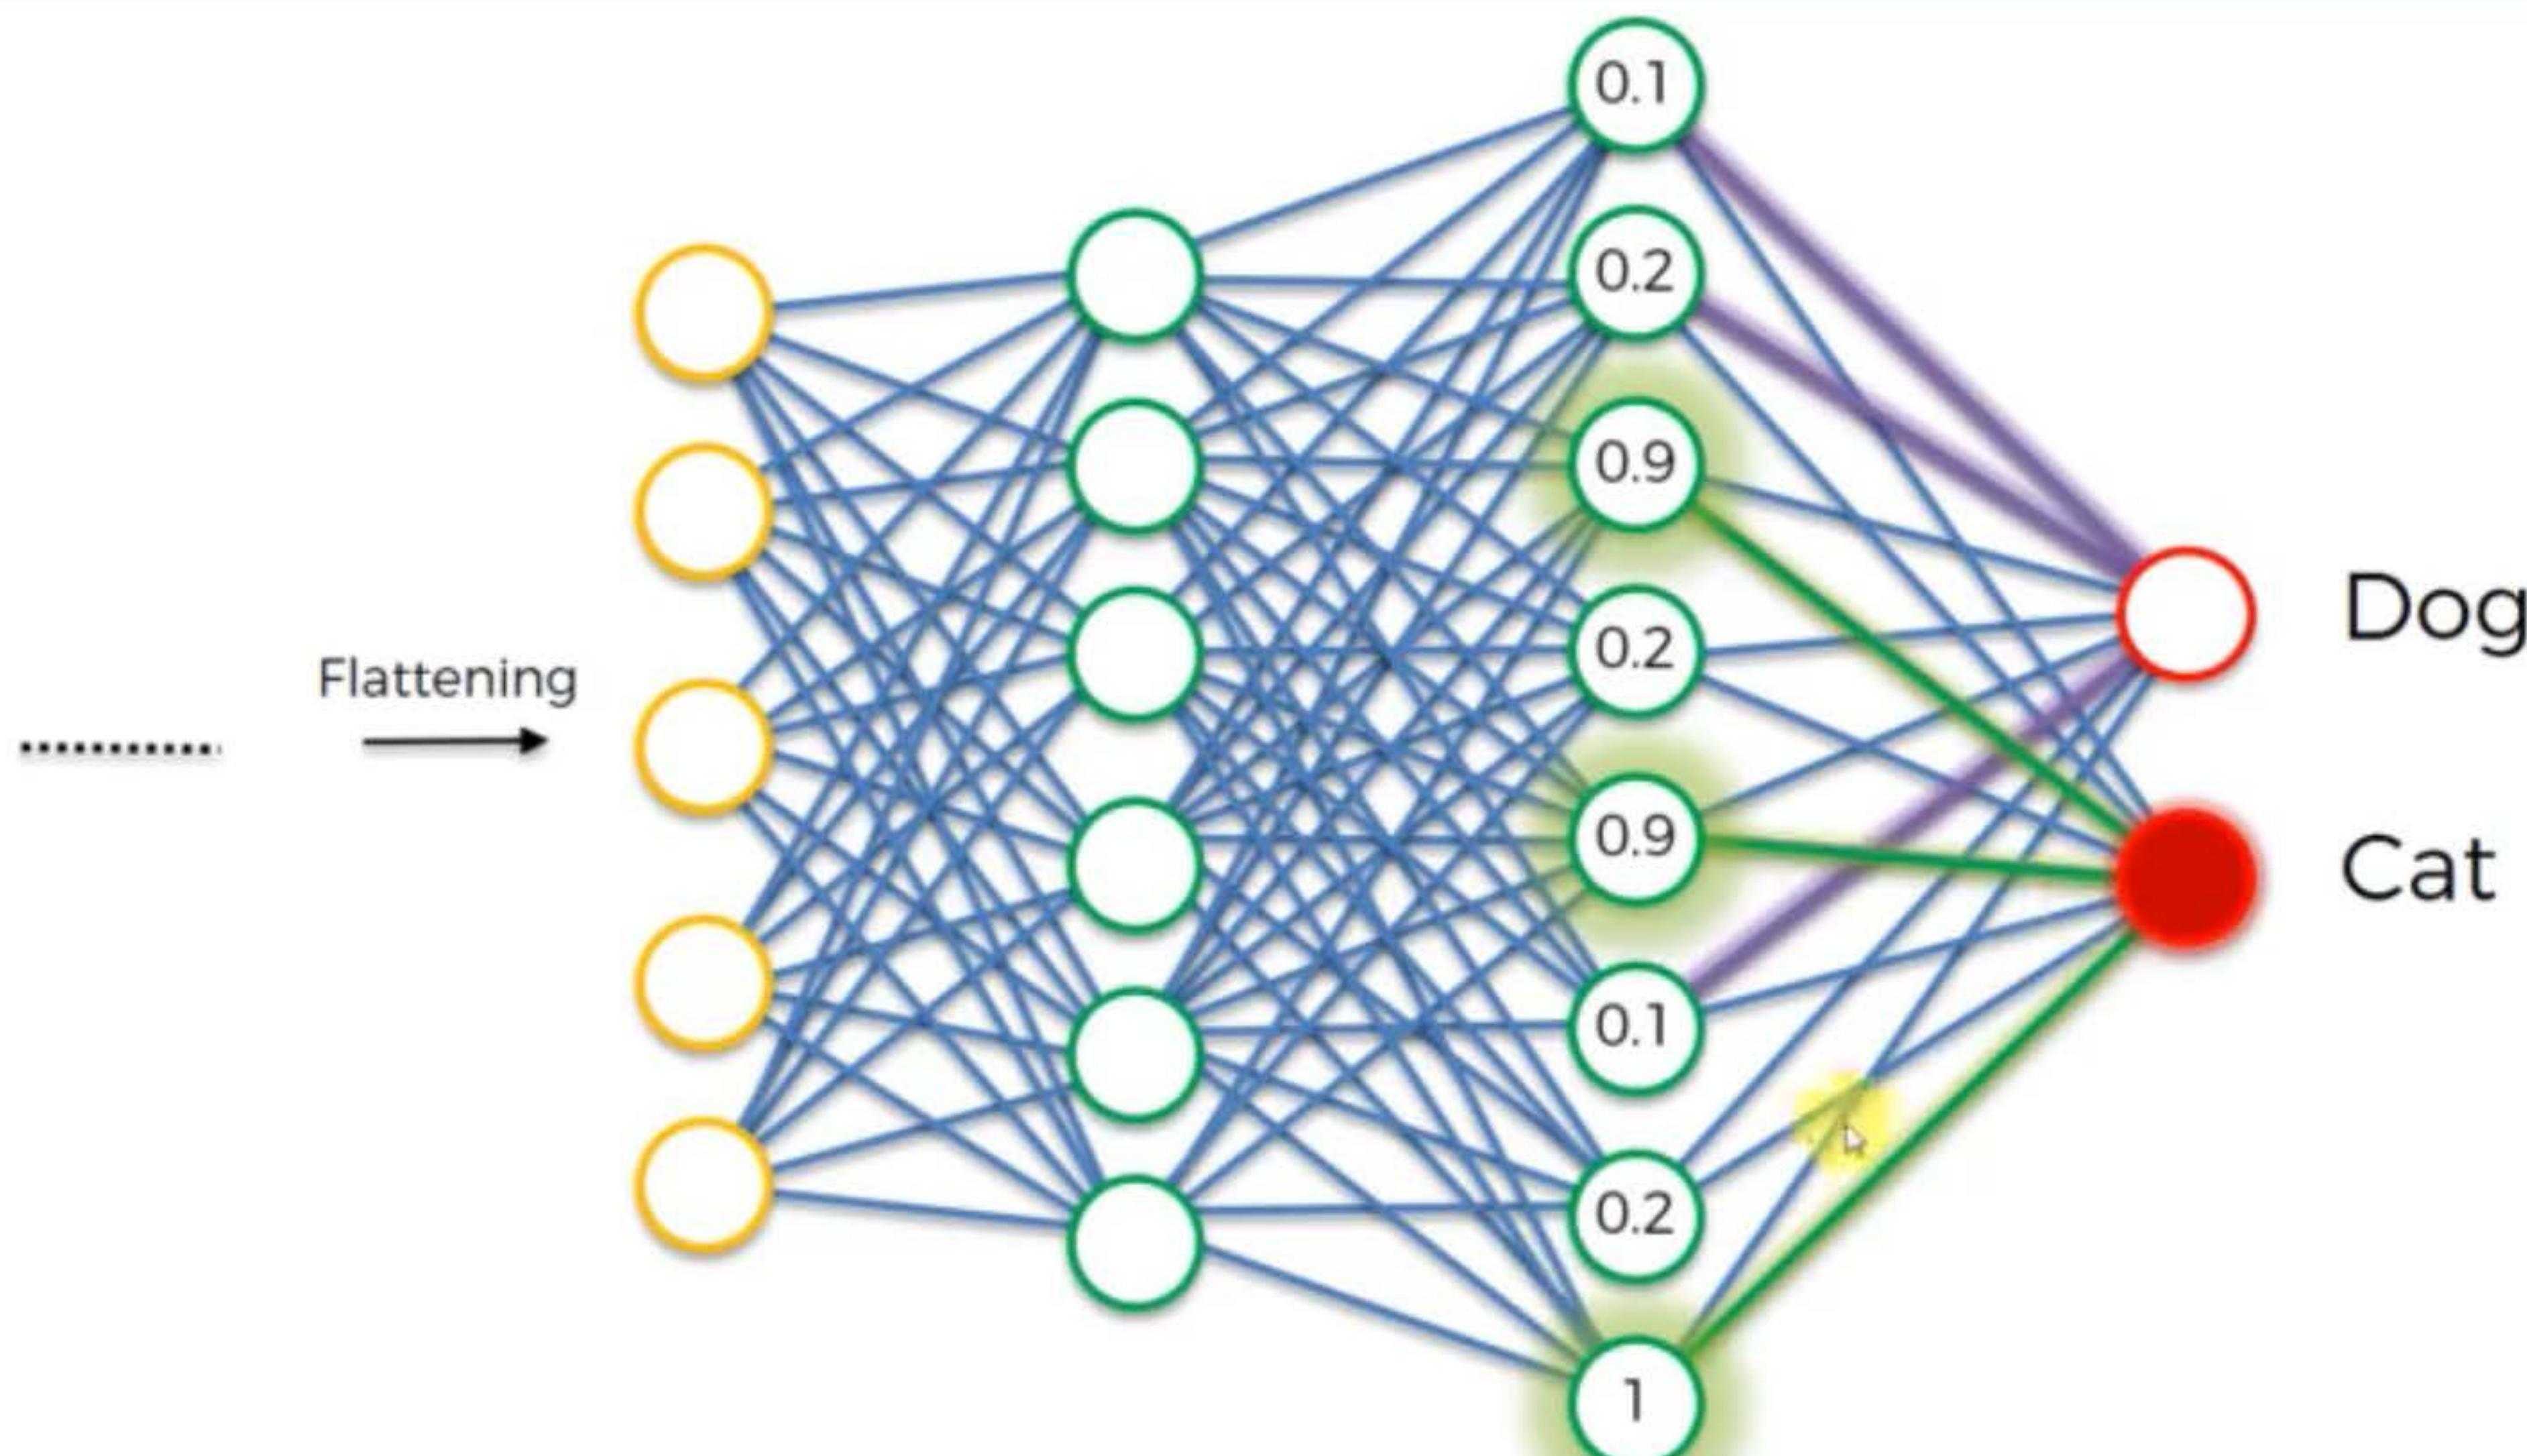
Step 4 - Full Connection



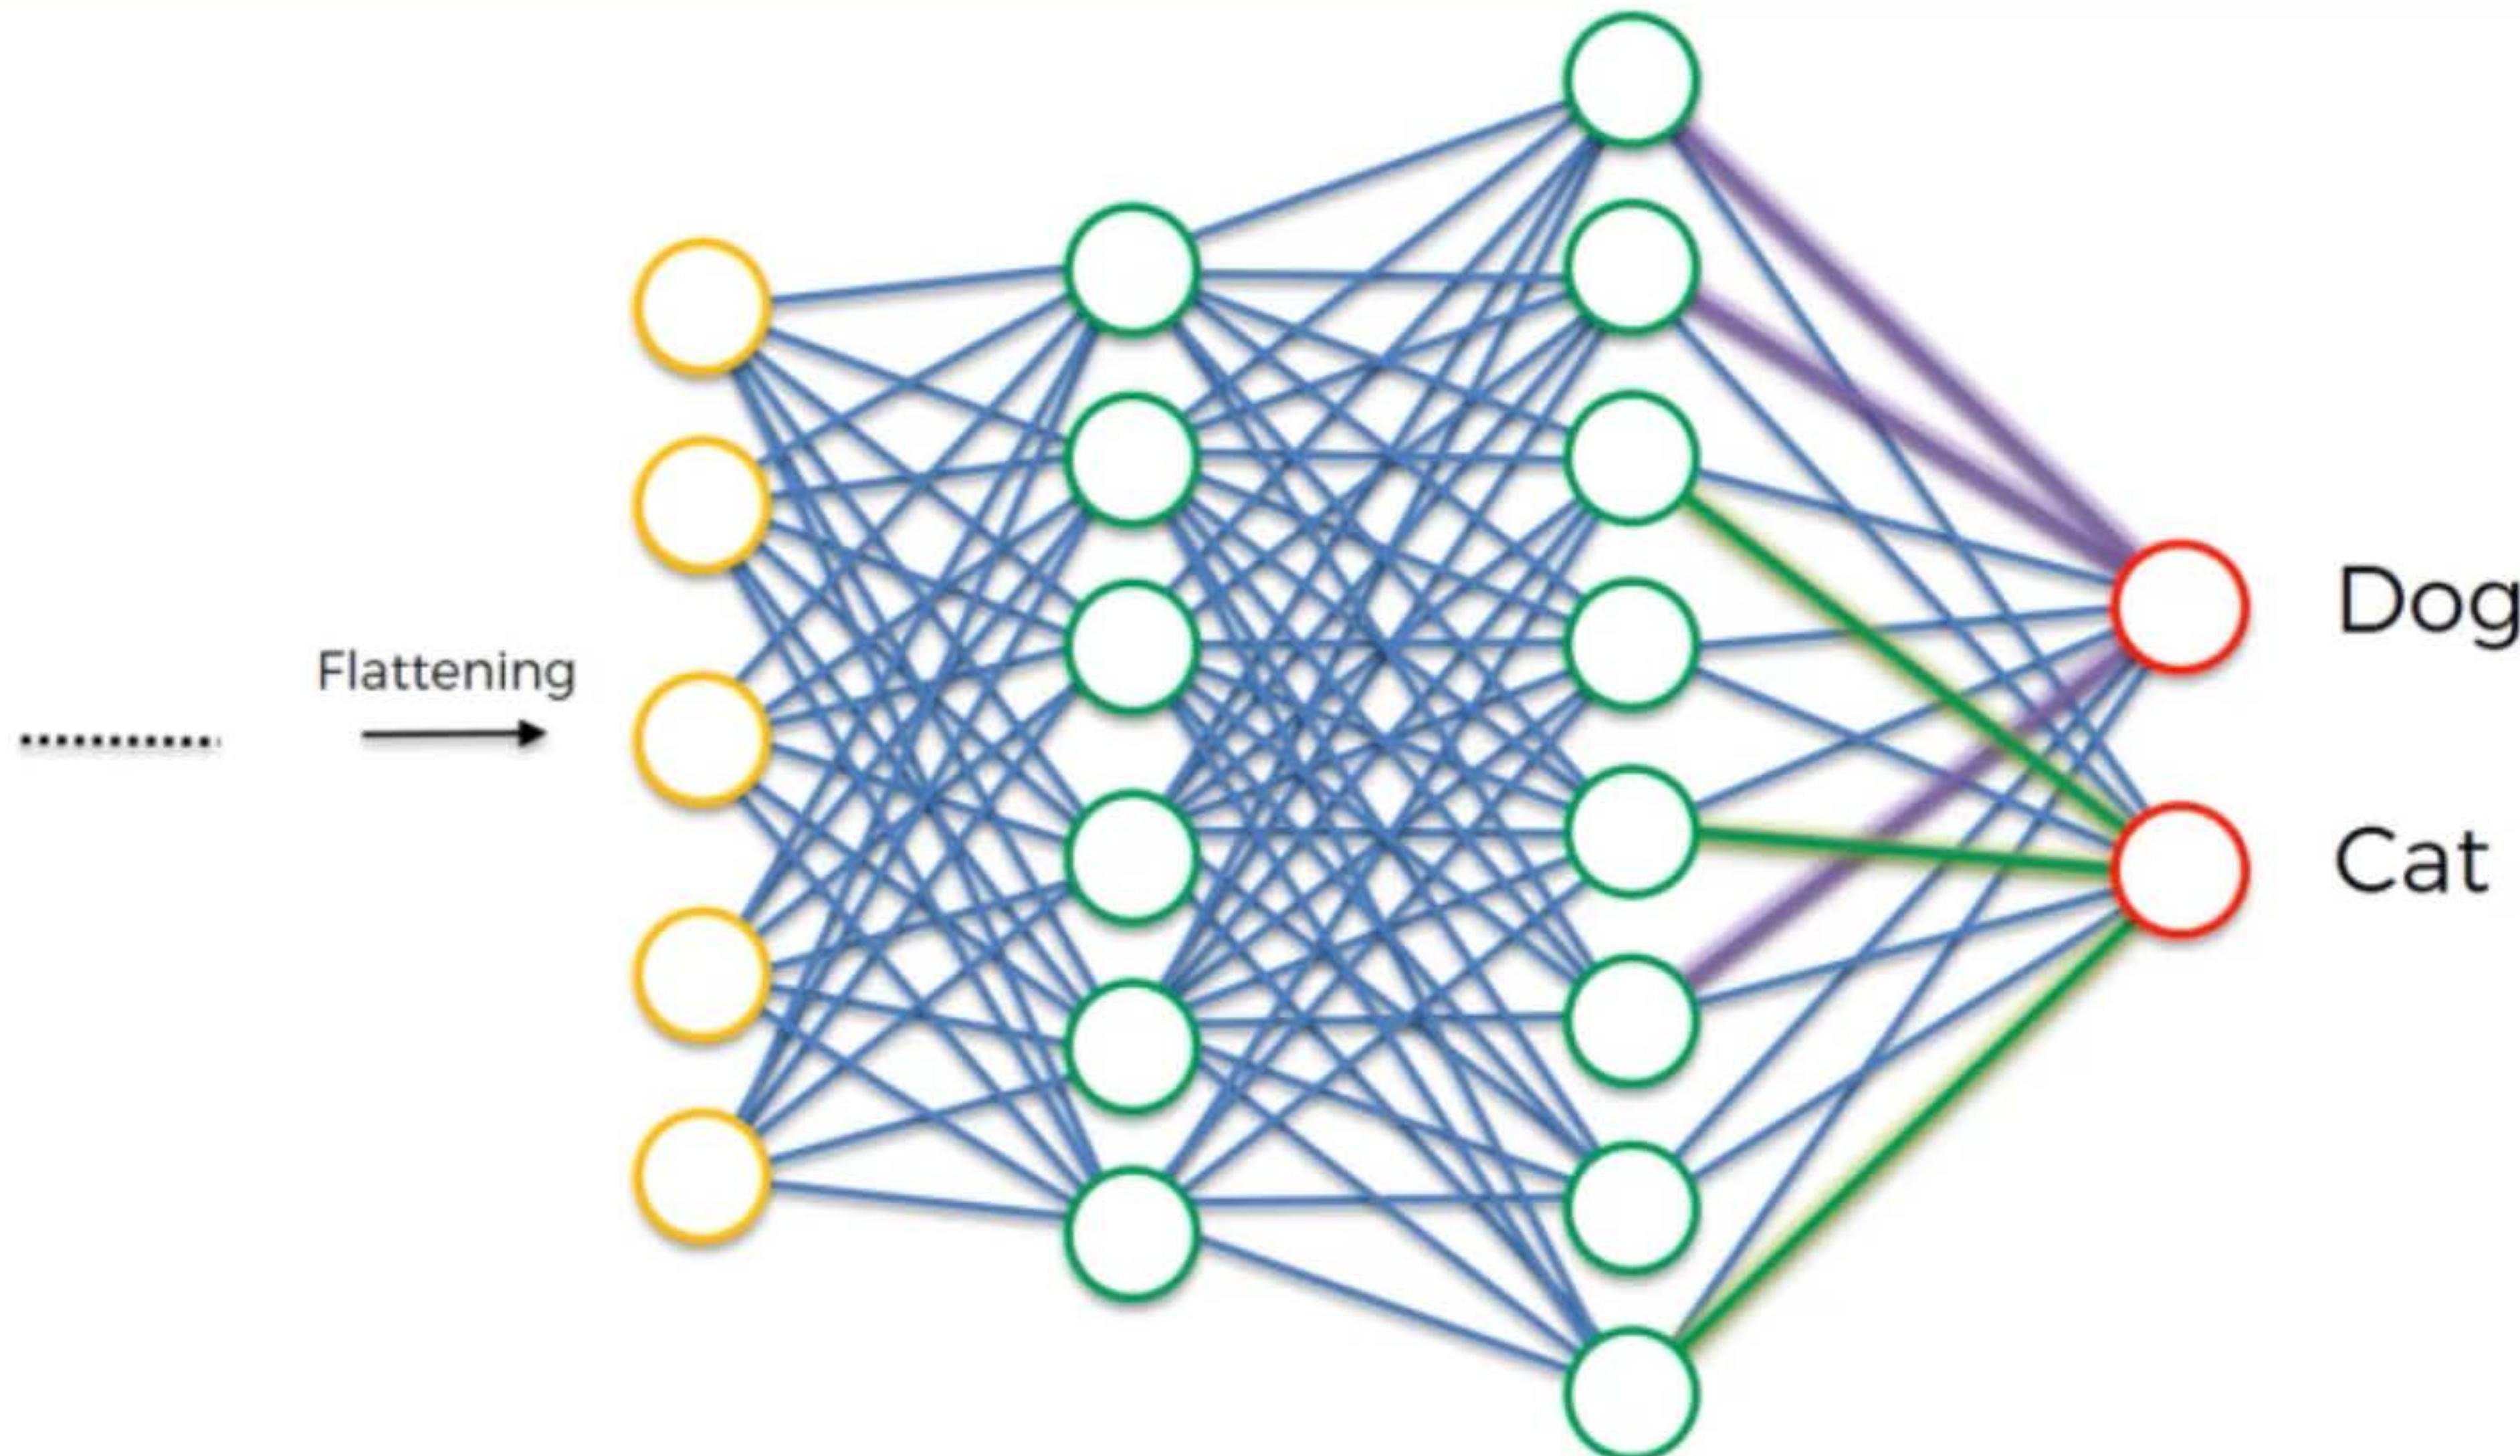
Step 4 - Full Connection



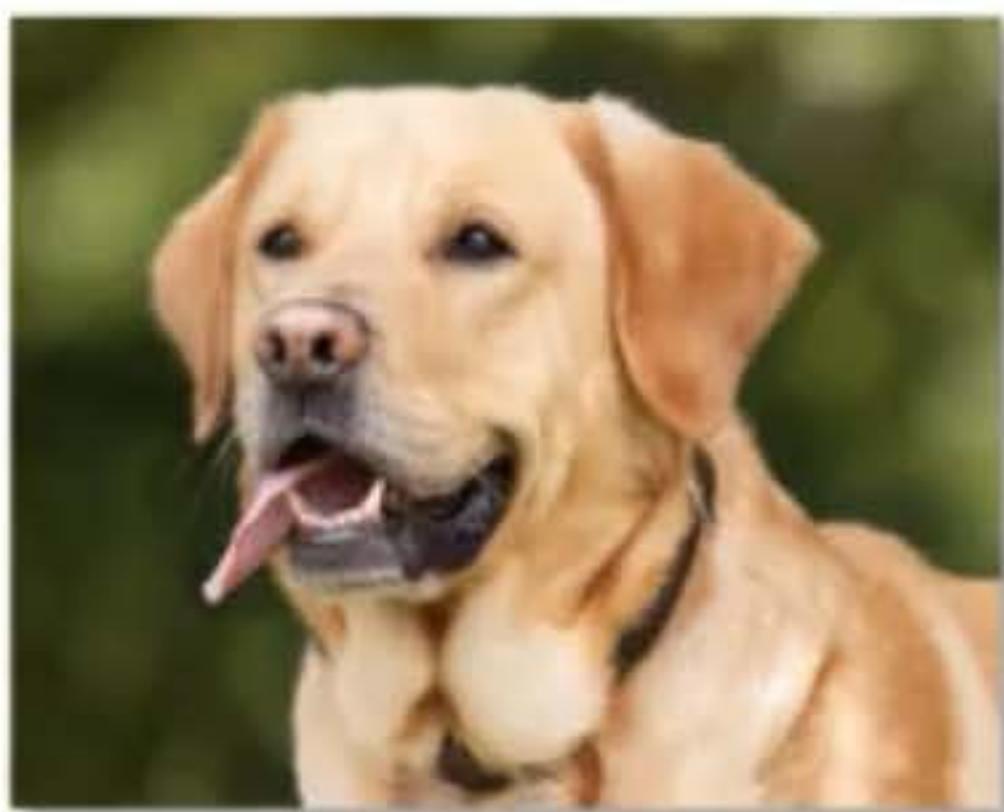
Step 4 - Full Connection



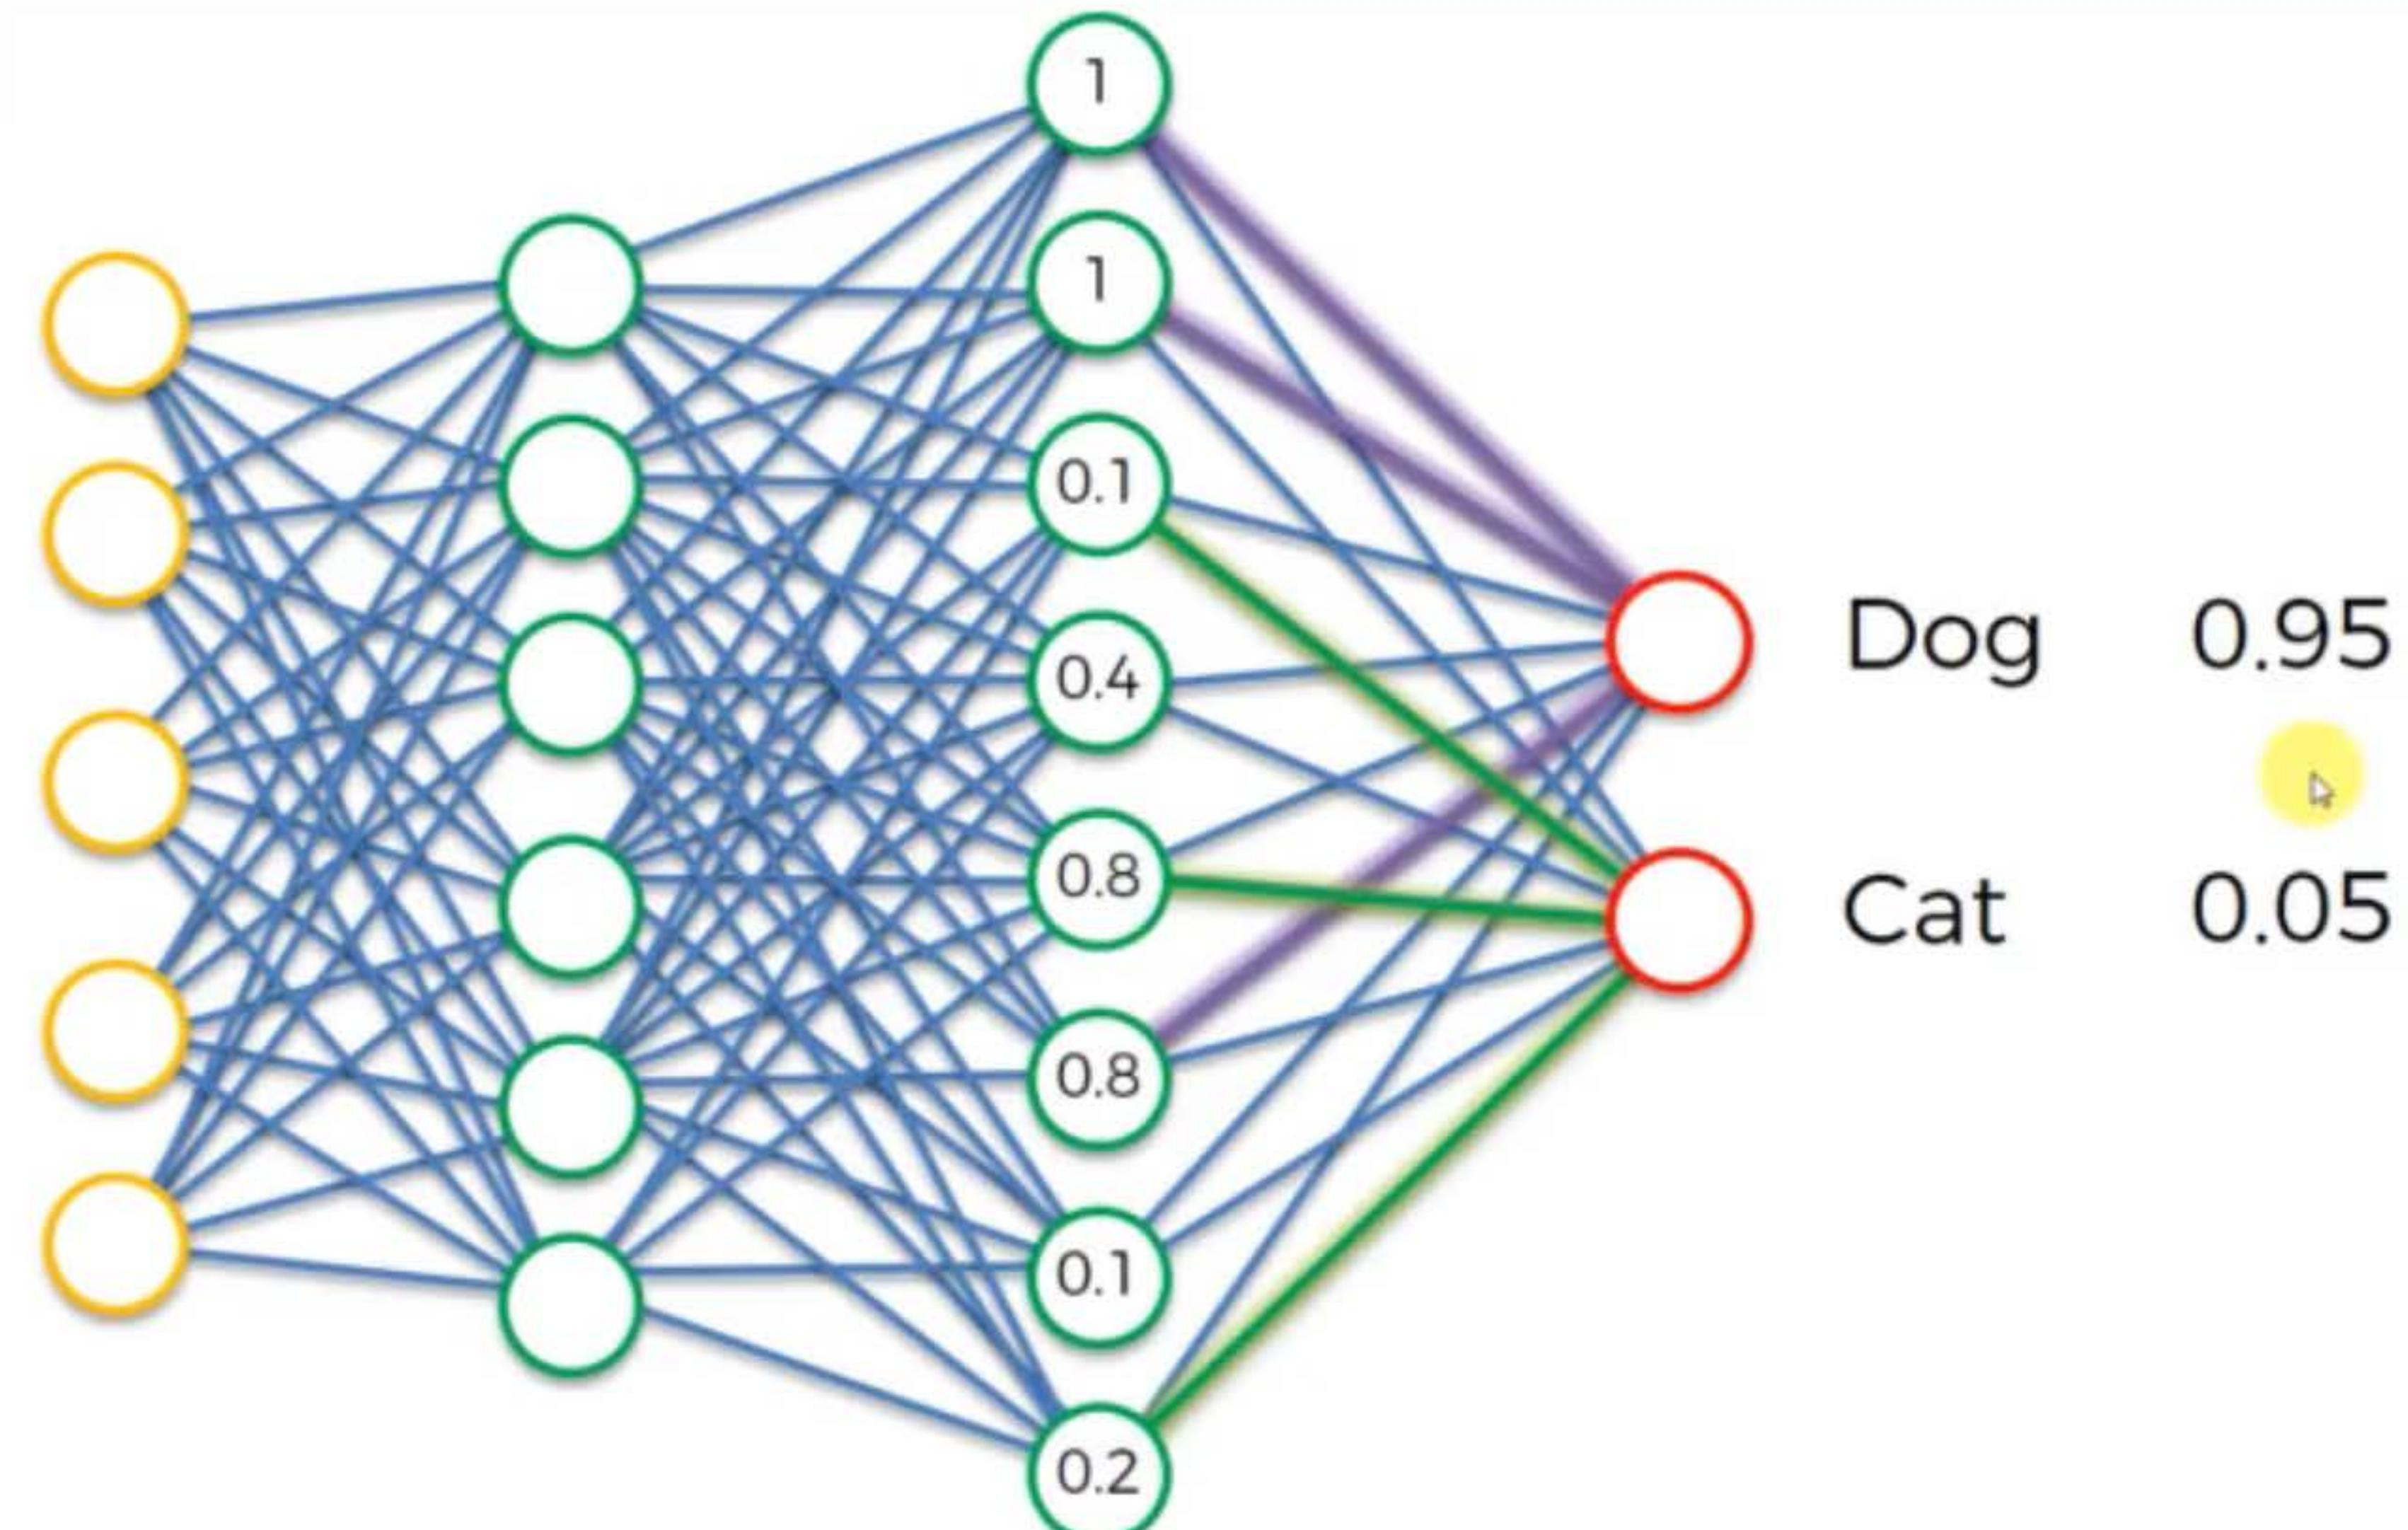
Step 4 - Full Connection



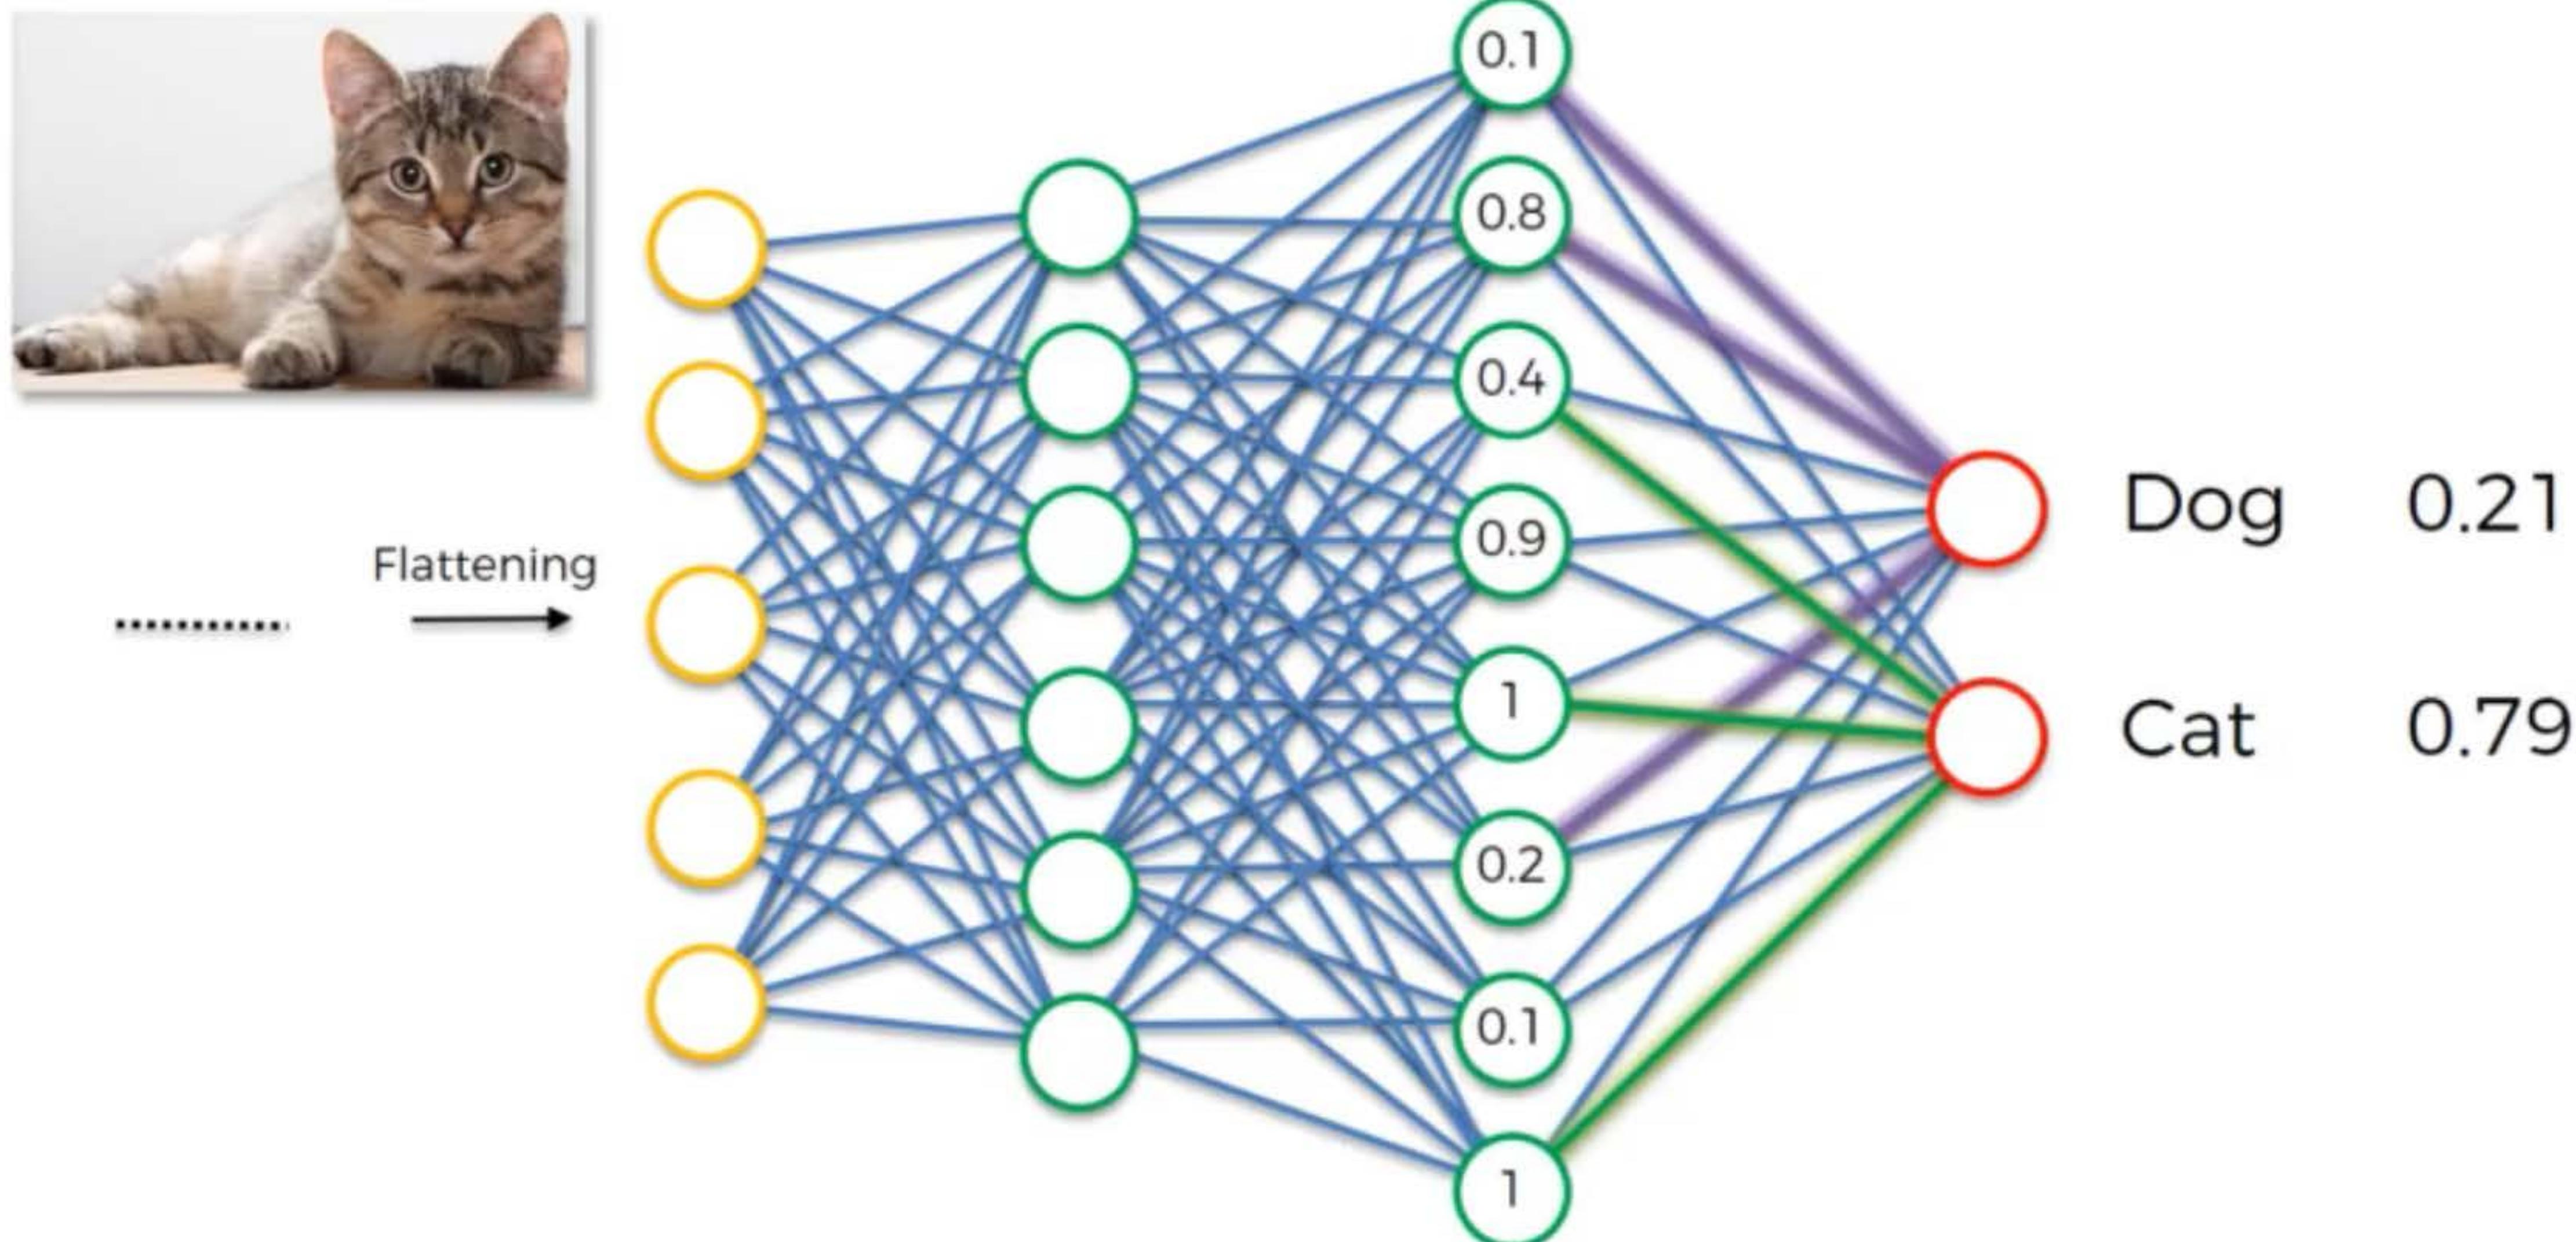
Step 4 - Full Connection



Flattening
..... →

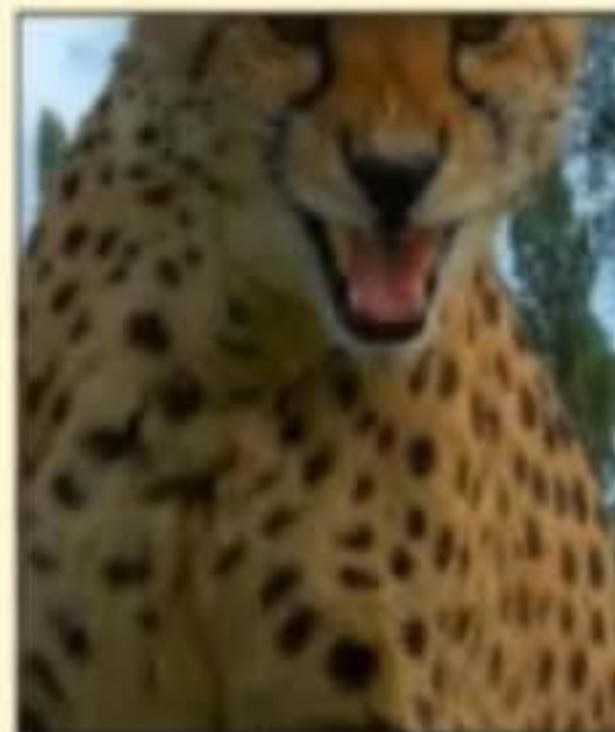


Step 4 - Full Connection



Step 4 - Full Connection

Examples from the test set
(with the network's guesses)



cheetah

cheetah
leopard
snow leopard
Egyptian cat



bullet train

bullet train
passenger car
subway train
electric locomotive

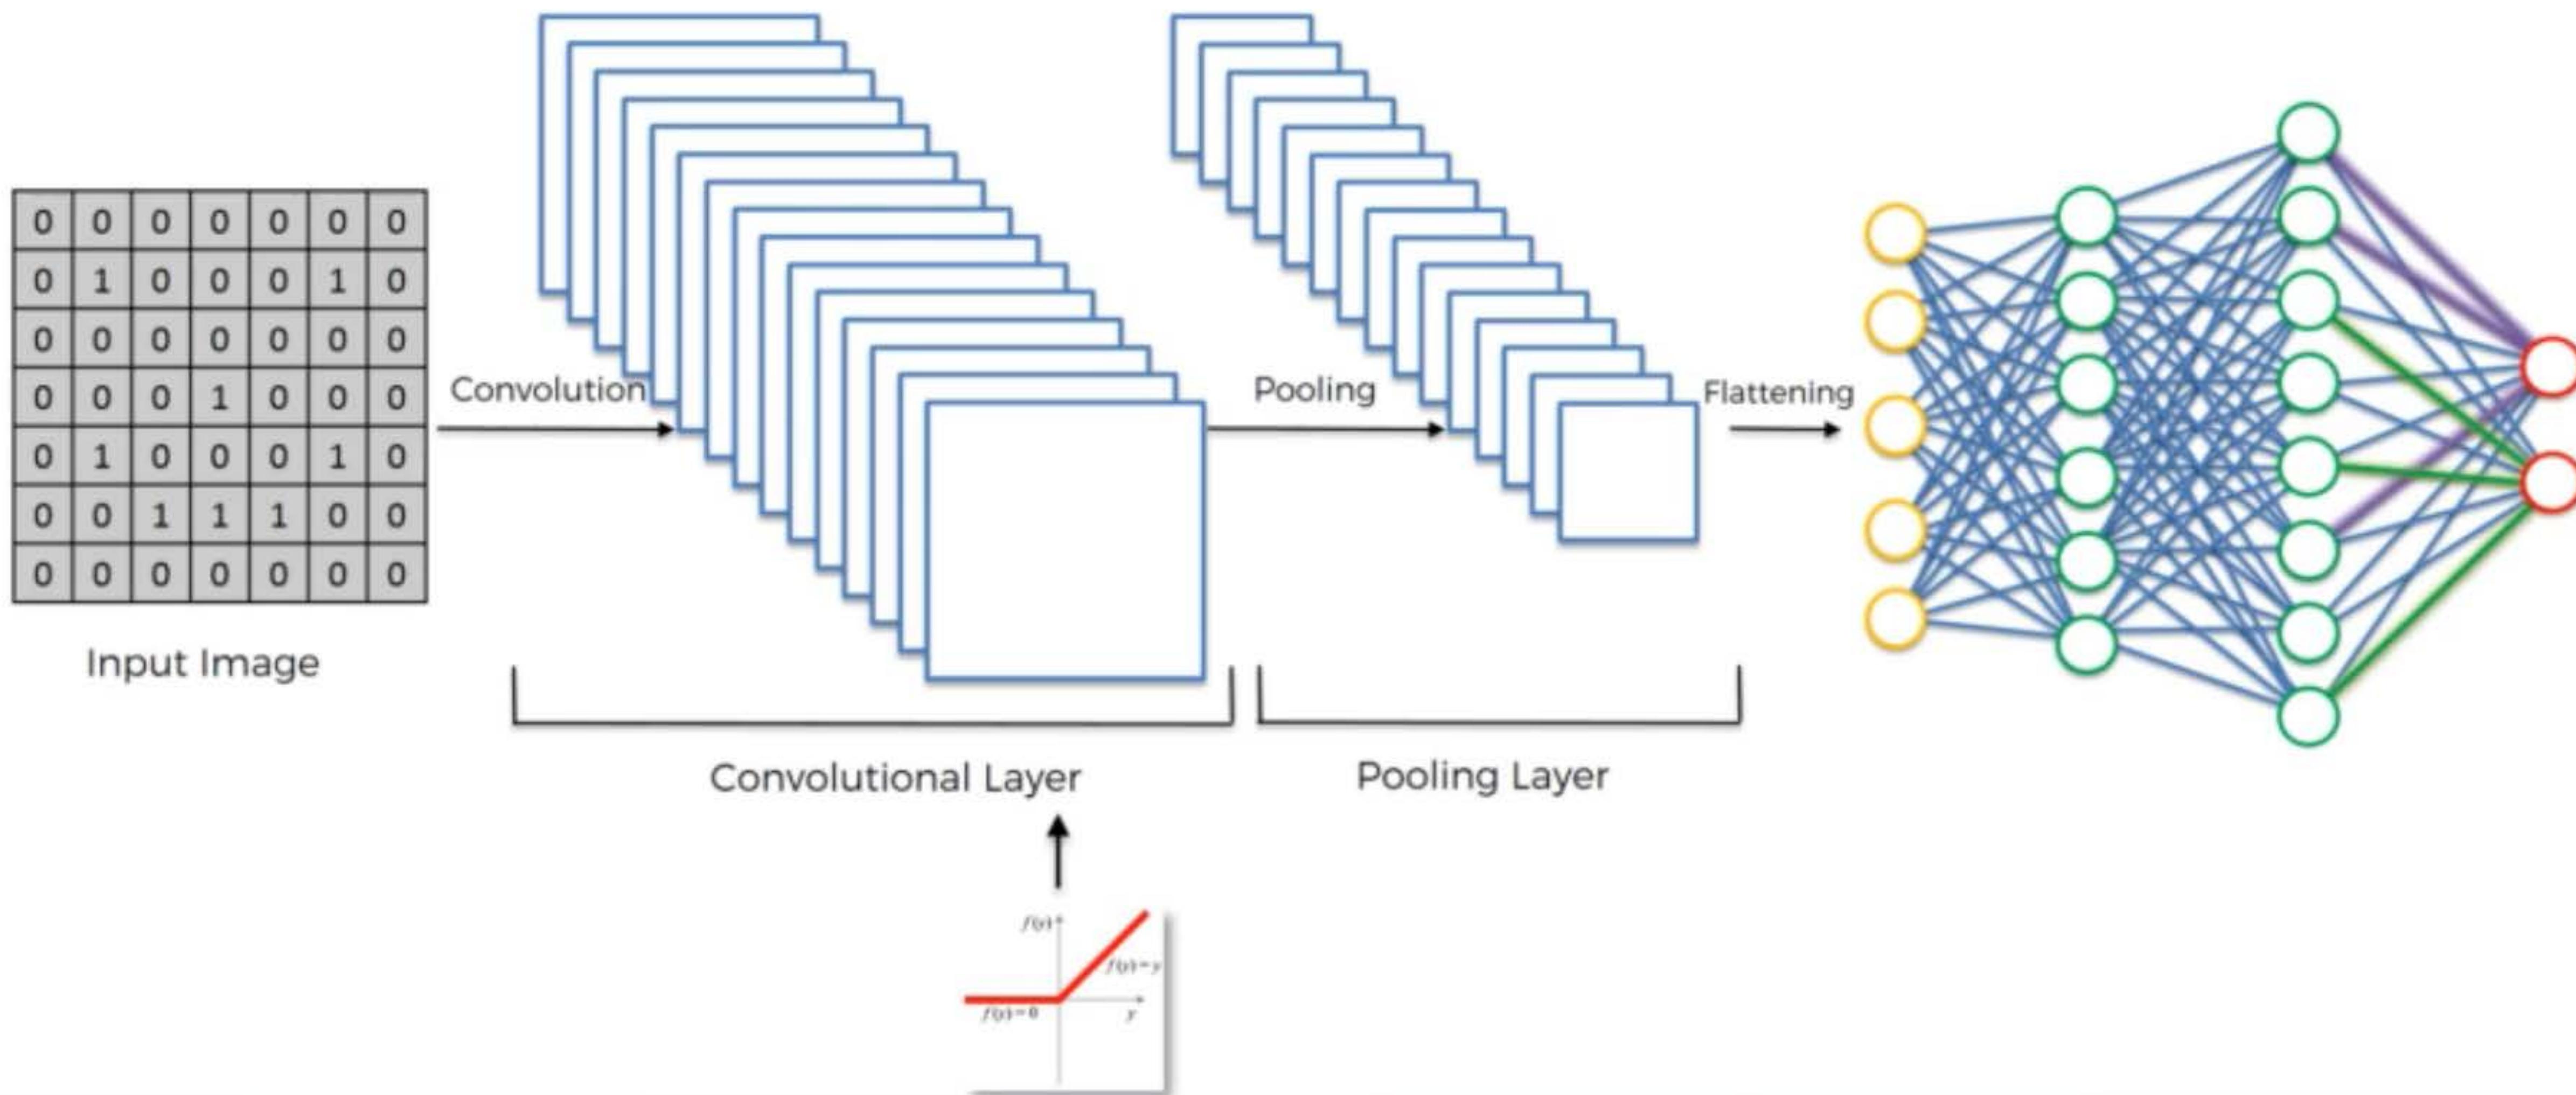


hand glass

scissors
hand glass
frying pan
stethoscope

Image Source: a talk by Geoffrey Hinton

Summary



Summary

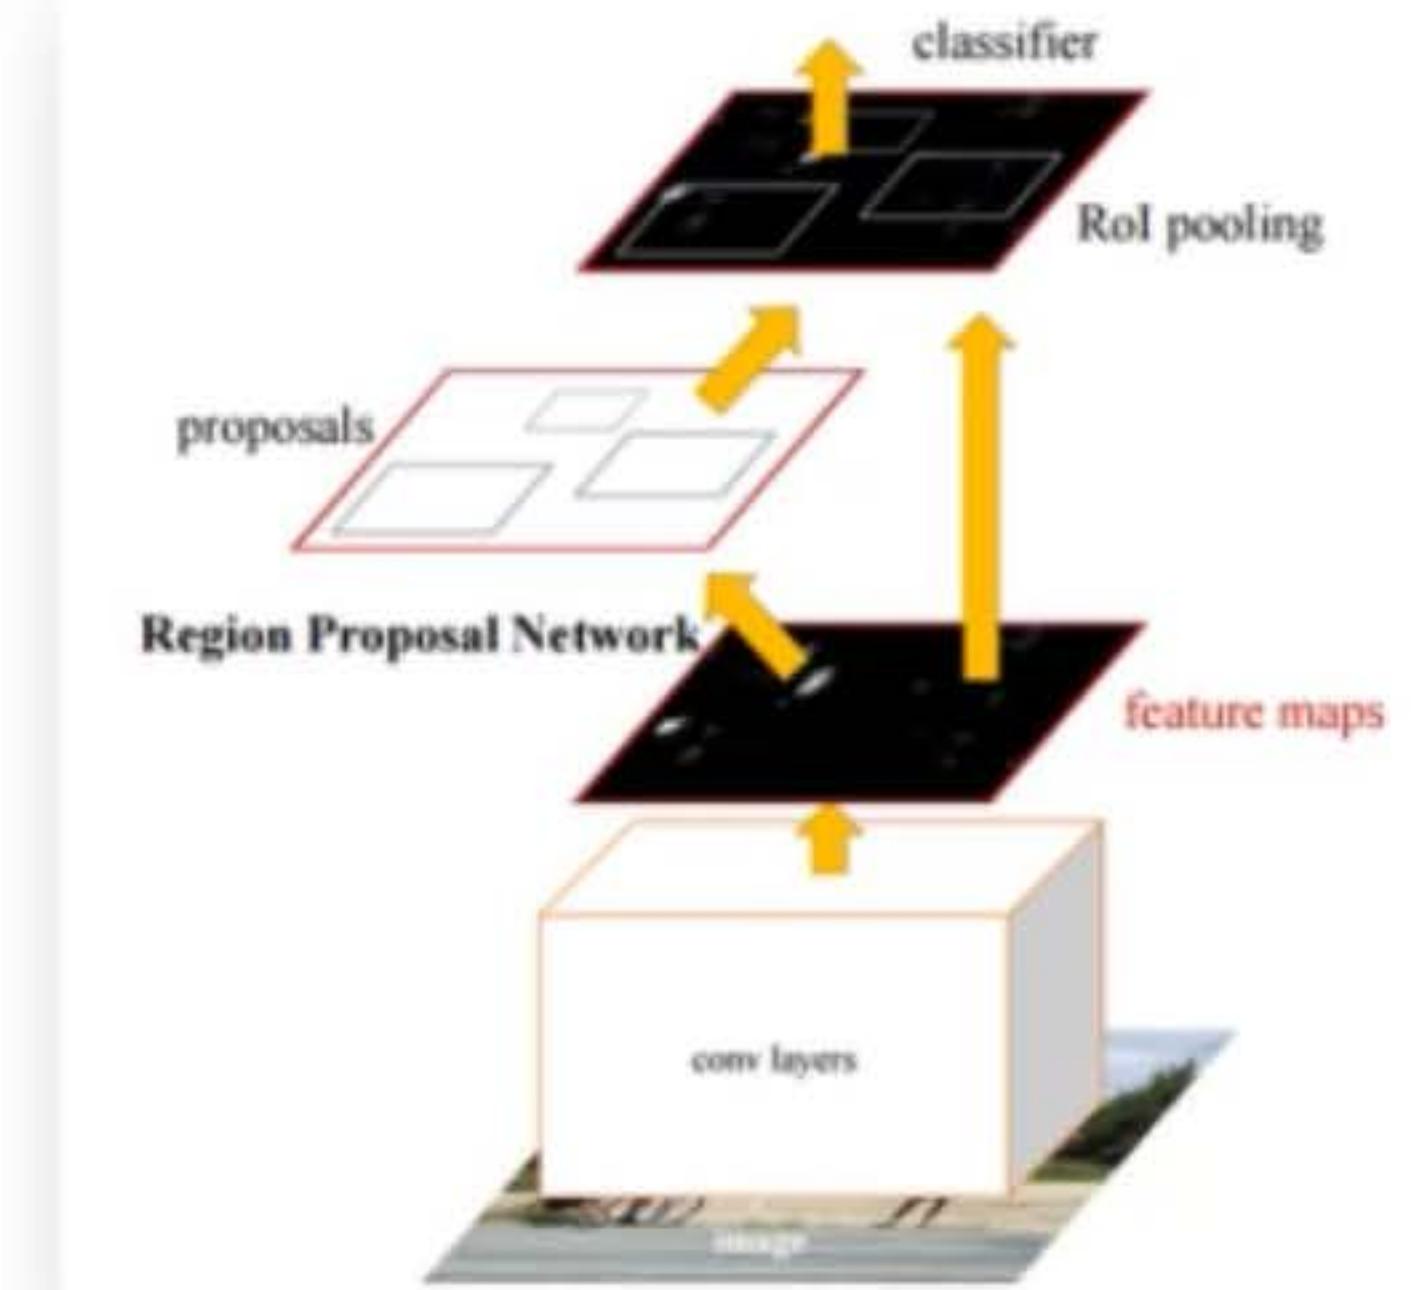
Additional Reading:

*The 9 Deep Learning Papers
You Need To Know About
(Understanding CNNs Part 3)*

Adit Deshpande (2016)

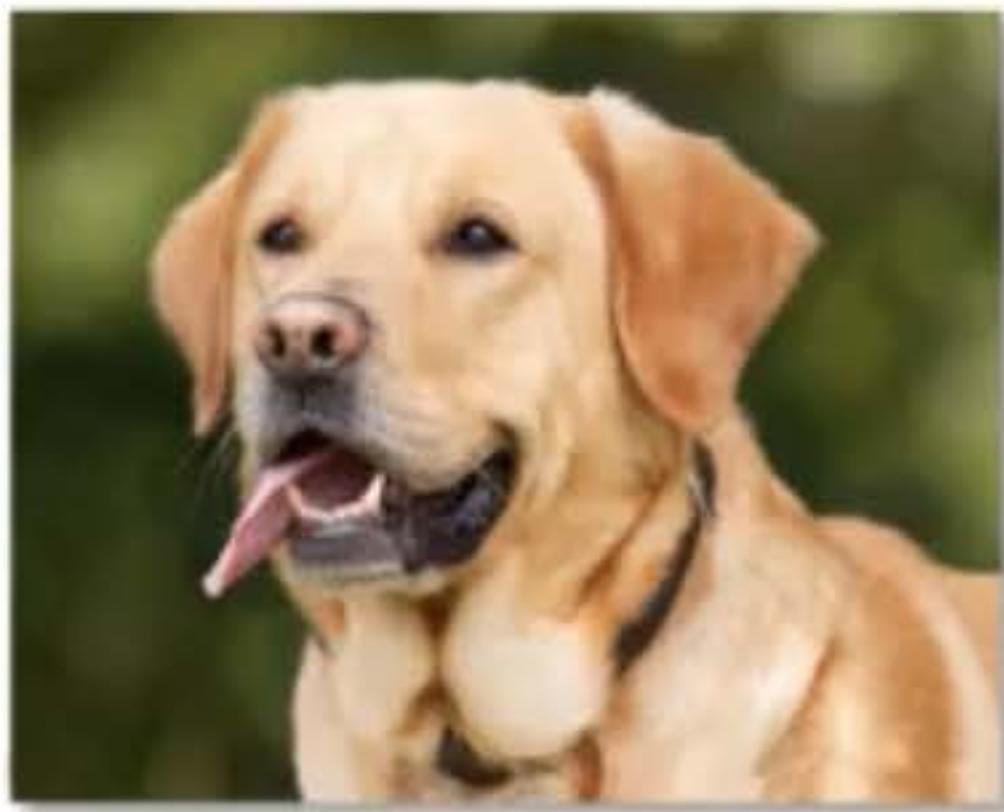
Link:

<https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Papers-You-Need-To-Know-About.html>

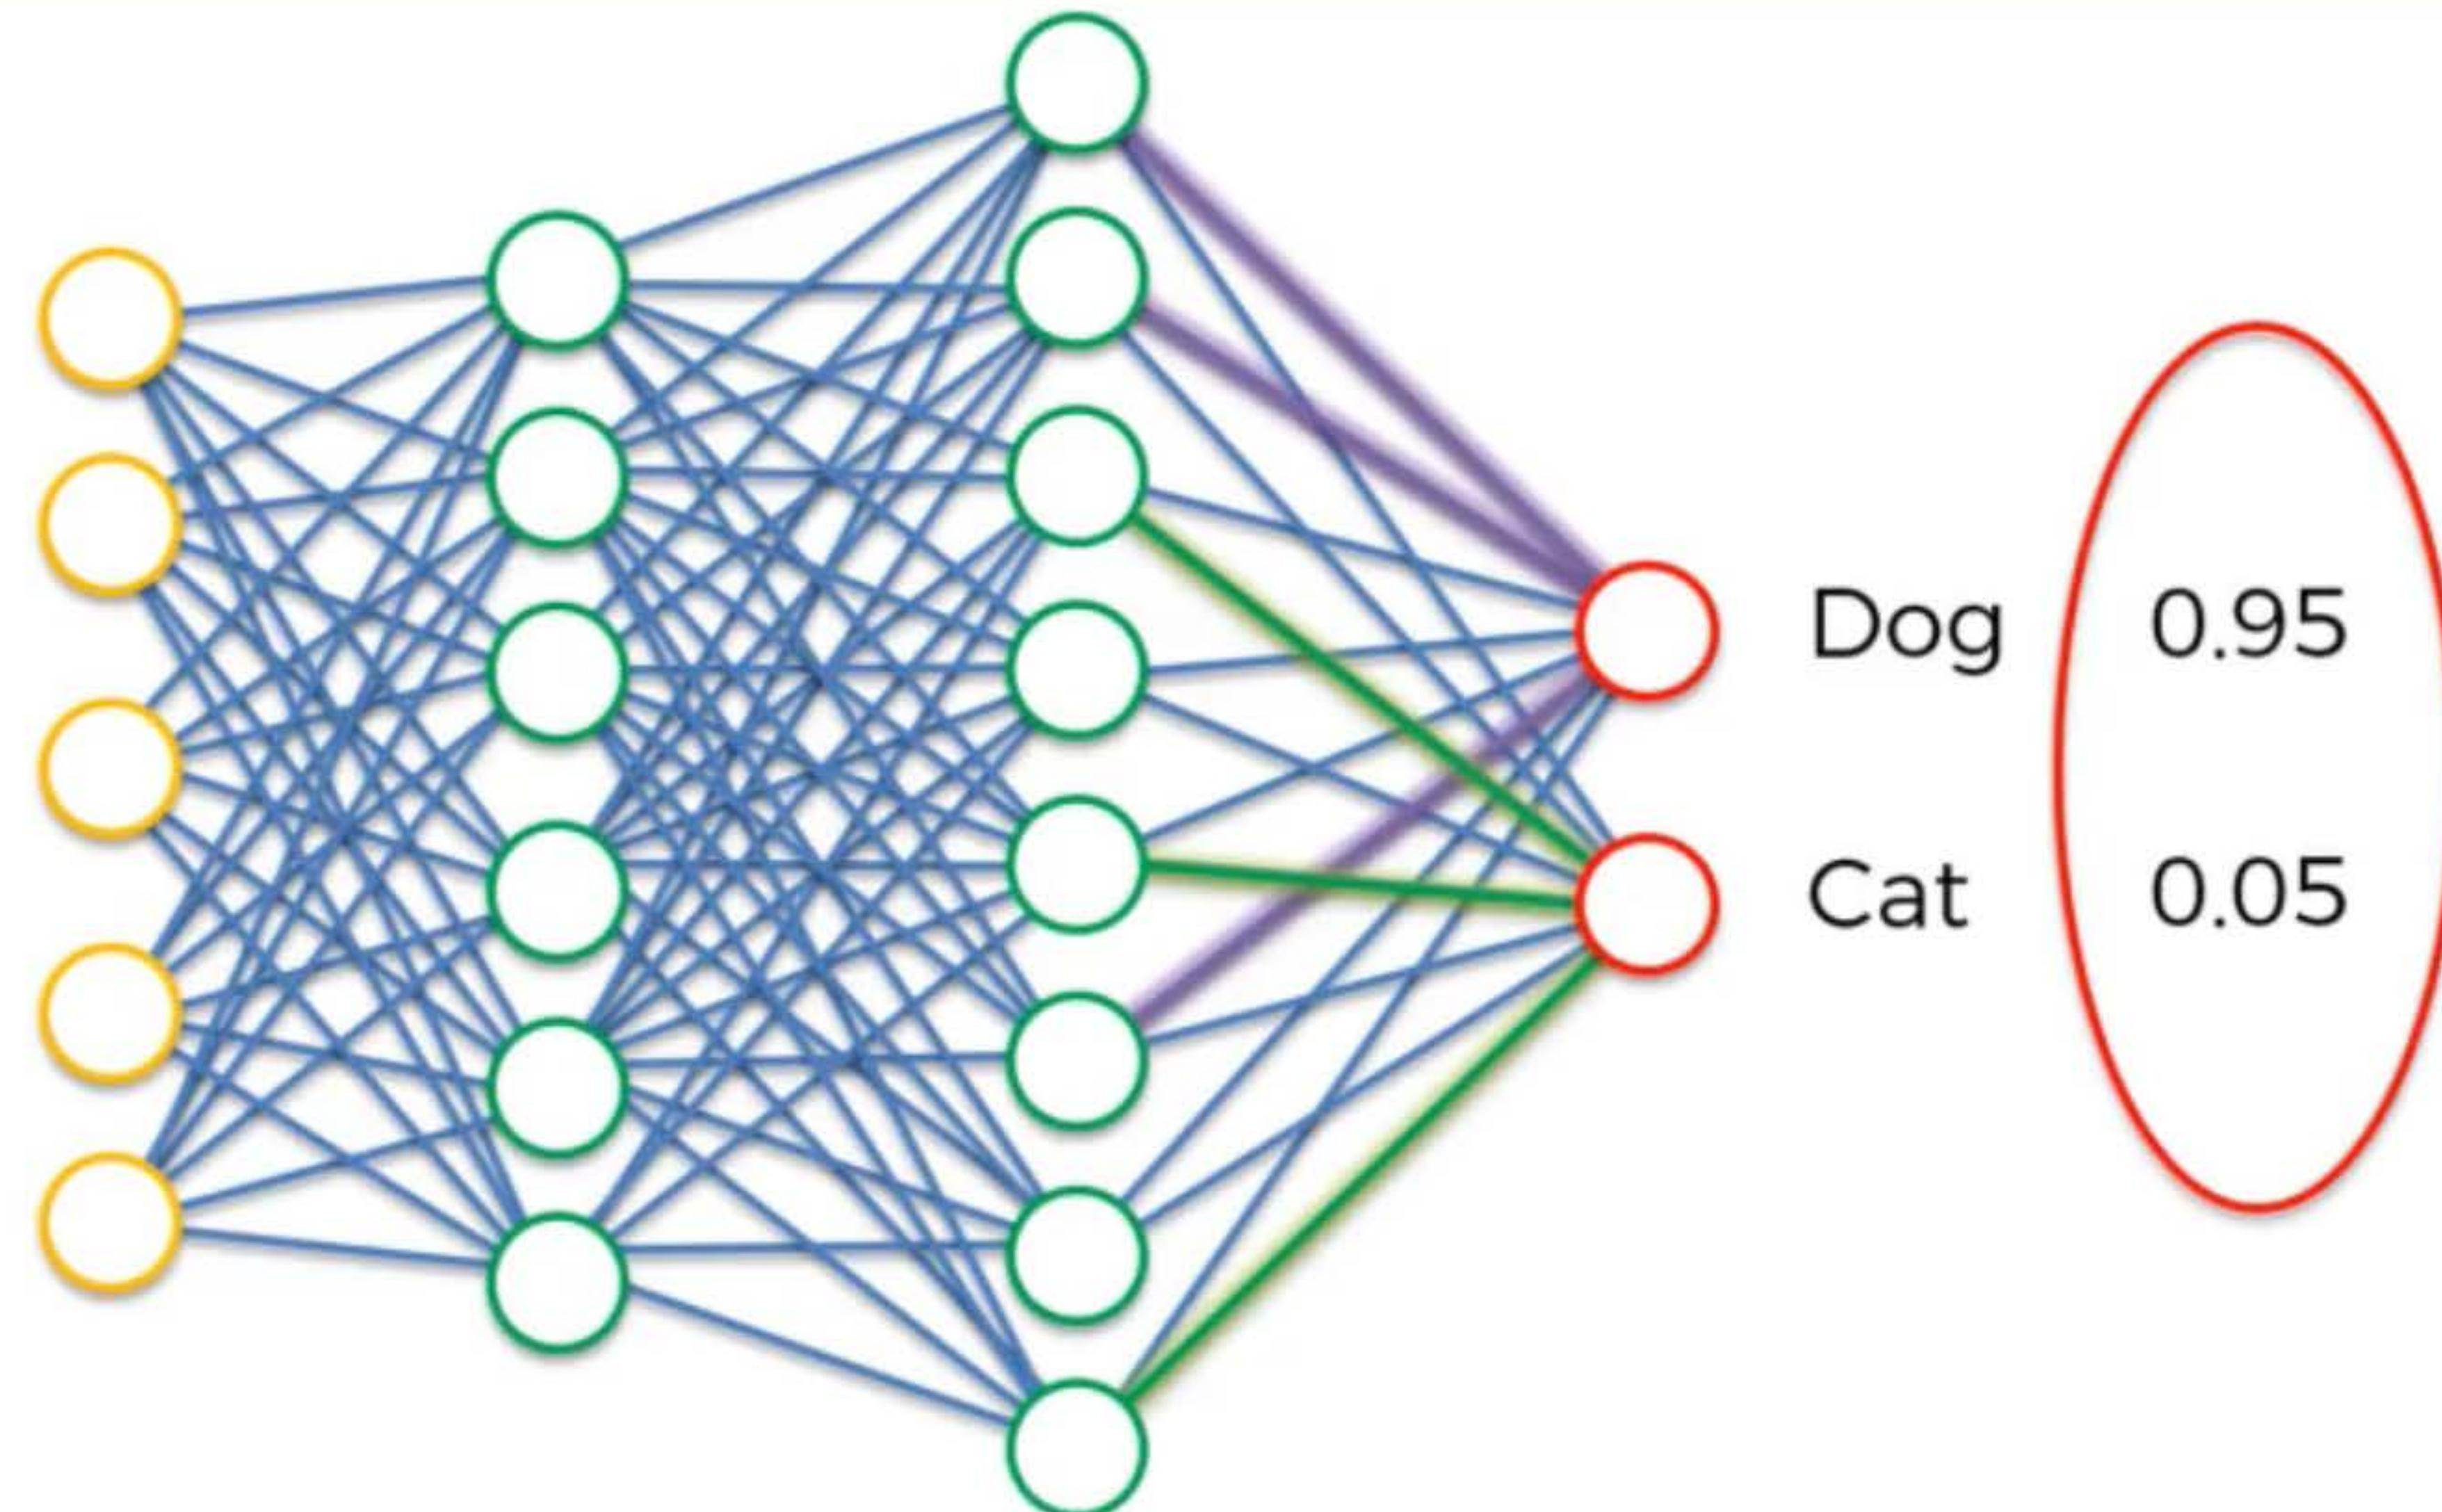


Softmax & Cross-Entropy

Softmax & Cross-Entropy



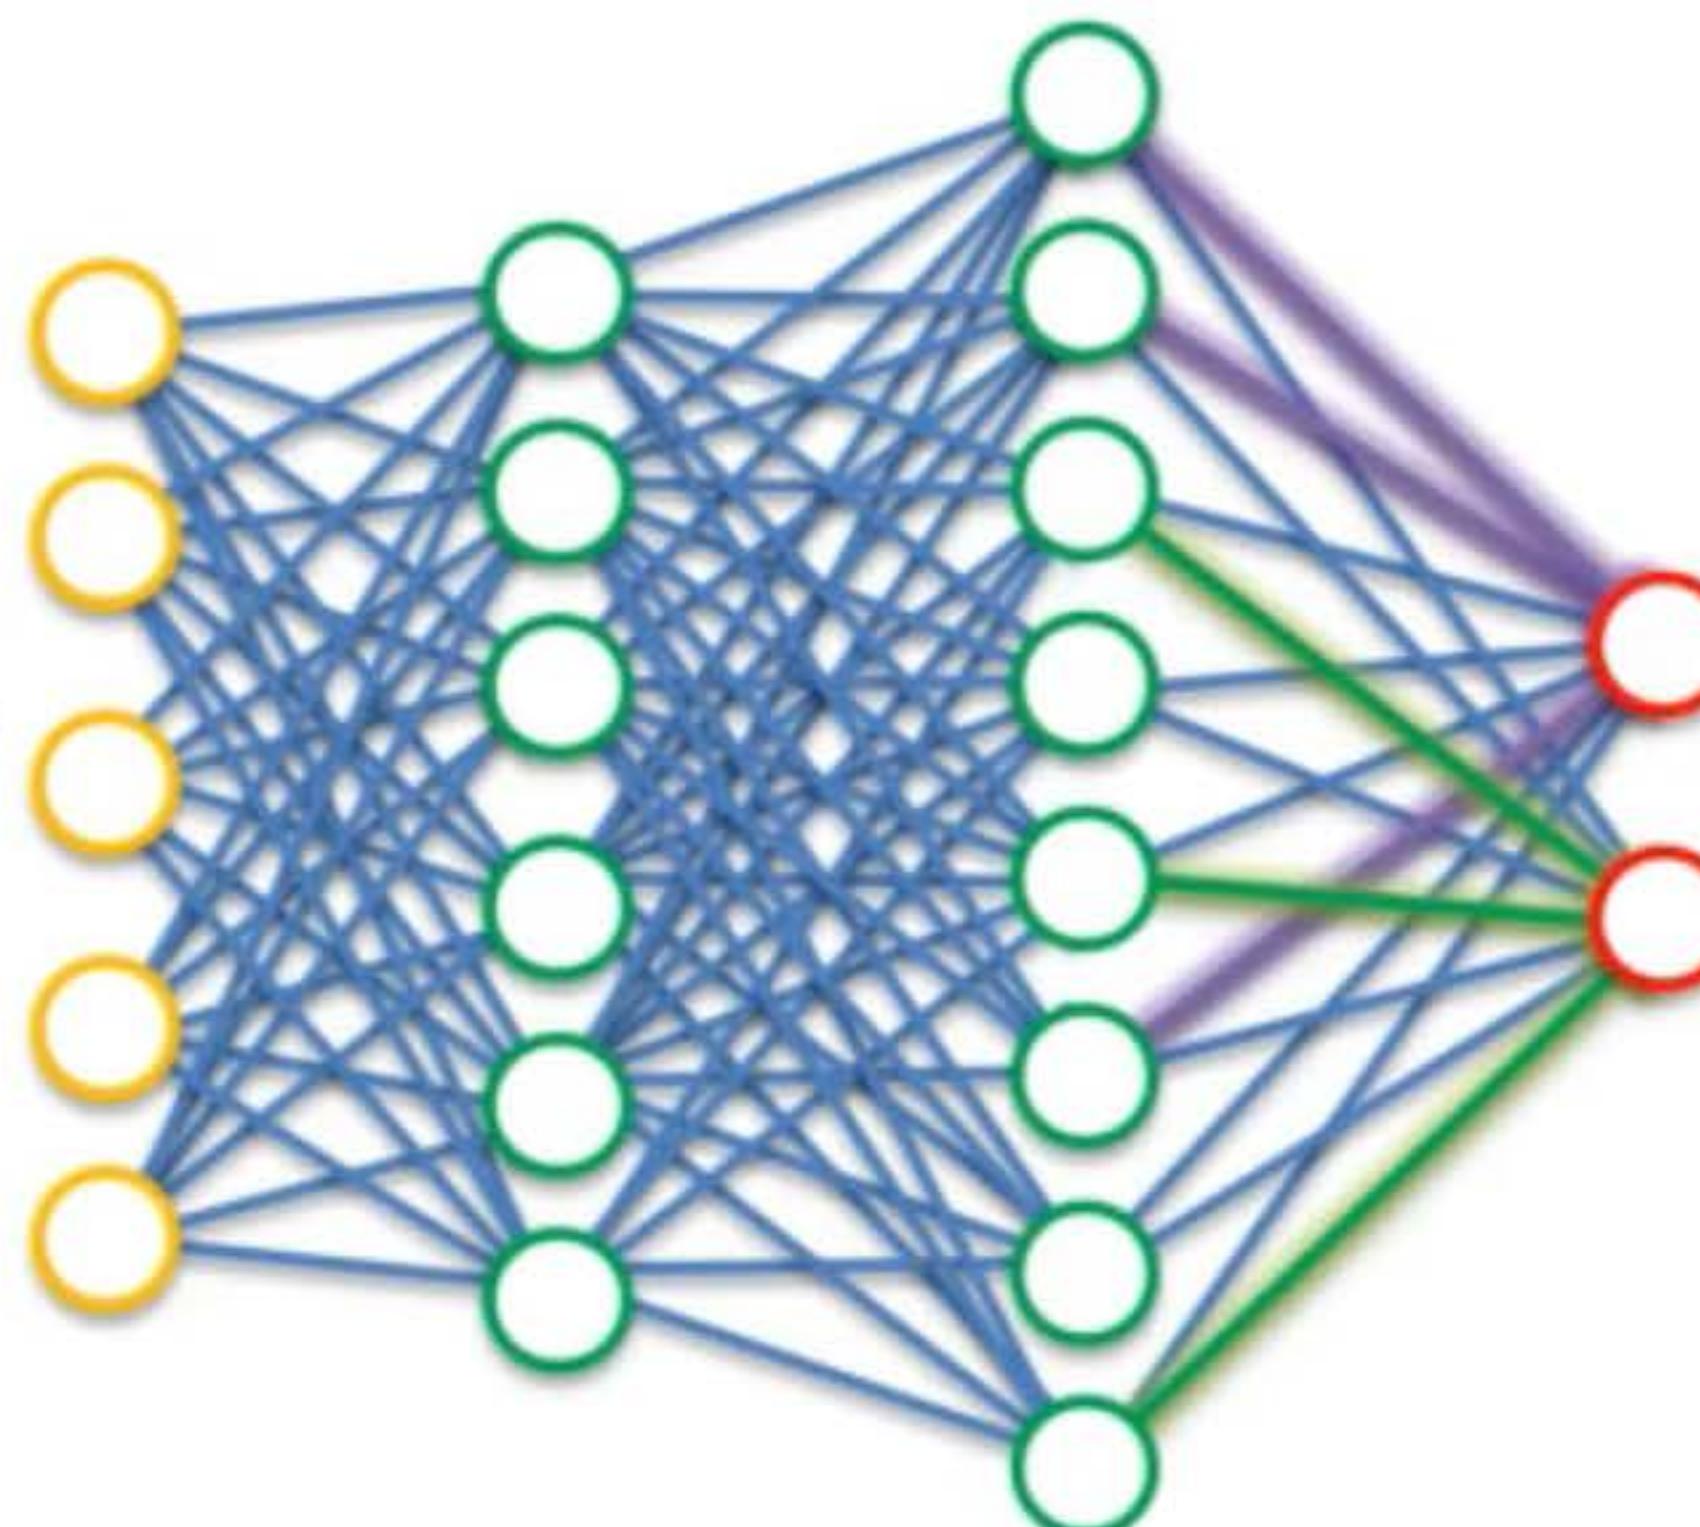
Flattening
..... →



Softmax & Cross-Entropy



.....
Flattening



$$f_j(z) = \frac{e^{z_j}}{\sum_k e^{z_k}}$$

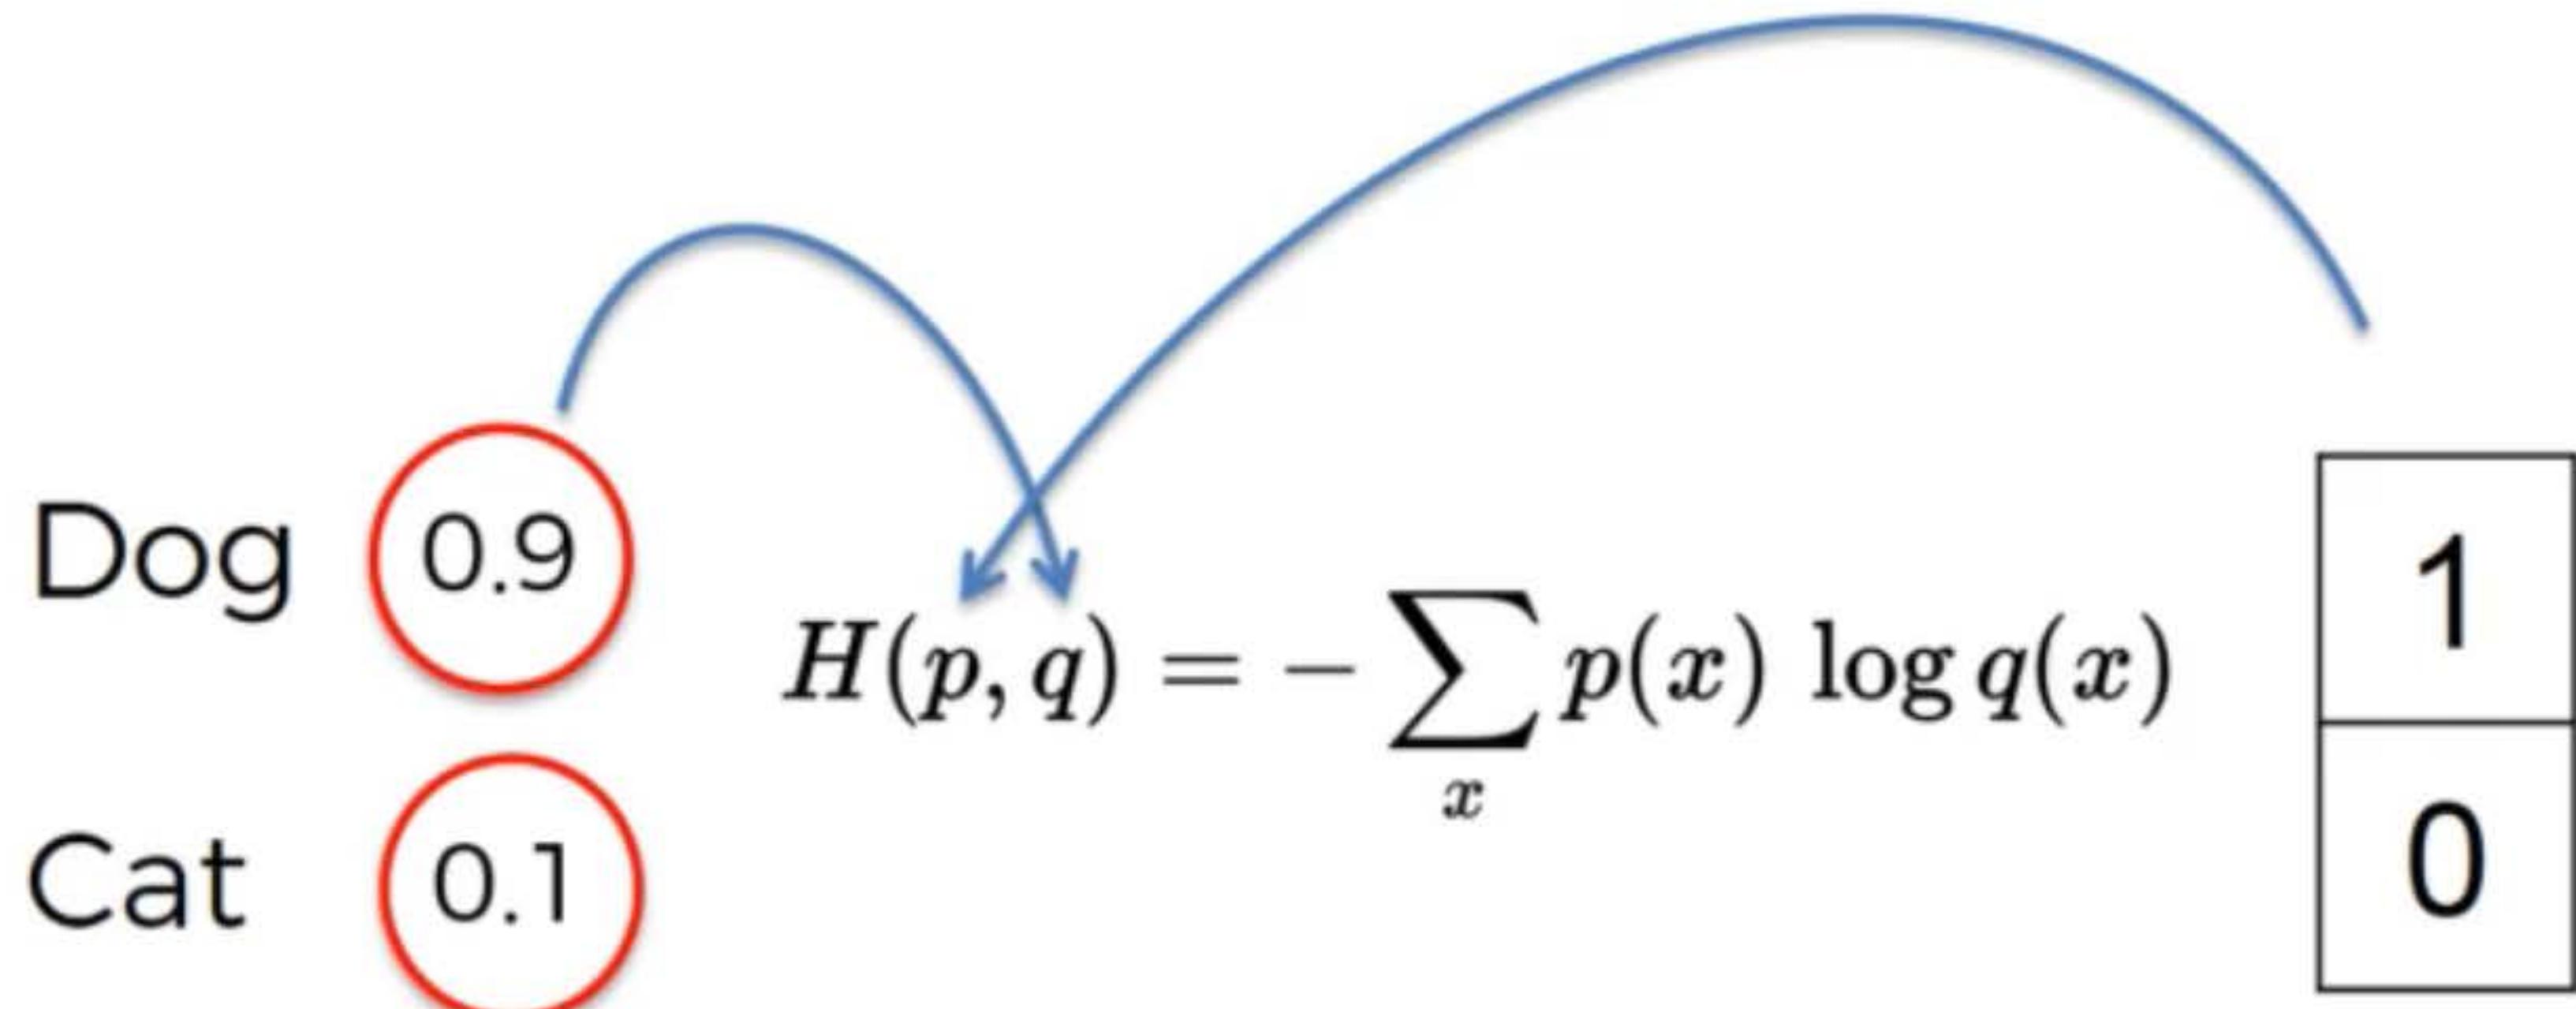
Dog $\rightarrow z_1 \rightarrow 0.95$
Cat $\rightarrow z_2 \rightarrow 0.05$

Softmax & Cross-Entropy

$$L_i = -\log \left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}} \right)$$

$$H(p, q) = - \sum_x p(x) \log q(x)$$

Softmax & Cross-Entropy

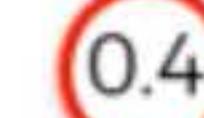


Softmax & Cross-Entropy

NN1 NN2



Dog	<table border="1"><tr><td>1</td></tr><tr><td>0</td></tr></table>	1	0
1			
0			
Cat	<table border="1"><tr><td>0</td></tr><tr><td>1</td></tr></table>	0	1
0			
1			



Dog	<table border="1"><tr><td>0</td></tr><tr><td>1</td></tr></table>	0	1
0			
1			
Cat	<table border="1"><tr><td>1</td></tr><tr><td>0</td></tr></table>	1	0
1			
0			



Dog	<table border="1"><tr><td>1</td></tr><tr><td>0</td></tr></table>	1	0
1			
0			
Cat	<table border="1"><tr><td>0</td></tr><tr><td>1</td></tr></table>	0	1
0			
1			



Softmax & Cross-Entropy

NN1

Row	Dog^	Cat^	Dog	Cat
#1	0.9	0.1	1	0
#2	0.1	0.9	0	1
#3	0.4	0.6	1	0

NN2

Row	Dog^	Cat^	Dog	Cat
#1	0.6	0.4	1	0
#2	0.3	0.7	0	1
#3	0.1	0.9	1	0

Classification Error

$$1/3 = 0.33$$

$$1/3 = 0.33$$

Mean Squared Error

$$0.25$$

$$0.71$$

Cross-Entropy

$$0.38$$

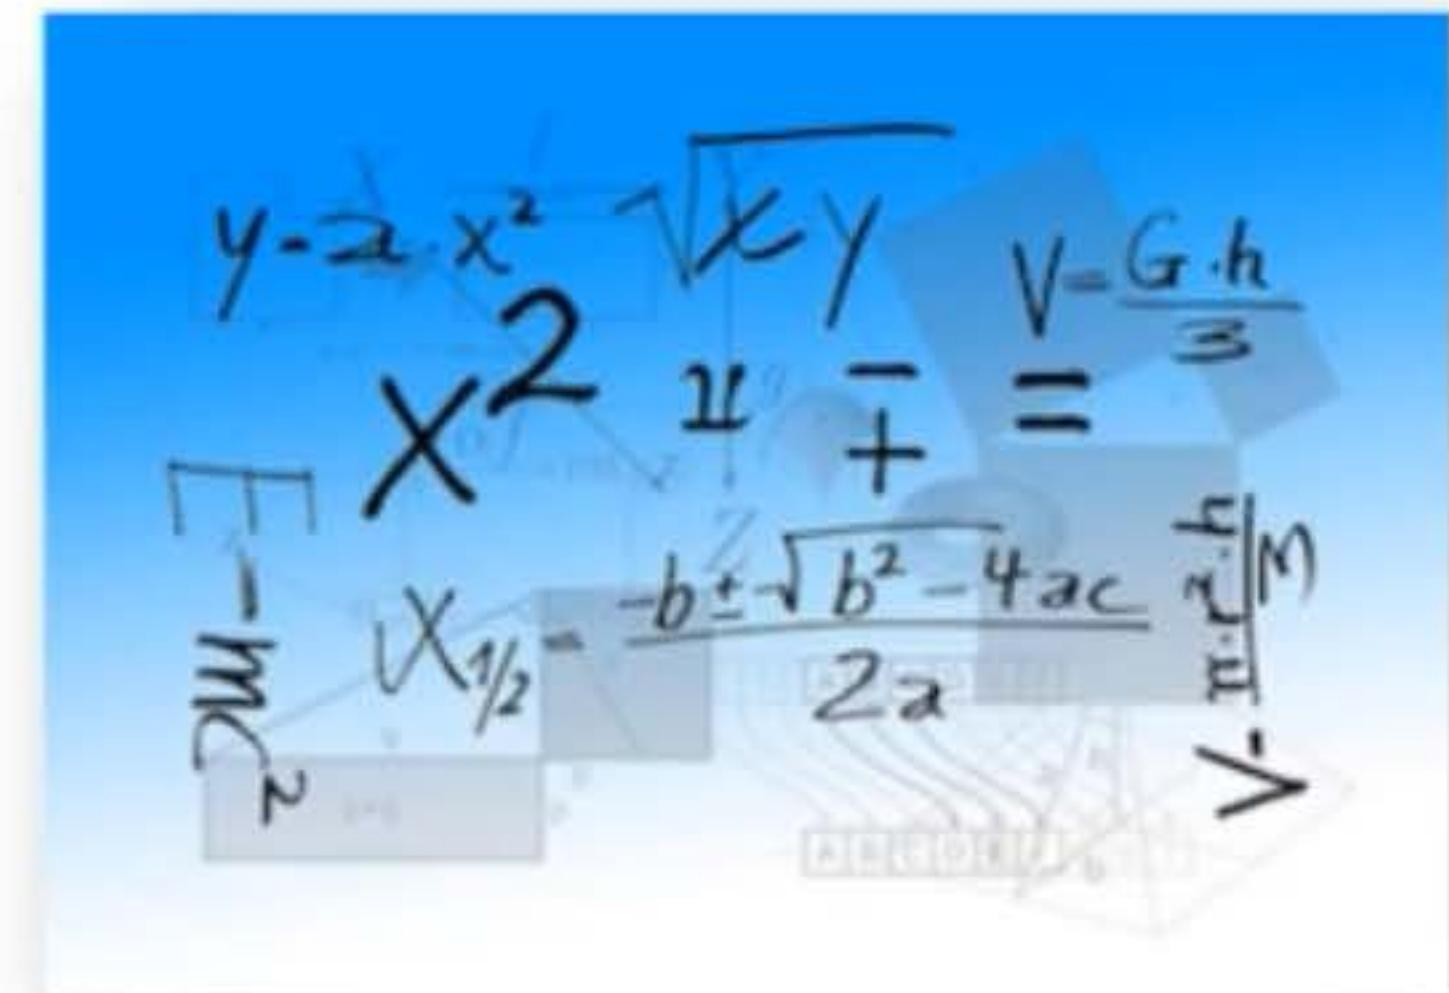
$$1.06$$

Softmax & Cross-Entropy

Additional Reading:

A Friendly Introduction to Cross-Entropy Loss

By Rob DiPietro (2016)



Link:

<https://rdipietro.github.io/friendly-intro-to-cross-entropy-loss/>

Softmax & Cross-Entropy

Additional Reading:

How to implement a neural network Intermezzo 2

By Peter Roelants (2016)

$$\begin{aligned}\frac{\partial \xi}{\partial z_i} &= - \sum_{j=1}^C \frac{\partial t_j \log(y_j)}{\partial z_i} = - \sum_{j=1}^C t_j \frac{\partial \log(y_j)}{\partial z_i} = - \sum_{j=1}^C t_j \frac{1}{y_j} \frac{\partial y_j}{\partial z_i} \\ &= - \frac{t_i}{y_i} \frac{\partial y_i}{\partial z_i} - \sum_{j \neq i}^C \frac{t_j}{y_j} \frac{\partial y_j}{\partial z_i} = - \frac{t_i}{y_i} y_i (1 - y_i) - \sum_{j \neq i}^C \frac{t_j}{y_j} (-y_j y_i) \\ &= -t_i + t_i y_i + \sum_{j \neq i}^C t_j y_i = -t_i + \sum_{j=1}^C t_j y_i = -t_i + y_i \sum_{j=1}^C t_j \\ &= y_i - t_i\end{aligned}$$

Link:

http://peterroelants.github.io/posts/neural_networkImplementation_intermezzo02/



Convolutional Neural Network

Importing the libraries

```
In [2]: import tensorflow as tf  
from keras.preprocessing.image import ImageDataGenerator
```

```
In [3]: tf.__version__
```

```
Out[3]: '2.3.0'
```

Part 1 - Data Preprocessing

Preprocessing the Training set

```
In [4]: train_datagen = ImageDataGenerator(rescale = 1./255,  
                                         shear_range = 0.2,  
                                         zoom_range = 0.2,  
                                         horizontal_flip = True)  
training_set = train_datagen.flow_from_directory('dataset/training_set',  
                                                target_size = (64, 64),  
                                                batch_size = 32,  
                                                class_mode = 'binary')
```

Found 8000 images belonging to 2 classes.

Preprocessing the Test set

```
In [5]: test_datagen = ImageDataGenerator(rescale = 1./255)  
test_set = test_datagen.flow_from_directory('dataset/test_set',  
                                             target_size = (64, 64),  
                                             batch_size = 32,  
                                             class_mode = 'binary')
```

Found 2000 images belonging to 2 classes.

Part 2 - Building the CNN

Initialising the CNN

```
In [6]: cnn = tf.keras.models.Sequential()
```

Step 1 - Convolution

```
In [7]: cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu', input_shape=[64, 64, 3]))
```

Step 2 - Pooling

```
In [8]: cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))
```

Adding a second convolutional layer

```
In [9]: cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))
```

Step 3 - Flattening

```
In [10]: cnn.add(tf.keras.layers.Flatten())
```

Step 4 - Full Connection

```
In [11]: cnn.add(tf.keras.layers.Dense(units=128, activation='relu'))
```

Step 5 - Output Layer

```
In [12]: cnn.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
```

Part 4 - Making a single prediction

```
In [15]: import numpy as np
from keras.preprocessing import image
test_image = image.load_img('dataset/single_prediction/cat_or_dog_1.jpg', target_size = (64, 64))
test_image = image.img_to_array(test_image)
#for the extra dimension of batch size 32
test_image = np.expand_dims(test_image, axis = 0)
result = cnn.predict(test_image)
#we get the right index for each class
training_set.class_indices
#we have only one batch and one image hence,[0][0]
if result[0][0] == 1:
    prediction = 'dog'
else:
    prediction = 'cat'
```

```
In [16]: print(prediction)
```

dog